

Lane pilots — fourteen API lanes and three CLIs on text and structured, one attempt each

PUBLIC RESULT · CAMPAIGN SUMMARY

13passes

Across 20 graded attempts; no cross-exam quality score

COVERAGE

2 exams · 17 lanes

COMPLETED ATTEMPTS

40

MODEL FAILURES

7

LANE ERRORS

15

WITHHELD / INVALID

3 / 2

Results by AI exam

Exam	Model	Passed / attempts	Errors
Record normalization and duplicate QA	@cf/openai/gpt-oss-20b	1/1	0
Intent and entity annotation	@cf/openai/gpt-oss-20b	1/1	0
Record normalization and duplicate QA	cli/claude-sonnet	1/1	0
Intent and entity annotation	cli/claude-sonnet	1/1	0
Record normalization and duplicate QA	cli/codex	1/1	0
Intent and entity annotation	cli/codex	1/1	0
Record normalization and duplicate QA	openai/gpt-oss-120b	1/1	0
Intent and entity annotation	openai/gpt-oss-120b	1/1	0

First 8 of 34 model/exam results in report order. The complete table is linked on page 2. Technical errors are counted in attempts but are not model failures.

6 verified records. Compare models within the same exam and input condition.

What was evaluated

INPUT CONDITION
Conditions differ by exam and model lane; the detailed matrix retains each cell without merging unlike tasks.
EXAM AND REVISION
2 exam revisions; each fingerprint is retained in the detailed report

Limits on the claim

- This summary is accounting, not a leaderboard or a broad capability score.
- Low sample cells support counts only; no ranking is inferred.
- Lane errors, withheld and invalid attempts are retained outside the model-failure denominator.
- Per-exam conditions and quality axes remain in the detailed report.

Evidence

CAMPAIGN	CANONICAL DIGEST
fc-0c7e7c67a231	6 sealed record digests in the detailed report

Open the full evidence matrix and audit report →

Public disclosure excludes submissions, raw answers, private references, exact geometry, frame timestamps, and hidden-test material.