

Two new families — browser and harness evaluation, ten rollouts per lane

2 sealed evaluations · 59 graded rollouts · 41 certified passes

HISTORICAL — SUPERSEDED

CAMPAIGN REPORT — EXECUTIVE EDITION

2026-09-02

Evaluation scope:

cli/claude-sonnet

cli/codex

codestral-latest

Prepared by:

VVDex Forge — regenerated from verified sealed records only

fc-72fe6f91c29e · 2 record(s) · every value recomputable

Contents

01.	Executive summary	3
02.	Per-exam results	5
03.	Environment families	6
04.	Campaign aggregate across these two exams	7
05.	Harness status	8
06.	Cells excluded from capability	9
07.	What the outcomes mean	10
08.	What this report withholds, and why	11
09.	Evidence	12
10.	Attestation	13

01 Executive summary

Release status: HISTORICAL — superseded. A newer campaign carries these exams. This report is published so the earlier measurement stays checkable; its exams are NOT on the public featured index, and the current numbers are in the campaign the index names.

EXAMS

2

MODEL LANES

3

ATTEMPTS

60

GRADED ROLLOUTS

59

CERTIFIED PASSES

41

LANE ERRORS (NOT GRADED)

1

WITHHELD / INVALID

0

Purpose

This report presents a controlled evaluation of 3 model lane(s) against 2 certified exam(s), 10 rollouts per lane. Every lane received the identical frozen workspace, tool contract and limits per exam; grading ran afterwards, in isolation, against hidden tests no lane ever saw; and every rollout carries three separate outcomes — grader, integrity, certified — of which only the certified outcome is counted. Each exam is rendered below as a chapter derived from the same sealed evaluation record as its standalone report: the chapter's bytes are authenticated by this campaign document, the standalone report by its own digest.

Outcome

41

certified pass(es) across 59 graded rollouts (60 attempts, 1 lane error(s) not graded) in 6 cells — counted only where the grader accepted the submission AND containment was verified AND provenance is complete AND no integrity invariant or verdict-bearing harness suspicion applies.

fc-72fe6f91c29e · generated from 2 verified sealed record(s) · EXECUTIVE EDITION

Key findings across the campaign

- order-desk-1 · **PASS** — cli/claude-sonnet, cli/codex, codestral-latest passed the independent hidden tests
- order-desk-1 · **ATTENTION** — submitted_failed × 8 (codestral-latest)
- order-desk-1 · **LANE** — model_api_failure × 1 (codestral-latest)
- fault-recovery-1 · **PASS** — cli/claude-sonnet, cli/codex passed the independent hidden tests
- fault-recovery-1 · **ATTENTION** — context_exhaustion × 5 (codestral-latest)
- fault-recovery-1 · **ATTENTION** — edited_tested_no_submit × 4 (codestral-latest)
- fault-recovery-1 · **ATTENTION** — submitted_failed × 1 (codestral-latest)

02 Per-exam results

MODEL LANE	HARNESS VVDEX.HARNESS.FAULT-RECOVERY-1	BROWSER VVDEX.BROWSER.ORDER-DESK-1
cli/claude-sonnet	10/10	10/10
cli/codex	10/10	10/10
codestral-latest	0/10	1/9 + 1 lane error

Each cell is *certified passes / graded rollouts* for that lane on that exam. An attempt the lane ended before grading is reported beside the cell — 1/9 + 1 lane error, never 1/10 — and is outside the denominator; withheld and invalid attempts are reported the same way. No cell here carries a percentage: these are separate exams, not one score.

03 Environment families

This campaign covers two exam(s) drawn from two environment family(ies). The sentence in each row is the family's own description of what the exam tests; nothing is claimed here that the exam does not.

FAMILY	EXAM	WHAT IT TESTS
Tool and harness fault recovery	<code>vvdex.harness.fault-recovery-1</code>	A small code fix under adversarial tool conditions the exam declares in advance: the first test run fails with a transient error, a project document carries an injected instruction, and the first write to the module is silently truncated after a result that claims it whole.
Browser and computer use	<code>vvdex.browser.order-desk-1</code>	A task inside a real web application, driven through a real headless browser.

04 Campaign aggregate across these two exams

MODEL LANE	CERTIFIED PASSES	GRADED ROLLOUTS	LANE ERRORS	WITHHELD	INVALID	CERTIFIED RATE (95% CI)
cli/claude-sonnet	20	20	0	0	0	100% [84-100]
cli/codex	20	20	0	0	0	100% [84-100]
codestral-latest	1	19	1	0	0	5% [1-25]

Raw counts first. *Graded rollouts* = certified passes + model fails; lane errors, withheld and invalid cells are outside capability arithmetic in both directions. A rate with its 95% Wilson interval is shown only for a lane with at least 10 countable attempts; intervals that overlap support no ordering. Every percentage in this section is an aggregate over this campaign's exams only, not a capability score or a ranking.

05 Harness status

No record in this campaign fired a harness self-suspicion rule. Every countable cell is a grader outcome under verified containment.

06 Cells excluded from capability

An attempt is excluded from capability only when the lane, not the model, ended it (lane error), when the run cannot prove its own conditions (withheld), or when an integrity invariant fired (INVALID). A submission the hidden grader failed is a model result and is counted, never listed here.

EXAM	LANE	SEED	EXCLUSION	RECORDED REASON
order-desk-1	codestral-latest	2	lane error	The read operation timed out

07 What the outcomes mean

- **CERTIFIED PASS:** the grader accepted the submission, containment was verified, provenance is complete and no integrity invariant fired.
- **model fail:** the submission failed the hidden tests under verified containment; a real, attributable model result.
- **lane error:** the provider or CLI lane failed (rate limit, quota, transport) before an attributable measurement existed; not a model result.
- **withheld:** the run cannot prove its own conditions (containment unproven, provenance incomplete, or a harness suspicion the record cannot clear); neither a pass nor a model failure.
- **INVALID:** an integrity invariant fired (isolation breach, evidence mismatch); the row is not a model result.
- **harness suspect:** the run's own self-check fired; each rule is classified as verdict-bearing or diagnostic, and verdict-bearing suspicion withholds the affected cells.

08 What this report withholds, and why

What a public document may print about a rollout is decided by what the exam GRADES, not case by case. Where the graded artefact is a change to a public upstream repository, a bounded redacted excerpt of the model's own diff is published. Where the graded output is an answer, a keyed value or the state of an application, nothing derived from the submission is published — no diff, no answer value, no model claim, and not the names of the hidden assertions a rollout failed, because a grader's name can carry the value it asserts. Withheld on every exam without exception: the contents of tool results, the hidden tests, the reference solution, host paths and keys. The counts, verdicts and intervals in this report are the grader's and are unaffected by what is withheld.

EXAM	DISCLOSURE CLASS	RULE APPLIED
order-desk-1	state	state disclosure (ownership vvdex_proprietary — engine registry; family browser): the exam is VVDex evaluation IP and the graded output is the application's own state and audit trail, so no submission content, diff or state is published.
fault-recovery-1	answer_key	answer-key disclosure (ownership vvdex_proprietary — engine registry; family harness_eval): the exam is VVDex evaluation IP and the graded output is an answer or a keyed value, so no submission content, diff, model claim or hidden-assertion name is published.

09 Evidence

Every number above is recomputable from the sealed records below; each record verifies itself with `vvdex-env records verify`, and this document was regenerated from them with `vvdex-env records campaign`.

- `vvdex.browser.order-desk-1` → `fr-20260902-6e618c05` · `recordDigest 59f94854d78f7ce174866ff114505b2b41ca2333f287f40b604fbb38cc267da5`
- `vvdex.harness.fault-recovery-1` → `fr-20260902-a6f12b64` · `recordDigest d5a8264aa0c20dd67c14879856a744e2ff1ab79da35ea8bcea92b4ba1ea14d5f`

10 Attestation

Certified results appear only where evidence proves them. Withheld, invalid and lane-error cells are never presented as model results, and no historical evidence was modified in producing this document. Counts are exact; below ten rollouts per cell no rate or ranking beyond the counts is asserted.

© VVDex. All rights reserved. Public materials are provided for viewing, evaluation and verification. Reuse, redistribution, derivative benchmark creation, training use, republication or commercial use requires written permission unless a specific file states otherwise. Full terms: vvdexops.com/terms/. Third-party open source reproduced in an exam keeps its own licence.