

Two new families — browser and harness evaluation, ten rollouts per lane

PUBLIC RESULT · CAMPAIGN SUMMARY

41 passes

Across 59 graded attempts; no cross-exam quality score

COVERAGE

2 exams · 3 lanes

COMPLETED ATTEMPTS

60

MODEL FAILURES

18

LANE ERRORS

1

WITHHELD / INVALID

0 / 0

Results by AI exam

Exam	Model	Passed / attempts	Errors
Harness recovery under a faulted tool loop	cli/claude-sonnet	10/10	0
Correct one order in an admin application, through a real browser	cli/claude-sonnet	10/10	0
Harness recovery under a faulted tool loop	cli/codex	10/10	0
Correct one order in an admin application, through a real browser	cli/codex	10/10	0
Harness recovery under a faulted tool loop	codestral-latest	0/10	0
Correct one order in an admin application, through a real browser	codestral-latest	1/9 + 1 lane error	1

Showing all 6 model/exam results. Technical errors are counted in attempts but are not model failures.

2 verified records. Compare models within the same exam and input condition.

What was evaluated

INPUT CONDITION

Conditions differ by exam and model lane; the detailed matrix retains each cell without merging unlike tasks.

EXAM AND REVISION

2 exam revisions; each fingerprint is retained in the detailed report

Limits on the claim

- This summary is accounting, not a leaderboard or a broad capability score.
- Low sample cells support counts only; no ranking is inferred.
- Lane errors, withheld and invalid attempts are retained outside the model-failure denominator.
- Per-exam conditions and quality axes remain in the detailed report.

Evidence

CAMPAIGN

fc-72fe6f91c29e

CANONICAL DIGEST

2 sealed record digests in the detailed report

Open the full evidence matrix and audit report →

Public disclosure excludes submissions, raw answers, private references, exact geometry, frame timestamps, and hidden-test material.