

Two new families — browser and harness evaluation, ten rollouts per lane

2 sealed evaluations · 59 graded rollouts · 41
certified passes

HISTORICAL — SUPERSEDED

CAMPAIGN REPORT

2026-09-02

Evaluation scope:

- cli/claude-sonnet
- cli/codex
- codestral-latest

Prepared by:

VVDex Forge — regenerated from verified sealed records only

fc-72fe6f91c29e · 2 record(s) · every value recomputable

Contents

01.	Executive summary	3
02.	Per-exam results	4
03.	Environment families	5
04.	Campaign aggregate across these two exams	6
05.	Harness status	7
06.	Cells excluded from capability	8
07.	What the outcomes mean	9
08.	What this report withholds, and why	10
09.	Exam · order-desk-1	11
10.	Exam · fault-recovery-1	48
11.	Evidence	82
12.	Attestation	83

01 Executive summary

Release status: HISTORICAL — superseded. A newer campaign carries these exams. This report is published so the earlier measurement stays checkable; its exams are NOT on the public featured index, and the current numbers are in the campaign the index names.

EXAMS 2	MODEL LANES 3	ATTEMPTS 60	GRADED ROLLOUTS 59
CERTIFIED PASSES 41	LANE ERRORS (NOT GRADED) 1	WITHHELD / INVALID 0	

Purpose

This report presents a controlled evaluation of 3 model lane(s) against 2 certified exam(s), 10 rollouts per lane. Every lane received the identical frozen workspace, tool contract and limits per exam; grading ran afterwards, in isolation, against hidden tests no lane ever saw; and every rollout carries three separate outcomes — grader, integrity, certified — of which only the certified outcome is counted. Each exam is rendered below as a chapter derived from the same sealed evaluation record as its standalone report: the chapter's bytes are authenticated by this campaign document, the standalone report by its own digest.

Outcome

41

certified pass(es) across 59 graded rollouts (60 attempts, 1 lane error(s) not graded) in 6 cells — counted only where the grader accepted the submission AND containment was verified AND provenance is complete AND no integrity invariant or verdict-bearing harness suspicion applies.

fc-72fe6f91c29e · generated from 2 verified sealed record(s)

Key findings across the campaign

- order-desk-1 · **PASS** — cli/claude-sonnet, cli/codex, codestral-latest passed the independent hidden tests
- order-desk-1 · **ATTENTION** — submitted_failed × 8 (codestral-latest)
- order-desk-1 · **LANE** — model_api_failure × 1 (codestral-latest)
- fault-recovery-1 · **PASS** — cli/claude-sonnet, cli/codex passed the independent hidden tests
- fault-recovery-1 · **ATTENTION** — context_exhaustion × 5 (codestral-latest)
- fault-recovery-1 · **ATTENTION** — edited_tested_no_submit × 4 (codestral-latest)
- fault-recovery-1 · **ATTENTION** — submitted_failed × 1 (codestral-latest)

02 Per-exam results

MODEL LANE	HARNESS VVDEX.HARNESS.FAULT-RECOVERY-1	BROWSER VVDEX.BROWSER.ORDER-DESK-1
cli/claude-sonnet	10/10	10/10
cli/codex	10/10	10/10
codestral-latest	0/10	1/9 + 1 lane error

Each cell is *certified passes / graded rollouts* for that lane on that exam. An attempt the lane ended before grading is reported beside the cell — 1/9 + 1 lane error, never 1/10 — and is outside the denominator; withheld and invalid attempts are reported the same way. No cell here carries a percentage: these are separate exams, not one score.

03 Environment families

This campaign covers two exam(s) drawn from two environment family(ies). The sentence in each row is the family's own description of what the exam tests; nothing is claimed here that the exam does not.

FAMILY	EXAM	WHAT IT TESTS
Tool and harness fault recovery	<code>vvdex.harness.fault-recovery-1</code>	A code-fix task carried out through a deliberately unreliable tool loop.
Browser and computer use	<code>vvdex.browser.order-desk-1</code>	A task inside a real web application, driven through a real headless browser.

04 Campaign aggregate across these two exams

MODEL LANE	CERTIFIED PASSES	GRADED ROLLOUTS	LANE ERRORS	WITHHELD	INVALID	CERTIFIED RATE (95% CI)
cli/claude-sonnet	20	20	0	0	0	100% [84-100]
cli/codex	20	20	0	0	0	100% [84-100]
codestral-latest	1	19	1	0	0	5% [1-25]

Raw counts first. *Graded rollouts* = certified passes + model fails; lane errors, withheld and invalid cells are outside capability arithmetic in both directions. A rate with its 95% Wilson interval is shown only for a lane with at least 10 countable attempts; intervals that overlap support no ordering. Every percentage in this section is an aggregate over this campaign's exams only, not a capability score or a ranking.

05 Harness status

No record in this campaign fired a harness self-suspicion rule. Every countable cell is a grader outcome under verified containment.

06 Cells excluded from capability

An attempt is excluded from capability only when the lane, not the model, ended it (lane error), when the run cannot prove its own conditions (withheld), or when an integrity invariant fired (INVALID). A submission the hidden grader failed is a model result and is counted, never listed here.

EXAM	LANE	SEED	EXCLUSION	RECORDED REASON
order-desk-1	codestral-latest	2	lane error	The read operation timed out

07 What the outcomes mean

- **CERTIFIED PASS:** the grader accepted the submission, containment was verified, provenance is complete and no integrity invariant fired.
- **model fail:** the submission failed the hidden tests under verified containment; a real, attributable model result.
- **lane error:** the provider or CLI lane failed (rate limit, quota, transport) before an attributable measurement existed; not a model result.
- **withheld:** the run cannot prove its own conditions (containment unproven, provenance incomplete, or a harness suspicion the record cannot clear); neither a pass nor a model failure.
- **INVALID:** an integrity invariant fired (isolation breach, evidence mismatch); the row is not a model result.
- **harness suspect:** the run's own self-check fired; each rule is classified as verdict-bearing or diagnostic, and verdict-bearing suspicion withholds the affected cells.

08 What this report withholds, and why

What a public document may print about a rollout is decided by what the exam GRADES, not case by case. Where the graded artefact is a change to a public upstream repository, a bounded redacted excerpt of the model's own diff is published. Where the graded output is an answer, a keyed value or the state of an application, nothing derived from the submission is published — no diff, no answer value, no model claim, and not the names of the hidden assertions a rollout failed, because a grader's name can carry the value it asserts. Withheld on every exam without exception: the contents of tool results, the hidden tests, the reference solution, host paths and keys. The counts, verdicts and intervals in this report are the grader's and are unaffected by what is withheld.

EXAM	DISCLOSURE CLASS	RULE APPLIED
order-desk-1	state	state disclosure (ownership vvdex_proprietary — engine registry; family browser): the exam is VVDex evaluation IP and the graded output is the application's own state and audit trail, so no submission content, diff or state is published.
fault-recovery-1	answer_key	answer-key disclosure (ownership vvdex_proprietary — engine registry; family harness_eval): the exam is VVDex evaluation IP and the graded output is an answer or a keyed value, so no submission content, diff, model claim or hidden-assertion name is published.

09 Exam · order-desk-1

vvdex.browser.order-desk-1 · record fr-20260902-6e618c05 · record digest 59f94854d78f7ce1... · harness clean · containme
nt clean

Executive summary

MODELS TESTED 3	CERTIFIED PASSES 21 of 29 graded rollouts	SUBMITTED 29 of 29 graded rollouts that called submit	NOT GRADED — LANE ERRORS 1 of 30 attempts
---------------------------------------	--	--	--

Purpose

This report presents a controlled evaluation of 3 model lane(s) against the certified exam vv
dex.browser.order-desk-1 (v0.1.0, family browser). Every lane received the identical
frozen workspace, tool contract and limits; grading ran afterwards, in isolation, against
hidden tests no lane ever saw. The purpose is to state, with reproducible evidence, what
each lane did and where it stopped.

Outcome

21 of 29

graded rollouts passed the independent hidden tests and were certified — 1 of 30 attempts not graded (lane errors)

21 of 29 graded rollouts passed. cli/claude-sonnet, cli/codex, codestral-latest completed
the required browser workflow, submitted, and were accepted by the independent grader.
The most common way to fail was submitted_failed — submitted a final application state
that the independent grader rejected.

Key findings

- F-01 cli/claude-sonnet, cli/codex, codestral-latest passed the independent hidden tests
- F-02 submitted_failed × 8 (codestral-latest)
- F-03 model_api_failure × 1 (codestral-latest)

Each finding is stated in full, with evidence, in the Findings section; recommendations for
acting on them follow it. Elapsed wall clock for the whole engagement: 3292.4s.

Exam definition

In scope. The ability of each configured model lane to complete this one certified task end to end — inspect and open the application, drive it through the declared browser actions to the required state, submit, and be accepted by the independent grader — within 40 tool steps and 9720s of wall clock, at 10 rollout(s) per lane. Behavioural measurements along the way (tool discipline, grounding, overclaim) are in scope because they are read from the recorded trace, not judged by another model.

Out of scope. General model quality, performance on other task families, prompt-tuning on the lane's behalf, and any comparison this run's sample size cannot support. A lane's API failure is reported as infrastructure, never as model capability.

The instrument. Exam `vvdex.browser.order-desk-1` was certified before this run: its reference solution passes, an empty submission fails, its grader distinguishes near-misses, and its sealed evidence is unchanged by this evaluation (see Certification evidence).

What was measured

EXAM	VERSION	FAMILY	ROLLOUTS
vvdex.browser.order-desk-1	0.1.0	browser	30 (10 per model)

Where the defect came from

This exam has no upstream repository to credit: the defect, the reference solution and the hidden suite are VVDex's own.

Certification evidence

4 of 5 certification pillars are recorded in this package, loaded from the receipts sealed into the exam and matching this run's exam fingerprint. What is recorded is shown with its real value; what is not is named below and is not claimed.

FROZEN BASELINE

FAIL

expected — an empty submission over 2 run(s) must not pass

→

GOLD / REFERENCE

PASS

expected — the reference solution over 2 run(s) must pass

GRADER CONTROLS

9 / 9

The grader accepted the reference and rejected every near-miss.

ATTACK PROBES

10 / 10 blocked

0 trivial exploit(s) found.

RUNTIME

sha256:679b36225a1f83238efba8c71b9cde3c24caaeacc9b682ae92819cb55eec328f

network policy: deny

Certification metadata is not included in this report package for:

- sealed evidence bundle

Benchmark results remain valid as run evidence; the pillar(s) above are not substantiated by this document and this report does not claim them.

What the package does carry: the exam fingerprint below is the contract digest this run was executed against — a holder of the certification bundle for this fingerprint can join the two independently.

EXAM FINGERPRINT

8d404ee5d5ab71456de3d3df0b8ca69271d8810767a694f71c6e45e72fc347bf

The contract digest this run was executed against.

Evaluation configuration

REPORT	fr-20260902-6e618c05
ISSUED	2026-09-02
SUBJECT EXAM	vvdex.browser.order-desk-1 · v0.1.0
FAMILY	browser
PREPARED BY	VVDex Forge engine vvdex-env 0.1.0
CLASSIFICATION	Independent model evaluation report — independently verifiable
CERTIFICATION	Exam certification: partially recorded (see Certification evidence)
STATUS	FINAL · report sealed against its own digest

Timeline

WHEN (UTC)	EVENT
2026-09-02 14:49:56	Evaluation run started
2026-09-02 15:44:48	Evaluation run finished
2026-09-02 15:44:48	Report generated from the sealed run evidence

Models under test

MODEL	PROVIDER	ENDPOINT	BILLING
Claude Code CLI (Sonnet, subscription)	cli	local runtime	operator subscription, flat — not billed per token
Codex CLI (GPT, subscription)	cli	local runtime	operator subscription, flat — not billed per token
codestral-latest	mistral	VVDex-configured	VVDex free-quota

A customer endpoint is recorded by origin only — never its port, its path, or its key.

Limits

RUNTIME docker	STEP LIMIT 40	WALL-CLOCK LIMIT 9720s	MAX OUTPUT TOKENS 2048
TEMPERATURE 0.2	RETRY POLICY none	COMMAND BUDGET 0	TOOLS OFFERED list_files, read_file, submit, browser_open, browser_snapshot, browser_click, browser_fill, browser_select, browser_press, browser_back

The evaluated model never saw the hidden tests or the reference solution, and could not change its result. Grading ran in a separate step, after submission, on a container with no network.

Model outcomes

Denominators. Attempts requested: 30. Graded (evaluable) rollouts: 29. Not graded — lane errors: 1. A lane error is a rollout the API or the runtime ended before the hidden grader ran: it is listed, excluded from every rate and pass@k, and never silently dropped. Passes are certified passes: hidden grader exit 0 with integrity verified.

MODEL	ATTEMPTS	GRADED	PASSED	MODEL FAILS	LANE ERRORS	SUBMITTED	MEAN TIME	TYPICAL DEEPEST STAGE
Claude Code CLI (Sonnet, subscription) cli/claude-sonnet	10	10	10 / 10	0	0	10 / 10	89.4s	Passed hidden tests
Codex CLI (GPT, subscription) cli/codex	10	10	10 / 10	0	0	10 / 10	125.2s	Passed hidden tests
codestral-latest codestral-latest	10	9	1 / 9	8	1	9 / 9	84.4s	Graded

Every rollout, with its own deepest stage and grade, is in Appendix A.

Success rate, confidence and pass@k

Every rate on this page is a measurement of a finite number of rollouts, so every rate carries the interval that measurement actually supports. n is the evaluable count — attempts minus rollouts the lane ended before grading; those are listed under "not graded" and sit in no rate, interval or pass@k.

MODEL	ATTEMPTS	N	NOT GRADED	PASSES	RATE	95% INTERVAL	SEEDS	PASS@1	PASS@3	PASS@5	PASS@10
cli/claude-sonnet	10	10	0	10	100.0%	[72%, 100%]	0–9	100.0%	100.0%	100.0%	100.0%
cli/codex	10	10	0	10	100.0%	[72%, 100%]	0–9	100.0%	100.0%	100.0%	100.0%
codestral-latest	10	9	1	1	11.1%	[2%, 43%]	0–9	11.1%	33.3%	55.6%	–

Interval and pass@k methods are stated in full in Appendix D. pass@k is shown for k = 1, 3, 5, 10; the sealed record carries every k up to n.

Model comparison

MODEL	PASS@1	95% INTERVAL	TOKENS	COST	MEAN ROLLOUT	OVERCLAIM RATE	TOP FAILURE CLASS
cli/claude-sonnet	100.0%	[72%, 100%]	n/a	n/a	89.4s	0.0%	none
cli/codex	100.0%	[72%, 100%]	n/a	n/a	125.2s	0.0%	none
codestral-latest	11.1%	[2%, 43%]	857 790	not billed	84.4s	0.0%	submitted_failed

Tokens and cost read n/a where a lane reports no telemetry (a flat-rate CLI) — an absent measurement, not zero. pass@1 and its interval are over evaluable rollouts only.

This is not a leaderboard. Separated by this run (intervals do not overlap): cli/claude-sonnet vs codestral-latest; cli/codex vs codestral-latest. Not separated (intervals overlap): cli/claude-sonnet vs cli/codex. Below the 10-rollout floor for reading a rate as a rate: codestral-latest (9) — read the count, not the percentage. The primary comparison in this report is the stage funnel below — where a model stops is a finding; a percentage over this many rollouts is not.

Stage funnel

The funnel is the primary reading of a run. Where a model stops says more than a single pass percentage does, especially at low N.

STAGE	CLI:CLI/CLAUDE-SONNET	CLI:CLI/CODEX	MISTRAL:CODESTRAL-LATEST
Inspected	10 / 10	10 / 10	9 / 9
Edited	N/A	N/A	N/A
Ran tests	N/A	N/A	N/A
Submitted	10 / 10	10 / 10	9 / 9
Graded	10 / 10	10 / 10	9 / 9
Passed hidden tests	10 / 10	10 / 10	1 / 9

Deepest stage reached

MODEL	DEEPEST (ANY ATTEMPT)	TYPICAL (MOST ATTEMPTS)	HIDDEN-TEST PASSES	NOT GRADED
cli:cli/claude-sonnet	Passed hidden tests	Passed hidden tests	10 / 10	0
cli:cli/codex	Passed hidden tests	Passed hidden tests	10 / 10	0
mistral:codestral-latest	Passed hidden tests	Graded	1 / 9	1

Stage definitions are in Appendix D. The matrix counts evaluable rollouts; lane errors are shown beside it, not inside it. The stages are not strictly nested: a model can submit without editing, and the matrix shows that rather than hiding it. Each rollout's own deepest stage is in Appendix A. N/A marks a stage this exam does not offer (Edited, Ran tests): not reached, not failed — not asked.

Behavioural analysis

Long-horizon behaviour

Counted off the recorded trace of each rollout. A dead end is a step that moved nothing: a re-read of a slice already returned, a repeat of the previous action, or a call to a tool this exam does not have.

MODEL	N	STEPS USED	RAN OUT OF STEPS	COMMANDS	CHANGED THE TREE	REFUSED (OVER BUDGET)	REVISITS	DEAD ENDS / ROLLOUT
cli:cli/claude-sonnet	10	18 / 40	0	0 / 0	0	0	0	3
cli:cli/codex	10	20 / 40	0	0 / 0	0	0	0	2
mistral:codestral-latest	10	22 / 40	0	0 / 0	0	0	0	1

Each column is defined in Appendix D.

Hallucination & overclaim

Three measurements, each with a stated definition. None of them is a judgement of the model's prose by another model — every number below comes from comparing what the model wrote against what was on disk, or against what the independent grader returned.

MODEL	N	GROUNDING FAILURES	KINDS	ROLLOUTS AFFECTED	OVERCLAIMS	OVERCLAIM RATE	RE-READS
cli:cli/claude-sonnet	10	0	none	0 / 10	0 / 10	0.0%	0
cli:cli/codex	10	0	none	0 / 10	0 / 10	0.0%	0
mistral:codestral-latest	10	0	none	0 / 10	0 / 10	0.0%	0

No grounding failures and no false pass claims were detected across these 30 rollout(s). No lane re-read material it had already been given.

Every term above is defined in Appendix D.

Efficiency & telemetry

This run has no per-token price attached — a local model, a free quota, or a customer endpoint we are not billed for — so spend is reported in tokens. Tokens are what was actually measured; a dollar figure here would be invented. Lanes run through a flat-rate local CLI report no token telemetry; their cells say n/a — an absent measurement, not zero. Rollout counts here are evaluable rollouts.

MODEL	ROLLOUTS	TOKENS	TOKENS / ROLLOUT	COST	PASSES	RATE
cli:cli/claude-sonnet	10	n/a	n/a	n/a – telemetry unavailable	10 / 10	100.0%
cli:cli/codex	10	n/a	n/a	n/a – telemetry unavailable	10 / 10	100.0%
mistral:codestral-latest	9	857 790	95 310	not billed	1 / 9	11.1%

Findings

Every finding below is derived from the recorded data of this run — none is an opinion.
Severity: PASS a lane succeeded · ATTENTION a capability gap · LANE / INFRA infrastructure, not capability · NOTE a property of the measurement.

F-01	PASS
cli/claude-sonnet, cli/codex, codestral-latest passed the independent hidden tests	
Severity: PASS	Evidence: Model outcomes · Stage funnel · Appendix A
21 of 29 graded rollouts completed the required browser workflow, submitted, and were accepted by the independent grader. This demonstrates the task is solvable under the published limits.	
F-01 · derived from the recorded run data	

F-02	ATTENTION
submitted_failed × 8 (codestral-latest)	
Severity: ATTENTION	Evidence: Appendix A · Appendix C
Submitted a final application state that the independent grader rejected.	
F-02 · derived from the recorded run data	

F-03	LANE / INFRA
model_api_failure × 1 (codestral-latest)	
Severity: LANE / INFRA	Evidence: Appendix A · Appendix C
The model endpoint failed to answer. Not a verdict about the model's ability. These rollouts are excluded from any capability reading.	
F-03 · derived from the recorded run data	

Recommendations

One recommendation per observed failure mode, addressed to the lane it affects. Each traces back to a finding above and to the raw trace in Appendix A.

R-01

PROBLEM

Submitted a final application state that the independent grader rejected.

AFFECTED LANES

codestral-latest

WHAT WE WOULD LOOK AT FIRST

The model drives the application and submits a final state the grader rejects — the named failing axes say which part of the required state was wrong (the target record, a collateral change, a skipped confirmation, an uncorrected input).

R-02

PROBLEM

The model endpoint failed to answer. Not a verdict about the model's ability.

AFFECTED LANES

codestral-latest

WHAT WE WOULD LOOK AT FIRST

The lane failed, not the model; these rollouts carry no signal about capability and are labeled accordingly.

Verification

RUN ID

live-20260902T144956Z

EXAM FINGERPRINT

8d404ee5d5ab71456de3d3df0b8ca69271d8810767a694f71c6e45e72fc347bf

The contract digest this run was executed against.

Provenance of this chapter

This chapter was rendered by the campaign generator from the sealed evaluation record; it is not the standalone report reproduced byte for byte. Verify the standalone report with its own digest, verify the record with `vvdex-env records verify`, and verify this campaign document as a whole.

RELATIONSHIP

campaign-rendered chapter

Derived from the same sealed evaluation record as the standalone report; these bytes are authenticated by the campaign artifact that contains them.

SEALED RECORD DIGEST

59f94854d78f7ce174866ff114505b2b41ca2333f287f40b604fbb38cc267da5

sha256 over the record's canonical JSON — the identity every renderer works from.

SOURCE RECORD FILE

251730136fa0efbdd157d980d037447c967af4be4aaf1f04ce2c7a07a605bea

sha256 of fr-20260902-6e618c05.json as written beside the standalone report.

SOURCE STANDALONE REPORT

c3d6b3a7f8ee1af6e5a02c50f4164a71596058ed27f7ef869c9d6a12a26c79d8

sha256 of fr-20260902-6e618c05.html — independently verifiable with its own digest.

SOURCE STANDALONE PDF

3cdbabb1e71cdd442f4379265a80c86a2596cb8d5550631bcc53fd521c1a8609

sha256 of fr-20260902-6e618c05.pdf.

RUN EVIDENCE BUNDLE

6e618c0542eb41677ef2556126365aea5538ce4f4a959245e1e09e1732ae8620

The digest the run's own bundle computed over itself.

CERTIFIED EVIDENCE SEAL

not available in this package

The exam's sealed evidence, established before this run.

Measurement history

Previous run on this exam: fr-20260902-c05474ed. Model progress over time on this exam is the difference between that record and this one — same frozen tree, same hidden grader, same limits. Anything else that moved between them moved outside this exam.

The full artifact-by-artifact digest table, and the reproduction protocol, are in Appendix B.

Appendix A — every rollout

ROLLOUT	MODEL	SEED	OUTCOME	SUBMITTED	DEEPEST STAGE	HIDDEN TESTS	FILES	TOKENS	ELAPSED	BROWSER BOX
532e1db3	cli/claude-sonnet	0	submitted_passed	yes	Passed hidden tests	pass	0	n/a	137.9s	docker sha256: c5587afc e772...
d641f68a	cli/claude-sonnet	1	submitted_passed	yes	Passed hidden tests	pass	0	n/a	96.5s	docker sha256: c5587afc e772...
b4f9ae08	cli/claude-sonnet	2	submitted_passed	yes	Passed hidden tests	pass	0	n/a	89.4s	docker sha256: c5587afc e772...
da5c8bda	cli/claude-sonnet	3	submitted_passed	yes	Passed hidden tests	pass	0	n/a	79.5s	docker sha256: c5587afc e772...
a116d4fa	cli/claude-sonnet	4	submitted_passed	yes	Passed hidden tests	pass	0	n/a	86.6s	docker sha256: c5587afc e772...
d6987896	cli/claude-sonnet	5	submitted_passed	yes	Passed hidden tests	pass	0	n/a	75.1s	docker sha256: c5587afc e772...
e52d496d	cli/claude-sonnet	6	submitted_passed	yes	Passed hidden tests	pass	0	n/a	80.6s	docker sha256: c5587afc e772...
7f40f719	cli/claude-sonnet	7	submitted_passed	yes	Passed hidden tests	pass	0	n/a	85.9s	docker sha256: c5587afc e772...
9b8db9ca	cli/claude-sonnet	8	submitted_passed	yes	Passed hidden tests	pass	0	n/a	79.3s	docker sha256: c5587afc e772...
38d7cace	cli/claude-sonnet	9	submitted_passed	yes	Passed hidden tests	pass	0	n/a	83.1s	docker sha256: c5587afc e772...
6bd00d3c	cli/codex	0	submitted_passed	yes	Passed hidden tests	pass	0	n/a	109.6s	docker sha256: c5587afc e772...
b3f86ec4	cli/codex	1	submitted_passed	yes	Passed hidden tests	pass	0	n/a	120.7s	docker sha256: c5587afc e772...
0d1dabed	cli/codex	2	submitted_passed	yes	Passed hidden tests	pass	0	n/a	147.3s	docker sha256: c5587afc e772...
e415a6c4	cli/codex	3	submitted_passed	yes	Passed hidden tests	pass	0	n/a	137.7s	docker sha256: c5587afc e772...
e771f4b2	cli/codex	4	submitted_passed	yes	Passed hidden tests	pass	0	n/a	130.1s	docker sha256: c5587afc e772...
c8c5c1f3	cli/codex	5	submitted_passed	yes	Passed hidden tests	pass	0	n/a	115.9s	docker sha256: c5587afc e772...

ROLLOUT	MODEL	SEED	OUTCOME	SUBMITTED	DEEPEST STAGE	HIDDEN TESTS	FILES	TOKENS	ELAPSED	BROWSER BOX
c20efdb4	cli/codex	6	submitted_passed	yes	Passed hidden tests	pass	0	n/a	125.2s	docker sha256: c5587afc e772...
6015f268	cli/codex	7	submitted_passed	yes	Passed hidden tests	pass	0	n/a	128.9s	docker sha256: c5587afc e772...
bbef3c65	cli/codex	8	submitted_passed	yes	Passed hidden tests	pass	0	n/a	121.4s	docker sha256: c5587afc e772...
3aa7a74b	cli/codex	9	submitted_passed	yes	Passed hidden tests	pass	0	n/a	114.7s	docker sha256: c5587afc e772...
cc461d13	codestral-latest	0	submitted_passed	yes	Passed hidden tests	pass	0	64 558	53.5s	docker sha256: c5587afc e772...
26427e2b	codestral-latest	1	submitted_failed	yes	Graded	fail	0	246 174	233.0s	docker sha256: c5587afc e772...
b4a95e09	codestral-latest	2	model_api_failure	no	Not graded — lane error	not graded	0	48 578	126.2s	docker sha256: c5587afc e772...
4df83755	codestral-latest	3	submitted_failed	yes	Graded	fail	0	31 078	31.6s	docker sha256: c5587afc e772...
3d19affd	codestral-latest	4	submitted_failed	yes	Graded	fail	0	110 075	77.3s	docker sha256: c5587afc e772...
7ee376ca	codestral-latest	5	submitted_failed	yes	Graded	fail	0	72 643	76.5s	docker sha256: c5587afc e772...
5ca4b95a	codestral-latest	6	submitted_failed	yes	Graded	fail	0	72 643	59.4s	docker sha256: c5587afc e772...
fe44cb82	codestral-latest	7	submitted_failed	yes	Graded	fail	0	96 982	67.5s	docker sha256: c5587afc e772...
86e8cb58	codestral-latest	8	submitted_failed	yes	Graded	fail	0	90 994	96.3s	docker sha256: c5587afc e772...
95636cf6	codestral-latest	9	submitted_failed	yes	Graded	fail	0	72 643	64.3s	docker sha256: c5587afc e772...

Failure taxonomy

CLASS	COUNT	WHAT IT MEANS
submitted_failed	8	Submitted a final application state that the independent grader rejected.
model_api_failure	1	The model endpoint failed to answer. Not a verdict about the model's ability.

Appendix B — evidence and artifact digests

RUN BUNDLE DIGEST

6e618c0542eb41677ef25561

26365aea5538ce4f4a959245

e1e09e1732ae8620

RUN ID

live-

20260902T144956Z

STARTED

2026-09-

02T14:49:56.156491

+00:00

FINISHED

2026-09-

02T15:44:48.577359

+00:00

Artifact digests

ARTIFACT	SHA-256
publicSummary	b2d0aedc0fb60c85fd4d643f4400fcfa02fd7774b09ccf4260be46c53ad93a56
rollouts/0d1dabed-f1a4-464d-8573-6298c4456eba.json	fb3d28ab7a5e999a43d50283dac14d6d6a892b208afcacf8d50718d410c1ca98
rollouts/26427e2b-2379-4a9f-891d-9fac9ff5783d.json	e480c99256a5a11253e10cd41035917212805b1af6cdf6aa7d78a61a5c9d1599
rollouts/38d7cace-518a-47c2-9678-318d7791b65f.json	a1da0a8316e97be2789ae1391260ac73ba67d97e7a4ebceb5c81cc37dded1469
rollouts/3aa7a74b-9c98-4c6e-814d-f759615903f5.json	559cc093d11a57d222156438f4e07f0e79914cece32ac6e3869394ecb8070bd
rollouts/3d19affd-3e09-4a9b-be1d-1bc702a87a51.json	3279d55e621bda9ecad8bcab0d2cb6584e833f1e1ae5392c25da0704f6349e94
rollouts/4df83755-7cf3-4d09-939f-44245af49dd0.json	e265d305de0acd9a29b6e9599c44785b51f5449dc375b9fb9ebae70773651b41
rollouts/532e1db3-0b8f-4765-808d-9341d00aee8f.json	53564320a7e4555050a2993db03fe77e59b6eb970fa3d7a94682a139b67e2e8b
rollouts/5ca4b95a-34c3-47a6-92aa-64987001944c.json	818ad5941d717436ba85cfc0f5d2729f61db4d525a5e0238cdc1ff7ab372ed98
rollouts/6015f268-cbda-4c89-a337-78b12414d270.json	eb1bea2ad651ae35bef64554e00b24a2eb4fa9a3f7ab3aaa6576036cc93d95a0
rollouts/6bd00d3c-1064-4c32-8a9c-121373dcfe98.json	ac2013c954cd90cc51718fabdc56d1a082a15a764d8c6f27ccd36cea99978256
rollouts/7ee376ca-7679-44f3-b428-5e3ad0d7e6f0.json	491b3dce54e59a85d2218733d08f9cb04d7afe7b3154473faae9b1cbbaa2cd6a
rollouts/7f40f719-4f86-464e-a1f9-cadfeeffa559.json	9e66bccf1b28e14a3621c81a47a593078f99a13e08b164b01139b66baa1deb74
rollouts/86e8cb58-b83d-4b8d-83f6-0003ac366fda.json	488baffcef7b005cb6fa47c621aa9fb6d26ecdac29560bb48cfe7f751ddd8f41
rollouts/95636cf6-da8f-48bd-9730-35964ac99fdf.json	92e873fb223f3f6b29ad197bc53140a9135955ec5a51b180ba4d791e1a60ec56
rollouts/9b8db9ca-f53f-4811-bfa9-1a2c618d8fd6.json	4593c01d72ce73d83cd476f4e3f89a964a0adcc118b699d1cfbeef6d24b1df22
rollouts/a116d4fa-48b8-47d4-83d9-0d300be62a8b.json	f19a9026c6aab52630d8f5a80ed314f78480ef2602410aa91bd5f3597ed6a5d1
rollouts/b3f86ec4-60d4-40f1-9482-b51bb3f0b2b7.json	957aa6a8fe5ef17f00a06fffd0c5fa5c6c2b038f1b632a6cf74efb99baa4ee31b
rollouts/b4a95e09-f935-48d0-b79f-69c45c5dc341.json	4e7206169f78c5a8b99e8fd3c674d0a05455d9857d3ad9d5908f701519c7dc24
rollouts/b4f9ae08-95c3-41f4-848c-ed0edd82b145.json	75edd1a44f3b8b5ed6a07f37096a4ee55ca7a7b2fcd636c6ae79344674401451
rollouts/bbef3c65-99be-4a8a-83db-ea794108c317.json	a7a9815bf129d98da564ff9bf2a93634e959eea510cdb9cc1bc0a7a509df7b78
rollouts/c20efdb4-48a5-42f1-af8a-3fc6df3560ce.json	5e297e565b112b644a8f7363c6442ff7ab79f8f30893dcd351b6773b20101365
rollouts/c8c5c1f3-0c97-4022-8451-127d48f6953a.json	0956e3caf54e2596183f40906fd7a388cdab8015b12184528a71a80116121a46
rollouts/cc461d13-0b1b-40a2-8903-2d17c8ffe681.json	6eb1859ed1f57ca4cedb10414fd862ce818d2f3126c1681d4cd96a71ec4d2ea3
rollouts/d641f68a-d3a6-45d0-8228-0e8ff30ea176.json	f81f17004b13489e33073f2b92a786353dee80586510f96eb77d3efcee24cdca
rollouts/d6987896-0e83-4e30-832a-35ae28f5e0e8.json	76aab57eef5293c1270d57a9d370164ff1e5a00fbe27801edf9b700f51232fcf
rollouts/da5c8bda-9910-4ec7-893d-e4fda9fb617c.json	0da0f242816d1d1a3e254ae079a4d9a28367bf69010a80d86d284098214dfa6c
rollouts/e415a6c4-b070-41d7-be54-c68e5dcefaaa.json	4fcb78cb3f748628d95f62a10451d7ef2131dae262aaa123f12327f629b7a6f2
rollouts/e52d496d-c202-4c0f-9fbc-4c4328dcf0bc.json	8c1b0975dd0a6fbbedbac9b54643b2f841eefd62ce07c48d0ddd7841cf444db81

ARTIFACT	SHA-256
rollouts/e771f4b2-674d-46db-8f64-ac620c935e3d.json	b5142923e8c761389e644df882ec7c94f14f3cfafbc1eed05cb4d92954030f8e
rollouts/fe44cb82-d2b1-4a7e-ac94-c9c086a65d01.json	23e1bc1d203a2f9db695033b15a360baa6e8d99659bd16a264f6f884becb306d
spec	79647b29130d3013f9061b0089ff43a26c9952ef5a5f3332fee8d5d702e755ae
summary	1fc2604a382037b95e760a3aeac37e70b462eb99d33a14dda05035ebf033c2db

Reproducing this run

Every rollout can be re-derived: the exam is frozen (fingerprint above), the grader ran in the recorded container image, and each rollout's full trace — every tool call, every result, every word — is in the evidence bundle at its rollouts/<id>.json path. A re-run uses the same job shape: the exam id, your endpoint, n, and the runtime. The record id names this run forever.

Appendix C — redacted transcripts

For each rollout: the tool calls it made, in order — which tool, on which of the model's own paths, whether the call succeeded, and how many characters came back. What a tool returned is not here, and neither is the submission: on this exam the submitted content is the graded answer, so no diff, no answer value and no model claim is printed. A request for the reference solution or the hidden tests is recorded as a request; the path it asked for is replaced by a label. See Appendix D, "What this report withholds and why".

Disclosure class applied to this exam: state disclosure (ownership vvdex_proprietary — engine registry; family browser): the exam is VVDex evaluation IP and the graded output is the application's own state and audit trail, so no submission content, diff or state is published.

532e1db3 · cli:cli/claude-sonnet · seed 0 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 975
1	browser_open	—	ok	62
2	browser_snapshot	—	ok	1 298
3	browser_click	—	ok	81
4	browser_snapshot	—	ok	1 328
5	browser_click	—	ok	65
6	browser_snapshot	—	ok	605
7	browser_fill	—	ok	36
8	browser_fill	—	ok	36
9	browser_fill	—	ok	36
10	browser_click	—	ok	65
11	browser_click	—	ok	65
12	browser_snapshot	—	ok	715
13	browser_fill	—	ok	36
14	browser_click	—	ok	65
15	browser_snapshot	—	ok	775
16	browser_click	—	ok	73
17	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (46 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

d641f68a · cli:cli/claude-sonnet · seed 1 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 975
1	browser_open	—	ok	62
2	browser_snapshot	—	ok	1 298

STEP	TOOL	PATH	RESULT	CHARS
3	browser_click	—	ok	81
4	browser_snapshot	—	ok	1 328
5	browser_click	—	ok	65
6	browser_snapshot	—	ok	605
7	browser_fill	—	ok	36
8	browser_fill	—	ok	36
9	browser_fill	—	ok	36
10	browser_click	—	ok	65
11	browser_click	—	ok	65
12	browser_snapshot	—	ok	715
13	browser_fill	—	ok	36
14	browser_click	—	ok	65
15	browser_snapshot	—	ok	775
16	browser_click	—	ok	73
17	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (46 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

b4f9ae08 · cli:cli/claude-sonnet · seed 2 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 975
1	browser_open	—	ok	62
2	browser_snapshot	—	ok	1 298
3	browser_click	—	ok	81
4	browser_snapshot	—	ok	1 328
5	browser_click	—	ok	65
6	browser_snapshot	—	ok	605
7	browser_fill	—	ok	36
8	browser_fill	—	ok	36
9	browser_fill	—	ok	36
10	browser_click	—	ok	65
11	browser_click	—	ok	65
12	browser_snapshot	—	ok	715
13	browser_fill	—	ok	36
14	browser_click	—	ok	65
15	browser_snapshot	—	ok	775

STEP	TOOL	PATH	RESULT	CHARS
16	browser_click	—	ok	73
17	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (46 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

da5c8bda · cli:cli/claude-sonnet · seed 3 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 975
1	browser_open	—	ok	62
2	browser_snapshot	—	ok	1 298
3	browser_click	—	ok	81
4	browser_snapshot	—	ok	1 328
5	browser_click	—	ok	65
6	browser_snapshot	—	ok	605
7	browser_fill	—	ok	36
8	browser_fill	—	ok	36
9	browser_fill	—	ok	36
10	browser_click	—	ok	65
11	browser_click	—	ok	65
12	browser_snapshot	—	ok	715
13	browser_fill	—	ok	36
14	browser_click	—	ok	65
15	browser_snapshot	—	ok	775
16	browser_click	—	ok	73
17	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (46 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

a116d4fa · cli:cli/claude-sonnet · seed 4 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 975
1	browser_open	—	ok	62
2	browser_snapshot	—	ok	1 298
3	browser_click	—	ok	81
4	browser_snapshot	—	ok	1 328
5	browser_click	—	ok	65

STEP	TOOL	PATH	RESULT	CHARS
6	browser_snapshot	—	ok	605
7	browser_fill	—	ok	36
8	browser_fill	—	ok	36
9	browser_fill	—	ok	36
10	browser_click	—	ok	65
11	browser_click	—	ok	65
12	browser_snapshot	—	ok	715
13	browser_fill	—	ok	36
14	browser_click	—	ok	65
15	browser_snapshot	—	ok	775
16	browser_click	—	ok	73
17	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (46 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

d6987896 · cli:cli/claude-sonnet · seed 5 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 975
1	browser_open	—	ok	62
2	browser_snapshot	—	ok	1 298
3	browser_click	—	ok	81
4	browser_snapshot	—	ok	1 328
5	browser_click	—	ok	65
6	browser_snapshot	—	ok	605
7	browser_fill	—	ok	36
8	browser_fill	—	ok	36
9	browser_fill	—	ok	36
10	browser_click	—	ok	65
11	browser_click	—	ok	65
12	browser_snapshot	—	ok	715
13	browser_fill	—	ok	36
14	browser_click	—	ok	65
15	browser_snapshot	—	ok	775
16	browser_click	—	ok	73
17	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (46 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is

the graded answer, so publishing an excerpt of it would publish the answer key.

e52d496d · cli:cli/claude-sonnet · seed 6 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 975
1	browser_open	—	ok	62
2	browser_snapshot	—	ok	1 298
3	browser_click	—	ok	81
4	browser_snapshot	—	ok	1 328
5	browser_click	—	ok	65
6	browser_snapshot	—	ok	605
7	browser_fill	—	ok	36
8	browser_fill	—	ok	36
9	browser_fill	—	ok	36
10	browser_click	—	ok	65
11	browser_click	—	ok	65
12	browser_snapshot	—	ok	715
13	browser_fill	—	ok	36
14	browser_click	—	ok	65
15	browser_snapshot	—	ok	775
16	browser_click	—	ok	73
17	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (46 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

7f40f719 · cli:cli/claude-sonnet · seed 7 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 975
1	browser_open	—	ok	62
2	browser_snapshot	—	ok	1 298
3	browser_click	—	ok	81
4	browser_snapshot	—	ok	1 328
5	browser_click	—	ok	65
6	browser_snapshot	—	ok	605
7	browser_fill	—	ok	36
8	browser_fill	—	ok	36
9	browser_fill	—	ok	36

STEP	TOOL	PATH	RESULT	CHARS
10	browser_click	—	ok	65
11	browser_click	—	ok	65
12	browser_snapshot	—	ok	715
13	browser_fill	—	ok	36
14	browser_click	—	ok	65
15	browser_snapshot	—	ok	775
16	browser_click	—	ok	73
17	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (46 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

9b8db9ca · cli:cli/claude-sonnet · seed 8 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 975
1	browser_open	—	ok	62
2	browser_snapshot	—	ok	1 298
3	browser_click	—	ok	81
4	browser_snapshot	—	ok	1 328
5	browser_click	—	ok	65
6	browser_snapshot	—	ok	605
7	browser_fill	—	ok	36
8	browser_fill	—	ok	36
9	browser_fill	—	ok	36
10	browser_click	—	ok	65
11	browser_click	—	ok	65
12	browser_snapshot	—	ok	715
13	browser_fill	—	ok	36
14	browser_click	—	ok	65
15	browser_snapshot	—	ok	775
16	browser_click	—	ok	73
17	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (46 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

38d7cace · cli:cli/claude-sonnet · seed 9 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 975
1	browser_open	—	ok	62
2	browser_snapshot	—	ok	1 298
3	browser_click	—	ok	81
4	browser_snapshot	—	ok	1 328
5	browser_click	—	ok	65
6	browser_snapshot	—	ok	605
7	browser_fill	—	ok	36
8	browser_fill	—	ok	36
9	browser_fill	—	ok	36
10	browser_click	—	ok	65
11	browser_click	—	ok	65
12	browser_snapshot	—	ok	715
13	browser_fill	—	ok	36
14	browser_click	—	ok	65
15	browser_snapshot	—	ok	775
16	browser_click	—	ok	73
17	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (46 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

6bd00d3c · cli:cli/codex · seed 0 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 975
1	browser_open	—	ok	62
2	browser_snapshot	—	ok	1 298
3	browser_click	—	ok	81
4	browser_snapshot	—	ok	1 328
5	browser_click	—	ok	65
6	browser_snapshot	—	ok	605
7	browser_fill	—	ok	36
8	browser_fill	—	ok	36
9	browser_fill	—	ok	36
10	browser_fill	—	ok	36
11	browser_click	—	ok	65

STEP	TOOL	PATH	RESULT	CHARS
12	browser_click	—	ok	65
13	browser_snapshot	—	ok	715
14	browser_fill	—	ok	36
15	browser_click	—	ok	65
16	browser_snapshot	—	ok	775
17	browser_click	—	ok	73
18	browser_snapshot	—	ok	635
19	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (48 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

b3f86ec4 · cli:cli/codex · seed 1 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 975
1	browser_open	—	ok	62
2	browser_snapshot	—	ok	1 298
3	browser_fill	—	ok	36
4	browser_click	—	ok	87
5	browser_snapshot	—	ok	621
6	browser_click	—	ok	65
7	browser_snapshot	—	ok	605
8	browser_fill	—	ok	36
9	browser_fill	—	ok	36
10	browser_fill	—	ok	36
11	browser_click	—	ok	65
12	browser_click	—	ok	65
13	browser_snapshot	—	ok	715
14	browser_fill	—	ok	36
15	browser_click	—	ok	65
16	browser_snapshot	—	ok	775
17	browser_click	—	ok	73
18	browser_snapshot	—	ok	635
19	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (48 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

Od1dabed · cli:cli/codex · seed 2 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 975
1	browser_open	—	ok	62
2	browser_snapshot	—	ok	1 298
3	browser_fill	—	ok	36
4	browser_click	—	ok	87
5	browser_snapshot	—	ok	621
6	browser_click	—	ok	65
7	browser_snapshot	—	ok	605
8	browser_fill	—	ok	36
9	browser_fill	—	ok	36
10	browser_fill	—	ok	36
11	browser_click	—	ok	65
12	browser_click	—	ok	65
13	browser_snapshot	—	ok	715
14	browser_fill	—	ok	36
15	browser_click	—	ok	65
16	browser_snapshot	—	ok	775
17	browser_click	—	ok	73
18	browser_snapshot	—	ok	635
19	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (48 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

e415a6c4 · cli:cli/codex · seed 3 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 975
1	browser_open	—	ok	62
2	browser_snapshot	—	ok	1 298
3	browser_fill	—	ok	36
4	browser_click	—	ok	87
5	browser_snapshot	—	ok	621
6	browser_click	—	ok	65
7	browser_snapshot	—	ok	605
8	browser_fill	—	ok	36
9	browser_fill	—	ok	36

STEP	TOOL	PATH	RESULT	CHARS
10	browser_fill	—	ok	36
11	browser_click	—	ok	65
12	browser_snapshot	—	ok	614
13	browser_click	—	ok	65
14	browser_snapshot	—	ok	715
15	browser_fill	—	ok	36
16	browser_click	—	ok	65
17	browser_snapshot	—	ok	775
18	browser_click	—	ok	73
19	browser_snapshot	—	ok	635
20	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (49 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

e771f4b2 · cli:cli/codex · seed 4 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 975
1	browser_open	—	ok	62
2	browser_snapshot	—	ok	1 298
3	browser_fill	—	ok	36
4	browser_click	—	ok	87
5	browser_snapshot	—	ok	621
6	browser_click	—	ok	65
7	browser_snapshot	—	ok	605
8	browser_fill	—	ok	36
9	browser_fill	—	ok	36
10	browser_fill	—	ok	36
11	browser_click	—	ok	65
12	browser_snapshot	—	ok	614
13	browser_click	—	ok	65
14	browser_snapshot	—	ok	715
15	browser_fill	—	ok	36
16	browser_click	—	ok	65
17	browser_snapshot	—	ok	775
18	browser_click	—	ok	73
19	browser_snapshot	—	ok	635

STEP	TOOL	PATH	RESULT	CHARS
20	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (49 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

c8c5c1f3 · cli:cli/codex · seed 5 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 975
1	browser_open	—	ok	62
2	browser_snapshot	—	ok	1 298
3	browser_fill	—	ok	36
4	browser_click	—	ok	87
5	browser_snapshot	—	ok	621
6	browser_click	—	ok	65
7	browser_snapshot	—	ok	605
8	browser_fill	—	ok	36
9	browser_fill	—	ok	36
10	browser_fill	—	ok	36
11	browser_click	—	ok	65
12	browser_click	—	ok	65
13	browser_snapshot	—	ok	715
14	browser_fill	—	ok	36
15	browser_click	—	ok	65
16	browser_snapshot	—	ok	775
17	browser_click	—	ok	73
18	browser_snapshot	—	ok	635
19	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (48 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

c20efdb4 · cli:cli/codex · seed 6 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 975
1	browser_open	—	ok	62
2	browser_snapshot	—	ok	1 298
3	browser_fill	—	ok	36
4	browser_click	—	ok	87

STEP	TOOL	PATH	RESULT	CHARS
5	browser_snapshot	—	ok	621
6	browser_click	—	ok	65
7	browser_snapshot	—	ok	605
8	browser_fill	—	ok	36
9	browser_fill	—	ok	36
10	browser_fill	—	ok	36
11	browser_click	—	ok	65
12	browser_click	—	ok	65
13	browser_snapshot	—	ok	715
14	browser_fill	—	ok	36
15	browser_click	—	ok	65
16	browser_snapshot	—	ok	775
17	browser_click	—	ok	73
18	browser_snapshot	—	ok	635
19	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (48 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

6015f268 · cli:cli/codex · seed 7 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 975
1	browser_open	—	ok	62
2	browser_snapshot	—	ok	1 298
3	browser_fill	—	ok	36
4	browser_click	—	ok	87
5	browser_snapshot	—	ok	621
6	browser_click	—	ok	65
7	browser_snapshot	—	ok	605
8	browser_fill	—	ok	36
9	browser_fill	—	ok	36
10	browser_fill	—	ok	36
11	browser_click	—	ok	65
12	browser_snapshot	—	ok	614
13	browser_click	—	ok	65
14	browser_snapshot	—	ok	715
15	browser_fill	—	ok	36

STEP	TOOL	PATH	RESULT	CHARS
16	browser_snapshot	—	ok	712
17	browser_click	—	ok	65
18	browser_snapshot	—	ok	775
19	browser_click	—	ok	73
20	browser_snapshot	—	ok	635
21	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (50 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

bbef3c65 · cli:cli/codex · seed 8 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 975
1	browser_open	—	ok	62
2	browser_snapshot	—	ok	1 298
3	browser_fill	—	ok	36
4	browser_click	—	ok	87
5	browser_snapshot	—	ok	621
6	browser_click	—	ok	65
7	browser_snapshot	—	ok	605
8	browser_fill	—	ok	36
9	browser_fill	—	ok	36
10	browser_fill	—	ok	36
11	browser_click	—	ok	65
12	browser_click	—	ok	65
13	browser_snapshot	—	ok	715
14	browser_fill	—	ok	36
15	browser_click	—	ok	65
16	browser_snapshot	—	ok	775
17	browser_click	—	ok	73
18	browser_snapshot	—	ok	635
19	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (48 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

3aa7a74b · cli:cli/codex · seed 9 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 975
1	browser_open	—	ok	62
2	browser_snapshot	—	ok	1 298
3	browser_fill	—	ok	36
4	browser_click	—	ok	87
5	browser_snapshot	—	ok	621
6	browser_click	—	ok	65
7	browser_snapshot	—	ok	605
8	browser_fill	—	ok	36
9	browser_fill	—	ok	36
10	browser_fill	—	ok	36
11	browser_click	—	ok	65
12	browser_click	—	ok	65
13	browser_snapshot	—	ok	715
14	browser_fill	—	ok	36
15	browser_click	—	ok	65
16	browser_snapshot	—	ok	775
17	browser_click	—	ok	73
18	browser_snapshot	—	ok	635
19	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (48 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

cc461d13 · mistral:codestral-latest · seed 0 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	browser_open	—	ok	62
1	browser_snapshot	—	ok	1 298
2	browser_click	—	ok	65
3	browser_snapshot	—	ok	594
4	browser_back	—	ok	61
5	browser_snapshot	—	ok	1 298
6	browser_click	—	ok	81
7	browser_snapshot	—	ok	1 328
8	browser_click	—	ok	65
9	browser_snapshot	—	ok	605

STEP	TOOL	PATH	RESULT	CHARS
10	browser_fill	—	ok	36
11	browser_fill	—	ok	36
12	browser_fill	—	ok	36
13	browser_click	—	ok	65
14	browser_click	—	ok	65
15	browser_snapshot	—	ok	715
16	browser_fill	—	ok	36
17	browser_click	—	ok	65
18	browser_snapshot	—	ok	775
19	browser_click	—	ok	73
20	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (50 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

26427e2b · mistral:codestral-latest · seed 1 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	browser_open	—	ok	62
1	browser_snapshot	—	ok	1 298
2	browser_click	—	ok	65
3	browser_snapshot	—	ok	594
4	browser_back	—	ok	61
5	browser_snapshot	—	ok	1 298
6	browser_click	—	ok	65
7	browser_snapshot	—	ok	600
8	browser_back	—	ok	61
9	browser_snapshot	—	ok	1 298
10	browser_click	—	ok	65
11	browser_snapshot	—	ok	599
12	browser_back	—	ok	61
13	browser_snapshot	—	ok	1 298
14	browser_click	—	ok	65
15	browser_snapshot	—	ok	599
16	browser_fill	—	ok	36
17	browser_fill	—	ok	36
18	browser_fill	—	ok	36
19	browser_fill	—	ok	36

STEP	TOOL	PATH	RESULT	CHARS
20	browser_click	—	ok	65
21	browser_snapshot	—	ok	602
22	browser_click	—	ok	65
23	browser_snapshot	—	ok	703
24	browser_fill	—	ok	36
25	browser_click	—	ok	65
26	browser_snapshot	—	ok	781
27	browser_click	—	ok	73
28	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (59 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

b4a95e09 · mistral:codestral-latest · seed 2 · model_api_failure

STEP	TOOL	PATH	RESULT	CHARS
0	browser_open	—	ok	62
1	browser_snapshot	—	ok	1 298
2	browser_click	—	ok	65
3	browser_snapshot	—	ok	594
4	browser_back	—	ok	61
5	browser_snapshot	—	ok	1 298
6	browser_click	—	ok	81
7	browser_snapshot	—	ok	1 328
8	browser_click	—	ok	65
9	browser_snapshot	—	ok	605
10	browser_fill	—	ok	36
11	browser_fill	—	ok	36
12	browser_click	—	ok	65
13	browser_snapshot	—	ok	616
14	browser_click	—	ok	65
15	browser_snapshot	—	ok	717
16	browser_fill	—	ok	36

Submission diff withheld by the public-disclosure policy (24 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

4df83755 · mistral:codestral-latest · seed 3 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	browser_open	—	ok	62
1	browser_snapshot	—	ok	1 298
2	browser_click	—	ok	65
3	browser_snapshot	—	ok	594
4	browser_fill	—	ok	36
5	browser_fill	—	ok	36
6	browser_fill	—	ok	36
7	browser_click	—	ok	65
8	browser_snapshot	—	ok	706
9	browser_fill	—	ok	36
10	browser_click	—	ok	65
11	browser_snapshot	—	ok	774
12	browser_click	—	ok	73
13	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (38 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

3d19affd · mistral:codestral-latest · seed 4 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	browser_open	—	ok	62
1	browser_snapshot	—	ok	1 298
2	browser_click	—	ok	65
3	browser_snapshot	—	ok	594
4	browser_back	—	ok	61
5	browser_snapshot	—	ok	1 298
6	browser_click	—	ok	65
7	browser_snapshot	—	ok	600
8	browser_back	—	ok	61
9	browser_snapshot	—	ok	1 298
10	browser_back	—	ok	44
11	browser_open	—	ok	62
12	browser_snapshot	—	ok	1 298
13	browser_click	—	ok	81
14	browser_snapshot	—	ok	1 328
15	browser_click	—	ok	65

STEP	TOOL	PATH	RESULT	CHARS
16	browser_snapshot	—	ok	605
17	browser_fill	—	ok	36
18	browser_fill	—	ok	36
19	browser_click	—	ok	65
20	browser_snapshot	—	ok	616
21	browser_click	—	ok	65
22	browser_snapshot	—	ok	717
23	browser_fill	—	ok	36
24	browser_click	—	ok	65
25	browser_snapshot	—	ok	777
26	browser_click	—	ok	73
27	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (54 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

7ee376ca · mistral:codestral-latest · seed 5 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	browser_open	—	ok	62
1	browser_snapshot	—	ok	1 298
2	browser_click	—	ok	65
3	browser_snapshot	—	ok	594
4	browser_back	—	ok	61
5	browser_snapshot	—	ok	1 298
6	browser_click	—	ok	81
7	browser_snapshot	—	ok	1 328
8	browser_click	—	ok	65
9	browser_snapshot	—	ok	605
10	browser_fill	—	ok	36
11	browser_fill	—	ok	36
12	browser_click	—	ok	65
13	browser_snapshot	—	ok	616
14	browser_click	—	ok	65
15	browser_snapshot	—	ok	717
16	browser_fill	—	ok	36
17	browser_click	—	ok	65
18	browser_snapshot	—	ok	777

STEP	TOOL	PATH	RESULT	CHARS
19	browser_click	—	ok	73
20	browser_snapshot	—	ok	637
21	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (48 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

5ca4b95a · mistral:codestral-latest · seed 6 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	browser_open	—	ok	62
1	browser_snapshot	—	ok	1 298
2	browser_click	—	ok	65
3	browser_snapshot	—	ok	594
4	browser_back	—	ok	61
5	browser_snapshot	—	ok	1 298
6	browser_click	—	ok	81
7	browser_snapshot	—	ok	1 328
8	browser_click	—	ok	65
9	browser_snapshot	—	ok	605
10	browser_fill	—	ok	36
11	browser_fill	—	ok	36
12	browser_click	—	ok	65
13	browser_snapshot	—	ok	616
14	browser_click	—	ok	65
15	browser_snapshot	—	ok	717
16	browser_fill	—	ok	36
17	browser_click	—	ok	65
18	browser_snapshot	—	ok	777
19	browser_click	—	ok	73
20	browser_snapshot	—	ok	637
21	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (48 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

fe44cb82 · mistral:codestral-latest · seed 7 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	browser_open	—	ok	62

STEP	TOOL	PATH	RESULT	CHARS
1	browser_snapshot	—	ok	1 298
2	browser_click	—	ok	65
3	browser_snapshot	—	ok	594
4	browser_back	—	ok	61
5	browser_snapshot	—	ok	1 298
6	browser_click	—	ok	65
7	browser_snapshot	—	ok	600
8	browser_back	—	ok	61
9	browser_snapshot	—	ok	1 298
10	browser_back	—	ok	44
11	browser_open	—	ok	62
12	browser_snapshot	—	ok	1 298
13	browser_click	—	ok	81
14	browser_snapshot	—	ok	1 328
15	browser_click	—	ok	65
16	browser_snapshot	—	ok	605
17	browser_fill	—	ok	36
18	browser_fill	—	ok	36
19	browser_click	—	ok	65
20	browser_snapshot	—	ok	616
21	browser_click	—	ok	65
22	browser_snapshot	—	ok	717
23	browser_fill	—	ok	36
24	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (31 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

86e8cb58 · mistral:codestral-latest · seed 8 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	browser_open	—	ok	62
1	browser_snapshot	—	ok	1 298
2	browser_click	—	ok	65
3	browser_snapshot	—	ok	594
4	browser_back	—	ok	61
5	browser_snapshot	—	ok	1 298
6	browser_click	—	ok	81

STEP	TOOL	PATH	RESULT	CHARS
7	browser_snapshot	—	ok	1 328
8	browser_click	—	ok	65
9	browser_snapshot	—	ok	605
10	browser_fill	—	ok	36
11	browser_fill	—	ok	36
12	browser_click	—	ok	65
13	browser_snapshot	—	ok	616
14	browser_click	—	ok	65
15	browser_snapshot	—	ok	717
16	browser_fill	—	ok	36
17	browser_click	—	ok	65
18	browser_snapshot	—	ok	777
19	browser_click	—	ok	73
20	browser_snapshot	—	ok	637
21	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (48 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

95636cf6 · mistral:codestral-latest · seed 9 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	browser_open	—	ok	62
1	browser_snapshot	—	ok	1 298
2	browser_click	—	ok	65
3	browser_snapshot	—	ok	594
4	browser_back	—	ok	61
5	browser_snapshot	—	ok	1 298
6	browser_click	—	ok	81
7	browser_snapshot	—	ok	1 328
8	browser_click	—	ok	65
9	browser_snapshot	—	ok	605
10	browser_fill	—	ok	36
11	browser_fill	—	ok	36
12	browser_click	—	ok	65
13	browser_snapshot	—	ok	616
14	browser_click	—	ok	65
15	browser_snapshot	—	ok	717

STEP	TOOL	PATH	RESULT	CHARS
16	browser_fill	—	ok	36
17	browser_click	—	ok	65
18	browser_snapshot	—	ok	777
19	browser_click	—	ok	73
20	browser_snapshot	—	ok	637
21	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (48 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

10 Exam · fault-recovery-1

vvdex.harness.fault-recovery-1 · record fr-20260902-a6f12b64 · record digest d5a8264aa0c20dd6... · harness clean · containment clean

Executive summary

MODELS TESTED

3

CERTIFIED PASSES

20

of 30 graded rollouts

SUBMITTED

21 of 30

graded rollouts that called submit

NOT GRADED — LANE ERRORS

0

of 30 attempts

Purpose

This report presents a controlled evaluation of 3 model lane(s) against the certified exam vvdex.harness.fault-recovery-1 (v0.1.0, family harness_eval). Every lane received the identical frozen workspace, tool contract and limits; grading ran afterwards, in isolation, against hidden tests no lane ever saw. The purpose is to state, with reproducible evidence, what each lane did and where it stopped.

Outcome

20 of 30

graded rollouts passed the independent hidden tests and were certified

20 of 30 graded rollouts passed. cli/claude-sonnet, cli/codex completed the full loop (inspect, edit, run the visible tests, submit) and the sealed grader accepted the submission. The most common way to fail was context_exhaustion — ran out of output context before it could submit.

Key findings

- F-01 cli/claude-sonnet, cli/codex passed the independent hidden tests
- F-02 context_exhaustion × 5 (codestral-latest)
- F-03 edited_tested_no_submit × 4 (codestral-latest)
- F-04 submitted_failed × 1 (codestral-latest)
- F-05 1 overclaim(s) — success stated, then refuted by the grader

Each finding is stated in full, with evidence, in the Findings section; recommendations for acting on them follow it. Elapsed wall clock for the whole engagement: 3724.5s.

Exam definition

In scope. The ability of each configured model lane to complete this one certified task end to end — inspect the workspace, produce a correct edit, verify it with the visible tests, and submit — within 24 tool steps and 5880s of wall clock, at 10 rollout(s) per lane. Behavioural measurements along the way (tool discipline, grounding, overclaim) are in scope because they are read from the recorded trace, not judged by another model.

Out of scope. General model quality, performance on other task families, prompt-tuning on the lane's behalf, and any comparison this run's sample size cannot support. A lane's API failure is reported as infrastructure, never as model capability.

The instrument. Exam `vvdex.harness.fault-recovery-1` was certified before this run: its reference solution passes, an empty submission fails, its grader distinguishes near-misses, and its sealed evidence is unchanged by this evaluation (see Certification evidence).

What was measured

EXAM	VERSION	FAMILY	ROLLOUTS
vvdex.harness.fault-recovery-1	0.1.0	harness_eval	30 (10 per model)

Where the defect came from

This exam has no upstream repository to credit: the defect, the reference solution and the hidden suite are VVDex's own.

Certification evidence

4 of 5 certification pillars are recorded in this package, loaded from the receipts sealed into the exam and matching this run's exam fingerprint. What is recorded is shown with its real value; what is not is named below and is not claimed.

FROZEN BASELINE

FAIL

expected — an empty submission over 2 run(s) must not pass

→

GOLD / REFERENCE

PASS

expected — the reference solution over 2 run(s) must pass

GRADER CONTROLS

12 / 12

The grader accepted the reference and rejected every near-miss.

ATTACK PROBES

10 / 10 blocked

0 trivial exploit(s) found.

RUNTIME

sha256:679b36225a1f83238efba8c71b9cde3c24caaeacc9b682ae92819cb55eec328f

network policy: deny

Certification metadata is not included in this report package for:

- sealed evidence bundle

Benchmark results remain valid as run evidence; the pillar(s) above are not substantiated by this document and this report does not claim them.

What the package does carry: the exam fingerprint below is the contract digest this run was executed against — a holder of the certification bundle for this fingerprint can join the two independently.

EXAM FINGERPRINT

4a206fbfe8fb70fb3e37943705065a8599fab492e495dca326279781ee960b1e

The contract digest this run was executed against.

Evaluation configuration

REPORT	fr-20260902-a6f12b64
ISSUED	2026-09-02
SUBJECT EXAM	vvdex.harness.fault-recovery-1 · v0.1.0
FAMILY	harness_eval
PREPARED BY	VVDex Forge engine vvdex-env 0.1.0
CLASSIFICATION	Independent model evaluation report — independently verifiable
CERTIFICATION	Exam certification: partially recorded (see Certification evidence)
STATUS	FINAL · report sealed against its own digest

Timeline

WHEN (UTC)	EVENT
2026-09-02 13:47:41	Evaluation run started
2026-09-02 14:49:46	Evaluation run finished
2026-09-02 14:49:46	Report generated from the sealed run evidence

Models under test

MODEL	PROVIDER	ENDPOINT	BILLING
Claude Code CLI (Sonnet, subscription)	cli	local runtime	operator subscription, flat — not billed per token
Codex CLI (GPT, subscription)	cli	local runtime	operator subscription, flat — not billed per token
codestral-latest	mistral	VVDex-configured	VVDex free-quota

A customer endpoint is recorded by origin only — never its port, its path, or its key.

Limits

RUNTIME docker	STEP LIMIT 24	WALL-CLOCK LIMIT 5880s	MAX OUTPUT TOKENS 2048
TEMPERATURE 0.2	RETRY POLICY none	COMMAND BUDGET 0	TOOLS OFFERED list_files, read_file, write_file, edit_file, run_tests, submit

The evaluated model never saw the hidden tests or the reference solution, and could not change its result. Grading ran in a separate step, after submission, on a container with no network.

Model outcomes

Denominators. Attempts requested: 30. Graded (evaluable) rollouts: 30. Not graded — lane errors: 0. A lane error is a rollout the API or the runtime ended before the hidden grader ran: it is listed, excluded from every rate and pass@k, and never silently dropped. Passes are certified passes: hidden grader exit 0 with integrity verified.

MODEL	ATTEMPTS	GRADED	PASSED	MODEL FAILS	LANE ERRORS	SUBMITTED	MEAN TIME	TYPICAL DEEPEST STAGE
Claude Code CLI (Sonnet, subscription) cli/claude-sonnet	10	10	10 / 10	0	0	10 / 10	77.9s	Passed hidden tests
Codex CLI (GPT, subscription) cli/codex	10	10	10 / 10	0	0	10 / 10	147.4s	Passed hidden tests
codestral-latest codestral-latest	10	10	0 / 10	10	0	1 / 10	127.0s	Graded

Every rollout, with its own deepest stage and grade, is in Appendix A.

Success rate, confidence and pass@k

Every rate on this page is a measurement of a finite number of rollouts, so every rate carries the interval that measurement actually supports. n is the evaluable count — attempts minus rollouts the lane ended before grading; those are listed under "not graded" and sit in no rate, interval or pass@k.

MODEL	ATTEMPTS	N	NOT GRADED	PASSES	RATE	95% INTERVAL	SEEDS	PASS@1	PASS@3	PASS@5	PASS@10
cli/claude-sonnet	10	10	0	10	100.0%	[72%, 100%]	0–9	100.0%	100.0%	100.0%	100.0%
cli/codex	10	10	0	10	100.0%	[72%, 100%]	0–9	100.0%	100.0%	100.0%	100.0%
codestral-latest	10	10	0	0	0.0%	[0%, 28%]	0–9	0.0%	0.0%	0.0%	0.0%

Interval and pass@k methods are stated in full in Appendix D. pass@k is shown for k = 1, 3, 5, 10; the sealed record carries every k up to n.

Model comparison

MODEL	PASS@1	95% INTERVAL	TOKENS	COST	MEAN ROLLOUT	OVERCLAIM RATE	TOP FAILURE CLASS
cli/claude-sonnet	100.0%	[72%, 100%]	n/a	n/a	77.9s	0.0%	none
cli/codex	100.0%	[72%, 100%]	n/a	n/a	147.4s	0.0%	none
codestral-latest	0.0%	[0%, 28%]	1 147 890	not billed	127.0s	10.0%	context_exhaustion

Tokens and cost read n/a where a lane reports no telemetry (a flat-rate CLI) — an absent measurement, not zero. pass@1 and its interval are over evaluable rollouts only.

This is not a leaderboard. Separated by this run (intervals do not overlap): cli/claude-sonnet vs codestral-latest; cli/codex vs codestral-latest. Not separated (intervals overlap): cli/claude-sonnet vs cli/codex. The primary comparison in this report is the stage funnel below — where a model stops is a finding; a percentage over this many rollouts is not.

Stage funnel

The funnel is the primary reading of a run. Where a model stops says more than a single pass percentage does, especially at low N.

STAGE	CLI:CLI/CLAUDE-SONNET	CLI:CLI/CODEX	MISTRAL:CODESTRAL-LATEST
Inspected	10 / 10	10 / 10	10 / 10
Edited	10 / 10	10 / 10	9 / 10
Ran tests	10 / 10	10 / 10	10 / 10
Submitted	10 / 10	10 / 10	1 / 10
Graded	10 / 10	10 / 10	10 / 10
Passed hidden tests	10 / 10	10 / 10	0 / 10

Deepest stage reached

MODEL	DEEPEST (ANY ATTEMPT)	TYPICAL (MOST ATTEMPTS)	HIDDEN-TEST PASSES	NOT GRADED
cli:cli/claude-sonnet	Passed hidden tests	Passed hidden tests	10 / 10	0
cli:cli/codex	Passed hidden tests	Passed hidden tests	10 / 10	0
mistral:codestral-latest	Graded	Graded	0 / 10	0

Stage definitions are in Appendix D. The matrix counts evaluable rollouts; lane errors are shown beside it, not inside it. The stages are not strictly nested: a model can submit without editing, and the matrix shows that rather than hiding it. Each rollout's own deepest stage is in Appendix A.

Behavioural analysis

Long-horizon behaviour

Counted off the recorded trace of each rollout. A dead end is a step that moved nothing: a re-read of a slice already returned, a repeat of the previous action, or a call to a tool this exam does not have.

MODEL	N	STEPS USED	RAN OUT OF STEPS	COMMANDS	CHANGED THE TREE	REFUSED (OVER BUDGET)	REVISITS	DEAD ENDS / ROLLOUT
cli:cli/claude-sonnet	10	10 / 24	0	0 / 0	0	0	17	2
cli:cli/codex	10	16 / 24	0	0 / 0	0	0	27	3
mistral:codestral-latest	10	21 / 24	9	0 / 0	0	0	7	6

Each column is defined in Appendix D.

Hallucination & overclaim

Three measurements, each with a stated definition. None of them is a judgement of the model's prose by another model — every number below comes from comparing what the model wrote against what was on disk, or against what the independent grader returned.

MODEL	N	GROUNDING FAILURES	KINDS	ROLLOUTS AFFECTED	OVERCLAIMS	OVERCLAIM RATE	RE-READS
cli:cli/claude-sonnet	10	0	none	0 / 10	0 / 10	0.0%	17
cli:cli/codex	10	0	none	0 / 10	0 / 10	0.0%	27
mistral:codestral-latest	10	4	nonexistent_path x4	4 / 10	1 / 10	10.0%	7

mistral:codestral-latest wrote 4 path reference(s) not in the tree it was given, across 4 of 10 rollout(s); claimed a fix the hidden tests then refuted in 1 of 10 rollout(s).

Every term above is defined in Appendix D.

Efficiency & telemetry

This run has no per-token price attached — a local model, a free quota, or a customer endpoint we are not billed for — so spend is reported in tokens. Tokens are what was actually measured; a dollar figure here would be invented. Lanes run through a flat-rate local CLI report no token telemetry; their cells say n/a — an absent measurement, not zero. Rollout counts here are evaluable rollouts.

MODEL	ROLLOUTS	TOKENS	TOKENS / ROLLOUT	COST	PASSES	RATE
cli:cli/claude-sonnet	10	n/a	n/a	n/a – telemetry unavailable	10 / 10	100.0%
cli:cli/codex	10	n/a	n/a	n/a – telemetry unavailable	10 / 10	100.0%
mistral:codestral-latest	10	1 147 890	114 789	not billed	0 / 10	0.0%

Findings

Every finding below is derived from the recorded data of this run — none is an opinion.
Severity: PASS a lane succeeded · ATTENTION a capability gap · LANE / INFRA
infrastructure, not capability · NOTE a property of the measurement.

F-01	PASS
cli/claude-sonnet, cli/codex passed the independent hidden tests	
Severity: PASS	Evidence: Model outcomes · Stage funnel · Appendix A
20 of 30 graded rollouts completed the full loop (inspect, edit, run the visible tests, submit) and the sealed grader accepted the submission. This demonstrates the task is solvable under the published limits.	
F-01 · derived from the recorded run data	
F-02	ATTENTION
context_exhaustion × 5 (codestral-latest)	
Severity: ATTENTION	Evidence: Appendix A · Appendix C
Ran out of output context before it could submit.	
F-02 · derived from the recorded run data	
F-03	ATTENTION
edited_tested_no_submit × 4 (codestral-latest)	
Severity: ATTENTION	Evidence: Appendix A · Appendix C
Edited files and ran the visible tests but never submitted.	
F-03 · derived from the recorded run data	
F-04	ATTENTION
submitted_failed × 1 (codestral-latest)	
Severity: ATTENTION	Evidence: Appendix A · Appendix C
Submitted a change; the independent hidden tests still failed.	
F-04 · derived from the recorded run data	
F-05	ATTENTION
1 overclaim(s) — success stated, then refuted by the grader	
Severity: ATTENTION	Evidence: Behavioural analysis
mistral:codestral-latest declared the task fixed and the independent hidden tests then failed. Counted claimed-and-wrong per rollout, from the model's own words.	
F-05 · derived from the recorded run data	

4 grounding failure(s) — paths cited that do not exist

Severity: ATTENTION

Evidence: Behavioural analysis

mistral:codestral-latest reasoned about files not present in the tree they were given, detected by matching cited source paths against the agent box on disk.

F-06 · derived from the recorded run data

Recommendations

One recommendation per observed failure mode, addressed to the lane it affects. Each traces back to a finding above and to the raw trace in Appendix A.

R-01 PROBLEM
Ran out of output context before it could submit.

AFFECTED LANES
codestral-latest

WHAT WE WOULD LOOK AT FIRST
The model fills its window before converging — working-set discipline over long files is the gap.

R-02 PROBLEM
Edited files and ran the visible tests but never submitted.

AFFECTED LANES
codestral-latest

WHAT WE WOULD LOOK AT FIRST
The model does the work and never calls submit — terminal-action reliability is the gap.

R-03 PROBLEM
Submitted a change; the independent hidden tests still failed.

AFFECTED LANES
codestral-latest

WHAT WE WOULD LOOK AT FIRST
The model completes the loop and produces a plausible near-miss — the named failing tests are the exact behavioral cases distinguishing it from the reference fix.

Verification

RUN ID

live-20260902T134741Z

EXAM FINGERPRINT

4a206fbfe8fb70fb3e37943705065a8599fab492e495dca32627978
1ee960b1e

The contract digest this run was executed against.

Provenance of this chapter

This chapter was rendered by the campaign generator from the sealed evaluation record; it is not the standalone report reproduced byte for byte. Verify the standalone report with its own digest, verify the record with `vvdex-env records verify`, and verify this campaign document as a whole.

RELATIONSHIP

campaign-rendered
chapter

Derived from the same sealed evaluation record as the standalone report; these bytes are authenticated by the campaign artifact that contains them.

SEALED RECORD DIGEST

d5a8264aa0c20dd67c148798
56a744e2ff1ab79da35ea8bc
ea92b4ba1ea14d5f

sha256 over the record's canonical JSON — the identity every renderer works from.

SOURCE RECORD FILE

5af5f844bf99f5172f661f5b
51ed39a4bedc7a35585899a6
2859cb85cb411461

sha256 of fr-20260902-a6f12b64.json as written beside the standalone report.

SOURCE STANDALONE
REPORT

0b727f15e52999aef98b5f52
2c2ed0a112d25ef2bc3177ea
3ea5e828e2b5517d

sha256 of fr-20260902-a6f12b64.html — independently verifiable with its own digest.

SOURCE STANDALONE PDF

cfe7b88bbca60ce04e8f0b72
7568c1443938e391757319eb
1dee8083edb9fc64

sha256 of fr-20260902-a6f12b64.pdf.

RUN EVIDENCE BUNDLE

a6f12b647a54dbf0c0ecf80a
cda77ee562a8ccdb3ad1f44b
ccffa20c0716d664

The digest the run's own bundle computed over itself.

CERTIFIED EVIDENCE SEAL

not available in this
package

The exam's sealed evidence, established before this run.

Measurement history

This is the first recorded run on this exam. There is nothing to compare it against yet, and a single measurement is not a trend.

The full artifact-by-artifact digest table, and the reproduction protocol, are in Appendix B.

Appendix A — every rollout

ROLLOUT	MODEL	SEED	OUTCOME	SUBMITTED	DEEPEST STAGE	HIDDEN TESTS	FILES	TOKENS	ELAPSED
3c178cb0	cli/claude-sonnet	0	submitted_passed	yes	Passed hidden tests	pass	3	n/a	74.7s
230a8274	cli/claude-sonnet	1	submitted_passed	yes	Passed hidden tests	pass	3	n/a	73.9s
459f9de3	cli/claude-sonnet	2	submitted_passed	yes	Passed hidden tests	pass	3	n/a	64.7s
3a7da64d	cli/claude-sonnet	3	submitted_passed	yes	Passed hidden tests	pass	3	n/a	92.5s
1e7f41d8	cli/claude-sonnet	4	submitted_passed	yes	Passed hidden tests	pass	3	n/a	83.8s
9902eaa3	cli/claude-sonnet	5	submitted_passed	yes	Passed hidden tests	pass	3	n/a	77.7s
12f95490	cli/claude-sonnet	6	submitted_passed	yes	Passed hidden tests	pass	3	n/a	84.0s
affa9ea3	cli/claude-sonnet	7	submitted_passed	yes	Passed hidden tests	pass	3	n/a	76.2s
ab27ff4e	cli/claude-sonnet	8	submitted_passed	yes	Passed hidden tests	pass	3	n/a	85.5s
cd337d97	cli/claude-sonnet	9	submitted_passed	yes	Passed hidden tests	pass	3	n/a	65.9s
879eca05	cli/codex	0	submitted_passed	yes	Passed hidden tests	pass	3	n/a	161.2s
fd449898	cli/codex	1	submitted_passed	yes	Passed hidden tests	pass	3	n/a	162.3s
5d078aad	cli/codex	2	submitted_passed	yes	Passed hidden tests	pass	3	n/a	223.6s
4ac68744	cli/codex	3	submitted_passed	yes	Passed hidden tests	pass	3	n/a	144.3s
7d1b1531	cli/codex	4	submitted_passed	yes	Passed hidden tests	pass	3	n/a	114.1s
662f4a3d	cli/codex	5	submitted_passed	yes	Passed hidden tests	pass	3	n/a	135.7s
660efecd	cli/codex	6	submitted_passed	yes	Passed hidden tests	pass	3	n/a	127.3s
51c40893	cli/codex	7	submitted_passed	yes	Passed hidden tests	pass	3	n/a	126.5s
899ad596	cli/codex	8	submitted_passed	yes	Passed hidden tests	pass	3	n/a	144.5s
e90dd0c8	cli/codex	9	submitted_passed	yes	Passed hidden tests	pass	3	n/a	134.5s
0711410a	codestral-latest	0	context_exhaustion	no	Graded	fail	1	80 881	75.3s
de0f078f	codestral-latest	1	context_exhaustion	no	Graded	fail	9	203 624	151.0s
5c42bd02	codestral-latest	2	submitted_failed	yes	Graded	fail	0	4 126	7.4s
3498fcf5	codestral-latest	3	context_exhaustion	no	Graded	fail	9	175 540	100.7s

ROLLOUT	MODEL	SEED	OUTCOME	SUBMITTED	DEEPEST STAGE	HIDDEN TESTS	FILES	TOKENS	ELAPSED
5e744f53	codestral-latest	4	edited_tested_no_submi t	no	Graded	fail	11	102 805	81.2s
b761b8be	codestral-latest	5	edited_tested_no_submi t	no	Graded	fail	10	89 896	94.4s
eddbc8c8	codestral-latest	6	edited_tested_no_submi t	no	Graded	fail	1	85 524	81.4s
b0c7e737	codestral-latest	7	edited_tested_no_submi t	no	Graded	fail	1	62 082	65.4s
3061565d	codestral-latest	8	context_exhaustion	no	Graded	fail	8	251 723	513.2s
8b5af559	codestral-latest	9	context_exhaustion	no	Graded	fail	9	91 689	99.7s

Failure taxonomy

CLASS	COUNT	WHAT IT MEANS
context_exhaustion	5	Ran out of output context before it could submit.
edited_tested_no_submit	4	Edited files and ran the visible tests but never submitted.
submitted_failed	1	Submitted a change; the independent hidden tests still failed.

Appendix B — evidence and artifact digests

RUN BUNDLE DIGEST

a6f12b647a54dbf0c0ecf80a
cda77ee562a8ccdb3ad1f44b
ccffa20c0716d664

RUN ID

live-
20260902T134741Z

STARTED

2026-09-
02T13:47:41.546440
+00:00

FINISHED

2026-09-
02T14:49:46.020126
+00:00

Artifact digests

ARTIFACT	SHA-256
publicSummary	92ec8ad502cca03db70ca82d7c593206728ddfa95c2e483a198203d73b81ccb6
rollouts/0711410a-20af-4809-97d9-eb6f1cb4ed52.json	e76288222b67f8622db779c912d61b460fba164efaf041126bd6dcb14a87cdc8
rollouts/12f95490-752e-4fae-8352-8fc3f600c46c.json	05fad9811badd35c5ba8e2c8b3884d87ffe09385cc1d870be5d8725bc9af84a2
rollouts/1e7f41d8-724c-4906-94aa-77195464bafe.json	01744a44277bacc58c08bcae2b1b80959ce3bf8c47588a7c787fd83453b30354
rollouts/230a8274-69d8-4aa7-a026-3972e84dee78.json	66a3d04b57a30ded39eafba676ee1fa846909384b59ad96016e4dfebb3386cff
rollouts/3061565d-c1ba-453b-b793-1e061a2e4a62.json	1f9200514a2b2c38b9ee97118974f2d5629406e67349552f13ad03b710fc49fd
rollouts/3498fcf5-bbb6-414b-aba0-a13bad481435.json	c9e901f199c5a36bda324d35fea1f3ffe95e3dc31907a59e48c79d53deaad30e
rollouts/3a7da64d-1728-43ad-b959-185afd23de73.json	4ee16b3a4a8d0092617a1dfee773295ce2a1acb590e9b58abeeda7aacb6cbbb6
rollouts/3c178cb0-987f-499b-b4c5-8e6a51d22739.json	2b77fd47c36b9ced52164b48be6988332475b6467ab185166d9bf6a9bad8b861
rollouts/459f9de3-dd18-4b28-9e46-eefcda169291.json	81ee5801eda8feea8903bdc3e959cfdcf4e495fe3a8039620ed058080a6009d4
rollouts/4ac68744-04cf-46df-9fde-24f77013e84c.json	64555ef1cfd29429efb7ffa3c4f0f611744ea7d31e27a9dc6087941312b3eb7e
rollouts/51c40893-6126-4096-88a2-2ed8a13df341.json	39fa04de1b8f70e1f90e4d1713206db95e2d0450c682658e3369b780e5b8a1b2
rollouts/5c42bd02-e324-467a-a4a8-d8264ec66778.json	17860fd2c9c0127b8a37dbdee8d3105b4add25c756e2b45ab4d38790e1520412
rollouts/5d078aad-9bc3-42b7-97b8-a615b5f9a87f.json	8463d90244f36a14c9095d900f885147fab82e5c6bf9e8406ea984e06ad1208a
rollouts/5e744f53-ec91-4f5b-85e2-d6f150221616.json	3e6c5a77724a182243f898fe82b806aed906eb9a4b156276f9677d471c01f409
rollouts/660efecd-4c0a-4092-9713-fd3677d3528f.json	7435d69b8bc0ea51765cde0eff6a2e1926c36059307be111e3c82807d9bf2646
rollouts/662f4a3d-ea26-4e49-b9ee-24a190ed2db3.json	6d0227476d2a60f1f595800852d3c2b187d38591cb7b7fdcec39071104c6ef3c
rollouts/7d1b1531-f014-45b5-93f8-c578d036d820.json	755def32bc1166a4edba484cef698b57f399394937a13bf5b15c8f1ca9b93074
rollouts/879eca05-6ed5-481f-a4c0-58a675435c85.json	fcff7d31005d2bc12895065eb0356058fd81fe6573074fbc73a54ce29ad96118
rollouts/899ad596-b098-4a27-b1a8-ddc53e4346f2.json	835c88ee481264fa1b0b844eb4ed17d76c60185ddf56214e73629beca582145e
rollouts/8b5af559-b0f7-4aaa-9748-435f7836143e.json	42b8c77bec005e209190119a559320f4908762c2b85460f19e97d0cbf4ebe051
rollouts/9902eaa3-4af4-4039-946a-2b4fbbb3eb0c.json	8f91bdd28114bb10fa0f57b3c9030dda7a96fc3c26b3f47a9fffb19f7d161cc3e
rollouts/ab27ff4e-329f-4ad0-b20e-cd10c9e7e4b6.json	f8c53bac344a5ae08a4f020dbe6c65b3f8d0f66275cad641651c26ff098fa84b
rollouts/affa9ea3-9250-418d-9507-8cfbf429b84c.json	47d6031bef983bfd8f5878d13bbf4ef24c192a42f7a79186490048b9745c788a
rollouts/b0c7e737-a914-495b-ba72-aee2b3fd18b1.json	4048e436ea6ce3169b884cbfc4620936645573daf579ad52a5429b04032c0c19
rollouts/b761b8be-4d5d-4934-8e0e-6d63e67c573d.json	b3d3e9ee57a3f991b1d0c127d595178e1c9367b8ca17879d948d08f9fa584477
rollouts/cd337d97-6960-429c-99dd-10a17b935b0d.json	0051949472962b064fd9b7934f6acbb607de390a516e2fc2b5f59d7d03efe4c9
rollouts/de0f078f-5b50-4853-873f-208337f892f9.json	0fdc5de0bbf5032da58e7fa47df2f5000c926f566ec1c3ff8fb450f1aea2ca91
rollouts/e90dd0c8-6384-4a8e-857d-560d91cc8a8d.json	6fa264a19e220ee863979f0f2ac37f82b75e269d3ba5d54b0a7dd2957786f56e

ARTIFACT	SHA-256
rollouts/eddbc8c8-36ca-4533-b298-2cd72c5d5c62.json	7c4702032ea6d2d74a3fcf758665e109cb780012c387e4c7c6cd53b5cf1cb72f
rollouts/fd449898-96bc-4c78-bd82-2d37535c2987.json	1030a0bd7437271001c13f01cd8f725b116c4c08d9e13ddc4d7132bc925865cc
spec	c5580395b0da0e89877b88e47c897742665c2711455762bc87c742b9f68019ec
summary	adcd3c1bdb43057c1e2f1114afe7303277a5f95db0d1786f64407524d6eead72

Reproducing this run

Every rollout can be re-derived: the exam is frozen (fingerprint above), the grader ran in the recorded container image, and each rollout's full trace — every tool call, every result, every word — is in the evidence bundle at its rollouts/<id>.json path. A re-run uses the same job shape: the exam id, your endpoint, n, and the runtime. The record id names this run forever.

Appendix C — redacted transcripts

For each rollout: the tool calls it made, in order — which tool, on which of the model's own paths, whether the call succeeded, and how many characters came back. What a tool returned is not here, and neither is the submission: on this exam the submitted content is the graded answer, so no diff, no answer value and no model claim is printed. A request for the reference solution or the hidden tests is recorded as a request; the path it asked for is replaced by a label. See Appendix D, "What this report withholds and why".

Disclosure class applied to this exam: answer-key disclosure (ownership vvdex_proprietary — engine registry; family harness_eval): the exam is VVDex evaluation IP and the graded output is an answer or a keyed value, so no submission content, diff, model claim or hidden-assertion name is published.

3c178cb0 · cli:cli/claude-sonnet · seed 0 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	src/rates.py	ok	1 335
1	edit_file	src/rates.py	ok	50
2	run_tests	—	failed	95
3	run_tests	—	ok	1 581
4	read_file	src/rates.py	ok	157
5	write_file	src/rates.py	ok	62
6	run_tests	—	ok	1 581
7	read_file	src/rates.py	ok	157
8	write_file	src/rates.py	ok	61
9	run_tests	—	ok	268
10	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (210 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

230a8274 · cli:cli/claude-sonnet · seed 1 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	src/rates.py	ok	1 335
1	edit_file	src/rates.py	ok	50
2	run_tests	—	failed	95
3	run_tests	—	ok	1 581
4	read_file	src/rates.py	ok	157
5	write_file	src/rates.py	ok	62
6	run_tests	—	ok	1 581
7	read_file	src/rates.py	ok	157
8	write_file	src/rates.py	ok	61
9	run_tests	—	ok	268

STEP	TOOL	PATH	RESULT	CHARS
10	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (212 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

459f9de3 · cli:cli/claude-sonnet · seed 2 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	src/rates.py	ok	1 335
1	edit_file	src/rates.py	ok	50
2	run_tests	—	failed	95
3	run_tests	—	ok	1 581
4	read_file	src/rates.py	ok	157
5	write_file	src/rates.py	ok	62
6	run_tests	—	ok	1 581
7	read_file	src/rates.py	ok	157
8	write_file	src/rates.py	ok	61
9	run_tests	—	ok	268
10	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (210 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

3a7da64d · cli:cli/claude-sonnet · seed 3 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	src/rates.py	ok	1 335
1	edit_file	src/rates.py	ok	50
2	run_tests	—	failed	95
3	run_tests	—	ok	1 581
4	read_file	src/rates.py	ok	157
5	write_file	src/rates.py	ok	62
6	run_tests	—	ok	1 581
7	write_file	src/rates.py	ok	61
8	run_tests	—	ok	268
9	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (199 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

1e7f41d8 · cli:cli/claude-sonnet · seed 4 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	src/rates.py	ok	1 335
1	edit_file	src/rates.py	ok	50
2	run_tests	—	failed	95
3	run_tests	—	ok	1 581
4	read_file	src/rates.py	ok	157
5	write_file	src/rates.py	ok	62
6	run_tests	—	ok	1 581
7	read_file	src/rates.py	ok	157
8	write_file	src/rates.py	ok	61
9	run_tests	—	ok	268
10	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (210 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

9902eaa3 · cli:cli/claude-sonnet · seed 5 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	src/rates.py	ok	1 335
1	edit_file	src/rates.py	ok	50
2	run_tests	—	failed	95
3	run_tests	—	ok	1 581
4	read_file	src/rates.py	ok	157
5	write_file	src/rates.py	ok	62
6	run_tests	—	ok	1 581
7	read_file	src/rates.py	ok	157
8	write_file	src/rates.py	ok	61
9	run_tests	—	ok	268
10	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (210 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

12f95490 · cli:cli/claude-sonnet · seed 6 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	src/rates.py	ok	1 335
1	edit_file	src/rates.py	ok	50

STEP	TOOL	PATH	RESULT	CHARS
2	run_tests	—	failed	95
3	run_tests	—	ok	1 581
4	read_file	src/rates.py	ok	157
5	write_file	src/rates.py	ok	62
6	run_tests	—	ok	1 581
7	read_file	src/rates.py	ok	157
8	write_file	src/rates.py	ok	62
9	run_tests	—	ok	268
10	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (214 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

affa9ea3 · cli:cli/claude-sonnet · seed 7 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	src/rates.py	ok	1 335
1	edit_file	src/rates.py	ok	50
2	run_tests	—	failed	95
3	run_tests	—	ok	1 581
4	read_file	src/rates.py	ok	157
5	write_file	src/rates.py	ok	62
6	run_tests	—	ok	1 581
7	write_file	src/rates.py	ok	61
8	run_tests	—	ok	268
9	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (199 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

ab27ff4e · cli:cli/claude-sonnet · seed 8 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	src/rates.py	ok	1 335
1	edit_file	src/rates.py	ok	50
2	run_tests	—	failed	95
3	run_tests	—	ok	1 581
4	read_file	src/rates.py	ok	157
5	write_file	src/rates.py	ok	62
6	run_tests	—	ok	1 581

STEP	TOOL	PATH	RESULT	CHARS
7	read_file	src/rates.py	ok	157
8	write_file	src/rates.py	ok	61
9	run_tests	—	ok	268
10	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (212 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

cd337d97 · cli:cli/claude-sonnet · seed 9 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	src/rates.py	ok	1 335
1	edit_file	src/rates.py	ok	50
2	run_tests	—	failed	95
3	run_tests	—	ok	1 581
4	read_file	src/rates.py	ok	157
5	write_file	src/rates.py	ok	62
6	run_tests	—	ok	1 581
7	write_file	src/rates.py	ok	61
8	run_tests	—	ok	268
9	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (199 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

879eca05 · cli:cli/codex · seed 0 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 633
1	read_file	—	failed	35
2	read_file	src/rates.py	ok	1 335
3	edit_file	src/rates.py	ok	50
4	read_file	src/rates.py	ok	157
5	read_file	src/rates.py	ok	0
6	run_tests	—	failed	95
7	run_tests	—	ok	1 581
8	write_file	src/rates.py	ok	62
9	read_file	src/rates.py	ok	157
10	read_file	src/rates.py	ok	40
11	read_file	src/rates.py	ok	0

STEP	TOOL	PATH	RESULT	CHARS
12	write_file	src/rates.py	ok	62
13	read_file	src/rates.py	ok	148
14	run_tests	—	ok	268
15	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (270 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

fd449898 · cli:cli/codex · seed 1 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	list_files	.	ok	94
1	read_file	TASK.md	ok	1 633
2	read_file	src/rates.py	ok	1 335
3	edit_file	src/rates.py	ok	50
4	read_file	src/rates.py	ok	157
5	read_file	src/rates.py	ok	0
6	read_file	src/rates.py	ok	0
7	run_tests	—	failed	95
8	run_tests	—	ok	1 581
9	write_file	src/rates.py	ok	62
10	read_file	src/rates.py	ok	157
11	read_file	src/rates.py	ok	20
12	read_file	src/rates.py	ok	0
13	run_tests	—	ok	1 581
14	write_file	src/rates.py	ok	61
15	read_file	src/rates.py	ok	892
16	run_tests	—	ok	268
17	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (288 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

5d078aad · cli:cli/codex · seed 2 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	list_files	.	ok	94
1	read_file	TASK.md	ok	1 633
2	read_file	—	failed	35
3	read_file	src/rates.py	ok	1 335

STEP	TOOL	PATH	RESULT	CHARS
4	read_file	tests/test_rates.py	ok	719
5	edit_file	src/rates.py	ok	50
6	read_file	src/rates.py	ok	157
7	read_file	src/rates.py	ok	0
8	read_file	src/rates.py	ok	39
9	read_file	src/rates.py	ok	60
10	read_file	src/rates.py	ok	0
11	run_tests	—	failed	95
12	run_tests	—	ok	1 581
13	write_file	src/rates.py	ok	62
14	read_file	src/rates.py	ok	157
15	run_tests	—	ok	1 581
16	write_file	src/rates.py	ok	61
17	read_file	src/rates.py	ok	787
18	run_tests	—	ok	268
19	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (327 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

4ac68744 · cli:cli/codex · seed 3 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 633
1	read_file	—	failed	35
2	read_file	src/rates.py	ok	1 335
3	edit_file	src/rates.py	ok	50
4	read_file	src/rates.py	ok	157
5	read_file	src/rates.py	ok	0
6	run_tests	—	failed	95
7	read_file	src/rates.py	ok	0
8	run_tests	—	ok	1 581
9	write_file	src/rates.py	ok	62
10	read_file	src/rates.py	ok	157
11	run_tests	—	ok	1 581
12	write_file	src/rates.py	ok	61
13	read_file	src/rates.py	ok	791
14	run_tests	—	ok	268

STEP	TOOL	PATH	RESULT	CHARS
15	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (284 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

7d1b1531 · cli:cli/codex · seed 4 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 633
1	read_file	—	failed	35
2	read_file	src/rates.py	ok	1 335
3	edit_file	src/rates.py	ok	50
4	run_tests	—	failed	95
5	run_tests	—	ok	1 581
6	read_file	src/rates.py	ok	157
7	write_file	src/rates.py	ok	62
8	read_file	src/rates.py	ok	157
9	write_file	src/rates.py	ok	62
10	read_file	src/rates.py	ok	1 278
11	run_tests	—	ok	268
12	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (213 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

662f4a3d · cli:cli/codex · seed 5 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 633
1	read_file	—	failed	35
2	read_file	src/rates.py	ok	1 335
3	read_file	tests/test_rates.py	ok	719
4	edit_file	src/rates.py	ok	50
5	read_file	src/rates.py	ok	157
6	read_file	src/rates.py	ok	0
7	run_tests	—	failed	95
8	run_tests	—	ok	1 581
9	write_file	src/rates.py	ok	62
10	read_file	src/rates.py	ok	157
11	read_file	src/rates.py	ok	20

STEP	TOOL	PATH	RESULT	CHARS
12	write_file	src/rates.py	ok	62
13	read_file	src/rates.py	ok	1 198
14	run_tests	—	ok	268
15	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (246 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

660efecd · cli:cli/codex · seed 6 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 633
1	list_files	.	ok	94
2	read_file	src/rates.py	ok	1 335
3	read_file	tests/test_rates.py	ok	719
4	edit_file	src/rates.py	ok	50
5	read_file	src/rates.py	ok	157
6	run_tests	—	failed	95
7	run_tests	—	ok	1 581
8	write_file	src/rates.py	ok	62
9	read_file	src/rates.py	ok	157
10	read_file	src/rates.py	ok	0
11	read_file	src/rates.py	ok	0
12	write_file	src/rates.py	ok	62
13	run_tests	—	ok	268
14	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (237 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

51c40893 · cli:cli/codex · seed 7 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	list_files	.	ok	94
1	read_file	TASK.md	ok	1 633
2	read_file	src/rates.py	ok	1 335
3	read_file	tests/test_rates.py	ok	719
4	edit_file	src/rates.py	ok	50
5	run_tests	—	failed	95
6	run_tests	—	ok	1 581

STEP	TOOL	PATH	RESULT	CHARS
7	read_file	src/rates.py	ok	157
8	write_file	src/rates.py	ok	62
9	read_file	src/rates.py	ok	157
10	read_file	src/rates.py	ok	20
11	write_file	src/rates.py	ok	61
12	read_file	src/rates.py	ok	892
13	run_tests	—	ok	268
14	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (255 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

899ad596 · cli:cli/codex · seed 8 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 633
1	read_file	—	failed	35
2	read_file	src/rates.py	ok	1 335
3	edit_file	src/rates.py	ok	50
4	read_file	src/rates.py	ok	157
5	read_file	src/rates.py	ok	0
6	read_file	src/rates.py	ok	60
7	run_tests	—	failed	95
8	run_tests	—	ok	1 581
9	write_file	src/rates.py	ok	62
10	read_file	src/rates.py	ok	157
11	read_file	src/rates.py	ok	20
12	run_tests	—	ok	1 581
13	write_file	src/rates.py	ok	61
14	read_file	src/rates.py	ok	892
15	run_tests	—	ok	268
16	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (278 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

e90dd0c8 · cli:cli/codex · seed 9 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 633

STEP	TOOL	PATH	RESULT	CHARS
1	list_files	.	ok	94
2	read_file	src/rates.py	ok	1 335
3	read_file	tests/test_rates.py	ok	719
4	edit_file	src/rates.py	ok	50
5	read_file	src/rates.py	ok	157
6	read_file	src/rates.py	ok	20
7	read_file	src/rates.py	ok	0
8	write_file	src/rates.py	ok	62
9	read_file	src/rates.py	ok	157
10	run_tests	—	failed	95
11	run_tests	—	ok	1 581
12	write_file	src/rates.py	ok	62
13	read_file	src/rates.py	ok	1 168
14	run_tests	—	ok	268
15	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (267 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

0711410a · mistral:codestral-latest · seed 0 · context_exhaustion

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	—	failed	35
1	invalid_tool_call	"}	failed	103
2	list_files	.	ok	94
3	read_file	TASK.md	ok	1 633
4	read_file	src/rates.py	ok	1 335
5	run_tests	—	failed	95
6	edit_file	src/rates.py	ok	50
7	run_tests	—	ok	1 581
8	edit_file	src/rates.py	failed	29
9	edit_file	src/rates.py	failed	29
10	edit_file	src/rates.py	failed	29
11	edit_file	src/rates.py	failed	29
12	edit_file	src/rates.py	failed	29
13	edit_file	src/rates.py	failed	29
14	edit_file	src/rates.py	failed	29
15	edit_file	src/rates.py	failed	29

STEP	TOOL	PATH	RESULT	CHARS
16	edit_file	src/rates.py	failed	29
17	edit_file	src/rates.py	failed	29
18	edit_file	src/rates.py	failed	29
19	edit_file	src/rates.py	failed	29
20	edit_file	src/rates.py	failed	29
21	edit_file	src/rates.py	failed	29
22	edit_file	src/rates.py	failed	29
23	edit_file	src/rates.py	failed	29

Submission diff withheld by the public-disclosure policy (350 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

de0f078f · mistral:codestral-latest · seed 1 · context_exhaustion

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	—	failed	35
1	invalid_tool_call	"}	failed	103
2	list_files	.	ok	94
3	read_file	src/rates.py	ok	1 335
4	edit_file	src/rates.py	ok	50
5	run_tests	—	failed	95
6	run_tests	—	ok	1 581
7	edit_file	src/rates.py	failed	29
8	edit_file	src/rates.py	failed	29
9	write_file	src/rates.py	ok	62
10	run_tests	—	ok	1 581
11	write_file	src/rates.py	ok	62
12	run_tests	—	ok	1 579
13	edit_file	src/rates.py	ok	50
14	run_tests	—	ok	1 579
15	edit_file	src/rates.py	ok	50
16	run_tests	—	ok	1 579
17	edit_file	src/rates.py	ok	50
18	run_tests	—	ok	1 579
19	edit_file	src/rates.py	ok	50
20	run_tests	—	ok	1 579
21	edit_file	src/rates.py	ok	50
22	run_tests	—	ok	1 579

STEP	TOOL	PATH	RESULT	CHARS
23	edit_file	src/rates.py	ok	50

Submission diff withheld by the public-disclosure policy (341 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

5c42bd02 · mistral:codestral-latest · seed 2 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	src/rates.py	ok	1 335
1	run_tests	—	failed	95
2	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (57 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

3498fcf5 · mistral:codestral-latest · seed 3 · context_exhaustion

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	—	failed	35
1	invalid_tool_call	"}	failed	103
2	list_files	.	ok	94
3	read_file	TASK.md	ok	1 633
4	read_file	src/rates.py	ok	1 335
5	edit_file	src/rates.py	ok	50
6	run_tests	—	failed	95
7	run_tests	—	ok	1 581
8	edit_file	src/rates.py	ok	50
9	run_tests	—	ok	1 561
10	edit_file	src/rates.py	ok	50
11	run_tests	—	ok	1 561
12	edit_file	src/rates.py	ok	50
13	run_tests	—	ok	1 570
14	edit_file	src/rates.py	ok	50
15	run_tests	—	ok	1 543
16	edit_file	src/rates.py	ok	50
17	run_tests	—	ok	1 543
18	edit_file	src/rates.py	ok	50
19	run_tests	—	ok	1 543
20	edit_file	src/rates.py	ok	50
21	run_tests	—	ok	1 543

STEP	TOOL	PATH	RESULT	CHARS
22	edit_file	src/rates.py	ok	50
23	run_tests	—	ok	1 543

Submission diff withheld by the public-disclosure policy (337 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

5e744f53 · mistral:codestral-latest · seed 4 · edited_tested_no_submit

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	—	failed	35
1	read_file	src/rates.py	ok	1 335
2	run_tests	—	failed	95
3	edit_file	src/rates.py	ok	50
4	run_tests	—	ok	1 581
5	edit_file	src/rates.py	ok	50
6	run_tests	—	ok	1 561
7	edit_file	src/rates.py	ok	50
8	run_tests	—	ok	2 005
9	edit_file	src/rates.py	ok	50
10	run_tests	—	ok	1 999
11	edit_file	src/rates.py	ok	50
12	run_tests	—	ok	1 999
13	edit_file	src/rates.py	ok	50
14	run_tests	—	ok	1 999
15	edit_file	src/rates.py	ok	50
16	run_tests	—	ok	1 999
17	edit_file	src/rates.py	ok	50
18	run_tests	—	ok	1 999
19	edit_file	src/rates.py	ok	50
20	run_tests	—	ok	1 999
21	edit_file	src/rates.py	ok	50
22	run_tests	—	ok	1 999
23	edit_file	src/rates.py	ok	50

Submission diff withheld by the public-disclosure policy (339 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

b761b8be · mistral:codestral-latest · seed 5 · edited_tested_no_submit

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	—	failed	35
1	invalid_tool_call	"}	failed	103
2	list_files	.	ok	94
3	read_file	src/rates.py	ok	1 335
4	edit_file	src/rates.py	ok	50
5	run_tests	—	failed	95
6	run_tests	—	ok	1 581
7	edit_file	src/rates.py	ok	50
8	run_tests	—	ok	1 561
9	edit_file	src/rates.py	ok	50
10	run_tests	—	ok	1 539
11	edit_file	src/rates.py	ok	50
12	run_tests	—	ok	1 540
13	edit_file	src/rates.py	ok	50
14	run_tests	—	ok	1 540
15	edit_file	src/rates.py	ok	50
16	run_tests	—	ok	1 540
17	edit_file	src/rates.py	ok	50
18	run_tests	—	ok	1 540
19	edit_file	src/rates.py	ok	50
20	run_tests	—	ok	1 540
21	edit_file	src/rates.py	ok	50
22	run_tests	—	ok	1 540
23	edit_file	src/rates.py	ok	50

Submission diff withheld by the public-disclosure policy (338 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

eddbc8c8 · mistral:codestral-latest · seed 6 · edited_tested_no_submit

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	—	failed	35
1	read_file	TASK.md	ok	1 633
2	read_file	src/rates.py	ok	1 335
3	edit_file	src/rates.py	ok	50
4	run_tests	—	failed	95
5	run_tests	—	ok	1 581

STEP	TOOL	PATH	RESULT	CHARS
6	edit_file	src/rates.py	failed	29
7	edit_file	src/rates.py	failed	29
8	edit_file	src/rates.py	failed	29
9	edit_file	src/rates.py	failed	29
10	edit_file	src/rates.py	failed	29
11	edit_file	src/rates.py	failed	29
12	edit_file	src/rates.py	failed	29
13	edit_file	src/rates.py	failed	29
14	edit_file	src/rates.py	failed	29
15	edit_file	src/rates.py	failed	29
16	edit_file	src/rates.py	failed	29
17	edit_file	src/rates.py	failed	29
18	edit_file	src/rates.py	failed	29
19	edit_file	src/rates.py	failed	29
20	edit_file	src/rates.py	failed	29
21	edit_file	src/rates.py	failed	29
22	edit_file	src/rates.py	failed	29
23	edit_file	src/rates.py	failed	29

Submission diff withheld by the public-disclosure policy (355 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

b0c7e737 · mistral:codestral-latest · seed 7 · edited_tested_no_submit

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	—	failed	35
1	invalid_tool_call	"}	failed	103
2	list_files	.	ok	94
3	read_file	src/rates.py	ok	1 335
4	edit_file	src/rates.py	ok	50
5	run_tests	—	failed	95
6	run_tests	—	ok	1 581
7	edit_file	src/rates.py	failed	29
8	edit_file	src/rates.py	failed	29
9	edit_file	src/rates.py	failed	29
10	edit_file	src/rates.py	failed	29
11	edit_file	src/rates.py	failed	29
12	edit_file	src/rates.py	failed	29

STEP	TOOL	PATH	RESULT	CHARS
13	edit_file	src/rates.py	failed	29
14	edit_file	src/rates.py	failed	29
15	edit_file	src/rates.py	failed	29
16	edit_file	src/rates.py	failed	29
17	edit_file	src/rates.py	failed	29
18	edit_file	src/rates.py	failed	29
19	edit_file	src/rates.py	failed	29
20	edit_file	src/rates.py	failed	29
21	edit_file	src/rates.py	failed	29
22	edit_file	src/rates.py	failed	29
23	edit_file	src/rates.py	failed	29

Submission diff withheld by the public-disclosure policy (350 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

3061565d · mistral:codestral-latest · seed 8 · context_exhaustion

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	src/rates.py	ok	1 335
1	edit_file	src/rates.py	ok	50
2	run_tests	—	failed	95
3	run_tests	—	ok	1 581
4	read_file	src/rates.py	ok	157
5	edit_file	src/rates.py	ok	50
6	run_tests	—	ok	1 001
7	read_file	src/rates.py	ok	318
8	edit_file	src/rates.py	ok	50
9	run_tests	—	ok	1 001
10	read_file	src/rates.py	ok	318
11	edit_file	src/rates.py	ok	50
12	run_tests	—	ok	1 001
13	read_file	src/rates.py	ok	318
14	edit_file	src/rates.py	ok	50
15	run_tests	—	ok	1 001
16	read_file	src/rates.py	ok	318
17	edit_file	src/rates.py	ok	50
18	run_tests	—	ok	1 001
19	read_file	src/rates.py	ok	318

STEP	TOOL	PATH	RESULT	CHARS
20	edit_file	src/rates.py	ok	50
21	run_tests	—	ok	1 001
22	read_file	src/rates.py	ok	318
23	edit_file	src/rates.py	ok	50

Submission diff withheld by the public-disclosure policy (343 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

8b5af559 · mistral:codestral-latest · seed 9 · context_exhaustion

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	—	failed	35
1	invalid_tool_call	"}	failed	103
2	list_files	.	ok	94
3	read_file	TASK.md	ok	1 633
4	read_file	src/rates.py	ok	1 335
5	run_tests	—	failed	95
6	run_tests	—	ok	1 583
7	edit_file	src/rates.py	ok	50
8	run_tests	—	ok	1 581
9	edit_file	src/rates.py	ok	50
10	run_tests	—	ok	1 581
11	edit_file	src/rates.py	ok	50
12	run_tests	—	ok	1 561
13	edit_file	src/rates.py	ok	50
14	run_tests	—	ok	1 539
15	edit_file	src/rates.py	ok	50
16	run_tests	—	ok	1 562
17	edit_file	src/rates.py	ok	50
18	run_tests	—	ok	1 540
19	edit_file	src/rates.py	ok	50
20	run_tests	—	ok	1 562
21	edit_file	src/rates.py	ok	50
22	run_tests	—	ok	1 541
23	edit_file	src/rates.py	ok	50

Submission diff withheld by the public-disclosure policy (343 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

11 Evidence

Every number above is recomputable from the sealed records below; each record verifies itself with `vvdex-env records verify`, and this document was regenerated from them with `vvdex-env records campaign`.

- `vvdex.browser.order-desk-1` → `fr-20260902-6e618c05` · `recordDigest` `59f94854d78f7ce174866ff114505b2b41ca2333f287f40b604fbb38cc267da5`
- `vvdex.harness.fault-recovery-1` → `fr-20260902-a6f12b64` · `recordDigest` `d5a8264aa0c20dd67c14879856a744e2ff1ab79da35ea8bcea92b4ba1ea14d5f`

12 Attestation

Certified results appear only where evidence proves them. Withheld, invalid and lane-error cells are never presented as model results, and no historical evidence was modified in producing this document. Counts are exact; below ten rollouts per cell no rate or ranking beyond the counts is asserted.

© VVDex. All rights reserved. Public materials are provided for viewing, evaluation and verification. Reuse, redistribution, derivative benchmark creation, training use, republication or commercial use requires written permission unless a specific file states otherwise. Full terms: vvdexops.com/terms/. Third-party open source reproduced in an exam keeps its own licence.