

Harness recovery under a faulted tool loop

EVALUATION REPORT

2026-09-02

Evaluation scope:

Claude Code CLI (Sonnet, subscription)

Codex CLI (GPT, subscription)

codestral-latest

Prepared by:

VVDex Forge — generated directly from sealed run evidence

fr-20260902-a6f12b64 · vvdex.harness.fault-recovery-1 v0.1.0 · Independent model evaluation report

About this report

About VVDex Forge

VVDex Forge builds and certifies evaluation environments for AI agents. Every exam is sealed before any model sees it: the reference solution must pass, an empty submission must fail, the grader must reject near-misses, and the hidden tests never enter the agent's tree. Reports and the public catalog live at **vvdexops.com**.

Report integrity

This report is generated from run evidence; its digests verify authenticity. The sole test of a copy is the three-line procedure in Verification — a copy whose digests do not reproduce is not this report.

Scope

One exam · 30 rollout(s) · 24 tool steps and 5880s wall clock per rollout, stated in full in Evaluation configuration.

Boundary

Not a leaderboard. Not a general capability score. A post-seal benchmark: this run stamped no certification gate and rewrote nothing — not the exam, not the reference solution, not the grader, not the certified evidence.

Contents

01.	Executive summary	4
02.	Exam definition	5
03.	Certification evidence	6
04.	Evaluation configuration	7
05.	Model outcomes	8
06.	Stage funnel	9
07.	Behavioural analysis	10
08.	Efficiency & telemetry	11
09.	Findings	12
10.	Recommendations	14
11.	Verification	15

Appendices: A — every rollout · B — evidence digests · C — redacted transcripts · D — metric definitions · Attestation

01 Executive summary

MODELS TESTED 3	CERTIFIED PASSES 20 of 30 graded rollouts	SUBMITTED 21 of 30 graded rollouts that called submit	NOT GRADED — LANE ERRORS 0 of 30 attempts
---------------------------------------	--	--	--

Purpose

This report presents a controlled evaluation of 3 model lane(s) against the certified exam `vv dex.harness.fault-recovery-1` (v0.1.0, family `harness_eval`). Every lane received the identical frozen workspace, tool contract and limits; grading ran afterwards, in isolation, against hidden tests no lane ever saw. The purpose is to state, with reproducible evidence, what each lane did and where it stopped.

Outcome

20 of 30

graded rollouts passed the independent hidden tests and were certified

20 of 30 graded rollouts passed. `cli/claude-sonnet`, `cli/codex` completed the full loop (inspect, edit, run the visible tests, submit) and the sealed grader accepted the submission. The most common way to fail was `context_exhaustion` — ran out of output context before it could submit.

Key findings

F-01	<code>cli/claude-sonnet</code> , <code>cli/codex</code> passed the independent hidden tests
F-02	<code>context_exhaustion</code> × 5 (<code>codestral-latest</code>)
F-03	<code>edited_tested_no_submit</code> × 4 (<code>codestral-latest</code>)
F-04	<code>submitted_failed</code> × 1 (<code>codestral-latest</code>)
F-05	1 overclaim(s) — success stated, then refuted by the grader

Each finding is stated in full, with evidence, in the Findings section; recommendations for acting on them follow it. Elapsed wall clock for the whole engagement: 3724.5s.

02 Exam definition

In scope. The ability of each configured model lane to complete this one certified task end to end — inspect the workspace, produce a correct edit, verify it with the visible tests, and submit — within 24 tool steps and 5880s of wall clock, at 10 rollout(s) per lane. Behavioural measurements along the way (tool discipline, grounding, overclaim) are in scope because they are read from the recorded trace, not judged by another model.

Out of scope. General model quality, performance on other task families, prompt-tuning on the lane's behalf, and any comparison this run's sample size cannot support. A lane's API failure is reported as infrastructure, never as model capability.

The instrument. Exam `vvdex.harness.fault-recovery-1` was certified before this run: its reference solution passes, an empty submission fails, its grader distinguishes near-misses, and its sealed evidence is unchanged by this evaluation (see Certification evidence).

What was measured

EXAM	VERSION	FAMILY	ROLLOUTS
vvdex.harness.fault-recovery-1	0.1.0	harness_eval	30 (10 per model)

Where the defect came from

This exam has no upstream repository to credit: the defect, the reference solution and the hidden suite are VVDex's own.

03 Certification evidence

4 of 5 certification pillars are recorded in this package, loaded from the receipts sealed into the exam and matching this run's exam fingerprint. What is recorded is shown with its real value; what is not is named below and is not claimed.

FROZEN BASELINE

FAIL

expected — an empty submission over 2 run(s) must not pass

→

GOLD / REFERENCE

PASS

expected — the reference solution over 2 run(s) must pass

GRADER CONTROLS

12 / 12

The grader accepted the reference and rejected every near-miss.

ATTACK PROBES

10 / 10 blocked

0 trivial exploit(s) found.

RUNTIME

sha256:679b36225a1f83238
efba8c71b9cde3c24caaeacc9
b682ae92819cb55eec328f

network policy: deny

Certification metadata is not included in this report package for:

- sealed evidence bundle

Benchmark results remain valid as run evidence; the pillar(s) above are not substantiated by this document and this report does not claim them.

What the package does carry: the exam fingerprint below is the contract digest this run was executed against — a holder of the certification bundle for this fingerprint can join the two independently.

EXAM FINGERPRINT

4a206fbfe8fb70fb3e37943705065a8599fab492e495dca326279781ee960b1e

The contract digest this run was executed against.

04 Evaluation configuration

REPORT	fr-20260902-a6f12b64
ISSUED	2026-09-02
SUBJECT EXAM	vvdex.harness.fault-recovery-1 · v0.1.0
FAMILY	harness_eval
PREPARED BY	VVDex Forge engine vvdex-env 0.1.0
CLASSIFICATION	Independent model evaluation report — independently verifiable
CERTIFICATION	Exam certification: partially recorded (see Certification evidence)
STATUS	FINAL · report sealed against its own digest

Timeline

WHEN (UTC)	EVENT
2026-09-02 13:47:41	Evaluation run started
2026-09-02 14:49:46	Evaluation run finished
2026-09-02 14:49:46	Report generated from the sealed run evidence

Models under test

MODEL	PROVIDER	ENDPOINT	BILLING
Claude Code CLI (Sonnet, subscription)	cli	local runtime	operator subscription, flat — not billed per token
Codex CLI (GPT, subscription)	cli	local runtime	operator subscription, flat — not billed per token
codestral-latest	mistral	VVDex-configured	VVDex free-quota

A customer endpoint is recorded by origin only — never its port, its path, or its key.

Limits

RUNTIME docker	STEP LIMIT 24	WALL-CLOCK LIMIT 5880s	MAX OUTPUT TOKENS 2048
TEMPERATURE 0.2	RETRY POLICY none	COMMAND BUDGET 0	TOOLS OFFERED list_files, read_file, write_file, edit_file, run_tests, submit

The evaluated model never saw the hidden tests or the reference solution, and could not change its result. Grading ran in a separate step, after submission, on a container with no network.

05 Model outcomes

Denominators. Attempts requested: 30. Graded (evaluable) rollouts: 30. Not graded — lane errors: 0. A lane error is a rollout the API or the runtime ended before the hidden grader ran: it is listed, excluded from every rate and pass@k, and never silently dropped. Passes are certified passes: hidden grader exit 0 with integrity verified.

MODEL	ATTEMPTS	GRADED	PASSED	MODEL FAILS	LANE ERRORS	SUBMITTED	MEAN TIME	TYPICAL DEEPEST STAGE
Claude Code CLI (Sonnet, subscription) cli/claude-sonnet	10	10	10 / 10	0	0	10 / 10	77.9s	Passed hidden tests
Codex CLI (GPT, subscription) cli/codex	10	10	10 / 10	0	0	10 / 10	147.4s	Passed hidden tests
codestral-latest codestral-latest	10	10	0 / 10	10	0	1 / 10	127.0s	Graded

Every rollout, with its own deepest stage and grade, is in Appendix A.

Success rate, confidence and pass@k

Every rate on this page is a measurement of a finite number of rollouts, so every rate carries the interval that measurement actually supports. n is the evaluable count — attempts minus rollouts the lane ended before grading; those are listed under "not graded" and sit in no rate, interval or pass@k.

MODEL	ATTEMPTS	N	NOT GRADED	PASSES	RATE	95% INTERVAL	SEEDS	PASS@1	PASS@3	PASS@5	PASS@10
cli/claude-sonnet	10	10	0	10	100.0%	[72%, 100%]	0–9	100.0%	100.0%	100.0%	100.0%
cli/codex	10	10	0	10	100.0%	[72%, 100%]	0–9	100.0%	100.0%	100.0%	100.0%
codestral-latest	10	10	0	0	0.0%	[0%, 28%]	0–9	0.0%	0.0%	0.0%	0.0%

Interval and pass@k methods are stated in full in Appendix D. pass@k is shown for k = 1, 3, 5, 10; the sealed record carries every k up to n.

Model comparison

MODEL	PASS@1	95% INTERVAL	TOKENS	COST	MEAN ROLLOUT	OVERCLAIM RATE	TOP FAILURE CLASS
cli/claude-sonnet	100.0%	[72%, 100%]	n/a	n/a	77.9s	0.0%	none
cli/codex	100.0%	[72%, 100%]	n/a	n/a	147.4s	0.0%	none
codestral-latest	0.0%	[0%, 28%]	1 147 890	not billed	127.0s	10.0%	context_exhaustion

Tokens and cost read n/a where a lane reports no telemetry (a flat-rate CLI) — an absent measurement, not zero. pass@1 and its interval are over evaluable rollouts only.

This is not a leaderboard. Separated by this run (intervals do not overlap): cli/claude-sonnet vs codestral-latest; cli/codex vs codestral-latest. Not separated (intervals overlap): cli/claude-sonnet vs cli/codex. The primary comparison in this report is the stage funnel below — where a model stops is a finding; a percentage over this many rollouts is not.

06 Stage funnel

The funnel is the primary reading of a run. Where a model stops says more than a single pass percentage does, especially at low N.

STAGE	CLI:CLI/CLAUDE-SONNET	CLI:CLI/CODEX	MISTRAL:CODESTRAL-LATEST
Inspected	10 / 10	10 / 10	10 / 10
Edited	10 / 10	10 / 10	9 / 10
Ran tests	10 / 10	10 / 10	10 / 10
Submitted	10 / 10	10 / 10	1 / 10
Graded	10 / 10	10 / 10	10 / 10
Passed hidden tests	10 / 10	10 / 10	0 / 10

Deepest stage reached

MODEL	DEEPEST (ANY ATTEMPT)	TYPICAL (MOST ATTEMPTS)	HIDDEN-TEST PASSES	NOT GRADED
cli:cli/claude-sonnet	Passed hidden tests	Passed hidden tests	10 / 10	0
cli:cli/codex	Passed hidden tests	Passed hidden tests	10 / 10	0
mistral:codestral-latest	Graded	Graded	0 / 10	0

Stage definitions are in Appendix D. The matrix counts evaluable rollouts; lane errors are shown beside it, not inside it. The stages are not strictly nested: a model can submit without editing, and the matrix shows that rather than hiding it. Each rollout's own deepest stage is in Appendix A.

07 Behavioural analysis

Long-horizon behaviour

Counted off the recorded trace of each rollout. A dead end is a step that moved nothing: a re-read of a slice already returned, a repeat of the previous action, or a call to a tool this exam does not have.

MODEL	N	STEPS USED	RAN OUT OF STEPS	COMMANDS	CHANGED THE TREE	REFUSED (OVER BUDGET)	REVISITS	DEAD ENDS / ROLLOUT
cli:cli/claude-sonnet	10	10 / 24	0	0 / 0	0	0	17	2
cli:cli/codex	10	16 / 24	0	0 / 0	0	0	27	3
mistral:codestral-latest	10	21 / 24	9	0 / 0	0	0	7	6

Each column is defined in Appendix D.

Hallucination & overclaim

Three measurements, each with a stated definition. None of them is a judgement of the model's prose by another model — every number below comes from comparing what the model wrote against what was on disk, or against what the independent grader returned.

MODEL	N	GROUNDING FAILURES	KINDS	ROLLOUTS AFFECTED	OVERCLAIMS	OVERCLAIM RATE	RE-READS
cli:cli/claude-sonnet	10	0	none	0 / 10	0 / 10	0.0%	17
cli:cli/codex	10	0	none	0 / 10	0 / 10	0.0%	27
mistral:codestral-latest	10	4	nonexistent_path x4	4 / 10	1 / 10	10.0%	7

mistral:codestral-latest wrote 4 path reference(s) not in the tree it was given, across 4 of 10 rollout(s); claimed a fix the hidden tests then refuted in 1 of 10 rollout(s).

Every term above is defined in Appendix D.

08 Efficiency & telemetry

This run has no per-token price attached — a local model, a free quota, or a customer endpoint we are not billed for — so spend is reported in tokens. Tokens are what was actually measured; a dollar figure here would be invented. Lanes run through a flat-rate local CLI report no token telemetry; their cells say n/a — an absent measurement, not zero. Rollout counts here are evaluable rollouts.

MODEL	ROLLOUTS	TOKENS	TOKENS / ROLLOUT	COST	PASSES	RATE
cli:cli/claude-sonnet	10	n/a	n/a	n/a – telemetry unavailable	10 / 10	100.0%
cli:cli/codex	10	n/a	n/a	n/a – telemetry unavailable	10 / 10	100.0%
mistral:codestral-latest	10	1 147 890	114 789	not billed	0 / 10	0.0%

09 Findings

Every finding below is derived from the recorded data of this run — none is an opinion.
Severity: PASS a lane succeeded · ATTENTION a capability gap · LANE / INFRA
infrastructure, not capability · NOTE a property of the measurement.

F-01	PASS
cli/claude-sonnet, cli/codex passed the independent hidden tests	
Severity: PASS	Evidence: Model outcomes · Stage funnel · Appendix A
20 of 30 graded rollouts completed the full loop (inspect, edit, run the visible tests, submit) and the sealed grader accepted the submission. This demonstrates the task is solvable under the published limits.	
F-01 · derived from the recorded run data	
F-02	ATTENTION
context_exhaustion × 5 (codestral-latest)	
Severity: ATTENTION	Evidence: Appendix A · Appendix C
Ran out of output context before it could submit.	
F-02 · derived from the recorded run data	
F-03	ATTENTION
edited_tested_no_submit × 4 (codestral-latest)	
Severity: ATTENTION	Evidence: Appendix A · Appendix C
Edited files and ran the visible tests but never submitted.	
F-03 · derived from the recorded run data	
F-04	ATTENTION
submitted_failed × 1 (codestral-latest)	
Severity: ATTENTION	Evidence: Appendix A · Appendix C
Submitted a change; the independent hidden tests still failed.	
F-04 · derived from the recorded run data	
F-05	ATTENTION
1 overclaim(s) — success stated, then refuted by the grader	
Severity: ATTENTION	Evidence: Behavioural analysis
mistral:codestral-latest declared the task fixed and the independent hidden tests then failed. Counted claimed-and-wrong per rollout, from the model's own words.	
F-05 · derived from the recorded run data	

4 grounding failure(s) — paths cited that do not exist

Severity: ATTENTION

Evidence: Behavioural analysis

mistral:codestral-latest reasoned about files not present in the tree they were given, detected by matching cited source paths against the agent box on disk.

F-06 · derived from the recorded run data

10 Recommendations

One recommendation per observed failure mode, addressed to the lane it affects. Each traces back to a finding above and to the raw trace in Appendix A.

R-01

PROBLEM

Ran out of output context before it could submit.

AFFECTED LANES

codestral-latest

WHAT WE WOULD LOOK AT FIRST

The model fills its window before converging — working-set discipline over long files is the gap.

R-02

PROBLEM

Edited files and ran the visible tests but never submitted.

AFFECTED LANES

codestral-latest

WHAT WE WOULD LOOK AT FIRST

The model does the work and never calls submit — terminal-action reliability is the gap.

R-03

PROBLEM

Submitted a change; the independent hidden tests still failed.

AFFECTED LANES

codestral-latest

WHAT WE WOULD LOOK AT FIRST

The model completes the loop and produces a plausible near-miss — the named failing tests are the exact behavioral cases distinguishing it from the reference fix.

11 Verification

RUN ID

live-20260902T134741Z

EXAM FINGERPRINT

4a206fbfe8fb70fb3e37943705065a8599fab492e495dca326279781ee960b1e

The contract digest this run was executed against.

Tamper-evidence

Three digests, and how to check each one. If any of them does not reproduce, this page is not the page that was generated for this run.

THIS REPORT

274a5f976193ed9f3794876982ed85549a4f52cd44b282ee8681433e21320398

sha256 over the report's UTF-8 bytes with every occurrence of the report digest itself replaced by 64 '0' characters

EVAL RECORD

5af5f844bf99f5172f661f5b51ed39a4bedc7a35585899a62859cb85cb411461

sha256 of fr-20260902-a6f12b64.json as written beside this file.

RUN EVIDENCE BUNDLE

a6f12b647a54dbf0c0ecf80acda77ee562a8ccdb3ad1f44bccffa20c0716d664

The digest the run's own bundle computed over itself.

CERTIFIED EVIDENCE SEAL

not recorded

The exam's sealed evidence, established before this run.

Measurement history

This is the first recorded run on this exam. There is nothing to compare it against yet, and a single measurement is not a trend.

Verifying it, in three lines

1. vvdex-env verify-report --report fr-20260902-a6f12b64.html

2. or by hand: replace every occurrence of the report digest above with 64 zeros, then sha256 the file – it must equal that digest

3. shasum -a 256 fr-20260902-a6f12b64.json must equal the eval-record digest, and the bundle digest must equal bundleSha256 in the run's bundle file

The full artifact-by-artifact digest table, and the reproduction protocol, are in Appendix B.

Appendix A — every rollout

ROLLOUT	MODEL	SEED	OUTCOME	SUBMITTED	DEEPEST STAGE	HIDDEN TESTS	FILES	TOKENS	ELAPSED
3c178cb0	cli/claude-sonnet	0	submitted_passed	yes	Passed hidden tests	pass	3	n/a	74.7s
230a8274	cli/claude-sonnet	1	submitted_passed	yes	Passed hidden tests	pass	3	n/a	73.9s
459f9de3	cli/claude-sonnet	2	submitted_passed	yes	Passed hidden tests	pass	3	n/a	64.7s
3a7da64d	cli/claude-sonnet	3	submitted_passed	yes	Passed hidden tests	pass	3	n/a	92.5s
1e7f41d8	cli/claude-sonnet	4	submitted_passed	yes	Passed hidden tests	pass	3	n/a	83.8s
9902eaa3	cli/claude-sonnet	5	submitted_passed	yes	Passed hidden tests	pass	3	n/a	77.7s
12f95490	cli/claude-sonnet	6	submitted_passed	yes	Passed hidden tests	pass	3	n/a	84.0s
affa9ea3	cli/claude-sonnet	7	submitted_passed	yes	Passed hidden tests	pass	3	n/a	76.2s
ab27ff4e	cli/claude-sonnet	8	submitted_passed	yes	Passed hidden tests	pass	3	n/a	85.5s
cd337d97	cli/claude-sonnet	9	submitted_passed	yes	Passed hidden tests	pass	3	n/a	65.9s
879eca05	cli/codex	0	submitted_passed	yes	Passed hidden tests	pass	3	n/a	161.2s
fd449898	cli/codex	1	submitted_passed	yes	Passed hidden tests	pass	3	n/a	162.3s
5d078aad	cli/codex	2	submitted_passed	yes	Passed hidden tests	pass	3	n/a	223.6s
4ac68744	cli/codex	3	submitted_passed	yes	Passed hidden tests	pass	3	n/a	144.3s
7d1b1531	cli/codex	4	submitted_passed	yes	Passed hidden tests	pass	3	n/a	114.1s
662f4a3d	cli/codex	5	submitted_passed	yes	Passed hidden tests	pass	3	n/a	135.7s
660efecd	cli/codex	6	submitted_passed	yes	Passed hidden tests	pass	3	n/a	127.3s
51c40893	cli/codex	7	submitted_passed	yes	Passed hidden tests	pass	3	n/a	126.5s
899ad596	cli/codex	8	submitted_passed	yes	Passed hidden tests	pass	3	n/a	144.5s
e90dd0c8	cli/codex	9	submitted_passed	yes	Passed hidden tests	pass	3	n/a	134.5s
0711410a	codestral-latest	0	context_exhaustion	no	Graded	fail	1	80 881	75.3s
de0f078f	codestral-latest	1	context_exhaustion	no	Graded	fail	9	203 624	151.0s

ROLLOUT	MODEL	SEED	OUTCOME	SUBMITTED	DEEPEST STAGE	HIDDEN TESTS	FILES	TOKENS	ELAPSED
5c42bd02	codestral-latest	2	submitted_failed	yes	Graded	fail	0	4 126	7.4s
3498fcf5	codestral-latest	3	context_exhaustion	no	Graded	fail	9	175 540	100.7s
5e744f53	codestral-latest	4	edited_tested_no_submi t	no	Graded	fail	11	102 805	81.2s
b761b8be	codestral-latest	5	edited_tested_no_submi t	no	Graded	fail	10	89 896	94.4s
eddbc8c8	codestral-latest	6	edited_tested_no_submi t	no	Graded	fail	1	85 524	81.4s
b0c7e737	codestral-latest	7	edited_tested_no_submi t	no	Graded	fail	1	62 082	65.4s
3061565d	codestral-latest	8	context_exhaustion	no	Graded	fail	8	251 723	513.2s
8b5af559	codestral-latest	9	context_exhaustion	no	Graded	fail	9	91 689	99.7s

Failure taxonomy

CLASS	COUNT	WHAT IT MEANS
context_exhaustion	5	Ran out of output context before it could submit.
edited_tested_no_submit	4	Edited files and ran the visible tests but never submitted.
submitted_failed	1	Submitted a change; the independent hidden tests still failed.

Appendix B — evidence and artifact digests

RUN BUNDLE DIGEST

a6f12b647a54dbf0c0ecf80a
cda77ee562a8ccdb3ad1f44b
ccffa20c0716d664

RUN ID

live-
20260902T134741Z

STARTED

2026-09-
02T13:47:41.546440
+00:00

FINISHED

2026-09-
02T14:49:46.020126
+00:00

Artifact digests

ARTIFACT	SHA-256
publicSummary	92ec8ad502cca03db70ca82d7c593206728ddfa95c2e483a198203d73b81ccb6
rollouts/0711410a-20af-4809-97d9-eb6f1cb4ed52.json	e76288222b67f8622db779c912d61b460fba164efaf041126bd6dcb14a87cdc8
rollouts/12f95490-752e-4fae-8352-8fc3f600c46c.json	05fad9811badd35c5ba8e2c8b3884d87ffe09385cc1d870be5d8725bc9af84a2
rollouts/1e7f41d8-724c-4906-94aa-77195464bafe.json	01744a44277bacc58c08bcae2b1b80959ce3bf8c47588a7c787fd83453b30354
rollouts/230a8274-69d8-4aa7-a026-3972e84dee78.json	66a3d04b57a30ded39eafb676ee1fa846909384b59ad96016e4dfebb3386cff
rollouts/3061565d-c1ba-453b-b793-1e061a2e4a62.json	1f9200514a2b2c38b9ee97118974f2d5629406e67349552f13ad03b710fc49fd
rollouts/3498fcf5-bbb6-414b-aba0-a13bad481435.json	c9e901f199c5a36bda324d35fea1f3ffe95e3dc31907a59e48c79d53deaad30e
rollouts/3a7da64d-1728-43ad-b959-185afd23de73.json	4ee16b3a4a8d0092617a1dfee773295ce2a1acb590e9b58abeeda7aacb6cbbb6
rollouts/3c178cb0-987f-499b-b4c5-8e6a51d22739.json	2b77fd47c36b9ced52164b48be6988332475b6467ab185166d9bf6a9bad8b861
rollouts/459f9de3-dd18-4b28-9e46-eeecd169291.json	81ee5801eda8feea8903bdc3e959cfd4c4e495fe3a8039620ed058080a6009d4
rollouts/4ac68744-04cf-46df-9fde-24f77013e84c.json	64555ef1cfd29429efb7ffa3c4f0f611744ea7d31e27a9dc6087941312b3eb7e
rollouts/51c40893-6126-4096-88a2-2ed8a13df341.json	39fa04de1b8f70e1f90e4d1713206db95e2d0450c682658e3369b780e5b8a1b2
rollouts/5c42bd02-e324-467a-a4a8-d8264ec66778.json	17860fd2c9c0127b8a37dbdee8d3105b4add25c756e2b45ab4d38790e1520412
rollouts/5d078aad-9bc3-42b7-97b8-a615b5f9a87f.json	8463d90244f36a14c9095d900f885147fab82e5c6bf9e8406ea984e06ad1208a
rollouts/5e744f53-ec91-4f5b-85e2-d6f150221616.json	3e6c5a77724a182243f898fe82b806aed906eb9a4b156276f9677d471c01f409
rollouts/660efecd-4c0a-4092-9713-fd3677d3528f.json	7435d69b8bc0ea51765cde0eff6a2e1926c36059307be111e3c82807d9bf2646
rollouts/662f4a3d-ea26-4e49-b9ee-24a190ed2db3.json	6d0227476d2a60f1f595800852d3c2b187d38591cb7b7fdcec39071104c6ef3c
rollouts/7d1b1531-f014-45b5-93f8-c578d036d820.json	755def32bc1166a4edba484cef698b57f399394937a13bf5b15c8f1ca9b93074
rollouts/879eca05-6ed5-481f-a4c0-58a675435c85.json	fcff7d31005d2bc12895065eb0356058fd81fe6573074fbc73a54ce29ad96118
rollouts/899ad596-b098-4a27-b1a8-ddc53e4346f2.json	835c88ee481264fa1b0b844eb4ed17d76c60185ddf56214e73629beca582145e
rollouts/8b5af559-b0f7-4aaa-9748-435f7836143e.json	42b8c77bec005e209190119a559320f4908762c2b85460f19e97d0cbf4ebe051
rollouts/9902eaa3-4af4-4039-946a-2b4fbbb3eb0c.json	8f91bdd28114bb10fa0f57b3c9030dda7a96fc3c26b3f47a9fffb19f7d161cc3e
rollouts/ab27ff4e-329f-4ad0-b20e-cd10c9e7e4b6.json	f8c53bac344a5ae08a4f020dbe6c65b3f8d0f66275cad641651c26ff098fa84b
rollouts/affa9ea3-9250-418d-9507-8cfbf429b84c.json	47d6031bef983bfd8f5878d13bbf4ef24c192a42f7a79186490048b9745c788a
rollouts/b0c7e737-a914-495b-ba72-aee2b3fd18b1.json	4048e436ea6ce3169b884cbfc4620936645573daf579ad52a5429b04032c0c19
rollouts/b761b8be-4d5d-4934-8e0e-6d63e67c573d.json	b3d3e9ee57a3f991b1d0c127d595178e1c9367b8ca17879d948d08f9fa584477
rollouts/cd337d97-6960-429c-99dd-10a17b935b0d.json	0051949472962b064fd9b7934f6acbb607de390a516e2fc2b5f59d7d03efe4c9

ARTIFACT	SHA-256
rollouts/de0f078f-5b50-4853-873f-208337f892f9.json	0fdc5de0bbf5032da58e7fa47df2f5000c926f566ec1c3ff8fb450f1aea2ca91
rollouts/e90dd0c8-6384-4a8e-857d-560d91cc8a8d.json	6fa264a19e220ee863979f0f2ac37f82b75e269d3ba5d54b0a7dd2957786f56e
rollouts/eddbc8c8-36ca-4533-b298-2cd72c5d5c62.json	7c4702032ea6d2d74a3fcf758665e109cb780012c387e4c7c6cd53b5cf1cb72f
rollouts/fd449898-96bc-4c78-bd82-2d37535c2987.json	1030a0bd7437271001c13f01cd8f725b116c4c08d9e13ddc4d7132bc925865cc
spec	c5580395b0da0e89877b88e47c897742665c2711455762bc87c742b9f68019ec
summary	adcd3c1bdb43057c1e2f1114afe7303277a5f95db0d1786f64407524d6eead72

Reproducing this run

Every rollout can be re-derived: the exam is frozen (fingerprint above), the grader ran in the recorded container image, and each rollout's full trace — every tool call, every result, every word — is in the evidence bundle at its rollouts/<id>.json path. A re-run uses the same job shape: the exam id, your endpoint, n, and the runtime. The record id names this run forever.

Appendix C — redacted transcripts

For each rollout: the tool calls it made, in order — which tool, on which of the model's own paths, whether the call succeeded, and how many characters came back. What a tool returned is not here, and neither is the submission: on this exam the submitted content is the graded answer, so no diff, no answer value and no model claim is printed. A request for the reference solution or the hidden tests is recorded as a request; the path it asked for is replaced by a label. See Appendix D, “What this report withholds and why”.

Disclosure class applied to this exam: answer-key disclosure (ownership vvdex_proprietary — engine registry; family harness_eval): the exam is VVDex evaluation IP and the graded output is an answer or a keyed value, so no submission content, diff, model claim or hidden-assertion name is published.

3c178cb0 · cli:cli/claude-sonnet · seed 0 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	src/rates.py	ok	1 335
1	edit_file	src/rates.py	ok	50
2	run_tests	—	failed	95
3	run_tests	—	ok	1 581
4	read_file	src/rates.py	ok	157
5	write_file	src/rates.py	ok	62
6	run_tests	—	ok	1 581
7	read_file	src/rates.py	ok	157
8	write_file	src/rates.py	ok	61
9	run_tests	—	ok	268
10	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (210 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

230a8274 · cli:cli/claude-sonnet · seed 1 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	src/rates.py	ok	1 335
1	edit_file	src/rates.py	ok	50
2	run_tests	—	failed	95
3	run_tests	—	ok	1 581
4	read_file	src/rates.py	ok	157
5	write_file	src/rates.py	ok	62
6	run_tests	—	ok	1 581
7	read_file	src/rates.py	ok	157
8	write_file	src/rates.py	ok	61

STEP	TOOL	PATH	RESULT	CHARS
9	run_tests	—	ok	268
10	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (212 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

459f9de3 · cli:cli/claude-sonnet · seed 2 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	src/rates.py	ok	1 335
1	edit_file	src/rates.py	ok	50
2	run_tests	—	failed	95
3	run_tests	—	ok	1 581
4	read_file	src/rates.py	ok	157
5	write_file	src/rates.py	ok	62
6	run_tests	—	ok	1 581
7	read_file	src/rates.py	ok	157
8	write_file	src/rates.py	ok	61
9	run_tests	—	ok	268
10	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (210 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

3a7da64d · cli:cli/claude-sonnet · seed 3 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	src/rates.py	ok	1 335
1	edit_file	src/rates.py	ok	50
2	run_tests	—	failed	95
3	run_tests	—	ok	1 581
4	read_file	src/rates.py	ok	157
5	write_file	src/rates.py	ok	62
6	run_tests	—	ok	1 581
7	write_file	src/rates.py	ok	61
8	run_tests	—	ok	268
9	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (199 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

1e7f41d8 · cli:cli/claude-sonnet · seed 4 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	src/rates.py	ok	1 335
1	edit_file	src/rates.py	ok	50
2	run_tests	—	failed	95
3	run_tests	—	ok	1 581
4	read_file	src/rates.py	ok	157
5	write_file	src/rates.py	ok	62
6	run_tests	—	ok	1 581
7	read_file	src/rates.py	ok	157
8	write_file	src/rates.py	ok	61
9	run_tests	—	ok	268
10	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (210 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

9902eaa3 · cli:cli/claude-sonnet · seed 5 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	src/rates.py	ok	1 335
1	edit_file	src/rates.py	ok	50
2	run_tests	—	failed	95
3	run_tests	—	ok	1 581
4	read_file	src/rates.py	ok	157
5	write_file	src/rates.py	ok	62
6	run_tests	—	ok	1 581
7	read_file	src/rates.py	ok	157
8	write_file	src/rates.py	ok	61
9	run_tests	—	ok	268
10	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (210 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

12f95490 · cli:cli/claude-sonnet · seed 6 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	src/rates.py	ok	1 335
1	edit_file	src/rates.py	ok	50
2	run_tests	—	failed	95

STEP	TOOL	PATH	RESULT	CHARS
3	run_tests	–	ok	1 581
4	read_file	src/rates.py	ok	157
5	write_file	src/rates.py	ok	62
6	run_tests	–	ok	1 581
7	read_file	src/rates.py	ok	157
8	write_file	src/rates.py	ok	62
9	run_tests	–	ok	268
10	submit	–	ok	85

Submission diff withheld by the public-disclosure policy (214 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

affa9ea3 · cli:cli/claude-sonnet · seed 7 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	src/rates.py	ok	1 335
1	edit_file	src/rates.py	ok	50
2	run_tests	–	failed	95
3	run_tests	–	ok	1 581
4	read_file	src/rates.py	ok	157
5	write_file	src/rates.py	ok	62
6	run_tests	–	ok	1 581
7	write_file	src/rates.py	ok	61
8	run_tests	–	ok	268
9	submit	–	ok	85

Submission diff withheld by the public-disclosure policy (199 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

ab27ff4e · cli:cli/claude-sonnet · seed 8 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	src/rates.py	ok	1 335
1	edit_file	src/rates.py	ok	50
2	run_tests	–	failed	95
3	run_tests	–	ok	1 581
4	read_file	src/rates.py	ok	157
5	write_file	src/rates.py	ok	62
6	run_tests	–	ok	1 581
7	read_file	src/rates.py	ok	157
8	write_file	src/rates.py	ok	61

STEP	TOOL	PATH	RESULT	CHARS
9	run_tests	—	ok	268
10	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (212 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

cd337d97 · cli:cli/claude-sonnet · seed 9 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	src/rates.py	ok	1 335
1	edit_file	src/rates.py	ok	50
2	run_tests	—	failed	95
3	run_tests	—	ok	1 581
4	read_file	src/rates.py	ok	157
5	write_file	src/rates.py	ok	62
6	run_tests	—	ok	1 581
7	write_file	src/rates.py	ok	61
8	run_tests	—	ok	268
9	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (199 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

879eca05 · cli:cli/codex · seed 0 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 633
1	read_file	—	failed	35
2	read_file	src/rates.py	ok	1 335
3	edit_file	src/rates.py	ok	50
4	read_file	src/rates.py	ok	157
5	read_file	src/rates.py	ok	0
6	run_tests	—	failed	95
7	run_tests	—	ok	1 581
8	write_file	src/rates.py	ok	62
9	read_file	src/rates.py	ok	157
10	read_file	src/rates.py	ok	40
11	read_file	src/rates.py	ok	0
12	write_file	src/rates.py	ok	62
13	read_file	src/rates.py	ok	148
14	run_tests	—	ok	268

STEP	TOOL	PATH	RESULT	CHARS
15	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (270 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

fd449898 · cli:cli/codex · seed 1 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	list_files	.	ok	94
1	read_file	TASK.md	ok	1 633
2	read_file	src/rates.py	ok	1 335
3	edit_file	src/rates.py	ok	50
4	read_file	src/rates.py	ok	157
5	read_file	src/rates.py	ok	0
6	read_file	src/rates.py	ok	0
7	run_tests	—	failed	95
8	run_tests	—	ok	1 581
9	write_file	src/rates.py	ok	62
10	read_file	src/rates.py	ok	157
11	read_file	src/rates.py	ok	20
12	read_file	src/rates.py	ok	0
13	run_tests	—	ok	1 581
14	write_file	src/rates.py	ok	61
15	read_file	src/rates.py	ok	892
16	run_tests	—	ok	268
17	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (288 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

5d078aad · cli:cli/codex · seed 2 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	list_files	.	ok	94
1	read_file	TASK.md	ok	1 633
2	read_file	—	failed	35
3	read_file	src/rates.py	ok	1 335
4	read_file	tests/test_rates.py	ok	719
5	edit_file	src/rates.py	ok	50
6	read_file	src/rates.py	ok	157
7	read_file	src/rates.py	ok	0

STEP	TOOL	PATH	RESULT	CHARS
8	read_file	src/rates.py	ok	39
9	read_file	src/rates.py	ok	60
10	read_file	src/rates.py	ok	0
11	run_tests	—	failed	95
12	run_tests	—	ok	1 581
13	write_file	src/rates.py	ok	62
14	read_file	src/rates.py	ok	157
15	run_tests	—	ok	1 581
16	write_file	src/rates.py	ok	61
17	read_file	src/rates.py	ok	787
18	run_tests	—	ok	268
19	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (327 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

4ac68744 · cli:cli/codex · seed 3 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 633
1	read_file	—	failed	35
2	read_file	src/rates.py	ok	1 335
3	edit_file	src/rates.py	ok	50
4	read_file	src/rates.py	ok	157
5	read_file	src/rates.py	ok	0
6	run_tests	—	failed	95
7	read_file	src/rates.py	ok	0
8	run_tests	—	ok	1 581
9	write_file	src/rates.py	ok	62
10	read_file	src/rates.py	ok	157
11	run_tests	—	ok	1 581
12	write_file	src/rates.py	ok	61
13	read_file	src/rates.py	ok	791
14	run_tests	—	ok	268
15	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (284 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

7d1b1531 · cli:cli/codex · seed 4 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 633
1	read_file	—	failed	35
2	read_file	src/rates.py	ok	1 335
3	edit_file	src/rates.py	ok	50
4	run_tests	—	failed	95
5	run_tests	—	ok	1 581
6	read_file	src/rates.py	ok	157
7	write_file	src/rates.py	ok	62
8	read_file	src/rates.py	ok	157
9	write_file	src/rates.py	ok	62
10	read_file	src/rates.py	ok	1 278
11	run_tests	—	ok	268
12	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (213 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

662f4a3d · cli:cli/codex · seed 5 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 633
1	read_file	—	failed	35
2	read_file	src/rates.py	ok	1 335
3	read_file	tests/test_rates.py	ok	719
4	edit_file	src/rates.py	ok	50
5	read_file	src/rates.py	ok	157
6	read_file	src/rates.py	ok	0
7	run_tests	—	failed	95
8	run_tests	—	ok	1 581
9	write_file	src/rates.py	ok	62
10	read_file	src/rates.py	ok	157
11	read_file	src/rates.py	ok	20
12	write_file	src/rates.py	ok	62
13	read_file	src/rates.py	ok	1 198
14	run_tests	—	ok	268
15	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (246 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

660efecd · cli:cli/codex · seed 6 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 633
1	list_files	.	ok	94
2	read_file	src/rates.py	ok	1 335
3	read_file	tests/test_rates.py	ok	719
4	edit_file	src/rates.py	ok	50
5	read_file	src/rates.py	ok	157
6	run_tests	—	failed	95
7	run_tests	—	ok	1 581
8	write_file	src/rates.py	ok	62
9	read_file	src/rates.py	ok	157
10	read_file	src/rates.py	ok	0
11	read_file	src/rates.py	ok	0
12	write_file	src/rates.py	ok	62
13	run_tests	—	ok	268
14	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (237 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

51c40893 · cli:cli/codex · seed 7 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	list_files	.	ok	94
1	read_file	TASK.md	ok	1 633
2	read_file	src/rates.py	ok	1 335
3	read_file	tests/test_rates.py	ok	719
4	edit_file	src/rates.py	ok	50
5	run_tests	—	failed	95
6	run_tests	—	ok	1 581
7	read_file	src/rates.py	ok	157
8	write_file	src/rates.py	ok	62
9	read_file	src/rates.py	ok	157
10	read_file	src/rates.py	ok	20
11	write_file	src/rates.py	ok	61
12	read_file	src/rates.py	ok	892

STEP	TOOL	PATH	RESULT	CHARS
13	run_tests	–	ok	268
14	submit	–	ok	85

Submission diff withheld by the public-disclosure policy (255 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

899ad596 · cli:cli/codex · seed 8 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 633
1	read_file	–	failed	35
2	read_file	src/rates.py	ok	1 335
3	edit_file	src/rates.py	ok	50
4	read_file	src/rates.py	ok	157
5	read_file	src/rates.py	ok	0
6	read_file	src/rates.py	ok	60
7	run_tests	–	failed	95
8	run_tests	–	ok	1 581
9	write_file	src/rates.py	ok	62
10	read_file	src/rates.py	ok	157
11	read_file	src/rates.py	ok	20
12	run_tests	–	ok	1 581
13	write_file	src/rates.py	ok	61
14	read_file	src/rates.py	ok	892
15	run_tests	–	ok	268
16	submit	–	ok	85

Submission diff withheld by the public-disclosure policy (278 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

e90dd0c8 · cli:cli/codex · seed 9 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 633
1	list_files	.	ok	94
2	read_file	src/rates.py	ok	1 335
3	read_file	tests/test_rates.py	ok	719
4	edit_file	src/rates.py	ok	50
5	read_file	src/rates.py	ok	157
6	read_file	src/rates.py	ok	20
7	read_file	src/rates.py	ok	0

STEP	TOOL	PATH	RESULT	CHARS
8	write_file	src/rates.py	ok	62
9	read_file	src/rates.py	ok	157
10	run_tests	—	failed	95
11	run_tests	—	ok	1 581
12	write_file	src/rates.py	ok	62
13	read_file	src/rates.py	ok	1 168
14	run_tests	—	ok	268
15	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (267 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

0711410a · mistral:codestral-latest · seed 0 · context_exhaustion

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	—	failed	35
1	invalid_tool_call	"}	failed	103
2	list_files	.	ok	94
3	read_file	TASK.md	ok	1 633
4	read_file	src/rates.py	ok	1 335
5	run_tests	—	failed	95
6	edit_file	src/rates.py	ok	50
7	run_tests	—	ok	1 581
8	edit_file	src/rates.py	failed	29
9	edit_file	src/rates.py	failed	29
10	edit_file	src/rates.py	failed	29
11	edit_file	src/rates.py	failed	29
12	edit_file	src/rates.py	failed	29
13	edit_file	src/rates.py	failed	29
14	edit_file	src/rates.py	failed	29
15	edit_file	src/rates.py	failed	29
16	edit_file	src/rates.py	failed	29
17	edit_file	src/rates.py	failed	29
18	edit_file	src/rates.py	failed	29
19	edit_file	src/rates.py	failed	29
20	edit_file	src/rates.py	failed	29
21	edit_file	src/rates.py	failed	29
22	edit_file	src/rates.py	failed	29
23	edit_file	src/rates.py	failed	29

Submission diff withheld by the public-disclosure policy (350 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

de0f078f · mistral:codestral-latest · seed 1 · context_exhaustion

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	—	failed	35
1	invalid_tool_call	"}	failed	103
2	list_files	.	ok	94
3	read_file	src/rates.py	ok	1 335
4	edit_file	src/rates.py	ok	50
5	run_tests	—	failed	95
6	run_tests	—	ok	1 581
7	edit_file	src/rates.py	failed	29
8	edit_file	src/rates.py	failed	29
9	write_file	src/rates.py	ok	62
10	run_tests	—	ok	1 581
11	write_file	src/rates.py	ok	62
12	run_tests	—	ok	1 579
13	edit_file	src/rates.py	ok	50
14	run_tests	—	ok	1 579
15	edit_file	src/rates.py	ok	50
16	run_tests	—	ok	1 579
17	edit_file	src/rates.py	ok	50
18	run_tests	—	ok	1 579
19	edit_file	src/rates.py	ok	50
20	run_tests	—	ok	1 579
21	edit_file	src/rates.py	ok	50
22	run_tests	—	ok	1 579
23	edit_file	src/rates.py	ok	50

Submission diff withheld by the public-disclosure policy (341 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

5c42bd02 · mistral:codestral-latest · seed 2 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	src/rates.py	ok	1 335
1	run_tests	—	failed	95
2	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (57 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

3498fcf5 · mistral:codestral-latest · seed 3 · context_exhaustion

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	—	failed	35
1	invalid_tool_call	"}	failed	103
2	list_files	.	ok	94
3	read_file	TASK.md	ok	1 633
4	read_file	src/rates.py	ok	1 335
5	edit_file	src/rates.py	ok	50
6	run_tests	—	failed	95
7	run_tests	—	ok	1 581
8	edit_file	src/rates.py	ok	50
9	run_tests	—	ok	1 561
10	edit_file	src/rates.py	ok	50
11	run_tests	—	ok	1 561
12	edit_file	src/rates.py	ok	50
13	run_tests	—	ok	1 570
14	edit_file	src/rates.py	ok	50
15	run_tests	—	ok	1 543
16	edit_file	src/rates.py	ok	50
17	run_tests	—	ok	1 543
18	edit_file	src/rates.py	ok	50
19	run_tests	—	ok	1 543
20	edit_file	src/rates.py	ok	50
21	run_tests	—	ok	1 543
22	edit_file	src/rates.py	ok	50
23	run_tests	—	ok	1 543

Submission diff withheld by the public-disclosure policy (337 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

5e744f53 · mistral:codestral-latest · seed 4 · edited_tested_no_submit

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	—	failed	35
1	read_file	src/rates.py	ok	1 335
2	run_tests	—	failed	95
3	edit_file	src/rates.py	ok	50

STEP	TOOL	PATH	RESULT	CHARS
4	run_tests	–	ok	1 581
5	edit_file	src/rates.py	ok	50
6	run_tests	–	ok	1 561
7	edit_file	src/rates.py	ok	50
8	run_tests	–	ok	2 005
9	edit_file	src/rates.py	ok	50
10	run_tests	–	ok	1 999
11	edit_file	src/rates.py	ok	50
12	run_tests	–	ok	1 999
13	edit_file	src/rates.py	ok	50
14	run_tests	–	ok	1 999
15	edit_file	src/rates.py	ok	50
16	run_tests	–	ok	1 999
17	edit_file	src/rates.py	ok	50
18	run_tests	–	ok	1 999
19	edit_file	src/rates.py	ok	50
20	run_tests	–	ok	1 999
21	edit_file	src/rates.py	ok	50
22	run_tests	–	ok	1 999
23	edit_file	src/rates.py	ok	50

Submission diff withheld by the public-disclosure policy (339 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

b761b8be · mistral:codestral-latest · seed 5 · edited_tested_no_submit

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	–	failed	35
1	invalid_tool_call	"}	failed	103
2	list_files	.	ok	94
3	read_file	src/rates.py	ok	1 335
4	edit_file	src/rates.py	ok	50
5	run_tests	–	failed	95
6	run_tests	–	ok	1 581
7	edit_file	src/rates.py	ok	50
8	run_tests	–	ok	1 561
9	edit_file	src/rates.py	ok	50
10	run_tests	–	ok	1 539
11	edit_file	src/rates.py	ok	50

STEP	TOOL	PATH	RESULT	CHARS
12	run_tests	—	ok	1 540
13	edit_file	src/rates.py	ok	50
14	run_tests	—	ok	1 540
15	edit_file	src/rates.py	ok	50
16	run_tests	—	ok	1 540
17	edit_file	src/rates.py	ok	50
18	run_tests	—	ok	1 540
19	edit_file	src/rates.py	ok	50
20	run_tests	—	ok	1 540
21	edit_file	src/rates.py	ok	50
22	run_tests	—	ok	1 540
23	edit_file	src/rates.py	ok	50

Submission diff withheld by the public-disclosure policy (338 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

eddbc8c8 · mistral:codestral-latest · seed 6 · edited_tested_no_submit

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	—	failed	35
1	read_file	TASK.md	ok	1 633
2	read_file	src/rates.py	ok	1 335
3	edit_file	src/rates.py	ok	50
4	run_tests	—	failed	95
5	run_tests	—	ok	1 581
6	edit_file	src/rates.py	failed	29
7	edit_file	src/rates.py	failed	29
8	edit_file	src/rates.py	failed	29
9	edit_file	src/rates.py	failed	29
10	edit_file	src/rates.py	failed	29
11	edit_file	src/rates.py	failed	29
12	edit_file	src/rates.py	failed	29
13	edit_file	src/rates.py	failed	29
14	edit_file	src/rates.py	failed	29
15	edit_file	src/rates.py	failed	29
16	edit_file	src/rates.py	failed	29
17	edit_file	src/rates.py	failed	29
18	edit_file	src/rates.py	failed	29
19	edit_file	src/rates.py	failed	29

STEP	TOOL	PATH	RESULT	CHARS
20	edit_file	src/rates.py	failed	29
21	edit_file	src/rates.py	failed	29
22	edit_file	src/rates.py	failed	29
23	edit_file	src/rates.py	failed	29

Submission diff withheld by the public-disclosure policy (355 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

b0c7e737 · mistral:codestral-latest · seed 7 · edited_tested_no_submit

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	—	failed	35
1	invalid_tool_call	"}	failed	103
2	list_files	.	ok	94
3	read_file	src/rates.py	ok	1 335
4	edit_file	src/rates.py	ok	50
5	run_tests	—	failed	95
6	run_tests	—	ok	1 581
7	edit_file	src/rates.py	failed	29
8	edit_file	src/rates.py	failed	29
9	edit_file	src/rates.py	failed	29
10	edit_file	src/rates.py	failed	29
11	edit_file	src/rates.py	failed	29
12	edit_file	src/rates.py	failed	29
13	edit_file	src/rates.py	failed	29
14	edit_file	src/rates.py	failed	29
15	edit_file	src/rates.py	failed	29
16	edit_file	src/rates.py	failed	29
17	edit_file	src/rates.py	failed	29
18	edit_file	src/rates.py	failed	29
19	edit_file	src/rates.py	failed	29
20	edit_file	src/rates.py	failed	29
21	edit_file	src/rates.py	failed	29
22	edit_file	src/rates.py	failed	29
23	edit_file	src/rates.py	failed	29

Submission diff withheld by the public-disclosure policy (350 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

3061565d · mistral:codestral-latest · seed 8 · context_exhaustion

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	src/rates.py	ok	1 335
1	edit_file	src/rates.py	ok	50
2	run_tests	—	failed	95
3	run_tests	—	ok	1 581
4	read_file	src/rates.py	ok	157
5	edit_file	src/rates.py	ok	50
6	run_tests	—	ok	1 001
7	read_file	src/rates.py	ok	318
8	edit_file	src/rates.py	ok	50
9	run_tests	—	ok	1 001
10	read_file	src/rates.py	ok	318
11	edit_file	src/rates.py	ok	50
12	run_tests	—	ok	1 001
13	read_file	src/rates.py	ok	318
14	edit_file	src/rates.py	ok	50
15	run_tests	—	ok	1 001
16	read_file	src/rates.py	ok	318
17	edit_file	src/rates.py	ok	50
18	run_tests	—	ok	1 001
19	read_file	src/rates.py	ok	318
20	edit_file	src/rates.py	ok	50
21	run_tests	—	ok	1 001
22	read_file	src/rates.py	ok	318
23	edit_file	src/rates.py	ok	50

Submission diff withheld by the public-disclosure policy (343 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

8b5af559 · mistral:codestral-latest · seed 9 · context_exhaustion

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	—	failed	35
1	invalid_tool_call	"}	failed	103
2	list_files	.	ok	94
3	read_file	TASK.md	ok	1 633
4	read_file	src/rates.py	ok	1 335
5	run_tests	—	failed	95
6	run_tests	—	ok	1 583

STEP	TOOL	PATH	RESULT	CHARS
7	edit_file	src/rates.py	ok	50
8	run_tests	—	ok	1 581
9	edit_file	src/rates.py	ok	50
10	run_tests	—	ok	1 581
11	edit_file	src/rates.py	ok	50
12	run_tests	—	ok	1 561
13	edit_file	src/rates.py	ok	50
14	run_tests	—	ok	1 539
15	edit_file	src/rates.py	ok	50
16	run_tests	—	ok	1 562
17	edit_file	src/rates.py	ok	50
18	run_tests	—	ok	1 540
19	edit_file	src/rates.py	ok	50
20	run_tests	—	ok	1 562
21	edit_file	src/rates.py	ok	50
22	run_tests	—	ok	1 541
23	edit_file	src/rates.py	ok	50

Submission diff withheld by the public-disclosure policy (343 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

Appendix D — metric definitions and methods

Statistics, stage funnel and efficiency in this document are derived at render time from the sealed rollout rows and their raw tool events: lane errors leave every capability denominator; a rollout's stages are booleans read off its own event sequence (inspected = a successful list or read, or a successful browser open/snapshot, edited = the workspace changed, tested = run_tests invoked, submitted = submit invoked, graded = the independent grader returned a verdict, passed = certified pass); absent telemetry is n/a. The record's own summary blocks are the values computed when it was sealed; the rows and events are the evidence.

Denominators

n counts evaluable rollouts only: attempts minus lane errors. A rollout the lane ended (API or runtime failure) is reported but never graded, so it is in no rate, interval or pass@k.

Interval method

Wilson score interval, not the normal approximation. The normal approximation collapses to a zero-width interval at 0 successes — it would print a 0/2 run as '0%, certain'. For k successes in n trials with $p = k/n$: centre = $(p + z^2/2n) / (1 + z^2/n)$, half-width = $z/(1 + z^2/n) \cdot \sqrt{(p(1-p)/n + z^2/4n^2)}$, clamped to [0, 1], $z = 1.9600$.

pass@k

Unbiased estimator: $\text{pass@k} = 1 - C(n-c, k) / C(n, k)$, for $k \leq n$ only — reported only up to the number of rollouts actually run. It answers "given k independent attempts at this task, how often would at least one pass" — nothing more.

Reading a failed rollout

A failed rollout is a measurement, not a defect report. Model API failures and runtime failures are classified separately precisely because they say nothing about the model's ability on this task.

Stage definitions

Inspected

Listed or read at least one file in the repository — or, on a browser exam, opened or snapshotted the page.

Edited

Wrote at least one file that actually changed on disk.

Ran tests

Invoked the visible-test tool at least once.

Submitted

Called submit, ending the rollout on its own terms.

Graded

The independent hidden grader ran on the submission and returned a verdict.

Passed hidden tests

The hidden grader returned exit 0 and the rollout was certified.

Long-horizon terms

commandsRun

Shell commands executed inside the model's own container, against the command budget the contract declares. Only exams that advertise run_command have any.

deadEnds

Steps that moved nothing: a re-read, a repeat of the previous action, or a call to a tool this exam does not have. Counted, not judged.

hitStepLimit

Rollouts that ran out of steps without submitting. On a long task this is the number that separates 'could not finish' from 'finished wrong'.

revisits

Re-reads of a slice the rollout had already been given. On a long task, re-reading is how a model compensates for losing its own earlier context.

stepsUsed

Tool calls the rollout actually spent, out of the step limit the contract declares. A rollout that submitted early spent fewer, and that is a result, not a shortfall.

Grounding and overclaim terms

groundingFailure

A file path the model wrote in its own reasoning that does not exist in the tree it was given, or that names a directory it is not allowed to see. Detected by matching source-file paths in the model's text against the agent box on disk — not by judging the prose.

nonexistent_path

The model named a source file that is not in the repository it was given. It was reasoning about a file that does not exist.

overclaim

The model stated the tests pass or the issue is fixed, and the independent hidden tests then failed. Counted per rollout: claimed-and-wrong, or not.

reReadChurn

How many times the model read a file it had already read in the same rollout. It is a cost and a coherence signal, not a failure: re-reading is sometimes right.

requested_hidden_path

The model named a path under the reference solution, the hidden grader or the hidden test directory. Those are not in its box; the request could not be served.

Failure classes

context_exhaustion

Ran out of output context before it could submit.

edited_tested_no_submit

Edited files and ran the visible tests but never submitted.

submitted_failed

Submitted a change; the independent hidden tests still failed.

submitted_passed

Submitted a change that passed the independent hidden tests.

What this report withholds, and why

On this exam the model's submitted content IS the graded answer, so nothing derived from a submission is published: no diff excerpt, no answer value or answer JSON body, no reference or gold value, no expected hidden output, none of the model's own words, and not the names of the hidden assertions a rollout failed — a grader's name can carry the value it asserts. What Appendix C does publish for each rollout: its id, its lane, its seed, the tool it called at each step, the public-safe path it called the tool on, whether the call succeeded, how many characters came back, the stage it reached and its outcome. A rollout the grader accepted is reported as submitted and accepted — the numbers in this report are unaffected, because the verdict is the grader's, not the reader's.

Rule applied to this exam: answer-key disclosure (ownership vvdex_proprietary — engine registry; family harness_eval): the exam is VVDex evaluation IP and the graded output is an answer or a keyed value, so no submission content, diff, model claim or hidden-assertion name is published. The rule is a property of the exam family, not a judgement made per

document, and the same rule governs every edition of this report, the campaign report that embeds it as a chapter, this exam's public page and every published JSON derived from the same sealed record.

Attestation

This report was generated end to end by VVDex Forge from the sealed evidence of the evaluation run described above. Measurement values were populated from recorded artifacts, not manually entered. The digests in Verification let any holder of this file confirm that the report, the eval record and the run evidence are the ones this run produced — without contacting VVDex.

VVDex Forge · vvdexops.com	2026-09-02	vvdex-env 0.1.0
Issuer — certified evaluation environments	Date of issue	Engine · verifiers pin c521e4146026

No upstream repository: this exam is VVDex's own.
Post-seal benchmark. It does not stamp any certifying gate and did not rewrite the exam, the gold solution, the grader or the certified evidence.

— END OF REPORT · fr-20260902-a6f12b64 —