

Updating a cross-session fact

Answer from the newer authoritative roster, not
stale carried memory

EVALUATION REPORT

2026-09-01

Evaluation scope:

- Claude Code CLI (Sonnet, subscription)
- Codex CLI (GPT, subscription)
- codestral-latest

Prepared by:

VVDex Forge — generated directly from sealed run evidence

About this report

About VVDex Forge

VVDex Forge builds and certifies evaluation environments for AI agents. Every exam is sealed before any model sees it: the reference solution must pass, an empty submission must fail, the grader must reject near-misses, and the hidden tests never enter the agent's tree. Reports and the public catalog live at **vvdexops.com**.

Report integrity

This report is generated from run evidence; its digests verify authenticity. The sole test of a copy is the three-line procedure in Verification — a copy whose digests do not reproduce is not this report.

Scope

One exam · 30 rollout(s) · 20 tool steps and 4920s wall clock per rollout, stated in full in Evaluation configuration.

Boundary

Not a leaderboard. Not a general capability score. A post-seal benchmark: this run stamped no certification gate and rewrote nothing — not the exam, not the reference solution, not the grader, not the certified evidence.

Contents

01.	Executive summary	4
02.	Exam definition	5
03.	Certification evidence	6
04.	Evaluation configuration	7
05.	Model outcomes	8
06.	Stage funnel	9
07.	Behavioural analysis	10
08.	Efficiency & telemetry	11
09.	Findings	12
10.	Recommendations	13
11.	Verification	14

Appendices: A — every rollout · B — evidence digests · C — redacted transcripts · D — metric definitions · Attestation

01 Executive summary

MODELS TESTED 3	CERTIFIED PASSES 22 of 30 graded rollouts	SUBMITTED 30 of 30 graded rollouts that called submit	NOT GRADED — LANE ERRORS 0 of 30 attempts
---------------------------------------	--	--	--

Purpose

This report presents a controlled evaluation of 3 model lane(s) against the certified exam `vv dex.knowledge.memory-fact-update-1` (v0.1.0, family `long_term_memory`). Every lane received the identical frozen workspace, tool contract and limits; grading ran afterwards, in isolation, against hidden tests no lane ever saw. The purpose is to state, with reproducible evidence, what each lane did and where it stopped.

Outcome

22 of 30

graded rollouts passed the independent hidden tests and were certified

22 of 30 graded rollouts passed. `cli/claude-sonnet`, `cli/codex`, `codestral-latest` completed the full loop (inspect, edit, run the visible tests, submit) and the sealed grader accepted the submission. The most common way to fail was `submitted_failed` — submitted a change; the independent hidden tests still failed.

Key findings

F-01 `cli/claude-sonnet`, `cli/codex`, `codestral-latest` passed the independent hidden tests

F-02 `submitted_failed` × 8 (`codestral-latest`)

Each finding is stated in full, with evidence, in the Findings section; recommendations for acting on them follow it. Elapsed wall clock for the whole engagement: 695.7s.

02 Exam definition

In scope. The ability of each configured model lane to complete this one certified task end to end — inspect the workspace, produce a correct edit, verify it with the visible tests, and submit — within 20 tool steps and 4920s of wall clock, at 10 rollout(s) per lane. Behavioural measurements along the way (tool discipline, grounding, overclaim) are in scope because they are read from the recorded trace, not judged by another model.

Out of scope. General model quality, performance on other task families, prompt-tuning on the lane's behalf, and any comparison this run's sample size cannot support. A lane's API failure is reported as infrastructure, never as model capability.

The instrument. Exam `vvdex.knowledge.memory-fact-update-1` was certified before this run: its reference solution passes, an empty submission fails, its grader distinguishes near-misses, and its sealed evidence is unchanged by this evaluation (see Certification evidence).

What was measured

EXAM	VERSION	FAMILY	ROLLOUTS
vvdex.knowledge.memory-fact-update-1	0.1.0	long_term_memory	30 (10 per model)

Where the defect came from

This exam has no upstream repository to credit: the defect, the reference solution and the hidden suite are VVDex's own.

03 Certification evidence

4 of 5 certification pillars are recorded in this package, loaded from the receipts sealed into the exam and matching this run's exam fingerprint. What is recorded is shown with its real value; what is not is named below and is not claimed.

FROZEN BASELINE

FAIL

expected — an empty submission over 2 run(s) must not pass

→

GOLD / REFERENCE

PASS

expected — the reference solution over 2 run(s) must pass

GRADER CONTROLS

7 / 7

The grader accepted the reference and rejected every near-miss.

ATTACK PROBES

10 / 10 blocked

0 trivial exploit(s) found.

RUNTIME

sha256:679b36225a1f83238efba8c71b9cde3c24caaeacc9b682ae92819cb55eec328f

network policy: deny

Certification metadata is not included in this report package for:

- sealed evidence bundle

Benchmark results remain valid as run evidence; the pillar(s) above are not substantiated by this document and this report does not claim them.

What the package does carry: the exam fingerprint below is the contract digest this run was executed against — a holder of the certification bundle for this fingerprint can join the two independently.

EXAM FINGERPRINT

2052eec0ff2d3044dc269e3599ae1bfaa45d0d8cd5df9554796105538d226c21

The contract digest this run was executed against.

04 Evaluation configuration

REPORT	fr-20260901-46eb9548
ISSUED	2026-09-01
SUBJECT EXAM	vvdex.knowledge.memory-fact-update-1 · v0.1.0
FAMILY	long_term_memory
PREPARED BY	VVDex Forge engine vvdex-env 0.1.0
CLASSIFICATION	Independent model evaluation report — independently verifiable
CERTIFICATION	Exam certification: partially recorded (see Certification evidence)
STATUS	FINAL · report sealed against its own digest

Timeline

WHEN (UTC)	EVENT
2026-09-01 23:42:10	Evaluation run started
2026-09-01 23:53:46	Evaluation run finished
2026-09-01 23:53:46	Report generated from the sealed run evidence

Models under test

MODEL	PROVIDER	ENDPOINT	BILLING
Claude Code CLI (Sonnet, subscription)	cli	local runtime	operator subscription, flat — not billed per token
Codex CLI (GPT, subscription)	cli	local runtime	operator subscription, flat — not billed per token
codestral-latest	mistral	VVDex-configured	VVDex free-quota

A customer endpoint is recorded by origin only — never its port, its path, or its key.

Limits

RUNTIME docker	STEP LIMIT 20	WALL-CLOCK LIMIT 4920s	MAX OUTPUT TOKENS 2048
TEMPERATURE 0.2	RETRY POLICY none	COMMAND BUDGET 24	TOOLS OFFERED list_files, read_file, write_file, edit_file, run_tests, submit, run_command

The evaluated model never saw the hidden tests or the reference solution, and could not change its result. Grading ran in a separate step, after submission, on a container with no network.

05 Model outcomes

Denominators. Attempts requested: 30. Graded (evaluable) rollouts: 30. Not graded — lane errors: 0. A lane error is a rollout the API or the runtime ended before the hidden grader ran: it is listed, excluded from every rate and pass@k, and never silently dropped. Passes are certified passes: hidden grader exit 0 with integrity verified.

MODEL	ATTEMPTS	GRADED	PASSED	MODEL FAILS	LANE ERRORS	SUBMITTED	MEAN TIME	TYPICAL DEEPEST STAGE
Claude Code CLI (Sonnet, subscription) cli/claude-sonnet	10	10	10 / 10	0	0	10 / 10	17.0s	Passed hidden tests
Codex CLI (GPT, subscription) cli/codex	10	10	10 / 10	0	0	10 / 10	26.3s	Passed hidden tests
codestral-latest codestral-latest	10	10	2 / 10	8	0	10 / 10	7.4s	Graded

Every rollout, with its own deepest stage and grade, is in Appendix A.

Success rate, confidence and pass@k

Every rate on this page is a measurement of a finite number of rollouts, so every rate carries the interval that measurement actually supports. n is the evaluable count — attempts minus rollouts the lane ended before grading; those are listed under "not graded" and sit in no rate, interval or pass@k.

MODEL	ATTEMPTS	N	NOT GRADED	PASSES	RATE	95% INTERVAL	SEEDS	PASS@1	PASS@3	PASS@5	PASS@10
cli/claude-sonnet	10	10	0	10	100.0%	[72%, 100%]	0–9	100.0%	100.0%	100.0%	100.0%
cli/codex	10	10	0	10	100.0%	[72%, 100%]	0–9	100.0%	100.0%	100.0%	100.0%
codestral-latest	10	10	0	2	20.0%	[6%, 51%]	0–9	20.0%	53.3%	77.8%	100.0%

Interval and pass@k methods are stated in full in Appendix D. pass@k is shown for k = 1, 3, 5, 10; the sealed record carries every k up to n.

Model comparison

MODEL	PASS@1	95% INTERVAL	TOKENS	COST	MEAN ROLLOUT	OVERCLAIM RATE	TOP FAILURE CLASS
cli/claude-sonnet	100.0%	[72%, 100%]	n/a	n/a	17.0s	0.0%	none
cli/codex	100.0%	[72%, 100%]	n/a	n/a	26.3s	0.0%	none
codestral-latest	20.0%	[6%, 51%]	60 737	not billed	7.4s	0.0%	submitted_failed

Tokens and cost read n/a where a lane reports no telemetry (a flat-rate CLI) — an absent measurement, not zero. pass@1 and its interval are over evaluable rollouts only.

This is not a leaderboard. Separated by this run (intervals do not overlap): cli/claude-sonnet vs codestral-latest; cli/codex vs codestral-latest. Not separated (intervals overlap): cli/claude-sonnet vs cli/codex. The primary comparison in this report is the stage funnel below — where a model stops is a finding; a percentage over this many rollouts is not.

06 Stage funnel

The funnel is the primary reading of a run. Where a model stops says more than a single pass percentage does, especially at low N.

STAGE	CLI:CLI/CLAUDE-SONNET	CLI:CLI/CODEX	MISTRAL:CODESTRAL-LATEST
Inspected	10 / 10	10 / 10	10 / 10
Edited	10 / 10	10 / 10	10 / 10
Ran tests	10 / 10	10 / 10	10 / 10
Submitted	10 / 10	10 / 10	10 / 10
Graded	10 / 10	10 / 10	10 / 10
Passed hidden tests	10 / 10	10 / 10	2 / 10

Deepest stage reached

MODEL	DEEPEST (ANY ATTEMPT)	TYPICAL (MOST ATTEMPTS)	HIDDEN-TEST PASSES	NOT GRADED
cli:cli/claude-sonnet	Passed hidden tests	Passed hidden tests	10 / 10	0
cli:cli/codex	Passed hidden tests	Passed hidden tests	10 / 10	0
mistral:codestral-latest	Passed hidden tests	Graded	2 / 10	0

Stage definitions are in Appendix D. The matrix counts evaluable rollouts; lane errors are shown beside it, not inside it. The stages are not strictly nested: a model can submit without editing, and the matrix shows that rather than hiding it. Each rollout's own deepest stage is in Appendix A.

07 Behavioural analysis

Long-horizon behaviour

Counted off the recorded trace of each rollout. A dead end is a step that moved nothing: a re-read of a slice already returned, a repeat of the previous action, or a call to a tool this exam does not have.

MODEL	N	STEPS USED	RAN OUT OF STEPS	COMMANDS	CHANGED THE TREE	REFUSED (OVER BUDGET)	REVISITS	DEAD ENDS / ROLLOUT
cli:cli/claude-sonnet	10	5 / 20	0	0 / 24		0	0	0
cli:cli/codex	10	6 / 20	0	0 / 24		0	0	0
mistral:codestral-latest	10	4 / 20	0	0 / 24		0	0	0

Each column is defined in Appendix D.

Hallucination & overclaim

Three measurements, each with a stated definition. None of them is a judgement of the model's prose by another model — every number below comes from comparing what the model wrote against what was on disk, or against what the independent grader returned.

MODEL	N	GROUNDING FAILURES	KINDS	ROLLOUTS AFFECTED	OVERCLAIMS	OVERCLAIM RATE	RE-READS
cli:cli/claude-sonnet	10	0	none	0 / 10	0 / 10	0.0%	0
cli:cli/codex	10	0	none	0 / 10	0 / 10	0.0%	0
mistral:codestral-latest	10	0	none	0 / 10	0 / 10	0.0%	0

No grounding failures and no false pass claims were detected across these 30 rollout(s). No lane re-read material it had already been given.

Every term above is defined in Appendix D.

08 Efficiency & telemetry

This run has no per-token price attached — a local model, a free quota, or a customer endpoint we are not billed for — so spend is reported in tokens. Tokens are what was actually measured; a dollar figure here would be invented. Lanes run through a flat-rate local CLI report no token telemetry; their cells say n/a — an absent measurement, not zero. Rollout counts here are evaluable rollouts.

MODEL	ROLLOUTS	TOKENS	TOKENS / ROLLOUT	COST	PASSES	RATE
cli:cli/claude-sonnet	10	n/a	n/a	n/a – telemetry unavailable	10 / 10	100.0%
cli:cli/codex	10	n/a	n/a	n/a – telemetry unavailable	10 / 10	100.0%
mistral:codestral-latest	10	60 737	6 073	not billed	2 / 10	20.0%

09 Findings

Every finding below is derived from the recorded data of this run — none is an opinion.

Severity: PASS a lane succeeded · ATTENTION a capability gap · LANE / INFRA
infrastructure, not capability · NOTE a property of the measurement.

F-01

PASS

cli/claude-sonnet, cli/codex, codestral-latest passed the independent hidden tests

Severity: PASS

Evidence: Model outcomes · Stage funnel · Appendix A

22 of 30 graded rollouts completed the full loop (inspect, edit, run the visible tests, submit) and the sealed grader accepted the submission. This demonstrates the task is solvable under the published limits.

F-01 · derived from the recorded run data

F-02

ATTENTION

submitted_failed × 8 (codestral-latest)

Severity: ATTENTION

Evidence: Appendix A · Appendix C

Submitted a change; the independent hidden tests still failed.

F-02 · derived from the recorded run data

10 Recommendations

One recommendation per observed failure mode, addressed to the lane it affects. Each traces back to a finding above and to the raw trace in Appendix A.

R-01

PROBLEM

Submitted a change; the independent hidden tests still failed.

AFFECTED LANES

codestral-latest

WHAT WE WOULD LOOK AT FIRST

The model completes the loop and produces a plausible near-miss — the named failing tests are the exact behavioral cases distinguishing it from the reference fix.

11 Verification

RUN ID

live-20260901T234210Z

EXAM FINGERPRINT

2052eec0ff2d3044dc269e3599ae1bfaa45d0d8cd5df9554796105538d226c21

The contract digest this run was executed against.

Tamper-evidence

Three digests, and how to check each one. If any of them does not reproduce, this page is not the page that was generated for this run.

THIS REPORT

2faffd52a3924c70ff346f20cdec7690073de05cf1f96da886dd9d2e9815b969

sha256 over the report's UTF-8 bytes with every occurrence of the report digest itself replaced by 64 '0' characters

EVAL RECORD

d3f91221afd12e368b53c09a12eba41633c305d41855842ca6eef918de09e857

sha256 of fr-20260901-46eb9548.json as written beside this file.

RUN EVIDENCE BUNDLE

46eb954840000beb6e8466f409bb6104c91cd91a0d34efa18a41871246efd142

The digest the run's own bundle computed over itself.

CERTIFIED EVIDENCE SEAL

not recorded

The exam's sealed evidence, established before this run.

Measurement history

This is the first recorded run on this exam. There is nothing to compare it against yet, and a single measurement is not a trend.

Verifying it, in three lines

1. vvdex-env verify-report --report fr-20260901-46eb9548.html

2. or by hand: replace every occurrence of the report digest above with 64 zeros, then sha256 the file – it must equal that digest

3. shasum -a 256 fr-20260901-46eb9548.json must equal the eval-record digest, and the bundle digest must equal bundleSha256 in the run's bundle file

The full artifact-by-artifact digest table, and the reproduction protocol, are in Appendix B.

Appendix A — every rollout

ROLLOUT	MODEL	SEED	OUTCOME	SUBMITTED	DEEPEST STAGE	HIDDEN TESTS	FILES	TOKENS	ELAPSED
c3e59f84	cli/claude-sonnet	0	submitted_passed	yes	Passed hidden tests	pass	1	n/a	20.5s
3f5d33eb	cli/claude-sonnet	1	submitted_passed	yes	Passed hidden tests	pass	1	n/a	16.5s
91cd84ed	cli/claude-sonnet	2	submitted_passed	yes	Passed hidden tests	pass	1	n/a	15.7s
d3ccb200	cli/claude-sonnet	3	submitted_passed	yes	Passed hidden tests	pass	1	n/a	14.7s
a41f48cc	cli/claude-sonnet	4	submitted_passed	yes	Passed hidden tests	pass	1	n/a	18.9s
be0bb482	cli/claude-sonnet	5	submitted_passed	yes	Passed hidden tests	pass	1	n/a	15.6s
295bfb44	cli/claude-sonnet	6	submitted_passed	yes	Passed hidden tests	pass	1	n/a	15.8s
dd33c07a	cli/claude-sonnet	7	submitted_passed	yes	Passed hidden tests	pass	1	n/a	18.3s
c138c3a9	cli/claude-sonnet	8	submitted_passed	yes	Passed hidden tests	pass	1	n/a	15.1s
9d3daa5f	cli/claude-sonnet	9	submitted_passed	yes	Passed hidden tests	pass	1	n/a	18.8s
37ff761f	cli/codex	0	submitted_passed	yes	Passed hidden tests	pass	1	n/a	26.7s
7452d373	cli/codex	1	submitted_passed	yes	Passed hidden tests	pass	1	n/a	26.2s
9d7733ce	cli/codex	2	submitted_passed	yes	Passed hidden tests	pass	1	n/a	29.9s
83cc4801	cli/codex	3	submitted_passed	yes	Passed hidden tests	pass	1	n/a	26.9s
607bc653	cli/codex	4	submitted_passed	yes	Passed hidden tests	pass	1	n/a	26.5s
f495da98	cli/codex	5	submitted_passed	yes	Passed hidden tests	pass	1	n/a	25.9s
e880dfa5	cli/codex	6	submitted_passed	yes	Passed hidden tests	pass	1	n/a	25.0s
730a0c3f	cli/codex	7	submitted_passed	yes	Passed hidden tests	pass	1	n/a	24.3s
59a40ede	cli/codex	8	submitted_passed	yes	Passed hidden tests	pass	1	n/a	25.8s
1b564187	cli/codex	9	submitted_passed	yes	Passed hidden tests	pass	1	n/a	26.3s
d38a9234	codestral-latest	0	submitted_failed	yes	Graded	fail	1	5 764	7.2s
0f23b24f	codestral-latest	1	submitted_failed	yes	Graded	fail	1	5 764	7.2s
17fd7879	codestral-latest	2	submitted_failed	yes	Graded	fail	1	5 764	6.6s
cee5d0af	codestral-latest	3	submitted_passed	yes	Passed hidden tests	pass	1	7 279	8.1s
2e244621	codestral-latest	4	submitted_failed	yes	Graded	fail	1	5 800	7.1s
d85946ef	codestral-latest	5	submitted_failed	yes	Graded	fail	1	5 764	6.6s
1ba66ffd	codestral-latest	6	submitted_failed	yes	Graded	fail	1	5 749	6.4s
217e2513	codestral-latest	7	submitted_failed	yes	Graded	fail	1	5 735	6.9s
4f72002b	codestral-latest	8	submitted_passed	yes	Passed hidden tests	pass	1	7 354	10.0s
036f7a4e	codestral-latest	9	submitted_failed	yes	Graded	fail	1	5 764	7.9s

Failure taxonomy

CLASS	COUNT	WHAT IT MEANS
submitted_failed	8	Submitted a change; the independent hidden tests still failed.

Appendix B — evidence and artifact digests

RUN BUNDLE DIGEST

46eb954840000beb6e8466f409bb6104c91cd91a0d34efa18a41871246efd142

RUN ID

live-20260901T234210Z

STARTED

2026-09-01T23:42:10.844945+00:00

FINISHED

2026-09-01T23:53:46.554546+00:00

Artifact digests

ARTIFACT	SHA-256
publicSummary	f00eecfd8c1b1f59dba8b9a50ea2f63d453b4e0510a0c539a07f29b817333639
rollouts/036f7a4e-95a8-44f4-9416-4a630956bde5.json	d078da23ce6dd43330d34fbca2f32ae906d034258b3aa6ad54082119bffcbfd4
rollouts/0f23b24f-8799-4da2-8eed-9409735afdfc.json	1fcf3e129ae41a02f38af0d9a98f4363fa1440c365311bef7fac89e35c98d70f
rollouts/17fd7879-a548-4888-b365-393483b721f2.json	e96c42e097096a493d9b8a62c847c7ae08199ace0b64a9c0185742753a99e82b
rollouts/1b564187-de5c-4b9e-bba0-a3497ba2abc2.json	820522a65667acba036bb85c75e50022864897aeadf268460be1a69015d233a4
rollouts/1ba66ffd-229b-40e8-b3b0-2dae920c85b1.json	54313858ba041c81401c4e7730dca90280991e21d7014a191ba8a9dedba0ef2b
rollouts/217e2513-92b3-4f09-9148-da9988b90f7f.json	2405e3e91b701f9be6136075080b29913d9befbd38e5e7293664929941eff90c
rollouts/295bfb44-12e8-48f8-8bdc-f5df9bcecc9b.json	470e991bd1b3c34e333f75659a78b7766c90c4f98c1d01cdd5e61809ff2cb042
rollouts/2e244621-3ca6-446e-9945-ed98537cc81b.json	7b786e0eaa7fa5202436cd67418fc96bec6068b29d0b9a30f53b253a1920894c
rollouts/37ff761f-a2b3-41fa-9b42-26ed00fb4d26.json	0ab83298fae4aca11a763554a35211137c3e0d1fbe46d59466c5a04bfcacaf5c43
rollouts/3f5d33eb-55bb-4320-b67c-2c20616d2509.json	0c25ae80d7162cdc131aa3168c52798a84cd18c33ec974100028dcd9cd6d09a4
rollouts/4f72002b-67f3-42b3-928a-79d7022d5180.json	c08d218f7332e24c3008c9c55814e5e29923041a02082d90f7c80defe1e741b3
rollouts/59a40ede-e515-48e5-9aa7-cce5bdd309b1.json	b33d19ff66eac735813648cfd97eeda7dd2de54eb8b82ad3b415cec658cc8f0d
rollouts/607bc653-711f-45d6-a26b-7bdfd01b0142.json	6bba74788ab8b38728a1fb1a72f07888544667fc5052389a6db6add67d782f1f
rollouts/730a0c3f-2c1b-43a0-a6c7-f005a0c7a0d6.json	c31b768016f3e9260afc4c7c8478eecdd4406a35377cf4340a49ee3514eaa5eb8
rollouts/7452d373-a07f-4c66-951b-b8c82c82badd.json	16725edfff8953f43b792cf66a786bf628a0a0e1e062413fa2f6d8327357e420
rollouts/83cc4801-1d9e-432d-81b1-31caa05f6033.json	237bfa5642324e6eac921a7aa6740045e8159a3604c9349dfec9c3551f07ec0c
rollouts/91cd84ed-bdd4-4eb6-9de6-56238bcb1d17.json	05848a847f9a10cb04d4b458675838967c5f43d7ff3dd3f19b23358ff4e0b3ce
rollouts/9d3daa5f-a2fe-4469-87dc-539a6d2d3d45.json	0182e362994ab161f4cf7ce4ed1a5206c64ad1c21a847412635538c8864d8e1b
rollouts/9d7733ce-0c7c-4f9f-8491-2d58b27af9d7.json	872aca186137b7f25f1454c455310ab41a76481ea656e1c49dfccb7b9f21e97d
rollouts/a41f48cc-c287-42e1-81b0-152f4f2017e6.json	a16e127ec3f8e19a5a816a9ae53fd4f343336e65efbb0de94fe28991caacd5b5
rollouts/be0bb482-0aa0-4452-aebc-78d3267da5c6.json	e9969415dbc882c8c8d1a53b1cc7b23adf2525700c97cbdb046f6f291691f10e
rollouts/c138c3a9-d1fc-4396-8437-dfa9356e8f57.json	ffe83adba7a97ad3ef5b2f83a52968f7b382ad2e88f252464e060e78c319edfe
rollouts/c3e59f84-2de4-4f17-9165-32d58a341c89.json	6a323bc275c391f285dd1bb3c2c67cd504190eef3f19636da5339f2caccc16b2
rollouts/cee5d0af-4f81-49fa-9da2-3ec9f3f4c2b1.json	0941fecc01b8577f8aaa411cd47566e51749690ac3280ec036282c6d1753fd65
rollouts/d38a9234-5c00-4282-b6d1-8597b2cbd8b5.json	d89c7fa3b927cb83fac42b3cadf115e0232420df1dd9675804079c2251e5b883
rollouts/d3ccb200-349b-4f7d-83fb-8cdb6453a294.json	d8ffbb6990497de81d41d9c34e29bfcfb6665dba0dfe2ee00c4df4411b1db866f

ARTIFACT	SHA-256
rollouts/d85946ef-5125-4b10-806a-9d17237fed3f.json	2418001521dae6707aa7154745363f2e55f589a1fbcb04ff66eeda00c924a16a
rollouts/dd33c07a-e79c-4cfa-9367-eb1c99991173.json	62da7725d6f6ee0b5c479b7336dfa4d1aaee713977a6df765ac506f5e0585479
rollouts/e880dfa5-8f79-4a13-817f-5be17c6dfba7.json	b33a34cf8dc3133a80c45e33a78b1d8e386c03fe8d9f0f253143db986ae57ddc
rollouts/f495da98-b293-4bd7-8722-bfb4966bc3ce.json	f2a4ffc76691fa6e27d96a489f2410c0c649f91fbd0d41307ae46cbf5abf1d1f
spec	cdaaa1f547234167c055e49c22f93ad3965c87c32220a4bf8c4c7cf94ea5ce50
summary	d5fb8760400394f5e669eff1476bdfccb37e0140a48e65d4f6954001d4cafd99

Reproducing this run

Every rollout can be re-derived: the exam is frozen (fingerprint above), the grader ran in the recorded container image, and each rollout's full trace — every tool call, every result, every word — is in the evidence bundle at its rollouts/<id>.json path. A re-run uses the same job shape: the exam id, your endpoint, n, and the runtime. The record id names this run forever.

Appendix C — redacted transcripts

For each rollout: the tool calls it made, in order — which tool, on which of the model's own paths, whether the call succeeded, and how many characters came back. What a tool returned is not here, and neither is the submission: on this exam the submitted content is the graded answer, so no diff, no answer value and no model claim is printed. A request for the reference solution or the hidden tests is recorded as a request; the path it asked for is replaced by a label. See Appendix D, "What this report withholds and why".

Disclosure class applied to this exam: answer-key disclosure (ownership vvdex_proprietary — engine registry; family long_term_memory): the exam is VVDex evaluation IP and the graded output is an answer or a keyed value, so no submission content, diff, model claim or hidden-assertion name is published.

c3e59f84 · cli:cli/claude-sonnet · seed 0 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	memory/notes.md	ok	135
1	read_file	facts/current.json	ok	65
2	write_file	answers/answer.json	ok	68
3	run_tests	—	ok	181
4	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (9 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

3f5d33eb · cli:cli/claude-sonnet · seed 1 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	memory/notes.md	ok	135
1	read_file	facts/current.json	ok	65
2	write_file	answers/answer.json	ok	68
3	run_tests	—	ok	181
4	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (9 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

91cd84ed · cli:cli/claude-sonnet · seed 2 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	memory/notes.md	ok	135
1	read_file	facts/current.json	ok	65
2	write_file	answers/answer.json	ok	68
3	run_tests	—	ok	181
4	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (9 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

d3ccb200 · cli:cli/claude-sonnet · seed 3 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	memory/notes.md	ok	135
1	read_file	facts/current.json	ok	65
2	write_file	answers/answer.json	ok	96
3	run_tests	—	ok	181
4	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (9 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

a41f48cc · cli:cli/claude-sonnet · seed 4 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	memory/notes.md	ok	135
1	read_file	facts/current.json	ok	65
2	write_file	answers/answer.json	ok	68
3	run_tests	—	ok	181
4	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (9 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

be0bb482 · cli:cli/claude-sonnet · seed 5 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	memory/notes.md	ok	135
1	read_file	facts/current.json	ok	65
2	write_file	answers/answer.json	ok	68
3	run_tests	—	ok	181
4	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (9 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

295bfb44 · cli:cli/claude-sonnet · seed 6 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	memory/notes.md	ok	135
1	read_file	facts/current.json	ok	65

STEP	TOOL	PATH	RESULT	CHARS
2	write_file	answers/answer.json	ok	68
3	run_tests	—	ok	181
4	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (9 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

dd33c07a · cli:cli/claude-sonnet · seed 7 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	memory/notes.md	ok	135
1	read_file	facts/current.json	ok	65
2	write_file	answers/answer.json	ok	68
3	run_tests	—	ok	181
4	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (9 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

c138c3a9 · cli:cli/claude-sonnet · seed 8 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	memory/notes.md	ok	135
1	read_file	facts/current.json	ok	65
2	write_file	answers/answer.json	ok	68
3	run_tests	—	ok	181
4	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (9 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

9d3daa5f · cli:cli/claude-sonnet · seed 9 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	memory/notes.md	ok	135
1	read_file	facts/current.json	ok	65
2	write_file	answers/answer.json	ok	68
3	run_tests	—	ok	181
4	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (9 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

37ff761f · cli:cli/codex · seed 0 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	ISSUE.md	ok	1 220
1	read_file	facts/current.json	ok	65
2	read_file	memory/notes.md	ok	135
3	write_file	answers/answer.json	ok	68
4	run_tests	—	ok	181
5	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (9 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

7452d373 · cli:cli/codex · seed 1 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	ISSUE.md	ok	1 220
1	read_file	facts/current.json	ok	65
2	read_file	memory/notes.md	ok	135
3	write_file	answers/answer.json	ok	68
4	run_tests	—	ok	181
5	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (9 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

9d7733ce · cli:cli/codex · seed 2 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	ISSUE.md	ok	1 220
1	list_files	.	ok	57
2	read_file	facts/current.json	ok	65
3	read_file	memory/notes.md	ok	135
4	write_file	answers/answer.json	ok	68
5	run_tests	—	ok	181
6	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (9 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

83cc4801 · cli:cli/codex · seed 3 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	ISSUE.md	ok	1 220

STEP	TOOL	PATH	RESULT	CHARS
1	list_files	.	ok	57
2	read_file	facts/current.json	ok	65
3	read_file	memory/notes.md	ok	135
4	write_file	answers/answer.json	ok	68
5	run_tests	—	ok	181
6	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (9 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

607bc653 · cli:cli/codex · seed 4 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	ISSUE.md	ok	1 220
1	read_file	facts/current.json	ok	65
2	read_file	memory/notes.md	ok	135
3	write_file	answers/answer.json	ok	68
4	run_tests	—	ok	181
5	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (9 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

f495da98 · cli:cli/codex · seed 5 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	ISSUE.md	ok	1 220
1	read_file	memory/notes.md	ok	135
2	read_file	facts/current.json	ok	65
3	write_file	answers/answer.json	ok	68
4	run_tests	—	ok	181
5	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (9 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

e880dfa5 · cli:cli/codex · seed 6 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	ISSUE.md	ok	1 220
1	read_file	memory/notes.md	ok	135
2	read_file	facts/current.json	ok	65
3	write_file	answers/answer.json	ok	68

STEP	TOOL	PATH	RESULT	CHARS
4	run_tests	—	ok	181
5	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (9 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

730a0c3f · cli:cli/codex · seed 7 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	ISSUE.md	ok	1 220
1	read_file	memory/notes.md	ok	135
2	read_file	facts/current.json	ok	65
3	write_file	answers/answer.json	ok	68
4	run_tests	—	ok	181
5	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (9 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

59a40ede · cli:cli/codex · seed 8 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	ISSUE.md	ok	1 220
1	read_file	memory/notes.md	ok	135
2	read_file	facts/current.json	ok	65
3	write_file	answers/answer.json	ok	68
4	run_tests	—	ok	181
5	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (9 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

1b564187 · cli:cli/codex · seed 9 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	ISSUE.md	ok	1 220
1	read_file	facts/current.json	ok	65
2	read_file	memory/notes.md	ok	135
3	write_file	answers/answer.json	ok	68
4	run_tests	—	ok	181
5	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (9 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

d38a9234 · mistral:codestral-latest · seed 0 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	memory/notes.md	ok	135
1	write_file	answers/answer.json	ok	68
2	run_tests	—	ok	181
3	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (4 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

0f23b24f · mistral:codestral-latest · seed 1 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	memory/notes.md	ok	135
1	write_file	answers/answer.json	ok	68
2	run_tests	—	ok	181
3	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (4 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

17fd7879 · mistral:codestral-latest · seed 2 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	memory/notes.md	ok	135
1	write_file	answers/answer.json	ok	68
2	run_tests	—	ok	181
3	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (4 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

cee5d0af · mistral:codestral-latest · seed 3 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	facts/current.json	ok	65
1	read_file	memory/notes.md	ok	135
2	write_file	answers/answer.json	ok	68
3	run_tests	—	ok	181
4	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (4 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

2e244621 · mistral:codestral-latest · seed 4 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	memory/notes.md	ok	135
1	write_file	answers/answer.json	ok	68
2	run_tests	—	ok	181
3	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (9 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

d85946ef · mistral:codestral-latest · seed 5 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	memory/notes.md	ok	135
1	write_file	answers/answer.json	ok	68
2	run_tests	—	ok	181
3	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (4 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

1ba66ffd · mistral:codestral-latest · seed 6 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	memory/notes.md	ok	135
1	write_file	answers/answer.json	ok	68
2	run_tests	—	ok	181
3	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (4 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

217e2513 · mistral:codestral-latest · seed 7 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	facts/current.json	ok	65
1	write_file	answers/answer.json	ok	68
2	run_tests	—	ok	181
3	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (9 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is

the graded answer, so publishing an excerpt of it would publish the answer key.

4f72002b · mistral:codestral-latest · seed 8 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	facts/current.json	ok	65
1	read_file	memory/notes.md	ok	135
2	run_tests	—	ok	181
3	write_file	answers/answer.json	ok	68
4	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (9 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

036f7a4e · mistral:codestral-latest · seed 9 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	memory/notes.md	ok	135
1	write_file	answers/answer.json	ok	68
2	run_tests	—	ok	181
3	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (4 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

Appendix D — metric definitions and methods

Statistics, stage funnel and efficiency in this document are derived at render time from the sealed rollout rows and their raw tool events: lane errors leave every capability denominator; a rollout's stages are booleans read off its own event sequence (inspected = a successful list or read, or a successful browser open/snapshot, edited = the workspace changed, tested = run_tests invoked, submitted = submit invoked, graded = the independent grader returned a verdict, passed = certified pass); absent telemetry is n/a. The record's own summary blocks are the values computed when it was sealed; the rows and events are the evidence.

Denominators

n counts evaluable rollouts only: attempts minus lane errors. A rollout the lane ended (API or runtime failure) is reported but never graded, so it is in no rate, interval or pass@k.

Interval method

Wilson score interval, not the normal approximation. The normal approximation collapses to a zero-width interval at 0 successes — it would print a 0/2 run as '0%, certain'. For k successes in n trials with $p = k/n$: centre = $(p + z^2/2n) / (1 + z^2/n)$, half-width = $z/(1 + z^2/n) \cdot \sqrt{(p(1-p)/n + z^2/4n^2)}$, clamped to [0, 1], $z = 1.9600$.

pass@k

Unbiased estimator: $\text{pass@k} = 1 - C(n-c, k) / C(n, k)$, for $k \leq n$ only — reported only up to the number of rollouts actually run. It answers "given k independent attempts at this task, how often would at least one pass" — nothing more.

Reading a failed rollout

A failed rollout is a measurement, not a defect report. Model API failures and runtime failures are classified separately precisely because they say nothing about the model's ability on this task.

Stage definitions

Inspected

Listed or read at least one file in the repository — or, on a browser exam, opened or snapshotted the page.

Edited

Wrote at least one file that actually changed on disk.

Ran tests

Invoked the visible-test tool at least once.

Submitted

Called submit, ending the rollout on its own terms.

Graded

The independent hidden grader ran on the submission and returned a verdict.

Passed hidden tests

The hidden grader returned exit 0 and the rollout was certified.

Long-horizon terms

commandsRun

Shell commands executed inside the model's own container, against the command budget the contract declares. Only exams that advertise run_command have any.

deadEnds

Steps that moved nothing: a re-read, a repeat of the previous action, or a call to a tool this exam does not have. Counted, not judged.

hitStepLimit

Rollouts that ran out of steps without submitting. On a long task this is the number that separates 'could not finish' from 'finished wrong'.

revisits

Re-reads of a slice the rollout had already been given. On a long task, re-reading is how a model compensates for losing its own earlier context.

stepsUsed

Tool calls the rollout actually spent, out of the step limit the contract declares. A rollout that submitted early spent fewer, and that is a result, not a shortfall.

Grounding and overclaim terms

groundingFailure

A file path the model wrote in its own reasoning that does not exist in the tree it was given, or that names a directory it is not allowed to see. Detected by matching source-file paths in the model's text against the agent box on disk — not by judging the prose.

nonexistent_path

The model named a source file that is not in the repository it was given. It was reasoning about a file that does not exist.

overclaim

The model stated the tests pass or the issue is fixed, and the independent hidden tests then failed. Counted per rollout: claimed-and-wrong, or not.

reReadChurn

How many times the model read a file it had already read in the same rollout. It is a cost and a coherence signal, not a failure: re-reading is sometimes right.

requested_hidden_path

The model named a path under the reference solution, the hidden grader or the hidden test directory. Those are not in its box; the request could not be served.

Failure classes

submitted_failed

Submitted a change; the independent hidden tests still failed.

submitted_passed

Submitted a change that passed the independent hidden tests.

What this report withholds, and why

On this exam the model's submitted content IS the graded answer, so nothing derived from a submission is published: no diff excerpt, no answer value or answer JSON body, no reference or gold value, no expected hidden output, none of the model's own words, and not the names of the hidden assertions a rollout failed — a grader's name can carry the value it asserts. What Appendix C does publish for each rollout: its id, its lane, its seed, the tool it called at each step, the public-safe path it called the tool on, whether the call succeeded, how many characters came back, the stage it reached and its outcome. A rollout the grader accepted is reported as submitted and accepted — the numbers in this report are unaffected, because the verdict is the grader's, not the reader's.

Rule applied to this exam: answer-key disclosure (ownership vvdex_proprietary — engine registry; family long_term_memory): the exam is VVDex evaluation IP and the graded output is an answer or a keyed value, so no submission content, diff, model claim or hidden-assertion name is published. The rule is a property of the exam family, not a judgement made per document, and the same rule governs every edition of this report, the campaign report that embeds it as a chapter, this exam's public page and every published JSON derived from the same sealed record.

This exam revision is retired for future evaluation

Retired for future evaluation on 2026-09-02: post-run public disclosure of answer-bearing submission content. Historical sealed results on this revision are unaffected: the disclosure happened after those runs, so it could not have influenced them.

Attestation

This report was generated end to end by VVDex Forge from the sealed evidence of the evaluation run described above. Measurement values were populated from recorded artifacts, not manually entered. The digests in Verification let any holder of this file confirm that the report, the eval record and the run evidence are the ones this run produced — without contacting VVDex.

VVDex Forge · vvdexops.com	2026-09-01	vvdex-env 0.1.0
Issuer — certified evaluation environments	Date of issue	Engine · verifiers pin c521e4146026

No upstream repository: this exam is VVDex's own.
Post-seal benchmark. It does not stamp any certifying gate and did not rewrite the exam, the gold solution, the grader or the certified evidence.

— END OF REPORT · fr-20260901-46eb9548 —