

Streaming items must equal what the full parse would have built

EVALUATION REPORT

2026-09-02

Evaluation scope:

- Claude Code CLI (Sonnet, subscription)
- Codex CLI (GPT, subscription)
- codestral-latest

Prepared by:

VVDex Forge — generated directly from sealed run evidence

fr-20260902-6616b188 · public.swe.martinblech-xmltodict-issue-257 v0.1.0 · Independent model evaluation report

About this report

About VVDex Forge

VVDex Forge builds and certifies evaluation environments for AI agents. Every exam is sealed before any model sees it: the reference solution must pass, an empty submission must fail, the grader must reject near-misses, and the hidden tests never enter the agent's tree. Reports and the public catalog live at vvdexops.com.

Report integrity

This report is generated from run evidence; its digests verify authenticity. The sole test of a copy is the three-line procedure in Verification — a copy whose digests do not reproduce is not this report.

Scope

One exam · 30 rollout(s) · 40 tool steps and 9720s wall clock per rollout, stated in full in Evaluation configuration.

Boundary

Not a leaderboard. Not a general capability score. A post-seal benchmark: this run stamped no certification gate and rewrote nothing — not the exam, not the reference solution, not the grader, not the certified evidence.

Contents

01.	Executive summary	4
02.	Exam definition	5
03.	Certification evidence	6
04.	Evaluation configuration	7
05.	Model outcomes	8
06.	Stage funnel	9
07.	Behavioural analysis	10
08.	Efficiency & telemetry	11
09.	Findings	12
10.	Recommendations	13
11.	Verification	14

Appendices: A — every rollout · B — evidence digests · C — redacted transcripts · D — metric definitions · Attestation

01 Executive summary

MODELS TESTED 3	CERTIFIED PASSES 19 of 30 graded rollouts	SUBMITTED 30 of 30 graded rollouts that called submit	NOT GRADED — LANE ERRORS 0 of 30 attempts
---------------------------------------	--	--	--

Purpose

This report presents a controlled evaluation of 3 model lane(s) against the certified exam public.swe.martinblech-xmltodict-issue-257 (v0.1.0, family public_oss_swe). Every lane received the identical frozen workspace, tool contract and limits; grading ran afterwards, in isolation, against hidden tests no lane ever saw. The purpose is to state, with reproducible evidence, what each lane did and where it stopped.

Outcome

19 of 30

graded rollouts passed the independent hidden tests and were certified

19 of 30 graded rollouts passed. cli/claude-sonnet, cli/codex completed the full loop (inspect, edit, run the visible tests, submit) and the sealed grader accepted the submission. The most common way to fail was submitted_failed — submitted a change; the independent hidden tests still failed.

Key findings

F-01 cli/claude-sonnet, cli/codex passed the independent hidden tests

F-02 submitted_failed × 11 (cli/claude-sonnet, codestral-latest)

Each finding is stated in full, with evidence, in the Findings section; recommendations for acting on them follow it. Elapsed wall clock for the whole engagement: 2327.3s.

02 Exam definition

In scope. The ability of each configured model lane to complete this one certified task end to end — inspect the workspace, produce a correct edit, verify it with the visible tests, and submit — within 40 tool steps and 9720s of wall clock, at 10 rollout(s) per lane. Behavioural measurements along the way (tool discipline, grounding, overclaim) are in scope because they are read from the recorded trace, not judged by another model.

Out of scope. General model quality, performance on other task families, prompt-tuning on the lane's behalf, and any comparison this run's sample size cannot support. A lane's API failure is reported as infrastructure, never as model capability.

The instrument. Exam `public.swe.martinblech-xmltodict-issue-257` was certified before this run: its reference solution passes, an empty submission fails, its grader distinguishes near-misses, and its sealed evidence is unchanged by this evaluation (see Certification evidence).

What was measured

EXAM	VERSION	FAMILY	ROLLOUTS
public.swe.martinblech-xmltodict-issue-257	0.1.0	public_oss_swe	30 (10 per model)

Where the defect came from

Source: `martinblech/xmltodict (MIT)` · issue
<https://github.com/martinblech/xmltodict/issues/257> · upstream fix
<https://github.com/martinblech/xmltodict/pull/421>

The defect and its upstream fix are public and unmodified. The exam around them — the frozen parent tree, the independently written hidden suite, the control probes and the isolation — is VVDex's, and it is what this report measures.

03 Certification evidence

4 of 5 certification pillars are recorded in this package, loaded from the receipts sealed into the exam and matching this run's exam fingerprint. What is recorded is shown with its real value; what is not is named below and is not claimed.

FROZEN BASELINE

FAIL

expected — an empty submission over 2 run(s) must not pass

→

GOLD / REFERENCE

PASS

expected — the reference solution over 2 run(s) must pass

GRADER CONTROLS

7 / 7

The grader accepted the reference and rejected every near-miss.

ATTACK PROBES

10 / 10 blocked

0 trivial exploit(s) found.

RUNTIME

sha256:679b36225a1f83238
efba8c71b9cde3c24caaeacc9
b682ae92819cb55eec328f

network policy: deny

Certification metadata is not included in this report package for:

- sealed evidence bundle

Benchmark results remain valid as run evidence; the pillar(s) above are not substantiated by this document and this report does not claim them.

What the package does carry: the exam fingerprint below is the contract digest this run was executed against — a holder of the certification bundle for this fingerprint can join the two independently.

EXAM FINGERPRINT

aa41977120e6ecdddb5c4fc603e8ab23cb272726a86750119969ff4dfc90e297

The contract digest this run was executed against.

04 Evaluation configuration

REPORT	fr-20260902-6616b188
ISSUED	2026-09-02
SUBJECT EXAM	public.swe.martinblech-xmltodict-issue-257 · v0.1.0
FAMILY	public_oss_swe
PREPARED BY	VVDex Forge engine vvdex-env 0.1.0
CLASSIFICATION	Independent model evaluation report — independently verifiable
CERTIFICATION	Exam certification: partially recorded (see Certification evidence)
STATUS	FINAL · report sealed against its own digest

Timeline

WHEN (UTC)	EVENT
2026-09-02 00:18:47	Evaluation run started
2026-09-02 00:57:34	Evaluation run finished
2026-09-02 00:57:34	Report generated from the sealed run evidence

Models under test

MODEL	PROVIDER	ENDPOINT	BILLING
Claude Code CLI (Sonnet, subscription)	cli	local runtime	operator subscription, flat — not billed per token
Codex CLI (GPT, subscription)	cli	local runtime	operator subscription, flat — not billed per token
codestral-latest	mistral	VVDex-configured	VVDex free-quota

A customer endpoint is recorded by origin only — never its port, its path, or its key.

Limits

RUNTIME docker	STEP LIMIT 40	WALL-CLOCK LIMIT 9720s	MAX OUTPUT TOKENS 2048
TEMPERATURE 0.2	RETRY POLICY none	COMMAND BUDGET 0	TOOLS OFFERED list_file, read_file, write_file, edit_file, run_tests, submit

The evaluated model never saw the hidden tests or the reference solution, and could not change its result. Grading ran in a separate step, after submission, on a container with no network.

05 Model outcomes

Denominators. Attempts requested: 30. Graded (evaluable) rollouts: 30. Not graded — lane errors: 0. A lane error is a rollout the API or the runtime ended before the hidden grader ran: it is listed, excluded from every rate and pass@k, and never silently dropped. Passes are certified passes: hidden grader exit 0 with integrity verified.

MODEL	ATTEMPTS	GRADED	PASSED	MODEL FAILS	LANE ERRORS	SUBMITTED	MEAN TIME	TYPICAL DEEPEST STAGE
Claude Code CLI (Sonnet, subscription) cli/claude-sonnet	10	10	9 / 10	1	0	10 / 10	159.8s	Passed hidden tests
Codex CLI (GPT, subscription) cli/codex	10	10	10 / 10	0	0	10 / 10	42.4s	Passed hidden tests
codestral-latest codestral-latest	10	10	0 / 10	10	0	10 / 10	11.5s	Graded

Every rollout, with its own deepest stage and grade, is in Appendix A.

Success rate, confidence and pass@k

Every rate on this page is a measurement of a finite number of rollouts, so every rate carries the interval that measurement actually supports. n is the evaluable count — attempts minus rollouts the lane ended before grading; those are listed under "not graded" and sit in no rate, interval or pass@k.

MODEL	ATTEMPTS	N	NOT GRADED	PASSES	RATE	95% INTERVAL	SEEDS	PASS@1	PASS@3	PASS@5	PASS@10
cli/claude-sonnet	10	10	0	9	90.0%	[60%, 98%]	0–9	90.0%	100.0%	100.0%	100.0%
cli/codex	10	10	0	10	100.0%	[72%, 100%]	0–9	100.0%	100.0%	100.0%	100.0%
codestral-latest	10	10	0	0	0.0%	[0%, 28%]	0–9	0.0%	0.0%	0.0%	0.0%

Interval and pass@k methods are stated in full in Appendix D. pass@k is shown for k = 1, 3, 5, 10; the sealed record carries every k up to n.

Model comparison

MODEL	PASS@1	95% INTERVAL	TOKENS	COST	MEAN ROLLOUT	OVERCLAIM RATE	TOP FAILURE CLASS
cli/claude-sonnet	90.0%	[60%, 98%]	n/a	n/a	159.8s	0.0%	submitted_failed
cli/codex	100.0%	[72%, 100%]	n/a	n/a	42.4s	0.0%	none
codestral-latest	0.0%	[0%, 28%]	172 890	not billed	11.5s	0.0%	submitted_failed

Tokens and cost read n/a where a lane reports no telemetry (a flat-rate CLI) — an absent measurement, not zero. pass@1 and its interval are over evaluable rollouts only.

This is not a leaderboard. Separated by this run (intervals do not overlap): cli/claude-sonnet vs codestral-latest; cli/codex vs codestral-latest. Not separated (intervals overlap): cli/claude-sonnet vs cli/codex. The primary comparison in this report is the stage funnel below — where a model stops is a finding; a percentage over this many rollouts is not.

06 Stage funnel

The funnel is the primary reading of a run. Where a model stops says more than a single pass percentage does, especially at low N.

STAGE	CLI:CLI/CLAUDE-SONNET	CLI:CLI/CODEX	MISTRAL:CODESTRAL-LATEST
Inspected	10 / 10	10 / 10	10 / 10
Edited	9 / 10	10 / 10	10 / 10
Ran tests	10 / 10	10 / 10	10 / 10
Submitted	10 / 10	10 / 10	10 / 10
Graded	10 / 10	10 / 10	10 / 10
Passed hidden tests	9 / 10	10 / 10	0 / 10

Deepest stage reached

MODEL	DEEPEST (ANY ROLLOUT)	TYPICAL (MOST ROLLOUTS)	HIDDEN-TEST PASSES	NOT GRADED
cli:cli/claude-sonnet	Passed hidden tests	Passed hidden tests	9 / 10	0
cli:cli/codex	Passed hidden tests	Passed hidden tests	10 / 10	0
mistral:codestral-latest	Graded	Graded	0 / 10	0

Stage definitions are in Appendix D. The matrix counts evaluable rollouts; lane errors are shown beside it, not inside it. The stages are not strictly nested: a model can submit without editing, and the matrix shows that rather than hiding it. Each rollout's own deepest stage is in Appendix A.

07 Behavioural analysis

Long-horizon behaviour

Counted off the recorded trace of each rollout. A dead end is a step that moved nothing: a re-read of a slice already returned, a repeat of the previous action, or a call to a tool this exam does not have.

MODEL	N	STEPS USED	RAN OUT OF STEPS	COMMANDS	CHANGED THE TREE	REFUSED (OVER BUDGET)	REVISITS	DEAD ENDS / ROLLOUT
cli:cli/claude-sonnet	10	5 / 40	0	0 / 0		0	0	0
cli:cli/codex	10	6 / 40	0	0 / 0		0	0	0
mistral:codestral-latest	10	5 / 40	0	0 / 0		0	0	0

Each column is defined in Appendix D.

Hallucination & overclaim

Three measurements, each with a stated definition. None of them is a judgement of the model's prose by another model — every number below comes from comparing what the model wrote against what was on disk, or against what the independent grader returned.

MODEL	N	GROUNDING FAILURES	KINDS	ROLLOUTS AFFECTED	OVERCLAIMS	OVERCLAIM RATE	RE-READS
cli:cli/claude-sonnet	10	0	none	0 / 10	0 / 10	0.0%	0
cli:cli/codex	10	0	none	0 / 10	0 / 10	0.0%	0
mistral:codestral-latest	10	0	none	0 / 10	0 / 10	0.0%	0

No grounding failures and no false pass claims were detected across these 30 rollout(s). No lane re-read material it had already been given.

Every term above is defined in Appendix D.

08 Efficiency & telemetry

This run has no per-token price attached — a local model, a free quota, or a customer endpoint we are not billed for — so spend is reported in tokens. Tokens are what was actually measured; a dollar figure here would be invented. Lanes run through a flat-rate local CLI report no token telemetry; their cells say n/a — an absent measurement, not zero. Rollout counts here are evaluable rollouts.

MODEL	ROLLOUTS	TOKENS	TOKENS / ROLLOUT	COST	PASSES	RATE
cli:cli/claude-sonnet	10	n/a	n/a	n/a – telemetry unavailable	9 / 10	90.0%
cli:cli/codex	10	n/a	n/a	n/a – telemetry unavailable	10 / 10	100.0%
mistral:codestral-latest	10	172 890	17 289	not billed	0 / 10	0.0%

09 Findings

Every finding below is derived from the recorded data of this run — none is an opinion.
Severity: PASS a lane succeeded · ATTENTION a capability gap · LANE / INFRA
infrastructure, not capability · NOTE a property of the measurement.

F-01

PASS

cli/claude-sonnet, cli/codex passed the independent hidden tests

Severity: PASS

Evidence: Model outcomes · Stage funnel · Appendix A

19 of 30 graded rollouts completed the full loop (inspect, edit, run the visible tests, submit) and the sealed grader accepted the submission. This demonstrates the task is solvable under the published limits.

F-01 · derived from the recorded run data

F-02

ATTENTION

submitted_failed × 11 (cli/claude-sonnet, codestral-latest)

Severity: ATTENTION

Evidence: Appendix A · Appendix C

Submitted a change; the independent hidden tests still failed.

F-02 · derived from the recorded run data

10 Recommendations

One recommendation per observed failure mode, addressed to the lane it affects. Each traces back to a finding above and to the raw trace in Appendix A.

R-01

PROBLEM

Submitted a change; the independent hidden tests still failed.

AFFECTED LANES

cli/claude-sonnet, codestral-latest

WHAT WE WOULD LOOK AT FIRST

The model completes the loop and produces a plausible near-miss — the named failing tests are the exact behavioral cases distinguishing it from the reference fix.

11 Verification

RUN ID

live-20260902T001847Z

EXAM FINGERPRINT

aa41977120e6ecdddb5c4fc603e8ab23cb272726a86750119969ff4dfc90e297

The contract digest this run was executed against.

Tamper-evidence

Three digests, and how to check each one. If any of them does not reproduce, this page is not the page that was generated for this run.

THIS REPORT

cb211eceb1b757e9c2d748a607cd86e966d7d915d8f5f44cc390033555671e3f

sha256 over the report's UTF-8 bytes with every occurrence of the report digest itself replaced by 64 '0' characters

EVAL RECORD

02b4f83ee15251a5808a83153c9c4590936af23c58cc94d172556e9f983d1909

sha256 of fr-20260902-6616b188.json as written beside this file.

RUN EVIDENCE BUNDLE

6616b188f015f52e157bdfb257e8d0f819e1f6cb9bcd3b024b340ac0771e0fd8

The digest the run's own bundle computed over itself.

CERTIFIED EVIDENCE SEAL

not recorded

The exam's sealed evidence, established before this run.

Measurement history

This is the first recorded run on this exam. There is nothing to compare it against yet, and a single measurement is not a trend.

Verifying it, in three lines

1. vvdex-env verify-report --report fr-20260902-6616b188.html

2. or by hand: replace every occurrence of the report digest above with 64 zeros, then sha256 the file – it must equal that digest

3. shasum -a 256 fr-20260902-6616b188.json must equal the eval-record digest, and the bundle digest must equal bundleSha256 in the run's bundle file

The full artifact-by-artifact digest table, and the reproduction protocol, are in Appendix B.

Appendix A — every rollout

ROLLOUT	MODEL	SEED	OUTCOME	SUBMITTED	DEEPEST STAGE	HIDDEN TESTS	FILES	TOKENS	ELAPSED
3e6298d8	cli/claude-sonnet	0	submitted_passed	yes	Passed hidden tests	pass	1	n/a	180.1s
cb495fe6	cli/claude-sonnet	1	submitted_passed	yes	Passed hidden tests	pass	1	n/a	150.5s
e76e1c66	cli/claude-sonnet	2	submitted_passed	yes	Passed hidden tests	pass	1	n/a	133.3s
9e19f314	cli/claude-sonnet	3	submitted_passed	yes	Passed hidden tests	pass	1	n/a	147.9s
aaa748af	cli/claude-sonnet	4	submitted_passed	yes	Passed hidden tests	pass	1	n/a	225.9s
3e36cea5	cli/claude-sonnet	5	submitted_passed	yes	Passed hidden tests	pass	1	n/a	181.1s
6516049a	cli/claude-sonnet	6	submitted_failed	yes	Graded	fail	0	n/a	208.2s
8c122607	cli/claude-sonnet	7	submitted_passed	yes	Passed hidden tests	pass	1	n/a	110.6s
3161d1b1	cli/claude-sonnet	8	submitted_passed	yes	Passed hidden tests	pass	1	n/a	100.9s
6b722882	cli/claude-sonnet	9	submitted_passed	yes	Passed hidden tests	pass	1	n/a	159.4s
92c89eb1	cli/codex	0	submitted_passed	yes	Passed hidden tests	pass	1	n/a	46.0s
6d48c276	cli/codex	1	submitted_passed	yes	Passed hidden tests	pass	1	n/a	40.8s
1d32b64e	cli/codex	2	submitted_passed	yes	Passed hidden tests	pass	1	n/a	40.2s
3e53d590	cli/codex	3	submitted_passed	yes	Passed hidden tests	pass	1	n/a	40.7s
90405084	cli/codex	4	submitted_passed	yes	Passed hidden tests	pass	1	n/a	51.1s
48e359d8	cli/codex	5	submitted_passed	yes	Passed hidden tests	pass	1	n/a	36.9s
fd7561b3	cli/codex	6	submitted_passed	yes	Passed hidden tests	pass	1	n/a	39.5s
4fe31a46	cli/codex	7	submitted_passed	yes	Passed hidden tests	pass	1	n/a	43.4s
18660280	cli/codex	8	submitted_passed	yes	Passed hidden tests	pass	1	n/a	46.1s
195b8f72	cli/codex	9	submitted_passed	yes	Passed hidden tests	pass	1	n/a	39.5s
5d129402	codestral-latest	0	submitted_failed	yes	Graded	fail	1	19 250	11.5s
bd30cb85	codestral-latest	1	submitted_failed	yes	Graded	fail	1	16 813	11.2s
7cc71ab5	codestral-latest	2	submitted_failed	yes	Graded	fail	1	16 792	10.6s
d38fb2e2	codestral-latest	3	submitted_failed	yes	Graded	fail	1	16 861	14.9s
8c65b44c	codestral-latest	4	submitted_failed	yes	Graded	fail	1	16 876	11.2s
b8a625a2	codestral-latest	5	submitted_failed	yes	Graded	fail	1	16 762	10.7s
58abb85c	codestral-latest	6	submitted_failed	yes	Graded	fail	1	16 876	10.6s
94122fe4	codestral-latest	7	submitted_failed	yes	Graded	fail	1	16 759	13.5s
b0843fec	codestral-latest	8	submitted_failed	yes	Graded	fail	1	19 025	10.2s
cecbcc3b	codestral-latest	9	submitted_failed	yes	Graded	fail	1	16 876	10.2s

Failure taxonomy

CLASS	COUNT	WHAT IT MEANS
submitted_failed	11	Submitted a change; the independent hidden tests still failed.

Appendix B — evidence and artifact digests

RUN BUNDLE DIGEST

6616b188f015f52e157bdfb2

57e8d0f819e1f6cb9bcd3b02

4b340ac0771e0fd8

RUN ID

live-

20260902T001847Z

STARTED

2026-09-

02T00:18:47.313247

+00:00

FINISHED

2026-09-

02T00:57:34.65040

8+00:00

Artifact digests

ARTIFACT	SHA-256
publicSummary	059a37157968962044d98a133f9422c6831ff452419eb4b8729b856a79cad646
rollouts/18660280-d29a-4c7f-928c-0f37b96f18e3.json	7049943d0e9a63ddf108caa49598a38a76ccde4f29d1e54cb9f3098ca7aa2e37
rollouts/195b8f72-f390-486a-a369-b5047d3bf081.json	f199086fc48f883f348848c82b55fa6cdf2dd1e415001939c9176faf4ab48b84
rollouts/1d32b64e-b2fd-4a37-a04f-1230e38ad920.json	d7ad72ef466cadbccd862d52e8b0cad772ef2c9360ccb602ab332ce976bd9e44
rollouts/3161d1b1-1be0-4bcc-bb1a-11592c0931aa.json	7590bef63572befafe69f1120525598f6f03ce59869ec21a7e9d61cf2c081ced
rollouts/3e36cea5-b3a2-49a1-8227-83d3c1f6b123.json	df73c10cd29332e01062110ea33c6df2c7e592aba53fbb2b67f58e333bd1d601
rollouts/3e53d590-9fa2-40ec-9d8d-284a6663637c.json	fbaf104694620c1643a2341422af017c17122b02d6dca2599d6ea77e27d10851
rollouts/3e6298d8-5275-4323-8970-b9d0c7709aea.json	e2483eb5171491c61186b5edda6e992b7a8fd87c794a010f76557a5ae0402813
rollouts/48e359d8-3ed5-46b1-b6a1-d6aac402eef9.json	8934223b0170fdaa59fc24e801016a92df6672d4e493d17c37f96db41875ec2d
rollouts/4fe31a46-de93-4684-bc08-eda4db9ceab0.json	3058c5a8ef15ec2dea64cc04c5d4ae747c66c55de8d0b80ac7acfbcb8524554da
rollouts/58abb85c-b63a-447e-a07b-071cdcb65b4b.json	ed4c3b6913b3df2f9098a2d6e8f384cd5b6f516e0de8999923013d55ac3ae22d
rollouts/5d129402-a899-48be-b2bd-0101eeaa04c8.json	3f98b67607190f913c03acb489561a96b6ddf351c5b17ac1687ae4fc75405016
rollouts/6516049a-dc0f-4636-9df0-9f827bb6dd1b.json	6d9ba60e40f7e85a59635d2f75675caf7d330cede1bdd8146b12c3c2a4445037
rollouts/6b722882-8ce0-4eb7-8cdd-db66e413e80f.json	2960bb9ff1bd93c878aa87ad43159daca5c38a6428c2309e78c9b79d4265e408
rollouts/6d48c276-533c-469a-92d4-738279c6027d.json	0ff66e69b4d90d13bac4ce41cae996babd1bf018df5002de15d5eb0cbc33daa7
rollouts/7cc71ab5-3ba2-4a7a-8417-7e1ac353afab.json	3d7170d35dd63c1803aaa4a090f6551d23a62fb350baabe62bf396d7b317fad8
rollouts/8c122607-d35f-45fc-ba31-68bb9a296352.json	2a699f21909127d47ac15af29455549e6e62d22635fa77162bd93d485a119c42
rollouts/8c65b44c-7a1b-4da4-9194-a66dcb79f193.json	a25b7c2c5e0e41c78185730582cc68d4aaf8142ba3812c6cef418651bfc0b6ed
rollouts/90405084-4394-484b-8d91-78909aa09787.json	f5d78417af80d56b68fb13cb84831e89b3c9d76f17e3264135cf5cc8065eff4f
rollouts/92c89eb1-4f77-4165-9b73-0b53e3bff694.json	b6432a4e97b4652b098fd7b50fc4420c5a1fd530974e40ab67aeebf088351447
rollouts/94122fe4-5f7e-453e-8ef5-41250ff398d5.json	585128f0e325a96995e9acace25db8f364a264fad91f227f2071e8dd7e60255
rollouts/9e19f314-772f-4a29-8bee-8405ef9196b5.json	12754796c393b4883e68e4b21b205b19042c045174b3c75795f4f0cef0bcf45a
rollouts/aaa748af-8efc-491a-98f0-ad6565cf7ab5.json	17b22a74714f16384f42d2e92609d770535a376f3cdf490a0df82024b4caed8d
rollouts/b0843fec-30f3-491b-bfb0-9b352dea69da.json	be256e7d715aa6ef3f50fcc4a0024a1e6224d536f4f76d6b1fd165cb17b9723e
rollouts/b8a625a2-6238-476f-9951-1f7b9ed2565c.json	41e5b0047d43b493b175b0518edbbd82c8e5b99cd9160d3941abed1e4d22251a
rollouts/bd30cb85-1647-4722-9b99-b469583dce5e.json	a2fe27d8be1fd255d92b707bdb439075a56ec040ff728b634774b435b9e3a652
rollouts/cb495fe6-4d25-457d-9e5a-2d830c2d1e87.json	28e67f4869f6a390518199f6dc27bbd78033140b3bcb58f49c8db39ca98f0c83

ARTIFACT	SHA-256
rollouts/cecbcc3b-ea8a-4b61-8b95-91bc9c831eea.json	9dec1bd6a4fb9af4c240dc72a6d1d974129f9a9c699ddf371148e2c80910cc2
rollouts/d38fb2e2-3dea-4276-abe9-5c44ef17994e.json	a0f91309e405bfb8cd93ebce214a8e5ef098ebbe7994d3e0d6972aa59bcfdfb1
rollouts/e76e1c66-ea48-4852-920a-dc077f09b4ab.json	a323881e38c7b2c9edf6ea34a992d305abc4ad23c77bdb3d8c934c9e3f622332
rollouts/fd7561b3-afd5-4de2-91a4-11262c5b065e.json	aec6d21e216d0d84d43edc7343309a41265f82880cd6aecb9289bbcd894d0ffc
spec	26f9634cb3113984cabf0c0f5675880645449731e9d8ed3baeb124c35e1e60f5
summary	2a5e9b681ab1c80119fb4922de9f9d7b74a1e58fb7594ceebf3f6be296cf4ba5

Reproducing this run

Every rollout can be re-derived: the exam is frozen (fingerprint above), the grader ran in the recorded container image, and each rollout's full trace — every tool call, every result, every word — is in the evidence bundle at its rollouts/<id>.json path. A re-run uses the same job shape: the exam id, your endpoint, n, and the runtime. The record id names this run forever.

Appendix C — redacted transcripts

For each rollout: the tool calls it made, in order, and an excerpt of the diff it produced. What a tool returned is not here — only which tool, on which of the model's own paths, whether it succeeded, and how many characters came back. A request for the reference solution or the hidden tests is recorded as a request; the path it asked for is replaced by a label, because printing it would put the hidden tree's layout in a file that leaves this machine. The diff excerpt is shown because this exam's disclosure class is public-source: the change is a patch against a public repository under the licence named in the exam definition.

Disclosure class applied to this exam: public-source disclosure (ownership public_source declared by engine registry, upstream <https://github.com/martinblech/xmltodict> under MIT; family public_oss_swe): the submitted change is a patch against a public repository, so a redacted excerpt of it is published.

3e6298d8 · cli:cli/claude-sonnet · seed 0 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	xmltodict.py	ok	12 166
1	read_file	xmltodict.py	ok	11 706
2	edit_file	xmltodict.py	ok	50
3	run_tests	—	ok	183
4	submit	—	ok	85

```
--- a/xmltodict.py
+++ b/xmltodict.py
@@ -107,12 +107,28 @@
     # without attaching it back to its parent. This avoids accumulating all
     # streamed items in memory when using item_depth > 0.
     if len(self.path) == self.item_depth:
+
+         data = (None if not self.data
+                else self.cdata_separator.join(self.data))
+         item = self.item
+
+         if item is None:
+             item = (None if not self.data
+                    else self.cdata_separator.join(self.data))
+
+         should_continue = self.item_callback(self.path, item)
+         if self.strip_whitespace and data:
+             data = data.strip() or None
+         if data and self._should_force_cdata(name, data) and item is None:
+             item = self.dict_constructor()
+         if item is not None:
+             if data:
+                 self.push_data(item, self.cdata_key, data)
+                 item_value = item
+             else:
+                 item_value = data
+             # Apply postprocessor as push_data would when storing in parent
+             if self.postprocessor is not None:
+                 result = self.postprocessor(self.path, name, item_value)
+                 if result is not None:
+                     _, item_value = result
+             else:
+                 item_value = None
+
+         should_continue = self.item_callback(self.path, item_value)
+         if not should_continue:
+             raise ParsingInterrupted
+         # Reset state for the parent context without keeping a reference to
```

Showing 36 of 36 diff line(s).

cb495fe6 · cli:cli/claude-sonnet · seed 1 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	xmltodict.py	ok	12 166
1	read_file	xmltodict.py	ok	11 706
2	edit_file	xmltodict.py	ok	50
3	run_tests	—	ok	183
4	submit	—	ok	85

```
--- a/xmltodict.py
+++ b/xmltodict.py
@@ -107,10 +107,28 @@
     # without attaching it back to its parent. This avoids accumulating all
     # streamed items in memory when using item_depth > 0.
     if len(self.path) == self.item_depth:
+
+         data = (None if not self.data
+                 else self.cdata_separator.join(self.data))
+
+         item = self.item
+
+         if item is None:
+             item = (None if not self.data
+                     else self.cdata_separator.join(self.data))
+
+         if self.strip_whitespace and data:
+             data = data.strip() or None
+
+         if data and self._should_force_cdata(name, data) and item is None:
+             item = self.dict_constructor()
+
+         if item is not None:
+             if data:
+                 self.push_data(item, self.cdata_key, data)
+
+         else:
+             item = data
+
+         # Simulate push_data(parent, name, item) to get the value
+         # that would be stored in the full (non-streaming) parse
+         if self.postprocessor is not None:
+             result = self.postprocessor(self.path, name, item)
+             if result is not None:
+                 _, item = result
+             else:
+                 item = None
+
+         if self._should_force_list(name, item):
+             item = [item]
+
+
+         should_continue = self.item_callback(self.path, item)
+         if not should_continue:
```

Showing 34 of 34 diff line(s).

e76e1c66 · cli:cli/claude-sonnet · seed 2 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	xmltodict.py	ok	12 166
1	read_file	xmltodict.py	ok	11 706
2	edit_file	xmltodict.py	ok	50
3	run_tests	—	ok	183
4	submit	—	ok	85

```

--- a/xmltodict.py
+++ b/xmltodict.py
@@ -107,21 +107,41 @@
     # without attaching it back to its parent. This avoids accumulating all
     # streamed items in memory when using item_depth > 0.
     if len(self.path) == self.item_depth:
+
+         data = (None if not self.data
+                 else self.cdata_separator.join(self.data))
         item = self.item
-
-         if item is None:
-             item = (None if not self.data
-                     else self.cdata_separator.join(self.data))
-
-         should_continue = self.item_callback(self.path, item)
-         if not should_continue:
-             raise ParsingInterrupted
-         # Reset state for the parent context without keeping a reference to
-         # the emitted item.
+
+         # Restore parent state (stack.pop mirrors non-streaming branch)
         if self.stack:
             self.item, self.data = self.stack.pop()
         else:
             self.item = None
             self.data = []
+
+         # Apply the same transformations the non-streaming branch uses
+         if self.strip_whitespace and data:
+             data = data.strip() or None
+         if data and self._should_force_cdata(name, data) and item is None:
+             item = self.dict_constructor()
+         if item is not None:
+             if data:
+                 self.push_data(item, self.cdata_key, data)
+                 # Simulate what push_data(parent, name, item) would do to the value
+                 if self.postprocessor is not None:
+                     result = self.postprocessor(self.path, name, item)
+                     if result is not None:
+                         _, item = result
+             else:
+                 # text-only (or empty) path
+                 if self.postprocessor is not None:

```

Showing 40 of 53 diff line(s); truncated.

9e19f314 · cli:cli/claude-sonnet · seed 3 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	xmltodict.py	ok	12 166
1	read_file	xmltodict.py	ok	11 706
2	edit_file	xmltodict.py	ok	50
3	run_tests	—	ok	183
4	submit	—	ok	85

```

--- a/xmltodict.py
+++ b/xmltodict.py
@@ -107,12 +107,21 @@
     # without attaching it back to its parent. This avoids accumulating all
     # streamed items in memory when using item_depth > 0.
     if len(self.path) == self.item_depth:
+
+         data = (None if not self.data
+                 else self.cdata_separator.join(self.data))
         item = self.item
-
-         if item is None:
-             item = (None if not self.data
-                     else self.cdata_separator.join(self.data))
-

```

```

-
-     should_continue = self.item_callback(self.path, item)
+     if self.strip_whitespace and data:
+         data = data.strip() or None
+     if data and self._should_force_cdata(name, data) and item is None:
+         item = self.dict_constructor()
+     if item is not None:
+         if data:
+             self.push_data(item, self.cdata_key, data)
+             result = item
+         else:
+             result = data
+
+     should_continue = self.item_callback(self.path, result)
+     if not should_continue:
+         raise ParsingInterrupted
+     # Reset state for the parent context without keeping a reference to

```

Showing 29 of 29 diff line(s).

aaa748af · cli:cli/claude-sonnet · seed 4 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	xmltodict.py	ok	12 166
1	read_file	xmltodict.py	ok	11 706
2	edit_file	xmltodict.py	ok	50
3	run_tests	—	ok	183
4	submit	—	ok	85

```

--- a/xmltodict.py
+++ b/xmltodict.py
@@ -107,21 +107,38 @@
     # without attaching it back to its parent. This avoids accumulating all
     # streamed items in memory when using item_depth > 0.
     if len(self.path) == self.item_depth:
+
+         data = (None if not self.data
+                 else self.cdata_separator.join(self.data))
+         item = self.item
+
+         if item is None:
+             item = (None if not self.data
+                     else self.cdata_separator.join(self.data))
+
+         should_continue = self.item_callback(self.path, item)
+         if not should_continue:
+             raise ParsingInterrupted
+         # Reset state for the parent context without keeping a reference to
+         # the emitted item.
+         if self.stack:
+             self.item, self.data = self.stack.pop()
+         else:
+             self.item = None
+             self.data = []
+
+         # Apply the same transformations as the non-streaming path
+         if self.strip_whitespace and data:
+             data = data.strip() or None
+         if data and self._should_force_cdata(name, data) and item is None:
+             item = self.dict_constructor()
+         if item is not None:
+             if data:
+                 self.push_data(item, self.cdata_key, data)
+                 callback_item = item
+             else:
+                 callback_item = data

```

```

+         # Apply postprocessor for the element name (as push_data would do for parent)
+         if self.postprocessor is not None:
+             result = self.postprocessor(self.path, name, callback_item)
+             if result is not None:

```

Showing 40 of 50 diff line(s); truncated.

3e36cea5 · cli:cli/claude-sonnet · seed 5 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	xmltodict.py	ok	12 166
1	read_file	xmltodict.py	ok	11 706
2	edit_file	xmltodict.py	ok	50
3	run_tests	—	ok	183
4	submit	—	ok	85

```

--- a/xmltodict.py
+++ b/xmltodict.py
@@ -107,21 +107,38 @@
     # without attaching it back to its parent. This avoids accumulating all
     # streamed items in memory when using item_depth > 0.
     if len(self.path) == self.item_depth:
+
+         data = (None if not self.data
+                 else self.cdata_separator.join(self.data))
+         item = self.item
+
+         if item is None:
+             item = (None if not self.data
+                     else self.cdata_separator.join(self.data))
+
+         should_continue = self.item_callback(self.path, item)
+         if not should_continue:
+             raise ParsingInterrupted
+         # Reset state for the parent context without keeping a reference to
+         # the emitted item.
+         if self.stack:
+             self.item, self.data = self.stack.pop()
+         else:
+             self.item = None
+             self.data = []
+
+         # Apply the same processing as the non-streaming branch
+         if self.strip_whitespace and data:
+             data = data.strip() or None
+         if data and self._should_force_cdata(name, data) and item is None:
+             item = self.dict_constructor()
+         if item is not None:
+             if data:
+                 self.push_data(item, self.cdata_key, data)
+                 final_value = item
+             else:
+                 final_value = data
+
+         # Replicate what push_data(parent, name, final_value) would compute
+         # (postprocessor + force_list) without actually storing in parent
+         if self.postprocessor is not None:
+             result = self.postprocessor(self.path, name, final_value)
+             if result is None:
+                 self.path.pop()

```

Showing 40 of 50 diff line(s); truncated.

6516049a · cli:cli/claude-sonnet · seed 6 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	xmldict.py	ok	12 166
1	read_file	xmldict.py	ok	11 706
2	invalid_tool_call	(none)	failed	107
3	run_tests	—	ok	183
4	submit	—	ok	85

This rollout changed no file, so there is no diff to show.

8c122607 · cli:cli/claude-sonnet · seed 7 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	xmldict.py	ok	12 166
1	edit_file	xmldict.py	ok	50
2	run_tests	—	ok	183
3	submit	—	ok	85

```
--- a/xmldict.py
+++ b/xmldict.py
@@ -107,21 +107,37 @@
     # without attaching it back to its parent. This avoids accumulating all
     # streamed items in memory when using item_depth > 0.
     if len(self.path) == self.item_depth:
+
+         data = (None if not self.data
+                 else self.cdata_separator.join(self.data))
+         item = self.item
+         if item is None:
+             item = (None if not self.data
+                     else self.cdata_separator.join(self.data))
+
+         should_continue = self.item_callback(self.path, item)
+         if not should_continue:
+             raise ParsingInterrupted
+         # Reset state for the parent context without keeping a reference to
+         # the emitted item.
+         # Restore parent state (same ordering as non-streaming branch)
+         if self.stack:
+             self.item, self.data = self.stack.pop()
+         else:
+             self.item = None
+             self.data = []
+
+         # Apply the same processing as the non-streaming branch
+         if self.strip_whitespace and data:
+             data = data.strip() or None
+         if data and self._should_force_cdata(name, data) and item is None:
+             item = self.dict_constructor()
+         if item is not None:
+             if data:
+                 self.push_data(item, self.cdata_key, data)
+                 result = item
+             else:
+                 result = data
+         # Apply postprocessor for the element itself
+         if self.postprocessor is not None:
+             proc_result = self.postprocessor(self.path, name, result)
+             if proc_result is not None:
+                 _, result = proc_result
```

Showing 40 of 49 diff line(s); truncated.

3161d1b1 · cli:cli/claude-sonnet · seed 8 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	xmltodict.py	ok	12 166
1	read_file	xmltodict.py	ok	11 706
2	edit_file	xmltodict.py	ok	50
3	run_tests	—	ok	183
4	submit	—	ok	85

```
--- a/xmltodict.py
+++ b/xmltodict.py
@@ -107,21 +107,28 @@
     # without attaching it back to its parent. This avoids accumulating all
     # streamed items in memory when using item_depth > 0.
     if len(self.path) == self.item_depth:
+
+         data = (None if not self.data
+                else self.cdata_separator.join(self.data))
         item = self.item
-
-         if item is None:
-             item = (None if not self.data
-                    else self.cdata_separator.join(self.data))
-
-         should_continue = self.item_callback(self.path, item)
-         if not should_continue:
-             raise ParsingInterrupted
-         # Reset state for the parent context without keeping a reference to
-         # the emitted item.
         if self.stack:
             self.item, self.data = self.stack.pop()
         else:
             self.item = None
             self.data = []
+
+         # Apply the same processing as the non-streaming branch
+         if self.strip_whitespace and data:
+             data = data.strip() or None
+         if data and self._should_force_cdata(name, data) and item is None:
+             item = self.dict_constructor()
+         if item is not None:
+             if data:
+                 self.push_data(item, self.cdata_key, data)
+                 value = item
+             else:
+                 value = data
+         should_continue = self.item_callback(self.path, value)
+         if not should_continue:
+             raise ParsingInterrupted
+         self.path.pop()
+         return
     if self.stack:
```

Showing 40 of 40 diff line(s).

6b722882 · cli:cli/claude-sonnet · seed 9 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	xmltodict.py	ok	12 166
1	read_file	xmltodict.py	ok	11 706
2	invalid_tool_call	(none)	failed	107

STEP	TOOL	PATH	RESULT	CHARS
3	edit_file	xmltodict.py	ok	50
4	run_tests	—	ok	183
5	submit	—	ok	85

```

--- a/xmltodict.py
+++ b/xmltodict.py
@@ -107,21 +107,39 @@
     # without attaching it back to its parent. This avoids accumulating all
     # streamed items in memory when using item_depth > 0.
     if len(self.path) == self.item_depth:
+
+         data = (None if not self.data
+                 else self.cdata_separator.join(self.data))
         item = self.item
-         if item is None:
-             item = (None if not self.data
-                     else self.cdata_separator.join(self.data))
-
-         should_continue = self.item_callback(self.path, item)
-         if not should_continue:
-             raise ParsingInterrupted
-         # Reset state for the parent context without keeping a reference to
-         # the emitted item.
+
+         # Restore parent state
         if self.stack:
             self.item, self.data = self.stack.pop()
         else:
             self.item = None
             self.data = []

+
+         # Apply the same transformations as the non-streaming branch
+         if self.strip_whitespace and data:
+             data = data.strip() or None
+         if data and self._should_force_cdata(name, data) and item is None:
+             item = self.dict_constructor()
+         if item is not None:
+             if data:
+                 self.push_data(item, self.cdata_key, data)
+             result = item
+         else:
+             result = data
+
+         # Apply postprocessor for the element itself (mirrors push_data)
+         if self.postprocessor is not None:
+             entry = self.postprocessor(self.path, name, result)

```

Showing 40 of 51 diff line(s); truncated.

92c89eb1 · cli:cli/codex · seed 0 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	list_files	.	ok	39
1	read_file	ISSUE.md	ok	1 282
2	read_file	xmltodict.py	ok	12 166
3	edit_file	xmltodict.py	ok	50
4	run_tests	—	ok	183
5	submit	—	ok	85

```

--- a/xmltodict.py
+++ b/xmltodict.py

```

```

@@ -107,10 +107,28 @@
    # without attaching it back to its parent. This avoids accumulating all
    # streamed items in memory when using item_depth > 0.
    if len(self.path) == self.item_depth:
+       data = (None if not self.data
+       else self.cdata_separator.join(self.data))
        item = self.item
        if item is None:
-           item = (None if not self.data
-           else self.cdata_separator.join(self.data))
+       if self.strip_whitespace and data:
+           data = data.strip() or None
+       if data and self._should_force_cdata(name, data) and item is None:
+           item = self.dict_constructor()
+       if item is not None:
+           if data:
+               self.push_data(item, self.cdata_key, data)
+       else:
+           item = data
+
+       # Apply the same element-level postprocessing and force-list logic
+       # that push_data() would apply when attaching this value to its
+       # parent, but use a temporary container so streaming still drops
+       # the completed element instead of retaining it in the document.
+       container = self.push_data(None, name, item)
+       if container:
+           item = container[next(iter(container))]
+       else:
+           item = None

        should_continue = self.item_callback(self.path, item)
        if not should_continue:

```

Showing 34 of 34 diff line(s).

6d48c276 · cli:cli/codex · seed 1 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	list_files	.	ok	39
1	read_file	ISSUE.md	ok	1 282
2	read_file	xmltodict.py	ok	12 166
3	edit_file	xmltodict.py	ok	50
4	run_tests	—	ok	183
5	submit	—	ok	85

```

--- a/xmltodict.py
+++ b/xmltodict.py
@@ -107,12 +107,34 @@
    # without attaching it back to its parent. This avoids accumulating all
    # streamed items in memory when using item_depth > 0.
    if len(self.path) == self.item_depth:
+       # Finalize the streamed element exactly as the ordinary branch
+       # would, but use a temporary parent so the document is not
+       # accumulated in streaming mode.
        data = (None if not self.data
        else self.cdata_separator.join(self.data))
        item = self.item
-       if item is None:
-           item = (None if not self.data
-           else self.cdata_separator.join(self.data))
-
-       should_continue = self.item_callback(self.path, item)
+       if self.strip_whitespace and data:
+           data = data.strip() or None

```

```

+         if data and self._should_force_cdata(name, data) and item is None:
+             item = self.dict_constructor()
+         if item is not None:
+             if data:
+                 self.push_data(item, self.cdata_key, data)
+             value = item
+         else:
+             value = data
+
+         # In a full parse, the completed value is passed through
+         # push_data when it is installed under its element name. That is
+         # where the element-level postprocessor and force_list handling
+         # occur. Build only that one-entry wrapper, then discard it.
+         wrapper = self.push_data(None, name, value)
+         if wrapper is None:
+             value = None
+         else:
+             value = next(iter(wrapper.values()))
+
+         should_continue = self.item_callback(self.path, value)
+         if not should_continue:

```

Showing 40 of 42 diff line(s); truncated.

1d32b64e · cli:cli/codex · seed 2 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	list_files	.	ok	39
1	read_file	ISSUE.md	ok	1 282
2	read_file	xmltodict.py	ok	12 166
3	edit_file	xmltodict.py	ok	50
4	run_tests	—	ok	183
5	submit	—	ok	85

```

--- a/xmltodict.py
+++ b/xmltodict.py
@@ -107,10 +107,27 @@
     # without attaching it back to its parent. This avoids accumulating all
     # streamed items in memory when using item_depth > 0.
     if len(self.path) == self.item_depth:
+
+         data = (None if not self.data
+                 else self.cdata_separator.join(self.data))
+         item = self.item
+
+         if item is None:
+             item = (None if not self.data
+                     else self.cdata_separator.join(self.data))
+
+         if self.strip_whitespace and data:
+             data = data.strip() or None
+         if data and self._should_force_cdata(name, data) and item is None:
+             item = self.dict_constructor()
+         if item is not None:
+             if data:
+                 self.push_data(item, self.cdata_key, data)
+         else:
+             item = data
+
+         # Apply the same element-level postprocessing and force-list logic
+         # as attaching the element to its parent, but use a temporary
+         # container so streamed items are still dropped after the callback.
+         container = self.push_data(None, name, item)
+         if container is None:
+             item = None
+         else:
+             item = next(iter(container.values()))

```

```
should_continue = self.item_callback(self.path, item)
if not should_continue:
```

Showing 33 of 33 diff line(s).

3e53d590 · cli:cli/codex · seed 3 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	list_files	.	ok	39
1	read_file	ISSUE.md	ok	1 282
2	read_file	xmltodict.py	ok	12 166
3	edit_file	xmltodict.py	ok	50
4	run_tests	—	ok	183
5	submit	—	ok	85

```
--- a/xmltodict.py
+++ b/xmltodict.py
@@ -107,10 +107,28 @@
     # without attaching it back to its parent. This avoids accumulating all
     # streamed items in memory when using item_depth > 0.
     if len(self.path) == self.item_depth:
+        data = (None if not self.data
+                else self.cdata_separator.join(self.data))
         item = self.item
-        if item is None:
-            item = (None if not self.data
-                  else self.cdata_separator.join(self.data))
+        if self.strip_whitespace and data:
+            data = data.strip() or None
+        if data and self._should_force_cdata(name, data) and item is None:
+            item = self.dict_constructor()
+        if item is not None:
+            if data:
+                self.push_data(item, self.cdata_key, data)
+        else:
+            item = data
+
+        # Apply the same element-level postprocessing and force-list logic
+        # that push_data applies when attaching an ordinary element to its
+        # parent, but keep the temporary wrapper local so streaming does
+        # not retain the completed item.
+        wrapped_item = self.push_data(None, name, item)
+        if wrapped_item is None:
+            item = None
+        else:
+            item = next(iter(wrapped_item.values()))

     should_continue = self.item_callback(self.path, item)
     if not should_continue:
```

Showing 34 of 34 diff line(s).

90405084 · cli:cli/codex · seed 4 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	list_files	.	ok	39
1	read_file	ISSUE.md	ok	1 282
2	read_file	xmltodict.py	ok	12 166

STEP	TOOL	PATH	RESULT	CHARS
3	edit_file	xmltodict.py	ok	50
4	run_tests	—	ok	183
5	submit	—	ok	85

```

--- a/xmltodict.py
+++ b/xmltodict.py
@@ -107,21 +107,36 @@
     # without attaching it back to its parent. This avoids accumulating all
     # streamed items in memory when using item_depth > 0.
     if len(self.path) == self.item_depth:
+
+         data = (None if not self.data
+                 else self.cdata_separator.join(self.data))
         item = self.item
-         if item is None:
-             item = (None if not self.data
-                     else self.cdata_separator.join(self.data))
+
+         if self.stack:
+             self.item, self.data = self.stack.pop()
+         else:
+             self.item = None
+             self.data = []
+
+         if self.strip_whitespace and data:
+             data = data.strip() or None
+         if data and self._should_force_cdata(name, data) and item is None:
+             item = self.dict_constructor()
+         if item is not None:
+             if data:
+                 self.push_data(item, self.cdata_key, data)
+         else:
+             item = data
+
+         # Apply the same element-level postprocessing and force-list logic
+         # as push_data would when attaching this value to its parent, but
+         # use a temporary container so the streamed document is not kept.
+         emitted = self.push_data(None, name, item)
+         if emitted is None:
+             item = None
+         else:
+             item = next(iter(emitted.values()))
+
+         should_continue = self.item_callback(self.path, item)
+         if not should_continue:
+             raise ParsingInterrupted
-         # Reset state for the parent context without keeping a reference to

```

Showing 40 of 49 diff line(s); truncated.

48e359d8 · cli:cli/codex · seed 5 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	list_files	.	ok	39
1	read_file	ISSUE.md	ok	1 282
2	read_file	xmltodict.py	ok	12 166
3	edit_file	xmltodict.py	ok	50
4	run_tests	—	ok	183
5	submit	—	ok	85

```

--- a/xmltodict.py
+++ b/xmltodict.py

```

```

@@ -107,11 +107,25 @@
    # without attaching it back to its parent. This avoids accumulating all
    # streamed items in memory when using item_depth > 0.
    if len(self.path) == self.item_depth:
+       data = (None if not self.data
+       else self.cdata_separator.join(self.data))
        item = self.item
        if item is None:
            item = (None if not self.data
                    else self.cdata_separator.join(self.data))
-
+       if self.strip_whitespace and data:
+           data = data.strip() or None
+       if data and self._should_force_cdata(name, data) and item is None:
+           item = self.dict_constructor()
+       if item is not None:
+           if data:
+               self.push_data(item, self.cdata_key, data)
+               emitted = self.push_data(None, name, item)
+           else:
+               emitted = self.push_data(None, name, data)
+
+       # Build the value through the same push_data path used by ordinary
+       # parsing, so element-level postprocessing and force_list apply,
+       # but do not retain the temporary containing mapping.
+       item = (None if emitted is None
+               else next(iter(emitted.values())))
        should_continue = self.item_callback(self.path, item)
        if not should_continue:
            raise ParsingInterrupted

```

Showing 32 of 32 diff line(s).

fd7561b3 · cli:cli/codex · seed 6 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	list_files	.	ok	39
1	read_file	ISSUE.md	ok	1 282
2	read_file	xmltodict.py	ok	12 166
3	edit_file	xmltodict.py	ok	50
4	run_tests	—	ok	183
5	submit	—	ok	85

```

--- a/xmltodict.py
+++ b/xmltodict.py
@@ -107,10 +107,28 @@
    # without attaching it back to its parent. This avoids accumulating all
    # streamed items in memory when using item_depth > 0.
    if len(self.path) == self.item_depth:
+       data = (None if not self.data
+       else self.cdata_separator.join(self.data))
        item = self.item
        if item is None:
            item = (None if not self.data
                    else self.cdata_separator.join(self.data))
+       if self.strip_whitespace and data:
+           data = data.strip() or None
+       if data and self._should_force_cdata(name, data) and item is None:
+           item = self.dict_constructor()
+       if item is not None:
+           if data:
+               item = self.push_data(item, self.cdata_key, data)
+           else:
+               item = data

```

```

+
+     # Apply the same element-level postprocessing and force-list logic
+     # that push_data applies when attaching a completed element to its
+     # parent, but keep it in a temporary container so streaming does
+     # not retain the completed item.
+     container = self.push_data(None, name, item)
+     if container:
+         item = next(iter(container.values()))
+     else:
+         item = None
+
+     should_continue = self.item_callback(self.path, item)
+     if not should_continue:

```

Showing 34 of 34 diff line(s).

4fe31a46 · cli:cli/codex · seed 7 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	list_files	.	ok	39
1	read_file	ISSUE.md	ok	1 282
2	read_file	xmltodict.py	ok	12 166
3	edit_file	xmltodict.py	ok	50
4	run_tests	—	ok	183
5	submit	—	ok	85

```

--- a/xmltodict.py
+++ b/xmltodict.py
@@ -107,10 +107,24 @@
     # without attaching it back to its parent. This avoids accumulating all
     # streamed items in memory when using item_depth > 0.
     if len(self.path) == self.item_depth:
+
+         data = (None if not self.data
+                 else self.cdata_separator.join(self.data))
+         item = self.item
+
+         if item is None:
+             item = (None if not self.data
+                     else self.cdata_separator.join(self.data))
+
+         if self.strip_whitespace and data:
+             data = data.strip() or None
+         if data and self._should_force_cdata(name, data) and item is None:
+             item = self.dict_constructor()
+         if item is not None:
+             if data:
+                 self.push_data(item, self.cdata_key, data)
+         else:
+             item = data
+
+         # Apply the element-level postprocessor and force_list handling in
+         # exactly the same way as the ordinary parent-attachment path, but
+         # use a temporary container so streamed items are still discarded.
+         container = self.push_data(None, name, item)
+         item = next(iter(container.values()), None)
+
+         should_continue = self.item_callback(self.path, item)
+         if not should_continue:

```

Showing 30 of 30 diff line(s).

STEP	TOOL	PATH	RESULT	CHARS
0	list_files	.	ok	39
1	read_file	ISSUE.md	ok	1 282
2	read_file	xmltodict.py	ok	12 166
3	edit_file	xmltodict.py	ok	50
4	run_tests	—	ok	183
5	submit	—	ok	85

```
--- a/xmltodict.py
+++ b/xmltodict.py
@@ -103,27 +103,6 @@

def endElement(self, full_name):
    name = self._build_name(full_name)
    - # If we just closed an item at the streaming depth, emit it and drop it
    - # without attaching it back to its parent. This avoids accumulating all
    - # streamed items in memory when using item_depth > 0.
    - if len(self.path) == self.item_depth:
    -     item = self.item
    -     if item is None:
    -         item = (None if not self.data
    -                else self.cdata_separator.join(self.data))
    -
    -     should_continue = self.item_callback(self.path, item)
    -     if not should_continue:
    -         raise ParsingInterrupted
    -     # Reset state for the parent context without keeping a reference to
    -     # the emitted item.
    -     if self.stack:
    -         self.item, self.data = self.stack.pop()
    -     else:
    -         self.item = None
    -         self.data = []
    -         self.path.pop()
    -     return
    if self.stack:
        data = (None if not self.data
               else self.cdata_separator.join(self.data))
@@ -135,10 +114,26 @@
        item = self.dict_constructor()
        if item is not None:
            if data:
    -             self.push_data(item, self.cdata_key, data)
    +             item = self.push_data(item, self.cdata_key, data)
    +         else:
    +             item = data
    +         # Build the same value that push_data would place under this
```

Showing 40 of 60 diff line(s); truncated.

STEP	TOOL	PATH	RESULT	CHARS
0	list_files	.	ok	39
1	read_file	ISSUE.md	ok	1 282
2	read_file	xmltodict.py	ok	12 166

STEP	TOOL	PATH	RESULT	CHARS
3	edit_file	xmltodict.py	ok	50
4	run_tests	—	ok	183
5	submit	—	ok	85

```

--- a/xmltodict.py
+++ b/xmltodict.py
@@ -107,10 +107,23 @@
     # without attaching it back to its parent. This avoids accumulating all
     # streamed items in memory when using item_depth > 0.
     if len(self.path) == self.item_depth:
+
+         data = (None if not self.data
+                 else self.cdata_separator.join(self.data))
+         item = self.item
+         if item is None:
+             item = (None if not self.data
+                     else self.cdata_separator.join(self.data))
+         if self.strip_whitespace and data:
+             data = data.strip() or None
+         if data and self._should_force_cdata(name, data) and item is None:
+             item = self.dict_constructor()
+         if item is not None:
+             if data:
+                 item = self.push_data(item, self.cdata_key, data)
+                 emitted = self.push_data(None, name, item)
+             else:
+                 emitted = self.push_data(None, name, data)
+             if emitted is None:
+                 item = None
+             else:
+                 item = next(iter(emitted.values()))

        should_continue = self.item_callback(self.path, item)
        if not should_continue:

```

Showing 29 of 29 diff line(s).

5d129402 · mistral:codestral-latest · seed 0 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	ISSUE.md	ok	1 282
1	read_file	xmltodict.py	ok	12 166
2	edit_file	xmltodict.py	ok	50
3	run_tests	—	ok	183
4	submit	—	ok	85

```

--- a/xmltodict.py
+++ b/xmltodict.py
@@ -111,6 +111,17 @@
     if item is None:
         item = (None if not self.data
                 else self.cdata_separator.join(self.data))
+
+     else:
+         if self.data:
+             if self.strip_whitespace:
+                 text = self.cdata_separator.join(self.data).strip()
+             else:
+                 text = self.cdata_separator.join(self.data)
+             if self._should_force_cdata(name, text):
+                 if self.cdata_key not in item:
+                     item[self.cdata_key] = text

```

```

+         else:
+             item[self.cdata_key] = text

        should_continue = self.item_callback(self.path, item)
        if not should_continue:

```

Showing 20 of 20 diff line(s).

bd30cb85 · mistral:codestral-latest · seed 1 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	list_files	.	ok	39
1	read_file	xmltodict.py	ok	12 166
2	edit_file	xmltodict.py	ok	50
3	run_tests	—	ok	183
4	submit	—	ok	85

```

--- a/xmltodict.py
+++ b/xmltodict.py
@@ -111,6 +111,16 @@
     if item is None:
         item = (None if not self.data
                 else self.cdata_separator.join(self.data))

+
+     # Apply postprocessor and other transformations
+     if self.postprocessor:
+         item = self.postprocessor(self.path, self.path[-1][0], item)
+     if self.strip_whitespace and isinstance(item, str):
+         item = item.strip() or None
+     if self.force_cdata and isinstance(item, str):
+         item = {self.cdata_key: item}
+     if self.force_list and isinstance(item, str):
+         item = [item]

        should_continue = self.item_callback(self.path, item)
        if not should_continue:

```

Showing 19 of 19 diff line(s).

7cc71ab5 · mistral:codestral-latest · seed 2 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	list_files	.	ok	39
1	read_file	xmltodict.py	ok	12 166
2	edit_file	xmltodict.py	ok	50
3	run_tests	—	ok	183
4	submit	—	ok	85

```

--- a/xmltodict.py
+++ b/xmltodict.py
@@ -111,6 +111,17 @@
     if item is None:
         item = (None if not self.data
                 else self.cdata_separator.join(self.data))

+     else:
+         if self.data:
+             if self.strip_whitespace:
+                 text = self.cdata_separator.join(self.data).strip()

```

```

+         else:
+             text = self.cdata_separator.join(self.data)
+         if text:
+             if self._should_force_cdata(self.path[-1][0], text):
+                 item[self.cdata_key] = text
+             else:
+                 item[self.cdata_key] = text

        should_continue = self.item_callback(self.path, item)
        if not should_continue:

```

Showing 20 of 20 diff line(s).

d38fb2e2 · mistral:codestral-latest · seed 3 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	list_files	.	ok	39
1	read_file	xmltodict.py	ok	12 166
2	edit_file	xmltodict.py	ok	50
3	run_tests	—	ok	183
4	submit	—	ok	85

```

--- a/xmltodict.py
+++ b/xmltodict.py
@@ -111,6 +111,18 @@
         if item is None:
             item = (None if not self.data
                     else self.cdata_separator.join(self.data))

+
+         # Apply postprocessor and other transformations
+         if self.postprocessor:
+             item = self.postprocessor(self.path, name, item)
+         if self.strip_whitespace and isinstance(item, str):
+             item = item.strip() or None
+         if self.force_cdata and isinstance(item, str):
+             item = {self.cdata_key: item}
+         if self.force_list and isinstance(item, dict):
+             for key in self.force_list:
+                 if key in item:
+                     item[key] = [item[key]]

        should_continue = self.item_callback(self.path, item)
        if not should_continue:

```

Showing 21 of 21 diff line(s).

8c65b44c · mistral:codestral-latest · seed 4 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	list_files	.	ok	39
1	read_file	xmltodict.py	ok	12 166
2	edit_file	xmltodict.py	ok	50
3	run_tests	—	ok	183
4	submit	—	ok	85

```

--- a/xmltodict.py
+++ b/xmltodict.py
@@ -111,6 +111,18 @@

```

Showing 21 of 21 diff line(s).

Showing 19 of 19 diff line(s).

STEP	TOOL	PATH	RESULT	CHARS
0	list_files	.	ok	39
1	read_file	xmlltodict.py	ok	12 166
2	edit_file	xmlltodict.py	ok	50
3	run_tests	—	ok	183

STEP	TOOL	PATH	RESULT	CHARS
4	submit	-	ok	85

```

--- a/xmltodict.py
+++ b/xmltodict.py
@@ -111,6 +111,18 @@
     if item is None:
         item = (None if not self.data
                 else self.cdata_separator.join(self.data))

+
+
+     # Apply postprocessor and other transformations
+     if self.postprocessor:
+         item = self.postprocessor(self.path, name, item)
+     if self.strip_whitespace and isinstance(item, str):
+         item = item.strip() or None
+     if self.force_cdata and isinstance(item, str):
+         item = {self.cdata_key: item}
+     if self.force_list and isinstance(item, dict):
+         for key in item:
+             if self._should_force_list(key, item[key]):
+                 item[key] = [item[key]]

         should_continue = self.item_callback(self.path, item)
         if not should_continue:

```

Showing 21 of 21 diff line(s).

94122fe4 · mistral:codestral-latest · seed 7 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	list_files	.	ok	39
1	read_file	xmltodict.py	ok	12 166
2	edit_file	xmltodict.py	ok	50
3	run_tests	-	ok	183
4	submit	-	ok	85

```

--- a/xmltodict.py
+++ b/xmltodict.py
@@ -111,6 +111,16 @@
     if item is None:
         item = (None if not self.data
                 else self.cdata_separator.join(self.data))

+
+     else:
+         if self.data:
+             data = self.cdata_separator.join(self.data)
+             if self.strip_whitespace:
+                 data = data.strip() or None
+             if data:
+                 if self._should_force_cdata(self.path[-1][0], data):
+                     item[self.cdata_key] = data
+                 else:
+                     item[self.cdata_key] = data

         should_continue = self.item_callback(self.path, item)
         if not should_continue:

```

Showing 19 of 19 diff line(s).

b0843fec · mistral:codestral-latest · seed 8 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	ISSUE.md	ok	1 282
1	read_file	xmltodict.py	ok	12 166
2	edit_file	xmltodict.py	ok	50
3	run_tests	—	ok	183
4	submit	—	ok	85

```
--- a/xmltodict.py
+++ b/xmltodict.py
@@ -102,6 +102,36 @@
     self.data = []

    def endElement(self, full_name):
+        name = self._build_name(full_name)
+        # If we just closed an item at the streaming depth, emit it and drop it
+        # without attaching it back to its parent. This avoids accumulating all
+        # streamed items in memory when using item_depth > 0.
+        if len(self.path) == self.item_depth:
+            item = self.item
+            if item is None:
+                item = (None if not self.data
+                       else self.cdata_separator.join(self.data))
+
+            # Apply the same processing as in non-streaming mode
+            if self.strip_whitespace and item:
+                item = item.strip() or None
+            if item and self._should_force_cdata(name, item):
+                item = {self.cdata_key: item}
+            if self.postprocessor:
+                item = self.postprocessor(self.path, name, item)
+
+            should_continue = self.item_callback(self.path, item)
+            if not should_continue:
+                raise ParsingInterrupted
+            # Reset state for the parent context without keeping a reference to
+            # the emitted item.
+            if self.stack:
+                self.item, self.data = self.stack.pop()
+            else:
+                self.item = None
+                self.data = []
+            self.path.pop()
+            return
+        name = self._build_name(full_name)
+        # If we just closed an item at the streaming depth, emit it and drop it
+        # without attaching it back to its parent. This avoids accumulating all
```

Showing 39 of 39 diff line(s).

cecbcc3b · mistral:codestral-latest · seed 9 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	list_files	.	ok	39
1	read_file	xmltodict.py	ok	12 166
2	edit_file	xmltodict.py	ok	50
3	run_tests	—	ok	183
4	submit	—	ok	85

```

--- a/xmltodict.py
+++ b/xmltodict.py
@@ -111,6 +111,18 @@
     if item is None:
         item = (None if not self.data
                 else self.cdata_separator.join(self.data))

+
+    # Apply postprocessor and other transformations
+    if self.postprocessor:
+        item = self.postprocessor(self.path, self.path[-1][0], item)
+    if self.strip_whitespace and isinstance(item, str):
+        item = item.strip() or None
+    if self.force_cdata and isinstance(item, str):
+        item = {self.cdata_key: item}
+    if self.force_list and isinstance(item, dict):
+        for key in self.force_list:
+            if key in item:
+                item[key] = [item[key]]

    should_continue = self.item_callback(self.path, item)
    if not should_continue:

```

Showing 21 of 21 diff line(s).

Appendix D — metric definitions and methods

Statistics, stage funnel and efficiency in this document are derived at render time from the sealed rollout rows and their raw tool events: lane errors leave every capability denominator; a rollout's stages are booleans read off its own event sequence (inspected = a successful list or read, or a successful browser open/snapshot, edited = the workspace changed, tested = run_tests invoked, submitted = submit invoked, graded = the independent grader returned a verdict, passed = certified pass); absent telemetry is n/a. The record's own summary blocks are the values computed when it was sealed; the rows and events are the evidence.

Denominators

n counts evaluable rollouts only: attempts minus lane errors. A rollout the lane ended (API or runtime failure) is reported but never graded, so it is in no rate, interval or pass@k.

Interval method

Wilson score interval, not the normal approximation. The normal approximation collapses to a zero-width interval at 0 successes — it would print a 0/2 run as '0%, certain'. For k successes in n trials with $p = k/n$: centre = $(p + z^2/2n) / (1 + z^2/n)$, half-width = $z/(1 + z^2/n) \cdot \sqrt{(p(1-p))/n + z^2/4n^2}$, clamped to [0, 1], $z = 1.9600$.

pass@k

Unbiased estimator: $\text{pass@k} = 1 - C(n-c, k) / C(n, k)$, for $k \leq n$ only — reported only up to the number of rollouts actually run. It answers "given k independent attempts at this task, how often would at least one pass" — nothing more.

Reading a failed rollout

A failed rollout is a measurement, not a defect report. Model API failures and runtime failures are classified separately precisely because they say nothing about the model's ability on this task.

Stage definitions

Inspected

Listed or read at least one file in the repository — or, on a browser exam, opened or snapshotted the page.

Edited

Wrote at least one file that actually changed on disk.

Ran tests

Invoked the visible-test tool at least once.

Submitted

Called submit, ending the rollout on its own terms.

Graded

The independent hidden grader ran on the submission and returned a verdict.

Passed hidden tests

The hidden grader returned exit 0 and the rollout was certified.

Long-horizon terms

commandsRun

Shell commands executed inside the model's own container, against the command budget the contract declares. Only exams that advertise run_command have any.

deadEnds

Steps that moved nothing: a re-read, a repeat of the previous action, or a call to a tool this exam does not have. Counted, not judged.

hitStepLimit

Rollouts that ran out of steps without submitting. On a long task this is the number that separates 'could not finish' from 'finished wrong'.

revisits

Re-reads of a slice the rollout had already been given. On a long task, re-reading is how a model compensates for losing its own earlier context.

stepsUsed

Tool calls the rollout actually spent, out of the step limit the contract declares. A rollout that submitted early spent fewer, and that is a result, not a shortfall.

Grounding and overclaim terms

groundingFailure

A file path the model wrote in its own reasoning that does not exist in the tree it was given, or that names a directory it is not allowed to see. Detected by matching source-file paths in the model's text against the agent box on disk — not by judging the prose.

nonexistent_path

The model named a source file that is not in the repository it was given. It was reasoning about a file that does not exist.

overclaim

The model stated the tests pass or the issue is fixed, and the independent hidden tests then failed. Counted per rollout: claimed-and-wrong, or not.

reReadChurn

How many times the model read a file it had already read in the same rollout. It is a cost and a coherence signal, not a failure: re-reading is sometimes right.

requested_hidden_path

The model named a path under the reference solution, the hidden grader or the hidden test directory. Those are not in its box; the request could not be served.

Failure classes

submitted_failed

Submitted a change; the independent hidden tests still failed.

submitted_passed

Submitted a change that passed the independent hidden tests.

What this report withholds, and why

This exam's graded artefact is a change to a public upstream repository under the licence named in the exam definition, so a bounded, redacted excerpt of each rollout's own diff is published in Appendix C. Withheld everywhere, on every exam: the contents of tool results, the hidden tests, the reference solution, host paths, keys, and any line matching the redaction rule.

Rule applied to this exam: public-source disclosure (ownership public_source declared by engine registry, upstream <https://github.com/martinblech/xmltodict> under MIT; family public_oss_swe): the submitted change is a patch against a public repository, so a redacted excerpt of it is published. The rule is a property of the exam family, not a judgement made per document, and the same rule governs every edition of this report, the campaign report that embeds it as a chapter, this exam's public page and every published JSON derived from the same sealed record.

Attestation

This report was generated end to end by VVDex Forge from the sealed evidence of the evaluation run described above. Measurement values were populated from recorded artifacts, not manually entered. The digests in Verification let any holder of this file confirm that the report, the eval record and the run evidence are the ones this run produced — without contacting VVDex.

VVDex Forge · vvdexops.com	2026-09-02	vvdex-env 0.1.0
Issuer — certified evaluation environments	Date of issue	Engine · verifiers pin c521e4146026

Source: martinblech/xmltodict (MIT)
Post-seal benchmark. It does not stamp any certifying gate and did not rewrite the exam, the gold solution, the grader or the certified evidence.

— END OF REPORT · fr-20260902-6616b188 —