

Free API lanes — text and structured, ten attempts each

2 fixed tasks · 4 sealed attempt records · 124
graded attempts · 67 certified passes

CURRENT

CAMPAIGN REPORT — EXECUTIVE EDITION

2026-09-07

Evaluation scope:

openai/gpt-oss-20b
openai/gpt-oss-120b
@cf/openai/gpt-oss-20b
codestral-latest
mistral-code-latest
ministral-14b-latest
ministral-8b-latest

Prepared by:

VVDex Forge — regenerated from verified sealed records only

fc-3377c49ce724 · 4 record(s) · every value recomputable

Contents

01.	Executive summary	3
02.	Per-exam results	5
03.	Environment families	6
04.	Campaign aggregate across these two exams	7
05.	Harness status	8
06.	Cells excluded from capability	9
07.	What the outcomes mean	10
08.	What this report withholds, and why	11
09.	Evidence	12
10.	Attestation	13

01 Executive summary

Release status: CURRENT. This is the campaign the published featured index carries; its exams' public pages state these numbers.

FIXED TASKS 2	SEALED RECORDS 4	MODEL LANES 7
ATTEMPTS 136	GRADED ATTEMPTS 124	CERTIFIED PASSES 67
LANE ERRORS (NOT GRADED) 12	WITHHELD / INVALID 0	

Purpose

This report presents a controlled evaluation of 7 model lane(s) against 2 fixed certified annotation tasks, two sealed attempt records on each task. For each task, every lane received that task's frozen workspace, tool contract and limits; grading ran afterwards, in isolation, against hidden tests no lane ever saw; and every rollout carries three separate outcomes — grader, integrity, certified — of which only the certified outcome is counted. The executive edition omits record-level chapters; the full report renders each sealed attempt record as a chapter derived from the same record as its standalone report. The full report authenticates the chapter bytes, while the standalone report carries its own digest.

Annotation campaign scope. This campaign evaluates two fixed annotation tasks with two sealed attempt records on each task (4 sealed attempts in total). The sealed tasks cover structured, text inputs. Audio is N/A because no audio task appears in this campaign's sealed records. Quality measurements are reported per task and remain separate from the strict certified-pass gate.

67

certified pass(es) across 124 graded attempts (136 attempts, 12 lane error(s) not graded) in 14 cells — counted only where the grader accepted the submission AND containment was verified AND provenance is complete AND no integrity invariant or verdict-bearing harness suspicion applies.

fc-3377c49ce724 · generated from 4 verified sealed record(s) · EXECUTIVE EDITION

Key findings across the campaign

- structured-data-1 · openai/gpt-oss-20b · **PASS** — 10 certified passes / 10 graded attempts; 0 model failures; 0 excluded (0 lane errors, 0 withheld, 0 invalid).
- structured-data-1 · openai/gpt-oss-120b · **ATTENTION** — 7 certified passes / 10 graded attempts; 3 model failures; 0 excluded (0 lane errors, 0 withheld, 0 invalid).
- structured-data-1 · @cf/openai/gpt-oss-20b · **ATTENTION** — 9 certified passes / 10 graded attempts; 1 model failure; 0 excluded (0 lane errors, 0 withheld, 0 invalid).
- structured-data-1 · codestral-latest · **ATTENTION** — 9 certified passes / 10 graded attempts; 1 model failure; 0 excluded (0 lane errors, 0 withheld, 0 invalid).
- structured-data-1 · mistral-code-latest · **ATTENTION** — 9 certified passes / 10 graded attempts; 1 model failure; 0 excluded (0 lane errors, 0 withheld, 0 invalid).
- structured-data-1 · ministral-14b-latest · **ATTENTION** — 8 certified passes / 10 graded attempts; 2 model failures; 0 excluded (0 lane errors, 0 withheld, 0 invalid).
- structured-data-1 · ministral-8b-latest · **ATTENTION** — 1 certified passes / 10 graded attempts; 9 model failures; 0 excluded (0 lane errors, 0 withheld, 0 invalid).
- text-labeling-1 · openai/gpt-oss-20b · **ATTENTION** — 1 certified passes / 1 graded attempts; 0 model failures; 9 excluded (9 lane errors, 0 withheld, 0 invalid).
- text-labeling-1 · openai/gpt-oss-120b · **ATTENTION** — 3 certified passes / 3 graded attempts; 0 model failures; 3 excluded (3 lane errors, 0 withheld, 0 invalid).
- text-labeling-1 · @cf/openai/gpt-oss-20b · **PASS** — 10 certified passes / 10 graded attempts; 0 model failures; 0 excluded (0 lane errors, 0 withheld, 0 invalid).
- text-labeling-1 · codestral-latest · **ATTENTION** — 0 certified passes / 10 graded attempts; 10 model failures; 0 excluded (0 lane errors, 0 withheld, 0 invalid).
- text-labeling-1 · mistral-code-latest · **ATTENTION** — 0 certified passes / 10 graded attempts; 10 model failures; 0 excluded (0 lane errors, 0 withheld, 0 invalid).
- text-labeling-1 · ministral-14b-latest · **ATTENTION** — 0 certified passes / 10 graded attempts; 10 model failures; 0 excluded (0 lane errors, 0 withheld, 0 invalid).
- text-labeling-1 · ministral-8b-latest · **ATTENTION** — 0 certified passes / 10 graded attempts; 10 model failures; 0 excluded (0 lane errors, 0 withheld, 0 invalid).

02 Per-exam results

MODEL LANE	ANNOTATION · STRUCTURED-DATA-1 VVDEX.ANNOTATION.STRUCTURED-DATA-1	ANNOTATION · TEXT-LABELING-1 VVDEX.ANNOTATION.TEXT-LABELING-1
@cf/openai/gpt-oss-20b	9/10	10/10
openai/gpt-oss-20b	10/10	1/1 + 9 lane errors
openai/gpt-oss-120b	7/10	3/3 + 3 lane errors
codestral-latest	9/10	0/10
mistral-code-latest	9/10	0/10
ministral-14b-latest	8/10	0/10
ministral-8b-latest	1/10	0/10

Each cell is *certified passes / graded attempts* for that lane on that exam. An attempt the lane ended before grading is reported beside the cell — 1/9 + 1 lane error, never 1/10 — and is outside the denominator; withheld and invalid attempts are reported the same way. A cell whose graded attempts all passed but which also holds such an attempt is shown as *passes with attempts outside the denominator*, not as a clean sweep — the clean-sweep state is reserved for a cell where every attempt was graded and every graded attempt certified. No cell here carries a percentage: these are separate exams, not one score.

03 Environment families

This campaign covers two exam(s) drawn from one environment family(ies). The sentence in each row is the family's own description of what the exam tests; nothing is claimed here that the exam does not.

FAMILY	EXAM	WHAT IT TESTS
Multimodal annotation	<code>vvdex.annotation.structured-data-1</code>	Invented tabular records are reviewed for normalization, validity, duplicates, anomalies and relationships.
Multimodal annotation	<code>vvdex.annotation.text-labeling-1</code>	An invented message is reviewed for classification and entity spans with consistent offsets.

04 Campaign aggregate across these two exams

MODEL LANE	CERTIFIED PASSES	GRADED ATTEMPTS	LANE ERRORS	WITHHELD	INVALID
@cf/openai/gpt-oss-20b	19	20	0	0	0
openai/gpt-oss-20b	11	11	9	0	0
openai/gpt-oss-120b	10	13	3	0	0
codestral-latest	9	20	0	0	0
mistral-code-latest	9	20	0	0	0
ministral-14b-latest	8	20	0	0	0
ministral-8b-latest	1	20	0	0	0

Raw counts first. *Graded attempts* = certified passes + model fails; lane errors, withheld and invalid cells are outside capability arithmetic in both directions. These two fixed annotation tasks measure different work, so no combined rate or confidence interval is reported. The totals are accounting counts only; the per-task results above carry the quality claims.

05 Harness status

harness SUSPECT. 1 of 2 exams include harness-suspect records (text-labeling-1). Each fired rule is classified below; a verdict-bearing rule the record cannot clear from its own evidence withholds the affected cells. **0 cell(s) withheld** on that basis.

Here R1 records provider/API interruption as a diagnostic; it does not assert a grader defect or turn an ungraded lane error into a model result.

text-labeling-1 · R1 · diagnostic

an API failure is a per-rollout lane event, already recorded on that rollout as lane_error; it cannot alter another rollout's grading. It lowers coverage, not verdicts.

06 Cells excluded from capability

An attempt is excluded from capability only when the lane, not the model, ended it (lane error), when the run cannot prove its own conditions (withheld), or when an integrity invariant fired (INVALID). A submission the hidden grader failed is a model result and is counted, never listed here.

EXAM	LANE	SEED	EXCLUSION	RECORDED REASON
text-labeling-1	openai/gpt-oss-120b	3	lane error	HTTP 429: {"error":{"message":"Rate limit reached for model `openai/gpt-oss-120b` in organization [redacted] service tier `on_demand` on tokens per day (TPD): ...
text-labeling-1	openai/gpt-oss-120b	4	lane error	HTTP 429: {"error":{"message":"Rate limit reached for model `openai/gpt-oss-120b` in organization [redacted] service tier `on_demand` on tokens per day (TPD): ...
text-labeling-1	openai/gpt-oss-120b	5	lane error	HTTP 429: {"error":{"message":"Rate limit reached for model `openai/gpt-oss-120b` in organization [redacted] service tier `on_demand` on tokens per day (TPD): ...
text-labeling-1	openai/gpt-oss-20b	0	lane error	HTTP 400: Failed to parse tool call arguments as JSON
text-labeling-1	openai/gpt-oss-20b	1	lane error	HTTP 400: Failed to parse tool call arguments as JSON
text-labeling-1	openai/gpt-oss-20b	3	lane error	HTTP 400: Failed to parse tool call arguments as JSON
text-labeling-1	openai/gpt-oss-20b	4	lane error	HTTP 400: Failed to parse tool call arguments as JSON
text-labeling-1	openai/gpt-oss-20b	5	lane error	HTTP 400: Failed to parse tool call arguments as JSON
text-labeling-1	openai/gpt-oss-20b	6	lane error	HTTP 400: Failed to parse tool call arguments as JSON
text-labeling-1	openai/gpt-oss-20b	7	lane error	HTTP 400: Failed to parse tool call arguments as JSON
text-labeling-1	openai/gpt-oss-20b	8	lane error	HTTP 400: Failed to parse tool call arguments as JSON
text-labeling-1	openai/gpt-oss-20b	9	lane error	HTTP 429: {"error":{"message":"Rate limit reached for model `openai/gpt-oss-20b` in organization [redacted] service tier `on_demand` on tokens per day (TPD): ...

07 What the outcomes mean

- **CERTIFIED PASS:** the grader accepted the submission, containment was verified, provenance is complete and no integrity invariant fired.
- **model fail:** the submission failed the hidden tests under verified containment; a real, attributable model result.
- **lane error:** the provider or CLI lane failed (rate limit, quota, transport) before an attributable measurement existed; not a model result.
- **withheld:** the run cannot prove its own conditions (containment unproven, provenance incomplete, or a harness suspicion the record cannot clear); neither a pass nor a model failure.
- **INVALID:** an integrity invariant fired (isolation breach, evidence mismatch); the row is not a model result.
- **harness suspect:** the run's own self-check fired; each rule is classified as verdict-bearing or diagnostic, and verdict-bearing suspicion withholds the affected cells.

08 What this report withholds, and why

What a public document may print about a rollout is decided by what the exam GRADES, not case by case. Where the graded artefact is a change to a public upstream repository, a bounded redacted excerpt of the model's own diff is published. Where the graded output is an answer, a keyed value or the state of an application, nothing derived from the submission is published — no diff, no answer value, no model claim, and not the names of the hidden assertions a rollout failed, because a grader's name can carry the value it asserts. Withheld on every exam without exception: the contents of tool results, the hidden tests, the reference solution, host paths and keys. The counts, verdicts and intervals in this report are the grader's and are unaffected by what is withheld.

EXAM	DISCLOSURE CLASS	RULE APPLIED
structured-data-1	answer_key	Applied consistently to 2 sealed records. answer-key disclosure (ownership vvdex_proprietary — no declaration — proprietary by default; family annotation): the exam is VVDex evaluation IP and the graded output is an answer or a keyed value, so no submission content, diff, model claim or hidden-assertion name is published.
text-labeling-1	answer_key	Applied consistently to 2 sealed records. answer-key disclosure (ownership vvdex_proprietary — no declaration — proprietary by default; family annotation): the exam is VVDex evaluation IP and the graded output is an answer or a keyed value, so no submission content, diff, model claim or hidden-assertion name is published.

09 Evidence

Every number above is recomputable from the sealed records below; each record verifies itself with `vvdex-env records verify`, and this document was regenerated from them with `vvdex-env records campaign`.

- `vvdex.annotation.structured-data-1` → `fr-20260907-811ca1e7` · `recordDigest` `50cbe9a820d64927b58510039c580bbbed3bd161565d085363f841eb867181fab`
- `vvdex.annotation.structured-data-1` → `fr-20260907-aa28bed9` · `recordDigest` `b5bf2458f67d7fef214fedfc66497d235bb436dfd5ca1a13080b1eb7f14e00dd`
- `vvdex.annotation.text-labeling-1` → `fr-20260907-d7a267e3` · `recordDigest` `7286bf674cfcf3482e06f76daaf1f68a1db1cbf89725833b7c3eb0f9f76e8009`
- `vvdex.annotation.text-labeling-1` → `fr-20260907-dad35536` · `recordDigest` `d98a068bbbb888bb6215a3e1a72a7d079fc043fd8b98f52d14be0a58303da22e`

10 Attestation

Certified results appear only where evidence proves them. Withheld, invalid and lane-error cells are never presented as model results, and no historical evidence was modified in producing this document. Counts are exact; below ten rollouts per cell no rate or ranking beyond the counts is asserted.

© VVDex. All rights reserved. Public materials are provided for viewing, evaluation and verification. Reuse, redistribution, derivative benchmark creation, training use, republication or commercial use requires written permission unless a specific file states otherwise. Full terms: vvdexops.com/terms/. Third-party open source reproduced in an exam keeps its own licence.