

# Certified campaign — one rollout per lane

5 sealed evaluations · 16 graded rollouts · 8  
certified passes

HISTORICAL — SUPERSEDED  
CAMPAIGN REPORT

2026-09-01

Evaluation scope:  
cli/claude-sonnet  
cli/codex  
codestral-latest  
openai/gpt-oss-120b

Prepared by:  
VVDex Forge — regenerated from verified sealed records only

fc-8626f712e26f · 5 record(s) · every value recomputable

# Contents

|     |                                            |    |
|-----|--------------------------------------------|----|
| 01. | Executive summary                          | 3  |
| 02. | Per-exam results                           | 5  |
| 03. | Environment families                       | 6  |
| 04. | Campaign aggregate across these five exams | 7  |
| 05. | Harness status                             | 8  |
| 06. | Cells excluded from capability             | 9  |
| 07. | What the outcomes mean                     | 10 |
| 08. | What this report withholds, and why        | 11 |
| 09. | Exam · martinblech-xmltodict-issue-257     | 12 |
| 10. | Exam · pytoolz-toolz-issue-602             | 28 |
| 11. | Exam · r1chardj0n3s-parse-issue-102        | 43 |
| 12. | Exam · grounded-rag-1                      | 58 |
| 13. | Exam · memory-fact-update-1                | 73 |
| 14. | Evidence                                   | 88 |
| 15. | Attestation                                | 89 |

01 Executive summary

**Release status: HISTORICAL — superseded.** A newer campaign carries these exams. This report is published so the earlier measurement stays checkable; its exams are NOT on the public featured index, and the current numbers are in the campaign the index names.

|                       |                       |                               |                         |
|-----------------------|-----------------------|-------------------------------|-------------------------|
| EXAMS<br>5            | SEALED RECORDS<br>5   | MODEL LANES<br>4              | ATTEMPTS<br>20          |
| GRADED ROLLOUTS<br>16 | CERTIFIED PASSES<br>8 | LANE ERRORS (NOT GRADED)<br>4 | WITHHELD / INVALID<br>0 |

Purpose

This report presents a controlled evaluation of 4 model lane(s) against 5 fixed certified exams, one rollout per lane. For each exam, every lane received that exam's frozen workspace, tool contract and limits; grading ran afterwards, in isolation, against hidden tests no lane ever saw; and every rollout carries three separate outcomes — grader, integrity, certified — of which only the certified outcome is counted. Each sealed evaluation record is rendered below as a chapter derived from the same record as its standalone report: the chapter's bytes are authenticated by this campaign document, the standalone report by its own digest.

Outcome

8

certified pass(es) across 16 graded rollouts (20 attempts, 4 lane error(s) not graded) in 20 cells — counted only where the grader accepted the submission AND containment was verified AND provenance is complete AND no integrity invariant or verdict-bearing harness suspicion applies.

fc-8626f712e26f · generated from 5 verified sealed record(s)

Key findings across the campaign

- martinblech-xmltodict-issue-257 · cli/claude-sonnet · **PASS** — 1 certified passes / 1 graded rollouts; 0 model failures; 0 excluded (0 lane errors, 0 withheld, 0 invalid).
- martinblech-xmltodict-issue-257 · cli/codex · **PASS** — 1 certified passes / 1 graded rollouts; 0 model failures; 0 excluded (0 lane errors, 0 withheld, 0 invalid).
- martinblech-xmltodict-issue-257 · codestral-latest · **ATTENTION** — 0 certified passes / 1 graded rollouts; 1 model failure; 0 excluded (0 lane errors, 0 withheld, 0 invalid).
- martinblech-xmltodict-issue-257 · openai/gpt-oss-120b · **NO RESULT** — 0 certified passes / 0 graded rollouts; 0 model failures; 1 excluded (1 lane error, 0 withheld, 0 invalid).
- pytoolz-toolz-issue-602 · cli/claude-sonnet · **PASS** — 1 certified passes / 1 graded rollouts; 0 model failures; 0 excluded (0 lane errors, 0 withheld, 0 invalid).
- pytoolz-toolz-issue-602 · cli/codex · **ATTENTION** — 0 certified passes / 1 graded rollouts; 1 model failure; 0 excluded (0 lane errors, 0 withheld, 0 invalid).
- pytoolz-toolz-issue-602 · codestral-latest · **ATTENTION** — 0 certified passes / 1 graded rollouts; 1 model failure; 0 excluded (0 lane errors, 0 withheld, 0 invalid).
- pytoolz-toolz-issue-602 · openai/gpt-oss-120b · **NO RESULT** — 0 certified passes / 0 graded rollouts; 0 model failures; 1 excluded (1 lane error, 0 withheld, 0 invalid).
- r1chardj0n3s-parse-issue-102 · cli/claude-sonnet · **PASS** — 1 certified passes / 1 graded rollouts; 0 model failures; 0 excluded (0 lane errors, 0 withheld, 0 invalid).

- r1chardj0n3s-parse-issue-102 · cli/codex · **ATTENTION** — 0 certified passes / 1 graded rollouts; 1 model failure; 0 excluded (0 lane errors, 0 withheld, 0 invalid).
- r1chardj0n3s-parse-issue-102 · codestral-latest · **ATTENTION** — 0 certified passes / 1 graded rollouts; 1 model failure; 0 excluded (0 lane errors, 0 withheld, 0 invalid).
- r1chardj0n3s-parse-issue-102 · openai/gpt-oss-120b · **NO RESULT** — 0 certified passes / 0 graded rollouts; 0 model failures; 1 excluded (1 lane error, 0 withheld, 0 invalid).
- grounded-rag-1 · cli/claude-sonnet · **ATTENTION** — 0 certified passes / 1 graded rollouts; 1 model failure; 0 excluded (0 lane errors, 0 withheld, 0 invalid).
- grounded-rag-1 · cli/codex · **PASS** — 1 certified passes / 1 graded rollouts; 0 model failures; 0 excluded (0 lane errors, 0 withheld, 0 invalid).
- grounded-rag-1 · codestral-latest · **ATTENTION** — 0 certified passes / 1 graded rollouts; 1 model failure; 0 excluded (0 lane errors, 0 withheld, 0 invalid).
- grounded-rag-1 · openai/gpt-oss-120b · **NO RESULT** — 0 certified passes / 0 graded rollouts; 0 model failures; 1 excluded (1 lane error, 0 withheld, 0 invalid).
- memory-fact-update-1 · cli/claude-sonnet · **PASS** — 1 certified passes / 1 graded rollouts; 0 model failures; 0 excluded (0 lane errors, 0 withheld, 0 invalid).
- memory-fact-update-1 · cli/codex · **PASS** — 1 certified passes / 1 graded rollouts; 0 model failures; 0 excluded (0 lane errors, 0 withheld, 0 invalid).
- memory-fact-update-1 · codestral-latest · **ATTENTION** — 0 certified passes / 1 graded rollouts; 1 model failure; 0 excluded (0 lane errors, 0 withheld, 0 invalid).
- memory-fact-update-1 · openai/gpt-oss-120b · **PASS** — 1 certified passes / 1 graded rollouts; 0 model failures; 0 excluded (0 lane errors, 0 withheld, 0 invalid).

02 Per-exam results

| MODEL LANE          | SWE · XMLTODICT-257<br>PUBLIC.SWE.MARTIN<br>BLECH-XMLTODICT-I<br>SSUE-257 | SWE · TOOLZ-602<br>PUBLIC.SWE.PYTO<br>OLZ-T00LZ-ISSUE<br>-602 | SWE · PARSE-102<br>PUBLIC.SWE.R1CHA<br>RDJ0N3S-PARSE-IS<br>SUE-102 | RAG<br>VVDEX.KNOWLEDG<br>E.GROUNDED-RAG<br>-1 | MEMORY<br>VVDEX.KNOWL<br>EDGE.MEMORY<br>-FACT-UPDAT<br>E-1 |
|---------------------|---------------------------------------------------------------------------|---------------------------------------------------------------|--------------------------------------------------------------------|-----------------------------------------------|------------------------------------------------------------|
| cli/claude-sonnet   | 1/1                                                                       | 1/1                                                           | 1/1                                                                | 0/1                                           | 1/1                                                        |
| cli/codex           | 1/1                                                                       | 0/1                                                           | 0/1                                                                | 1/1                                           | 1/1                                                        |
| openai/gpt-oss-120b | 0/0 + 1 lane error                                                        | 0/0 + 1 lane error                                            | 0/0 + 1 lane error                                                 | 0/0 + 1 lane error                            | 1/1                                                        |
| codestral-latest    | 0/1                                                                       | 0/1                                                           | 0/1                                                                | 0/1                                           | 0/1                                                        |

Each cell is *certified passes / graded rollouts* for that lane on that exam. An attempt the lane ended before grading is reported beside the cell — 1/9 + 1 lane error, never 1/10 — and is outside the denominator; withheld and invalid attempts are reported the same way. No cell here carries a percentage: these are separate exams, not one score.

03 Environment families

This campaign covers five exam(s) drawn from three environment family(ies). The sentence in each row is the family's own description of what the exam tests; nothing is claimed here that the exam does not.

| FAMILY                           | EXAM                                       | WHAT IT TESTS                                                                                                                                                                                                                                       |
|----------------------------------|--------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Software repair (public OSS SWE) | public.swe.martinblech-xmltodict-issue-257 | A real software-engineering task taken from a public open-source repository: the model is given the codebase and a written issue, and must change the code so that the behaviour described in the issue is fixed.                                   |
| Software repair (public OSS SWE) | public.swe.pytoolz-toolz-issue-602         | A real software-engineering task taken from a public open-source repository: the model is given the codebase and a written issue, and must change the code so that the behaviour described in the issue is fixed.                                   |
| Software repair (public OSS SWE) | public.swe.r1chardj0n3s-parse-issue-102    | A real software-engineering task taken from a public open-source repository: the model is given the codebase and a written issue, and must change the code so that the behaviour described in the issue is fixed.                                   |
| Grounded retrieval (RAG)         | vvdex.knowledge.grounded-rag-1             | Grounded question answering over a small document collection.                                                                                                                                                                                       |
| Cross-session memory update      | vvdex.knowledge.memory-fact-update-1       | A state-and-memory task across sessions: a fact was recorded earlier, an authoritative source has since changed it, and the model must answer from the newer authoritative roster rather than the stale carried memory, and acknowledge the change. |

04 Campaign aggregate across these five exams

| MODEL LANE          | CERTIFIED PASSES | GRADED ROLLOUTS | LANE ERRORS | WITHHELD | INVALID | CERTIFIED RATE (95% CI) |
|---------------------|------------------|-----------------|-------------|----------|---------|-------------------------|
| cli/claude-sonnet   | 4                | 5               | 0           | 0        | 0       | —                       |
| cli/codex           | 3                | 5               | 0           | 0        | 0       | —                       |
| openai/gpt-oss-120b | 1                | 1               | 4           | 0        | 0       | —                       |
| codestral-latest    | 0                | 5               | 0           | 0        | 0       | —                       |

Raw counts first. *Graded rollouts* = certified passes + model fails; lane errors, withheld and invalid cells are outside capability arithmetic in both directions. Denominators differ between lanes, so no rate is shown and none should be inferred — a lane with one countable attempt has one data point, not a percentage.

## 05 Harness status

---

No record in this campaign fired a harness self-suspicion rule. Every countable cell is a grader outcome under verified containment.



06 Cells excluded from capability

An attempt is excluded from capability only when the lane, not the model, ended it (lane error), when the run cannot prove its own conditions (withheld), or when an integrity invariant fired (INVALID). A submission the hidden grader failed is a model result and is counted, never listed here.

| EXAM                            | LANE                | SEED | EXCLUSION  | RECORDED REASON                                                                                                                                                    |
|---------------------------------|---------------------|------|------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| martinblech-xmltodict-issue-257 | openai/gpt-oss-120b | 0    | lane error | HTTP 429: {"error":{"message":"Rate limit reached for model `openai/gpt-oss-120b` in organization [redacted] service tier `on_demand` on tokens per day (TPD): ... |
| pytoolz-toolz-issue-602         | openai/gpt-oss-120b | 0    | lane error | HTTP 429: {"error":{"message":"Rate limit reached for model `openai/gpt-oss-120b` in organization [redacted] service tier `on_demand` on tokens per day (TPD): ... |
| r1chardj0n3s-parse-issue-102    | openai/gpt-oss-120b | 0    | lane error | HTTP 429: {"error":{"message":"Rate limit reached for model `openai/gpt-oss-120b` in organization [redacted] service tier `on_demand` on tokens per day (TPD): ... |
| grounded-rag-1                  | openai/gpt-oss-120b | 0    | lane error | HTTP 429: {"error":{"message":"Rate limit reached for model `openai/gpt-oss-120b` in organization [redacted] service tier `on_demand` on tokens per day (TPD): ... |

## 07 What the outcomes mean

- **CERTIFIED PASS:** the grader accepted the submission, containment was verified, provenance is complete and no integrity invariant fired.
- **model fail:** the submission failed the hidden tests under verified containment; a real, attributable model result.
- **lane error:** the provider or CLI lane failed (rate limit, quota, transport) before an attributable measurement existed; not a model result.
- **withheld:** the run cannot prove its own conditions (containment unproven, provenance incomplete, or a harness suspicion the record cannot clear); neither a pass nor a model failure.
- **INVALID:** an integrity invariant fired (isolation breach, evidence mismatch); the row is not a model result.
- **harness suspect:** the run's own self-check fired; each rule is classified as verdict-bearing or diagnostic, and verdict-bearing suspicion withholds the affected cells.

08 What this report withholds, and why

What a public document may print about a rollout is decided by what the exam GRADES, not case by case. Where the graded artefact is a change to a public upstream repository, a bounded redacted excerpt of the model's own diff is published. Where the graded output is an answer, a keyed value or the state of an application, nothing derived from the submission is published — no diff, no answer value, no model claim, and not the names of the hidden assertions a rollout failed, because a grader's name can carry the value it asserts. Withheld on every exam without exception: the contents of tool results, the hidden tests, the reference solution, host paths and keys. The counts, verdicts and intervals in this report are the grader's and are unaffected by what is withheld.

| EXAM                          | DISCLOSURE CLASS | RULE APPLIED                                                                                                                                                                                                                                                                                                                            |
|-------------------------------|------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| martinblech-xmldict-issue-257 | public_source    | Applied consistently to 1 sealed records. public-source disclosure (ownership public_source declared by engine registry, upstream https://github.com/martinblech/xmldict under MIT; family public_oss_swe): the submitted change is a patch against a public repository, so a redacted excerpt of it is published.                      |
| pytoolz-toolz-issue-602       | answer_key       | Applied consistently to 1 sealed records. answer-key disclosure (ownership vvdex_proprietary — no declaration — proprietary by default; family public_oss_swe): the exam is VVDex evaluation IP and the graded output is an answer or a keyed value, so no submission content, diff, model claim or hidden-assertion name is published. |
| r1chardj0n3s-parse-issue-102  | answer_key       | Applied consistently to 1 sealed records. answer-key disclosure (ownership vvdex_proprietary — no declaration — proprietary by default; family public_oss_swe): the exam is VVDex evaluation IP and the graded output is an answer or a keyed value, so no submission content, diff, model claim or hidden-assertion name is published. |
| grounded-rag-1                | answer_key       | Applied consistently to 1 sealed records. answer-key disclosure (ownership vvdex_proprietary — engine registry; family rag_grounding): the exam is VVDex evaluation IP and the graded output is an answer or a keyed value, so no submission content, diff, model claim or hidden-assertion name is published.                          |
| memory-fact-update-1          | answer_key       | Applied consistently to 1 sealed records. answer-key disclosure (ownership vvdex_proprietary — engine registry; family long_term_memory): the exam is VVDex evaluation IP and the graded output is an answer or a keyed value, so no submission content, diff, model claim or hidden-assertion name is published.                       |

Exam revisions retired for future evaluation

- vvdex.knowledge.grounded-rag-1 0.1.0 — Retired for future evaluation on 2026-09-02: post-run public disclosure of answer-bearing submission content. Historical sealed results on this revision are unaffected: the disclosure happened after those runs, so it could not have influenced them.
- vvdex.knowledge.memory-fact-update-1 0.1.0 — Retired for future evaluation on 2026-09-02: post-run public disclosure of answer-bearing submission content. Historical sealed results on this revision are unaffected: the disclosure happened after those runs, so it could not have influenced them.

09 Exam · martinblech-xmltodict-issue-257

public.swe.martinblech-xmltodict-issue-257 · record fr-20260901-a27316a3 · record digest 96f8b631c6d65f03... · harness c  
lean · containment clean

Executive summary

|                                       |                                                                          |                                                                                      |                                                                           |
|---------------------------------------|--------------------------------------------------------------------------|--------------------------------------------------------------------------------------|---------------------------------------------------------------------------|
| <div>MODELS TESTED</div> <div>4</div> | <div>CERTIFIED PASSES</div> <div>2</div> <div>of 3 graded rollouts</div> | <div>SUBMITTED</div> <div>3 of 3</div> <div>graded rollouts that called submit</div> | <div>NOT GRADED — LANE ERRORS</div> <div>1</div> <div>of 4 attempts</div> |
|---------------------------------------|--------------------------------------------------------------------------|--------------------------------------------------------------------------------------|---------------------------------------------------------------------------|

Purpose

This report presents a controlled evaluation of 4 model lane(s) against the certified exam pu  
blic.swe.martinblech-xmltodict-issue-257 (v0.1.0, family public\_oss\_swe). Every  
lane received the identical frozen workspace, tool contract and limits; grading ran  
afterwards, in isolation, against hidden tests no lane ever saw. The purpose is to state, with  
reproducible evidence, what each lane did and where it stopped.

Outcome

2 of 3

graded rollouts passed the independent hidden tests and were certified — 1 of 4 attempts not graded (lane errors)

2 of 3 graded rollouts passed. cli/claude-sonnet, cli/codex completed the full loop (inspect,  
edit, run the visible tests, submit) and the sealed grader accepted the submission. **Low-N:**  
4 rollout(s) is below the 10-rollout floor for reading a rate as a rate — read the per-rollout  
outcomes, not the percentages. The most common way to fail was model\_api\_failure —  
the model endpoint failed to answer. not a verdict about the model's ability.

Key findings

F-01 cli/claude-sonnet, cli/codex passed the independent hidden tests

F-02 model\_api\_failure × 1 (openai/gpt-oss-120b)

F-03 submitted\_failed × 1 (codestral-latest)

F-04 Statistical floor — 4 rollout(s) is below the 10-rollout rate floor

Each finding is stated in full, with evidence, in the Findings section; recommendations for  
acting on them follow it. Elapsed wall clock for the whole engagement: 1398.7s.

Exam definition

**In scope.** The ability of each configured model lane to complete this one certified task end to end — inspect the workspace, produce a correct edit, verify it with the visible tests, and submit — within 40 tool steps and 9720s of wall clock, at 1 rollout(s) per lane. Behavioural measurements along the way (tool discipline, grounding, overclaim) are in scope because they are read from the recorded trace, not judged by another model.

**Out of scope.** General model quality, performance on other task families, prompt-tuning on the lane's behalf, and any comparison this run's sample size cannot support. A lane's API failure is reported as infrastructure, never as model capability.

**The instrument.** Exam `public.swe.martinblech-xmltodict-issue-257` was certified before this run: its reference solution passes, an empty submission fails, its grader distinguishes near-misses, and its sealed evidence is unchanged by this evaluation (see Certification evidence).

What was measured

|                                            |         |                |                 |
|--------------------------------------------|---------|----------------|-----------------|
| EXAM                                       | VERSION | FAMILY         | ROLLOUTS        |
| public.swe.martinblech-xmltodict-issue-257 | 0.1.0   | public_oss_swe | 4 (1 per model) |

Where the defect came from

**Source:** `martinblech/xmltodict (MIT)` · issue  
<https://github.com/martinblech/xmltodict/issues/257> · upstream fix  
<https://github.com/martinblech/xmltodict/pull/421>

The defect and its upstream fix are public and unmodified. The exam around them — the frozen parent tree, the independently written hidden suite, the control probes and the isolation — is VVDex's, and it is what this report measures.

Certification evidence

4 of 5 certification pillars are recorded in this package, loaded from the receipts sealed into the exam and matching this run's exam fingerprint. What is recorded is shown with its real value; what is not is named below and is not claimed.

FROZEN BASELINE

FAIL

expected — an empty submission over 2 run(s) must not pass

→

GOLD / REFERENCE

PASS

expected — the reference solution over 2 run(s) must pass

GRADER CONTROLS

7 / 7

The grader accepted the reference and rejected every near-miss.

ATTACK PROBES

10 / 10 blocked

0 trivial exploit(s) found.

RUNTIME

sha256:679b36225a1f83238efba8c71b9cde3c24caaeacc9b682ae92819cb55eec328f

network policy: deny

Certification metadata is not included in this report package for:

- sealed evidence bundle

Benchmark results remain valid as run evidence; the pillar(s) above are not substantiated by this document and this report does not claim them.

What the package does carry: the exam fingerprint below is the contract digest this run was executed against — a holder of the certification bundle for this fingerprint can join the two independently.

EXAM FINGERPRINT

aa41977120e6ecddb5c4fc603e8ab23cb272726a86750119969ff4dfc90e297

The contract digest this run was executed against.

Evaluation configuration

|                |                                                                     |
|----------------|---------------------------------------------------------------------|
| REPORT         | fr-20260901-a27316a3                                                |
| ISSUED         | 2026-09-01                                                          |
| SUBJECT EXAM   | public.swe.martinblech-xmltodict-issue-257 · v0.1.0                 |
| FAMILY         | public_oss_swe                                                      |
| PREPARED BY    | VVDex Forge engine vvdex-env 0.1.0                                  |
| CLASSIFICATION | Independent model evaluation report — independently verifiable      |
| CERTIFICATION  | Exam certification: partially recorded (see Certification evidence) |
| STATUS         | FINAL · report sealed against its own digest                        |

Timeline

| WHEN (UTC)          | EVENT                                         |
|---------------------|-----------------------------------------------|
| 2026-09-01 22:35:32 | Evaluation run started                        |
| 2026-09-01 22:58:51 | Evaluation run finished                       |
| 2026-09-01 22:58:51 | Report generated from the sealed run evidence |

Models under test

| MODEL                                  | PROVIDER | ENDPOINT         | BILLING                                            |
|----------------------------------------|----------|------------------|----------------------------------------------------|
| Claude Code CLI (Sonnet, subscription) | cli      | local runtime    | operator subscription, flat — not billed per token |
| Codex CLI (GPT, subscription)          | cli      | local runtime    | operator subscription, flat — not billed per token |
| codestral-latest                       | mistral  | VVDex-configured | VVDex free-quota                                   |
| openai/gpt-oss-120b                    | groq     | VVDex-configured | VVDex free-quota                                   |

A customer endpoint is recorded by origin only — never its port, its path, or its key.

Limits

|                    |                      |                           |                                                                                  |
|--------------------|----------------------|---------------------------|----------------------------------------------------------------------------------|
| RUNTIME<br>docker  | STEP LIMIT<br>40     | WALL-CLOCK LIMIT<br>9720s | MAX OUTPUT TOKENS<br>2048                                                        |
| TEMPERATURE<br>0.2 | RETRY POLICY<br>none | COMMAND BUDGET<br>0       | TOOLS OFFERED<br>list_files, read_file, write_file, edit_file, run_tests, submit |

The evaluated model never saw the hidden tests or the reference solution, and could not change its result. Grading ran in a separate step, after submission, on a container with no network.

Model outcomes

**Denominators.** Attempts requested: 4. Graded (evaluable) rollouts: 3. Not graded — lane errors: 1. A lane error is a rollout the API or the runtime ended before the hidden grader ran: it is listed, excluded from every rate and pass@k, and never silently dropped. Passes are certified passes: hidden grader exit 0 with integrity verified.

| MODEL                                                       | ATTEMPTS | GRADED | PASSED     | MODEL FAILS | LANE ERRORS | SUBMITTED | MEAN TIME | TYPICAL DEEPEST STAGE |
|-------------------------------------------------------------|----------|--------|------------|-------------|-------------|-----------|-----------|-----------------------|
| Claude Code CLI (Sonnet, subscription)<br>cli/claude-sonnet | 1        | 1      | 1 / 1      | 0           | 0           | 1 / 1     | 110.2s    | Passed hidden tests   |
| Codex CLI (GPT, subscription)<br>cli/codex                  | 1        | 1      | 1 / 1      | 0           | 0           | 1 / 1     | 40.1s     | Passed hidden tests   |
| codestral-latest<br>codestral-latest                        | 1        | 1      | 0 / 1      | 1           | 0           | 1 / 1     | 12.5s     | Graded                |
| openai/gpt-oss-120b<br>openai/gpt-oss-120b                  | 1        | 0      | not graded | 0           | 1           | —         | —         | nothing               |

Every rollout, with its own deepest stage and grade, is in Appendix A.

Success rate, confidence and pass@k

Every rate on this page is a measurement of a finite number of rollouts, so every rate carries the interval that measurement actually supports. n is the evaluable count — attempts minus rollouts the lane ended before grading; those are listed under "not graded" and sit in no rate, interval or pass@k.

| MODEL               | ATTEMPTS | N | NOT GRADED | PASSES | RATE         | 95% INTERVAL | SEEDS | PASS@1 |
|---------------------|----------|---|------------|--------|--------------|--------------|-------|--------|
| cli/claude-sonnet   | 1        | 1 | 0          | 1      | 100.0%       | [21%, 100%]  | 0     | 100.0% |
| cli/codex           | 1        | 1 | 0          | 1      | 100.0%       | [21%, 100%]  | 0     | 100.0% |
| codestral-latest    | 1        | 1 | 0          | 0      | 0.0%         | [0%, 79%]    | 0     | 0.0%   |
| openai/gpt-oss-120b | 1        | 0 | 1          | 0      | not recorded | [0%, 100%]   | 0     | —      |

Interval and pass@k methods are stated in full in Appendix D. pass@k is shown for k = 1; the sealed record carries every k up to n.

Model comparison

| MODEL               | PASS@1       | 95% INTERVAL | TOKENS | COST       | MEAN ROLLOUT | OVERCLAIM RATE | TOP FAILURE CLASS |
|---------------------|--------------|--------------|--------|------------|--------------|----------------|-------------------|
| cli/claude-sonnet   | 100.0%       | [21%, 100%]  | n/a    | n/a        | 110.2s       | 0.0%           | none              |
| cli/codex           | 100.0%       | [21%, 100%]  | n/a    | n/a        | 40.1s        | 0.0%           | none              |
| codestral-latest    | 0.0%         | [0%, 79%]    | 19 214 | not billed | 12.5s        | 0.0%           | submitted_failed  |
| openai/gpt-oss-120b | not recorded | [0%, 100%]   | n/a    | n/a        | 0ms          | 0.0%           | model_api_failure |

Tokens and cost read n/a where a lane reports no telemetry (a flat-rate CLI) — an absent measurement, not zero. pass@1 and its interval are over evaluable rollouts only.

This is not a leaderboard. Not separated (intervals overlap): cli/claude-sonnet vs cli/codex; cli/claude-sonnet vs codestral-latest; cli/claude-sonnet vs openai/gpt-oss-120b; cli/codex vs codestral-latest; cli/codex vs openai/gpt-oss-120b; codestral-latest vs openai/gpt-oss-120b. Below the 10-rollout floor for reading a rate as a rate: cli/claude-sonnet (1), cli/codex (1), codestral-latest (1), openai/gpt-oss-120b (0) — read the count, not the percentage. The primary comparison in this report is the stage funnel below — where a model stops is a finding; a percentage over this many rollouts is not.



Stage funnel

The funnel is the primary reading of a run. Where a model stops says more than a single pass percentage does, especially at low N.

| STAGE               | CLI:CLI/CLAUDE-SONNET | CLI:CLI/CODEX | MISTRAL:CODESTRAL-LATEST | GROQ:OPENAI/GPT-OSS-120B |
|---------------------|-----------------------|---------------|--------------------------|--------------------------|
| Inspected           | 1 / 1                 | 1 / 1         | 1 / 1                    | 0 / 0                    |
| Edited              | 1 / 1                 | 1 / 1         | 1 / 1                    | 0 / 0                    |
| Ran tests           | 1 / 1                 | 1 / 1         | 1 / 1                    | 0 / 0                    |
| Submitted           | 1 / 1                 | 1 / 1         | 1 / 1                    | 0 / 0                    |
| Graded              | 1 / 1                 | 1 / 1         | 1 / 1                    | 0 / 0                    |
| Passed hidden tests | 1 / 1                 | 1 / 1         | 0 / 1                    | 0 / 0                    |

Deepest stage reached

| MODEL                    | DEEPEST (ANY ATTEMPT) | TYPICAL (MOST ATTEMPTS) | HIDDEN-TEST PASSES | NOT GRADED |
|--------------------------|-----------------------|-------------------------|--------------------|------------|
| cli:cli/claude-sonnet    | Passed hidden tests   | Passed hidden tests     | 1 / 1              | 0          |
| cli:cli/codex            | Passed hidden tests   | Passed hidden tests     | 1 / 1              | 0          |
| mistral:codestral-latest | Graded                | Graded                  | 0 / 1              | 0          |
| groq:openai/gpt-oss-120b | nothing               | nothing                 | 0 / 0              | 1          |

Stage definitions are in Appendix D. The matrix counts evaluable rollouts; lane errors are shown beside it, not inside it. The stages are not strictly nested: a model can submit without editing, and the matrix shows that rather than hiding it. Each rollout's own deepest stage is in Appendix A.

Behavioural analysis

Long-horizon behaviour

Counted off the recorded trace of each rollout. A dead end is a step that moved nothing: a re-read of a slice already returned, a repeat of the previous action, or a call to a tool this exam does not have.

| MODEL                    | N | STEPS USED | RAN OUT OF STEPS | COMMANDS | CHANGED THE TREE | REFUSED (OVER BUDGET) | REVISITS | DEAD ENDS / ROLLOUT |
|--------------------------|---|------------|------------------|----------|------------------|-----------------------|----------|---------------------|
| cli:cli/claude-sonnet    | 1 | 5 / 40     | 0                | 0 / 0    | 0                | 0                     | 0        | 0                   |
| cli:cli/codex            | 1 | 6 / 40     | 0                | 0 / 0    | 0                | 0                     | 0        | 0                   |
| mistral:codestral-latest | 1 | 5 / 40     | 0                | 0 / 0    | 0                | 0                     | 0        | 0                   |
| groq:openai/gpt-oss-120b | 1 | 1 / 40     | 0                | 0 / 0    | 0                | 0                     | 0        | 0                   |

Each column is defined in Appendix D.

Hallucination & overclaim

Three measurements, each with a stated definition. None of them is a judgement of the model's prose by another model — every number below comes from comparing what the model wrote against what was on disk, or against what the independent grader returned.

| MODEL                    | N | GROUNDING FAILURES | KINDS | ROLLOUTS AFFECTED | OVERCLAIMS | OVERCLAIM RATE | RE-READS |
|--------------------------|---|--------------------|-------|-------------------|------------|----------------|----------|
| cli:cli/claude-sonnet    | 1 | 0                  | none  | 0 / 1             | 0 / 1      | 0.0%           | 0        |
| cli:cli/codex            | 1 | 0                  | none  | 0 / 1             | 0 / 1      | 0.0%           | 0        |
| mistral:codestral-latest | 1 | 0                  | none  | 0 / 1             | 0 / 1      | 0.0%           | 0        |
| groq:openai/gpt-oss-120b | 1 | 0                  | none  | 0 / 1             | 0 / 1      | 0.0%           | 0        |

No grounding failures and no false pass claims were detected across these 4 rollout(s). No lane re-read material it had already been given.

Every term above is defined in Appendix D.

Efficiency & telemetry

This run has no per-token price attached — a local model, a free quota, or a customer endpoint we are not billed for — so spend is reported in tokens. Tokens are what was actually measured; a dollar figure here would be invented. Lanes run through a flat-rate local CLI report no token telemetry; their cells say n/a — an absent measurement, not zero. Rollout counts here are evaluable rollouts.

| MODEL                    | ROLLOUTS | TOKENS | TOKENS / ROLLOUT | COST                        | PASSES | RATE   |
|--------------------------|----------|--------|------------------|-----------------------------|--------|--------|
| cli:cli/claude-sonnet    | 1        | n/a    | n/a              | n/a — telemetry unavailable | 1 / 1  | 100.0% |
| cli:cli/codex            | 1        | n/a    | n/a              | n/a — telemetry unavailable | 1 / 1  | 100.0% |
| mistral:codestral-latest | 1        | 19 214 | 19 214           | not billed                  | 0 / 1  | 0.0%   |
| groq:openai/gpt-oss-120b | 0        | n/a    | n/a              | n/a — telemetry unavailable | 0 / 0  | 0.0%   |

Findings

Every finding below is derived from the recorded data of this run — none is an opinion.  
Severity: PASS a lane succeeded · ATTENTION a capability gap · LANE / INFRA infrastructure, not capability · NOTE a property of the measurement.

F-01

PASS

cli/claude-sonnet, cli/codex passed the independent hidden tests

Severity: PASSEvidence: Model outcomes · Stage funnel · Appendix A

2 of 3 graded rollouts completed the full loop (inspect, edit, run the visible tests, submit) and the sealed grader accepted the submission. This demonstrates the task is solvable under the published limits.  
F-01 · derived from the recorded run data

F-02

LANE / INFRA

model\_api\_failure × 1 (openai/gpt-oss-120b)

Severity: LANE / INFRAEvidence: Appendix A · Appendix C

The model endpoint failed to answer. Not a verdict about the model's ability. These rollouts are excluded from any capability reading.  
F-02 · derived from the recorded run data

F-03

ATTENTION

submitted\_failed × 1 (codestral-latest)

Severity: ATTENTIONEvidence: Appendix A · Appendix C

Submitted a change; the independent hidden tests still failed.  
F-03 · derived from the recorded run data

F-04

NOTE

Statistical floor — 4 rollout(s) is below the 10-rollout rate floor

Severity: NOTEEvidence: Appendix D

The counts are exact; the percentages are anecdotes with decimal points. Per-rollout outcomes, not rates, are the reliable reading of this run.  
F-04 · derived from the recorded run data

Recommendations

One recommendation per observed failure mode, addressed to the lane it affects. Each traces back to a finding above and to the raw trace in Appendix A.

**R-01** PROBLEM  
**The model endpoint failed to answer. Not a verdict about the model's ability.**

AFFECTED LANES  
**openai/gpt-oss-120b**

WHAT WE WOULD LOOK AT FIRST  
The lane failed, not the model; these rollouts carry no signal about capability and are labeled accordingly.

**R-02** PROBLEM  
**Submitted a change; the independent hidden tests still failed.**

AFFECTED LANES  
**codestral-latest**

WHAT WE WOULD LOOK AT FIRST  
The model completes the loop and produces a plausible near-miss — the named failing tests are the exact behavioral cases distinguishing it from the reference fix.

**R-03** For anyone reading the percentages

Commission a run at  $n \geq 10$  rollouts per lane before treating any rate here as a rate; this run's per-rollout outcomes are exact, its percentages are not yet stable.

Verification

RUN ID

**live-20260901T223532Z**

EXAM FINGERPRINT

**aa41977120e6ecddb5c4fc603e8ab23cb272726a86750119969ff4dfc90e297**

The contract digest this run was executed against.

Provenance of this chapter

This chapter was rendered by the campaign generator from the sealed evaluation record; it is not the standalone report reproduced byte for byte. Verify the standalone report with its own digest, verify the record with `vvdex-env records verify`, and verify this campaign document as a whole.

RELATIONSHIP

**campaign-rendered chapter**

Derived from the same sealed evaluation record as the standalone report; these bytes are authenticated by the campaign artifact that contains them.

SEALED RECORD DIGEST

**96f8b631c6d65f03b21e4a9a94e86c82874e1506fac38c9a093d2f8c618dd2dc**

sha256 over the record's canonical JSON — the identity every renderer works from.

SOURCE RECORD FILE

**cac23cf090b72e918199a0968244efe9871320c3a2a5a2b9863e1a0963a4c5b0**

sha256 of fr-20260901-a27316a3.json as written beside the standalone report.

SOURCE STANDALONE REPORT

**b72719ef41d6a6d15f78c00ffcbfb2c6fa0d3646faf6593d72c87a23e4574880**

sha256 of fr-20260901-a27316a3.html — independently verifiable with its own digest.

SOURCE STANDALONE PDF

**fcc2bb4c0484075ca2aa80d21f3b7661b07703df0876c45915704e5b5ba172e4**

sha256 of fr-20260901-a27316a3.pdf.

RUN EVIDENCE BUNDLE

**a27316a347e66d960235e2f5d5880c5254a556000866fcb91eb0bb85ed4af5be**

The digest the run's own bundle computed over itself.

CERTIFIED EVIDENCE SEAL

**not available in this package**

The exam's sealed evidence, established before this run.

Measurement history

This is the first recorded run on this exam. There is nothing to compare it against yet, and a single measurement is not a trend.

The full artifact-by-artifact digest table, and the reproduction protocol, are in Appendix B.

Appendix A — every rollout

| ROLLOUT  | MODEL               | SEED | OUTCOME           | SUBMITTED | DEEPEST STAGE           | HIDDEN TESTS | FILES | TOKENS | ELAPSED |
|----------|---------------------|------|-------------------|-----------|-------------------------|--------------|-------|--------|---------|
| 65ba9fe6 | cli/claude-sonnet   | 0    | submitted_passed  | yes       | Passed hidden tests     | pass         | 1     | n/a    | 110.2s  |
| 3b5cf81e | cli/codex           | 0    | submitted_passed  | yes       | Passed hidden tests     | pass         | 1     | n/a    | 40.1s   |
| 1f03a051 | codestral-latest    | 0    | submitted_failed  | yes       | Graded                  | fail         | 1     | 19 214 | 12.5s   |
| 43ea93d0 | openai/gpt-oss-120b | 0    | model_api_failure | no        | Not graded — lane error | not graded   | 0     | 1 217  | 1214.5s |

Failure taxonomy

| CLASS             | COUNT | WHAT IT MEANS                                                                 |
|-------------------|-------|-------------------------------------------------------------------------------|
| model_api_failure | 1     | The model endpoint failed to answer. Not a verdict about the model's ability. |
| submitted_failed  | 1     | Submitted a change; the independent hidden tests still failed.                |

Appendix B — evidence and artifact digests

RUN BUNDLE DIGEST

a27316a347e66d960235e2f5d5880c5254a556000866fcb91eb0bb85ed4af5be

RUN ID

live-20260901T223532Z

STARTED

2026-09-01T22:35:32.913870+00:00

FINISHED

2026-09-01T22:58:51.597668+00:00

Artifact digests

| ARTIFACT                                           | SHA-256                                                           |
|----------------------------------------------------|-------------------------------------------------------------------|
| publicSummary                                      | d8d74d56cb9e61d158de398f5b616841217b544c8385eade746f8a2783cb5904  |
| rollouts/1f03a051-d5ed-46f8-b421-3b08eff2ec50.json | feec4f7afbfb85a7a77a5105fd5afd24c305e1b771f56076460b4fcd55fa61eeb |
| rollouts/3b5cf81e-0754-435f-9d06-74d408b02016.json | b735926f1987113807c6bd342fb658cabf8830f4772fda0fac7e2112b07576ed  |
| rollouts/43ea93d0-f7b9-401b-913d-5094ad34ade8.json | 4468f2fa8ec26459d52c476cfd8f9acb4069d67601c0233cda2e808807866f45  |
| rollouts/65ba9fe6-0de3-4733-9f6a-e444ce5e84e4.json | a02851de9dec2c935ad19eb603f209279f0f63bd748076438ab84a57a85d09f5  |
| spec                                               | b32fc5d6978a10421f3e522abfe530c59b462251a2e6c58e69b2eb6ea2a84a30  |
| summary                                            | 584de4b70cfc06667ca9462857141a67d4928902d13aa3c3a60544442b20e681  |

Reproducing this run

Every rollout can be re-derived: the exam is frozen (fingerprint above), the grader ran in the recorded container image, and each rollout's full trace — every tool call, every result, every word — is in the evidence bundle at its rollouts/<id>.json path. A re-run uses the same job shape: the exam id, your endpoint, n, and the runtime. The record id names this run forever.



Appendix C — redacted transcripts

For each rollout: the tool calls it made, in order, and an excerpt of the diff it produced. What a tool returned is not here — only which tool, on which of the model's own paths, whether it succeeded, and how many characters came back. A request for the reference solution or the hidden tests is recorded as a request; the path it asked for is replaced by a label, because printing it would put the hidden tree's layout in a file that leaves this machine. The diff excerpt is shown because this exam's disclosure class is public-source: the change is a patch against a public repository under the licence named in the exam definition.

Disclosure class applied to this exam: public-source disclosure (ownership public\_source declared by engine registry, upstream <https://github.com/martinblech/xmltodict> under MIT; family public\_oss\_swe): the submitted change is a patch against a public repository, so a redacted excerpt of it is published.

65ba9fe6 · cli:cli/claude-sonnet · seed 0 · submitted\_passed

| STEP | TOOL      | PATH         | RESULT | CHARS  |
|------|-----------|--------------|--------|--------|
| 0    | read_file | xmltodict.py | ok     | 12 166 |
| 1    | read_file | xmltodict.py | ok     | 11 706 |
| 2    | edit_file | xmltodict.py | ok     | 50     |
| 3    | run_tests | —            | ok     | 183    |
| 4    | submit    | —            | ok     | 85     |

```
--- a/xmltodict.py
+++ b/xmltodict.py
@@ -107,21 +107,41 @@
     # without attaching it back to its parent. This avoids accumulating all
     # streamed items in memory when using item_depth > 0.
     if len(self.path) == self.item_depth:
+
+         data = (None if not self.data
+                else self.cdata_separator.join(self.data))
+         item = self.item
+         if item is None:
+             item = (None if not self.data
+                    else self.cdata_separator.join(self.data))
+
+         should_continue = self.item_callback(self.path, item)
+         if not should_continue:
+             raise ParsingInterrupted
+         # Reset state for the parent context without keeping a reference to
+         # the emitted item.
+         if self.stack:
+             self.item, self.data = self.stack.pop()
+         else:
+             self.item = None
+             self.data = []
+
+         # Apply the same processing as the non-streaming branch so that
+         # the value handed to the callback equals what the element would
+         # have in a full (non-streaming) parse.
+         if self.strip_whitespace and data:
+             data = data.strip() or None
+         if data and self._should_force_cdata(name, data) and item is None:
+             item = self.dict_constructor()
+         if item is not None:
+             if data:
+                 self.push_data(item, self.cdata_key, data)
+             # Determine the value that would be stored in the parent
+             # (applying postprocessor the same way push_data would).
+             if self.postprocessor is not None:
+                 result = self.postprocessor(self.path, name, item)
+                 callback_value = item if result is None else result[1]
```

```
+         else:
+             callback_value = item
```

Showing 40 of 53 diff line(s); truncated.

### 3b5cf81e · cli:cli/codex · seed 0 · submitted\_passed

| STEP | TOOL       | PATH         | RESULT | CHARS  |
|------|------------|--------------|--------|--------|
| 0    | list_files | .            | ok     | 39     |
| 1    | read_file  | ISSUE.md     | ok     | 1 282  |
| 2    | read_file  | xmltodict.py | ok     | 12 166 |
| 3    | edit_file  | xmltodict.py | ok     | 50     |
| 4    | run_tests  | —            | ok     | 183    |
| 5    | submit     | —            | ok     | 85     |

```
--- a/xmltodict.py
+++ b/xmltodict.py
@@ -107,10 +107,27 @@
     # without attaching it back to its parent. This avoids accumulating all
     # streamed items in memory when using item_depth > 0.
     if len(self.path) == self.item_depth:
+
+         data = (None if not self.data
+                 else self.cdata_separator.join(self.data))
+         item = self.item
+         if item is None:
+             item = (None if not self.data
+                     else self.cdata_separator.join(self.data))
+         if self.strip_whitespace and data:
+             data = data.strip() or None
+         if data and self._should_force_cdata(name, data) and item is None:
+             item = self.dict_constructor()
+         if item is not None:
+             if data:
+                 self.push_data(item, self.cdata_key, data)
+         else:
+             item = data
+
+         # Apply the same element-level postprocessing and force-list logic
+         # as push_data() would when attaching this value to its parent, but
+         # keep it in a temporary container so streaming does not retain it.
+         emitted = self.push_data(None, name, item)
+         if emitted is None:
+             item = None
+         else:
+             item = next(iter(emitted.values()))
+
     should_continue = self.item_callback(self.path, item)
     if not should_continue:
```

Showing 33 of 33 diff line(s).

### 1f03a051 · mistral:codestral-latest · seed 0 · submitted\_failed

| STEP | TOOL      | PATH         | RESULT | CHARS  |
|------|-----------|--------------|--------|--------|
| 0    | read_file | ISSUE.md     | ok     | 1 282  |
| 1    | read_file | xmltodict.py | ok     | 12 166 |
| 2    | edit_file | xmltodict.py | ok     | 50     |

| STEP | TOOL      | PATH | RESULT | CHARS |
|------|-----------|------|--------|-------|
| 3    | run_tests | -    | ok     | 183   |
| 4    | submit    | -    | ok     | 85    |

```
--- a/xmltodict.py
+++ b/xmltodict.py
@@ -111,6 +111,16 @@
     if item is None:
         item = (None if not self.data
                 else self.cdata_separator.join(self.data))
+
+    else:
+        if self.data:
+            data = self.cdata_separator.join(self.data)
+            if self.strip_whitespace:
+                data = data.strip() or None
+            if data:
+                if self._should_force_cdata(self.path[-1][0], data):
+                    item[self.cdata_key] = data
+                else:
+                    item[self.cdata_key] = data
+
     should_continue = self.item_callback(self.path, item)
     if not should_continue:
```

Showing 19 of 19 diff line(s).

43ea93d0 · groq:openai/gpt-oss-120b · seed 0 · model\_api\_failure

| STEP | TOOL       | PATH | RESULT | CHARS |
|------|------------|------|--------|-------|
| 0    | list_files | .    | ok     | 39    |

This rollout changed no file, so there is no diff to show.

10 Exam · pytoolz-toolz-issue-602

public.swe.pytoolz-toolz-issue-602 · record fr-20260901-88c3cdf1 · record digest f49a31336e4865c1... · harness clean · c  
ontainment clean

Executive summary

|                                       |                                                                          |                                                                                      |                                                                           |
|---------------------------------------|--------------------------------------------------------------------------|--------------------------------------------------------------------------------------|---------------------------------------------------------------------------|
| <div>MODELS TESTED</div> <div>4</div> | <div>CERTIFIED PASSES</div> <div>1</div> <div>of 3 graded rollouts</div> | <div>SUBMITTED</div> <div>3 of 3</div> <div>graded rollouts that called submit</div> | <div>NOT GRADED — LANE ERRORS</div> <div>1</div> <div>of 4 attempts</div> |
|---------------------------------------|--------------------------------------------------------------------------|--------------------------------------------------------------------------------------|---------------------------------------------------------------------------|

Purpose

This report presents a controlled evaluation of 4 model lane(s) against the certified exam pu  
blic.swe.pytoolz-toolz-issue-602 (v0.1.0, family public\_oss\_swe). Every lane received  
the identical frozen workspace, tool contract and limits; grading ran afterwards, in isolation,  
against hidden tests no lane ever saw. The purpose is to state, with reproducible evidence,  
what each lane did and where it stopped.

Outcome

1 of 3

graded rollouts passed the independent hidden tests and were certified — 1 of 4 attempts not graded (lane errors)

1 of 3 graded rollouts passed. cli/claude-sonnet completed the full loop (inspect, edit, run  
the visible tests, submit) and the sealed grader accepted the submission. **Low-N: 4**  
rollout(s) is below the 10-rollout floor for reading a rate as a rate — read the per-rollout  
outcomes, not the percentages. The most common way to fail was submitted\_failed —  
submitted a change; the independent hidden tests still failed.

Key findings

- F-01 cli/claude-sonnet passed the independent hidden tests
- F-02 submitted\_failed × 2 (cli/codex, codestral-latest)
- F-03 model\_api\_failure × 1 (openai/gpt-oss-120b)
- F-04 Statistical floor — 4 rollout(s) is below the 10-rollout rate floor

Each finding is stated in full, with evidence, in the Findings section; recommendations for  
acting on them follow it. Elapsed wall clock for the whole engagement: 1315.4s.

Exam definition

**In scope.** The ability of each configured model lane to complete this one certified task end to end — inspect the workspace, produce a correct edit, verify it with the visible tests, and submit — within 40 tool steps and 9720s of wall clock, at 1 rollout(s) per lane. Behavioural measurements along the way (tool discipline, grounding, overclaim) are in scope because they are read from the recorded trace, not judged by another model.

**Out of scope.** General model quality, performance on other task families, prompt-tuning on the lane's behalf, and any comparison this run's sample size cannot support. A lane's API failure is reported as infrastructure, never as model capability.

**The instrument.** Exam `public.swe.pytoolz-toolz-issue-602` was certified before this run: its reference solution passes, an empty submission fails, its grader distinguishes near-misses, and its sealed evidence is unchanged by this evaluation (see Certification evidence).

What was measured

|                                    |         |                |                 |
|------------------------------------|---------|----------------|-----------------|
| EXAM                               | VERSION | FAMILY         | ROLLOUTS        |
| public.swe.pytoolz-toolz-issue-602 | 0.1.0   | public_oss_swe | 4 (1 per model) |

Where the defect came from

**Source:** `pytoolz/toolz (BSD-3-Clause)` · issue <https://github.com/pytoolz/toolz/issues/602>  
· upstream fix <https://github.com/pytoolz/toolz/pull/603>

The defect and its upstream fix are public and unmodified. The exam around them — the frozen parent tree, the independently written hidden suite, the control probes and the isolation — is VVDex's, and it is what this report measures.

Certification evidence

4 of 5 certification pillars are recorded in this package, loaded from the receipts sealed into the exam and matching this run's exam fingerprint. What is recorded is shown with its real value; what is not is named below and is not claimed.

FROZEN BASELINE

FAIL

expected — an empty submission over 2 run(s) must not pass

→

GOLD / REFERENCE

PASS

expected — the reference solution over 2 run(s) must pass

GRADER CONTROLS

7 / 7

The grader accepted the reference and rejected every near-miss.

ATTACK PROBES

10 / 10 blocked

0 trivial exploit(s) found.

RUNTIME

sha256:679b36225a1f83238efba8c71b9cde3c24caaeacc9b682ae92819cb55eec328f

network policy: deny

Certification metadata is not included in this report package for:

- sealed evidence bundle

Benchmark results remain valid as run evidence; the pillar(s) above are not substantiated by this document and this report does not claim them.

What the package does carry: the exam fingerprint below is the contract digest this run was executed against — a holder of the certification bundle for this fingerprint can join the two independently.

EXAM FINGERPRINT

e8312ce98b82558f9a1c5593e7d431f409fb1024b3013b7dbd934f1398a8937f

The contract digest this run was executed against.

Evaluation configuration

|                |                                                                     |
|----------------|---------------------------------------------------------------------|
| REPORT         | fr-20260901-88c3cdf1                                                |
| ISSUED         | 2026-09-01                                                          |
| SUBJECT EXAM   | public.swe.pytoolz-toolz-issue-602 · v0.1.0                         |
| FAMILY         | public_oss_swe                                                      |
| PREPARED BY    | VVDex Forge engine vvdex-env 0.1.0                                  |
| CLASSIFICATION | Independent model evaluation report — independently verifiable      |
| CERTIFICATION  | Exam certification: partially recorded (see Certification evidence) |
| STATUS         | FINAL · report sealed against its own digest                        |

Timeline

| WHEN (UTC)          | EVENT                                         |
|---------------------|-----------------------------------------------|
| 2026-09-01 21:49:07 | Evaluation run started                        |
| 2026-09-01 22:11:03 | Evaluation run finished                       |
| 2026-09-01 22:11:03 | Report generated from the sealed run evidence |

Models under test

| MODEL                                  | PROVIDER | ENDPOINT         | BILLING                                            |
|----------------------------------------|----------|------------------|----------------------------------------------------|
| Claude Code CLI (Sonnet, subscription) | cli      | local runtime    | operator subscription, flat — not billed per token |
| Codex CLI (GPT, subscription)          | cli      | local runtime    | operator subscription, flat — not billed per token |
| codestral-latest                       | mistral  | VVDex-configured | VVDex free-quota                                   |
| openai/gpt-oss-120b                    | groq     | VVDex-configured | VVDex free-quota                                   |

A customer endpoint is recorded by origin only — never its port, its path, or its key.

Limits

|                                       |                                         |                                              |                                                                                                     |
|---------------------------------------|-----------------------------------------|----------------------------------------------|-----------------------------------------------------------------------------------------------------|
| <div>RUNTIME</div> <div>docker</div>  | <div>STEP LIMIT</div> <div>40</div>     | <div>WALL-CLOCK LIMIT</div> <div>9720s</div> | <div>MAX OUTPUT TOKENS</div> <div>2048</div>                                                        |
| <div>TEMPERATURE</div> <div>0.2</div> | <div>RETRY POLICY</div> <div>none</div> | <div>COMMAND BUDGET</div> <div>0</div>       | <div>TOOLS OFFERED</div> <div>list_files, read_file, write_file, edit_file, run_tests, submit</div> |

The evaluated model never saw the hidden tests or the reference solution, and could not change its result. Grading ran in a separate step, after submission, on a container with no network.

Model outcomes

**Denominators.** Attempts requested: 4. Graded (evaluable) rollouts: 3. Not graded — lane errors: 1. A lane error is a rollout the API or the runtime ended before the hidden grader ran: it is listed, excluded from every rate and pass@k, and never silently dropped. Passes are certified passes: hidden grader exit 0 with integrity verified.

| MODEL                                                       | ATTEMPTS | GRADED | PASSED     | MODEL FAILS | LANE ERRORS | SUBMITTED | MEAN TIME | TYPICAL DEEPEST STAGE |
|-------------------------------------------------------------|----------|--------|------------|-------------|-------------|-----------|-----------|-----------------------|
| Claude Code CLI (Sonnet, subscription)<br>cli/claude-sonnet | 1        | 1      | 1 / 1      | 0           | 0           | 1 / 1     | 93.2s     | Passed hidden tests   |
| Codex CLI (GPT, subscription)<br>cli/codex                  | 1        | 1      | 0 / 1      | 1           | 0           | 1 / 1     | 54.8s     | Graded                |
| codestral-latest<br>codestral-latest                        | 1        | 1      | 0 / 1      | 1           | 0           | 1 / 1     | 16.7s     | Graded                |
| openai/gpt-oss-120b<br>openai/gpt-oss-120b                  | 1        | 0      | not graded | 0           | 1           | —         | —         | nothing               |

Every rollout, with its own deepest stage and grade, is in Appendix A.

Success rate, confidence and pass@k

Every rate on this page is a measurement of a finite number of rollouts, so every rate carries the interval that measurement actually supports. n is the evaluable count — attempts minus rollouts the lane ended before grading; those are listed under "not graded" and sit in no rate, interval or pass@k.

| MODEL               | ATTEMPTS | N | NOT GRADED | PASSES | RATE         | 95% INTERVAL | SEEDS | PASS@1 |
|---------------------|----------|---|------------|--------|--------------|--------------|-------|--------|
| cli/claude-sonnet   | 1        | 1 | 0          | 1      | 100.0%       | [21%, 100%]  | 0     | 100.0% |
| cli/codex           | 1        | 1 | 0          | 0      | 0.0%         | [0%, 79%]    | 0     | 0.0%   |
| codestral-latest    | 1        | 1 | 0          | 0      | 0.0%         | [0%, 79%]    | 0     | 0.0%   |
| openai/gpt-oss-120b | 1        | 0 | 1          | 0      | not recorded | [0%, 100%]   | 0     | —      |

Interval and pass@k methods are stated in full in Appendix D. pass@k is shown for k = 1; the sealed record carries every k up to n.

Model comparison

| MODEL               | PASS@1       | 95% INTERVAL | TOKENS | COST       | MEAN ROLLOUT | OVERCLAIM RATE | TOP FAILURE CLASS |
|---------------------|--------------|--------------|--------|------------|--------------|----------------|-------------------|
| cli/claude-sonnet   | 100.0%       | [21%, 100%]  | n/a    | n/a        | 93.2s        | 0.0%           | none              |
| cli/codex           | 0.0%         | [0%, 79%]    | n/a    | n/a        | 54.8s        | 0.0%           | submitted_failed  |
| codestral-latest    | 0.0%         | [0%, 79%]    | 40 054 | not billed | 16.7s        | 0.0%           | submitted_failed  |
| openai/gpt-oss-120b | not recorded | [0%, 100%]   | n/a    | n/a        | 0ms          | 0.0%           | model_api_failure |

Tokens and cost read n/a where a lane reports no telemetry (a flat-rate CLI) — an absent measurement, not zero. pass@1 and its interval are over evaluable rollouts only.

This is not a leaderboard. Not separated (intervals overlap): cli/claude-sonnet vs cli/codex; cli/claude-sonnet vs codestral-latest; cli/claude-sonnet vs openai/gpt-oss-120b; cli/codex vs codestral-latest; cli/codex vs openai/gpt-oss-120b; codestral-latest vs openai/gpt-oss-120b. Below the 10-rollout floor for reading a rate as a rate: cli/claude-sonnet (1), cli/codex (1), codestral-latest (1), openai/gpt-oss-120b (0) — read the count, not the percentage. The primary comparison in this report is the stage funnel below — where a model stops is a finding; a percentage over this many rollouts is not.



Stage funnel

The funnel is the primary reading of a run. Where a model stops says more than a single pass percentage does, especially at low N.

| STAGE               | CLI:CLI/CLAUDE-SONNET | CLI:CLI/CODEX | MISTRAL:CODESTRAL-LATEST | GROQ:OPENAI/GPT-OSS-120B |
|---------------------|-----------------------|---------------|--------------------------|--------------------------|
| Inspected           | 1 / 1                 | 1 / 1         | 1 / 1                    | 0 / 0                    |
| Edited              | 1 / 1                 | 1 / 1         | 0 / 1                    | 0 / 0                    |
| Ran tests           | 1 / 1                 | 1 / 1         | 1 / 1                    | 0 / 0                    |
| Submitted           | 1 / 1                 | 1 / 1         | 1 / 1                    | 0 / 0                    |
| Graded              | 1 / 1                 | 1 / 1         | 1 / 1                    | 0 / 0                    |
| Passed hidden tests | 1 / 1                 | 0 / 1         | 0 / 1                    | 0 / 0                    |

Deepest stage reached

| MODEL                    | DEEPEST (ANY ATTEMPT) | TYPICAL (MOST ATTEMPTS) | HIDDEN-TEST PASSES | NOT GRADED |
|--------------------------|-----------------------|-------------------------|--------------------|------------|
| cli:cli/claude-sonnet    | Passed hidden tests   | Passed hidden tests     | 1 / 1              | 0          |
| cli:cli/codex            | Graded                | Graded                  | 0 / 1              | 0          |
| mistral:codestral-latest | Graded                | Graded                  | 0 / 1              | 0          |
| groq:openai/gpt-oss-120b | nothing               | nothing                 | 0 / 0              | 1          |

Stage definitions are in Appendix D. The matrix counts evaluable rollouts; lane errors are shown beside it, not inside it. The stages are not strictly nested: a model can submit without editing, and the matrix shows that rather than hiding it. Each rollout's own deepest stage is in Appendix A.

Behavioural analysis

Long-horizon behaviour

Counted off the recorded trace of each rollout. A dead end is a step that moved nothing: a re-read of a slice already returned, a repeat of the previous action, or a call to a tool this exam does not have.

| MODEL                    | N | STEPS USED | RAN OUT OF STEPS | COMMANDS | CHANGED THE TREE | REFUSED (OVER BUDGET) | REVISITS | DEAD ENDS / ROLLOUT |
|--------------------------|---|------------|------------------|----------|------------------|-----------------------|----------|---------------------|
| cli:cli/claude-sonnet    | 1 | 9 / 40     | 0                | 0 / 0    | 0                | 0                     | 0        | 3                   |
| cli:cli/codex            | 1 | 7 / 40     | 0                | 0 / 0    | 0                | 0                     | 0        | 0                   |
| mistral:codestral-latest | 1 | 6 / 40     | 0                | 0 / 0    | 0                | 0                     | 0        | 0                   |
| groq:openai/gpt-oss-120b | 1 | 1 / 40     | 0                | 0 / 0    | 0                | 0                     | 0        | 0                   |

Each column is defined in Appendix D.

Hallucination & overclaim

Three measurements, each with a stated definition. None of them is a judgement of the model's prose by another model — every number below comes from comparing what the model wrote against what was on disk, or against what the independent grader returned.

| MODEL                    | N | GROUNDING FAILURES | KINDS | ROLLOUTS AFFECTED | OVERCLAIMS | OVERCLAIM RATE | RE-READS |
|--------------------------|---|--------------------|-------|-------------------|------------|----------------|----------|
| cli:cli/claude-sonnet    | 1 | 0                  | none  | 0 / 1             | 0 / 1      | 0.0%           | 0        |
| cli:cli/codex            | 1 | 0                  | none  | 0 / 1             | 0 / 1      | 0.0%           | 0        |
| mistral:codestral-latest | 1 | 0                  | none  | 0 / 1             | 0 / 1      | 0.0%           | 0        |
| groq:openai/gpt-oss-120b | 1 | 0                  | none  | 0 / 1             | 0 / 1      | 0.0%           | 0        |

No grounding failures and no false pass claims were detected across these 4 rollout(s). No lane re-read material it had already been given.

Every term above is defined in Appendix D.

Efficiency & telemetry

This run has no per-token price attached — a local model, a free quota, or a customer endpoint we are not billed for — so spend is reported in tokens. Tokens are what was actually measured; a dollar figure here would be invented. Lanes run through a flat-rate local CLI report no token telemetry; their cells say n/a — an absent measurement, not zero. Rollout counts here are evaluable rollouts.

| MODEL                    | ROLLOUTS | TOKENS | TOKENS / ROLLOUT | COST                        | PASSES | RATE   |
|--------------------------|----------|--------|------------------|-----------------------------|--------|--------|
| cli:cli/claude-sonnet    | 1        | n/a    | n/a              | n/a — telemetry unavailable | 1 / 1  | 100.0% |
| cli:cli/codex            | 1        | n/a    | n/a              | n/a — telemetry unavailable | 0 / 1  | 0.0%   |
| mistral:codestral-latest | 1        | 40 054 | 40 054           | not billed                  | 0 / 1  | 0.0%   |
| groq:openai/gpt-oss-120b | 0        | n/a    | n/a              | n/a — telemetry unavailable | 0 / 0  | 0.0%   |

Findings

Every finding below is derived from the recorded data of this run — none is an opinion.  
Severity: PASS a lane succeeded · ATTENTION a capability gap · LANE / INFRA infrastructure, not capability · NOTE a property of the measurement.

F-01

PASS

cli/claude-sonnet passed the independent hidden tests

Severity: PASS

Evidence: Model outcomes · Stage funnel · Appendix A

1 of 3 graded rollouts completed the full loop (inspect, edit, run the visible tests, submit) and the sealed grader accepted the submission. This demonstrates the task is solvable under the published limits.

F-01 · derived from the recorded run data

F-02

ATTENTION

submitted\_failed × 2 (cli/codex, codestral-latest)

Severity: ATTENTION

Evidence: Appendix A · Appendix C

Submitted a change; the independent hidden tests still failed.

F-02 · derived from the recorded run data

F-03

LANE / INFRA

model\_api\_failure × 1 (openai/gpt-oss-120b)

Severity: LANE / INFRA

Evidence: Appendix A · Appendix C

The model endpoint failed to answer. Not a verdict about the model's ability. These rollouts are excluded from any capability reading.

F-03 · derived from the recorded run data

F-04

NOTE

Statistical floor — 4 rollout(s) is below the 10-rollout rate floor

Severity: NOTE

Evidence: Appendix D

The counts are exact; the percentages are anecdotes with decimal points. Per-rollout outcomes, not rates, are the reliable reading of this run.

F-04 · derived from the recorded run data

Recommendations

One recommendation per observed failure mode, addressed to the lane it affects. Each traces back to a finding above and to the raw trace in Appendix A.

**R-01** PROBLEM  
**Submitted a change; the independent hidden tests still failed.**

AFFECTED LANES  
**cli/codex, codestral-latest**

WHAT WE WOULD LOOK AT FIRST  
The model completes the loop and produces a plausible near-miss — the named failing tests are the exact behavioral cases distinguishing it from the reference fix.

**R-02** PROBLEM  
**The model endpoint failed to answer. Not a verdict about the model's ability.**

AFFECTED LANES  
**openai/gpt-oss-120b**

WHAT WE WOULD LOOK AT FIRST  
The lane failed, not the model; these rollouts carry no signal about capability and are labeled accordingly.

**R-03** **For anyone reading the percentages**

Commission a run at  $n \geq 10$  rollouts per lane before treating any rate here as a rate; this run's per-rollout outcomes are exact, its percentages are not yet stable.

Verification

RUN ID

live-20260901T214907Z

EXAM FINGERPRINT

e8312ce98b82558f9a1c5593e7d431f409fb1024b3013b7dbd934f1398a8937f

The contract digest this run was executed against.

Provenance of this chapter

This chapter was rendered by the campaign generator from the sealed evaluation record; it is not the standalone report reproduced byte for byte. Verify the standalone report with its own digest, verify the record with `vvdex-env records verify`, and verify this campaign document as a whole.

RELATIONSHIP

**campaign-rendered chapter**

Derived from the same sealed evaluation record as the standalone report; these bytes are authenticated by the campaign artifact that contains them.

SEALED RECORD DIGEST

**f49a31336e4865c1dcefccf5141e2f6926110c72f03d6cb3d645a1761fa09ac8**

sha256 over the record's canonical JSON — the identity every renderer works from.

SOURCE RECORD FILE

**6bafd70a8214b1e782b25d64d4733b0b2223a2cffe518820df0a030ae5f4673b**

sha256 of `fr-20260901-88c3cdf1.json` as written beside the standalone report.

SOURCE STANDALONE REPORT

**bfb1b0dcbb7574e53573907812a76f0134b47a36ebbeb9d04ff543d189ffe58e**

sha256 of `fr-20260901-88c3cdf1.html` — independently verifiable with its own digest.

SOURCE STANDALONE PDF

**e85ac24764612e3343a739a09e26bcad398b003352b9877b03d4a5a79bdbd973**

sha256 of `fr-20260901-88c3cdf1.pdf`.

RUN EVIDENCE BUNDLE

**88c3cdf1c1b6d606e81360be7fcb06bd12c6340699acd7e3737cd37ebfce226c**

The digest the run's own bundle computed over itself.

CERTIFIED EVIDENCE SEAL

**not available in this package**

The exam's sealed evidence, established before this run.

Measurement history

This is the first recorded run on this exam. There is nothing to compare it against yet, and a single measurement is not a trend.

The full artifact-by-artifact digest table, and the reproduction protocol, are in Appendix B.

Appendix A — every rollout

| ROLLOUT  | MODEL               | SEED | OUTCOME           | SUBMITTED | DEEPEST STAGE           | HIDDEN TESTS | FILES | TOKENS | ELAPSED |
|----------|---------------------|------|-------------------|-----------|-------------------------|--------------|-------|--------|---------|
| ea3e4b9b | cli/claude-sonnet   | 0    | submitted_passed  | yes       | Passed hidden tests     | pass         | 1     | n/a    | 93.2s   |
| bbf9607d | cli/codex           | 0    | submitted_failed  | yes       | Graded                  | fail         | 1     | n/a    | 54.8s   |
| 441654d0 | codestral-latest    | 0    | submitted_failed  | yes       | Graded                  | fail         | 0     | 40 054 | 16.7s   |
| 5f24c149 | openai/gpt-oss-120b | 0    | model_api_failure | no        | Not graded — lane error | not graded   | 0     | 1 164  | 1129.5s |

Failure taxonomy

| CLASS             | COUNT | WHAT IT MEANS                                                                 |
|-------------------|-------|-------------------------------------------------------------------------------|
| submitted_failed  | 2     | Submitted a change; the independent hidden tests still failed.                |
| model_api_failure | 1     | The model endpoint failed to answer. Not a verdict about the model's ability. |

Appendix B — evidence and artifact digests

RUN BUNDLE DIGEST

88c3cdf1c1b6d606e81360be7fcb06bd12c6340699acd7e3737cd37ebfce226c

RUN ID

live-20260901T214907Z

STARTED

2026-09-01T21:49:07.619617+00:00

FINISHED

2026-09-01T22:11:03.052288+00:00

Artifact digests

| ARTIFACT                                           | SHA-256                                                          |
|----------------------------------------------------|------------------------------------------------------------------|
| publicSummary                                      | a0a7dfe8cfc33cf167972072583fd76cf2fbc9fa0e3635c077b5cfcc2d907abd |
| rollouts/441654d0-e828-4779-91df-e35763a02c85.json | f4d362a1e4dac96705b00d459b69dd4350b888e5c02e99199e633e5eb7cc6483 |
| rollouts/5f24c149-f696-4b36-861f-26fd7281b1bc.json | 981516b5a142693648fb800d9dc7ccda9b470362261d74211e5d9afb0ff4664c |
| rollouts/bbf9607d-548f-4d11-9aae-777105292a62.json | 9bf2824267ce3d005aa8bc66ff581cc5f433ad43bc4b241dac4218f1c19eb918 |
| rollouts/ea3e4b9b-fe43-4913-bcbe-0c61631dd1b5.json | 6c1c277aafcd3bc8378613b18967cf3d3cb7d37bdc786296494821cc430eeec2 |
| spec                                               | a29231f530943bca94c1b1f15da405051ea0d695923657aa869c2c5d8716de27 |
| summary                                            | dd09a0bb2c237da0371d9f3ef41236f0be9b6647c31303d70d6991cf19d53fd0 |

Reproducing this run

Every rollout can be re-derived: the exam is frozen (fingerprint above), the grader ran in the recorded container image, and each rollout's full trace — every tool call, every result, every word — is in the evidence bundle at its rollouts/<id>.json path. A re-run uses the same job shape: the exam id, your endpoint, n, and the runtime. The record id names this run forever.



Appendix C — redacted transcripts

For each rollout: the tool calls it made, in order — which tool, on which of the model's own paths, whether the call succeeded, and how many characters came back. What a tool returned is not here, and neither is the submission: on this exam the submitted content is the graded answer, so no diff, no answer value and no model claim is printed. A request for the reference solution or the hidden tests is recorded as a request; the path it asked for is replaced by a label. See Appendix D, "What this report withholds and why".

Disclosure class applied to this exam: answer-key disclosure (ownership vvdex\_proprietary — no declaration — proprietary by default; family public\_oss\_swe): the exam is VVDex evaluation IP and the graded output is an answer or a keyed value, so no submission content, diff, model claim or hidden-assertion name is published.

ea3e4b9b · cli:cli/claude-sonnet · seed 0 · submitted\_passed

| STEP | TOOL              | PATH               | RESULT | CHARS  |
|------|-------------------|--------------------|--------|--------|
| 0    | read_file         | ISSUE.md           | ok     | 1 371  |
| 1    | list_files        | toolz              | ok     | 304    |
| 2    | read_file         | toolz/itertools.py | ok     | 12 166 |
| 3    | read_file         | toolz/itertools.py | ok     | 12 170 |
| 4    | invalid_tool_call | (none)             | failed | 107    |
| 5    | invalid_tool_call | (none)             | failed | 107    |
| 6    | edit_file         | toolz/itertools.py | ok     | 56     |
| 7    | run_tests         | —                  | ok     | 179    |
| 8    | submit            | —                  | ok     | 85     |

Submission diff withheld by the public-disclosure policy (58 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

bbf9607d · cli:cli/codex · seed 0 · submitted\_failed

| STEP | TOOL       | PATH               | RESULT | CHARS  |
|------|------------|--------------------|--------|--------|
| 0    | list_files | .                  | ok     | 336    |
| 1    | read_file  | ISSUE.md           | ok     | 1 371  |
| 2    | read_file  | toolz/itertools.py | ok     | 12 166 |
| 3    | read_file  | toolz/itertools.py | ok     | 12 170 |
| 4    | edit_file  | toolz/itertools.py | ok     | 56     |
| 5    | run_tests  | —                  | ok     | 179    |
| 6    | submit     | —                  | ok     | 85     |

Submission diff withheld by the public-disclosure policy (69 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

441654d0 · mistral:codestral-latest · seed 0 · submitted\_failed

| STEP | TOOL      | PATH               | RESULT | CHARS  |
|------|-----------|--------------------|--------|--------|
| 0    | read_file | ISSUE.md           | ok     | 1 371  |
| 1    | read_file | toolz/itertools.py | ok     | 12 166 |
| 2    | read_file | toolz/itertools.py | ok     | 12 170 |
| 3    | read_file | toolz/itertools.py | ok     | 3 618  |
| 4    | run_tests | —                  | ok     | 179    |
| 5    | submit    | —                  | ok     | 85     |

Submission diff withheld by the public-disclosure policy. Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

5f24c149 · groq:openai/gpt-oss-120b · seed 0 · model\_api\_failure

| STEP | TOOL       | PATH | RESULT | CHARS |
|------|------------|------|--------|-------|
| 0    | list_files | .    | ok     | 336   |

Submission diff withheld by the public-disclosure policy. Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

11 Exam · r1chardj0n3s-parse-issue-102

public.swe.r1chardj0n3s-parse-issue-102 · record fr-20260901-4b1e3251 · record digest d94c313144f8b36d... · harness clean · containment clean

Executive summary

MODELS TESTED

4

CERTIFIED PASSES

1

of 3 graded rollouts

SUBMITTED

3 of 3

graded rollouts that called submit

NOT GRADED — LANE ERRORS

1

of 4 attempts

Purpose

This report presents a controlled evaluation of 4 model lane(s) against the certified exam public.swe.r1chardj0n3s-parse-issue-102 (v0.1.0, family public\_oss\_swe). Every lane received the identical frozen workspace, tool contract and limits; grading ran afterwards, in isolation, against hidden tests no lane ever saw. The purpose is to state, with reproducible evidence, what each lane did and where it stopped.

Outcome

1 of 3

graded rollouts passed the independent hidden tests and were certified — 1 of 4 attempts not graded (lane errors)

1 of 3 graded rollouts passed. cli/claude-sonnet completed the full loop (inspect, edit, run the visible tests, submit) and the sealed grader accepted the submission. **Low-N:** 4 rollout(s) is below the 10-rollout floor for reading a rate as a rate — read the per-rollout outcomes, not the percentages. The most common way to fail was submitted\_failed — submitted a change; the independent hidden tests still failed.

Key findings

- F-01 cli/claude-sonnet passed the independent hidden tests
- F-02 submitted\_failed × 2 (cli/codex, codestral-latest)
- F-03 model\_api\_failure × 1 (openai/gpt-oss-120b)
- F-04 Statistical floor — 4 rollout(s) is below the 10-rollout rate floor

Each finding is stated in full, with evidence, in the Findings section; recommendations for acting on them follow it. Elapsed wall clock for the whole engagement: 1469.0s.

Exam definition

**In scope.** The ability of each configured model lane to complete this one certified task end to end — inspect the workspace, produce a correct edit, verify it with the visible tests, and submit — within 40 tool steps and 9720s of wall clock, at 1 rollout(s) per lane. Behavioural measurements along the way (tool discipline, grounding, overclaim) are in scope because they are read from the recorded trace, not judged by another model.

**Out of scope.** General model quality, performance on other task families, prompt-tuning on the lane's behalf, and any comparison this run's sample size cannot support. A lane's API failure is reported as infrastructure, never as model capability.

**The instrument.** Exam `public.swe.r1chardj0n3s-parse-issue-102` was certified before this run: its reference solution passes, an empty submission fails, its grader distinguishes near-misses, and its sealed evidence is unchanged by this evaluation (see Certification evidence).

What was measured

|                                         |         |                |                 |
|-----------------------------------------|---------|----------------|-----------------|
| EXAM                                    | VERSION | FAMILY         | ROLLOUTS        |
| public.swe.r1chardj0n3s-parse-issue-102 | 0.1.0   | public_oss_swe | 4 (1 per model) |

Where the defect came from

**Source:** `r1chardj0n3s/parse (MIT)` · issue  
<https://github.com/r1chardj0n3s/parse/issues/102> · upstream fix  
<https://github.com/r1chardj0n3s/parse/pull/103>

The defect and its upstream fix are public and unmodified. The exam around them — the frozen parent tree, the independently written hidden suite, the control probes and the isolation — is VVDex's, and it is what this report measures.

Certification evidence

4 of 5 certification pillars are recorded in this package, loaded from the receipts sealed into the exam and matching this run's exam fingerprint. What is recorded is shown with its real value; what is not is named below and is not claimed.

FROZEN BASELINE

FAIL

expected — an empty submission over 2 run(s) must not pass

→

GOLD / REFERENCE

PASS

expected — the reference solution over 2 run(s) must pass

GRADER CONTROLS

7 / 7

The grader accepted the reference and rejected every near-miss.

ATTACK PROBES

10 / 10 blocked

0 trivial exploit(s) found.

RUNTIME

sha256:679b36225a1f83238efba8c71b9cde3c24caaeacc9b682ae92819cb55eec328f

network policy: deny

Certification metadata is not included in this report package for:

- sealed evidence bundle

Benchmark results remain valid as run evidence; the pillar(s) above are not substantiated by this document and this report does not claim them.

What the package does carry: the exam fingerprint below is the contract digest this run was executed against — a holder of the certification bundle for this fingerprint can join the two independently.

EXAM FINGERPRINT

483b68b383c851060781196403325de0e00e5f6327fd52d24eb581005be38812

The contract digest this run was executed against.

Evaluation configuration

|                |                                                                     |
|----------------|---------------------------------------------------------------------|
| REPORT         | fr-20260901-4b1e3251                                                |
| ISSUED         | 2026-09-01                                                          |
| SUBJECT EXAM   | public.swe.r1chardj0n3s-parse-issue-102 · v0.1.0                    |
| FAMILY         | public_oss_swe                                                      |
| PREPARED BY    | VVDex Forge engine vvdex-env 0.1.0                                  |
| CLASSIFICATION | Independent model evaluation report — independently verifiable      |
| CERTIFICATION  | Exam certification: partially recorded (see Certification evidence) |
| STATUS         | FINAL · report sealed against its own digest                        |

Timeline

| WHEN (UTC)          | EVENT                                         |
|---------------------|-----------------------------------------------|
| 2026-09-01 22:11:03 | Evaluation run started                        |
| 2026-09-01 22:35:32 | Evaluation run finished                       |
| 2026-09-01 22:35:32 | Report generated from the sealed run evidence |

Models under test

| MODEL                                  | PROVIDER | ENDPOINT         | BILLING                                            |
|----------------------------------------|----------|------------------|----------------------------------------------------|
| Claude Code CLI (Sonnet, subscription) | cli      | local runtime    | operator subscription, flat — not billed per token |
| Codex CLI (GPT, subscription)          | cli      | local runtime    | operator subscription, flat — not billed per token |
| codestral-latest                       | mistral  | VVDex-configured | VVDex free-quota                                   |
| openai/gpt-oss-120b                    | groq     | VVDex-configured | VVDex free-quota                                   |

A customer endpoint is recorded by origin only — never its port, its path, or its key.

Limits

|                    |                      |                           |                                                                                  |
|--------------------|----------------------|---------------------------|----------------------------------------------------------------------------------|
| RUNTIME<br>docker  | STEP LIMIT<br>40     | WALL-CLOCK LIMIT<br>9720s | MAX OUTPUT TOKENS<br>2048                                                        |
| TEMPERATURE<br>0.2 | RETRY POLICY<br>none | COMMAND BUDGET<br>0       | TOOLS OFFERED<br>list_files, read_file, write_file, edit_file, run_tests, submit |

The evaluated model never saw the hidden tests or the reference solution, and could not change its result. Grading ran in a separate step, after submission, on a container with no network.

Model outcomes

**Denominators.** Attempts requested: 4. Graded (evaluable) rollouts: 3. Not graded — lane errors: 1. A lane error is a rollout the API or the runtime ended before the hidden grader ran: it is listed, excluded from every rate and pass@k, and never silently dropped. Passes are certified passes: hidden grader exit 0 with integrity verified.

| MODEL                                                              | ATTEMPTS | GRADED | PASSED     | MODEL FAILS | LANE ERRORS | SUBMITTED | MEAN TIME | TYPICAL DEEPEST STAGE |
|--------------------------------------------------------------------|----------|--------|------------|-------------|-------------|-----------|-----------|-----------------------|
| <b>Claude Code CLI (Sonnet, subscription)</b><br>cli/claude-sonnet | 1        | 1      | 1 / 1      | 0           | 0           | 1 / 1     | 255.5s    | Passed hidden tests   |
| <b>Codex CLI (GPT, subscription)</b><br>cli/codex                  | 1        | 1      | 0 / 1      | 1           | 0           | 1 / 1     | 72.9s     | Graded                |
| <b>codestral-latest</b><br>codestral-latest                        | 1        | 1      | 0 / 1      | 1           | 0           | 1 / 1     | 22.3s     | Graded                |
| <b>openai/gpt-oss-120b</b><br>openai/gpt-oss-120b                  | 1        | 0      | not graded | 0           | 1           | —         | —         | nothing               |

Every rollout, with its own deepest stage and grade, is in Appendix A.

Success rate, confidence and pass@k

Every rate on this page is a measurement of a finite number of rollouts, so every rate carries the interval that measurement actually supports. n is the evaluable count — attempts minus rollouts the lane ended before grading; those are listed under "not graded" and sit in no rate, interval or pass@k.

| MODEL               | ATTEMPTS | N | NOT GRADED | PASSES | RATE         | 95% INTERVAL | SEEDS | PASS@1 |
|---------------------|----------|---|------------|--------|--------------|--------------|-------|--------|
| cli/claude-sonnet   | 1        | 1 | 0          | 1      | 100.0%       | [21%, 100%]  | 0     | 100.0% |
| cli/codex           | 1        | 1 | 0          | 0      | 0.0%         | [0%, 79%]    | 0     | 0.0%   |
| codestral-latest    | 1        | 1 | 0          | 0      | 0.0%         | [0%, 79%]    | 0     | 0.0%   |
| openai/gpt-oss-120b | 1        | 0 | 1          | 0      | not recorded | [0%, 100%]   | 0     | —      |

Interval and pass@k methods are stated in full in Appendix D. pass@k is shown for k = 1; the sealed record carries every k up to n.

Model comparison

| MODEL               | PASS@1       | 95% INTERVAL | TOKENS | COST       | MEAN ROLLOUT | OVERCLAIM RATE | TOP FAILURE CLASS |
|---------------------|--------------|--------------|--------|------------|--------------|----------------|-------------------|
| cli/claude-sonnet   | 100.0%       | [21%, 100%]  | n/a    | n/a        | 255.5s       | 0.0%           | none              |
| cli/codex           | 0.0%         | [0%, 79%]    | n/a    | n/a        | 72.9s        | 0.0%           | submitted_failed  |
| codestral-latest    | 0.0%         | [0%, 79%]    | 73 260 | not billed | 22.3s        | 0.0%           | submitted_failed  |
| openai/gpt-oss-120b | not recorded | [0%, 100%]   | n/a    | n/a        | 0ms          | 0.0%           | model_api_failure |

Tokens and cost read n/a where a lane reports no telemetry (a flat-rate CLI) — an absent measurement, not zero. pass@1 and its interval are over evaluable rollouts only.

This is not a leaderboard. Not separated (intervals overlap): cli/claude-sonnet vs cli/codex; cli/claude-sonnet vs codestral-latest; cli/claude-sonnet vs openai/gpt-oss-120b; cli/codex vs codestral-latest; cli/codex vs openai/gpt-oss-120b; codestral-latest vs openai/gpt-oss-120b. Below the 10-rollout floor for reading a rate as a rate: cli/claude-sonnet (1), cli/codex (1), codestral-latest (1), openai/gpt-oss-120b (0) — read the count, not the percentage. The primary comparison in this report is the stage funnel below — where a model stops is a finding; a percentage over this many rollouts is not.

Stage funnel

The funnel is the primary reading of a run. Where a model stops says more than a single pass percentage does, especially at low N.

| STAGE               | CLI:CLI/CLAUDE-SONNET | CLI:CLI/CODEX | MISTRAL:CODESTRAL-LATEST | GROQ:OPENAI/GPT-OSS-120B |
|---------------------|-----------------------|---------------|--------------------------|--------------------------|
| Inspected           | 1 / 1                 | 1 / 1         | 1 / 1                    | 0 / 0                    |
| Edited              | 1 / 1                 | 1 / 1         | 1 / 1                    | 0 / 0                    |
| Ran tests           | 1 / 1                 | 1 / 1         | 1 / 1                    | 0 / 0                    |
| Submitted           | 1 / 1                 | 1 / 1         | 1 / 1                    | 0 / 0                    |
| Graded              | 1 / 1                 | 1 / 1         | 1 / 1                    | 0 / 0                    |
| Passed hidden tests | 1 / 1                 | 0 / 1         | 0 / 1                    | 0 / 0                    |

Deepest stage reached

| MODEL                    | DEEPEST (ANY ATTEMPT) | TYPICAL (MOST ATTEMPTS) | HIDDEN-TEST PASSES | NOT GRADED |
|--------------------------|-----------------------|-------------------------|--------------------|------------|
| cli:cli/claude-sonnet    | Passed hidden tests   | Passed hidden tests     | 1 / 1              | 0          |
| cli:cli/codex            | Graded                | Graded                  | 0 / 1              | 0          |
| mistral:codestral-latest | Graded                | Graded                  | 0 / 1              | 0          |
| groq:openai/gpt-oss-120b | nothing               | nothing                 | 0 / 0              | 1          |

Stage definitions are in Appendix D. The matrix counts evaluable rollouts; lane errors are shown beside it, not inside it. The stages are not strictly nested: a model can submit without editing, and the matrix shows that rather than hiding it. Each rollout's own deepest stage is in Appendix A.



Behavioural analysis

Long-horizon behaviour

Counted off the recorded trace of each rollout. A dead end is a step that moved nothing: a re-read of a slice already returned, a repeat of the previous action, or a call to a tool this exam does not have.

| MODEL                    | N | STEPS USED | RAN OUT OF STEPS | COMMANDS | CHANGED THE TREE | REFUSED (OVER BUDGET) | REVISITS | DEAD ENDS / ROLLOUT |
|--------------------------|---|------------|------------------|----------|------------------|-----------------------|----------|---------------------|
| cli:cli/claude-sonnet    | 1 | 11 / 40    | 0                | 0 / 0    | 0                | 0                     | 0        | 3                   |
| cli:cli/codex            | 1 | 11 / 40    | 0                | 0 / 0    | 0                | 0                     | 0        | 2                   |
| mistral:codestral-latest | 1 | 9 / 40     | 0                | 0 / 0    | 0                | 0                     | 0        | 1                   |
| groq:openai/gpt-oss-120b | 1 | 1 / 40     | 0                | 0 / 0    | 0                | 0                     | 0        | 0                   |

Each column is defined in Appendix D.

Hallucination & overclaim

Three measurements, each with a stated definition. None of them is a judgement of the model's prose by another model — every number below comes from comparing what the model wrote against what was on disk, or against what the independent grader returned.

| MODEL                    | N | GROUNDING FAILURES | KINDS | ROLLOUTS AFFECTED | OVERCLAIMS | OVERCLAIM RATE | RE-READS |
|--------------------------|---|--------------------|-------|-------------------|------------|----------------|----------|
| cli:cli/claude-sonnet    | 1 | 0                  | none  | 0 / 1             | 0 / 1      | 0.0%           | 0        |
| cli:cli/codex            | 1 | 0                  | none  | 0 / 1             | 0 / 1      | 0.0%           | 0        |
| mistral:codestral-latest | 1 | 0                  | none  | 0 / 1             | 0 / 1      | 0.0%           | 0        |
| groq:openai/gpt-oss-120b | 1 | 0                  | none  | 0 / 1             | 0 / 1      | 0.0%           | 0        |

No grounding failures and no false pass claims were detected across these 4 rollout(s). No lane re-read material it had already been given.

Every term above is defined in Appendix D.

Efficiency & telemetry

This run has no per-token price attached — a local model, a free quota, or a customer endpoint we are not billed for — so spend is reported in tokens. Tokens are what was actually measured; a dollar figure here would be invented. Lanes run through a flat-rate local CLI report no token telemetry; their cells say n/a — an absent measurement, not zero. Rollout counts here are evaluable rollouts.

| MODEL                    | ROLLOUTS | TOKENS | TOKENS / ROLLOUT | COST                        | PASSES | RATE   |
|--------------------------|----------|--------|------------------|-----------------------------|--------|--------|
| cli:cli/claude-sonnet    | 1        | n/a    | n/a              | n/a — telemetry unavailable | 1 / 1  | 100.0% |
| cli:cli/codex            | 1        | n/a    | n/a              | n/a — telemetry unavailable | 0 / 1  | 0.0%   |
| mistral:codestral-latest | 1        | 73 260 | 73 260           | not billed                  | 0 / 1  | 0.0%   |
| groq:openai/gpt-oss-120b | 0        | n/a    | n/a              | n/a — telemetry unavailable | 0 / 0  | 0.0%   |

Findings

Every finding below is derived from the recorded data of this run — none is an opinion.  
Severity: PASS a lane succeeded · ATTENTION a capability gap · LANE / INFRA infrastructure, not capability · NOTE a property of the measurement.

F-01

PASS

cli/claude-sonnet passed the independent hidden tests

Severity: PASS

Evidence: Model outcomes · Stage funnel · Appendix A

1 of 3 graded rollouts completed the full loop (inspect, edit, run the visible tests, submit) and the sealed grader accepted the submission. This demonstrates the task is solvable under the published limits.

F-01 · derived from the recorded run data

F-02

ATTENTION

submitted\_failed × 2 (cli/codex, codestral-latest)

Severity: ATTENTION

Evidence: Appendix A · Appendix C

Submitted a change; the independent hidden tests still failed.

F-02 · derived from the recorded run data

F-03

LANE / INFRA

model\_api\_failure × 1 (openai/gpt-oss-120b)

Severity: LANE / INFRA

Evidence: Appendix A · Appendix C

The model endpoint failed to answer. Not a verdict about the model's ability. These rollouts are excluded from any capability reading.

F-03 · derived from the recorded run data

F-04

NOTE

Statistical floor — 4 rollout(s) is below the 10-rollout rate floor

Severity: NOTE

Evidence: Appendix D

The counts are exact; the percentages are anecdotes with decimal points. Per-rollout outcomes, not rates, are the reliable reading of this run.

F-04 · derived from the recorded run data

Recommendations

One recommendation per observed failure mode, addressed to the lane it affects. Each traces back to a finding above and to the raw trace in Appendix A.

**R-01** PROBLEM  
**Submitted a change; the independent hidden tests still failed.**

AFFECTED LANES  
**cli/codex, codestral-latest**

WHAT WE WOULD LOOK AT FIRST  
The model completes the loop and produces a plausible near-miss — the named failing tests are the exact behavioral cases distinguishing it from the reference fix.

**R-02** PROBLEM  
**The model endpoint failed to answer. Not a verdict about the model's ability.**

AFFECTED LANES  
**openai/gpt-oss-120b**

WHAT WE WOULD LOOK AT FIRST  
The lane failed, not the model; these rollouts carry no signal about capability and are labeled accordingly.

**R-03** **For anyone reading the percentages**

Commission a run at  $n \geq 10$  rollouts per lane before treating any rate here as a rate; this run's per-rollout outcomes are exact, its percentages are not yet stable.

Verification

RUN ID

**live-20260901T221103Z**

EXAM FINGERPRINT

**483b68b383c851060781196403325de0e00e5f6327fd52d24eb581005be38812**

The contract digest this run was executed against.

Provenance of this chapter

This chapter was rendered by the campaign generator from the sealed evaluation record; it is not the standalone report reproduced byte for byte. Verify the standalone report with its own digest, verify the record with `vvdex-env records verify`, and verify this campaign document as a whole.

RELATIONSHIP

**campaign-rendered chapter**

Derived from the same sealed evaluation record as the standalone report; these bytes are authenticated by the campaign artifact that contains them.

SEALED RECORD DIGEST

**d94c313144f8b36db4448440b873462bc9372e1783b1164451695b7b8497816a**

sha256 over the record's canonical JSON — the identity every renderer works from.

SOURCE RECORD FILE

**31dd2cfcaf563ac9ae7c31d50b5d8adf9d8ebbc97646767d922b02731ad62f88**

sha256 of fr-20260901-4b1e3251.json as written beside the standalone report.

SOURCE STANDALONE REPORT

**3410a4fc424b582f2fd03119f69ff337cd9785ec620503730755c9b78379987b**

sha256 of fr-20260901-4b1e3251.html — independently verifiable with its own digest.

SOURCE STANDALONE PDF

**bddb2882e46b9be29bdfdfef30188294705be407c601284a7b2caef6c7d697155**

sha256 of fr-20260901-4b1e3251.pdf.

RUN EVIDENCE BUNDLE

**4b1e3251d3e5a146a99f2e348bb71c76953245e7f85d35c8f5a1a24011faa5ed**

The digest the run's own bundle computed over itself.

CERTIFIED EVIDENCE SEAL

**not available in this package**

The exam's sealed evidence, established before this run.

Measurement history

This is the first recorded run on this exam. There is nothing to compare it against yet, and a single measurement is not a trend.

The full artifact-by-artifact digest table, and the reproduction protocol, are in Appendix B.

Appendix A — every rollout

| ROLLOUT  | MODEL               | SEED | OUTCOME           | SUBMITTED | DEEPEST STAGE           | HIDDEN TESTS | FILES | TOKENS | ELAPSED |
|----------|---------------------|------|-------------------|-----------|-------------------------|--------------|-------|--------|---------|
| 14c94f8c | cli/claude-sonnet   | 0    | submitted_passed  | yes       | Passed hidden tests     | pass         | 3     | n/a    | 255.5s  |
| 1f2eb2de | cli/codex           | 0    | submitted_failed  | yes       | Graded                  | fail         | 3     | n/a    | 72.9s   |
| 890eafa7 | codestral-latest    | 0    | submitted_failed  | yes       | Graded                  | fail         | 1     | 73 260 | 22.3s   |
| c4fb19d4 | openai/gpt-oss-120b | 0    | model_api_failure | no        | Not graded — lane error | not graded   | 0     | 1 381  | 1097.0s |

Failure taxonomy

| CLASS             | COUNT | WHAT IT MEANS                                                                 |
|-------------------|-------|-------------------------------------------------------------------------------|
| submitted_failed  | 2     | Submitted a change; the independent hidden tests still failed.                |
| model_api_failure | 1     | The model endpoint failed to answer. Not a verdict about the model's ability. |

Appendix B — evidence and artifact digests

RUN BUNDLE DIGEST

4b1e3251d3e5a146a99f2e348bb71c76953245e7f85d35c8f5a1a24011faa5ed

RUN ID

live-20260901T221103Z

STARTED

2026-09-01T22:11:03.440919+00:00

FINISHED

2026-09-01T22:35:32.433884+00:00

Artifact digests

| ARTIFACT                                           | SHA-256                                                           |
|----------------------------------------------------|-------------------------------------------------------------------|
| publicSummary                                      | d782d1af3bde0176db198da9b48bc877239d6988eaf07e48540e16fd4be4bbce  |
| rollouts/14c94f8c-9a21-4fab-9d6d-4b3d379e5fa1.json | be97be4950dccc8cf96c33573f8b3ce5dc5387afcac804f422d6c0ac137c4b20  |
| rollouts/1f2eb2de-ddf7-4ca8-88b9-392f2a829f01.json | 881c103264c5bc44358ab0fcb4229d0f8387f97d453c16ac2c762250e045189f  |
| rollouts/890eafa7-f9d7-4805-818b-8003d1157013.json | 4cb38d26a39c49f2799928f1f2209d5a1f31ea71f402dc84b35e832dcfc943bd  |
| rollouts/c4fb19d4-b5b5-4cc2-a6cf-0c9e2c815b04.json | b77ccd2450443c6339c9e99d4ec89af7e9bd83a7e2edaca195967dfdac9f03c4  |
| spec                                               | b15c867013aeb4603011ad5b63f4244fa12213d3a3bc8e3cb944efebbb47b2d81 |
| summary                                            | 1a300c4fdd48b1dedc2b7632e5a495d39413e29f668ade037463d474421b7273  |

Reproducing this run

Every rollout can be re-derived: the exam is frozen (fingerprint above), the grader ran in the recorded container image, and each rollout's full trace — every tool call, every result, every word — is in the evidence bundle at its rollouts/<id>.json path. A re-run uses the same job shape: the exam id, your endpoint, n, and the runtime. The record id names this run forever.

Appendix C — redacted transcripts

For each rollout: the tool calls it made, in order — which tool, on which of the model's own paths, whether the call succeeded, and how many characters came back. What a tool returned is not here, and neither is the submission: on this exam the submitted content is the graded answer, so no diff, no answer value and no model claim is printed. A request for the reference solution or the hidden tests is recorded as a request; the path it asked for is replaced by a label. See Appendix D, "What this report withholds and why".

Disclosure class applied to this exam: answer-key disclosure (ownership vvdex\_proprietary — no declaration — proprietary by default; family public\_oss\_swe): the exam is VVDex evaluation IP and the graded output is an answer or a keyed value, so no submission content, diff, model claim or hidden-assertion name is published.

14c94f8c · cli:cli/claude-sonnet · seed 0 · submitted\_passed

| STEP | TOOL              | PATH     | RESULT | CHARS  |
|------|-------------------|----------|--------|--------|
| 0    | read_file         | ISSUE.md | ok     | 1 181  |
| 1    | read_file         | parse.py | ok     | 12 166 |
| 2    | read_file         | parse.py | ok     | 12 170 |
| 3    | read_file         | parse.py | ok     | 12 170 |
| 4    | read_file         | parse.py | ok     | 12 170 |
| 5    | invalid_tool_call | (none)   | failed | 107    |
| 6    | edit_file         | parse.py | ok     | 46     |
| 7    | edit_file         | parse.py | ok     | 46     |
| 8    | edit_file         | parse.py | ok     | 46     |
| 9    | run_tests         | —        | ok     | 179    |
| 10   | submit            | —        | ok     | 85     |

Submission diff withheld by the public-disclosure policy (101 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

1f2eb2de · cli:cli/codex · seed 0 · submitted\_failed

| STEP | TOOL       | PATH     | RESULT | CHARS  |
|------|------------|----------|--------|--------|
| 0    | list_files | .        | ok     | 36     |
| 1    | read_file  | ISSUE.md | ok     | 1 181  |
| 2    | read_file  | parse.py | ok     | 12 166 |
| 3    | read_file  | parse.py | ok     | 12 170 |
| 4    | read_file  | parse.py | ok     | 12 170 |
| 5    | read_file  | parse.py | ok     | 12 170 |
| 6    | edit_file  | parse.py | ok     | 46     |
| 7    | edit_file  | parse.py | ok     | 46     |
| 8    | edit_file  | parse.py | ok     | 46     |
| 9    | run_tests  | —        | ok     | 179    |



| STEP | TOOL   | PATH | RESULT | CHARS |
|------|--------|------|--------|-------|
| 10   | submit | —    | ok     | 85    |

Submission diff withheld by the public-disclosure policy (74 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

890eafa7 · mistral:codestral-latest · seed 0 · submitted\_failed

| STEP | TOOL              | PATH     | RESULT | CHARS  |
|------|-------------------|----------|--------|--------|
| 0    | list_files        | .        | ok     | 36     |
| 1    | read_file         | parse.py | ok     | 12 166 |
| 2    | read_file         | parse.py | ok     | 12 170 |
| 3    | read_file         | parse.py | ok     | 12 170 |
| 4    | read_file         | parse.py | ok     | 924    |
| 5    | invalid_tool_call | (none)   | failed | 107    |
| 6    | edit_file         | parse.py | ok     | 46     |
| 7    | run_tests         | —        | ok     | 179    |
| 8    | submit            | —        | ok     | 85     |

Submission diff withheld by the public-disclosure policy (12 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

c4fb19d4 · groq:openai/gpt-oss-120b · seed 0 · model\_api\_failure

| STEP | TOOL       | PATH | RESULT | CHARS |
|------|------------|------|--------|-------|
| 0    | list_files | .    | ok     | 36    |

Submission diff withheld by the public-disclosure policy. Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

12 Exam · grounded-rag-1

vvdex.knowledge.grounded-rag-1 · record fr-20260901-718db345 · record digest a6426b0b7932b6c0... · harness clean · containment clean

Executive summary

|                                       |                                                                          |                                                                                      |                                                                           |
|---------------------------------------|--------------------------------------------------------------------------|--------------------------------------------------------------------------------------|---------------------------------------------------------------------------|
| <div>MODELS TESTED</div> <div>4</div> | <div>CERTIFIED PASSES</div> <div>1</div> <div>of 3 graded rollouts</div> | <div>SUBMITTED</div> <div>3 of 3</div> <div>graded rollouts that called submit</div> | <div>NOT GRADED — LANE ERRORS</div> <div>1</div> <div>of 4 attempts</div> |
|---------------------------------------|--------------------------------------------------------------------------|--------------------------------------------------------------------------------------|---------------------------------------------------------------------------|

Purpose

This report presents a controlled evaluation of 4 model lane(s) against the certified exam vvdex.knowledge.grounded-rag-1 (v0.1.0, family rag\_grounding). Every lane received the identical frozen workspace, tool contract and limits; grading ran afterwards, in isolation, against hidden tests no lane ever saw. The purpose is to state, with reproducible evidence, what each lane did and where it stopped.

Outcome

1 of 3

graded rollouts passed the independent hidden tests and were certified — 1 of 4 attempts not graded (lane errors)

1 of 3 graded rollouts passed. cli/codex completed the full loop (inspect, edit, run the visible tests, submit) and the sealed grader accepted the submission. **Low-N:** 4 rollout(s) is below the 10-rollout floor for reading a rate as a rate — read the per-rollout outcomes, not the percentages. The most common way to fail was submitted\_failed — submitted a change; the independent hidden tests still failed.

Key findings

- F-01 cli/codex passed the independent hidden tests
- F-02 submitted\_failed × 2 (cli/claude-sonnet, codestral-latest)
- F-03 model\_api\_failure × 1 (openai/gpt-oss-120b)
- F-04 Statistical floor — 4 rollout(s) is below the 10-rollout rate floor

Each finding is stated in full, with evidence, in the Findings section; recommendations for acting on them follow it. Elapsed wall clock for the whole engagement: 1136.8s.

Exam definition

**In scope.** The ability of each configured model lane to complete this one certified task end to end — inspect the workspace, produce a correct edit, verify it with the visible tests, and submit — within 30 tool steps and 7320s of wall clock, at 1 rollout(s) per lane. Behavioural measurements along the way (tool discipline, grounding, overclaim) are in scope because they are read from the recorded trace, not judged by another model.

**Out of scope.** General model quality, performance on other task families, prompt-tuning on the lane's behalf, and any comparison this run's sample size cannot support. A lane's API failure is reported as infrastructure, never as model capability.

**The instrument.** Exam `vvdex.knowledge.grounded-rag-1` was certified before this run: its reference solution passes, an empty submission fails, its grader distinguishes near-misses, and its sealed evidence is unchanged by this evaluation (see Certification evidence).

What was measured

|                                |         |               |                 |
|--------------------------------|---------|---------------|-----------------|
| EXAM                           | VERSION | FAMILY        | ROLLOUTS        |
| vvdex.knowledge.grounded-rag-1 | 0.1.0   | rag_grounding | 4 (1 per model) |

Where the defect came from

This exam has no upstream repository to credit: the defect, the reference solution and the hidden suite are VVDex's own.

Certification evidence

4 of 5 certification pillars are recorded in this package, loaded from the receipts sealed into the exam and matching this run's exam fingerprint. What is recorded is shown with its real value; what is not is named below and is not claimed.

FROZEN BASELINE

FAIL

expected — an empty submission over 2 run(s) must not pass

→

GOLD / REFERENCE

PASS

expected — the reference solution over 2 run(s) must pass

GRADER CONTROLS

7 / 7

The grader accepted the reference and rejected every near-miss.

ATTACK PROBES

10 / 10 blocked

0 trivial exploit(s) found.

RUNTIME

sha256:679b36225a1f83238efba8c71b9cde3c24caaeacc9b682ae92819cb55eec328f

network policy: deny

Certification metadata is not included in this report package for:

- sealed evidence bundle

Benchmark results remain valid as run evidence; the pillar(s) above are not substantiated by this document and this report does not claim them.

What the package does carry: the exam fingerprint below is the contract digest this run was executed against — a holder of the certification bundle for this fingerprint can join the two independently.

EXAM FINGERPRINT

4880a527ef97ca2f3fe6b57be026c3dc1be51e9e104d79fc13a3e558331dda92

The contract digest this run was executed against.

Evaluation configuration

|                |                                                                     |
|----------------|---------------------------------------------------------------------|
| REPORT         | fr-20260901-718db345                                                |
| ISSUED         | 2026-09-01                                                          |
| SUBJECT EXAM   | vvdex.knowledge.grounded-rag-1 · v0.1.0                             |
| FAMILY         | rag_grounding                                                       |
| PREPARED BY    | VVDex Forge engine vvdex-env 0.1.0                                  |
| CLASSIFICATION | Independent model evaluation report — independently verifiable      |
| CERTIFICATION  | Exam certification: partially recorded (see Certification evidence) |
| STATUS         | FINAL · report sealed against its own digest                        |

Timeline

| WHEN (UTC)          | EVENT                                         |
|---------------------|-----------------------------------------------|
| 2026-09-01 21:30:10 | Evaluation run started                        |
| 2026-09-01 21:49:07 | Evaluation run finished                       |
| 2026-09-01 21:49:07 | Report generated from the sealed run evidence |

Models under test

| MODEL                                  | PROVIDER | ENDPOINT         | BILLING                                            |
|----------------------------------------|----------|------------------|----------------------------------------------------|
| Claude Code CLI (Sonnet, subscription) | cli      | local runtime    | operator subscription, flat — not billed per token |
| Codex CLI (GPT, subscription)          | cli      | local runtime    | operator subscription, flat — not billed per token |
| codestral-latest                       | mistral  | VVDex-configured | VVDex free-quota                                   |
| openai/gpt-oss-120b                    | groq     | VVDex-configured | VVDex free-quota                                   |

A customer endpoint is recorded by origin only — never its port, its path, or its key.

Limits

|                    |                      |                           |                                                                                                        |
|--------------------|----------------------|---------------------------|--------------------------------------------------------------------------------------------------------|
| RUNTIME<br>docker  | STEP LIMIT<br>30     | WALL-CLOCK LIMIT<br>7320s | MAX OUTPUT TOKENS<br>2048                                                                              |
| TEMPERATURE<br>0.2 | RETRY POLICY<br>none | COMMAND BUDGET<br>40      | TOOLS OFFERED<br>list_files, read_file,<br>write_file, edit_file,<br>run_tests, submit,<br>run_command |

The evaluated model never saw the hidden tests or the reference solution, and could not change its result. Grading ran in a separate step, after submission, on a container with no network.

Model outcomes

**Denominators.** Attempts requested: 4. Graded (evaluable) rollouts: 3. Not graded — lane errors: 1. A lane error is a rollout the API or the runtime ended before the hidden grader ran: it is listed, excluded from every rate and pass@k, and never silently dropped. Passes are certified passes: hidden grader exit 0 with integrity verified.

| MODEL                                                              | ATTEMPTS | GRADED | PASSED     | MODEL FAILS | LANE ERRORS | SUBMITTED | MEAN TIME | TYPICAL DEEPEST STAGE |
|--------------------------------------------------------------------|----------|--------|------------|-------------|-------------|-----------|-----------|-----------------------|
| <b>Claude Code CLI (Sonnet, subscription)</b><br>cli/claude-sonnet | 1        | 1      | 0 / 1      | 1           | 0           | 1 / 1     | 71.8s     | Graded                |
| <b>Codex CLI (GPT, subscription)</b><br>cli/codex                  | 1        | 1      | 1 / 1      | 0           | 0           | 1 / 1     | 68.5s     | Passed hidden tests   |
| <b>codestral-latest</b><br>codestral-latest                        | 1        | 1      | 0 / 1      | 1           | 0           | 1 / 1     | 18.7s     | Graded                |
| <b>openai/gpt-oss-120b</b><br>openai/gpt-oss-120b                  | 1        | 0      | not graded | 0           | 1           | —         | —         | nothing               |

Every rollout, with its own deepest stage and grade, is in Appendix A.

Success rate, confidence and pass@k

Every rate on this page is a measurement of a finite number of rollouts, so every rate carries the interval that measurement actually supports. n is the evaluable count — attempts minus rollouts the lane ended before grading; those are listed under "not graded" and sit in no rate, interval or pass@k.

| MODEL               | ATTEMPTS | N | NOT GRADED | PASSES | RATE         | 95% INTERVAL | SEEDS | PASS@1 |
|---------------------|----------|---|------------|--------|--------------|--------------|-------|--------|
| cli/claude-sonnet   | 1        | 1 | 0          | 0      | 0.0%         | [0%, 79%]    | 0     | 0.0%   |
| cli/codex           | 1        | 1 | 0          | 1      | 100.0%       | [21%, 100%]  | 0     | 100.0% |
| codestral-latest    | 1        | 1 | 0          | 0      | 0.0%         | [0%, 79%]    | 0     | 0.0%   |
| openai/gpt-oss-120b | 1        | 0 | 1          | 0      | not recorded | [0%, 100%]   | 0     | —      |

Interval and pass@k methods are stated in full in Appendix D. pass@k is shown for k = 1; the sealed record carries every k up to n.

Model comparison

| MODEL               | PASS@1       | 95% INTERVAL | TOKENS | COST       | MEAN ROLLOUT | OVERCLAIM RATE | TOP FAILURE CLASS |
|---------------------|--------------|--------------|--------|------------|--------------|----------------|-------------------|
| cli/claude-sonnet   | 0.0%         | [0%, 79%]    | n/a    | n/a        | 71.8s        | 0.0%           | submitted_failed  |
| cli/codex           | 100.0%       | [21%, 100%]  | n/a    | n/a        | 68.5s        | 0.0%           | none              |
| codestral-latest    | 0.0%         | [0%, 79%]    | 27 288 | not billed | 18.7s        | 0.0%           | submitted_failed  |
| openai/gpt-oss-120b | not recorded | [0%, 100%]   | n/a    | n/a        | 0ms          | 0.0%           | model_api_failure |

Tokens and cost read n/a where a lane reports no telemetry (a flat-rate CLI) — an absent measurement, not zero. pass@1 and its interval are over evaluable rollouts only.

This is not a leaderboard. Not separated (intervals overlap): cli/claude-sonnet vs cli/codex; cli/claude-sonnet vs codestral-latest; cli/claude-sonnet vs openai/gpt-oss-120b; cli/codex vs codestral-latest; cli/codex vs openai/gpt-oss-120b; codestral-latest vs openai/gpt-oss-120b. Below the 10-rollout floor for reading a rate as a rate: cli/claude-sonnet (1), cli/codex (1), codestral-latest (1), openai/gpt-oss-120b (0) — read the count, not the percentage. The primary comparison in this report is the stage funnel below — where a model stops is a finding; a percentage over this many rollouts is not.

Stage funnel

The funnel is the primary reading of a run. Where a model stops says more than a single pass percentage does, especially at low N.

| STAGE               | CLI:CLI/CLAUDE-SONNET | CLI:CLI/CODEX | MISTRAL:CODESTRAL-LATEST | GROQ:OPENAI/GPT-OSS-120B |
|---------------------|-----------------------|---------------|--------------------------|--------------------------|
| Inspected           | 1 / 1                 | 1 / 1         | 1 / 1                    | 0 / 0                    |
| Edited              | 1 / 1                 | 1 / 1         | 1 / 1                    | 0 / 0                    |
| Ran tests           | 1 / 1                 | 1 / 1         | 1 / 1                    | 0 / 0                    |
| Submitted           | 1 / 1                 | 1 / 1         | 1 / 1                    | 0 / 0                    |
| Graded              | 1 / 1                 | 1 / 1         | 1 / 1                    | 0 / 0                    |
| Passed hidden tests | 0 / 1                 | 1 / 1         | 0 / 1                    | 0 / 0                    |

Deepest stage reached

| MODEL                    | DEEPEST (ANY ATTEMPT) | TYPICAL (MOST ATTEMPTS) | HIDDEN-TEST PASSES | NOT GRADED |
|--------------------------|-----------------------|-------------------------|--------------------|------------|
| cli:cli/claude-sonnet    | Graded                | Graded                  | 0 / 1              | 0          |
| cli:cli/codex            | Passed hidden tests   | Passed hidden tests     | 1 / 1              | 0          |
| mistral:codestral-latest | Graded                | Graded                  | 0 / 1              | 0          |
| groq:openai/gpt-oss-120b | nothing               | nothing                 | 0 / 0              | 1          |

Stage definitions are in Appendix D. The matrix counts evaluable rollouts; lane errors are shown beside it, not inside it. The stages are not strictly nested: a model can submit without editing, and the matrix shows that rather than hiding it. Each rollout's own deepest stage is in Appendix A.

Behavioural analysis

Long-horizon behaviour

Counted off the recorded trace of each rollout. A dead end is a step that moved nothing: a re-read of a slice already returned, a repeat of the previous action, or a call to a tool this exam does not have.

| MODEL                    | N | STEPS USED | RAN OUT OF STEPS | COMMANDS | CHANGED THE TREE | REFUSED (OVER BUDGET) | REVISITS | DEAD ENDS / ROLLOUT |
|--------------------------|---|------------|------------------|----------|------------------|-----------------------|----------|---------------------|
| cli:cli/claude-sonnet    | 1 | 12 / 30    | 0                | 0 / 40   | 0                | 0                     | 0        | 0                   |
| cli:cli/codex            | 1 | 14 / 30    | 0                | 0 / 40   | 0                | 0                     | 0        | 0                   |
| mistral:codestral-latest | 1 | 10 / 30    | 0                | 0 / 40   | 0                | 0                     | 0        | 1                   |
| groq:openai/gpt-oss-120b | 1 | 4 / 30     | 0                | 0 / 40   | 0                | 0                     | 0        | 0                   |

Each column is defined in Appendix D.

Hallucination & overclaim

Three measurements, each with a stated definition. None of them is a judgement of the model's prose by another model — every number below comes from comparing what the model wrote against what was on disk, or against what the independent grader returned.

| MODEL                    | N | GROUNDING FAILURES | KINDS | ROLLOUTS AFFECTED | OVERCLAIMS | OVERCLAIM RATE | RE-READS |
|--------------------------|---|--------------------|-------|-------------------|------------|----------------|----------|
| cli:cli/claude-sonnet    | 1 | 0                  | none  | 0 / 1             | 0 / 1      | 0.0%           | 0        |
| cli:cli/codex            | 1 | 0                  | none  | 0 / 1             | 0 / 1      | 0.0%           | 0        |
| mistral:codestral-latest | 1 | 0                  | none  | 0 / 1             | 0 / 1      | 0.0%           | 0        |
| groq:openai/gpt-oss-120b | 1 | 0                  | none  | 0 / 1             | 0 / 1      | 0.0%           | 0        |

No grounding failures and no false pass claims were detected across these 4 rollout(s). No lane re-read material it had already been given.

Every term above is defined in Appendix D.



Efficiency & telemetry

This run has no per-token price attached — a local model, a free quota, or a customer endpoint we are not billed for — so spend is reported in tokens. Tokens are what was actually measured; a dollar figure here would be invented. Lanes run through a flat-rate local CLI report no token telemetry; their cells say n/a — an absent measurement, not zero. Rollout counts here are evaluable rollouts.

| MODEL                    | ROLLOUTS | TOKENS | TOKENS / ROLLOUT | COST                        | PASSES | RATE   |
|--------------------------|----------|--------|------------------|-----------------------------|--------|--------|
| cli:cli/claude-sonnet    | 1        | n/a    | n/a              | n/a — telemetry unavailable | 0 / 1  | 0.0%   |
| cli:cli/codex            | 1        | n/a    | n/a              | n/a — telemetry unavailable | 1 / 1  | 100.0% |
| mistral:codestral-latest | 1        | 27 288 | 27 288           | not billed                  | 0 / 1  | 0.0%   |
| groq:openai/gpt-oss-120b | 0        | n/a    | n/a              | n/a — telemetry unavailable | 0 / 0  | 0.0%   |

Findings

Every finding below is derived from the recorded data of this run — none is an opinion.  
Severity: PASS a lane succeeded · ATTENTION a capability gap · LANE / INFRA infrastructure, not capability · NOTE a property of the measurement.

F-01

PASS

cli/codex passed the independent hidden tests

Severity: PASS

Evidence: Model outcomes · Stage funnel · Appendix A

1 of 3 graded rollouts completed the full loop (inspect, edit, run the visible tests, submit) and the sealed grader accepted the submission. This demonstrates the task is solvable under the published limits.

F-01 · derived from the recorded run data

F-02

ATTENTION

submitted\_failed × 2 (cli/claude-sonnet, codestral-latest)

Severity: ATTENTION

Evidence: Appendix A · Appendix C

Submitted a change; the independent hidden tests still failed.

F-02 · derived from the recorded run data

F-03

LANE / INFRA

model\_api\_failure × 1 (openai/gpt-oss-120b)

Severity: LANE / INFRA

Evidence: Appendix A · Appendix C

The model endpoint failed to answer. Not a verdict about the model's ability. These rollouts are excluded from any capability reading.

F-03 · derived from the recorded run data

F-04

NOTE

Statistical floor — 4 rollout(s) is below the 10-rollout rate floor

Severity: NOTE

Evidence: Appendix D

The counts are exact; the percentages are anecdotes with decimal points. Per-rollout outcomes, not rates, are the reliable reading of this run.

F-04 · derived from the recorded run data

Recommendations

One recommendation per observed failure mode, addressed to the lane it affects. Each traces back to a finding above and to the raw trace in Appendix A.

R-01

PROBLEM

Submitted a change; the independent hidden tests still failed.

AFFECTED LANES

cli/claude-sonnet, codestral-latest

WHAT WE WOULD LOOK AT FIRST

The model completes the loop and produces a plausible near-miss — the named failing tests are the exact behavioral cases distinguishing it from the reference fix.

R-02

PROBLEM

The model endpoint failed to answer. Not a verdict about the model's ability.

AFFECTED LANES

openai/gpt-oss-120b

WHAT WE WOULD LOOK AT FIRST

The lane failed, not the model; these rollouts carry no signal about capability and are labeled accordingly.

R-03

For anyone reading the percentages

Commission a run at  $n \geq 10$  rollouts per lane before treating any rate here as a rate; this run's per-rollout outcomes are exact, its percentages are not yet stable.

Verification

RUN ID

**live-20260901T213010Z**

EXAM FINGERPRINT

**4880a527ef97ca2f3fe6b57be026c3dc1be51e9e104d79fc13a3e558331dda92**

The contract digest this run was executed against.

Provenance of this chapter

This chapter was rendered by the campaign generator from the sealed evaluation record; it is not the standalone report reproduced byte for byte. Verify the standalone report with its own digest, verify the record with `vvdex-env records verify`, and verify this campaign document as a whole.

RELATIONSHIP

**campaign-rendered chapter**

Derived from the same sealed evaluation record as the standalone report; these bytes are authenticated by the campaign artifact that contains them.

SEALED RECORD DIGEST

**a6426b0b7932b6c0f679f230c44c4b75a30df10a17f7e145b567ac9849d63a36**

sha256 over the record's canonical JSON — the identity every renderer works from.

SOURCE RECORD FILE

**1d818470e53ea53ee6497e9038e19517a4d827d382d3f335018284b31c0e8002**

sha256 of fr-20260901-718db345.json as written beside the standalone report.

SOURCE STANDALONE REPORT

**b287b1728db1ec9c26f6c714a3a7afa404faa12832740c8a9cbd56d0807a9c36**

sha256 of fr-20260901-718db345.html — independently verifiable with its own digest.

SOURCE STANDALONE PDF

**079770731b29ce711d1b3ea3cef453fcb13807cc7318f06647d5b85c0e12de**

sha256 of fr-20260901-718db345.pdf.

RUN EVIDENCE BUNDLE

**718db345e9c729728f9b9272bf0b2ede5774a680273099943b6a7d7bdfce9108**

The digest the run's own bundle computed over itself.

CERTIFIED EVIDENCE SEAL

**not available in this package**

The exam's sealed evidence, established before this run.

Measurement history

This is the first recorded run on this exam. There is nothing to compare it against yet, and a single measurement is not a trend.

The full artifact-by-artifact digest table, and the reproduction protocol, are in Appendix B.

Appendix A — every rollout

| ROLLOUT  | MODEL               | SEED | OUTCOME           | SUBMITTED | DEEPEST STAGE           | HIDDEN TESTS | FILES | TOKENS | ELAPSED |
|----------|---------------------|------|-------------------|-----------|-------------------------|--------------|-------|--------|---------|
| a8b6ca0b | cli/claude-sonnet   | 0    | submitted_failed  | yes       | Graded                  | fail         | 1     | n/a    | 71.8s   |
| f9d23573 | cli/codex           | 0    | submitted_passed  | yes       | Passed hidden tests     | pass         | 1     | n/a    | 68.5s   |
| 2c0b85d4 | codestral-latest    | 0    | submitted_failed  | yes       | Graded                  | fail         | 1     | 27 288 | 18.7s   |
| 97723e61 | openai/gpt-oss-120b | 0    | model_api_failure | no        | Not graded — lane error | not graded   | 0     | 7 583  | 956.6s  |

Failure taxonomy

| CLASS             | COUNT | WHAT IT MEANS                                                                 |
|-------------------|-------|-------------------------------------------------------------------------------|
| submitted_failed  | 2     | Submitted a change; the independent hidden tests still failed.                |
| model_api_failure | 1     | The model endpoint failed to answer. Not a verdict about the model's ability. |

Appendix B — evidence and artifact digests

RUN BUNDLE DIGEST

718db345e9c729728f9b9272

bf0b2ede5774a68027309994

3b6a7d7bdfce9108

RUN ID

live-

20260901T213010Z

STARTED

2026-09-

01T21:30:10.375515+

00:00

FINISHED

2026-09-

01T21:49:07.174353+

00:00

Artifact digests

| ARTIFACT                                           | SHA-256                                                          |
|----------------------------------------------------|------------------------------------------------------------------|
| publicSummary                                      | 0cf64390cffbfc2dd80444c3bfda61950565191526629d2c3793739421d4d2c6 |
| rollouts/2c0b85d4-ba0d-4c0c-8216-ea243bfbaab9.json | 0abbbfe6f3b34f1eb205a6a3b3dbfbe512347c36a6af75aed38d37a69d8c0218 |
| rollouts/97723e61-d56d-4aed-8b61-dee14d643db8.json | efcf0a5acc2b90fe8ec1305a21e14b08ddeee9c076a44e03bf6bad6f10dacbd3 |
| rollouts/a8b6ca0b-1ed7-4a52-a051-1055108b710b.json | 3dda83791f608a375d207c90d6d5a078feb5723d64da05d46f42ff0a518cab6a |
| rollouts/f9d23573-24ee-4f45-89e9-2771b0bafc68.json | 890b8ef9a1178d3c7ecb764f5b06a28aa57a5d3e509c3ed2d3d62e5397571638 |
| spec                                               | 0b760c0ff8444d7143b455fd1603165f7f4adfc0dec6f1c792d573624333536f |
| summary                                            | 19ec6555d12e20cc823c18b06ae11821257eac563b184d022824e4198d81141d |

Reproducing this run

Every rollout can be re-derived: the exam is frozen (fingerprint above), the grader ran in the recorded container image, and each rollout's full trace — every tool call, every result, every word — is in the evidence bundle at its rollouts/<id>.json path. A re-run uses the same job shape: the exam id, your endpoint, n, and the runtime. The record id names this run forever.

Appendix C — redacted transcripts

For each rollout: the tool calls it made, in order — which tool, on which of the model's own paths, whether the call succeeded, and how many characters came back. What a tool returned is not here, and neither is the submission: on this exam the submitted content is the graded answer, so no diff, no answer value and no model claim is printed. A request for the reference solution or the hidden tests is recorded as a request; the path it asked for is replaced by a label. See Appendix D, "What this report withholds and why".

Disclosure class applied to this exam: answer-key disclosure (ownership vvdex\_proprietary — engine registry; family rag\_grounding): the exam is VVDex evaluation IP and the graded output is an answer or a keyed value, so no submission content, diff, model claim or hidden-assertion name is published.

a8b6ca0b · cli:cli/claude-sonnet · seed 0 · submitted\_failed

| STEP | TOOL       | PATH                                 | RESULT | CHARS |
|------|------------|--------------------------------------|--------|-------|
| 0    | list_files | .                                    | ok     | 287   |
| 1    | read_file  | questions.json                       | ok     | 883   |
| 2    | read_file  | corpus/data-protection-policy-v3.txt | ok     | 293   |
| 3    | read_file  | corpus/retention-addendum-2026.txt   | ok     | 142   |
| 4    | read_file  | corpus/org-chart-2026.txt            | ok     | 183   |
| 5    | read_file  | corpus/incident-report-2026.txt      | ok     | 169   |
| 6    | read_file  | corpus/security-whitepaper.txt       | ok     | 142   |
| 7    | read_file  | corpus/employee-handbook.txt         | ok     | 167   |
| 8    | read_file  | corpus/vendor-list.txt               | ok     | 153   |
| 9    | write_file | answers/answers.json                 | ok     | 69    |
| 10   | run_tests  | —                                    | ok     | 181   |
| 11   | submit     | —                                    | ok     | 85    |

Submission diff withheld by the public-disclosure policy (45 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

f9d23573 · cli:cli/codex · seed 0 · submitted\_passed

| STEP | TOOL       | PATH                                 | RESULT | CHARS |
|------|------------|--------------------------------------|--------|-------|
| 0    | read_file  | ISSUE.md                             | ok     | 1 953 |
| 1    | list_files | .                                    | ok     | 287   |
| 2    | read_file  | questions.json                       | ok     | 883   |
| 3    | read_file  | corpus/data-protection-policy-v1.txt | ok     | 239   |
| 4    | read_file  | corpus/data-protection-policy-v3.txt | ok     | 293   |
| 5    | read_file  | corpus/employee-handbook.txt         | ok     | 167   |
| 6    | read_file  | corpus/incident-report-2026.txt      | ok     | 169   |
| 7    | read_file  | corpus/org-chart-2026.txt            | ok     | 183   |
| 8    | read_file  | corpus/retention-addendum-2026.txt   | ok     | 142   |

| STEP | TOOL       | PATH                           | RESULT | CHARS |
|------|------------|--------------------------------|--------|-------|
| 9    | read_file  | corpus/security-whitepaper.txt | ok     | 142   |
| 10   | read_file  | corpus/vendor-list.txt         | ok     | 153   |
| 11   | write_file | answers/answers.json           | ok     | 69    |
| 12   | run_tests  | —                              | ok     | 181   |
| 13   | submit     | —                              | ok     | 85    |

Submission diff withheld by the public-disclosure policy (45 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

#### 2c0b85d4 · mistral:codestral-latest · seed 0 · submitted\_failed

| STEP | TOOL       | PATH                                 | RESULT | CHARS |
|------|------------|--------------------------------------|--------|-------|
| 0    | read_file  | ISSUE.md                             | ok     | 1 953 |
| 1    | read_file  | questions.json                       | ok     | 883   |
| 2    | read_file  | —                                    | failed | 35    |
| 3    | read_file  | —                                    | failed | 35    |
| 4    | list_files | corpus                               | ok     | 249   |
| 5    | read_file  | corpus/data-protection-policy-v3.txt | ok     | 293   |
| 6    | read_file  | corpus/retention-addendum-2026.txt   | ok     | 142   |
| 7    | write_file | answers/answers.json                 | ok     | 69    |
| 8    | run_tests  | —                                    | ok     | 181   |
| 9    | submit     | —                                    | ok     | 85    |

Submission diff withheld by the public-disclosure policy (45 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

#### 97723e61 · groq:openai/gpt-oss-120b · seed 0 · model\_api\_failure

| STEP | TOOL       | PATH           | RESULT | CHARS |
|------|------------|----------------|--------|-------|
| 0    | list_files | .              | ok     | 287   |
| 1    | read_file  | ISSUE.md       | ok     | 1 953 |
| 2    | list_files | .              | ok     | 287   |
| 3    | read_file  | questions.json | ok     | 883   |

Submission diff withheld by the public-disclosure policy. Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.



13 Exam · memory-fact-update-1

vvdex.knowledge.memory-fact-update-1 · record fr-20260901-109cfa3c · record digest 6354cff6fc47eca1... · harness clean · containment clean

Executive summary

|                                       |                                                                          |                                                                                      |                                                                           |
|---------------------------------------|--------------------------------------------------------------------------|--------------------------------------------------------------------------------------|---------------------------------------------------------------------------|
| <div>MODELS TESTED</div> <div>4</div> | <div>CERTIFIED PASSES</div> <div>3</div> <div>of 4 graded rollouts</div> | <div>SUBMITTED</div> <div>4 of 4</div> <div>graded rollouts that called submit</div> | <div>NOT GRADED — LANE ERRORS</div> <div>0</div> <div>of 4 attempts</div> |
|---------------------------------------|--------------------------------------------------------------------------|--------------------------------------------------------------------------------------|---------------------------------------------------------------------------|

Purpose

This report presents a controlled evaluation of 4 model lane(s) against the certified exam vvdex.knowledge.memory-fact-update-1 (v0.1.0, family long\_term\_memory). Every lane received the identical frozen workspace, tool contract and limits; grading ran afterwards, in isolation, against hidden tests no lane ever saw. The purpose is to state, with reproducible evidence, what each lane did and where it stopped.

Outcome

3 of 4

graded rollouts passed the independent hidden tests and were certified

3 of 4 graded rollouts passed. cli/claude-sonnet, cli/codex, openai/gpt-oss-120b completed the full loop (inspect, edit, run the visible tests, submit) and the sealed grader accepted the submission. **Low-N:** 4 rollout(s) is below the 10-rollout floor for reading a rate as a rate — read the per-rollout outcomes, not the percentages. The most common way to fail was submitted\_failed — submitted a change; the independent hidden tests still failed.

Key findings

- F-01 cli/claude-sonnet, cli/codex, openai/gpt-oss-120b passed the independent hidden tests
- F-02 submitted\_failed × 1 (codestral-latest)
- F-03 Statistical floor — 4 rollout(s) is below the 10-rollout rate floor

Each finding is stated in full, with evidence, in the Findings section; recommendations for acting on them follow it. Elapsed wall clock for the whole engagement: 122.3s.

Exam definition

**In scope.** The ability of each configured model lane to complete this one certified task end to end — inspect the workspace, produce a correct edit, verify it with the visible tests, and submit — within 20 tool steps and 4920s of wall clock, at 1 rollout(s) per lane. Behavioural measurements along the way (tool discipline, grounding, overclaim) are in scope because they are read from the recorded trace, not judged by another model.

**Out of scope.** General model quality, performance on other task families, prompt-tuning on the lane's behalf, and any comparison this run's sample size cannot support. A lane's API failure is reported as infrastructure, never as model capability.

**The instrument.** Exam `vvdex.knowledge.memory-fact-update-1` was certified before this run: its reference solution passes, an empty submission fails, its grader distinguishes near-misses, and its sealed evidence is unchanged by this evaluation (see Certification evidence).

What was measured

|                                      |         |                  |                 |
|--------------------------------------|---------|------------------|-----------------|
| EXAM                                 | VERSION | FAMILY           | ROLLOUTS        |
| vvdex.knowledge.memory-fact-update-1 | 0.1.0   | long_term_memory | 4 (1 per model) |

Where the defect came from

This exam has no upstream repository to credit: the defect, the reference solution and the hidden suite are VVDex's own.

Certification evidence

4 of 5 certification pillars are recorded in this package, loaded from the receipts sealed into the exam and matching this run's exam fingerprint. What is recorded is shown with its real value; what is not is named below and is not claimed.

FROZEN BASELINE

FAIL

expected — an empty submission over 2 run(s) must not pass

→

GOLD / REFERENCE

PASS

expected — the reference solution over 2 run(s) must pass

GRADER CONTROLS

7 / 7

The grader accepted the reference and rejected every near-miss.

ATTACK PROBES

10 / 10 blocked

0 trivial exploit(s) found.

RUNTIME

sha256:679b36225a1f83238efba8c71b9cde3c24caaeacc9b682ae92819cb55eec328f

network policy: deny

Certification metadata is not included in this report package for:

- sealed evidence bundle

Benchmark results remain valid as run evidence; the pillar(s) above are not substantiated by this document and this report does not claim them.

What the package does carry: the exam fingerprint below is the contract digest this run was executed against — a holder of the certification bundle for this fingerprint can join the two independently.

EXAM FINGERPRINT

2052eec0ff2d3044dc269e3599ae1bfaa45d0d8cd5df9554796105538d226c21

The contract digest this run was executed against.

Evaluation configuration

|                |                                                                     |
|----------------|---------------------------------------------------------------------|
| REPORT         | fr-20260901-109cfa3c                                                |
| ISSUED         | 2026-09-01                                                          |
| SUBJECT EXAM   | vvdex.knowledge.memory-fact-update-1 · v0.1.0                       |
| FAMILY         | long_term_memory                                                    |
| PREPARED BY    | VVDex Forge engine vvdex-env 0.1.0                                  |
| CLASSIFICATION | Independent model evaluation report — independently verifiable      |
| CERTIFICATION  | Exam certification: partially recorded (see Certification evidence) |
| STATUS         | FINAL · report sealed against its own digest                        |

Timeline

|                     |                                               |
|---------------------|-----------------------------------------------|
| WHEN (UTC)          | EVENT                                         |
| 2026-09-01 21:28:07 | Evaluation run started                        |
| 2026-09-01 21:30:10 | Evaluation run finished                       |
| 2026-09-01 21:30:10 | Report generated from the sealed run evidence |

Models under test

|                                        |          |                  |                                                    |
|----------------------------------------|----------|------------------|----------------------------------------------------|
| MODEL                                  | PROVIDER | ENDPOINT         | BILLING                                            |
| Claude Code CLI (Sonnet, subscription) | cli      | local runtime    | operator subscription, flat — not billed per token |
| Codex CLI (GPT, subscription)          | cli      | local runtime    | operator subscription, flat — not billed per token |
| codestral-latest                       | mistral  | VVDex-configured | VVDex free-quota                                   |
| openai/gpt-oss-120b                    | groq     | VVDex-configured | VVDex free-quota                                   |

A customer endpoint is recorded by origin only — never its port, its path, or its key.

Limits

|                    |                      |                           |                                                                                               |
|--------------------|----------------------|---------------------------|-----------------------------------------------------------------------------------------------|
| RUNTIME<br>docker  | STEP LIMIT<br>20     | WALL-CLOCK LIMIT<br>4920s | MAX OUTPUT TOKENS<br>2048                                                                     |
| TEMPERATURE<br>0.2 | RETRY POLICY<br>none | COMMAND BUDGET<br>24      | TOOLS OFFERED<br>list_files, read_file, write_file, edit_file, run_tests, submit, run_command |

The evaluated model never saw the hidden tests or the reference solution, and could not change its result. Grading ran in a separate step, after submission, on a container with no network.

Model outcomes

**Denominators.** Attempts requested: 4. Graded (evaluable) rollouts: 4. Not graded — lane errors: 0. A lane error is a rollout the API or the runtime ended before the hidden grader ran: it is listed, excluded from every rate and pass@k, and never silently dropped. Passes are certified passes: hidden grader exit 0 with integrity verified.

| MODEL                                                       | ATTEMPTS | GRADED | PASSED | MODEL FAILS | LANE ERRORS | SUBMITTED | MEAN TIME | TYPICAL DEEPEST STAGE |
|-------------------------------------------------------------|----------|--------|--------|-------------|-------------|-----------|-----------|-----------------------|
| Claude Code CLI (Sonnet, subscription)<br>cli/claude-sonnet | 1        | 1      | 1 / 1  | 0           | 0           | 1 / 1     | 18.2s     | Passed hidden tests   |
| Codex CLI (GPT, subscription)<br>cli/codex                  | 1        | 1      | 1 / 1  | 0           | 0           | 1 / 1     | 27.8s     | Passed hidden tests   |
| codestral-latest<br>codestral-latest                        | 1        | 1      | 0 / 1  | 1           | 0           | 1 / 1     | 6.8s      | Graded                |
| openai/gpt-oss-120b<br>openai/gpt-oss-120b                  | 1        | 1      | 1 / 1  | 0           | 0           | 1 / 1     | 48.4s     | Passed hidden tests   |

Every rollout, with its own deepest stage and grade, is in Appendix A.

Success rate, confidence and pass@k

Every rate on this page is a measurement of a finite number of rollouts, so every rate carries the interval that measurement actually supports. n is the evaluable count — attempts minus rollouts the lane ended before grading; those are listed under "not graded" and sit in no rate, interval or pass@k.

| MODEL               | ATTEMPTS | N | NOT GRADED | PASSES | RATE   | 95% INTERVAL | SEEDS | PASS@1 |
|---------------------|----------|---|------------|--------|--------|--------------|-------|--------|
| cli/claude-sonnet   | 1        | 1 | 0          | 1      | 100.0% | [21%, 100%]  | 0     | 100.0% |
| cli/codex           | 1        | 1 | 0          | 1      | 100.0% | [21%, 100%]  | 0     | 100.0% |
| codestral-latest    | 1        | 1 | 0          | 0      | 0.0%   | [0%, 79%]    | 0     | 0.0%   |
| openai/gpt-oss-120b | 1        | 1 | 0          | 1      | 100.0% | [21%, 100%]  | 0     | 100.0% |

Interval and pass@k methods are stated in full in Appendix D. pass@k is shown for k = 1; the sealed record carries every k up to n.

Model comparison

| MODEL               | PASS@1 | 95% INTERVAL | TOKENS | COST       | MEAN ROLLOUT | OVERCLAIM RATE | TOP FAILURE CLASS |
|---------------------|--------|--------------|--------|------------|--------------|----------------|-------------------|
| cli/claude-sonnet   | 100.0% | [21%, 100%]  | n/a    | n/a        | 18.2s        | 0.0%           | none              |
| cli/codex           | 100.0% | [21%, 100%]  | n/a    | n/a        | 27.8s        | 0.0%           | none              |
| codestral-latest    | 0.0%   | [0%, 79%]    | 5 785  | not billed | 6.8s         | 0.0%           | submitted_failed  |
| openai/gpt-oss-120b | 100.0% | [21%, 100%]  | 7 980  | not billed | 48.4s        | 0.0%           | none              |

Tokens and cost read n/a where a lane reports no telemetry (a flat-rate CLI) — an absent measurement, not zero. pass@1 and its interval are over evaluable rollouts only.

This is not a leaderboard. Not separated (intervals overlap): cli/claude-sonnet vs cli/codex; cli/claude-sonnet vs codestral-latest; cli/claude-sonnet vs openai/gpt-oss-120b; cli/codex vs codestral-latest; cli/codex vs openai/gpt-oss-120b; codestral-latest vs openai/gpt-oss-120b. Below the 10-rollout floor for reading a rate as a rate: cli/claude-sonnet (1), cli/codex (1), codestral-latest (1), openai/gpt-oss-120b (1) — read the count, not the percentage. The primary comparison in this report is the stage funnel below — where a model stops is a finding; a percentage over this many rollouts is not.

Stage funnel

The funnel is the primary reading of a run. Where a model stops says more than a single pass percentage does, especially at low N.

| STAGE               | CLI:CLI/CLAUDE-SONNET | CLI:CLI/CODEX | MISTRAL:CODESTRAL-LATEST | GROQ:OPENAI/GPT-OSS-120B |
|---------------------|-----------------------|---------------|--------------------------|--------------------------|
| Inspected           | 1 / 1                 | 1 / 1         | 1 / 1                    | 1 / 1                    |
| Edited              | 1 / 1                 | 1 / 1         | 1 / 1                    | 1 / 1                    |
| Ran tests           | 1 / 1                 | 1 / 1         | 1 / 1                    | 1 / 1                    |
| Submitted           | 1 / 1                 | 1 / 1         | 1 / 1                    | 1 / 1                    |
| Graded              | 1 / 1                 | 1 / 1         | 1 / 1                    | 1 / 1                    |
| Passed hidden tests | 1 / 1                 | 1 / 1         | 0 / 1                    | 1 / 1                    |

Deepest stage reached

| MODEL                    | DEEPEST (ANY ATTEMPT) | TYPICAL (MOST ATTEMPTS) | HIDDEN-TEST PASSES | NOT GRADED |
|--------------------------|-----------------------|-------------------------|--------------------|------------|
| cli:cli/claude-sonnet    | Passed hidden tests   | Passed hidden tests     | 1 / 1              | 0          |
| cli:cli/codex            | Passed hidden tests   | Passed hidden tests     | 1 / 1              | 0          |
| mistral:codestral-latest | Graded                | Graded                  | 0 / 1              | 0          |
| groq:openai/gpt-oss-120b | Passed hidden tests   | Passed hidden tests     | 1 / 1              | 0          |

Stage definitions are in Appendix D. The matrix counts evaluable rollouts; lane errors are shown beside it, not inside it. The stages are not strictly nested: a model can submit without editing, and the matrix shows that rather than hiding it. Each rollout's own deepest stage is in Appendix A.

Behavioural analysis

Long-horizon behaviour

Counted off the recorded trace of each rollout. A dead end is a step that moved nothing: a re-read of a slice already returned, a repeat of the previous action, or a call to a tool this exam does not have.

| MODEL                    | N | STEPS USED | RAN OUT OF STEPS | COMMANDS | CHANGED THE TREE | REFUSED (OVER BUDGET) | REVISITS | DEAD ENDS / ROLLOUT |
|--------------------------|---|------------|------------------|----------|------------------|-----------------------|----------|---------------------|
| cli:cli/claude-sonnet    | 1 | 5 / 20     | 0                | 0 / 24   | 0                | 0                     | 0        | 0                   |
| cli:cli/codex            | 1 | 6 / 20     | 0                | 0 / 24   | 0                | 0                     | 0        | 0                   |
| mistral:codestral-latest | 1 | 4 / 20     | 0                | 0 / 24   | 0                | 0                     | 0        | 0                   |
| groq:openai/gpt-oss-120b | 1 | 6 / 20     | 0                | 0 / 24   | 0                | 0                     | 0        | 0                   |

Each column is defined in Appendix D.

Hallucination & overclaim

Three measurements, each with a stated definition. None of them is a judgement of the model's prose by another model — every number below comes from comparing what the model wrote against what was on disk, or against what the independent grader returned.


| MODEL                    | N | GROUNDING FAILURES | KINDS | ROLLOUTS AFFECTED | OVERCLAIMS | OVERCLAIM RATE | RE-READS |
|--------------------------|---|--------------------|-------|-------------------|------------|----------------|----------|
| cli:cli/claude-sonnet    | 1 | 0                  | none  | 0 / 1             | 0 / 1      | 0.0%           | 0        |
| cli:cli/codex            | 1 | 0                  | none  | 0 / 1             | 0 / 1      | 0.0%           | 0        |
| mistral:codestral-latest | 1 | 0                  | none  | 0 / 1             | 0 / 1      | 0.0%           | 0        |
| groq:openai/gpt-oss-120b | 1 | 0                  | none  | 0 / 1             | 0 / 1      | 0.0%           | 0        |

No grounding failures and no false pass claims were detected across these 4 rollout(s). No lane re-read material it had already been given.

Every term above is defined in Appendix D.

Efficiency & telemetry

This run has no per-token price attached — a local model, a free quota, or a customer endpoint we are not billed for — so spend is reported in tokens. Tokens are what was actually measured; a dollar figure here would be invented. Lanes run through a flat-rate local CLI report no token telemetry; their cells say n/a — an absent measurement, not zero. Rollout counts here are evaluable rollouts.



Only lanes that report token telemetry are charted. The vertical bar through each point is that model's 95% interval, drawn to the same scale as the axis — where it spans most of the axis, the point is one measurement with a wide interval, not a position on a frontier.

| MODEL                    | ROLLOUTS | TOKENS | TOKENS / ROLLOUT | COST                        | PASSES | RATE   |
|--------------------------|----------|--------|------------------|-----------------------------|--------|--------|
| cli:cli/claude-sonnet    | 1        | n/a    | n/a              | n/a – telemetry unavailable | 1 / 1  | 100.0% |
| cli:cli/codex            | 1        | n/a    | n/a              | n/a – telemetry unavailable | 1 / 1  | 100.0% |
| mistral:codestral-latest | 1        | 5 785  | 5 785            | not billed                  | 0 / 1  | 0.0%   |
| groq:openai/gpt-oss-120b | 1        | 7 980  | 7 980            | not billed                  | 1 / 1  | 100.0% |



Findings

Every finding below is derived from the recorded data of this run — none is an opinion.

Severity: PASS a lane succeeded · ATTENTION a capability gap · LANE / INFRA  
infrastructure, not capability · NOTE a property of the measurement.

|                                                                                                                                                                                                                 |                                                      |
|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------|
| F-01                                                                                                                                                                                                            | PASS                                                 |
| cli/claude-sonnet, cli/codex, openai/gpt-oss-120b passed the independent hidden tests                                                                                                                           |                                                      |
| Severity: PASS                                                                                                                                                                                                  | Evidence: Model outcomes · Stage funnel · Appendix A |
| 3 of 4 graded rollouts completed the full loop (inspect, edit, run the visible tests, submit) and the sealed grader accepted the submission. This demonstrates the task is solvable under the published limits. |                                                      |
| F-01 · derived from the recorded run data                                                                                                                                                                       |                                                      |

|                                                                |                                   |
|----------------------------------------------------------------|-----------------------------------|
| F-02                                                           | ATTENTION                         |
| submitted_failed × 1 (codestral-latest)                        |                                   |
| Severity: ATTENTION                                            | Evidence: Appendix A · Appendix C |
| Submitted a change; the independent hidden tests still failed. |                                   |
| F-02 · derived from the recorded run data                      |                                   |

|                                                                                                                                                 |                      |
|-------------------------------------------------------------------------------------------------------------------------------------------------|----------------------|
| F-03                                                                                                                                            | NOTE                 |
| Statistical floor — 4 rollout(s) is below the 10-rollout rate floor                                                                             |                      |
| Severity: NOTE                                                                                                                                  | Evidence: Appendix D |
| The counts are exact; the percentages are anecdotes with decimal points. Per-rollout outcomes, not rates, are the reliable reading of this run. |                      |
| F-03 · derived from the recorded run data                                                                                                       |                      |

Recommendations

One recommendation per observed failure mode, addressed to the lane it affects. Each traces back to a finding above and to the raw trace in Appendix A.

**R-01** PROBLEM  
**Submitted a change; the independent hidden tests still failed.**

AFFECTED LANES  
**codestral-latest**

WHAT WE WOULD LOOK AT FIRST  
The model completes the loop and produces a plausible near-miss — the named failing tests are the exact behavioral cases distinguishing it from the reference fix.

**R-02** For anyone reading the percentages  
Commission a run at  $n \geq 10$  rollouts per lane before treating any rate here as a rate; this run's per-rollout outcomes are exact, its percentages are not yet stable.

Verification

RUN ID

**live-20260901T212807Z**

EXAM FINGERPRINT

**2052eec0ff2d3044dc269e3599ae1bfaa45d0d8cd5df9554796105538d226c21**

The contract digest this run was executed against.

Provenance of this chapter

This chapter was rendered by the campaign generator from the sealed evaluation record; it is not the standalone report reproduced byte for byte. Verify the standalone report with its own digest, verify the record with `vvdex-env records verify`, and verify this campaign document as a whole.

RELATIONSHIP

**campaign-rendered chapter**

Derived from the same sealed evaluation record as the standalone report; these bytes are authenticated by the campaign artifact that contains them.

SEALED RECORD DIGEST

**6354cff6fc47eca13b596d285f805a9f3994d048d69b446166acfeb7f8659929**

sha256 over the record's canonical JSON — the identity every renderer works from.

SOURCE RECORD FILE

**57be625c483af8a4fd88cfae9aaa8a5e6f8678cd89edf7fc6d4bce490c5f405**

sha256 of `fr-20260901-109cfa3c.json` as written beside the standalone report.

SOURCE STANDALONE REPORT

**9c2db3fff15cc6edb9ee1cc71c15a45911a5fdc7e9b88984220898a5f71eb12f**

sha256 of `fr-20260901-109cfa3c.html` — independently verifiable with its own digest.

SOURCE STANDALONE PDF

**1419c72eff2ac1b20c29b732009b25783591a2234a6f5281c6bbfcbce85b77**

sha256 of `fr-20260901-109cfa3c.pdf`.

RUN EVIDENCE BUNDLE

**109cfa3cb6a27c4876a0b1e77b7655db20cbf3d7d5a4bc821f9b5544568c5ff7**

The digest the run's own bundle computed over itself.

CERTIFIED EVIDENCE SEAL

**not available in this package**

The exam's sealed evidence, established before this run.

Measurement history

This is the first recorded run on this exam. There is nothing to compare it against yet, and a single measurement is not a trend.

The full artifact-by-artifact digest table, and the reproduction protocol, are in Appendix B.

Appendix A — every rollout

| ROLLOUT  | MODEL               | SEED | OUTCOME          | SUBMITTED | DEEPEST STAGE       | HIDDEN TESTS | FILES | TOKENS | ELAPSED |
|----------|---------------------|------|------------------|-----------|---------------------|--------------|-------|--------|---------|
| 267a40b8 | cli/claude-sonnet   | 0    | submitted_passed | yes       | Passed hidden tests | pass         | 1     | n/a    | 18.2s   |
| 59ba63e4 | cli/codex           | 0    | submitted_passed | yes       | Passed hidden tests | pass         | 1     | n/a    | 27.8s   |
| 468a0f6c | codestral-latest    | 0    | submitted_failed | yes       | Graded              | fail         | 1     | 5 785  | 6.8s    |
| a2aa98b9 | openai/gpt-oss-120b | 0    | submitted_passed | yes       | Passed hidden tests | pass         | 1     | 7 980  | 48.4s   |

Failure taxonomy

| CLASS            | COUNT | WHAT IT MEANS                                                  |
|------------------|-------|----------------------------------------------------------------|
| submitted_failed | 1     | Submitted a change; the independent hidden tests still failed. |

Appendix B — evidence and artifact digests

RUN BUNDLE DIGEST

109cfa3cb6a27c4876a0b1e77b7655db20cbf3d7d5a4bc821f9b5544568c5ff7

RUN ID

live-20260901T212807Z

STARTED

2026-09-01T21:28:07.881781+00:00

FINISHED

2026-09-01T21:30:10.198797+00:00

Artifact digests

| ARTIFACT                                           | SHA-256                                                          |
|----------------------------------------------------|------------------------------------------------------------------|
| publicSummary                                      | 21cb3b6801e9ba0fbfde7aef29d35f722640e917b4f2e55f25a831f886216304 |
| rollouts/267a40b8-e381-4e92-a3d0-98520dac5c21.json | c288894288ae41ccdcda9328e61981581917d3e3b7825460e473e745b00e6698 |
| rollouts/468a0f6c-89d2-4e79-a529-9e65f466d232.json | a665188e839b85232fe95320387767209c135ed09ab9a1e47919722e92f737e2 |
| rollouts/59ba63e4-1d43-44e2-9e3a-9de7ee90b848.json | 59ac9e2ec584ad5b69bc9268bf9c764416c8e960ea22da3541b721ca8f089f38 |
| rollouts/a2aa98b9-d5c5-4971-8689-b4f8561ac5ff.json | 5720d045e7ff25175a10ade657e521b9e648190596a51ec0ff529f173dffa5ac |
| spec                                               | 672ae47bc67fb407b31949f6e69d215424e4d6d2c03ec838a55924b88776d0e3 |
| summary                                            | d424d32f949d89fca759517f140b1af38e709799822fab8d571768b74c6b93ac |

Reproducing this run

Every rollout can be re-derived: the exam is frozen (fingerprint above), the grader ran in the recorded container image, and each rollout's full trace — every tool call, every result, every word — is in the evidence bundle at its rollouts/<id>.json path. A re-run uses the same job shape: the exam id, your endpoint, n, and the runtime. The record id names this run forever.

Appendix C — redacted transcripts

For each rollout: the tool calls it made, in order — which tool, on which of the model's own paths, whether the call succeeded, and how many characters came back. What a tool returned is not here, and neither is the submission: on this exam the submitted content is the graded answer, so no diff, no answer value and no model claim is printed. A request for the reference solution or the hidden tests is recorded as a request; the path it asked for is replaced by a label. See Appendix D, "What this report withholds and why".

Disclosure class applied to this exam: answer-key disclosure (ownership vvdex\_proprietary — engine registry; family long\_term\_memory): the exam is VVDex evaluation IP and the graded output is an answer or a keyed value, so no submission content, diff, model claim or hidden-assertion name is published.

267a40b8 · cli:cli/claude-sonnet · seed 0 · submitted\_passed

| STEP | TOOL       | PATH                | RESULT | CHARS |
|------|------------|---------------------|--------|-------|
| 0    | read_file  | memory/notes.md     | ok     | 135   |
| 1    | read_file  | facts/current.json  | ok     | 65    |
| 2    | write_file | answers/answer.json | ok     | 68    |
| 3    | run_tests  | —                   | ok     | 181   |
| 4    | submit     | —                   | ok     | 85    |

Submission diff withheld by the public-disclosure policy (9 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

59ba63e4 · cli:cli/codex · seed 0 · submitted\_passed

| STEP | TOOL       | PATH                | RESULT | CHARS |
|------|------------|---------------------|--------|-------|
| 0    | read_file  | ISSUE.md            | ok     | 1 220 |
| 1    | read_file  | facts/current.json  | ok     | 65    |
| 2    | read_file  | memory/notes.md     | ok     | 135   |
| 3    | write_file | answers/answer.json | ok     | 68    |
| 4    | run_tests  | —                   | ok     | 181   |
| 5    | submit     | —                   | ok     | 85    |

Submission diff withheld by the public-disclosure policy (9 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

468a0f6c · mistral:codestral-latest · seed 0 · submitted\_failed

| STEP | TOOL       | PATH                | RESULT | CHARS |
|------|------------|---------------------|--------|-------|
| 0    | read_file  | memory/notes.md     | ok     | 135   |
| 1    | write_file | answers/answer.json | ok     | 68    |
| 2    | run_tests  | —                   | ok     | 181   |
| 3    | submit     | —                   | ok     | 85    |

Submission diff withheld by the public-disclosure policy (9 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

a2aa98b9 · groq:openai/gpt-oss-120b · seed 0 · submitted\_passed

| STEP | TOOL       | PATH                | RESULT | CHARS |
|------|------------|---------------------|--------|-------|
| 0    | list_files | .                   | ok     | 57    |
| 1    | read_file  | facts/current.json  | ok     | 65    |
| 2    | read_file  | memory/notes.md     | ok     | 135   |
| 3    | write_file | answers/answer.json | ok     | 68    |
| 4    | run_tests  | —                   | ok     | 181   |
| 5    | submit     | —                   | ok     | 85    |

Submission diff withheld by the public-disclosure policy (9 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

## 14 Evidence

Every number above is recomputable from the sealed records below; each record verifies itself with `vvdex-env records verify`, and this document was regenerated from them with `vvdex-env records campaign`.

- `public.swe.martinblech-xmltodict-issue-257` → `fr-20260901-a27316a3` · `recordDigest` `96f8b631c6d65f03b21e4a9a94e86c82874e1506fac38c9a093d2f8c618dd2dc`
- `public.swe.pytoolz-toolz-issue-602` → `fr-20260901-88c3cdf1` · `recordDigest` `f49a31336e4865c1dcefc5141e2f6926110c72f03d6cb3d645a1761fa09ac8`
- `public.swe.r1chardj0n3s-parse-issue-102` → `fr-20260901-4b1e3251` · `recordDigest` `d94c313144f8b36db4448440b873462bc9372e1783b1164451695b7b8497816a`
- `vvdex.knowledge.grounded-rag-1` → `fr-20260901-718db345` · `recordDigest` `a6426b0b7932b6c0f679f230c44c4b75a30df10a17f7e145b567ac9849d63a36`
- `vvdex.knowledge.memory-fact-update-1` → `fr-20260901-109cfa3c` · `recordDigest` `6354cff6fc47eca13b596d285f805a9f3994d048d69b446166acfeb7f8659929`



---

## 15 Attestation

---

Certified results appear only where evidence proves them. Withheld, invalid and lane-error cells are never presented as model results, and no historical evidence was modified in producing this document. Counts are exact; below ten rollouts per cell no rate or ranking beyond the counts is asserted.

© VVDex. All rights reserved. Public materials are provided for viewing, evaluation and verification. Reuse, redistribution, derivative benchmark creation, training use, republication or commercial use requires written permission unless a specific file states otherwise. Full terms: [vvdexops.com/terms/](https://vvdexops.com/terms/). Third-party open source reproduced in an exam keeps its own licence.