VVDEX FORGE

EVALUATION ENVIRONMENTS

# Updating a cross-session fact

Answer from the newer authoritative roster, not stale carried memory

EVALUATION REPORT

**2026-09-01**

*Evaluation scope:*

**Claude Code CLI (Sonnet, subscription)**

**Codex CLI (GPT, subscription)**

**codestral-latest**

**openai/gpt-oss-120b**

*Prepared by:*

**VVDex Forge** — generated directly from sealed run evidence

`fr-20260901-109cfa3c` · vvdex.knowledge.memory-fact-update-1 v0.1.0 · Independent model evaluation report

# About this report

## About VVDex Forge

VVDex Forge builds and certifies evaluation environments for AI agents. Every exam is sealed before any model sees it: the reference solution must pass, an empty submission must fail, the grader must reject near-misses, and the hidden tests never enter the agent's tree. Reports and the public catalog live at **vvdexops.com**.

## Report integrity

This report is generated from run evidence; its digests verify authenticity. The sole test of a copy is the three-line procedure in Verification — a copy whose digests do not reproduce is not this report.

## Scope

One exam · 4 rollout(s) · 20 tool steps and 4920s wall clock per rollout, stated in full in Evaluation configuration.

## Boundary

Not a leaderboard. Not a general capability score. A post-seal benchmark: this run stamped no certification gate and rewrote nothing — not the exam, not the reference solution, not the grader, not the certified evidence.

# Contents

## 01 Executive summary

| MODELS TESTED | CERTIFIED PASSES | SUBMITTED | NOT GRADED — LANE ERRORS |
|---|---|---|---|
| 4 | 3<br>of 4 graded rollouts | 4 of 4<br>graded rollouts that called submit | 0<br>of 4 attempts |

### Purpose

This report presents a controlled evaluation of 4 model lane(s) against the certified exam `vvdex.knowledge.memory-fact-update-1` (v0.1.0, family long_term_memory). Every lane received the identical frozen workspace, tool contract and limits; grading ran afterwards, in isolation, against hidden tests no lane ever saw. The purpose is to state, with reproducible evidence, what each lane did and where it stopped.

### Outcome

**3 of 4** graded rollouts passed the independent hidden tests and were certified

3 of 4 graded rollouts passed. cli/claude-sonnet, cli/codex, openai/gpt-oss-120b completed the full loop (inspect, edit, run the visible tests, submit) and the sealed grader accepted the submission. **Low-N:** 4 rollout(s) is below the 10-rollout floor for reading a rate as a rate — read the per-rollout outcomes, not the percentages. The most common way to fail was `submitted_failed` — submitted a change; the independent hidden tests still failed.

### Key findings

F-01 cli/claude-sonnet, cli/codex, openai/gpt-oss-120b passed the independent hidden tests

F-02 submitted_failed × 1 (codestral-latest)

F-03 Statistical floor — 4 rollout(s) is below the 10-rollout rate floor

Each finding is stated in full, with evidence, in the Findings section; recommendations for acting on them follow it. Elapsed wall clock for the whole engagement: 122.3s.

## 02 Exam definition

**In scope.** The ability of each configured model lane to complete this one certified task end to end — inspect the workspace, produce a correct edit, verify it with the visible tests, and submit — within 20 tool steps and 4920s of wall clock, at 1 rollout(s) per lane. Behavioural measurements along the way (tool discipline, grounding, overclaim) are in scope because they are read from the recorded trace, not judged by another model.

**Out of scope.** General model quality, performance on other task families, prompt-tuning on the lane's behalf, and any comparison this run's sample size cannot support. A lane's API failure is reported as infrastructure, never as model capability.

**The instrument.** Exam `vvdex.knowledge.memory-fact-update-1` was certified before this run: its reference solution passes, an empty submission fails, its grader distinguishes near-misses, and its sealed evidence is unchanged by this evaluation (see Certification evidence).

### What was measured

| EXAM | VERSION | FAMILY | ROLLOUTS |
|---|---|---|---|
| vvdex.knowledge.memory-fact-update-1 | 0.1.0 | long_term_memory | 4 (1 per model) |

### Where the defect came from

This exam has no upstream repository to credit: the defect, the reference solution and the hidden suite are VVDex's own.

## 03 Certification evidence

**4 of 5 certification pillars are recorded in this package**, loaded from the receipts sealed into the exam and matching this run's exam fingerprint. What is recorded is shown with its real value; what is not is named below and is not claimed.

FROZEN BASELINE

**FAIL**

expected — an empty submission over 2 run(s) must not pass

→

GOLD / REFERENCE

**PASS**

expected — the reference solution over 2 run(s) must pass

GRADER CONTROLS

**7 / 7**

The grader accepted the reference and rejected every near-miss.

ATTACK PROBES

**10 / 10 blocked**

0 trivial exploit(s) found.

RUNTIME

**sha256:679b36225a1f83238efba8c71b9cde3c24caaeacc9b682ae92819cb55eec328f**

network policy: deny

**Certification metadata is not included in this report package for:**

- sealed evidence bundle

Benchmark results remain valid as run evidence; the pillar(s) above are not substantiated by this document and this report does not claim them.

What the package does carry: the exam fingerprint below is the contract digest this run was executed against — a holder of the certification bundle for this fingerprint can join the two independently.

EXAM FINGERPRINT

`2052eec0ff2d3044dc269e3599ae1bfaa45d0d8cd5df9554796105538d226c21`

The contract digest this run was executed against.

## 04 Evaluation configuration

| | |
|---|---|
| REPORT | fr-20260901-109cfa3c |
| ISSUED | 2026-09-01 |
| SUBJECT EXAM | vvdex.knowledge.memory-fact-update-1 · v0.1.0 |
| FAMILY | long_term_memory |
| PREPARED BY | VVDex Forge engine vvdex-env 0.1.0 |
| CLASSIFICATION | Independent model evaluation report — independently verifiable |
| CERTIFICATION | Exam certification: partially recorded (see Certification evidence) |
| STATUS | FINAL · report sealed against its own digest |

### Timeline

| WHEN (UTC) | EVENT |
|---|---|
| 2026-09-01 21:28:07 | Evaluation run started |
| 2026-09-01 21:30:10 | Evaluation run finished |
| 2026-09-01 21:30:10 | Report generated from the sealed run evidence |

### Models under test

| MODEL | PROVIDER | ENDPOINT | BILLING |
|---|---|---|---|
| Claude Code CLI (Sonnet, subscription) | cli | `local runtime` | operator subscription, flat — not billed per token |
| Codex CLI (GPT, subscription) | cli | `local runtime` | operator subscription, flat — not billed per token |
| codestral-latest | mistral | `VVDex-configured` | VVDex free-quota |
| openai/gpt-oss-120b | groq | `VVDex-configured` | VVDex free-quota |

A customer endpoint is recorded by origin only — never its port, its path, or its key.

### Limits

| RUNTIME | STEP LIMIT | WALL-CLOCK LIMIT | MAX OUTPUT TOKENS |
|---|---|---|---|
| docker | 20 | 4920s | 2048 |

| TEMPERATURE | RETRY POLICY | COMMAND BUDGET | TOOLS OFFERED |
|---|---|---|---|
| 0.2 | none | 24 | list_files, read_file, write_file, edit_file, run_tests, submit, run_command |

The evaluated model never saw the hidden tests or the reference solution, and could not change its result. Grading ran in a separate step, after submission, on a container with no network.

## 05 Model outcomes

**Denominators.** Attempts requested: 4. Graded (evaluable) rollouts: 4. Not graded — lane errors: 0. A lane error is a rollout the API or the runtime ended before the hidden grader ran: it is listed, excluded from every rate and pass@k, and never silently dropped. Passes are certified passes: hidden grader exit 0 with integrity verified.

| MODEL | ATTEMPTS | GRADED | PASSED | MODEL FAILS | LANE ERRORS | SUBMITTED | MEAN TIME | TYPICAL DEEPEST STAGE |
|---|---|---|---|---|---|---|---|---|
| **Claude Code CLI (Sonnet, subscription)** `cli/claude-sonnet` | 1 | 1 | 1 / 1 | 0 | 0 | 1 / 1 | 18.2s | Passed hidden tests |
| **Codex CLI (GPT, subscription)** `cli/codex` | 1 | 1 | 1 / 1 | 0 | 0 | 1 / 1 | 27.8s | Passed hidden tests |
| **codestral-latest** `codestral-latest` | 1 | 1 | 0 / 1 | 1 | 0 | 1 / 1 | 6.8s | Graded |
| **openai/gpt-oss-120b** `openai/gpt-oss-120b` | 1 | 1 | 1 / 1 | 0 | 0 | 1 / 1 | 48.4s | Passed hidden tests |

Every rollout, with its own deepest stage and grade, is in Appendix A.

### Success rate, confidence and pass@k

Every rate on this page is a measurement of a finite number of rollouts, so every rate carries the interval that measurement actually supports. n is the evaluable count — attempts minus rollouts the lane ended before grading; those are listed under "not graded" and sit in no rate, interval or pass@k.

| MODEL | ATTEMPTS | N | NOT GRADED | PASSES | RATE | 95% INTERVAL | SEEDS | PASS@1 |
|---|---|---|---|---|---|---|---|---|
| cli/claude-sonnet | 1 | 1 | 0 | 1 | 100.0% | [21%, 100%] | 0 | 100.0% |
| cli/codex | 1 | 1 | 0 | 1 | 100.0% | [21%, 100%] | 0 | 100.0% |
| codestral-latest | 1 | 1 | 0 | 0 | 0.0% | [0%, 79%] | 0 | 0.0% |
| openai/gpt-oss-120b | 1 | 1 | 0 | 1 | 100.0% | [21%, 100%] | 0 | 100.0% |

Interval and pass@k methods are stated in full in Appendix D. pass@k is shown for k = 1; the sealed record carries every k up to n.

### Model comparison

| MODEL | PASS@1 | 95% INTERVAL | TOKENS | COST | MEAN ROLLOUT | OVERCLAIM RATE | TOP FAILURE CLASS |
|---|---|---|---|---|---|---|---|
| cli/claude-sonnet | 100.0% | [21%, 100%] | n/a | n/a | 18.2s | 0.0% | none |
| cli/codex | 100.0% | [21%, 100%] | n/a | n/a | 27.8s | 0.0% | none |
| codestral-latest | 0.0% | [0%, 79%] | 5 785 | not billed | 6.8s | 0.0% | submitted_failed |
| openai/gpt-oss-120b | 100.0% | [21%, 100%] | 7 980 | not billed | 48.4s | 0.0% | none |

Tokens and cost read n/a where a lane reports no telemetry (a flat-rate CLI) — an absent measurement, not zero. pass@1 and its interval are over evaluable rollouts only.

This is not a leaderboard. Not separated (intervals overlap): cli/claude-sonnet vs cli/codex; cli/claude-sonnet vs codestral-latest; cli/claude-sonnet vs openai/gpt-oss-120b; cli/codex vs codestral-latest; cli/codex vs openai/gpt-oss-120b; codestral-latest vs openai/gpt-oss-120b. Below the 10-rollout floor for reading a rate as a rate: cli/claude-sonnet (1), cli/codex (1), codestral-latest (1), openai/gpt-oss-120b (1) — read the count, not the percentage. The primary comparison in this report is the stage funnel below — where a model stops is a finding; a percentage over this many rollouts is not.

## 06 Stage funnel

The funnel is the primary reading of a run. Where a model stops says more than a single pass percentage does, especially at low N.

| STAGE | CLI:CLI/CLAUDE-SONNET | CLI:CLI/CODEX | MISTRAL:CODESTRAL-LATEST | GROQ:OPENAI/GPT-OSS-120B |
|---|---|---|---|---|
| Inspected | 1 / 1 | 1 / 1 | 1 / 1 | 1 / 1 |
| Edited | 1 / 1 | 1 / 1 | 1 / 1 | 1 / 1 |
| Ran tests | 1 / 1 | 1 / 1 | 1 / 1 | 1 / 1 |
| Submitted | 1 / 1 | 1 / 1 | 1 / 1 | 1 / 1 |
| Graded | 1 / 1 | 1 / 1 | 1 / 1 | 1 / 1 |
| Passed hidden tests | 1 / 1 | 1 / 1 | 0 / 1 | 1 / 1 |

### Deepest stage reached

| MODEL | DEEPEST (ANY ROLLOUT) | TYPICAL (MOST ROLLOUTS) | HIDDEN-TEST PASSES | NOT GRADED |
|---|---|---|---|---|
| cli:cli/claude-sonnet | Passed hidden tests | Passed hidden tests | 1 / 1 | 0 |
| cli:cli/codex | Passed hidden tests | Passed hidden tests | 1 / 1 | 0 |
| mistral:codestral-latest | Graded | Graded | 0 / 1 | 0 |
| groq:openai/gpt-oss-120b | Passed hidden tests | Passed hidden tests | 1 / 1 | 0 |

Stage definitions are in Appendix D. The matrix counts evaluable rollouts; lane errors are shown beside it, not inside it. The stages are not strictly nested: a model can submit without editing, and the matrix shows that rather than hiding it. Each rollout's own deepest stage is in Appendix A.

## 07 Behavioural analysis

### Long-horizon behaviour

Counted off the recorded trace of each rollout. A dead end is a step that moved nothing: a re-read of a slice already returned, a repeat of the previous action, or a call to a tool this exam does not have.

| MODEL | N | STEPS USED | RAN OUT OF STEPS | COMMANDS | CHANGED THE TREE | REFUSED (OVER BUDGET) | REVISITS | DEAD ENDS / ROLLOUT |
|---|---|---|---|---|---|---|---|---|
| cli:cli/claude-sonnet | 1 | 5 / 20 | 0 | 0 / 24 | 0 | 0 | 0 | 0 |
| cli:cli/codex | 1 | 6 / 20 | 0 | 0 / 24 | 0 | 0 | 0 | 0 |
| mistral:codestral-latest | 1 | 4 / 20 | 0 | 0 / 24 | 0 | 0 | 0 | 0 |
| groq:openai/gpt-oss-120b | 1 | 6 / 20 | 0 | 0 / 24 | 0 | 0 | 0 | 0 |

Each column is defined in Appendix D.

### Hallucination & overclaim

Three measurements, each with a stated definition. None of them is a judgement of the model's prose by another model — every number below comes from comparing what the model wrote against what was on disk, or against what the independent grader returned.

| MODEL | N | GROUNDING FAILURES | KINDS | ROLLOUTS AFFECTED | OVERCLAIMS | OVERCLAIM RATE | RE-READS |
|---|---|---|---|---|---|---|---|
| cli:cli/claude-sonnet | 1 | 0 | none | 0 / 1 | 0 / 1 | 0.0% | 0 |
| cli:cli/codex | 1 | 0 | none | 0 / 1 | 0 / 1 | 0.0% | 0 |
| mistral:codestral-latest | 1 | 0 | none | 0 / 1 | 0 / 1 | 0.0% | 0 |
| groq:openai/gpt-oss-120b | 1 | 0 | none | 0 / 1 | 0 / 1 | 0.0% | 0 |

> No grounding failures and no false pass claims were detected across these 4 rollout(s). No lane re-read material it had already been given.

Every term above is defined in Appendix D.

# 08 Efficiency & telemetry

This run has no per-token price attached — a local model, a free quota, or a customer endpoint we are not billed for — so spend is reported in tokens. Tokens are what was actually measured; a dollar figure here would be invented. Lanes run through a flat-rate local CLI report no token telemetry; their cells say n/a — an absent measurement, not zero. Rollout counts here are evaluable rollouts.


Only lanes that report token telemetry are charted. The vertical bar through each point is that model's 95% interval, drawn to the same scale as the axis — where it spans most of the axis, the point is one measurement with a wide interval, not a position on a frontier.

| MODEL | ROLLOUTS | TOKENS | TOKENS / ROLLOUT | COST | PASSES | RATE |
|---|---|---|---|---|---|---|
| cli:cli/claude-sonnet | 1 | n/a | n/a | n/a — telemetry unavailable | 1 / 1 | 100.0% |
| cli:cli/codex | 1 | n/a | n/a | n/a — telemetry unavailable | 1 / 1 | 100.0% |
| mistral:codestral-latest | 1 | 5 785 | 5 785 | not billed | 0 / 1 | 0.0% |
| groq:openai/gpt-oss-120b | 1 | 7 980 | 7 980 | not billed | 1 / 1 | 100.0% |

## 09 Findings

Every finding below is derived from the recorded data of this run — none is an opinion.
Severity: PASS a lane succeeded · ATTENTION a capability gap · LANE / INFRA infrastructure, not capability · NOTE a property of the measurement.

F-01 PASS

**cli/claude-sonnet, cli/codex, openai/gpt-oss-120b passed the independent hidden tests**

| Severity: PASS | Evidence: Model outcomes · Stage funnel · Appendix A |
|---|---|

3 of 4 graded rollouts completed the full loop (inspect, edit, run the visible tests, submit) and the sealed grader accepted the submission. This demonstrates the task is solvable under the published limits.

`F-01 · derived from the recorded run data`

F-02 ATTENTION

**submitted_failed × 1 (codestral-latest)**

| Severity: ATTENTION | Evidence: Appendix A · Appendix C |
|---|---|

Submitted a change; the independent hidden tests still failed.

`F-02 · derived from the recorded run data`

F-03 NOTE

**Statistical floor — 4 rollout(s) is below the 10-rollout rate floor**

| Severity: NOTE | Evidence: Appendix D |
|---|---|

The counts are exact; the percentages are anecdotes with decimal points. Per-rollout outcomes, not rates, are the reliable reading of this run.

`F-03 · derived from the recorded run data`

## 10 Recommendations

One recommendation per observed failure mode, addressed to the lane it affects. Each traces back to a finding above and to the raw trace in Appendix A.

R–01

PROBLEM

**Submitted a change; the independent hidden tests still failed.**

AFFECTED LANES

**codestral-latest**

WHAT WE WOULD LOOK AT FIRST

The model completes the loop and produces a plausible near-miss — the named failing tests are the exact behavioral cases distinguishing it from the reference fix.

R–02

**For anyone reading the percentages**

Commission a run at $n \geq 10$ rollouts per lane before treating any rate here as a rate; this run's per-rollout outcomes are exact, its percentages are not yet stable.

# 11 Verification

RUN ID

live-20260901T212807Z

EXAM FINGERPRINT

**2052eec0ff2d3044dc269e3599ae1bfaa45d0d8cd5df95547961055 38d226c21**

The contract digest this run was executed against.

## Tamper-evidence

Three digests, and how to check each one. If any of them does not reproduce, this page is not the page that was generated for this run.

THIS REPORT

**5a24a8e3f476b9d72ffcd1dc 731403633a165228e93e1bd8 eb401adef3d47373**

sha256 over the report's UTF-8 bytes with every occurrence of the report digest itself replaced by 64 '0' characters

EVAL RECORD

**57be625c483af8a4fd88cfae 9aaa8a5e6f8678cd89edf7fc c6d4bce490c5f405**

sha256 of fr-20260901-109cfa3c.json as written beside this file.

RUN EVIDENCE BUNDLE

**109cfa3cb6a27c4876a0b1e7 7b7655db20cbf3d7d5a4bc82 1f9b5544568c5ff7**

The digest the run's own bundle computed over itself.

CERTIFIED EVIDENCE SEAL

**not recorded**

The exam's sealed evidence, established before this run.

## Measurement history

This is the first recorded run on this exam. There is nothing to compare it against yet, and a single measurement is not a trend.

## Verifying it, in three lines

```
1. vvdex-env verify-report --report fr-20260901-109cfa3c.html
2. or by hand: replace every occurrence of the report digest above with 64 zeros, then sha256 the file — it must
   equal that digest
3. shasum -a 256 fr-20260901-109cfa3c.json must equal the eval-record digest, and the bundle digest must equal
   bundleSha256 in the run's bundle file
```

The full artifact-by-artifact digest table, and the reproduction protocol, are in Appendix B.

## Appendix A — every rollout

| ROLLOUT | MODEL | SEED | OUTCOME | SUBMITTED | DEEPEST STAGE | HIDDEN TESTS | FILES | TOKENS | ELAPSED |
|---|---|---|---|---|---|---|---|---|---|
| 267a40b8 | cli/claude-sonnet | 0 | submitted_passed | yes | Passed hidden tests | pass | 1 | n/a | 18.2s |
| 59ba63e4 | cli/codex | 0 | submitted_passed | yes | Passed hidden tests | pass | 1 | n/a | 27.8s |
| 468a0f6c | codestral-latest | 0 | submitted_failed | yes | Graded | fail | 1 | 5 785 | 6.8s |
| a2aa98b9 | openai/gpt-oss-120b | 0 | submitted_passed | yes | Passed hidden tests | pass | 1 | 7 980 | 48.4s |

### Failure taxonomy

| CLASS | COUNT | WHAT IT MEANS |
|---|---|---|
| submitted_failed | 1 | Submitted a change; the independent hidden tests still failed. |

## Appendix B — evidence and artifact digests

| RUN BUNDLE DIGEST | RUN ID | STARTED | FINISHED |
| --- | --- | --- | --- |
| **109cfa3cb6a27c4876a0b1e77b7655db20cbf3d7d5a4bc821f9b5544568c5ff7** | live-20260901T212807Z | 2026-09-01T21:28:07.881781+00:00 | 2026-09-01T21:30:10.198797+00:00 |

### Artifact digests

| ARTIFACT | SHA-256 |
| --- | --- |
| publicSummary | 21cb3b6801e9ba0fbfde7aef29d35f722640e917b4f2e55f25a831f886216304 |
| rollouts/267a40b8-e381-4e92-a3d0-98520dac5c21.json | c288894288ae41ccdcda9328e61981581917d3e3b7825460e473e745b00e6698 |
| rollouts/468a0f6c-89d2-4e79-a529-9e65f466d232.json | a665188e839b85232fe95320387767209c135ed09ab9a1e47919722e92f737e2 |
| rollouts/59ba63e4-1d43-44e2-9e3a-9de7ee90b848.json | 59ac9e2ec584ad5b69bc9268bf9c764416c8e960ea22da3541b721ca8f089f38 |
| rollouts/a2aa98b9-d5c5-4971-8689-b4f8561ac5ff.json | 5720d045e7ff25175a10ade657e521b9e648190596a51ec0ff529f173dffa5ac |
| spec | 672ae47bc67fb407b31949f6e69d215424e4d6d2c03ec838a55924b88776d0e3 |
| summary | d424d32f949d89fca759517f140b1af38e709799822fab8d571768b74c6b93ac |

### Reproducing this run

Every rollout can be re-derived: the exam is frozen (fingerprint above), the grader ran in the recorded container image, and each rollout's full trace — every tool call, every result, every word — is in the evidence bundle at its `rollouts/<id>.json` path. A re-run uses the same job shape: the exam id, your endpoint, n, and the runtime. The record id names this run forever.

## Appendix C — redacted transcripts

For each rollout: the tool calls it made, in order — which tool, on which of the model's own paths, whether the call succeeded, and how many characters came back. What a tool returned is not here, and neither is the submission: on this exam the submitted content is the graded answer, so no diff, no answer value and no model claim is printed. A request for the reference solution or the hidden tests is recorded as a request; the path it asked for is replaced by a label. See Appendix D, “What this report withholds and why”.

Disclosure class applied to this exam: answer-key disclosure (ownership vvdex_proprietary — engine registry; family long_term_memory): the exam is VVDex evaluation IP and the graded output is an answer or a keyed value, so no submission content, diff, model claim or hidden-assertion name is published.

**267a40b8 · cli:cli/claude-sonnet · seed 0 · submitted_passed**

| STEP | TOOL | PATH | RESULT | CHARS |
|---|---|---|---|---|
| 0 | read_file | memory/notes.md | ok | 135 |
| 1 | read_file | facts/current.json | ok | 65 |
| 2 | write_file | answers/answer.json | ok | 68 |
| 3 | run_tests | – | ok | 181 |
| 4 | submit | – | ok | 85 |

Submission diff withheld by the public-disclosure policy (9 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

**59ba63e4 · cli:cli/codex · seed 0 · submitted_passed**

| STEP | TOOL | PATH | RESULT | CHARS |
|---|---|---|---|---|
| 0 | read_file | ISSUE.md | ok | 1 220 |
| 1 | read_file | facts/current.json | ok | 65 |
| 2 | read_file | memory/notes.md | ok | 135 |
| 3 | write_file | answers/answer.json | ok | 68 |
| 4 | run_tests | – | ok | 181 |
| 5 | submit | – | ok | 85 |

Submission diff withheld by the public-disclosure policy (9 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

**468a0f6c · mistral:codestral-latest · seed 0 · submitted_failed**

| STEP | TOOL | PATH | RESULT | CHARS |
|---|---|---|---|---|
| 0 | read_file | memory/notes.md | ok | 135 |
| 1 | write_file | answers/answer.json | ok | 68 |
| 2 | run_tests | – | ok | 181 |
| 3 | submit | – | ok | 85 |

Submission diff withheld by the public-disclosure policy (9 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

**a2aa98b9 · groq:openai/gpt-oss-120b · seed 0 · submitted_passed**

| STEP | TOOL | PATH | RESULT | CHARS |
|---|---|---|---|---|
| 0 | list_files | . | ok | 57 |
| 1 | read_file | facts/current.json | ok | 65 |
| 2 | read_file | memory/notes.md | ok | 135 |
| 3 | write_file | answers/answer.json | ok | 68 |
| 4 | run_tests | – | ok | 181 |
| 5 | submit | – | ok | 85 |

Submission diff withheld by the public-disclosure policy (9 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

## Appendix D — metric definitions and methods

Statistics, stage funnel and efficiency in this document are derived at render time from the sealed rollout rows and their raw tool events: lane errors leave every capability denominator; a rollout's stages are booleans read off its own event sequence (inspected = a successful list or read, or a successful browser open/snapshot, edited = the workspace changed, tested = run_tests invoked, submitted = submit invoked, graded = the independent grader returned a verdict, passed = certified pass); absent telemetry is n/a. The record's own summary blocks are the values computed when it was sealed; the rows and events are the evidence.

### Denominators

n counts evaluable rollouts only: attempts minus lane errors. A rollout the lane ended (API or runtime failure) is reported but never graded, so it is in no rate, interval or pass@k.

### Interval method

Wilson score interval, not the normal approximation. The normal approximation collapses to a zero-width interval at 0 successes — it would print a 0/2 run as '0%, certain'. For k successes in n trials with $p = k/n$: centre = $(p + z^2/2n) / (1 + z^2/n)$, half-width = $z/(1 + z^2/n) \cdot \sqrt{(p(1-p)/n + z^2/4n^2)}$, clamped to [0, 1], z = 1.9600.

### pass@k

Unbiased estimator: pass@k = $1 - C(n-c, k) / C(n, k)$, for $k \le n$ only — reported only up to the number of rollouts actually run. It answers "given k independent attempts at this task, how often would at least one pass" — nothing more.

### Reading a failed rollout

A failed rollout is a measurement, not a defect report. Model API failures and runtime failures are classified separately precisely because they say nothing about the model's ability on this task.

This run has 4 rollout(s), below the 10 treated as the floor for reading a rate as a rate: the counts are exact, the percentages are anecdotes with decimal points.

### Stage definitions

#### Inspected
Listed or read at least one file in the repository — or, on a browser exam, opened or snapshotted the page.

#### Edited
Wrote at least one file that actually changed on disk.

#### Ran tests
Invoked the visible-test tool at least once.

#### Submitted
Called submit, ending the rollout on its own terms.

#### Graded
The independent hidden grader ran on the submission and returned a verdict.

#### Passed hidden tests
The hidden grader returned exit 0 and the rollout was certified.

### Long-horizon terms

#### commandsRun

Shell commands executed inside the model's own container, against the command budget the contract declares. Only exams that advertise run_command have any.

**deadEnds**

Steps that moved nothing: a re-read, a repeat of the previous action, or a call to a tool this exam does not have. Counted, not judged.

**hitStepLimit**

Rollouts that ran out of steps without submitting. On a long task this is the number that separates 'could not finish' from 'finished wrong'.

**revisits**

Re-reads of a slice the rollout had already been given. On a long task, re-reading is how a model compensates for losing its own earlier context.

**stepsUsed**

Tool calls the rollout actually spent, out of the step limit the contract declares. A rollout that submitted early spent fewer, and that is a result, not a shortfall.

### Grounding and overclaim terms

**groundingFailure**

A file path the model wrote in its own reasoning that does not exist in the tree it was given, or that names a directory it is not allowed to see. Detected by matching source-file paths in the model's text against the agent box on disk — not by judging the prose.

**nonexistent_path**

The model named a source file that is not in the repository it was given. It was reasoning about a file that does not exist.

**overclaim**

The model stated the tests pass or the issue is fixed, and the independent hidden tests then failed. Counted per rollout: claimed-and-wrong, or not.

**reReadChurn**

How many times the model read a file it had already read in the same rollout. It is a cost and a coherence signal, not a failure: re-reading is sometimes right.

**requested_hidden_path**

The model named a path under the reference solution, the hidden grader or the hidden test directory. Those are not in its box; the request could not be served.

### Failure classes

**submitted_failed**

Submitted a change; the independent hidden tests still failed.

**submitted_passed**

Submitted a change that passed the independent hidden tests.

### What this report withholds, and why

On this exam the model's submitted content IS the graded answer, so nothing derived from a submission is published: no diff excerpt, no answer value or answer JSON body, no reference or gold value, no expected hidden output, none of the model's own words, and not the names of the hidden assertions a rollout failed — a grader's name can carry the value it asserts. What Appendix C does publish for each rollout: its id, its lane, its seed, the tool it called at each step, the public-safe path it called the tool on, whether the call succeeded, how many characters came back, the stage it reached and its outcome. A rollout the grader accepted is reported as submitted and accepted — the numbers in this report are unaffected, because the verdict is the grader's, not the reader's.

Rule applied to this exam: answer-key disclosure (ownership vvdex_proprietary — engine registry; family long_term_memory): the exam is VVDex evaluation IP and the graded output is an answer or a keyed value, so no submission content, diff, model claim or hidden-assertion name is published. The rule is a property of the exam family, not a judgement made per document, and the same rule governs every edition of this report, the campaign report that embeds it as a chapter, this exam's public page and every published JSON derived from the same sealed record.

## This exam revision is retired for future evaluation

Retired for future evaluation on 2026-09-02: post-run public disclosure of answer-bearing submission content. Historical sealed results on this revision are unaffected: the disclosure happened after those runs, so it could not have influenced them.

## Attestation

This report was generated end to end by VVDex Forge from the sealed evidence of the evaluation run described above. Measurement values were populated from recorded artifacts, not manually entered. The digests in Verification let any holder of this file confirm that the report, the eval record and the run evidence are the ones this run produced — without contacting VVDex.

**VVDex Forge · vvdexops.com**
Issuer — certified evaluation environments

**2026-09-01**
Date of issue

**vvdex-env 0.1.0**
Engine · verifiers pin c521e4146026

No upstream repository: this exam is VVDex's own.

Post-seal benchmark. It does not stamp any certifying gate and did not rewrite the exam, the gold solution, the grader or the certified evidence.

— END OF REPORT · fr-20260901-109cfa3c —