VVDEX FORGE

EVALUATION ENVIRONMENTS

Streaming items must equal what the full parse would have built

EVALUATION REPORT

2026-09-01

Evaluation scope:

Claude Code CLI (Sonnet, subscription)

Codex CLI (GPT, subscription)

codestral-latest

openai/gpt-oss-120b

Prepared by:

VVDex Forge — generated directly from sealed run evidence

`fr-20260901-a27316a3` · public.swe.martinblech-xmltodict-issue-257 v0.1.0 · Independent model evaluation report

About this report

About VVDex Forge

VVDex Forge builds and certifies evaluation environments for AI agents. Every exam is sealed before any model sees it: the reference solution must pass, an empty submission must fail, the grader must reject near-misses, and the hidden tests never enter the agent's tree. Reports and the public catalog live at **vvdexops.com**.

Report integrity

This report is generated from run evidence; its digests verify authenticity. The sole test of a copy is the three-line procedure in Verification — a copy whose digests do not reproduce is not this report.

Scope

One exam · 4 rollout(s) · 40 tool steps and 9720s wall clock per rollout, stated in full in Evaluation configuration.

Boundary

Not a leaderboard. Not a general capability score. A post-seal benchmark: this run stamped no certification gate and rewrote nothing — not the exam, not the reference solution, not the grader, not the certified evidence.

Contents

01 Executive summary

MODELS TESTED	CERTIFIED PASSES	SUBMITTED
4	2 of 3 graded rollouts	3 of 3 graded rollouts that called submit

NOT GRADED — LANE ERRORS
1 of 4 attempts

Purpose

This report presents a controlled evaluation of 4 model lane(s) against the certified exam `public.swe.martinblech-xmltodict-issue-257` (v0.1.0, family public_oss_swe). Every lane received the identical frozen workspace, tool contract and limits; grading ran afterwards, in isolation, against hidden tests no lane ever saw. The purpose is to state, with reproducible evidence, what each lane did and where it stopped.

Outcome

2 of 3 graded rollouts passed the independent hidden tests and were certified — 1 of 4 attempts not graded (lane errors)

2 of 3 graded rollouts passed. cli/claude-sonnet, cli/codex completed the full loop (inspect, edit, run the visible tests, submit) and the sealed grader accepted the submission. **Low-N:** 4 rollout(s) is below the 10-rollout floor for reading a rate as a rate — read the per-rollout outcomes, not the percentages. The most common way to fail was `model_api_failure` — the model endpoint failed to answer. not a verdict about the model's ability.

Key findings

F-01	cli/claude-sonnet, cli/codex passed the independent hidden tests
F-02	model_api_failure × 1 (openai/gpt-oss-120b)
F-03	submitted_failed × 1 (codestral-latest)
F-04	Statistical floor — 4 rollout(s) is below the 10-rollout rate floor

Each finding is stated in full, with evidence, in the Findings section; recommendations for acting on them follow it. Elapsed wall clock for the whole engagement: 1398.7s.

02 Exam definition

In scope. The ability of each configured model lane to complete this one certified task end to end — inspect the workspace, produce a correct edit, verify it with the visible tests, and submit — within 40 tool steps and 9720s of wall clock, at 1 rollout(s) per lane. Behavioural measurements along the way (tool discipline, grounding, overclaim) are in scope because they are read from the recorded trace, not judged by another model.

Out of scope. General model quality, performance on other task families, prompt-tuning on the lane's behalf, and any comparison this run's sample size cannot support. A lane's API failure is reported as infrastructure, never as model capability.

The instrument. Exam `public.swe.martinblech-xmltodict-issue-257` was certified before this run: its reference solution passes, an empty submission fails, its grader distinguishes near-misses, and its sealed evidence is unchanged by this evaluation (see Certification evidence).

What was measured

EXAM	VERSION	FAMILY
public.swe.martinblech-xmltodict-issue-257	0.1.0	public_oss_swe

ROLLOUTS
4 (1 per model)

Where the defect came from

Source: martinblech/xmltodict (MIT) · issue https://github.com/martinblech/xmltodict/issues/257 · upstream fix https://github.com/martinblech/xmltodict/pull/421

The defect and its upstream fix are public and unmodified. The exam around them — the frozen parent tree, the independently written hidden suite, the control probes and the isolation — is VVDex's, and it is what this report measures.

03 Certification evidence

4 of 5 certification pillars are recorded in this package, loaded from the receipts sealed into the exam and matching this run's exam fingerprint. What is recorded is shown with its real value; what is not is named below and is not claimed.

FROZEN BASELINE

FAIL

expected — an empty submission over 2 run(s) must not pass

→

GOLD / REFERENCE

PASS

expected — the reference solution over 2 run(s) must pass

GRADER CONTROLS

7 / 7

The grader accepted the reference and rejected every near-miss.

ATTACK PROBES

10 / 10 blocked

0 trivial exploit(s) found.

RUNTIME

sha256:679b36225a1f83238efba8c71b9cde3c24caaeacc9b682ae92819cb55eec328f

network policy: deny

Certification metadata is not included in this report package for:

- sealed evidence bundle

Benchmark results remain valid as run evidence; the pillar(s) above are not substantiated by this document and this report does not claim them.

What the package does carry: the exam fingerprint below is the contract digest this run was executed against — a holder of the certification bundle for this fingerprint can join the two independently.

EXAM FINGERPRINT

aa41977120e6ecdddb5c4fc603e8ab23cb272726a86750119969ff4dfc90e297

The contract digest this run was executed against.

04 Model outcomes

Denominators. Attempts requested: 4. Graded (evaluable) rollouts: 3. Not graded — lane errors: 1. A lane error is a rollout the API or the runtime ended before the hidden grader ran: it is listed, excluded from every rate and pass@k, and never silently dropped. Passes are certified passes: hidden grader exit 0 with integrity verified.

MODEL	ATTEMPTS	GRADED	PASSED	MODEL FAILS	LANE ERRORS	SUBMITTED	MEAN TIME	TYPICAL DEEPEST STAGE
Claude Code CLI (Sonnet, subscription) cli/claude-sonnet	1	1	1 / 1	0	0	1 / 1	110.2s	Passed hidden tests

MODEL	ATTEMPTS	GRADED	PASSED	MODEL FAILS	LANE ERRORS	SUBMITTED	MEAN TIME	TYPICAL DEEPEST STAGE
Codex CLI (GPT, subscription) cli/codex	1	1	1 / 1	0	0	1 / 1	40.1s	Passed hidden tests
codestral-latest codestral-latest	1	1	0 / 1	1	0	1 / 1	12.5s	Graded
openai/gpt-oss-120b openai/gpt-oss-120b	1	0	not graded	0	1	–	–	nothing

Every rollout, with its own deepest stage and grade, is in Appendix A.

Success rate, confidence and pass@k

Every rate on this page is a measurement of a finite number of rollouts, so every rate carries the interval that measurement actually supports. n is the evaluable count — attempts minus rollouts the lane ended before grading; those are listed under "not graded" and sit in no rate, interval or pass@k.

MODEL	ATTEMPTS	N	NOT GRADED	PASSES	RATE	95% INTERVAL	SEEDS	PASS@1
cli/claude-sonnet	1	1	0	1	100.0%	[21%, 100%]	0	100.0%
cli/codex	1	1	0	1	100.0%	[21%, 100%]	0	100.0%
codestral-latest	1	1	0	0	0.0%	[0%, 79%]	0	0.0%
openai/gpt-oss-120b	1	0	1	0	not recorded	[0%, 100%]	0	–

Interval and pass@k methods are stated in full in Appendix D. pass@k is shown for k = 1; the sealed record carries every k up to n.

Model comparison

MODEL	PASS@1	95% INTERVAL	TOKENS	COST	MEAN ROLLOUT	OVERCLAIM RATE	TOP FAILURE CLASS
cli/claude-sonnet	100.0%	[21%, 100%]	n/a	n/a	110.2s	0.0%	none
cli/codex	100.0%	[21%, 100%]	n/a	n/a	40.1s	0.0%	none

MODEL	PASS@1	95% INTERVAL	TOKENS	COST	MEAN ROLLOUT	OVERCLAIM RATE	TOP FAILURE CLASS
codestral-latest	0.0%	[0%, 79%]	19 214	not billed	12.5s	0.0%	submitted_failed
openai/gpt-oss-120b	not recorded	[0%, 100%]	n/a	n/a	0ms	0.0%	model_api_failure

Tokens and cost read n/a where a lane reports no telemetry (a flat-rate CLI) — an absent measurement, not zero. pass@1 and its interval are over evaluable rollouts only.

This is not a leaderboard. Not separated (intervals overlap): cli/claude-sonnet vs cli/codex; cli/claude-sonnet vs codestral-latest; cli/claude-sonnet vs openai/gpt-oss-120b; cli/codex vs codestral-latest; cli/codex vs openai/gpt-oss-120b; codestral-latest vs openai/gpt-oss-120b. Below the 10-rollout floor for reading a rate as a rate: cli/claude-sonnet (1), cli/codex (1), codestral-latest (1), openai/gpt-oss-120b (0) — read the count, not the percentage. The primary comparison in this report is the stage funnel below — where a model stops is a finding; a percentage over this many rollouts is not.

05 Stage funnel

The funnel is the primary reading of a run. Where a model stops says more than a single pass percentage does, especially at low N.

STAGE	CLI:CLI/CLAUDE-SONNET	CLI:CLI/CODEX	MISTRAL:CODESTRAL-LATEST	GROQ:OPENAI/GPT-OSS-120B
Inspected	1 / 1	1 / 1	1 / 1	0 / 0
Edited	1 / 1	1 / 1	1 / 1	0 / 0
Ran tests	1 / 1	1 / 1	1 / 1	0 / 0
Submitted	1 / 1	1 / 1	1 / 1	0 / 0
Graded	1 / 1	1 / 1	1 / 1	0 / 0
Passed hidden tests	1 / 1	1 / 1	0 / 1	0 / 0

Deepest stage reached

MODEL	DEEPEST (ANY ROLLOUT)	TYPICAL (MOST ROLLOUTS)	HIDDEN-TEST PASSES	NOT GRADED
cli:cli/claude-sonnet	Passed hidden tests	Passed hidden tests	1 / 1	0
cli:cli/codex	Passed hidden tests	Passed hidden tests	1 / 1	0
mistral:codestral-latest	Graded	Graded	0 / 1	0

MODEL	DEEPEST (ANY ROLLOUT)	TYPICAL (MOST ROLLOUTS)	HIDDEN-TEST PASSES	NOT GRADED
groq:openai/gpt-oss-120b	nothing	nothing	0 / 0	1

Stage definitions are in Appendix D. The matrix counts evaluable rollouts; lane errors are shown beside it, not inside it. The stages are not strictly nested: a model can submit without editing, and the matrix shows that rather than hiding it. Each rollout's own deepest stage is in Appendix A.

06 Findings

Every finding below is derived from the recorded data of this run — none is an opinion. Severity: PASS a lane succeeded · ATTENTION a capability gap · LANE / INFRA infrastructure, not capability · NOTE a property of the measurement.

F-01 PASS

cli/claude-sonnet, cli/codex passed the independent hidden tests

Severity: PASS	Evidence: Model outcomes · Stage funnel · Appendix A

2 of 3 graded rollouts completed the full loop (inspect, edit, run the visible tests, submit) and the sealed grader accepted the submission. This demonstrates the task is solvable under the published limits.

F-01 · derived from the recorded run data

F-02 LANE / INFRA

model_api_failure × 1 (openai/gpt-oss-120b)

Severity: LANE / INFRA	Evidence: Appendix A · Appendix C

The model endpoint failed to answer. Not a verdict about the model's ability. These rollouts are excluded from any capability reading.

F-02 · derived from the recorded run data

F-03 ATTENTION

submitted_failed × 1 (codestral-latest)

Severity: ATTENTION	Evidence: Appendix A · Appendix C

Submitted a change; the independent hidden tests still failed.

F-03 · derived from the recorded run data

F-04 NOTE

Statistical floor — 4 rollout(s) is below the 10-rollout rate floor

Severity: NOTE	Evidence: Appendix D

The counts are exact; the percentages are anecdotes with decimal points. Per-rollout outcomes, not rates, are the reliable reading of this run.

F-04 · derived from the recorded run data

07 Recommendations

One recommendation per observed failure mode, addressed to the lane it affects. Each traces back to a finding above and to the raw trace in Appendix A.

R-01 PROBLEM

The model endpoint failed to answer. Not a verdict about the model's ability.

AFFECTED LANES

openai/gpt-oss-120b

WHAT WE WOULD LOOK AT FIRST

The lane failed, not the model; these rollouts carry no signal about capability and are labeled accordingly.

R-02 PROBLEM

Submitted a change; the independent hidden tests still failed.

AFFECTED LANES

codestral-latest

WHAT WE WOULD LOOK AT FIRST

The model completes the loop and produces a plausible near-miss — the named failing tests are the exact behavioral cases distinguishing it from the reference fix.

R-03 **For anyone reading the percentages**

Commission a run at $n \geq 10$ rollouts per lane before treating any rate here as a rate; this run's per-rollout outcomes are exact, its percentages are not yet stable.

08 Verification

RUN ID

live-20260901T223532Z

EXAM FINGERPRINT

aa41977120e6ecdddb5c4fc603e8ab23cb272726a8

6750119969ff4dfc90e297

The contract digest this run was executed against.

Tamper-evidence

Three digests, and how to check each one. If any of them does not reproduce, this page is not the page that was generated for this run.

THIS REPORT

aa05da417141a6dc4c7fc26059

a031ce7b9a851ee6a40d420f79

f61a089def2c

sha256 over the report's UTF-8 bytes with every occurrence of the report digest itself replaced by 64 '0' characters

EVAL RECORD

cac23cf090b72e918199a09682

44efe9871320c3a2a5a2b9863e

1a0963a4c5b0

sha256 of fr-20260901-a27316a3.json as written beside this file.

RUN EVIDENCE BUNDLE

a27316a347e66d960235e2f5d5

880c5254a556000866fcb91eb0

bb85ed4af5be

The digest the run's own bundle computed over itself.

CERTIFIED EVIDENCE SEAL

not recorded

The exam's sealed evidence, established before this run.

Measurement history

This is the first recorded run on this exam. There is nothing to compare it against yet, and a single measurement is not a trend.

Verifying it, in three lines

1. `vvdex-env verify-report --report fr-20260901-a27316a3.html`
2. `or by hand: replace every occurrence of the report digest above with 64 zeros, then sha256 the file — it must equal that digest`
3. `shasum -a 256 fr-20260901-a27316a3.json must equal the eval-record digest, and the bundle digest must equal bundleSha256 in the run's bundle file`

The full artifact-by-artifact digest table, and the reproduction protocol, are in Appendix B.

This is the executive report. The full verifiable evidence report for the same run — every rollout, artifact digests, redacted transcripts and methods — is `fr-20260901-a27316a3.pdf`, generated from exactly the same sealed record.

Attestation

This report was generated end to end by VVDex Forge from the sealed evidence of the evaluation run described above. Measurement values were populated from recorded artifacts, not manually entered. The digests in Verification let any holder of this file confirm that the report, the eval record and the run evidence are the ones this run produced — without contacting VVDex.

VVDex Forge · vvdexops.com
Issuer — certified evaluation environments

2026-09-01
Date of issue

vvdex-env 0.1.0
Engine · verifiers pin c521e4146026

Source: martinblech/xmltodict (MIT)

Post-seal benchmark. It does not stamp any certifying gate and did not rewrite the exam, the gold solution, the grader or the certified evidence.

— END OF REPORT · fr-20260901-a27316a3 —