

Robot catch temporal and tracking annotation

EVALUATION REPORT

2026-09-07

Evaluation scope:

meta/muse-spark-1.3

Prepared by:

VVDex Forge — generated directly from sealed run evidence

fr-20260907-023b5cdb · vvdex.annotation.robot-video-1 v0.2.0 · Independent model evaluation report

About this report

About VVDex Forge

VVDex Forge builds and certifies evaluation environments for AI agents. Every exam is sealed before any model sees it: the reference solution must pass, an empty submission must fail, the grader must reject near-misses, and the hidden tests never enter the agent's tree. Reports and the public catalog live at **vvdexops.com**.

Report integrity

This report is generated from run evidence; its digests verify authenticity. The sole test of a copy is the three-line procedure in Verification — a copy whose digests do not reproduce is not this report.

Scope

One exam · 10 rollout(s) · 24 tool steps and 2280s wall clock per rollout, stated in full in Evaluation configuration.

Boundary

Not a leaderboard. Not a general capability score. A post-seal benchmark: this run stamped no certification gate and rewrote nothing — not the exam, not the reference solution, not the grader, not the certified evidence.

Contents

01.	Executive summary	4
02.	Exam definition	5
03.	Certification evidence	6
04.	Evaluation configuration	7
05.	Model outcomes	8
06.	Stage funnel	9
07.	Annotation evidence	10
08.	Quality axes	15
09.	Behavioural analysis	16
10.	Efficiency & telemetry	17
11.	Findings	18
12.	Recommendations	19
13.	Verification	20

Appendices: A — every rollout · B — evidence digests · C — redacted transcripts · D — metric definitions · Attestation

01 Executive summary

MODELS TESTED 1	CERTIFIED PASSES 3 of 8 graded rollouts	SUBMITTED 8 of 8 graded rollouts that called submit	NOT GRADED — LANE ERRORS 2 of 10 attempts
---------------------------------------	--	--	--

Purpose

This report presents a controlled evaluation of 1 model lane(s) against the certified exam vv dex.annotation.robot-video-1 (v0.2.0, family annotation). Every lane received the identical frozen workspace, tool contract and limits; grading ran afterwards, in isolation, against hidden tests no lane ever saw. The purpose is to state, with reproducible evidence, what each lane did and where it stopped.

Outcome

3 of 8

graded rollouts passed the independent hidden tests and were certified — 2 of 10 attempts not graded (lane errors)

3 of 8 graded rollouts passed. meta/muse-spark-1.3 completed the applicable annotation workflow and were accepted by the independent annotation grader. The most common way to fail was submitted_failed — submitted annotations rejected by the independent grader.

Key findings

F-01 meta/muse-spark-1.3 passed the independent hidden tests

F-02 submitted_failed × 5 (meta/muse-spark-1.3)

F-03 model_api_failure × 2 (meta/muse-spark-1.3)

Each finding is stated in full, with evidence, in the Findings section; recommendations for acting on them follow it. Elapsed wall clock for the whole engagement: 2824.7s.

02 Exam definition

In scope. The ability of each configured model lane to complete this one certified task end to end — inspect the asset, annotate against the declared rubric, validate, submit, and be accepted by the independent annotation grader — within 24 tool steps and 2280s of wall clock, at 10 rollout(s) per lane. Behavioural measurements along the way (tool discipline, grounding, overclaim) are in scope because they are read from the recorded trace, not judged by another model.

Out of scope. General model quality, performance on other task families, prompt-tuning on the lane's behalf, and any comparison this run's sample size cannot support. A lane's API failure is reported as infrastructure, never as model capability.

The instrument. Exam `vvdex.annotation.robot-video-1` was certified before this run: its reference solution passes, an empty submission fails, its grader distinguishes near-misses, and its sealed evidence is unchanged by this evaluation (see Certification evidence).

What was measured

EXAM	VERSION	FAMILY	ROLLOUTS
vvdex.annotation.robot-video-1	0.2.0	annotation	10 (10 per model)

Where the defect came from

This exam has no upstream repository to credit: the defect, the reference solution and the hidden suite are VVDex's own.

03 Certification evidence

Why this exam is trustworthy. Every pillar below was established before any model saw the exam, loaded from the sealed evidence matching this run's exam fingerprint, and unchanged by this evaluation.

FROZEN BASELINE

FAIL

expected — an empty submission over 2 run(s) must not pass

→

GOLD / REFERENCE

PASS

expected — the reference solution over 2 run(s) must pass

GRADER CONTROLS

19 / 19

The grader accepted the reference and rejected every near-miss.

ATTACK PROBES

9 / 9 blocked

0 trivial exploit(s) found.

SEALED EVIDENCE

verified

Sealed 2026-09-06 over 13 artifact(s).

RUNTIME

sha256:679b36225a1f83238efba8c71b9cde3c24caaeacc9b682ae92819cb55eec328f

network policy: deny

CERTIFICATION BUNDLE DIGEST

0a0a51360cde69737c9b5caef895b8b7c2a5e2d66da89617d8776ac724b019b9

EXAM FINGERPRINT

85e026007af17db692614084...

Full value in Appendix B; the contract digest this run was executed against.

04 Evaluation configuration

REPORT	fr-20260907-023b5cdb
ISSUED	2026-09-07
SUBJECT EXAM	vvdex.annotation.robot-video-1 · v0.2.0
FAMILY	annotation
PREPARED BY	VVDex Forge engine vvdex-env 0.1.0
CLASSIFICATION	Independent model evaluation report — independently verifiable
CERTIFICATION	Exam certification: recorded and passing (see Certification evidence)
STATUS	FINAL · report sealed against its own digest

Timeline

WHEN (UTC)	EVENT
2026-09-06 21:19:43	Exam evidence sealed (before any model saw the exam)
2026-09-07 08:22:33	Evaluation run started
2026-09-07 09:09:38	Evaluation run finished
2026-09-07 09:09:38	Report generated from the sealed run evidence

Models under test

MODEL	PROVIDER	ENDPOINT	BILLING
meta/muse-spark-1.3	openrouter	VVDex-configured	VVDex free-quota

A customer endpoint is recorded by origin only — never its port, its path, or its key.

Limits

RUNTIME docker	STEP LIMIT 24	WALL-CLOCK LIMIT 2280s	MAX OUTPUT TOKENS 2048
TEMPERATURE 0.2	RETRY POLICY http-429	COMMAND BUDGET 0	TOOLS OFFERED list_files, read_file, write_file, edit_file, run_tests, submit

The evaluated model never saw the hidden tests or the reference solution, and could not change its result. Grading ran in a separate step, after submission, on a container with no network.

05 Model outcomes

Denominators. Attempts requested: 10. Graded (evaluable) rollouts: 8. Not graded — lane errors: 2. A lane error is a rollout the API or the runtime ended before the hidden grader ran: it is listed, excluded from every rate and pass@k, and never silently dropped. Passes are certified passes: hidden grader exit 0 with integrity verified.

MODEL	ATTEMPTS	GRADED	PASSED	MODEL FAILS	LANE ERRORS	SUBMITTED	MEAN TIME	TYPICAL DEEPEST STAGE
meta/muse-spark-1.3	10	8	3 / 8	5	2	8 / 8	267.1s	Graded

Every rollout, with its own deepest stage and grade, is in Appendix A.

Success rate, confidence and pass@k

Every rate on this page is a measurement of a finite number of rollouts, so every rate carries the interval that measurement actually supports. n is the evaluable count — attempts minus rollouts the lane ended before grading; those are listed under "not graded" and sit in no rate, interval or pass@k.

MODEL	ATTEMPTS	N	NOT GRADED	PASSES	RATE	95% INTERVAL	SEEDS	PASS@1	PASS@3	PASS@5	PASS@8
meta/muse-spark-1.3	10	8	2	3	37.5%	[14%, 69%]	0–9	37.5%	82.1%	98.2%	100.0%

Interval and pass@k methods are stated in full in Appendix D. pass@k is shown for k = 1, 3, 5, 8; the sealed record carries every k up to n.

Model comparison

One model ran in this record, so there is nothing to compare it with. A comparison table across models appears here when a run covers more than one.

06 Stage funnel

Recorded annotation actions and independent grader verdicts. Missing capabilities are N/A; each stage is measured independently.

STAGE	OPENROUTER:META/MUSE-SPARK-1.3
Inspect	8 / 8
Annotate	8 / 8
Validate	8 / 8
Submit	8 / 8
Graded	8 / 8
Passed	3 / 8

Deepest stage reached

MODEL	DEEPEST (ANY ATTEMPT)	TYPICAL (MOST ATTEMPTS)	ANNOTATION QA PASSES	NOT GRADED
openrouter:meta/muse-spark-1.3	Passed	Graded	3 / 8	2

The matrix counts evaluable annotation attempts; lane errors are shown separately. Successful recorded actions establish each stage independently.

07 Annotation evidence

Modality: video. **Rubric:** 1.1.

A robot-catching video is reviewed for temporal events, frame-level object boxes, track identity and occlusion.

Submission dd120ba7-a5d4-42fd-8c4a-7c312d18a602

Score	Items	Valid / invalid	Types
0.7638	6	6 / 0	box, classification, event, segment

Measured metrics

annotationCount	6	boxF1	1	boxIoU	0.579
boxPrecision	1	boxRecall	1	classificationAccuracy	1
classificationF1	1	classificationPrecision	1	classificationRecall	1
eventF1	1	eventPrecision	1	eventRecall	1
eventTemporalIoU	0.7125	matchedCount	6	mismatchCount	4
segmentF1	1	segmentPrecision	1	segmentRecall	1
segmentTemporalIoU	1	Track consistency on matched boxes	1		

Track consistency depends on successful box matching. A zero can reflect inaccurate boxes even when track IDs are correct.

Validation error codes

None recorded

Disagreement

Agreement reviewNot measured

Submission cc1b55f2-6d54-436d-8737-aa81aae3a220

Score	Items	Valid / invalid	Types
0.7003	6	6 / 0	box, classification, event, segment

Measured metrics

annotationCount	6	boxF1	0.5	boxIoU	0.5528
boxPrecision	0.5	boxRecall	0.5	classificationAccuracy	1
classificationF1	1	classificationPrecision	1	classificationRecall	1
eventF1	1	eventPrecision	1	eventRecall	1

eventTemporalIoU	0.5982	matchedCount	6	mismatchCount	5
segmentF1	1	segmentPrecision	1	segmentRecall	1
segmentTemporalIoU	0.9	Track consistency on matched boxes	1		

Track consistency depends on successful box matching. A zero can reflect inaccurate boxes even when track IDs are correct.

Validation error codes

None recorded

Disagreement

Agreement reviewNot measured

Submission 415869af-1240-4a10-bf95-d12fccfcb96c

Score	Items	Valid / invalid	Types
0.7802	6	6 / 0	box, classification, event, segment

Measured metrics

annotationCount	6	boxF1	0	boxIoU	0.4656
boxPrecision	0	boxRecall	0	classificationAccuracy	1
classificationF1	1	classificationPrecision	1	classificationRecall	1
eventF1	1	eventPrecision	1	eventRecall	1
eventTemporalIoU	0.875	matchedCount	6	mismatchCount	3
segmentF1	1	segmentPrecision	1	segmentRecall	1
segmentTemporalIoU	1	Track consistency on matched boxes	1		

Track consistency depends on successful box matching. A zero can reflect inaccurate boxes even when track IDs are correct.

Validation error codes

None recorded

Disagreement

Agreement reviewNot measured

Submission 8fe2d8d3-b887-4b2f-ae75-a28eda0bb8be

Score	Items	Valid / invalid	Types
0.6654	6	6 / 0	box, classification, event, segment

Measured metrics

annotationCount	6	boxF1	1	boxIoU	0.5797
boxPrecision	1	boxRecall	1	classificationAccuracy	1
classificationF1	1	classificationPrecision	1	classificationRecall	1
eventF1	0.5	eventPrecision	0.5	eventRecall	0.5
eventTemporalIoU	0.4667	matchedCount	6	mismatchCount	5
segmentF1	1	segmentPrecision	1	segmentRecall	1
segmentTemporalIoU	0.9	Track consistency on matched boxes	1		

Track consistency depends on successful box matching. A zero can reflect inaccurate boxes even when track IDs are correct.

Validation error codes

None recorded

Disagreement

Agreement review Not measured

Submission 0dbb95df-4a7f-438d-b3c3-74a1e68017eb

Score	Items	Valid / invalid	Types
0.7908	6	6 / 0	box, classification, event, segment

Measured metrics

annotationCount	6	boxF1	0.5	boxIoU	0.4975
boxPrecision	0.5	boxRecall	0.5	classificationAccuracy	1
classificationF1	1	classificationPrecision	1	classificationRecall	1
eventF1	1	eventPrecision	1	eventRecall	1
eventTemporalIoU	0.875	matchedCount	6	mismatchCount	3
segmentF1	1	segmentPrecision	1	segmentRecall	1
segmentTemporalIoU	1	Track consistency on matched boxes	1		

Track consistency depends on successful box matching. A zero can reflect inaccurate boxes even when track IDs are correct.

Validation error codes

None recorded

Disagreement

Agreement review Not measured

Submission 078d6f61-7dd8-4c2a-8f74-eb4049a5b391

Score	Items	Valid / invalid	Types
0.8217	6	6 / 0	box, classification, event, segment

Measured metrics

annotationCount	6	boxF1	1	boxIoU	0.5902
boxPrecision	1	boxRecall	1	classificationAccuracy	1
classificationF1	1	classificationPrecision	1	classificationRecall	1
eventF1	1	eventPrecision	1	eventRecall	1
eventTemporalIoU	0.875	matchedCount	6	mismatchCount	3
segmentF1	1	segmentPrecision	1	segmentRecall	1
segmentTemporalIoU	1	Track consistency on matched boxes	1		

Track consistency depends on successful box matching. A zero can reflect inaccurate boxes even when track IDs are correct.

Validation error codes

None recorded

Disagreement

Agreement review	Not measured
------------------	--------------

Submission 2efda1c5-b89e-4014-a4cf-161e5c35599f

Score	Items	Valid / invalid	Types
0.6826	6	6 / 0	box, classification, event, segment

Measured metrics

annotationCount	6	boxF1	0	boxIoU	0.459
boxPrecision	0	boxRecall	0	classificationAccuracy	1
classificationF1	1	classificationPrecision	1	classificationRecall	1
eventF1	1	eventPrecision	1	eventRecall	1
eventTemporalIoU	0.6125	matchedCount	6	mismatchCount	5
segmentF1	1	segmentPrecision	1	segmentRecall	1
segmentTemporalIoU	0.9524	Track consistency on matched boxes	1		

Track consistency depends on successful box matching. A zero can reflect inaccurate boxes even when track IDs are correct.

Validation error codes	
None recorded	
Disagreement	
Agreement review	Not measured

Submission be8a7b18-8bd0-4ebd-9588-0211ca846bc1

Score	Items	Valid / invalid	Types
0.7473	6	6 / 0	box, classification, event, segment

Measured metrics

annotationCount	6	boxF1	1	boxIoU	0.5295
boxPrecision	1	boxRecall	1	classificationAccuracy	1
classificationF1	1	classificationPrecision	1	classificationRecall	1
eventF1	1	eventPrecision	1	eventRecall	1
eventTemporalIoU	0.7125	matchedCount	6	mismatchCount	4
segmentF1	1	segmentPrecision	1	segmentRecall	1
segmentTemporalIoU	1	Track consistency on matched boxes	1		

Track consistency depends on successful box matching. A zero can reflect inaccurate boxes even when track IDs are correct.

Validation error codes	
None recorded	
Disagreement	
Agreement review	Not measured

Measured grader aggregates only. Missing metrics or disagreement are not measured, never zero. Reference answers and item-level labels are excluded. Certification and the evidence digest are recorded in this report's verification sections.

08 Quality axes

Each axis was measured inside the grade box, on the submitted workspace, by the same deterministic script for every rollout. The reward is the conjunction of the required axes; the score is a weighted reading beside it, not the gate.

AXIS	GATE	ROLLOUTS PASSING	MEASURED	BUDGET	WEIGHT	WHAT IT MEASURES
boxPrecision	recorded	N/A	0.625	–	–	
boxRecall	recorded	N/A	0.625	–	–	
boxF1	recorded	N/A	0.625	–	–	
boxIoU	recorded	N/A	0.5317	–	–	
classificationPrecision	recorded	N/A	1	–	–	
classificationRecall	recorded	N/A	1	–	–	
classificationF1	recorded	N/A	1	–	–	
classificationAccuracy	recorded	N/A	1	–	–	
eventPrecision	recorded	N/A	0.9375	–	–	
eventRecall	recorded	N/A	0.9375	–	–	
eventF1	recorded	N/A	0.9375	–	–	
eventTemporalIoU	recorded	N/A	0.7159	–	–	
segmentPrecision	recorded	N/A	1	–	–	
segmentRecall	recorded	N/A	1	–	–	
segmentF1	recorded	N/A	1	–	–	
segmentTemporalIoU	recorded	N/A	0.969	–	–	
trackingConsistency	recorded	N/A	1	–	–	

QUALITY SCORE (MEAN)

0.744

Weighted reading over the axes. It is not the reward.

REWARD COMPOSITION

Withheld from the public report.

09 Behavioural analysis

Long-horizon behaviour

Counted off the recorded trace of each rollout. A dead end is a step that moved nothing: a re-read of a slice already returned, a repeat of the previous action, or a call to a tool this exam does not have.

MODEL	N	STEPS USED	RAN OUT OF STEPS	COMMANDS	CHANGED THE TREE	REFUSED (OVER BUDGET)	REVISITS	DEAD ENDS / ROLLOUT
openrouter:meta/muse-spark-1.3	10	10 / 24	0	0 / 0	0	0	4	1

Each column is defined in Appendix D.

Hallucination & overclaim

Three measurements, each with a stated definition. None of them is a judgement of the model's prose by another model — every number below comes from comparing what the model wrote against what was on disk, or against what the independent grader returned.

MODEL	N	GROUNDING FAILURES	KINDS	ROLLOUTS AFFECTED	OVERCLAIMS	OVERCLAIM RATE	RE-READS
openrouter:meta/muse-spark-1.3	10	0	none	0 / 10	0 / 10	0.0%	4

No grounding failures and no false pass claims were detected across these 10 rollout(s). openrouter:meta/muse-spark-1.3 re-read previously returned material 4 time(s); other lanes did not.

Every term above is defined in Appendix D.

10 Efficiency & telemetry

This run has no per-token price attached — a local model, a free quota, or a customer endpoint we are not billed for — so spend is reported in tokens. Tokens are what was actually measured; a dollar figure here would be invented. Lanes run through a flat-rate local CLI report no token telemetry; their cells say n/a — an absent measurement, not zero. Rollout counts here are evaluable rollouts.

MODEL	ROLLOUTS	TOKENS	TOKENS / ROLLOUT	COST	PASSES	RATE
openrouter:meta/muse-spark-1.3	8	1 073 893	134 236	not billed	3 / 8	37.5%

11 Findings

Every finding below is derived from the recorded data of this run — none is an opinion.
Severity: PASS a lane succeeded · ATTENTION a capability gap · LANE / INFRA
infrastructure, not capability · NOTE a property of the measurement.

F-01

PASS

meta/muse-spark-1.3 passed the independent hidden tests

Severity: PASS

Evidence: Model outcomes · Stage funnel · Appendix A

3 of 8 graded rollouts completed the applicable annotation workflow and were accepted by the independent annotation grader. This demonstrates the task is solvable under the published limits.

F-01 · derived from the recorded run data

F-02

ATTENTION

submitted_failed × 5 (meta/muse-spark-1.3)

Severity: ATTENTION

Evidence: Appendix A · Appendix C

Submitted annotations rejected by the independent grader.

F-02 · derived from the recorded run data

F-03

LANE / INFRA

model_api_failure × 2 (meta/muse-spark-1.3)

Severity: LANE / INFRA

Evidence: Appendix A · Appendix C

The model endpoint failed to answer. Not a verdict about the model's ability. These rollouts are excluded from any capability reading.

F-03 · derived from the recorded run data

12 Recommendations

One recommendation per observed failure mode, addressed to the lane it affects. Each traces back to a finding above and to the raw trace in Appendix A.

R-01

PROBLEM

Submitted annotations rejected by the independent grader.

AFFECTED LANES

meta/muse-spark-1.3

WHAT WE WOULD LOOK AT FIRST

Review the measured annotation metrics and validator error codes against the declared rubric.

R-02

PROBLEM

The model endpoint failed to answer. Not a verdict about the model's ability.

AFFECTED LANES

meta/muse-spark-1.3

WHAT WE WOULD LOOK AT FIRST

The lane failed, not the model; these rollouts carry no signal about capability and are labeled accordingly.

13 Verification

RUN ID

live-20260907T082233Z

EXAM FINGERPRINT

85e026007af17db692614084f12ff40c5bbf2348bcfda0a01deb95aad8928f4c

The contract digest this run was executed against.

Tamper-evidence

Three digests, and how to check each one. If any of them does not reproduce, this page is not the page that was generated for this run.

THIS REPORT

8e62ff0688553d21ef90f4e30642e5a1d2bfdd02cd4182126eda9587bc0f1432

sha256 over the report's UTF-8 bytes with every occurrence of the report digest itself replaced by 64 '0' characters

EVAL RECORD

079d54cc556ec585dc806c66514f025eca04c9325bf64854e9421669b3f09056

sha256 of fr-20260907-023b5cdb.json as written beside this file.

RUN EVIDENCE BUNDLE

023b5cdb0bec7cdb5402199e372974bbbaacea8cfea639dc9fbf118bbcef9d1b

The digest the run's own bundle computed over itself.

CERTIFIED EVIDENCE SEAL

0a0a51360cde69737c9b5caef895b8b7c2a5e2d66da89617d8776ac724b019b9

The exam's sealed evidence, established before this run.

Measurement history

This is the first recorded run on this exam. There is nothing to compare it against yet, and a single measurement is not a trend.

Verifying it, in three lines

1. vvdex-env verify-report --report fr-20260907-023b5cdb.html

2. or by hand: replace every occurrence of the report digest above with 64 zeros, then sha256 the file – it must equal that digest

3. shasum -a 256 fr-20260907-023b5cdb.json must equal the eval-record digest, and the bundle digest must equal bundleSha256 in the run's bundle file

The full artifact-by-artifact digest table, and the reproduction protocol, are in Appendix B.

Appendix A — every rollout

ROLLOUT	MODEL	SEED	OUTCOME	SUBMITTED	DEEPEST STAGE	HIDDEN TESTS	FILES	TOKENS	ELAPSED
dd120ba7	meta/muse-spark-1.3	0	submitted_passed	yes	Passed	pass	3	129 862	279.4s
cc1b55f2	meta/muse-spark-1.3	1	submitted_failed	yes	Graded	fail	4	178 087	333.7s
415869af	meta/muse-spark-1.3	2	submitted_failed	yes	Graded	fail	3	133 056	244.8s
8fe2d8d3	meta/muse-spark-1.3	3	submitted_failed	yes	Graded	fail	3	172 970	290.5s
da93e045	meta/muse-spark-1.3	4	model_api_failure	no	Not graded — lane error	not graded	6	n/a	373.3s
0dbb95df	meta/muse-spark-1.3	5	submitted_failed	yes	Graded	fail	2	102 569	250.4s
078d6f61	meta/muse-spark-1.3	6	submitted_passed	yes	Passed	pass	2	112 954	213.4s
2efda1c5	meta/muse-spark-1.3	7	submitted_failed	yes	Graded	fail	3	129 912	254.6s
1eb52e4b	meta/muse-spark-1.3	8	model_api_failure	no	Not graded — lane error	not graded	5	n/a	311.8s
be8a7b18	meta/muse-spark-1.3	9	submitted_passed	yes	Passed	pass	3	114 483	270.3s

Failure taxonomy

CLASS	COUNT	WHAT IT MEANS
submitted_failed	5	Submitted annotations rejected by the independent grader.
model_api_failure	2	The model endpoint failed to answer. Not a verdict about the model's ability.

Appendix B — evidence and artifact digests

RUN BUNDLE DIGEST

023b5cdb0bec7cdb5402199e

372974bbbaacea8cfea639dc

9fbf118bbcef9d1b

RUN ID

live-

20260907T082233Z

STARTED

2026-09-

07T08:22:33.271069

+00:00

FINISHED

2026-09-

07T09:09:38.009136

+00:00

Artifact digests

ARTIFACT	SHA-256
publicSummary	c50ed23ef6b4c129b4232edc2fd8661c684e14943f9775b07eb60fdf1235beda
rollouts/078d6f61-7dd8-4c2a-8f74-eb4049a5b391.json	21d1d4ac0afff1bdcd7f1ece175225fc0b05e3f29a9263ab785ac8d7c57f6e45
rollouts/0dbb95df-4a7f-438d-b3c3-74a1e68017eb.json	2d255e71772a144234665a7fb3b8d7c1dce99fdc37f7827b56493b04d13f1159
rollouts/1eb52e4b-3662-4fe9-80e5-445d7222a8cf.json	efa0138630ce73d80c14d0913a32003d43cb3d5bfce89d90f0d66b4f0b32ae96
rollouts/2efda1c5-b89e-4014-a4cf-161e5c35599f.json	92a6ccd12bff499e88854d3cb385d6f23cd9d51cc5616759c239389cb7221112
rollouts/415869af-1240-4a10-bf95-d12fccfcb96c.json	a9ab8a3ebfaf30005278931ed8b3750b765989fa5fc38ac9a1f4e57a36bbacd0
rollouts/8fe2d8d3-b887-4b2f-ae75-a28eda0bb8be.json	5e7ff9c2e01d5e1410ad76fbdcce0a9ad12f3d78db7414e8c79374850f05c528
rollouts/be8a7b18-8bd0-4ebd-9588-0211ca846bc1.json	4f4731b37f62645d23028f5e25536acbbd43fd281337aa2eb0d9ba7641090765
rollouts/cc1b55f2-6d54-436d-8737-aa81aae3a220.json	ee29ada9d5ba47c6b5f7c0424aab712f492b09a545a09bfe8b3de65f9bfe1b74
rollouts/da93e045-0855-4ec9-a6fc-5ea7a6963d22.json	6a0deba95998a940dff94c35ae0838f458b4f707c149033078d8e2cc12ba5c99
rollouts/dd120ba7-a5d4-42fd-8c4a-7c312d18a602.json	79727bf726edd7458694df404ffac925fba97bf93ea83134ae9c24e68390ccfb
spec	13a78198f4184d167af717ac09b899f9c9b07899b2a73f4f0c68fb4f4005ed6e
summary	5db612d49a6f66d4f0ee16b77f60e86f6c1f4df368ddd0dc1fe7f05ff01a7c42

Reproducing this run

Every rollout can be re-derived: the exam is frozen (fingerprint above), the grader ran in the recorded container image, and each rollout's full trace — every tool call, every result, every word — is in the evidence bundle at its rollouts/<id>.json path. A re-run uses the same job shape: the exam id, your endpoint, n, and the runtime. The record id names this run forever.

Appendix C — redacted transcripts

For each rollout: the tool calls it made, in order — which tool, on which of the model's own paths, whether the call succeeded, and how many characters came back. What a tool returned is not here, and neither is the submission: on this exam the submitted content is the graded answer, so no diff, no answer value and no model claim is printed. A request for the reference solution or the hidden tests is recorded as a request; the path it asked for is replaced by a label. See Appendix D, "What this report withholds and why".

Disclosure class applied to this exam: answer-key disclosure (ownership vvdex_proprietary — no declaration — proprietary by default; family annotation): the exam is VVDex evaluation IP and the graded output is an answer or a keyed value, so no submission content, diff, model claim or hidden-assertion name is published.

dd120ba7 · openrouter:meta/muse-spark-1.3 · seed 0 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	list_files	.	ok	79
1	read_file	TASK.md	ok	1 119
2	read_file	annotation-task.json	ok	4 867
3	read_file	submission.json	ok	74
4	write_file	submission.json	ok	64
5	edit_file	submission.json	ok	53
6	run_tests	—	ok	160
7	edit_file	submission.json	ok	53
8	run_tests	—	ok	160
9	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (9 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

cc1b55f2 · openrouter:meta/muse-spark-1.3 · seed 1 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	list_files	.	ok	79
1	read_file	TASK.md	ok	1 119
2	read_file	annotation-task.json	ok	4 867
3	read_file	submission.json	ok	74
4	write_file	submission.json	ok	64
5	edit_file	submission.json	ok	53
6	edit_file	submission.json	ok	53
7	run_tests	—	ok	160
8	list_files	assets	ok	17
9	edit_file	submission.json	ok	53

STEP	TOOL	PATH	RESULT	CHARS
10	run_tests	—	ok	160
11	run_tests	—	ok	160
12	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (9 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

415869af · openrouter:meta/muse-spark-1.3 · seed 2 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	list_files	.	ok	79
1	read_file	TASK.md	ok	1 119
2	read_file	annotation-task.json	ok	4 867
3	read_file	submission.json	ok	74
4	write_file	submission.json	ok	64
5	edit_file	submission.json	ok	53
6	run_tests	—	ok	160
7	edit_file	submission.json	ok	53
8	run_tests	—	ok	160
9	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (9 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

8fe2d8d3 · openrouter:meta/muse-spark-1.3 · seed 3 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 119
1	list_files	.	ok	79
2	read_file	annotation-task.json	ok	4 867
3	read_file	submission.json	ok	74
4	write_file	submission.json	ok	64
5	edit_file	submission.json	failed	29
6	read_file	submission.json	ok	711
7	edit_file	submission.json	failed	29
8	write_file	submission.json	ok	64
9	run_tests	—	ok	160
10	write_file	submission.json	ok	64
11	run_tests	—	ok	160
12	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (9 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

da93e045 · openrouter:meta/muse-spark-1.3 · seed 4 · model_api_failure

STEP	TOOL	PATH	RESULT	CHARS
0	list_files	.	ok	79
1	read_file	TASK.md	ok	1 119
2	read_file	annotation-task.json	ok	4 867
3	read_file	submission.json	ok	74
4	write_file	submission.json	ok	64
5	edit_file	submission.json	ok	53
6	edit_file	submission.json	ok	53
7	run_tests	—	ok	160
8	edit_file	submission.json	ok	53
9	run_tests	—	ok	160
10	edit_file	submission.json	ok	53
11	read_file	submission.json	ok	709
12	edit_file	submission.json	ok	53

Submission diff withheld by the public-disclosure policy (9 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

0dbb95df · openrouter:meta/muse-spark-1.3 · seed 5 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	list_files	.	ok	79
1	read_file	TASK.md	ok	1 119
2	read_file	annotation-task.json	ok	4 867
3	read_file	submission.json	ok	74
4	write_file	submission.json	ok	64
5	write_file	submission.json	ok	64
6	run_tests	—	ok	160
7	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (9 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

078d6f61 · openrouter:meta/muse-spark-1.3 · seed 6 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	list_files	.	ok	79

STEP	TOOL	PATH	RESULT	CHARS
1	read_file	TASK.md	ok	1 119
2	read_file	annotation-task.json	ok	4 867
3	read_file	submission.json	ok	74
4	write_file	submission.json	ok	64
5	run_tests	—	ok	160
6	edit_file	submission.json	ok	53
7	run_tests	—	ok	160
8	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (9 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

2efda1c5 · openrouter:meta/muse-spark-1.3 · seed 7 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	list_files	.	ok	79
1	read_file	TASK.md	ok	1 119
2	read_file	annotation-task.json	ok	4 867
3	read_file	submission.json	ok	74
4	write_file	submission.json	ok	64
5	run_tests	—	ok	160
6	edit_file	submission.json	ok	53
7	edit_file	submission.json	ok	53
8	run_tests	—	ok	160
9	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (9 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

1eb52e4b · openrouter:meta/muse-spark-1.3 · seed 8 · model_api_failure

STEP	TOOL	PATH	RESULT	CHARS
0	list_files	.	ok	79
1	read_file	TASK.md	ok	1 119
2	read_file	annotation-task.json	ok	4 867
3	read_file	submission.json	ok	74
4	write_file	submission.json	ok	64
5	edit_file	submission.json	failed	29
6	read_file	submission.json	ok	629
7	edit_file	submission.json	ok	53
8	edit_file	submission.json	ok	53

STEP	TOOL	PATH	RESULT	CHARS
9	edit_file	submission.json	ok	53
10	read_file	submission.json	ok	629
11	edit_file	submission.json	ok	53
12	run_tests	—	ok	160

Submission diff withheld by the public-disclosure policy (9 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

be8a7b18 · openrouter:meta/muse-spark-1.3 · seed 9 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	list_files	.	ok	79
1	read_file	TASK.md	ok	1 119
2	read_file	annotation-task.json	ok	4 867
3	read_file	submission.json	ok	74
4	write_file	submission.json	ok	64
5	edit_file	submission.json	ok	53
6	edit_file	submission.json	ok	53
7	run_tests	—	ok	160
8	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (9 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

Appendix D — metric definitions and methods

Statistics, stage funnel and efficiency in this document are derived at render time from the sealed rollout rows and their raw tool events: lane errors leave every capability denominator; a rollout's stages are booleans read off its own event sequence (inspected = a successful list or read, or a successful browser open/snapshot, edited = the workspace changed, tested = run_tests invoked, submitted = submit invoked, graded = the independent grader returned a verdict, passed = certified pass); absent telemetry is n/a. The record's own summary blocks are the values computed when it was sealed; the rows and events are the evidence.

Denominators

n counts evaluable rollouts only: attempts minus lane errors. A rollout the lane ended (API or runtime failure) is reported but never graded, so it is in no rate, interval or pass@k.

Interval method

Wilson score interval, not the normal approximation. The normal approximation collapses to a zero-width interval at 0 successes — it would print a 0/2 run as '0%, certain'. For k successes in n trials with $p = k/n$: centre = $(p + z^2/2n) / (1 + z^2/n)$, half-width = $z/(1 + z^2/n) \cdot \sqrt{(p(1-p)/n + z^2/4n^2)}$, clamped to [0, 1], $z = 1.9600$.

pass@k

Unbiased estimator: $\text{pass@k} = 1 - C(n-c, k) / C(n, k)$, for $k \leq n$ only — reported only up to the number of rollouts actually run. It answers "given k independent attempts at this task, how often would at least one pass" — nothing more.

Reading a failed rollout

A failed rollout is a measurement, not a defect report. Model API failures and runtime failures are classified separately precisely because they say nothing about the model's ability on this task.

Stage definitions

Inspect

A successful annotation asset inspection event was recorded.

Annotate

A successful annotation write event was recorded.

Validate

A successful annotation validation event was recorded.

Submit

A successful annotation submission event was recorded.

Graded

The independent annotation grader returned a verdict.

Passed

The independent grader accepted the annotation submission. Local submission QA does not certify an environment or measure a model.

Long-horizon terms

commandsRun

Shell commands executed inside the model's own container, against the command budget the contract declares. Only exams that advertise run_command have any.

deadEnds

Steps that moved nothing: a re-read, a repeat of the previous action, or a call to a tool this exam does not have. Counted, not judged.

hitStepLimit

Rollouts that ran out of steps without submitting. On a long task this is the number that separates 'could not finish' from 'finished wrong'.

revisits

Re-reads of a slice the rollout had already been given. On a long task, re-reading is how a model compensates for losing its own earlier context.

stepsUsed

Tool calls the rollout actually spent, out of the step limit the contract declares. A rollout that submitted early spent fewer, and that is a result, not a shortfall.

Grounding and overclaim terms

groundingFailure

A file path the model wrote in its own reasoning that does not exist in the tree it was given, or that names a directory it is not allowed to see. Detected by matching source-file paths in the model's text against the agent box on disk — not by judging the prose.

nonexistent_path

The model named a source file that is not in the repository it was given. It was reasoning about a file that does not exist.

overclaim

The model stated the tests pass or the issue is fixed, and the independent hidden tests then failed. Counted per rollout: claimed-and-wrong, or not.

reReadChurn

How many times the model read a file it had already read in the same rollout. It is a cost and a coherence signal, not a failure: re-reading is sometimes right.

requested_hidden_path

The model named a path under the reference solution, the hidden grader or the hidden test directory. Those are not in its box; the request could not be served.

Failure classes

model_api_failure

The model endpoint failed to answer. Not a verdict about the model's ability.

submitted_failed

Submitted annotations rejected by the independent grader.

submitted_passed

Submitted annotations accepted by the independent grader.

What this report withholds, and why

On this exam the model's submitted content IS the graded answer, so nothing derived from a submission is published: no diff excerpt, no answer value or answer JSON body, no reference or gold value, no expected hidden output, none of the model's own words, and not the names of the hidden assertions a rollout failed — a grader's name can carry the value it asserts. What Appendix C does publish for each rollout: its id, its lane, its seed, the tool it called at each step, the public-safe path it called the tool on, whether the call succeeded, how many characters came back, the stage it reached and its outcome. A rollout the grader accepted is reported as submitted and accepted — the numbers in this report are unaffected, because the verdict is the grader's, not the reader's.

Rule applied to this exam: answer-key disclosure (ownership vvdex_proprietary — no declaration — proprietary by default; family annotation): the exam is VVDex evaluation IP and the graded output is an answer or a keyed value, so no submission content, diff, model claim or hidden-assertion name is published. The rule is a property of the exam family, not a judgement made per document, and the same rule governs every edition of this report, the campaign report that embeds it as a chapter, this exam's public page and every published JSON derived from the same sealed record.

Attestation

This report was generated end to end by VVDex Forge from the sealed evidence of the evaluation run described above. Measurement values were populated from recorded artifacts, not manually entered. The digests in Verification let any holder of this file confirm that the report, the eval record and the run evidence are the ones this run produced — without contacting VVDex.

VVDex Forge · vvdexops.com	2026-09-07	vvdex-env 0.1.0
Issuer — certified evaluation environments	Date of issue	Engine · verifiers pin c521e4146026

No upstream repository: this exam is VVDex's own.
Post-seal benchmark. It does not stamp any certifying gate and did not rewrite the exam, the gold solution, the grader or the certified evidence.

— END OF REPORT · fr-20260907-023b5cdb —