

Five exams, five families — ten rollouts per lane

5 sealed evaluations · 148 graded rollouts · 93
certified passes

CURRENT
CAMPAIGN REPORT
2026-09-02

Evaluation scope:
cli/claude-sonnet
cli/codex
codestral-latest

Prepared by:
VVDex Forge — regenerated from verified sealed records only

fc-d89e429d2781 · 5 record(s) · every value recomputable

Contents

01.	Executive summary	3
02.	Per-exam results	4
03.	Environment families	5
04.	Campaign aggregate across these five exams	6
05.	Harness status	7
06.	Cells excluded from capability	8
07.	What the outcomes mean	9
08.	What this report withholds, and why	10
09.	Exam · martinblech-xmltodict-issue-257	11
10.	Exam · order-desk-1	47
11.	Exam · fault-recovery-1	84
12.	Exam · grounded-rag-1	118
13.	Exam · memory-fact-update-1	146
14.	Evidence	169
15.	Attestation	170

01 Executive summary

Release status: CURRENT. This is the campaign the published featured index carries; its exams' public pages state these numbers.

EXAMS 5	MODEL LANES 3	ATTEMPTS 150	GRADED ROLLOUTS 148
CERTIFIED PASSES 93	LANE ERRORS (NOT GRADED) 2	WITHHELD / INVALID 0	

Purpose

This report presents a controlled evaluation of 3 model lane(s) against 5 certified exam(s), 10 rollouts per lane. Every lane received the identical frozen workspace, tool contract and limits per exam; grading ran afterwards, in isolation, against hidden tests no lane ever saw; and every rollout carries three separate outcomes — grader, integrity, certified — of which only the certified outcome is counted. Each exam is rendered below as a chapter derived from the same sealed evaluation record as its standalone report: the chapter's bytes are authenticated by this campaign document, the standalone report by its own digest.

Outcome

93

certified pass(es) across 148 graded rollouts (150 attempts, 2 lane error(s) not graded) in 15 cells — counted only where the grader accepted the submission AND containment was verified AND provenance is complete AND no integrity invariant or verdict-bearing harness suspicion applies.

fc-d89e429d2781 · generated from 5 verified sealed record(s)

Key findings across the campaign

- martinblech-xmltodict-issue-257 · **PASS** — cli/claude-sonnet, cli/codex passed the independent hidden tests
- martinblech-xmltodict-issue-257 · **ATTENTION** — submitted_failed × 11 (cli/claude-sonnet, codestral-latest)
- order-desk-1 · **PASS** — cli/claude-sonnet, cli/codex, codestral-latest passed the independent hidden tests
- order-desk-1 · **ATTENTION** — submitted_failed × 8 (codestral-latest)
- order-desk-1 · **LANE** — model_api_failure × 1 (codestral-latest)
- fault-recovery-1 · **PASS** — cli/claude-sonnet, cli/codex passed the independent hidden tests
- fault-recovery-1 · **ATTENTION** — context_exhaustion × 5 (codestral-latest)
- fault-recovery-1 · **ATTENTION** — edited_tested_no_submit × 4 (codestral-latest)
- fault-recovery-1 · **ATTENTION** — submitted_failed × 1 (codestral-latest)
- grounded-rag-1 · **PASS** — cli/claude-sonnet, cli/codex passed the independent hidden tests
- grounded-rag-1 · **ATTENTION** — submitted_failed × 18 (cli/claude-sonnet, codestral-latest)
- grounded-rag-1 · **LANE** — model_api_failure × 1 (codestral-latest)
- memory-fact-update-1 · **PASS** — cli/claude-sonnet, cli/codex, codestral-latest passed the independent hidden tests
- memory-fact-update-1 · **ATTENTION** — submitted_failed × 8 (codestral-latest)

02 Per-exam results

MODEL LANE	SWE PUBLIC.SWE.MART INBLECH-XMLTODI CT-ISSUE-257	RAG VVDEX.KNOWLEDGE.G ROUNDED-RAG-1	MEMORY VVDEX.KNOWLEDG E.MEMORY-FACT-U PDATE-1	HARNESS VVDEX.HARNES S.FAULT-RECOV ERY-1	BROWSER VVDEX.BROWSER.0 RDER-DESK-1
cli/codex	10/10	10/10	10/10	10/10	10/10
cli/claude-sonnet	9/10	1/10	10/10	10/10	10/10
codestral-latest	0/10	0/9 + 1 lane error	2/10	0/10	1/9 + 1 lane error

Each cell is *certified passes / graded rollouts* for that lane on that exam. An attempt the lane ended before grading is reported beside the cell — 1/9 + 1 lane error, never 1/10 — and is outside the denominator; withheld and invalid attempts are reported the same way. No cell here carries a percentage: these are separate exams, not one score.

03 Environment families

This campaign covers five exam(s) drawn from five environment family(ies). The sentence in each row is the family's own description of what the exam tests; nothing is claimed here that the exam does not.

FAMILY	EXAM	WHAT IT TESTS
Software repair (public OSS SWE)	<code>public.swe.martinblech-xmltodict-issue-257</code>	A real software-engineering task taken from a public open-source repository: the model is given the codebase and a written issue, and must change the code so that the behaviour described in the issue is fixed.
Grounded retrieval (RAG)	<code>vvdex.knowledge.grounded-rag-1</code>	Grounded question answering over a small document collection.
Cross-session memory update	<code>vvdex.knowledge.memory-fact-update-1</code>	A state-and-memory task across sessions: a fact was recorded earlier, an authoritative source has since changed it, and the model must answer from the newer authoritative roster rather than the stale carried memory, and acknowledge the change.
Tool and harness fault recovery	<code>vvdex.harness.fault-recovery-1</code>	A code-fix task carried out through a deliberately unreliable tool loop.
Browser and computer use	<code>vvdex.browser.order-desk-1</code>	A task inside a real web application, driven through a real headless browser.

04 Campaign aggregate across these five exams

MODEL LANE	CERTIFIED PASSES	GRADED ROLLOUTS	LANE ERRORS	WITHHELD	INVALID	CERTIFIED RATE (95% CI)
cli/codex	50	50	0	0	0	100% [93-100]
cli/claude-sonnet	40	50	0	0	0	80% [67-89]
codestral-latest	3	48	2	0	0	6% [2-17]

Raw counts first. *Graded rollouts* = certified passes + model fails; lane errors, withheld and invalid cells are outside capability arithmetic in both directions. A rate with its 95% Wilson interval is shown only for a lane with at least 10 countable attempts; intervals that overlap support no ordering. Every percentage in this section is an aggregate over this campaign's exams only, not a capability score or a ranking.

05 Harness status

No record in this campaign fired a harness self-suspicion rule. Every countable cell is a grader outcome under verified containment.

06 Cells excluded from capability

An attempt is excluded from capability only when the lane, not the model, ended it (lane error), when the run cannot prove its own conditions (withheld), or when an integrity invariant fired (INVALID). A submission the hidden grader failed is a model result and is counted, never listed here.

EXAM	LANE	SEED	EXCLUSION	RECORDED REASON
order-desk-1	codestral-latest	2	lane error	The read operation timed out
grounded-rag-1	codestral-latest	4	lane error	HTTP 400: Assistant message must have either content or tool_calls, but not none.

07 What the outcomes mean

- **CERTIFIED PASS:** the grader accepted the submission, containment was verified, provenance is complete and no integrity invariant fired.
- **model fail:** the submission failed the hidden tests under verified containment; a real, attributable model result.
- **lane error:** the provider or CLI lane failed (rate limit, quota, transport) before an attributable measurement existed; not a model result.
- **withheld:** the run cannot prove its own conditions (containment unproven, provenance incomplete, or a harness suspicion the record cannot clear); neither a pass nor a model failure.
- **INVALID:** an integrity invariant fired (isolation breach, evidence mismatch); the row is not a model result.
- **harness suspect:** the run's own self-check fired; each rule is classified as verdict-bearing or diagnostic, and verdict-bearing suspicion withholds the affected cells.

08 What this report withholds, and why

What a public document may print about a rollout is decided by what the exam GRADES, not case by case. Where the graded artefact is a change to a public upstream repository, a bounded redacted excerpt of the model's own diff is published. Where the graded output is an answer, a keyed value or the state of an application, nothing derived from the submission is published — no diff, no answer value, no model claim, and not the names of the hidden assertions a rollout failed, because a grader's name can carry the value it asserts. Withheld on every exam without exception: the contents of tool results, the hidden tests, the reference solution, host paths and keys. The counts, verdicts and intervals in this report are the grader's and are unaffected by what is withheld.

EXAM	DISCLOSURE CLASS	RULE APPLIED
<code>martinblech-xmltodict-issue-257</code>	<code>public_source</code>	public-source disclosure (ownership <code>public_source</code> declared by engine registry, upstream https://github.com/martinblech/xmltodict under MIT; family <code>public_oss_swe</code>): the submitted change is a patch against a public repository, so a redacted excerpt of it is published.
<code>order-desk-1</code>	<code>state</code>	state disclosure (ownership <code>vvdex_proprietary</code> — engine registry; family <code>browser</code>): the exam is VVDex evaluation IP and the graded output is the application's own state and audit trail, so no submission content, diff or state is published.
<code>fault-recovery-1</code>	<code>answer_key</code>	answer-key disclosure (ownership <code>vvdex_proprietary</code> — engine registry; family <code>harness_eval</code>): the exam is VVDex evaluation IP and the graded output is an answer or a keyed value, so no submission content, diff, model claim or hidden-assertion name is published.
<code>grounded-rag-1</code>	<code>answer_key</code>	answer-key disclosure (ownership <code>vvdex_proprietary</code> — engine registry; family <code>rag_grounding</code>): the exam is VVDex evaluation IP and the graded output is an answer or a keyed value, so no submission content, diff, model claim or hidden-assertion name is published.
<code>memory-fact-update-1</code>	<code>answer_key</code>	answer-key disclosure (ownership <code>vvdex_proprietary</code> — engine registry; family <code>long_term_memory</code>): the exam is VVDex evaluation IP and the graded output is an answer or a keyed value, so no submission content, diff, model claim or hidden-assertion name is published.

Exam revisions retired for future evaluation

- `vvdex.knowledge.grounded-rag-1 0.1.0` — Retired for future evaluation on 2026-09-02: post-run public disclosure of answer-bearing submission content. Historical sealed results on this revision are unaffected: the disclosure happened after those runs, so it could not have influenced them.
- `vvdex.knowledge.memory-fact-update-1 0.1.0` — Retired for future evaluation on 2026-09-02: post-run public disclosure of answer-bearing submission content. Historical sealed results on this revision are unaffected: the disclosure happened after those runs, so it could not have influenced them.

09 Exam · martinblech-xmltodict-issue-257

public.swe.martinblech-xmltodict-issue-257 · record fr-20260902-6616b188 · record digest 138425da59aabc18... · harness clean · containment clean

Executive summary

MODELS TESTED 3	CERTIFIED PASSES 19 of 30 graded rollouts	SUBMITTED 30 of 30 graded rollouts that called submit	NOT GRADED — LANE ERRORS 0 of 30 attempts
---------------------------------------	--	--	--

Purpose

This report presents a controlled evaluation of 3 model lane(s) against the certified exam public.swe.martinblech-xmltodict-issue-257 (v0.1.0, family public_oss_swe). Every lane received the identical frozen workspace, tool contract and limits; grading ran afterwards, in isolation, against hidden tests no lane ever saw. The purpose is to state, with reproducible evidence, what each lane did and where it stopped.

Outcome

19 of 30

graded rollouts passed the independent hidden tests and were certified

19 of 30 graded rollouts passed. cli/claude-sonnet, cli/codex completed the full loop (inspect, edit, run the visible tests, submit) and the sealed grader accepted the submission. The most common way to fail was submitted_failed — submitted a change; the independent hidden tests still failed.

Key findings

F-01 cli/claude-sonnet, cli/codex passed the independent hidden tests

F-02 submitted_failed × 11 (cli/claude-sonnet, codestral-latest)

Each finding is stated in full, with evidence, in the Findings section; recommendations for acting on them follow it. Elapsed wall clock for the whole engagement: 2327.3s.

Exam definition

In scope. The ability of each configured model lane to complete this one certified task end to end — inspect the workspace, produce a correct edit, verify it with the visible tests, and submit — within 40 tool steps and 9720s of wall clock, at 10 rollout(s) per lane. Behavioural measurements along the way (tool discipline, grounding, overclaim) are in scope because they are read from the recorded trace, not judged by another model.

Out of scope. General model quality, performance on other task families, prompt-tuning on the lane's behalf, and any comparison this run's sample size cannot support. A lane's API failure is reported as infrastructure, never as model capability.

The instrument. Exam `public.swe.martinblech-xmltodict-issue-257` was certified before this run: its reference solution passes, an empty submission fails, its grader distinguishes near-misses, and its sealed evidence is unchanged by this evaluation (see Certification evidence).

What was measured

EXAM	VERSION	FAMILY	ROLLOUTS
public.swe.martinblech-xmltodict-issue-257	0.1.0	public_oss_swe	30 (10 per model)

Where the defect came from

Source: `martinblech/xmltodict (MIT)` · issue
<https://github.com/martinblech/xmltodict/issues/257> · upstream fix
<https://github.com/martinblech/xmltodict/pull/421>

The defect and its upstream fix are public and unmodified. The exam around them — the frozen parent tree, the independently written hidden suite, the control probes and the isolation — is VVDex's, and it is what this report measures.

Certification evidence

4 of 5 certification pillars are recorded in this package, loaded from the receipts sealed into the exam and matching this run's exam fingerprint. What is recorded is shown with its real value; what is not is named below and is not claimed.

FROZEN BASELINE

FAIL

expected — an empty submission over 2 run(s) must not pass

→

GOLD / REFERENCE

PASS

expected — the reference solution over 2 run(s) must pass

GRADER CONTROLS

7 / 7

The grader accepted the reference and rejected every near-miss.

ATTACK PROBES

10 / 10 blocked

0 trivial exploit(s) found.

RUNTIME

sha256:679b36225a1f83238efba8c71b9cde3c24caaeacc9b682ae92819cb55eec328f

network policy: deny

Certification metadata is not included in this report package for:

- sealed evidence bundle

Benchmark results remain valid as run evidence; the pillar(s) above are not substantiated by this document and this report does not claim them.

What the package does carry: the exam fingerprint below is the contract digest this run was executed against — a holder of the certification bundle for this fingerprint can join the two independently.

EXAM FINGERPRINT

aa41977120e6ecddb5c4fc603e8ab23cb272726a86750119969ff4dfc90e297

The contract digest this run was executed against.

Evaluation configuration

REPORT	fr-20260902-6616b188
ISSUED	2026-09-02
SUBJECT EXAM	public.swe.martinblech-xmltodict-issue-257 · v0.1.0
FAMILY	public_oss_swe
PREPARED BY	VVDex Forge engine vvdex-env 0.1.0
CLASSIFICATION	Independent model evaluation report — independently verifiable
CERTIFICATION	Exam certification: partially recorded (see Certification evidence)
STATUS	FINAL · report sealed against its own digest

Timeline

WHEN (UTC)	EVENT
2026-09-02 00:18:47	Evaluation run started
2026-09-02 00:57:34	Evaluation run finished
2026-09-02 00:57:34	Report generated from the sealed run evidence

Models under test

MODEL	PROVIDER	ENDPOINT	BILLING
Claude Code CLI (Sonnet, subscription)	cli	local runtime	operator subscription, flat — not billed per token
Codex CLI (GPT, subscription)	cli	local runtime	operator subscription, flat — not billed per token
codestral-latest	mistral	VVDex-configured	VVDex free-quota

A customer endpoint is recorded by origin only — never its port, its path, or its key.

Limits

RUNTIME docker	STEP LIMIT 40	WALL-CLOCK LIMIT 9720s	MAX OUTPUT TOKENS 2048
TEMPERATURE 0.2	RETRY POLICY none	COMMAND BUDGET 0	TOOLS OFFERED list_files, read_file, write_file, edit_file, run_tests, submit

The evaluated model never saw the hidden tests or the reference solution, and could not change its result. Grading ran in a separate step, after submission, on a container with no network.

Model outcomes

Denominators. Attempts requested: 30. Graded (evaluable) rollouts: 30. Not graded — lane errors: 0. A lane error is a rollout the API or the runtime ended before the hidden grader ran: it is listed, excluded from every rate and pass@k, and never silently dropped. Passes are certified passes: hidden grader exit 0 with integrity verified.

MODEL	ATTEMPTS	GRADED	PASSED	MODEL FAILS	LANE ERRORS	SUBMITTED	MEAN TIME	TYPICAL DEEPEST STAGE
Claude Code CLI (Sonnet, subscription) cli/claude-sonnet	10	10	9 / 10	1	0	10 / 10	159.8s	Passed hidden tests
Codex CLI (GPT, subscription) cli/codex	10	10	10 / 10	0	0	10 / 10	42.4s	Passed hidden tests
codestral-latest codestral-latest	10	10	0 / 10	10	0	10 / 10	11.5s	Graded

Every rollout, with its own deepest stage and grade, is in Appendix A.

Success rate, confidence and pass@k

Every rate on this page is a measurement of a finite number of rollouts, so every rate carries the interval that measurement actually supports. n is the evaluable count — attempts minus rollouts the lane ended before grading; those are listed under "not graded" and sit in no rate, interval or pass@k.

MODEL	ATTEMPTS	N	NOT GRADED	PASSES	RATE	95% INTERVAL	SEEDS	PASS@1	PASS@3	PASS@5	PASS@10
cli/claude-sonnet	10	10	0	9	90.0%	[60%, 98%]	0–9	90.0%	100.0%	100.0%	100.0%
cli/codex	10	10	0	10	100.0%	[72%, 100%]	0–9	100.0%	100.0%	100.0%	100.0%
codestral-latest	10	10	0	0	0.0%	[0%, 28%]	0–9	0.0%	0.0%	0.0%	0.0%

Interval and pass@k methods are stated in full in Appendix D. pass@k is shown for k = 1, 3, 5, 10; the sealed record carries every k up to n.

Model comparison

MODEL	PASS@1	95% INTERVAL	TOKENS	COST	MEAN ROLLOUT	OVERCLAIM RATE	TOP FAILURE CLASS
cli/claude-sonnet	90.0%	[60%, 98%]	n/a	n/a	159.8s	0.0%	submitted_failed
cli/codex	100.0%	[72%, 100%]	n/a	n/a	42.4s	0.0%	none
codestral-latest	0.0%	[0%, 28%]	172 890	not billed	11.5s	0.0%	submitted_failed

Tokens and cost read n/a where a lane reports no telemetry (a flat-rate CLI) — an absent measurement, not zero. pass@1 and its interval are over evaluable rollouts only.

This is not a leaderboard. Separated by this run (intervals do not overlap): cli/claude-sonnet vs codestral-latest; cli/codex vs codestral-latest. Not separated (intervals overlap): cli/claude-sonnet vs cli/codex. The primary comparison in this report is the stage funnel below — where a model stops is a finding; a percentage over this many rollouts is not.

Stage funnel

The funnel is the primary reading of a run. Where a model stops says more than a single pass percentage does, especially at low N.

STAGE	CLI:CLI/CLAUDE-SONNET	CLI:CLI/CODEX	MISTRAL:CODESTRAL-LATEST
Inspected	10 / 10	10 / 10	10 / 10
Edited	9 / 10	10 / 10	10 / 10
Ran tests	10 / 10	10 / 10	10 / 10
Submitted	10 / 10	10 / 10	10 / 10
Graded	10 / 10	10 / 10	10 / 10
Passed hidden tests	9 / 10	10 / 10	0 / 10

Deepest stage reached

MODEL	DEEPEST (ANY ROLLOUT)	TYPICAL (MOST ROLLOUTS)	HIDDEN-TEST PASSES	NOT GRADED
cli:cli/claude-sonnet	Passed hidden tests	Passed hidden tests	9 / 10	0
cli:cli/codex	Passed hidden tests	Passed hidden tests	10 / 10	0
mistral:codestral-latest	Graded	Graded	0 / 10	0

Stage definitions are in Appendix D. The matrix counts evaluable rollouts; lane errors are shown beside it, not inside it. The stages are not strictly nested: a model can submit without editing, and the matrix shows that rather than hiding it. Each rollout's own deepest stage is in Appendix A.

Behavioural analysis

Long-horizon behaviour

Counted off the recorded trace of each rollout. A dead end is a step that moved nothing: a re-read of a slice already returned, a repeat of the previous action, or a call to a tool this exam does not have.

MODEL	N	STEPS USED	RAN OUT OF STEPS	COMMANDS	CHANGED THE TREE	REFUSED (OVER BUDGET)	REVISITS	DEAD ENDS / ROLLOUT
cli:cli/claude-sonnet	10	5 / 40	0	0 / 0	0	0	0	0
cli:cli/codex	10	6 / 40	0	0 / 0	0	0	0	0
mistral:codestral-latest	10	5 / 40	0	0 / 0	0	0	0	0

Each column is defined in Appendix D.

Hallucination & overclaim

Three measurements, each with a stated definition. None of them is a judgement of the model's prose by another model — every number below comes from comparing what the model wrote against what was on disk, or against what the independent grader returned.

MODEL	N	GROUNDING FAILURES	KINDS	ROLLOUTS AFFECTED	OVERCLAIMS	OVERCLAIM RATE	RE-READS
cli:cli/claude-sonnet	10	0	none	0 / 10	0 / 10	0.0%	0
cli:cli/codex	10	0	none	0 / 10	0 / 10	0.0%	0
mistral:codestral-latest	10	0	none	0 / 10	0 / 10	0.0%	0

No grounding failures and no false pass claims were detected across these 30 rollout(s). No lane re-read material it had already been given.

Every term above is defined in Appendix D.

Efficiency & telemetry

This run has no per-token price attached — a local model, a free quota, or a customer endpoint we are not billed for — so spend is reported in tokens. Tokens are what was actually measured; a dollar figure here would be invented. Lanes run through a flat-rate local CLI report no token telemetry; their cells say n/a — an absent measurement, not zero. Rollout counts here are evaluable rollouts.

MODEL	ROLLOUTS	TOKENS	TOKENS / ROLLOUT	COST	PASSES	RATE
cli:cli/claude-sonnet	10	n/a	n/a	n/a – telemetry unavailable	9 / 10	90.0%
cli:cli/codex	10	n/a	n/a	n/a – telemetry unavailable	10 / 10	100.0%
mistral:codestral-latest	10	172 890	17 289	not billed	0 / 10	0.0%

Findings

Every finding below is derived from the recorded data of this run — none is an opinion.
Severity: PASS a lane succeeded · ATTENTION a capability gap · LANE / INFRA
infrastructure, not capability · NOTE a property of the measurement.

F-01

PASS

cli/claude-sonnet, cli/codex passed the independent hidden tests

Severity: PASS

Evidence: Model outcomes · Stage funnel · Appendix A

19 of 30 graded rollouts completed the full loop (inspect, edit, run the visible tests, submit) and the sealed grader accepted the submission. This demonstrates the task is solvable under the published limits.

F-01 · derived from the recorded run data

F-02

ATTENTION

submitted_failed × 11 (cli/claude-sonnet, codestral-latest)

Severity: ATTENTION

Evidence: Appendix A · Appendix C

Submitted a change; the independent hidden tests still failed.

F-02 · derived from the recorded run data

Recommendations

One recommendation per observed failure mode, addressed to the lane it affects. Each traces back to a finding above and to the raw trace in Appendix A.

R-01

PROBLEM

Submitted a change; the independent hidden tests still failed.

AFFECTED LANES

cli/claude-sonnet, codestral-latest

WHAT WE WOULD LOOK AT FIRST

The model completes the loop and produces a plausible near-miss — the named failing tests are the exact behavioral cases distinguishing it from the reference fix.

Verification

RUN ID

live-20260902T001847Z

EXAM FINGERPRINT

aa41977120e6ecdddb5c4fc603e8ab23cb272726a86750119969ff4dfc90e297

The contract digest this run was executed against.

Provenance of this chapter

This chapter was rendered by the campaign generator from the sealed evaluation record; it is not the standalone report reproduced byte for byte. Verify the standalone report with its own digest, verify the record with `vvdex-env records verify`, and verify this campaign document as a whole.

RELATIONSHIP

campaign-rendered chapter

Derived from the same sealed evaluation record as the standalone report; these bytes are authenticated by the campaign artifact that contains them.

SEALED RECORD DIGEST

138425da59aabc18aa208c339aa41919f387e1ecb80fb725cfed8f7d6ba78713

sha256 over the record's canonical JSON — the identity every renderer works from.

SOURCE RECORD FILE

02b4f83ee15251a5808a83153c9c4590936af23c58cc94d172556e9f983d1909

sha256 of fr-20260902-6616b188.json as written beside the standalone report.

SOURCE STANDALONE REPORT

fc395492a83645c6f5cd4360e53359fe85eca4fd1a00c005055989ba286694fa

sha256 of fr-20260902-6616b188.html — independently verifiable with its own digest.

SOURCE STANDALONE PDF

c6a3f43aeaf0686dec3d34e298ffd1f936c3ffe328f62b8dc2601c3b9d2ced4b

sha256 of fr-20260902-6616b188.pdf.

RUN EVIDENCE BUNDLE

6616b188f015f52e157bdfb257e8d0f819e1f6cb9bcd3b024b340ac0771e0fd8

The digest the run's own bundle computed over itself.

CERTIFIED EVIDENCE SEAL

not available in this package

The exam's sealed evidence, established before this run.

Measurement history

This is the first recorded run on this exam. There is nothing to compare it against yet, and a single measurement is not a trend.

The full artifact-by-artifact digest table, and the reproduction protocol, are in Appendix B.

Appendix A — every rollout

ROLLOUT	MODEL	SEED	OUTCOME	SUBMITTED	DEEPEST STAGE	HIDDEN TESTS	FILES	TOKENS	ELAPSED
3e6298d8	cli/claude-sonnet	0	submitted_passed	yes	Passed hidden tests	pass	1	n/a	180.1s
cb495fe6	cli/claude-sonnet	1	submitted_passed	yes	Passed hidden tests	pass	1	n/a	150.5s
e76e1c66	cli/claude-sonnet	2	submitted_passed	yes	Passed hidden tests	pass	1	n/a	133.3s
9e19f314	cli/claude-sonnet	3	submitted_passed	yes	Passed hidden tests	pass	1	n/a	147.9s
aaa748af	cli/claude-sonnet	4	submitted_passed	yes	Passed hidden tests	pass	1	n/a	225.9s
3e36cea5	cli/claude-sonnet	5	submitted_passed	yes	Passed hidden tests	pass	1	n/a	181.1s
6516049a	cli/claude-sonnet	6	submitted_failed	yes	Graded	fail	0	n/a	208.2s
8c122607	cli/claude-sonnet	7	submitted_passed	yes	Passed hidden tests	pass	1	n/a	110.6s
3161d1b1	cli/claude-sonnet	8	submitted_passed	yes	Passed hidden tests	pass	1	n/a	100.9s
6b722882	cli/claude-sonnet	9	submitted_passed	yes	Passed hidden tests	pass	1	n/a	159.4s
92c89eb1	cli/codex	0	submitted_passed	yes	Passed hidden tests	pass	1	n/a	46.0s
6d48c276	cli/codex	1	submitted_passed	yes	Passed hidden tests	pass	1	n/a	40.8s
1d32b64e	cli/codex	2	submitted_passed	yes	Passed hidden tests	pass	1	n/a	40.2s
3e53d590	cli/codex	3	submitted_passed	yes	Passed hidden tests	pass	1	n/a	40.7s
90405084	cli/codex	4	submitted_passed	yes	Passed hidden tests	pass	1	n/a	51.1s
48e359d8	cli/codex	5	submitted_passed	yes	Passed hidden tests	pass	1	n/a	36.9s
fd7561b3	cli/codex	6	submitted_passed	yes	Passed hidden tests	pass	1	n/a	39.5s
4fe31a46	cli/codex	7	submitted_passed	yes	Passed hidden tests	pass	1	n/a	43.4s
18660280	cli/codex	8	submitted_passed	yes	Passed hidden tests	pass	1	n/a	46.1s
195b8f72	cli/codex	9	submitted_passed	yes	Passed hidden tests	pass	1	n/a	39.5s
5d129402	codestral-latest	0	submitted_failed	yes	Graded	fail	1	19 250	11.5s
bd30cb85	codestral-latest	1	submitted_failed	yes	Graded	fail	1	16 813	11.2s
7cc71ab5	codestral-latest	2	submitted_failed	yes	Graded	fail	1	16 792	10.6s
d38fb2e2	codestral-latest	3	submitted_failed	yes	Graded	fail	1	16 861	14.9s
8c65b44c	codestral-latest	4	submitted_failed	yes	Graded	fail	1	16 876	11.2s
b8a625a2	codestral-latest	5	submitted_failed	yes	Graded	fail	1	16 762	10.7s
58abb85c	codestral-latest	6	submitted_failed	yes	Graded	fail	1	16 876	10.6s
94122fe4	codestral-latest	7	submitted_failed	yes	Graded	fail	1	16 759	13.5s
b0843fec	codestral-latest	8	submitted_failed	yes	Graded	fail	1	19 025	10.2s
cecbcc3b	codestral-latest	9	submitted_failed	yes	Graded	fail	1	16 876	10.2s

Failure taxonomy

CLASS	COUNT	WHAT IT MEANS
submitted_failed	11	Submitted a change; the independent hidden tests still failed.

Appendix B — evidence and artifact digests

RUN BUNDLE DIGEST 6616b188f015f52e157bdfb2 57e8d0f819e1f6cb9bcd3b02 4b340ac0771e0fd8	RUN ID live- 20260902T001847Z	STARTED 2026-09- 02T00:18:47.313247 +00:00	FINISHED 2026-09- 02T00:57:34.65040 8+00:00
--	--	--	---

Artifact digests

ARTIFACT	SHA-256
publicSummary	059a37157968962044d98a133f9422c6831ff452419eb4b8729b856a79cad646
rollouts/18660280-d29a-4c7f-928c-0f37b96f18e3.json	7049943d0e9a63ddf108caa49598a38a76ccde4f29d1e54cb9f3098ca7aa2e37
rollouts/195b8f72-f390-486a-a369-b5047d3bf081.json	f199086fc48f883f348848c82b55fa6cdf2dd1e415001939c9176faf4ab48b84
rollouts/1d32b64e-b2fd-4a37-a04f-1230e38ad920.json	d7ad72ef466cadbccd862d52e8b0cad772ef2c9360ccb602ab332ce976bd9e44
rollouts/3161d1b1-1be0-4bcc-bb1a-11592c0931aa.json	7590bef63572befafe69f1120525598f6f03ce59869ec21a7e9d61cf2c081ced
rollouts/3e36cea5-b3a2-49a1-8227-83d3c1f6b123.json	df73c10cd29332e01062110ea33c6df2c7e592aba53fbb2b67f58e333bd1d601
rollouts/3e53d590-9fa2-40ec-9d8d-284a6663637c.json	fbaf104694620c1643a2341422af017c17122b02d6dca2599d6ea77e27d10851
rollouts/3e6298d8-5275-4323-8970-b9d0c7709aea.json	e2483eb5171491c61186b5edda6e992b7a8fd87c794a010f76557a5ae0402813
rollouts/48e359d8-3ed5-46b1-b6a1-d6aac402eef9.json	8934223b0170fdaa59fc24e801016a92df6672d4e493d17c37f96db41875ec2d
rollouts/4fe31a46-de93-4684-bc08-eda4db9ceab0.json	3058c5a8ef15ec2dea64cc04c5d4ae747c66c55de8d0b80ac7acfbcb8524554da
rollouts/58abb85c-b63a-447e-a07b-071cdcb65b4b.json	ed4c3b6913b3df2f9098a2d6e8f384cd5b6f516e0de8999923013d55ac3ae22d
rollouts/5d129402-a899-48be-b2bd-0101eeaa04c8.json	3f98b67607190f913c03acb489561a96b6ddf351c5b17ac1687ae4fc75405016
rollouts/6516049a-dc0f-4636-9df0-9f827bb6dd1b.json	6d9ba60e40f7e85a59635d2f75675caf7d330cede1bdd8146b12c3c2a4445037
rollouts/6b722882-8ce0-4eb7-8cdd-db66e413e80f.json	2960bb9ff1bd93c878aa87ad43159daca5c38a6428c2309e78c9b79d4265e408
rollouts/6d48c276-533c-469a-92d4-738279c6027d.json	0ff66e69b4d90d13bac4ce41cae996babd1bf018df5002de15d5eb0cbc33daa7
rollouts/7cc71ab5-3ba2-4a7a-8417-7e1ac353afab.json	3d7170d35dd63c1803aaa4a090f6551d23a62fb350baabe62bf396d7b317fad8
rollouts/8c122607-d35f-45fc-ba31-68bb9a296352.json	2a699f21909127d47ac15af29455549e6e62d22635fa77162bd93d485a119c42
rollouts/8c65b44c-7a1b-4da4-9194-a66dcb79f193.json	a25b7c2c5e0e41c78185730582cc68d4aaf8142ba3812c6cef418651bfc0b6ed
rollouts/90405084-4394-484b-8d91-78909aa09787.json	f5d78417af80d56b68fb13cb84831e89b3c9d76f17e3264135cf5cc8065eff4f
rollouts/92c89eb1-4f77-4165-9b73-0b53e3bff694.json	b6432a4e97b4652b098fd7b50fc4420c5a1fd530974e40ab67aeebf088351447
rollouts/94122fe4-5f7e-453e-8ef5-41250ff398d5.json	585128f0e325a96995e9acace25db8f364a264fad9c91f227f2071e8dd7e60255
rollouts/9e19f314-772f-4a29-8bee-8405ef9196b5.json	12754796c393b4883e68e4b21b205b19042c045174b3c75795f4f0cef0bcf45a
rollouts/aaa748af-8efc-491a-98f0-ad6565cf7ab5.json	17b22a74714f16384f42d2e92609d770535a376f3cdf490a0df82024b4caed8d
rollouts/b0843fec-30f3-491b-bfb0-9b352dea69da.json	be256e7d715aa6ef3f50fcc4a0024a1e6224d536f4f76d6b1fd165cb17b9723e
rollouts/b8a625a2-6238-476f-9951-1f7b9ed2565c.json	41e5b0047d43b493b175b0518edbbd82c8e5b99cd9160d3941abed1e4d22251a
rollouts/bd30cb85-1647-4722-9b99-b469583dce5e.json	a2fe27d8be1fd255d92b707bdb439075a56ec040fff728b634774b435b9e3a652
rollouts/cb495fe6-4d25-457d-9e5a-2d830c2d1e87.json	28e67f4869f6a390518199f6dc27bbd78033140b3bcb58f49c8db39ca98f0c83
rollouts/cecbcc3b-ea8a-4b61-8b95-91bc9c831eea.json	9dec1bd6a4fb9af4c240dc72a6d1d974129f9a9c699ddf371148e2c80910cc2
rollouts/d38fb2e2-3dea-4276-abe9-5c44ef17994e.json	a0f91309e405bfb8cd93ebce214a8e5ef098ebbe7994d3e0d6972aa59bcfdfb1

ARTIFACT	SHA-256
rollouts/e76e1c66-ea48-4852-920a-dc077f09b4ab.json	a323881e38c7b2c9edf6ea34a992d305abc4ad23c77bdb3d8c934c9e3f622332
rollouts/fd7561b3-afd5-4de2-91a4-11262c5b065e.json	aec6d21e216d0d84d43edc7343309a41265f82880cd6aecb9289bbcd894d0ffc
spec	26f9634cb3113984cabf0c0f5675880645449731e9d8ed3baeb124c35e1e60f5
summary	2a5e9b681ab1c80119fb4922de9f9d7b74a1e58fb7594ceebf3f6be296cf4ba5

Reproducing this run

Every rollout can be re-derived: the exam is frozen (fingerprint above), the grader ran in the recorded container image, and each rollout's full trace — every tool call, every result, every word — is in the evidence bundle at its rollouts/<id>.json path. A re-run uses the same job shape: the exam id, your endpoint, n, and the runtime. The record id names this run forever.

Appendix C — redacted transcripts

For each rollout: the tool calls it made, in order, and an excerpt of the diff it produced. What a tool returned is not here — only which tool, on which of the model's own paths, whether it succeeded, and how many characters came back. A request for the reference solution or the hidden tests is recorded as a request; the path it asked for is replaced by a label, because printing it would put the hidden tree's layout in a file that leaves this machine. The diff excerpt is shown because this exam's disclosure class is public-source: the change is a patch against a public repository under the licence named in the exam definition.

Disclosure class applied to this exam: public-source disclosure (ownership public_source declared by engine registry, upstream <https://github.com/martinblech/xmltodict> under MIT; family public_oss_swe): the submitted change is a patch against a public repository, so a redacted excerpt of it is published.

3e6298d8 · cli:cli/claude-sonnet · seed 0 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	xmltodict.py	ok	12 166
1	read_file	xmltodict.py	ok	11 706
2	edit_file	xmltodict.py	ok	50
3	run_tests	—	ok	183
4	submit	—	ok	85

```
--- a/xmltodict.py
+++ b/xmltodict.py
@@ -107,12 +107,28 @@
     # without attaching it back to its parent. This avoids accumulating all
     # streamed items in memory when using item_depth > 0.
     if len(self.path) == self.item_depth:
+
+         data = (None if not self.data
+                else self.cdata_separator.join(self.data))
+         item = self.item
+         if item is None:
+             item = (None if not self.data
+                    else self.cdata_separator.join(self.data))
+
+         should_continue = self.item_callback(self.path, item)
+         if self.strip_whitespace and data:
+             data = data.strip() or None
+         if data and self._should_force_cdata(name, data) and item is None:
+             item = self.dict_constructor()
+         if item is not None:
+             if data:
+                 self.push_data(item, self.cdata_key, data)
+                 item_value = item
+             else:
+                 item_value = data
+         # Apply postprocessor as push_data would when storing in parent
+         if self.postprocessor is not None:
+             result = self.postprocessor(self.path, name, item_value)
+             if result is not None:
+                 _, item_value = result
+             else:
+                 item_value = None
+
+         should_continue = self.item_callback(self.path, item_value)
+         if not should_continue:
+             raise ParsingInterrupted
+         # Reset state for the parent context without keeping a reference to
```

Showing 36 of 36 diff line(s).

cb495fe6 · cli:cli/claude-sonnet · seed 1 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	xmltodict.py	ok	12 166
1	read_file	xmltodict.py	ok	11 706
2	edit_file	xmltodict.py	ok	50
3	run_tests	—	ok	183
4	submit	—	ok	85

```
--- a/xmltodict.py
+++ b/xmltodict.py
@@ -107,10 +107,28 @@
     # without attaching it back to its parent. This avoids accumulating all
     # streamed items in memory when using item_depth > 0.
     if len(self.path) == self.item_depth:
+
+         data = (None if not self.data
+                 else self.cdata_separator.join(self.data))
+         item = self.item
-         if item is None:
-             item = (None if not self.data
+                 else self.cdata_separator.join(self.data))
+         if self.strip_whitespace and data:
+             data = data.strip() or None
+         if data and self._should_force_cdata(name, data) and item is None:
+             item = self.dict_constructor()
+         if item is not None:
+             if data:
+                 self.push_data(item, self.cdata_key, data)
+         else:
+             item = data
+         # Simulate push_data(parent, name, item) to get the value
+         # that would be stored in the full (non-streaming) parse
+         if self.postprocessor is not None:
+             result = self.postprocessor(self.path, name, item)
+             if result is not None:
+                 _, item = result
+             else:
+                 item = None
+         if self._should_force_list(name, item):
+             item = [item]

     should_continue = self.item_callback(self.path, item)
     if not should_continue:
```

Showing 34 of 34 diff line(s).

e76e1c66 · cli:cli/claude-sonnet · seed 2 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	xmltodict.py	ok	12 166
1	read_file	xmltodict.py	ok	11 706
2	edit_file	xmltodict.py	ok	50
3	run_tests	—	ok	183
4	submit	—	ok	85

```

--- a/xmltodict.py
+++ b/xmltodict.py
@@ -107,21 +107,41 @@
     # without attaching it back to its parent. This avoids accumulating all
     # streamed items in memory when using item_depth > 0.
     if len(self.path) == self.item_depth:
+
+         data = (None if not self.data
+                 else self.cdata_separator.join(self.data))
         item = self.item
-
-         if item is None:
-             item = (None if not self.data
-                     else self.cdata_separator.join(self.data))
-
-         should_continue = self.item_callback(self.path, item)
-         if not should_continue:
-             raise ParsingInterrupted
-         # Reset state for the parent context without keeping a reference to
-         # the emitted item.
+
+         # Restore parent state (stack.pop mirrors non-streaming branch)
         if self.stack:
             self.item, self.data = self.stack.pop()
         else:
             self.item = None
             self.data = []
+
+         # Apply the same transformations the non-streaming branch uses
+         if self.strip_whitespace and data:
+             data = data.strip() or None
+         if data and self._should_force_cdata(name, data) and item is None:
+             item = self.dict_constructor()
+         if item is not None:
+             if data:
+                 self.push_data(item, self.cdata_key, data)
+                 # Simulate what push_data(parent, name, item) would do to the value
+                 if self.postprocessor is not None:
+                     result = self.postprocessor(self.path, name, item)
+                     if result is not None:
+                         _, item = result
+             else:
+                 # text-only (or empty) path
+                 if self.postprocessor is not None:

```

Showing 40 of 53 diff line(s); truncated.

9e19f314 · cli:cli/claude-sonnet · seed 3 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	xmltodict.py	ok	12 166
1	read_file	xmltodict.py	ok	11 706
2	edit_file	xmltodict.py	ok	50
3	run_tests	—	ok	183
4	submit	—	ok	85

```

--- a/xmltodict.py
+++ b/xmltodict.py
@@ -107,12 +107,21 @@
     # without attaching it back to its parent. This avoids accumulating all
     # streamed items in memory when using item_depth > 0.
     if len(self.path) == self.item_depth:
+
+         data = (None if not self.data
+                 else self.cdata_separator.join(self.data))
         item = self.item
-
-         if item is None:
-             item = (None if not self.data

```

```

-         else self.cdata_separator.join(self.data))
-
-     should_continue = self.item_callback(self.path, item)
+     if self.strip_whitespace and data:
+         data = data.strip() or None
+     if data and self._should_force_cdata(name, data) and item is None:
+         item = self.dict_constructor()
+     if item is not None:
+         if data:
+             self.push_data(item, self.cdata_key, data)
+             result = item
+     else:
+         result = data
+
+     should_continue = self.item_callback(self.path, result)
+     if not should_continue:
+         raise ParsingInterrupted
+     # Reset state for the parent context without keeping a reference to

```

Showing 29 of 29 diff line(s).

aaa748af · cli:cli/claude-sonnet · seed 4 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	xmldict.py	ok	12 166
1	read_file	xmldict.py	ok	11 706
2	edit_file	xmldict.py	ok	50
3	run_tests	—	ok	183
4	submit	—	ok	85

```

--- a/xmldict.py
+++ b/xmldict.py
@@ -107,21 +107,38 @@
     # without attaching it back to its parent. This avoids accumulating all
     # streamed items in memory when using item_depth > 0.
     if len(self.path) == self.item_depth:
+
+         data = (None if not self.data
+                 else self.cdata_separator.join(self.data))
+         item = self.item
-         if item is None:
-             item = (None if not self.data
-                     else self.cdata_separator.join(self.data))
-
-         should_continue = self.item_callback(self.path, item)
-         if not should_continue:
-             raise ParsingInterrupted
-         # Reset state for the parent context without keeping a reference to
-         # the emitted item.
-         if self.stack:
-             self.item, self.data = self.stack.pop()
-         else:
-             self.item = None
-             self.data = []
+
+         # Apply the same transformations as the non-streaming path
+         if self.strip_whitespace and data:
+             data = data.strip() or None
+         if data and self._should_force_cdata(name, data) and item is None:
+             item = self.dict_constructor()
+         if item is not None:
+             if data:
+                 self.push_data(item, self.cdata_key, data)
+                 callback_item = item
+         else:

```

```

+         callback_item = data
+
+         # Apply postprocessor for the element name (as push_data would do for parent)
+         if self.postprocessor is not None:
+             result = self.postprocessor(self.path, name, callback_item)
+             if result is not None:

```

Showing 40 of 50 diff line(s); truncated.

3e36cea5 · cli:cli/claude-sonnet · seed 5 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	xmltodict.py	ok	12 166
1	read_file	xmltodict.py	ok	11 706
2	edit_file	xmltodict.py	ok	50
3	run_tests	—	ok	183
4	submit	—	ok	85

```

--- a/xmltodict.py
+++ b/xmltodict.py
@@ -107,21 +107,38 @@
     # without attaching it back to its parent. This avoids accumulating all
     # streamed items in memory when using item_depth > 0.
     if len(self.path) == self.item_depth:
+
+         data = (None if not self.data
+                 else self.cdata_separator.join(self.data))
+         item = self.item
+         if item is None:
+             item = (None if not self.data
+                     else self.cdata_separator.join(self.data))
+
+         should_continue = self.item_callback(self.path, item)
+         if not should_continue:
+             raise ParsingInterrupted
+         # Reset state for the parent context without keeping a reference to
+         # the emitted item.
+         if self.stack:
+             self.item, self.data = self.stack.pop()
+         else:
+             self.item = None
+             self.data = []
+
+         # Apply the same processing as the non-streaming branch
+         if self.strip_whitespace and data:
+             data = data.strip() or None
+         if data and self._should_force_cdata(name, data) and item is None:
+             item = self.dict_constructor()
+         if item is not None:
+             if data:
+                 self.push_data(item, self.cdata_key, data)
+                 final_value = item
+             else:
+                 final_value = data
+
+         # Replicate what push_data(parent, name, final_value) would compute
+         # (postprocessor + force_list) without actually storing in parent
+         if self.postprocessor is not None:
+             result = self.postprocessor(self.path, name, final_value)
+             if result is None:
+                 self.path.pop()

```

Showing 40 of 50 diff line(s); truncated.

6516049a · cli:cli/claude-sonnet · seed 6 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	xmltodict.py	ok	12 166
1	read_file	xmltodict.py	ok	11 706
2	invalid_tool_call	(none)	failed	107
3	run_tests	—	ok	183
4	submit	—	ok	85

This rollout changed no file, so there is no diff to show.

8c122607 · cli:cli/claude-sonnet · seed 7 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	xmltodict.py	ok	12 166
1	edit_file	xmltodict.py	ok	50
2	run_tests	—	ok	183
3	submit	—	ok	85

```
--- a/xmltodict.py
+++ b/xmltodict.py
@@ -107,21 +107,37 @@
     # without attaching it back to its parent. This avoids accumulating all
     # streamed items in memory when using item_depth > 0.
     if len(self.path) == self.item_depth:
+
+         data = (None if not self.data
+                 else self.cdata_separator.join(self.data))
         item = self.item
-
-         if item is None:
-             item = (None if not self.data
-                     else self.cdata_separator.join(self.data))
-
-         should_continue = self.item_callback(self.path, item)
-         if not should_continue:
-             raise ParsingInterrupted
-         # Reset state for the parent context without keeping a reference to
-         # the emitted item.
+         # Restore parent state (same ordering as non-streaming branch)
         if self.stack:
             self.item, self.data = self.stack.pop()
         else:
             self.item = None
             self.data = []
+
+         # Apply the same processing as the non-streaming branch
+         if self.strip_whitespace and data:
+             data = data.strip() or None
+         if data and self._should_force_cdata(name, data) and item is None:
+             item = self.dict_constructor()
+         if item is not None:
+             if data:
+                 self.push_data(item, self.cdata_key, data)
+             result = item
+         else:
+             result = data
+         # Apply postprocessor for the element itself
+         if self.postprocessor is not None:
+             proc_result = self.postprocessor(self.path, name, result)
```

```

+         if proc_result is not None:
+             _, result = proc_result

```

Showing 40 of 49 diff line(s); truncated.

3161d1b1 · cli:cli/claude-sonnet · seed 8 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	xmltodict.py	ok	12 166
1	read_file	xmltodict.py	ok	11 706
2	edit_file	xmltodict.py	ok	50
3	run_tests	—	ok	183
4	submit	—	ok	85

```

--- a/xmltodict.py
+++ b/xmltodict.py
@@ -107,21 +107,28 @@
     # without attaching it back to its parent. This avoids accumulating all
     # streamed items in memory when using item_depth > 0.
     if len(self.path) == self.item_depth:
+
+         data = (None if not self.data
+                 else self.cdata_separator.join(self.data))
+
         item = self.item
-
-         if item is None:
-             item = (None if not self.data
-                     else self.cdata_separator.join(self.data))
-
-         should_continue = self.item_callback(self.path, item)
-         if not should_continue:
-             raise ParsingInterrupted
-         # Reset state for the parent context without keeping a reference to
-         # the emitted item.
-         if self.stack:
-             self.item, self.data = self.stack.pop()
-         else:
-             self.item = None
-             self.data = []
+
+         # Apply the same processing as the non-streaming branch
+         if self.strip_whitespace and data:
+             data = data.strip() or None
+         if data and self._should_force_cdata(name, data) and item is None:
+             item = self.dict_constructor()
+         if item is not None:
+             if data:
+                 self.push_data(item, self.cdata_key, data)
+                 value = item
+             else:
+                 value = data
+             should_continue = self.item_callback(self.path, value)
+             if not should_continue:
+                 raise ParsingInterrupted
+             self.path.pop()
+             return
+         if self.stack:

```

Showing 40 of 40 diff line(s).

6b722882 · cli:cli/claude-sonnet · seed 9 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	xmldict.py	ok	12 166
1	read_file	xmldict.py	ok	11 706
2	invalid_tool_call	(none)	failed	107
3	edit_file	xmldict.py	ok	50
4	run_tests	—	ok	183
5	submit	—	ok	85

```
--- a/xmldict.py
+++ b/xmldict.py
@@ -107,21 +107,39 @@
     # without attaching it back to its parent. This avoids accumulating all
     # streamed items in memory when using item_depth > 0.
     if len(self.path) == self.item_depth:
+
+         data = (None if not self.data
+                 else self.cdata_separator.join(self.data))
         item = self.item
-         if item is None:
-             item = (None if not self.data
-                     else self.cdata_separator.join(self.data))
-
-         should_continue = self.item_callback(self.path, item)
-         if not should_continue:
-             raise ParsingInterrupted
-         # Reset state for the parent context without keeping a reference to
-         # the emitted item.
+
+         # Restore parent state
         if self.stack:
             self.item, self.data = self.stack.pop()
         else:
             self.item = None
             self.data = []

+
+         # Apply the same transformations as the non-streaming branch
+         if self.strip_whitespace and data:
+             data = data.strip() or None
+         if data and self._should_force_cdata(name, data) and item is None:
+             item = self.dict_constructor()
+         if item is not None:
+             if data:
+                 self.push_data(item, self.cdata_key, data)
+             result = item
+         else:
+             result = data
+
+         # Apply postprocessor for the element itself (mirrors push_data)
+         if self.postprocessor is not None:
+             entry = self.postprocessor(self.path, name, result)
```

Showing 40 of 51 diff line(s); truncated.

92c89eb1 · cli:cli/codex · seed 0 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	list_files	.	ok	39
1	read_file	ISSUE.md	ok	1 282

STEP	TOOL	PATH	RESULT	CHARS
2	read_file	xmltodict.py	ok	12 166
3	edit_file	xmltodict.py	ok	50
4	run_tests	—	ok	183
5	submit	—	ok	85

```

--- a/xmltodict.py
+++ b/xmltodict.py
@@ -107,10 +107,28 @@
     # without attaching it back to its parent. This avoids accumulating all
     # streamed items in memory when using item_depth > 0.
     if len(self.path) == self.item_depth:
+
+         data = (None if not self.data
+                 else self.cdata_separator.join(self.data))
+         item = self.item
-         if item is None:
-             item = (None if not self.data
+                     else self.cdata_separator.join(self.data))
+         if self.strip_whitespace and data:
+             data = data.strip() or None
+         if data and self._should_force_cdata(name, data) and item is None:
+             item = self.dict_constructor()
+         if item is not None:
+             if data:
+                 self.push_data(item, self.cdata_key, data)
+         else:
+             item = data
+
+         # Apply the same element-level postprocessing and force-list logic
+         # that push_data() would apply when attaching this value to its
+         # parent, but use a temporary container so streaming still drops
+         # the completed element instead of retaining it in the document.
+         container = self.push_data(None, name, item)
+         if container:
+             item = container[next(iter(container))]
+         else:
+             item = None
+
+         should_continue = self.item_callback(self.path, item)
+         if not should_continue:

```

Showing 34 of 34 diff line(s).

6d48c276 · cli:cli/codex · seed 1 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	list_files	.	ok	39
1	read_file	ISSUE.md	ok	1 282
2	read_file	xmltodict.py	ok	12 166
3	edit_file	xmltodict.py	ok	50
4	run_tests	—	ok	183
5	submit	—	ok	85

```

--- a/xmltodict.py
+++ b/xmltodict.py
@@ -107,12 +107,34 @@
     # without attaching it back to its parent. This avoids accumulating all
     # streamed items in memory when using item_depth > 0.
     if len(self.path) == self.item_depth:

```

```

+         # Finalize the streamed element exactly as the ordinary branch
+         # would, but use a temporary parent so the document is not
+         # accumulated in streaming mode.
+         data = (None if not self.data
+                 else self.cdata_separator.join(self.data))
+         item = self.item
-         if item is None:
-             item = (None if not self.data
-                     else self.cdata_separator.join(self.data))
-
-         should_continue = self.item_callback(self.path, item)
+         if self.strip_whitespace and data:
+             data = data.strip() or None
+         if data and self._should_force_cdata(name, data) and item is None:
+             item = self.dict_constructor()
+         if item is not None:
+             if data:
+                 self.push_data(item, self.cdata_key, data)
+                 value = item
+             else:
+                 value = data
+
+         # In a full parse, the completed value is passed through
+         # push_data when it is installed under its element name. That is
+         # where the element-level postprocessor and force_list handling
+         # occur. Build only that one-entry wrapper, then discard it.
+         wrapper = self.push_data(None, name, value)
+         if wrapper is None:
+             value = None
+         else:
+             value = next(iter(wrapper.values()))
+
+         should_continue = self.item_callback(self.path, value)
+         if not should_continue:

```

Showing 40 of 42 diff line(s); truncated.

1d32b64e · cli:cli/codex · seed 2 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	list_files	.	ok	39
1	read_file	ISSUE.md	ok	1 282
2	read_file	xmltodict.py	ok	12 166
3	edit_file	xmltodict.py	ok	50
4	run_tests	—	ok	183
5	submit	—	ok	85

```

--- a/xmltodict.py
+++ b/xmltodict.py
@@ -107,10 +107,27 @@
     # without attaching it back to its parent. This avoids accumulating all
     # streamed items in memory when using item_depth > 0.
     if len(self.path) == self.item_depth:
+         data = (None if not self.data
+                 else self.cdata_separator.join(self.data))
+         item = self.item
-         if item is None:
-             item = (None if not self.data
-                     else self.cdata_separator.join(self.data))
+         if self.strip_whitespace and data:
+             data = data.strip() or None
+         if data and self._should_force_cdata(name, data) and item is None:
+             item = self.dict_constructor()

```

```

+         if item is not None:
+             if data:
+                 self.push_data(item, self.cdata_key, data)
+         else:
+             item = data
+
+         # Apply the same element-level postprocessing and force-list logic
+         # as attaching the element to its parent, but use a temporary
+         # container so streamed items are still dropped after the callback.
+         container = self.push_data(None, name, item)
+         if container is None:
+             item = None
+         else:
+             item = next(iter(container.values()))
+
+         should_continue = self.item_callback(self.path, item)
+         if not should_continue:

```

Showing 33 of 33 diff line(s).

3e53d590 · cli:cli/codex · seed 3 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	list_files	.	ok	39
1	read_file	ISSUE.md	ok	1 282
2	read_file	xmltodict.py	ok	12 166
3	edit_file	xmltodict.py	ok	50
4	run_tests	—	ok	183
5	submit	—	ok	85

```

--- a/xmltodict.py
+++ b/xmltodict.py
@@ -107,10 +107,28 @@
     # without attaching it back to its parent. This avoids accumulating all
     # streamed items in memory when using item_depth > 0.
     if len(self.path) == self.item_depth:
+         data = (None if not self.data
+                 else self.cdata_separator.join(self.data))
+         item = self.item
-         if item is None:
-             item = (None if not self.data
-                     else self.cdata_separator.join(self.data))
+         if self.strip_whitespace and data:
+             data = data.strip() or None
+         if data and self._should_force_cdata(name, data) and item is None:
+             item = self.dict_constructor()
+         if item is not None:
+             if data:
+                 self.push_data(item, self.cdata_key, data)
+         else:
+             item = data
+
+         # Apply the same element-level postprocessing and force-list logic
+         # that push_data applies when attaching an ordinary element to its
+         # parent, but keep the temporary wrapper local so streaming does
+         # not retain the completed item.
+         wrapped_item = self.push_data(None, name, item)
+         if wrapped_item is None:
+             item = None
+         else:
+             item = next(iter(wrapped_item.values()))

```

```
should_continue = self.item_callback(self.path, item)
if not should_continue:
```

Showing 34 of 34 diff line(s).

90405084 · cli:cli/codex · seed 4 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	list_files	.	ok	39
1	read_file	ISSUE.md	ok	1 282
2	read_file	xmltodict.py	ok	12 166
3	edit_file	xmltodict.py	ok	50
4	run_tests	—	ok	183
5	submit	—	ok	85

```
--- a/xmltodict.py
+++ b/xmltodict.py
@@ -107,21 +107,36 @@
     # without attaching it back to its parent. This avoids accumulating all
     # streamed items in memory when using item_depth > 0.
     if len(self.path) == self.item_depth:
+
+         data = (None if not self.data
+                 else self.cdata_separator.join(self.data))
+         item = self.item
+
+         if item is None:
+             item = (None if not self.data
+                     else self.cdata_separator.join(self.data))
+
+         if self.stack:
+             self.item, self.data = self.stack.pop()
+         else:
+             self.item = None
+             self.data = []
+
+         if self.strip_whitespace and data:
+             data = data.strip() or None
+
+         if data and self._should_force_cdata(name, data) and item is None:
+             item = self.dict_constructor()
+
+         if item is not None:
+             if data:
+                 self.push_data(item, self.cdata_key, data)
+         else:
+             item = data
+
+         # Apply the same element-level postprocessing and force-list logic
+         # as push_data would when attaching this value to its parent, but
+         # use a temporary container so the streamed document is not kept.
+         emitted = self.push_data(None, name, item)
+         if emitted is None:
+             item = None
+         else:
+             item = next(iter(emitted.values()))

        should_continue = self.item_callback(self.path, item)
        if not should_continue:
            raise ParsingInterrupted
-
-     # Reset state for the parent context without keeping a reference to
```

Showing 40 of 49 diff line(s); truncated.

48e359d8 · cli:cli/codex · seed 5 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	list_files	.	ok	39
1	read_file	ISSUE.md	ok	1 282
2	read_file	xmltodict.py	ok	12 166
3	edit_file	xmltodict.py	ok	50
4	run_tests	—	ok	183
5	submit	—	ok	85

```
--- a/xmltodict.py
+++ b/xmltodict.py
@@ -107,11 +107,25 @@
     # without attaching it back to its parent. This avoids accumulating all
     # streamed items in memory when using item_depth > 0.
     if len(self.path) == self.item_depth:
+
+         data = (None if not self.data
+                 else self.cdata_separator.join(self.data))
         item = self.item
-         if item is None:
-             item = (None if not self.data
-                     else self.cdata_separator.join(self.data))
-
+         if self.strip_whitespace and data:
+             data = data.strip() or None
+         if data and self._should_force_cdata(name, data) and item is None:
+             item = self.dict_constructor()
+         if item is not None:
+             if data:
+                 self.push_data(item, self.cdata_key, data)
+                 emitted = self.push_data(None, name, item)
+             else:
+                 emitted = self.push_data(None, name, data)
+
+         # Build the value through the same push_data path used by ordinary
+         # parsing, so element-level postprocessing and force_list apply,
+         # but do not retain the temporary containing mapping.
+         item = (None if emitted is None
+                 else next(iter(emitted.values())))
        should_continue = self.item_callback(self.path, item)
        if not should_continue:
            raise ParsingInterrupted
```

Showing 32 of 32 diff line(s).

fd7561b3 · cli:cli/codex · seed 6 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	list_files	.	ok	39
1	read_file	ISSUE.md	ok	1 282
2	read_file	xmltodict.py	ok	12 166
3	edit_file	xmltodict.py	ok	50
4	run_tests	—	ok	183
5	submit	—	ok	85

```

--- a/xmltodict.py
+++ b/xmltodict.py
@@ -107,10 +107,28 @@
     # without attaching it back to its parent. This avoids accumulating all
     # streamed items in memory when using item_depth > 0.
     if len(self.path) == self.item_depth:
+
+         data = (None if not self.data
+                 else self.cdata_separator.join(self.data))
         item = self.item
-
-         if item is None:
-             item = (None if not self.data
-                     else self.cdata_separator.join(self.data))
+
+         if self.strip_whitespace and data:
+             data = data.strip() or None
+         if data and self._should_force_cdata(name, data) and item is None:
+             item = self.dict_constructor()
+         if item is not None:
+             if data:
+                 item = self.push_data(item, self.cdata_key, data)
+
+         else:
+             item = data
+
+         # Apply the same element-level postprocessing and force-list logic
+         # that push_data applies when attaching a completed element to its
+         # parent, but keep it in a temporary container so streaming does
+         # not retain the completed item.
+         container = self.push_data(None, name, item)
+         if container:
+             item = next(iter(container.values()))
+         else:
+             item = None

        should_continue = self.item_callback(self.path, item)
        if not should_continue:

```

Showing 34 of 34 diff line(s).

4fe31a46 · cli:cli/codex · seed 7 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	list_files	.	ok	39
1	read_file	ISSUE.md	ok	1 282
2	read_file	xmltodict.py	ok	12 166
3	edit_file	xmltodict.py	ok	50
4	run_tests	—	ok	183
5	submit	—	ok	85

```

--- a/xmltodict.py
+++ b/xmltodict.py
@@ -107,10 +107,24 @@
     # without attaching it back to its parent. This avoids accumulating all
     # streamed items in memory when using item_depth > 0.
     if len(self.path) == self.item_depth:
+
+         data = (None if not self.data
+                 else self.cdata_separator.join(self.data))
         item = self.item
-
-         if item is None:
-             item = (None if not self.data
-                     else self.cdata_separator.join(self.data))
+
+         if self.strip_whitespace and data:
+             data = data.strip() or None
+         if data and self._should_force_cdata(name, data) and item is None:

```

```

+         item = self.dict_constructor()
+         if item is not None:
+             if data:
+                 self.push_data(item, self.cdata_key, data)
+         else:
+             item = data
+
+         # Apply the element-level postprocessor and force_list handling in
+         # exactly the same way as the ordinary parent-attachment path, but
+         # use a temporary container so streamed items are still discarded.
+         container = self.push_data(None, name, item)
+         item = next(iter(container.values()), None)
+
+         should_continue = self.item_callback(self.path, item)
+         if not should_continue:

```

Showing 30 of 30 diff line(s).

18660280 · cli:cli/codex · seed 8 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	list_files	.	ok	39
1	read_file	ISSUE.md	ok	1 282
2	read_file	xmltodict.py	ok	12 166
3	edit_file	xmltodict.py	ok	50
4	run_tests	—	ok	183
5	submit	—	ok	85

```

--- a/xmltodict.py
+++ b/xmltodict.py
@@ -103,27 +103,6 @@

def endElement(self, full_name):
    name = self._build_name(full_name)
-    # If we just closed an item at the streaming depth, emit it and drop it
-    # without attaching it back to its parent. This avoids accumulating all
-    # streamed items in memory when using item_depth > 0.
-    if len(self.path) == self.item_depth:
-        item = self.item
-        if item is None:
-            item = (None if not self.data
-                   else self.cdata_separator.join(self.data))
-
-        should_continue = self.item_callback(self.path, item)
-        if not should_continue:
-            raise ParsingInterrupted
-        # Reset state for the parent context without keeping a reference to
-        # the emitted item.
-        if self.stack:
-            self.item, self.data = self.stack.pop()
-        else:
-            self.item = None
-            self.data = []
-        self.path.pop()
-        return
    if self.stack:
        data = (None if not self.data
               else self.cdata_separator.join(self.data))
@@ -135,10 +114,26 @@
        item = self.dict_constructor()
        if item is not None:
            if data:
-                self.push_data(item, self.cdata_key, data)

```

```

+         item = self.push_data(item, self.cdata_key, data)
+     else:
+         item = data
+
+     # Build the same value that push_data would place under this

```

Showing 40 of 60 diff line(s); truncated.

195b8f72 · cli:cli/codex · seed 9 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	list_files	.	ok	39
1	read_file	ISSUE.md	ok	1 282
2	read_file	xmltodict.py	ok	12 166
3	edit_file	xmltodict.py	ok	50
4	run_tests	—	ok	183
5	submit	—	ok	85

```

--- a/xmltodict.py
+++ b/xmltodict.py
@@ -107,10 +107,23 @@
     # without attaching it back to its parent. This avoids accumulating all
     # streamed items in memory when using item_depth > 0.
     if len(self.path) == self.item_depth:
+
+         data = (None if not self.data
+                 else self.cdata_separator.join(self.data))
+         item = self.item
+         if item is None:
+             item = (None if not self.data
+                     else self.cdata_separator.join(self.data))
+         if self.strip_whitespace and data:
+             data = data.strip() or None
+         if data and self._should_force_cdata(name, data) and item is None:
+             item = self.dict_constructor()
+         if item is not None:
+             if data:
+                 item = self.push_data(item, self.cdata_key, data)
+                 emitted = self.push_data(None, name, item)
+             else:
+                 emitted = self.push_data(None, name, data)
+             if emitted is None:
+                 item = None
+             else:
+                 item = next(iter(emitted.values()))
+
         should_continue = self.item_callback(self.path, item)
         if not should_continue:

```

Showing 29 of 29 diff line(s).

5d129402 · mistral:codestral-latest · seed 0 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	ISSUE.md	ok	1 282
1	read_file	xmltodict.py	ok	12 166
2	edit_file	xmltodict.py	ok	50
3	run_tests	—	ok	183

STEP	TOOL	PATH	RESULT	CHARS
4	submit	—	ok	85

```

--- a/xmltodict.py
+++ b/xmltodict.py
@@ -111,6 +111,17 @@
     if item is None:
         item = (None if not self.data
                 else self.cdata_separator.join(self.data))
+
+    else:
+        if self.data:
+            if self.strip_whitespace:
+                text = self.cdata_separator.join(self.data).strip()
+            else:
+                text = self.cdata_separator.join(self.data)
+            if self._should_force_cdata(name, text):
+                if self.cdata_key not in item:
+                    item[self.cdata_key] = text
+            else:
+                item[self.cdata_key] = text
+
     should_continue = self.item_callback(self.path, item)
     if not should_continue:

```

Showing 20 of 20 diff line(s).

bd30cb85 · mistral:codestral-latest · seed 1 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	list_files	.	ok	39
1	read_file	xmltodict.py	ok	12 166
2	edit_file	xmltodict.py	ok	50
3	run_tests	—	ok	183
4	submit	—	ok	85

```

--- a/xmltodict.py
+++ b/xmltodict.py
@@ -111,6 +111,16 @@
     if item is None:
         item = (None if not self.data
                 else self.cdata_separator.join(self.data))
+
+    # Apply postprocessor and other transformations
+    if self.postprocessor:
+        item = self.postprocessor(self.path, self.path[-1][0], item)
+    if self.strip_whitespace and isinstance(item, str):
+        item = item.strip() or None
+    if self.force_cdata and isinstance(item, str):
+        item = {self.cdata_key: item}
+    if self.force_list and isinstance(item, str):
+        item = [item]
+
     should_continue = self.item_callback(self.path, item)
     if not should_continue:

```

Showing 19 of 19 diff line(s).

7cc71ab5 · mistral:codestral-latest · seed 2 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	list_files	.	ok	39
1	read_file	xmltodict.py	ok	12 166
2	edit_file	xmltodict.py	ok	50
3	run_tests	—	ok	183
4	submit	—	ok	85

```
--- a/xmltodict.py
+++ b/xmltodict.py
@@ -111,6 +111,17 @@
     if item is None:
         item = (None if not self.data
                 else self.cdata_separator.join(self.data))
+
+    else:
+        if self.data:
+            if self.strip_whitespace:
+                text = self.cdata_separator.join(self.data).strip()
+            else:
+                text = self.cdata_separator.join(self.data)
+            if text:
+                if self._should_force_cdata(self.path[-1][0], text):
+                    item[self.cdata_key] = text
+                else:
+                    item[self.cdata_key] = text
+
     should_continue = self.item_callback(self.path, item)
     if not should_continue:
```

Showing 20 of 20 diff line(s).

d38fb2e2 · mistral:codestral-latest · seed 3 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	list_files	.	ok	39
1	read_file	xmltodict.py	ok	12 166
2	edit_file	xmltodict.py	ok	50
3	run_tests	—	ok	183
4	submit	—	ok	85

```
--- a/xmltodict.py
+++ b/xmltodict.py
@@ -111,6 +111,18 @@
     if item is None:
         item = (None if not self.data
                 else self.cdata_separator.join(self.data))
+
+    # Apply postprocessor and other transformations
+    if self.postprocessor:
+        item = self.postprocessor(self.path, name, item)
+    if self.strip_whitespace and isinstance(item, str):
+        item = item.strip() or None
+    if self.force_cdata and isinstance(item, str):
+        item = {self.cdata_key: item}
+    if self.force_list and isinstance(item, dict):
+        for key in self.force_list:
```

```

+         if key in item:
+             item[key] = [item[key]]

should_continue = self.item_callback(self.path, item)
if not should_continue:

```

Showing 21 of 21 diff line(s).

8c65b44c · mistral:codestral-latest · seed 4 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	list_files	.	ok	39
1	read_file	xmldict.py	ok	12 166
2	edit_file	xmldict.py	ok	50
3	run_tests	—	ok	183
4	submit	—	ok	85

```

--- a/xmldict.py
+++ b/xmldict.py
@@ -111,6 +111,18 @@
     if item is None:
         item = (None if not self.data
                 else self.cdata_separator.join(self.data))

+
+
+    # Apply postprocessor and other transformations
+    if self.postprocessor:
+        item = self.postprocessor(self.path, self.path[-1][0], item)
+    if self.strip_whitespace and isinstance(item, str):
+        item = item.strip() or None
+    if self.force_cdata and isinstance(item, str):
+        item = {self.cdata_key: item}
+    if self.force_list and isinstance(item, dict):
+        for key in self.force_list:
+            if key in item:
+                item[key] = [item[key]]

should_continue = self.item_callback(self.path, item)
if not should_continue:

```

Showing 21 of 21 diff line(s).

b8a625a2 · mistral:codestral-latest · seed 5 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	list_files	.	ok	39
1	read_file	xmldict.py	ok	12 166
2	edit_file	xmldict.py	ok	50
3	run_tests	—	ok	183
4	submit	—	ok	85

```

--- a/xmldict.py
+++ b/xmldict.py
@@ -111,6 +111,16 @@
     if item is None:
         item = (None if not self.data
                 else self.cdata_separator.join(self.data))

```

```

+         else:
+             if self.data:
+                 if self.strip_whitespace:
+                     text = self.cdata_separator.join(self.data).strip()
+                 else:
+                     text = self.cdata_separator.join(self.data)
+                 if self._should_force_cdata(name, text):
+                     item[self.cdata_key] = text
+                 else:
+                     item[self.cdata_key] = text

        should_continue = self.item_callback(self.path, item)
        if not should_continue:

```

Showing 19 of 19 diff line(s).

58abb85c · mistral:codestral-latest · seed 6 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	list_files	.	ok	39
1	read_file	xmltodict.py	ok	12 166
2	edit_file	xmltodict.py	ok	50
3	run_tests	—	ok	183
4	submit	—	ok	85

```

--- a/xmltodict.py
+++ b/xmltodict.py
@@ -111,6 +111,18 @@
         if item is None:
             item = (None if not self.data
                     else self.cdata_separator.join(self.data))

+
+         # Apply postprocessor and other transformations
+         if self.postprocessor:
+             item = self.postprocessor(self.path, name, item)
+         if self.strip_whitespace and isinstance(item, str):
+             item = item.strip() or None
+         if self.force_cdata and isinstance(item, str):
+             item = {self.cdata_key: item}
+         if self.force_list and isinstance(item, dict):
+             for key in item:
+                 if self._should_force_list(key, item[key]):
+                     item[key] = [item[key]]

        should_continue = self.item_callback(self.path, item)
        if not should_continue:

```

Showing 21 of 21 diff line(s).

94122fe4 · mistral:codestral-latest · seed 7 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	list_files	.	ok	39
1	read_file	xmltodict.py	ok	12 166
2	edit_file	xmltodict.py	ok	50
3	run_tests	—	ok	183
4	submit	—	ok	85

```

--- a/xmltodict.py
+++ b/xmltodict.py
@@ -111,6 +111,16 @@
     if item is None:
         item = (None if not self.data
                 else self.cdata_separator.join(self.data))
+
+    else:
+        if self.data:
+            data = self.cdata_separator.join(self.data)
+            if self.strip_whitespace:
+                data = data.strip() or None
+            if data:
+                if self._should_force_cdata(self.path[-1][0], data):
+                    item[self.cdata_key] = data
+                else:
+                    item[self.cdata_key] = data

    should_continue = self.item_callback(self.path, item)
    if not should_continue:

```

Showing 19 of 19 diff line(s).

b0843fec · mistral:codestral-latest · seed 8 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	ISSUE.md	ok	1 282
1	read_file	xmltodict.py	ok	12 166
2	edit_file	xmltodict.py	ok	50
3	run_tests	—	ok	183
4	submit	—	ok	85

```

--- a/xmltodict.py
+++ b/xmltodict.py
@@ -102,6 +102,36 @@
     self.data = []

    def endElement(self, full_name):
+
+        name = self._build_name(full_name)
+        # If we just closed an item at the streaming depth, emit it and drop it
+        # without attaching it back to its parent. This avoids accumulating all
+        # streamed items in memory when using item_depth > 0.
+        if len(self.path) == self.item_depth:
+            item = self.item
+            if item is None:
+                item = (None if not self.data
+                       else self.cdata_separator.join(self.data))
+
+            # Apply the same processing as in non-streaming mode
+            if self.strip_whitespace and item:
+                item = item.strip() or None
+            if item and self._should_force_cdata(name, item):
+                item = {self.cdata_key: item}
+            if self.postprocessor:
+                item = self.postprocessor(self.path, name, item)
+
+            should_continue = self.item_callback(self.path, item)
+            if not should_continue:
+                raise ParsingInterrupted
+            # Reset state for the parent context without keeping a reference to
+            # the emitted item.
+            if self.stack:
+                self.item, self.data = self.stack.pop()
+            else:

```

```

+         self.item = None
+         self.data = []
+         self.path.pop()
+         return
name = self._build_name(full_name)
# If we just closed an item at the streaming depth, emit it and drop it
# without attaching it back to its parent. This avoids accumulating all

```

Showing 39 of 39 diff line(s).

cecbcc3b · mistral:codestral-latest · seed 9 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	list_files	.	ok	39
1	read_file	xmltodict.py	ok	12 166
2	edit_file	xmltodict.py	ok	50
3	run_tests	—	ok	183
4	submit	—	ok	85

```

--- a/xmltodict.py
+++ b/xmltodict.py
@@ -111,6 +111,18 @@
     if item is None:
         item = (None if not self.data
                 else self.cdata_separator.join(self.data))

+
+     # Apply postprocessor and other transformations
+     if self.postprocessor:
+         item = self.postprocessor(self.path, self.path[-1][0], item)
+     if self.strip_whitespace and isinstance(item, str):
+         item = item.strip() or None
+     if self.force_cdata and isinstance(item, str):
+         item = {self.cdata_key: item}
+     if self.force_list and isinstance(item, dict):
+         for key in self.force_list:
+             if key in item:
+                 item[key] = [item[key]]

     should_continue = self.item_callback(self.path, item)
     if not should_continue:

```

Showing 21 of 21 diff line(s).

10 Exam · order-desk-1

vvdex.browser.order-desk-1 · record fr-20260902-6e618c05 · record digest 59f94854d78f7ce1... · harness clean · containme
nt clean

Executive summary

MODELS TESTED 3	CERTIFIED PASSES 21 of 29 graded rollouts	SUBMITTED 29 of 29 graded rollouts that called submit	NOT GRADED — LANE ERRORS 1 of 30 attempts
---------------------------------------	--	--	--

Purpose

This report presents a controlled evaluation of 3 model lane(s) against the certified exam vv
dex.browser.order-desk-1 (v0.1.0, family browser). Every lane received the identical
frozen workspace, tool contract and limits; grading ran afterwards, in isolation, against
hidden tests no lane ever saw. The purpose is to state, with reproducible evidence, what
each lane did and where it stopped.

Outcome

21 of 29

graded rollouts passed the independent hidden tests and were certified — 1 of 30 attempts not graded (lane errors)

21 of 29 graded rollouts passed. cli/claude-sonnet, cli/codex, codestral-latest completed
the required browser workflow, submitted, and were accepted by the independent grader.
The most common way to fail was submitted_failed — submitted a final application state
that the independent grader rejected.

Key findings

- F-01 cli/claude-sonnet, cli/codex, codestral-latest passed the independent hidden tests
- F-02 submitted_failed × 8 (codestral-latest)
- F-03 model_api_failure × 1 (codestral-latest)

Each finding is stated in full, with evidence, in the Findings section; recommendations for
acting on them follow it. Elapsed wall clock for the whole engagement: 3292.4s.

Exam definition

In scope. The ability of each configured model lane to complete this one certified task end to end — inspect and open the application, drive it through the declared browser actions to the required state, submit, and be accepted by the independent grader — within 40 tool steps and 9720s of wall clock, at 10 rollout(s) per lane. Behavioural measurements along the way (tool discipline, grounding, overclaim) are in scope because they are read from the recorded trace, not judged by another model.

Out of scope. General model quality, performance on other task families, prompt-tuning on the lane's behalf, and any comparison this run's sample size cannot support. A lane's API failure is reported as infrastructure, never as model capability.

The instrument. Exam `vvdex.browser.order-desk-1` was certified before this run: its reference solution passes, an empty submission fails, its grader distinguishes near-misses, and its sealed evidence is unchanged by this evaluation (see Certification evidence).

What was measured

EXAM	VERSION	FAMILY	ROLLOUTS
vvdex.browser.order-desk-1	0.1.0	browser	30 (10 per model)

Where the defect came from

This exam has no upstream repository to credit: the defect, the reference solution and the hidden suite are VVDex's own.

Certification evidence

4 of 5 certification pillars are recorded in this package, loaded from the receipts sealed into the exam and matching this run's exam fingerprint. What is recorded is shown with its real value; what is not is named below and is not claimed.

FROZEN BASELINE

FAIL

expected — an empty submission over 2 run(s) must not pass

→

GOLD / REFERENCE

PASS

expected — the reference solution over 2 run(s) must pass

GRADER CONTROLS

9 / 9

The grader accepted the reference and rejected every near-miss.

ATTACK PROBES

10 / 10 blocked

0 trivial exploit(s) found.

RUNTIME

sha256:679b36225a1f83238efba8c71b9cde3c24caaeacc9b682ae92819cb55eec328f

network policy: deny

Certification metadata is not included in this report package for:

- sealed evidence bundle

Benchmark results remain valid as run evidence; the pillar(s) above are not substantiated by this document and this report does not claim them.

What the package does carry: the exam fingerprint below is the contract digest this run was executed against — a holder of the certification bundle for this fingerprint can join the two independently.

EXAM FINGERPRINT

8d404ee5d5ab71456de3d3df0b8ca69271d8810767a694f71c6e45e72fc347bf

The contract digest this run was executed against.

Evaluation configuration

REPORT	fr-20260902-6e618c05
ISSUED	2026-09-02
SUBJECT EXAM	vvdex.browser.order-desk-1 · v0.1.0
FAMILY	browser
PREPARED BY	VVDex Forge engine vvdex-env 0.1.0
CLASSIFICATION	Independent model evaluation report — independently verifiable
CERTIFICATION	Exam certification: partially recorded (see Certification evidence)
STATUS	FINAL · report sealed against its own digest

Timeline

WHEN (UTC)	EVENT
2026-09-02 14:49:56	Evaluation run started
2026-09-02 15:44:48	Evaluation run finished
2026-09-02 15:44:48	Report generated from the sealed run evidence

Models under test

MODEL	PROVIDER	ENDPOINT	BILLING
Claude Code CLI (Sonnet, subscription)	cli	local runtime	operator subscription, flat — not billed per token
Codex CLI (GPT, subscription)	cli	local runtime	operator subscription, flat — not billed per token
codestral-latest	mistral	VVDex-configured	VVDex free-quota

A customer endpoint is recorded by origin only — never its port, its path, or its key.

Limits

RUNTIME docker	STEP LIMIT 40	WALL-CLOCK LIMIT 9720s	MAX OUTPUT TOKENS 2048
TEMPERATURE 0.2	RETRY POLICY none	COMMAND BUDGET 0	TOOLS OFFERED list_files, read_file, submit, browser_open, browser_snapshot, browser_click, browser_fill, browser_select, browser_press, browser_back

The evaluated model never saw the hidden tests or the reference solution, and could not change its result. Grading ran in a separate step, after submission, on a container with no network.

Model outcomes

Denominators. Attempts requested: 30. Graded (evaluable) rollouts: 29. Not graded — lane errors: 1. A lane error is a rollout the API or the runtime ended before the hidden grader ran: it is listed, excluded from every rate and pass@k, and never silently dropped. Passes are certified passes: hidden grader exit 0 with integrity verified.

MODEL	ATTEMPTS	GRADED	PASSED	MODEL FAILS	LANE ERRORS	SUBMITTED	MEAN TIME	TYPICAL DEEPEST STAGE
Claude Code CLI (Sonnet, subscription) cli/claude-sonnet	10	10	10 / 10	0	0	10 / 10	89.4s	Passed hidden tests
Codex CLI (GPT, subscription) cli/codex	10	10	10 / 10	0	0	10 / 10	125.2s	Passed hidden tests
codestral-latest codestral-latest	10	9	1 / 9	8	1	9 / 9	84.4s	Graded

Every rollout, with its own deepest stage and grade, is in Appendix A.

Success rate, confidence and pass@k

Every rate on this page is a measurement of a finite number of rollouts, so every rate carries the interval that measurement actually supports. n is the evaluable count — attempts minus rollouts the lane ended before grading; those are listed under "not graded" and sit in no rate, interval or pass@k.

MODEL	ATTEMPTS	N	NOT GRADED	PASSES	RATE	95% INTERVAL	SEEDS	PASS@1	PASS@3	PASS@5	PASS@10
cli/claude-sonnet	10	10	0	10	100.0%	[72%, 100%]	0–9	100.0%	100.0%	100.0%	100.0%
cli/codex	10	10	0	10	100.0%	[72%, 100%]	0–9	100.0%	100.0%	100.0%	100.0%
codestral-latest	10	9	1	1	11.1%	[2%, 43%]	0–9	11.1%	33.3%	55.6%	–

Interval and pass@k methods are stated in full in Appendix D. pass@k is shown for k = 1, 3, 5, 10; the sealed record carries every k up to n.

Model comparison

MODEL	PASS@1	95% INTERVAL	TOKENS	COST	MEAN ROLLOUT	OVERCLAIM RATE	TOP FAILURE CLASS
cli/claude-sonnet	100.0%	[72%, 100%]	n/a	n/a	89.4s	0.0%	none
cli/codex	100.0%	[72%, 100%]	n/a	n/a	125.2s	0.0%	none
codestral-latest	11.1%	[2%, 43%]	857 790	not billed	84.4s	0.0%	submitted_failed

Tokens and cost read n/a where a lane reports no telemetry (a flat-rate CLI) — an absent measurement, not zero. pass@1 and its interval are over evaluable rollouts only.

This is not a leaderboard. Separated by this run (intervals do not overlap): cli/claude-sonnet vs codestral-latest; cli/codex vs codestral-latest. Not separated (intervals overlap): cli/claude-sonnet vs cli/codex. Below the 10-rollout floor for reading a rate as a rate: codestral-latest (9) — read the count, not the percentage. The primary comparison in this report is the stage funnel below — where a model stops is a finding; a percentage over this many rollouts is not.

Stage funnel

The funnel is the primary reading of a run. Where a model stops says more than a single pass percentage does, especially at low N.

STAGE	CLI:CLI/CLAUDE-SONNET	CLI:CLI/CODEX	MISTRAL:CODESTRAL-LATEST
Inspected	10 / 10	10 / 10	9 / 9
Edited	N/A	N/A	N/A
Ran tests	N/A	N/A	N/A
Submitted	10 / 10	10 / 10	9 / 9
Graded	10 / 10	10 / 10	9 / 9
Passed hidden tests	10 / 10	10 / 10	1 / 9

Deepest stage reached

MODEL	DEEPEST (ANY ROLLOUT)	TYPICAL (MOST ROLLOUTS)	HIDDEN-TEST PASSES	NOT GRADED
cli:cli/claude-sonnet	Passed hidden tests	Passed hidden tests	10 / 10	0
cli:cli/codex	Passed hidden tests	Passed hidden tests	10 / 10	0
mistral:codestral-latest	Passed hidden tests	Graded	1 / 9	1

Stage definitions are in Appendix D. The matrix counts evaluable rollouts; lane errors are shown beside it, not inside it. The stages are not strictly nested: a model can submit without editing, and the matrix shows that rather than hiding it. N/A marks a stage this exam does not offer (Edited, Ran tests): not reached, not failed — not asked. Each rollout's own deepest stage is in Appendix A.

Behavioural analysis

Long-horizon behaviour

Counted off the recorded trace of each rollout. A dead end is a step that moved nothing: a re-read of a slice already returned, a repeat of the previous action, or a call to a tool this exam does not have.

MODEL	N	STEPS USED	RAN OUT OF STEPS	COMMANDS	CHANGED THE TREE	REFUSED (OVER BUDGET)	REVISITS	DEAD ENDS / ROLLOUT
cli:cli/claude-sonnet	10	18 / 40	0	0 / 0	0	0	0	3
cli:cli/codex	10	20 / 40	0	0 / 0	0	0	0	2
mistral:codestral-latest	10	22 / 40	0	0 / 0	0	0	0	1

Each column is defined in Appendix D.

Hallucination & overclaim

Three measurements, each with a stated definition. None of them is a judgement of the model's prose by another model — every number below comes from comparing what the model wrote against what was on disk, or against what the independent grader returned.

MODEL	N	GROUNDING FAILURES	KINDS	ROLLOUTS AFFECTED	OVERCLAIMS	OVERCLAIM RATE	RE-READS
cli:cli/claude-sonnet	10	0	none	0 / 10	0 / 10	0.0%	0
cli:cli/codex	10	0	none	0 / 10	0 / 10	0.0%	0
mistral:codestral-latest	10	0	none	0 / 10	0 / 10	0.0%	0

No grounding failures and no false pass claims were detected across these 30 rollout(s). No lane re-read material it had already been given.

Every term above is defined in Appendix D.

Efficiency & telemetry

This run has no per-token price attached — a local model, a free quota, or a customer endpoint we are not billed for — so spend is reported in tokens. Tokens are what was actually measured; a dollar figure here would be invented. Lanes run through a flat-rate local CLI report no token telemetry; their cells say n/a — an absent measurement, not zero. Rollout counts here are evaluable rollouts.

MODEL	ROLLOUTS	TOKENS	TOKENS / ROLLOUT	COST	PASSES	RATE
cli:cli/claude-sonnet	10	n/a	n/a	n/a – telemetry unavailable	10 / 10	100.0%
cli:cli/codex	10	n/a	n/a	n/a – telemetry unavailable	10 / 10	100.0%
mistral:codestral-latest	9	857 790	95 310	not billed	1 / 9	11.1%

Findings

Every finding below is derived from the recorded data of this run — none is an opinion.
Severity: PASS a lane succeeded · ATTENTION a capability gap · LANE / INFRA infrastructure, not capability · NOTE a property of the measurement.

F-01	PASS
cli/claude-sonnet, cli/codex, codestral-latest passed the independent hidden tests	
Severity: PASS	Evidence: Model outcomes · Stage funnel · Appendix A
21 of 29 graded rollouts completed the required browser workflow, submitted, and were accepted by the independent grader. This demonstrates the task is solvable under the published limits.	
F-01 · derived from the recorded run data	

F-02	ATTENTION
submitted_failed × 8 (codestral-latest)	
Severity: ATTENTION	Evidence: Appendix A · Appendix C
Submitted a final application state that the independent grader rejected.	
F-02 · derived from the recorded run data	

F-03	LANE / INFRA
model_api_failure × 1 (codestral-latest)	
Severity: LANE / INFRA	Evidence: Appendix A · Appendix C
The model endpoint failed to answer. Not a verdict about the model's ability. These rollouts are excluded from any capability reading.	
F-03 · derived from the recorded run data	

Recommendations

One recommendation per observed failure mode, addressed to the lane it affects. Each traces back to a finding above and to the raw trace in Appendix A.

R-01

PROBLEM

Submitted a final application state that the independent grader rejected.

AFFECTED LANES

codestral-latest

WHAT WE WOULD LOOK AT FIRST

The model drives the application and submits a final state the grader rejects — the named failing axes say which part of the required state was wrong (the target record, a collateral change, a skipped confirmation, an uncorrected input).

R-02

PROBLEM

The model endpoint failed to answer. Not a verdict about the model's ability.

AFFECTED LANES

codestral-latest

WHAT WE WOULD LOOK AT FIRST

The lane failed, not the model; these rollouts carry no signal about capability and are labeled accordingly.

Verification

RUN ID

live-20260902T144956Z

EXAM FINGERPRINT

8d404ee5d5ab71456de3d3df0b8ca69271d8810767a694f71c6e45e72fc347bf

The contract digest this run was executed against.

Provenance of this chapter

This chapter was rendered by the campaign generator from the sealed evaluation record; it is not the standalone report reproduced byte for byte. Verify the standalone report with its own digest, verify the record with `vvdex-env records verify`, and verify this campaign document as a whole.

RELATIONSHIP

campaign-rendered chapter

Derived from the same sealed evaluation record as the standalone report; these bytes are authenticated by the campaign artifact that contains them.

SEALED RECORD DIGEST

59f94854d78f7ce174866ff114505b2b41ca2333f287f40b604fbb38cc267da5

sha256 over the record's canonical JSON — the identity every renderer works from.

SOURCE RECORD FILE

251730136fa0efbdd157d980d037447c967af4be4aaf1f04ce2c7a07a605bea

sha256 of `fr-20260902-6e618c05.json` as written beside the standalone report.

SOURCE STANDALONE REPORT

1f79c3ac34a62b68661217a9bc94e0d94fc42f68dcecf59d7919e9fd97f71585

sha256 of `fr-20260902-6e618c05.html` — independently verifiable with its own digest.

SOURCE STANDALONE PDF

522ca583bd5fe165c7a410e3b66dc4fe13dd18740fe2111ca27f72d9e92f09e0

sha256 of `fr-20260902-6e618c05.pdf`.

RUN EVIDENCE BUNDLE

6e618c0542eb41677ef2556126365aea5538ce4f4a959245e1e09e1732ae8620

The digest the run's own bundle computed over itself.

CERTIFIED EVIDENCE SEAL

not available in this package

The exam's sealed evidence, established before this run.

Measurement history

Previous run on this exam: `fr-20260902-c05474ed`. Model progress over time on this exam is the difference between that record and this one — same frozen tree, same hidden grader, same limits. Anything else that moved between them moved outside this exam.

The full artifact-by-artifact digest table, and the reproduction protocol, are in Appendix B.

Appendix A — every rollout

ROLLOUT	MODEL	SEED	OUTCOME	SUBMITTED	DEEPEST STAGE	HIDDEN TESTS	FILES	TOKENS	ELAPSED	BROWSER BOX
532e1db3	cli/claude-sonnet	0	submitted_passed	yes	Passed hidden tests	pass	0	n/a	137.9s	docker sha256:c5587afc e772...
d641f68a	cli/claude-sonnet	1	submitted_passed	yes	Passed hidden tests	pass	0	n/a	96.5s	docker sha256:c5587afc e772...
b4f9ae08	cli/claude-sonnet	2	submitted_passed	yes	Passed hidden tests	pass	0	n/a	89.4s	docker sha256:c5587afc e772...
da5c8bda	cli/claude-sonnet	3	submitted_passed	yes	Passed hidden tests	pass	0	n/a	79.5s	docker sha256:c5587afc e772...
a116d4fa	cli/claude-sonnet	4	submitted_passed	yes	Passed hidden tests	pass	0	n/a	86.6s	docker sha256:c5587afc e772...
d6987896	cli/claude-sonnet	5	submitted_passed	yes	Passed hidden tests	pass	0	n/a	75.1s	docker sha256:c5587afc e772...
e52d496d	cli/claude-sonnet	6	submitted_passed	yes	Passed hidden tests	pass	0	n/a	80.6s	docker sha256:c5587afc e772...
7f40f719	cli/claude-sonnet	7	submitted_passed	yes	Passed hidden tests	pass	0	n/a	85.9s	docker sha256:c5587afc e772...
9b8db9ca	cli/claude-sonnet	8	submitted_passed	yes	Passed hidden tests	pass	0	n/a	79.3s	docker sha256:c5587afc e772...
38d7cace	cli/claude-sonnet	9	submitted_passed	yes	Passed hidden tests	pass	0	n/a	83.1s	docker sha256:c5587afc e772...
6bd00d3c	cli/codex	0	submitted_passed	yes	Passed hidden tests	pass	0	n/a	109.6s	docker sha256:c5587afc e772...
b3f86ec4	cli/codex	1	submitted_passed	yes	Passed hidden tests	pass	0	n/a	120.7s	docker sha256:c5587afc e772...
0d1dabed	cli/codex	2	submitted_passed	yes	Passed hidden tests	pass	0	n/a	147.3s	docker sha256:c5587afc e772...
e415a6c4	cli/codex	3	submitted_passed	yes	Passed hidden tests	pass	0	n/a	137.7s	docker sha256:c5587afc e772...
e771f4b2	cli/codex	4	submitted_passed	yes	Passed hidden tests	pass	0	n/a	130.1s	docker sha256:c5587afc e772...
c8c5c1f3	cli/codex	5	submitted_passed	yes	Passed hidden tests	pass	0	n/a	115.9s	docker sha256:c5587afc e772...

ROLLOUT	MODEL	SEED	OUTCOME	SUBMITTED	DEEPEST STAGE	HIDDEN TESTS	FILES	TOKENS	ELAPSED	BROWSER BOX
c20efdb4	cli/codex	6	submitted_passed	yes	Passed hidden tests	pass	0	n/a	125.2s	docker sha256: c5587afc e772...
6015f268	cli/codex	7	submitted_passed	yes	Passed hidden tests	pass	0	n/a	128.9s	docker sha256: c5587afc e772...
bbef3c65	cli/codex	8	submitted_passed	yes	Passed hidden tests	pass	0	n/a	121.4s	docker sha256: c5587afc e772...
3aa7a74b	cli/codex	9	submitted_passed	yes	Passed hidden tests	pass	0	n/a	114.7s	docker sha256: c5587afc e772...
cc461d13	codestral-latest	0	submitted_passed	yes	Passed hidden tests	pass	0	64 558	53.5s	docker sha256: c5587afc e772...
26427e2b	codestral-latest	1	submitted_failed	yes	Graded	fail	0	246 174	233.0s	docker sha256: c5587afc e772...
b4a95e09	codestral-latest	2	model_api_failure	no	Not graded — lane error	not graded	0	48 578	126.2s	docker sha256: c5587afc e772...
4df83755	codestral-latest	3	submitted_failed	yes	Graded	fail	0	31 078	31.6s	docker sha256: c5587afc e772...
3d19affd	codestral-latest	4	submitted_failed	yes	Graded	fail	0	110 075	77.3s	docker sha256: c5587afc e772...
7ee376ca	codestral-latest	5	submitted_failed	yes	Graded	fail	0	72 643	76.5s	docker sha256: c5587afc e772...
5ca4b95a	codestral-latest	6	submitted_failed	yes	Graded	fail	0	72 643	59.4s	docker sha256: c5587afc e772...
fe44cb82	codestral-latest	7	submitted_failed	yes	Graded	fail	0	96 982	67.5s	docker sha256: c5587afc e772...
86e8cb58	codestral-latest	8	submitted_failed	yes	Graded	fail	0	90 994	96.3s	docker sha256: c5587afc e772...
95636cf6	codestral-latest	9	submitted_failed	yes	Graded	fail	0	72 643	64.3s	docker sha256: c5587afc e772...

Failure taxonomy

CLASS	COUNT	WHAT IT MEANS
submitted_failed	8	Submitted a final application state that the independent grader rejected.
model_api_failure	1	The model endpoint failed to answer. Not a verdict about the model's ability.

Appendix B — evidence and artifact digests

RUN BUNDLE DIGEST

6e618c0542eb41677ef25561

26365aea5538ce4f4a959245

e1e09e1732ae8620

RUN ID

live-

20260902T144956Z

STARTED

2026-09-

02T14:49:56.156491

+00:00

FINISHED

2026-09-

02T15:44:48.577359

+00:00

Artifact digests

ARTIFACT	SHA-256
publicSummary	b2d0aedc0fb60c85fd4d643f4400fcfa02fd7774b09ccf4260be46c53ad93a56
rollouts/0d1dabed-f1a4-464d-8573-6298c4456eba.json	fb3d28ab7a5e999a43d50283dac14d6d6a892b208afcacf8d50718d410c1ca98
rollouts/26427e2b-2379-4a9f-891d-9fac9ff5783d.json	e480c99256a5a11253e10cd41035917212805b1af6cdf6aa7d78a61a5c9d1599
rollouts/38d7cace-518a-47c2-9678-318d7791b65f.json	a1da0a8316e97be2789ae1391260ac73ba67d97e7a4ebceb5c81cc37dded1469
rollouts/3aa7a74b-9c98-4c6e-814d-f759615903f5.json	559cc093d11a57d222156438f4e07f0e79914cece32ac6e3869394ecb8070bd
rollouts/3d19affd-3e09-4a9b-be1d-1bc702a87a51.json	3279d55e621bda9ecad8bcab0d2cb6584e833f1e1ae5392c25da0704f6349e94
rollouts/4df83755-7cf3-4d09-939f-44245af49dd0.json	e265d305de0acd9a29b6e9599c44785b51f5449dc375b9fb9ebae70773651b41
rollouts/532e1db3-0b8f-4765-808d-9341d00aee8f.json	53564320a7e4555050a2993db03fe77e59b6eb970fa3d7a94682a139b67e2e8b
rollouts/5ca4b95a-34c3-47a6-92aa-64987001944c.json	818ad5941d717436ba85cfc0f5d2729f61db4d525a5e0238cdc1ff7ab372ed98
rollouts/6015f268-cbda-4c89-a337-78b12414d270.json	eb1bea2ad651ae35bef64554e00b24a2eb4fa9a3f7ab3aaa6576036cc93d95a0
rollouts/6bd00d3c-1064-4c32-8a9c-121373dcfe98.json	ac2013c954cd90cc51718fabdc56d1a082a15a764d8c6f27ccd36cea99978256
rollouts/7ee376ca-7679-44f3-b428-5e3ad0d7e6f0.json	491b3dce54e59a85d2218733d08f9cb04d7afe7b3154473faae9b1cbbaa2cd6a
rollouts/7f40f719-4f86-464e-a1f9-cadfeeffa559.json	9e66bccf1b28e14a3621c81a47a593078f99a13e08b164b01139b66baa1deb74
rollouts/86e8cb58-b83d-4b8d-83f6-0003ac366fda.json	488baffcef7b005cb6fa47c621aa9fb6d26ecdac29560bb48cfe7f751ddd8f41
rollouts/95636cf6-da8f-48bd-9730-35964ac99fdf.json	92e873fb223f3f6b29ad197bc53140a9135955ec5a51b180ba4d791e1a60ec56
rollouts/9b8db9ca-f53f-4811-bfa9-1a2c618d8fd6.json	4593c01d72ce73d83cd476f4e3f89a964a0adcc118b699d1cfbeef6d24b1df22
rollouts/a116d4fa-48b8-47d4-83d9-0d300be62a8b.json	f19a9026c6aab52630d8f5a80ed314f78480ef2602410aa91bd5f3597ed6a5d1
rollouts/b3f86ec4-60d4-40f1-9482-b51bb3f0b2b7.json	957aa6a8fe5ef17f00a06fffd0c5fa5c6c2b038f1b632a6cf74efb99baa4ee31b
rollouts/b4a95e09-f935-48d0-b79f-69c45c5dc341.json	4e7206169f78c5a8b99e8fd3c674d0a05455d9857d3ad9d5908f701519c7dc24
rollouts/b4f9ae08-95c3-41f4-848c-ed0edd82b145.json	75edd1a44f3b8b5ed6a07f37096a4ee55ca7a7b2fcd636c6ae79344674401451
rollouts/bbef3c65-99be-4a8a-83db-ea794108c317.json	a7a9815bf129d98da564ff9bf2a93634e959eea510cdb9cc1bc0a7a509df7b78
rollouts/c20efdb4-48a5-42f1-af8a-3fc6df3560ce.json	5e297e565b112b644a8f7363c6442ff7ab79f8f30893dcd351b6773b20101365
rollouts/c8c5c1f3-0c97-4022-8451-127d48f6953a.json	0956e3caf54e2596183f40906fd7a388cdab8015b12184528a71a80116121a46
rollouts/cc461d13-0b1b-40a2-8903-2d17c8ffe681.json	6eb1859ed1f57ca4cedb10414fd862ce818d2f3126c1681d4cd96a71ec4d2ea3
rollouts/d641f68a-d3a6-45d0-8228-0e8ff30ea176.json	f81f17004b13489e33073f2b92a786353dee80586510f96eb77d3efcee24cdca
rollouts/d6987896-0e83-4e30-832a-35ae28f5e0e8.json	76aab57eef5293c1270d57a9d370164ff1e5a00fbe27801edf9b700f51232fcf
rollouts/da5c8bda-9910-4ec7-893d-e4fda9fb617c.json	0da0f242816d1d1a3e254ae079a4d9a28367bf69010a80d86d284098214dfa6c
rollouts/e415a6c4-b070-41d7-be54-c68e5dcefaaa.json	4fcb78cb3f748628d95f62a10451d7ef2131dae262aaa123f12327f629b7a6f2
rollouts/e52d496d-c202-4c0f-9fbc-4c4328dcf0bc.json	8c1b0975dd0a6fbbedbac9b54643b2f841eefd62ce07c48d0ddd7841cf444db81

ARTIFACT	SHA-256
rollouts/e771f4b2-674d-46db-8f64-ac620c935e3d.json	b5142923e8c761389e644df882ec7c94f14f3cfafbc1eed05cb4d92954030f8e
rollouts/fe44cb82-d2b1-4a7e-ac94-c9c086a65d01.json	23e1bc1d203a2f9db695033b15a360baa6e8d99659bd16a264f6f884becb306d
spec	79647b29130d3013f9061b0089ff43a26c9952ef5a5f3332fee8d5d702e755ae
summary	1fc2604a382037b95e760a3aeac37e70b462eb99d33a14dda05035ebf033c2db

Reproducing this run

Every rollout can be re-derived: the exam is frozen (fingerprint above), the grader ran in the recorded container image, and each rollout's full trace — every tool call, every result, every word — is in the evidence bundle at its rollouts/<id>.json path. A re-run uses the same job shape: the exam id, your endpoint, n, and the runtime. The record id names this run forever.

Appendix C — redacted transcripts

For each rollout: the tool calls it made, in order — which tool, on which of the model's own paths, whether the call succeeded, and how many characters came back. What a tool returned is not here, and neither is the submission: on this exam the submitted content is the graded answer, so no diff, no answer value and no model claim is printed. A request for the reference solution or the hidden tests is recorded as a request; the path it asked for is replaced by a label. See Appendix D, "What this report withholds and why".

Disclosure class applied to this exam: state disclosure (ownership vvdex_proprietary — engine registry; family browser): the exam is VVDex evaluation IP and the graded output is the application's own state and audit trail, so no submission content, diff or state is published.

532e1db3 · cli:cli/claude-sonnet · seed 0 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 975
1	browser_open	—	ok	62
2	browser_snapshot	—	ok	1 298
3	browser_click	—	ok	81
4	browser_snapshot	—	ok	1 328
5	browser_click	—	ok	65
6	browser_snapshot	—	ok	605
7	browser_fill	—	ok	36
8	browser_fill	—	ok	36
9	browser_fill	—	ok	36
10	browser_click	—	ok	65
11	browser_click	—	ok	65
12	browser_snapshot	—	ok	715
13	browser_fill	—	ok	36
14	browser_click	—	ok	65
15	browser_snapshot	—	ok	775
16	browser_click	—	ok	73
17	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (46 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

d641f68a · cli:cli/claude-sonnet · seed 1 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 975
1	browser_open	—	ok	62
2	browser_snapshot	—	ok	1 298

STEP	TOOL	PATH	RESULT	CHARS
3	browser_click	—	ok	81
4	browser_snapshot	—	ok	1 328
5	browser_click	—	ok	65
6	browser_snapshot	—	ok	605
7	browser_fill	—	ok	36
8	browser_fill	—	ok	36
9	browser_fill	—	ok	36
10	browser_click	—	ok	65
11	browser_click	—	ok	65
12	browser_snapshot	—	ok	715
13	browser_fill	—	ok	36
14	browser_click	—	ok	65
15	browser_snapshot	—	ok	775
16	browser_click	—	ok	73
17	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (46 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

b4f9ae08 · cli:cli/claude-sonnet · seed 2 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 975
1	browser_open	—	ok	62
2	browser_snapshot	—	ok	1 298
3	browser_click	—	ok	81
4	browser_snapshot	—	ok	1 328
5	browser_click	—	ok	65
6	browser_snapshot	—	ok	605
7	browser_fill	—	ok	36
8	browser_fill	—	ok	36
9	browser_fill	—	ok	36
10	browser_click	—	ok	65
11	browser_click	—	ok	65
12	browser_snapshot	—	ok	715
13	browser_fill	—	ok	36
14	browser_click	—	ok	65
15	browser_snapshot	—	ok	775

STEP	TOOL	PATH	RESULT	CHARS
16	browser_click	—	ok	73
17	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (46 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

da5c8bda · cli:cli/claude-sonnet · seed 3 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 975
1	browser_open	—	ok	62
2	browser_snapshot	—	ok	1 298
3	browser_click	—	ok	81
4	browser_snapshot	—	ok	1 328
5	browser_click	—	ok	65
6	browser_snapshot	—	ok	605
7	browser_fill	—	ok	36
8	browser_fill	—	ok	36
9	browser_fill	—	ok	36
10	browser_click	—	ok	65
11	browser_click	—	ok	65
12	browser_snapshot	—	ok	715
13	browser_fill	—	ok	36
14	browser_click	—	ok	65
15	browser_snapshot	—	ok	775
16	browser_click	—	ok	73
17	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (46 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

a116d4fa · cli:cli/claude-sonnet · seed 4 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 975
1	browser_open	—	ok	62
2	browser_snapshot	—	ok	1 298
3	browser_click	—	ok	81
4	browser_snapshot	—	ok	1 328
5	browser_click	—	ok	65

STEP	TOOL	PATH	RESULT	CHARS
6	browser_snapshot	—	ok	605
7	browser_fill	—	ok	36
8	browser_fill	—	ok	36
9	browser_fill	—	ok	36
10	browser_click	—	ok	65
11	browser_click	—	ok	65
12	browser_snapshot	—	ok	715
13	browser_fill	—	ok	36
14	browser_click	—	ok	65
15	browser_snapshot	—	ok	775
16	browser_click	—	ok	73
17	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (46 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

d6987896 · cli:cli/claude-sonnet · seed 5 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 975
1	browser_open	—	ok	62
2	browser_snapshot	—	ok	1 298
3	browser_click	—	ok	81
4	browser_snapshot	—	ok	1 328
5	browser_click	—	ok	65
6	browser_snapshot	—	ok	605
7	browser_fill	—	ok	36
8	browser_fill	—	ok	36
9	browser_fill	—	ok	36
10	browser_click	—	ok	65
11	browser_click	—	ok	65
12	browser_snapshot	—	ok	715
13	browser_fill	—	ok	36
14	browser_click	—	ok	65
15	browser_snapshot	—	ok	775
16	browser_click	—	ok	73
17	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (46 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is

the graded answer, so publishing an excerpt of it would publish the answer key.

e52d496d · cli:cli/claude-sonnet · seed 6 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 975
1	browser_open	—	ok	62
2	browser_snapshot	—	ok	1 298
3	browser_click	—	ok	81
4	browser_snapshot	—	ok	1 328
5	browser_click	—	ok	65
6	browser_snapshot	—	ok	605
7	browser_fill	—	ok	36
8	browser_fill	—	ok	36
9	browser_fill	—	ok	36
10	browser_click	—	ok	65
11	browser_click	—	ok	65
12	browser_snapshot	—	ok	715
13	browser_fill	—	ok	36
14	browser_click	—	ok	65
15	browser_snapshot	—	ok	775
16	browser_click	—	ok	73
17	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (46 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

7f40f719 · cli:cli/claude-sonnet · seed 7 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 975
1	browser_open	—	ok	62
2	browser_snapshot	—	ok	1 298
3	browser_click	—	ok	81
4	browser_snapshot	—	ok	1 328
5	browser_click	—	ok	65
6	browser_snapshot	—	ok	605
7	browser_fill	—	ok	36
8	browser_fill	—	ok	36
9	browser_fill	—	ok	36

STEP	TOOL	PATH	RESULT	CHARS
10	browser_click	—	ok	65
11	browser_click	—	ok	65
12	browser_snapshot	—	ok	715
13	browser_fill	—	ok	36
14	browser_click	—	ok	65
15	browser_snapshot	—	ok	775
16	browser_click	—	ok	73
17	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (46 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

9b8db9ca · cli:cli/claude-sonnet · seed 8 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 975
1	browser_open	—	ok	62
2	browser_snapshot	—	ok	1 298
3	browser_click	—	ok	81
4	browser_snapshot	—	ok	1 328
5	browser_click	—	ok	65
6	browser_snapshot	—	ok	605
7	browser_fill	—	ok	36
8	browser_fill	—	ok	36
9	browser_fill	—	ok	36
10	browser_click	—	ok	65
11	browser_click	—	ok	65
12	browser_snapshot	—	ok	715
13	browser_fill	—	ok	36
14	browser_click	—	ok	65
15	browser_snapshot	—	ok	775
16	browser_click	—	ok	73
17	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (46 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

38d7cace · cli:cli/claude-sonnet · seed 9 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 975
1	browser_open	—	ok	62
2	browser_snapshot	—	ok	1 298
3	browser_click	—	ok	81
4	browser_snapshot	—	ok	1 328
5	browser_click	—	ok	65
6	browser_snapshot	—	ok	605
7	browser_fill	—	ok	36
8	browser_fill	—	ok	36
9	browser_fill	—	ok	36
10	browser_click	—	ok	65
11	browser_click	—	ok	65
12	browser_snapshot	—	ok	715
13	browser_fill	—	ok	36
14	browser_click	—	ok	65
15	browser_snapshot	—	ok	775
16	browser_click	—	ok	73
17	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (46 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

6bd00d3c · cli:cli/codex · seed 0 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 975
1	browser_open	—	ok	62
2	browser_snapshot	—	ok	1 298
3	browser_click	—	ok	81
4	browser_snapshot	—	ok	1 328
5	browser_click	—	ok	65
6	browser_snapshot	—	ok	605
7	browser_fill	—	ok	36
8	browser_fill	—	ok	36
9	browser_fill	—	ok	36
10	browser_fill	—	ok	36
11	browser_click	—	ok	65

STEP	TOOL	PATH	RESULT	CHARS
12	browser_click	—	ok	65
13	browser_snapshot	—	ok	715
14	browser_fill	—	ok	36
15	browser_click	—	ok	65
16	browser_snapshot	—	ok	775
17	browser_click	—	ok	73
18	browser_snapshot	—	ok	635
19	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (48 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

b3f86ec4 · cli:cli/codex · seed 1 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 975
1	browser_open	—	ok	62
2	browser_snapshot	—	ok	1 298
3	browser_fill	—	ok	36
4	browser_click	—	ok	87
5	browser_snapshot	—	ok	621
6	browser_click	—	ok	65
7	browser_snapshot	—	ok	605
8	browser_fill	—	ok	36
9	browser_fill	—	ok	36
10	browser_fill	—	ok	36
11	browser_click	—	ok	65
12	browser_click	—	ok	65
13	browser_snapshot	—	ok	715
14	browser_fill	—	ok	36
15	browser_click	—	ok	65
16	browser_snapshot	—	ok	775
17	browser_click	—	ok	73
18	browser_snapshot	—	ok	635
19	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (48 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

Od1dabed · cli:cli/codex · seed 2 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 975
1	browser_open	—	ok	62
2	browser_snapshot	—	ok	1 298
3	browser_fill	—	ok	36
4	browser_click	—	ok	87
5	browser_snapshot	—	ok	621
6	browser_click	—	ok	65
7	browser_snapshot	—	ok	605
8	browser_fill	—	ok	36
9	browser_fill	—	ok	36
10	browser_fill	—	ok	36
11	browser_click	—	ok	65
12	browser_click	—	ok	65
13	browser_snapshot	—	ok	715
14	browser_fill	—	ok	36
15	browser_click	—	ok	65
16	browser_snapshot	—	ok	775
17	browser_click	—	ok	73
18	browser_snapshot	—	ok	635
19	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (48 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

e415a6c4 · cli:cli/codex · seed 3 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 975
1	browser_open	—	ok	62
2	browser_snapshot	—	ok	1 298
3	browser_fill	—	ok	36
4	browser_click	—	ok	87
5	browser_snapshot	—	ok	621
6	browser_click	—	ok	65
7	browser_snapshot	—	ok	605
8	browser_fill	—	ok	36
9	browser_fill	—	ok	36

STEP	TOOL	PATH	RESULT	CHARS
10	browser_fill	—	ok	36
11	browser_click	—	ok	65
12	browser_snapshot	—	ok	614
13	browser_click	—	ok	65
14	browser_snapshot	—	ok	715
15	browser_fill	—	ok	36
16	browser_click	—	ok	65
17	browser_snapshot	—	ok	775
18	browser_click	—	ok	73
19	browser_snapshot	—	ok	635
20	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (49 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

e771f4b2 · cli:cli/codex · seed 4 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 975
1	browser_open	—	ok	62
2	browser_snapshot	—	ok	1 298
3	browser_fill	—	ok	36
4	browser_click	—	ok	87
5	browser_snapshot	—	ok	621
6	browser_click	—	ok	65
7	browser_snapshot	—	ok	605
8	browser_fill	—	ok	36
9	browser_fill	—	ok	36
10	browser_fill	—	ok	36
11	browser_click	—	ok	65
12	browser_snapshot	—	ok	614
13	browser_click	—	ok	65
14	browser_snapshot	—	ok	715
15	browser_fill	—	ok	36
16	browser_click	—	ok	65
17	browser_snapshot	—	ok	775
18	browser_click	—	ok	73
19	browser_snapshot	—	ok	635

STEP	TOOL	PATH	RESULT	CHARS
20	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (49 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

c8c5c1f3 · cli:cli/codex · seed 5 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 975
1	browser_open	—	ok	62
2	browser_snapshot	—	ok	1 298
3	browser_fill	—	ok	36
4	browser_click	—	ok	87
5	browser_snapshot	—	ok	621
6	browser_click	—	ok	65
7	browser_snapshot	—	ok	605
8	browser_fill	—	ok	36
9	browser_fill	—	ok	36
10	browser_fill	—	ok	36
11	browser_click	—	ok	65
12	browser_click	—	ok	65
13	browser_snapshot	—	ok	715
14	browser_fill	—	ok	36
15	browser_click	—	ok	65
16	browser_snapshot	—	ok	775
17	browser_click	—	ok	73
18	browser_snapshot	—	ok	635
19	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (48 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

c20efdb4 · cli:cli/codex · seed 6 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 975
1	browser_open	—	ok	62
2	browser_snapshot	—	ok	1 298
3	browser_fill	—	ok	36
4	browser_click	—	ok	87

STEP	TOOL	PATH	RESULT	CHARS
5	browser_snapshot	—	ok	621
6	browser_click	—	ok	65
7	browser_snapshot	—	ok	605
8	browser_fill	—	ok	36
9	browser_fill	—	ok	36
10	browser_fill	—	ok	36
11	browser_click	—	ok	65
12	browser_click	—	ok	65
13	browser_snapshot	—	ok	715
14	browser_fill	—	ok	36
15	browser_click	—	ok	65
16	browser_snapshot	—	ok	775
17	browser_click	—	ok	73
18	browser_snapshot	—	ok	635
19	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (48 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

6015f268 · cli:cli/codex · seed 7 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 975
1	browser_open	—	ok	62
2	browser_snapshot	—	ok	1 298
3	browser_fill	—	ok	36
4	browser_click	—	ok	87
5	browser_snapshot	—	ok	621
6	browser_click	—	ok	65
7	browser_snapshot	—	ok	605
8	browser_fill	—	ok	36
9	browser_fill	—	ok	36
10	browser_fill	—	ok	36
11	browser_click	—	ok	65
12	browser_snapshot	—	ok	614
13	browser_click	—	ok	65
14	browser_snapshot	—	ok	715
15	browser_fill	—	ok	36

STEP	TOOL	PATH	RESULT	CHARS
16	browser_snapshot	—	ok	712
17	browser_click	—	ok	65
18	browser_snapshot	—	ok	775
19	browser_click	—	ok	73
20	browser_snapshot	—	ok	635
21	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (50 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

bbef3c65 · cli:cli/codex · seed 8 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 975
1	browser_open	—	ok	62
2	browser_snapshot	—	ok	1 298
3	browser_fill	—	ok	36
4	browser_click	—	ok	87
5	browser_snapshot	—	ok	621
6	browser_click	—	ok	65
7	browser_snapshot	—	ok	605
8	browser_fill	—	ok	36
9	browser_fill	—	ok	36
10	browser_fill	—	ok	36
11	browser_click	—	ok	65
12	browser_click	—	ok	65
13	browser_snapshot	—	ok	715
14	browser_fill	—	ok	36
15	browser_click	—	ok	65
16	browser_snapshot	—	ok	775
17	browser_click	—	ok	73
18	browser_snapshot	—	ok	635
19	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (48 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

3aa7a74b · cli:cli/codex · seed 9 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 975
1	browser_open	—	ok	62
2	browser_snapshot	—	ok	1 298
3	browser_fill	—	ok	36
4	browser_click	—	ok	87
5	browser_snapshot	—	ok	621
6	browser_click	—	ok	65
7	browser_snapshot	—	ok	605
8	browser_fill	—	ok	36
9	browser_fill	—	ok	36
10	browser_fill	—	ok	36
11	browser_click	—	ok	65
12	browser_click	—	ok	65
13	browser_snapshot	—	ok	715
14	browser_fill	—	ok	36
15	browser_click	—	ok	65
16	browser_snapshot	—	ok	775
17	browser_click	—	ok	73
18	browser_snapshot	—	ok	635
19	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (48 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

cc461d13 · mistral:codestral-latest · seed 0 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	browser_open	—	ok	62
1	browser_snapshot	—	ok	1 298
2	browser_click	—	ok	65
3	browser_snapshot	—	ok	594
4	browser_back	—	ok	61
5	browser_snapshot	—	ok	1 298
6	browser_click	—	ok	81
7	browser_snapshot	—	ok	1 328
8	browser_click	—	ok	65
9	browser_snapshot	—	ok	605

STEP	TOOL	PATH	RESULT	CHARS
10	browser_fill	—	ok	36
11	browser_fill	—	ok	36
12	browser_fill	—	ok	36
13	browser_click	—	ok	65
14	browser_click	—	ok	65
15	browser_snapshot	—	ok	715
16	browser_fill	—	ok	36
17	browser_click	—	ok	65
18	browser_snapshot	—	ok	775
19	browser_click	—	ok	73
20	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (50 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

26427e2b · mistral:codestral-latest · seed 1 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	browser_open	—	ok	62
1	browser_snapshot	—	ok	1 298
2	browser_click	—	ok	65
3	browser_snapshot	—	ok	594
4	browser_back	—	ok	61
5	browser_snapshot	—	ok	1 298
6	browser_click	—	ok	65
7	browser_snapshot	—	ok	600
8	browser_back	—	ok	61
9	browser_snapshot	—	ok	1 298
10	browser_click	—	ok	65
11	browser_snapshot	—	ok	599
12	browser_back	—	ok	61
13	browser_snapshot	—	ok	1 298
14	browser_click	—	ok	65
15	browser_snapshot	—	ok	599
16	browser_fill	—	ok	36
17	browser_fill	—	ok	36
18	browser_fill	—	ok	36
19	browser_fill	—	ok	36

STEP	TOOL	PATH	RESULT	CHARS
20	browser_click	—	ok	65
21	browser_snapshot	—	ok	602
22	browser_click	—	ok	65
23	browser_snapshot	—	ok	703
24	browser_fill	—	ok	36
25	browser_click	—	ok	65
26	browser_snapshot	—	ok	781
27	browser_click	—	ok	73
28	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (59 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

b4a95e09 · mistral:codestral-latest · seed 2 · model_api_failure

STEP	TOOL	PATH	RESULT	CHARS
0	browser_open	—	ok	62
1	browser_snapshot	—	ok	1 298
2	browser_click	—	ok	65
3	browser_snapshot	—	ok	594
4	browser_back	—	ok	61
5	browser_snapshot	—	ok	1 298
6	browser_click	—	ok	81
7	browser_snapshot	—	ok	1 328
8	browser_click	—	ok	65
9	browser_snapshot	—	ok	605
10	browser_fill	—	ok	36
11	browser_fill	—	ok	36
12	browser_click	—	ok	65
13	browser_snapshot	—	ok	616
14	browser_click	—	ok	65
15	browser_snapshot	—	ok	717
16	browser_fill	—	ok	36

Submission diff withheld by the public-disclosure policy (24 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

4df83755 · mistral:codestral-latest · seed 3 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	browser_open	—	ok	62
1	browser_snapshot	—	ok	1 298
2	browser_click	—	ok	65
3	browser_snapshot	—	ok	594
4	browser_fill	—	ok	36
5	browser_fill	—	ok	36
6	browser_fill	—	ok	36
7	browser_click	—	ok	65
8	browser_snapshot	—	ok	706
9	browser_fill	—	ok	36
10	browser_click	—	ok	65
11	browser_snapshot	—	ok	774
12	browser_click	—	ok	73
13	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (38 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

3d19affd · mistral:codestral-latest · seed 4 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	browser_open	—	ok	62
1	browser_snapshot	—	ok	1 298
2	browser_click	—	ok	65
3	browser_snapshot	—	ok	594
4	browser_back	—	ok	61
5	browser_snapshot	—	ok	1 298
6	browser_click	—	ok	65
7	browser_snapshot	—	ok	600
8	browser_back	—	ok	61
9	browser_snapshot	—	ok	1 298
10	browser_back	—	ok	44
11	browser_open	—	ok	62
12	browser_snapshot	—	ok	1 298
13	browser_click	—	ok	81
14	browser_snapshot	—	ok	1 328
15	browser_click	—	ok	65

STEP	TOOL	PATH	RESULT	CHARS
16	browser_snapshot	—	ok	605
17	browser_fill	—	ok	36
18	browser_fill	—	ok	36
19	browser_click	—	ok	65
20	browser_snapshot	—	ok	616
21	browser_click	—	ok	65
22	browser_snapshot	—	ok	717
23	browser_fill	—	ok	36
24	browser_click	—	ok	65
25	browser_snapshot	—	ok	777
26	browser_click	—	ok	73
27	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (54 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

7ee376ca · mistral:codestral-latest · seed 5 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	browser_open	—	ok	62
1	browser_snapshot	—	ok	1 298
2	browser_click	—	ok	65
3	browser_snapshot	—	ok	594
4	browser_back	—	ok	61
5	browser_snapshot	—	ok	1 298
6	browser_click	—	ok	81
7	browser_snapshot	—	ok	1 328
8	browser_click	—	ok	65
9	browser_snapshot	—	ok	605
10	browser_fill	—	ok	36
11	browser_fill	—	ok	36
12	browser_click	—	ok	65
13	browser_snapshot	—	ok	616
14	browser_click	—	ok	65
15	browser_snapshot	—	ok	717
16	browser_fill	—	ok	36
17	browser_click	—	ok	65
18	browser_snapshot	—	ok	777

STEP	TOOL	PATH	RESULT	CHARS
19	browser_click	—	ok	73
20	browser_snapshot	—	ok	637
21	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (48 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

5ca4b95a · mistral:codestral-latest · seed 6 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	browser_open	—	ok	62
1	browser_snapshot	—	ok	1 298
2	browser_click	—	ok	65
3	browser_snapshot	—	ok	594
4	browser_back	—	ok	61
5	browser_snapshot	—	ok	1 298
6	browser_click	—	ok	81
7	browser_snapshot	—	ok	1 328
8	browser_click	—	ok	65
9	browser_snapshot	—	ok	605
10	browser_fill	—	ok	36
11	browser_fill	—	ok	36
12	browser_click	—	ok	65
13	browser_snapshot	—	ok	616
14	browser_click	—	ok	65
15	browser_snapshot	—	ok	717
16	browser_fill	—	ok	36
17	browser_click	—	ok	65
18	browser_snapshot	—	ok	777
19	browser_click	—	ok	73
20	browser_snapshot	—	ok	637
21	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (48 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

fe44cb82 · mistral:codestral-latest · seed 7 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	browser_open	—	ok	62

STEP	TOOL	PATH	RESULT	CHARS
1	browser_snapshot	—	ok	1 298
2	browser_click	—	ok	65
3	browser_snapshot	—	ok	594
4	browser_back	—	ok	61
5	browser_snapshot	—	ok	1 298
6	browser_click	—	ok	65
7	browser_snapshot	—	ok	600
8	browser_back	—	ok	61
9	browser_snapshot	—	ok	1 298
10	browser_back	—	ok	44
11	browser_open	—	ok	62
12	browser_snapshot	—	ok	1 298
13	browser_click	—	ok	81
14	browser_snapshot	—	ok	1 328
15	browser_click	—	ok	65
16	browser_snapshot	—	ok	605
17	browser_fill	—	ok	36
18	browser_fill	—	ok	36
19	browser_click	—	ok	65
20	browser_snapshot	—	ok	616
21	browser_click	—	ok	65
22	browser_snapshot	—	ok	717
23	browser_fill	—	ok	36
24	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (31 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

86e8cb58 · mistral:codestral-latest · seed 8 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	browser_open	—	ok	62
1	browser_snapshot	—	ok	1 298
2	browser_click	—	ok	65
3	browser_snapshot	—	ok	594
4	browser_back	—	ok	61
5	browser_snapshot	—	ok	1 298
6	browser_click	—	ok	81

STEP	TOOL	PATH	RESULT	CHARS
7	browser_snapshot	—	ok	1 328
8	browser_click	—	ok	65
9	browser_snapshot	—	ok	605
10	browser_fill	—	ok	36
11	browser_fill	—	ok	36
12	browser_click	—	ok	65
13	browser_snapshot	—	ok	616
14	browser_click	—	ok	65
15	browser_snapshot	—	ok	717
16	browser_fill	—	ok	36
17	browser_click	—	ok	65
18	browser_snapshot	—	ok	777
19	browser_click	—	ok	73
20	browser_snapshot	—	ok	637
21	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (48 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

95636cf6 · mistral:codestral-latest · seed 9 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	browser_open	—	ok	62
1	browser_snapshot	—	ok	1 298
2	browser_click	—	ok	65
3	browser_snapshot	—	ok	594
4	browser_back	—	ok	61
5	browser_snapshot	—	ok	1 298
6	browser_click	—	ok	81
7	browser_snapshot	—	ok	1 328
8	browser_click	—	ok	65
9	browser_snapshot	—	ok	605
10	browser_fill	—	ok	36
11	browser_fill	—	ok	36
12	browser_click	—	ok	65
13	browser_snapshot	—	ok	616
14	browser_click	—	ok	65
15	browser_snapshot	—	ok	717

STEP	TOOL	PATH	RESULT	CHARS
16	browser_fill	—	ok	36
17	browser_click	—	ok	65
18	browser_snapshot	—	ok	777
19	browser_click	—	ok	73
20	browser_snapshot	—	ok	637
21	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (48 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

11 Exam · fault-recovery-1

vvdex.harness.fault-recovery-1 · record fr-20260902-a6f12b64 · record digest d5a8264aa0c20dd6... · harness clean · containment clean

Executive summary

MODELS TESTED

3

CERTIFIED PASSES

20

of 30 graded rollouts

SUBMITTED

21 of 30

graded rollouts that called submit

NOT GRADED — LANE ERRORS

0

of 30 attempts

Purpose

This report presents a controlled evaluation of 3 model lane(s) against the certified exam vvdex.harness.fault-recovery-1 (v0.1.0, family harness_eval). Every lane received the identical frozen workspace, tool contract and limits; grading ran afterwards, in isolation, against hidden tests no lane ever saw. The purpose is to state, with reproducible evidence, what each lane did and where it stopped.

Outcome

20 of 30

graded rollouts passed the independent hidden tests and were certified

20 of 30 graded rollouts passed. cli/claude-sonnet, cli/codex completed the full loop (inspect, edit, run the visible tests, submit) and the sealed grader accepted the submission. The most common way to fail was context_exhaustion — ran out of output context before it could submit.

Key findings

- F-01 cli/claude-sonnet, cli/codex passed the independent hidden tests
- F-02 context_exhaustion × 5 (codestral-latest)
- F-03 edited_tested_no_submit × 4 (codestral-latest)
- F-04 submitted_failed × 1 (codestral-latest)
- F-05 1 overclaim(s) — success stated, then refuted by the grader

Each finding is stated in full, with evidence, in the Findings section; recommendations for acting on them follow it. Elapsed wall clock for the whole engagement: 3724.5s.

Exam definition

In scope. The ability of each configured model lane to complete this one certified task end to end — inspect the workspace, produce a correct edit, verify it with the visible tests, and submit — within 24 tool steps and 5880s of wall clock, at 10 rollout(s) per lane. Behavioural measurements along the way (tool discipline, grounding, overclaim) are in scope because they are read from the recorded trace, not judged by another model.

Out of scope. General model quality, performance on other task families, prompt-tuning on the lane's behalf, and any comparison this run's sample size cannot support. A lane's API failure is reported as infrastructure, never as model capability.

The instrument. Exam `vvdex.harness.fault-recovery-1` was certified before this run: its reference solution passes, an empty submission fails, its grader distinguishes near-misses, and its sealed evidence is unchanged by this evaluation (see Certification evidence).

What was measured

EXAM	VERSION	FAMILY	ROLLOUTS
vvdex.harness.fault-recovery-1	0.1.0	harness_eval	30 (10 per model)

Where the defect came from

This exam has no upstream repository to credit: the defect, the reference solution and the hidden suite are VVDex's own.

Certification evidence

4 of 5 certification pillars are recorded in this package, loaded from the receipts sealed into the exam and matching this run's exam fingerprint. What is recorded is shown with its real value; what is not is named below and is not claimed.

FROZEN BASELINE

FAIL

expected — an empty submission over 2 run(s) must not pass

→

GOLD / REFERENCE

PASS

expected — the reference solution over 2 run(s) must pass

GRADER CONTROLS

12 / 12

The grader accepted the reference and rejected every near-miss.

ATTACK PROBES

10 / 10 blocked

0 trivial exploit(s) found.

RUNTIME

sha256:679b36225a1f83238efba8c71b9cde3c24caaeacc9b682ae92819cb55eec328f

network policy: deny

Certification metadata is not included in this report package for:

- sealed evidence bundle

Benchmark results remain valid as run evidence; the pillar(s) above are not substantiated by this document and this report does not claim them.

What the package does carry: the exam fingerprint below is the contract digest this run was executed against — a holder of the certification bundle for this fingerprint can join the two independently.

EXAM FINGERPRINT

4a206fbfe8fb70fb3e37943705065a8599fab492e495dca326279781ee960b1e

The contract digest this run was executed against.

Evaluation configuration

REPORT	fr-20260902-a6f12b64
ISSUED	2026-09-02
SUBJECT EXAM	vvdex.harness.fault-recovery-1 · v0.1.0
FAMILY	harness_eval
PREPARED BY	VVDex Forge engine vvdex-env 0.1.0
CLASSIFICATION	Independent model evaluation report — independently verifiable
CERTIFICATION	Exam certification: partially recorded (see Certification evidence)
STATUS	FINAL · report sealed against its own digest

Timeline

WHEN (UTC)	EVENT
2026-09-02 13:47:41	Evaluation run started
2026-09-02 14:49:46	Evaluation run finished
2026-09-02 14:49:46	Report generated from the sealed run evidence

Models under test

MODEL	PROVIDER	ENDPOINT	BILLING
Claude Code CLI (Sonnet, subscription)	cli	local runtime	operator subscription, flat — not billed per token
Codex CLI (GPT, subscription)	cli	local runtime	operator subscription, flat — not billed per token
codestral-latest	mistral	VVDex-configured	VVDex free-quota

A customer endpoint is recorded by origin only — never its port, its path, or its key.

Limits

RUNTIME docker	STEP LIMIT 24	WALL-CLOCK LIMIT 5880s	MAX OUTPUT TOKENS 2048
TEMPERATURE 0.2	RETRY POLICY none	COMMAND BUDGET 0	TOOLS OFFERED list_files, read_file, write_file, edit_file, run_tests, submit

The evaluated model never saw the hidden tests or the reference solution, and could not change its result. Grading ran in a separate step, after submission, on a container with no network.

Model outcomes

Denominators. Attempts requested: 30. Graded (evaluable) rollouts: 30. Not graded — lane errors: 0. A lane error is a rollout the API or the runtime ended before the hidden grader ran: it is listed, excluded from every rate and pass@k, and never silently dropped. Passes are certified passes: hidden grader exit 0 with integrity verified.

MODEL	ATTEMPTS	GRADED	PASSED	MODEL FAILS	LANE ERRORS	SUBMITTED	MEAN TIME	TYPICAL DEEPEST STAGE
Claude Code CLI (Sonnet, subscription) cli/claude-sonnet	10	10	10 / 10	0	0	10 / 10	77.9s	Passed hidden tests
Codex CLI (GPT, subscription) cli/codex	10	10	10 / 10	0	0	10 / 10	147.4s	Passed hidden tests
codestral-latest codestral-latest	10	10	0 / 10	10	0	1 / 10	127.0s	Graded

Every rollout, with its own deepest stage and grade, is in Appendix A.

Success rate, confidence and pass@k

Every rate on this page is a measurement of a finite number of rollouts, so every rate carries the interval that measurement actually supports. n is the evaluable count — attempts minus rollouts the lane ended before grading; those are listed under "not graded" and sit in no rate, interval or pass@k.

MODEL	ATTEMPTS	N	NOT GRADED	PASSES	RATE	95% INTERVAL	SEEDS	PASS@1	PASS@3	PASS@5	PASS@10
cli/claude-sonnet	10	10	0	10	100.0%	[72%, 100%]	0–9	100.0%	100.0%	100.0%	100.0%
cli/codex	10	10	0	10	100.0%	[72%, 100%]	0–9	100.0%	100.0%	100.0%	100.0%
codestral-latest	10	10	0	0	0.0%	[0%, 28%]	0–9	0.0%	0.0%	0.0%	0.0%

Interval and pass@k methods are stated in full in Appendix D. pass@k is shown for k = 1, 3, 5, 10; the sealed record carries every k up to n.

Model comparison

MODEL	PASS@1	95% INTERVAL	TOKENS	COST	MEAN ROLLOUT	OVERCLAIM RATE	TOP FAILURE CLASS
cli/claude-sonnet	100.0%	[72%, 100%]	n/a	n/a	77.9s	0.0%	none
cli/codex	100.0%	[72%, 100%]	n/a	n/a	147.4s	0.0%	none
codestral-latest	0.0%	[0%, 28%]	1 147 890	not billed	127.0s	10.0%	context_exhaustion

Tokens and cost read n/a where a lane reports no telemetry (a flat-rate CLI) — an absent measurement, not zero. pass@1 and its interval are over evaluable rollouts only.

This is not a leaderboard. Separated by this run (intervals do not overlap): cli/claude-sonnet vs codestral-latest; cli/codex vs codestral-latest. Not separated (intervals overlap): cli/claude-sonnet vs cli/codex. The primary comparison in this report is the stage funnel below — where a model stops is a finding; a percentage over this many rollouts is not.

Stage funnel

The funnel is the primary reading of a run. Where a model stops says more than a single pass percentage does, especially at low N.

STAGE	CLI:CLI/CLAUDE-SONNET	CLI:CLI/CODEX	MISTRAL:CODESTRAL-LATEST
Inspected	10 / 10	10 / 10	10 / 10
Edited	10 / 10	10 / 10	9 / 10
Ran tests	10 / 10	10 / 10	10 / 10
Submitted	10 / 10	10 / 10	1 / 10
Graded	10 / 10	10 / 10	10 / 10
Passed hidden tests	10 / 10	10 / 10	0 / 10

Deepest stage reached

MODEL	DEEPEST (ANY ROLLOUT)	TYPICAL (MOST ROLLOUTS)	HIDDEN-TEST PASSES	NOT GRADED
cli:cli/claude-sonnet	Passed hidden tests	Passed hidden tests	10 / 10	0
cli:cli/codex	Passed hidden tests	Passed hidden tests	10 / 10	0
mistral:codestral-latest	Graded	Graded	0 / 10	0

Stage definitions are in Appendix D. The matrix counts evaluable rollouts; lane errors are shown beside it, not inside it. The stages are not strictly nested: a model can submit without editing, and the matrix shows that rather than hiding it. Each rollout's own deepest stage is in Appendix A.

Behavioural analysis

Long-horizon behaviour

Counted off the recorded trace of each rollout. A dead end is a step that moved nothing: a re-read of a slice already returned, a repeat of the previous action, or a call to a tool this exam does not have.

MODEL	N	STEPS USED	RAN OUT OF STEPS	COMMANDS	CHANGED THE TREE	REFUSED (OVER BUDGET)	REVISITS	DEAD ENDS / ROLLOUT
cli:cli/claude-sonnet	10	10 / 24	0	0 / 0	0	0	17	2
cli:cli/codex	10	16 / 24	0	0 / 0	0	0	27	3
mistral:codestral-latest	10	21 / 24	9	0 / 0	0	0	7	6

Each column is defined in Appendix D.

Hallucination & overclaim

Three measurements, each with a stated definition. None of them is a judgement of the model's prose by another model — every number below comes from comparing what the model wrote against what was on disk, or against what the independent grader returned.

MODEL	N	GROUNDING FAILURES	KINDS	ROLLOUTS AFFECTED	OVERCLAIMS	OVERCLAIM RATE	RE-READS
cli:cli/claude-sonnet	10	0	none	0 / 10	0 / 10	0.0%	17
cli:cli/codex	10	0	none	0 / 10	0 / 10	0.0%	27
mistral:codestral-latest	10	4	nonexistent_path x4	4 / 10	1 / 10	10.0%	7

mistral:codestral-latest wrote 4 path reference(s) not in the tree it was given, across 4 of 10 rollout(s); claimed a fix the hidden tests then refuted in 1 of 10 rollout(s).

Every term above is defined in Appendix D.

Efficiency & telemetry

This run has no per-token price attached — a local model, a free quota, or a customer endpoint we are not billed for — so spend is reported in tokens. Tokens are what was actually measured; a dollar figure here would be invented. Lanes run through a flat-rate local CLI report no token telemetry; their cells say n/a — an absent measurement, not zero. Rollout counts here are evaluable rollouts.

MODEL	ROLLOUTS	TOKENS	TOKENS / ROLLOUT	COST	PASSES	RATE
cli:cli/claude-sonnet	10	n/a	n/a	n/a – telemetry unavailable	10 / 10	100.0%
cli:cli/codex	10	n/a	n/a	n/a – telemetry unavailable	10 / 10	100.0%
mistral:codestral-latest	10	1 147 890	114 789	not billed	0 / 10	0.0%

Findings

Every finding below is derived from the recorded data of this run — none is an opinion.
Severity: PASS a lane succeeded · ATTENTION a capability gap · LANE / INFRA
infrastructure, not capability · NOTE a property of the measurement.

F-01	PASS
cli/claude-sonnet, cli/codex passed the independent hidden tests	
Severity: PASS	Evidence: Model outcomes · Stage funnel · Appendix A
20 of 30 graded rollouts completed the full loop (inspect, edit, run the visible tests, submit) and the sealed grader accepted the submission. This demonstrates the task is solvable under the published limits.	
F-01 · derived from the recorded run data	
F-02	ATTENTION
context_exhaustion × 5 (codestral-latest)	
Severity: ATTENTION	Evidence: Appendix A · Appendix C
Ran out of output context before it could submit.	
F-02 · derived from the recorded run data	
F-03	ATTENTION
edited_tested_no_submit × 4 (codestral-latest)	
Severity: ATTENTION	Evidence: Appendix A · Appendix C
Edited files and ran the visible tests but never submitted.	
F-03 · derived from the recorded run data	
F-04	ATTENTION
submitted_failed × 1 (codestral-latest)	
Severity: ATTENTION	Evidence: Appendix A · Appendix C
Submitted a change; the independent hidden tests still failed.	
F-04 · derived from the recorded run data	
F-05	ATTENTION
1 overclaim(s) — success stated, then refuted by the grader	
Severity: ATTENTION	Evidence: Behavioural analysis
mistral:codestral-latest declared the task fixed and the independent hidden tests then failed. Counted claimed-and-wrong per rollout, from the model's own words.	
F-05 · derived from the recorded run data	

4 grounding failure(s) — paths cited that do not exist

Severity: ATTENTION

Evidence: Behavioural analysis

mistral:codestral-latest reasoned about files not present in the tree they were given, detected by matching cited source paths against the agent box on disk.

F-06 · derived from the recorded run data

Recommendations

One recommendation per observed failure mode, addressed to the lane it affects. Each traces back to a finding above and to the raw trace in Appendix A.

R-01

PROBLEM

Ran out of output context before it could submit.

AFFECTED LANES

codestral-latest

WHAT WE WOULD LOOK AT FIRST

The model fills its window before converging — working-set discipline over long files is the gap.

R-02

PROBLEM

Edited files and ran the visible tests but never submitted.

AFFECTED LANES

codestral-latest

WHAT WE WOULD LOOK AT FIRST

The model does the work and never calls submit — terminal-action reliability is the gap.

R-03

PROBLEM

Submitted a change; the independent hidden tests still failed.

AFFECTED LANES

codestral-latest

WHAT WE WOULD LOOK AT FIRST

The model completes the loop and produces a plausible near-miss — the named failing tests are the exact behavioral cases distinguishing it from the reference fix.

Verification

RUN ID

live-20260902T134741Z

EXAM FINGERPRINT

4a206fbfe8fb70fb3e37943705065a8599fab492e495dca326279781ee960b1e

The contract digest this run was executed against.

Provenance of this chapter

This chapter was rendered by the campaign generator from the sealed evaluation record; it is not the standalone report reproduced byte for byte. Verify the standalone report with its own digest, verify the record with `vvdex-env records verify`, and verify this campaign document as a whole.

RELATIONSHIP

campaign-rendered chapter

Derived from the same sealed evaluation record as the standalone report; these bytes are authenticated by the campaign artifact that contains them.

SEALED RECORD DIGEST

d5a8264aa0c20dd67c14879856a744e2ff1ab79da35ea8bcea92b4ba1ea14d5f

sha256 over the record's canonical JSON — the identity every renderer works from.

SOURCE RECORD FILE

5af5f844bf99f5172f661f5b51ed39a4bedc7a35585899a62859cb85cb411461

sha256 of fr-20260902-a6f12b64.json as written beside the standalone report.

SOURCE STANDALONE REPORT

0f87ffc98b14bc815daa2e7a2a427c54fd8c876b18df91b1c4221279a601109d

sha256 of fr-20260902-a6f12b64.html — independently verifiable with its own digest.

SOURCE STANDALONE PDF

e5c2cbf4120b4e602d9978e5d9181716475a3d61523029bb3e21658ece808549

sha256 of fr-20260902-a6f12b64.pdf.

RUN EVIDENCE BUNDLE

a6f12b647a54dbf0c0ecf80acda77ee562a8ccdb3ad1f44bccffa20c0716d664

The digest the run's own bundle computed over itself.

CERTIFIED EVIDENCE SEAL

not available in this package

The exam's sealed evidence, established before this run.

Measurement history

This is the first recorded run on this exam. There is nothing to compare it against yet, and a single measurement is not a trend.

The full artifact-by-artifact digest table, and the reproduction protocol, are in Appendix B.

Appendix A — every rollout

ROLLOUT	MODEL	SEED	OUTCOME	SUBMITTED	DEEPEST STAGE	HIDDEN TESTS	FILES	TOKENS	ELAPSED
3c178cb0	cli/claude-sonnet	0	submitted_passed	yes	Passed hidden tests	pass	3	n/a	74.7s
230a8274	cli/claude-sonnet	1	submitted_passed	yes	Passed hidden tests	pass	3	n/a	73.9s
459f9de3	cli/claude-sonnet	2	submitted_passed	yes	Passed hidden tests	pass	3	n/a	64.7s
3a7da64d	cli/claude-sonnet	3	submitted_passed	yes	Passed hidden tests	pass	3	n/a	92.5s
1e7f41d8	cli/claude-sonnet	4	submitted_passed	yes	Passed hidden tests	pass	3	n/a	83.8s
9902eaa3	cli/claude-sonnet	5	submitted_passed	yes	Passed hidden tests	pass	3	n/a	77.7s
12f95490	cli/claude-sonnet	6	submitted_passed	yes	Passed hidden tests	pass	3	n/a	84.0s
affa9ea3	cli/claude-sonnet	7	submitted_passed	yes	Passed hidden tests	pass	3	n/a	76.2s
ab27ff4e	cli/claude-sonnet	8	submitted_passed	yes	Passed hidden tests	pass	3	n/a	85.5s
cd337d97	cli/claude-sonnet	9	submitted_passed	yes	Passed hidden tests	pass	3	n/a	65.9s
879eca05	cli/codex	0	submitted_passed	yes	Passed hidden tests	pass	3	n/a	161.2s
fd449898	cli/codex	1	submitted_passed	yes	Passed hidden tests	pass	3	n/a	162.3s
5d078aad	cli/codex	2	submitted_passed	yes	Passed hidden tests	pass	3	n/a	223.6s
4ac68744	cli/codex	3	submitted_passed	yes	Passed hidden tests	pass	3	n/a	144.3s
7d1b1531	cli/codex	4	submitted_passed	yes	Passed hidden tests	pass	3	n/a	114.1s
662f4a3d	cli/codex	5	submitted_passed	yes	Passed hidden tests	pass	3	n/a	135.7s
660efecd	cli/codex	6	submitted_passed	yes	Passed hidden tests	pass	3	n/a	127.3s
51c40893	cli/codex	7	submitted_passed	yes	Passed hidden tests	pass	3	n/a	126.5s
899ad596	cli/codex	8	submitted_passed	yes	Passed hidden tests	pass	3	n/a	144.5s
e90dd0c8	cli/codex	9	submitted_passed	yes	Passed hidden tests	pass	3	n/a	134.5s
0711410a	codestral-latest	0	context_exhaustion	no	Graded	fail	1	80 881	75.3s
de0f078f	codestral-latest	1	context_exhaustion	no	Graded	fail	9	203 624	151.0s
5c42bd02	codestral-latest	2	submitted_failed	yes	Graded	fail	0	4 126	7.4s
3498fcf5	codestral-latest	3	context_exhaustion	no	Graded	fail	9	175 540	100.7s

ROLLOUT	MODEL	SEED	OUTCOME	SUBMITTED	DEEPEST STAGE	HIDDEN TESTS	FILES	TOKENS	ELAPSED
5e744f53	codestral-latest	4	edited_tested_no_submi t	no	Graded	fail	11	102 805	81.2s
b761b8be	codestral-latest	5	edited_tested_no_submi t	no	Graded	fail	10	89 896	94.4s
eddbc8c8	codestral-latest	6	edited_tested_no_submi t	no	Graded	fail	1	85 524	81.4s
b0c7e737	codestral-latest	7	edited_tested_no_submi t	no	Graded	fail	1	62 082	65.4s
3061565d	codestral-latest	8	context_exhaustion	no	Graded	fail	8	251 723	513.2s
8b5af559	codestral-latest	9	context_exhaustion	no	Graded	fail	9	91 689	99.7s

Failure taxonomy

CLASS	COUNT	WHAT IT MEANS
context_exhaustion	5	Ran out of output context before it could submit.
edited_tested_no_submit	4	Edited files and ran the visible tests but never submitted.
submitted_failed	1	Submitted a change; the independent hidden tests still failed.

Appendix B — evidence and artifact digests

RUN BUNDLE DIGEST

a6f12b647a54dbf0c0ecf80a
cda77ee562a8ccdb3ad1f44b
ccffa20c0716d664

RUN ID

live-
20260902T134741Z

STARTED

2026-09-
02T13:47:41.546440
+00:00

FINISHED

2026-09-
02T14:49:46.020126
+00:00

Artifact digests

ARTIFACT	SHA-256
publicSummary	92ec8ad502cca03db70ca82d7c593206728ddfa95c2e483a198203d73b81ccb6
rollouts/0711410a-20af-4809-97d9-eb6f1cb4ed52.json	e76288222b67f8622db779c912d61b460fba164efaf041126bd6dcb14a87cdc8
rollouts/12f95490-752e-4fae-8352-8fc3f600c46c.json	05fad9811badd35c5ba8e2c8b3884d87ffe09385cc1d870be5d8725bc9af84a2
rollouts/1e7f41d8-724c-4906-94aa-77195464bafe.json	01744a44277bacc58c08bcae2b1b80959ce3bf8c47588a7c787fd83453b30354
rollouts/230a8274-69d8-4aa7-a026-3972e84dee78.json	66a3d04b57a30ded39eafba676ee1fa846909384b59ad96016e4dfebb3386cff
rollouts/3061565d-c1ba-453b-b793-1e061a2e4a62.json	1f9200514a2b2c38b9ee97118974f2d5629406e67349552f13ad03b710fc49fd
rollouts/3498fcf5-bbb6-414b-aba0-a13bad481435.json	c9e901f199c5a36bda324d35fea1f3ffe95e3dc31907a59e48c79d53deaad30e
rollouts/3a7da64d-1728-43ad-b959-185afd23de73.json	4ee16b3a4a8d0092617a1dfee773295ce2a1acb590e9b58abeeda7aacb6cbbb6
rollouts/3c178cb0-987f-499b-b4c5-8e6a51d22739.json	2b77fd47c36b9ced52164b48be6988332475b6467ab185166d9bf6a9bad8b861
rollouts/459f9de3-dd18-4b28-9e46-eefcda169291.json	81ee5801eda8feea8903bdc3e959cfdcf4e495fe3a8039620ed058080a6009d4
rollouts/4ac68744-04cf-46df-9fde-24f77013e84c.json	64555ef1cfd29429efb7ffa3c4f0f611744ea7d31e27a9dc6087941312b3eb7e
rollouts/51c40893-6126-4096-88a2-2ed8a13df341.json	39fa04de1b8f70e1f90e4d1713206db95e2d0450c682658e3369b780e5b8a1b2
rollouts/5c42bd02-e324-467a-a4a8-d8264ec66778.json	17860fd2c9c0127b8a37dbdee8d3105b4add25c756e2b45ab4d38790e1520412
rollouts/5d078aad-9bc3-42b7-97b8-a615b5f9a87f.json	8463d90244f36a14c9095d900f885147fab82e5c6bf9e8406ea984e06ad1208a
rollouts/5e744f53-ec91-4f5b-85e2-d6f150221616.json	3e6c5a77724a182243f898fe82b806aed906eb9a4b156276f9677d471c01f409
rollouts/660efecd-4c0a-4092-9713-fd3677d3528f.json	7435d69b8bc0ea51765cde0eff6a2e1926c36059307be111e3c82807d9bf2646
rollouts/662f4a3d-ea26-4e49-b9ee-24a190ed2db3.json	6d0227476d2a60f1f595800852d3c2b187d38591cb7b7fdcec39071104c6ef3c
rollouts/7d1b1531-f014-45b5-93f8-c578d036d820.json	755def32bc1166a4edba484cef698b57f399394937a13bf5b15c8f1ca9b93074
rollouts/879eca05-6ed5-481f-a4c0-58a675435c85.json	fcff7d31005d2bc12895065eb0356058fd81fe6573074fbc73a54ce29ad96118
rollouts/899ad596-b098-4a27-b1a8-ddc53e4346f2.json	835c88ee481264fa1b0b844eb4ed17d76c60185ddf56214e73629beca582145e
rollouts/8b5af559-b0f7-4aaa-9748-435f7836143e.json	42b8c77bec005e209190119a559320f4908762c2b85460f19e97d0cbf4ebe051
rollouts/9902eaa3-4af4-4039-946a-2b4fbbb3eb0c.json	8f91bdd28114bb10fa0f57b3c9030dda7a96fc3c26b3f47a9fffb19f7d161cc3e
rollouts/ab27ff4e-329f-4ad0-b20e-cd10c9e7e4b6.json	f8c53bac344a5ae08a4f020dbe6c65b3f8d0f66275cad641651c26ff098fa84b
rollouts/affa9ea3-9250-418d-9507-8cfbf429b84c.json	47d6031bef983bfd8f5878d13bbf4ef24c192a42f7a79186490048b9745c788a
rollouts/b0c7e737-a914-495b-ba72-aee2b3fd18b1.json	4048e436ea6ce3169b884cbfc4620936645573daf579ad52a5429b04032c0c19
rollouts/b761b8be-4d5d-4934-8e0e-6d63e67c573d.json	b3d3e9ee57a3f991b1d0c127d595178e1c9367b8ca17879d948d08f9fa584477
rollouts/cd337d97-6960-429c-99dd-10a17b935b0d.json	0051949472962b064fd9b7934f6acbb607de390a516e2fc2b5f59d7d03efe4c9
rollouts/de0f078f-5b50-4853-873f-208337f892f9.json	0fdc5de0bbf5032da58e7fa47df2f5000c926f566ec1c3ff8fb450f1aea2ca91
rollouts/e90dd0c8-6384-4a8e-857d-560d91cc8a8d.json	6fa264a19e220ee863979f0f2ac37f82b75e269d3ba5d54b0a7dd2957786f56e

ARTIFACT	SHA-256
rollouts/eddbc8c8-36ca-4533-b298-2cd72c5d5c62.json	7c4702032ea6d2d74a3fcf758665e109cb780012c387e4c7c6cd53b5cf1cb72f
rollouts/fd449898-96bc-4c78-bd82-2d37535c2987.json	1030a0bd7437271001c13f01cd8f725b116c4c08d9e13ddc4d7132bc925865cc
spec	c5580395b0da0e89877b88e47c897742665c2711455762bc87c742b9f68019ec
summary	adcd3c1bdb43057c1e2f1114afe7303277a5f95db0d1786f64407524d6eead72

Reproducing this run

Every rollout can be re-derived: the exam is frozen (fingerprint above), the grader ran in the recorded container image, and each rollout's full trace — every tool call, every result, every word — is in the evidence bundle at its rollouts/<id>.json path. A re-run uses the same job shape: the exam id, your endpoint, n, and the runtime. The record id names this run forever.

Appendix C — redacted transcripts

For each rollout: the tool calls it made, in order — which tool, on which of the model's own paths, whether the call succeeded, and how many characters came back. What a tool returned is not here, and neither is the submission: on this exam the submitted content is the graded answer, so no diff, no answer value and no model claim is printed. A request for the reference solution or the hidden tests is recorded as a request; the path it asked for is replaced by a label. See Appendix D, "What this report withholds and why".

Disclosure class applied to this exam: answer-key disclosure (ownership vvdex_proprietary — engine registry; family harness_eval): the exam is VVDex evaluation IP and the graded output is an answer or a keyed value, so no submission content, diff, model claim or hidden-assertion name is published.

3c178cb0 · cli:cli/claude-sonnet · seed 0 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	src/rates.py	ok	1 335
1	edit_file	src/rates.py	ok	50
2	run_tests	—	failed	95
3	run_tests	—	ok	1 581
4	read_file	src/rates.py	ok	157
5	write_file	src/rates.py	ok	62
6	run_tests	—	ok	1 581
7	read_file	src/rates.py	ok	157
8	write_file	src/rates.py	ok	61
9	run_tests	—	ok	268
10	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (210 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

230a8274 · cli:cli/claude-sonnet · seed 1 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	src/rates.py	ok	1 335
1	edit_file	src/rates.py	ok	50
2	run_tests	—	failed	95
3	run_tests	—	ok	1 581
4	read_file	src/rates.py	ok	157
5	write_file	src/rates.py	ok	62
6	run_tests	—	ok	1 581
7	read_file	src/rates.py	ok	157
8	write_file	src/rates.py	ok	61
9	run_tests	—	ok	268

STEP	TOOL	PATH	RESULT	CHARS
10	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (212 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

459f9de3 · cli:cli/claude-sonnet · seed 2 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	src/rates.py	ok	1 335
1	edit_file	src/rates.py	ok	50
2	run_tests	—	failed	95
3	run_tests	—	ok	1 581
4	read_file	src/rates.py	ok	157
5	write_file	src/rates.py	ok	62
6	run_tests	—	ok	1 581
7	read_file	src/rates.py	ok	157
8	write_file	src/rates.py	ok	61
9	run_tests	—	ok	268
10	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (210 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

3a7da64d · cli:cli/claude-sonnet · seed 3 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	src/rates.py	ok	1 335
1	edit_file	src/rates.py	ok	50
2	run_tests	—	failed	95
3	run_tests	—	ok	1 581
4	read_file	src/rates.py	ok	157
5	write_file	src/rates.py	ok	62
6	run_tests	—	ok	1 581
7	write_file	src/rates.py	ok	61
8	run_tests	—	ok	268
9	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (199 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

1e7f41d8 · cli:cli/claude-sonnet · seed 4 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	src/rates.py	ok	1 335
1	edit_file	src/rates.py	ok	50
2	run_tests	—	failed	95
3	run_tests	—	ok	1 581
4	read_file	src/rates.py	ok	157
5	write_file	src/rates.py	ok	62
6	run_tests	—	ok	1 581
7	read_file	src/rates.py	ok	157
8	write_file	src/rates.py	ok	61
9	run_tests	—	ok	268
10	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (210 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

9902eaa3 · cli:cli/claude-sonnet · seed 5 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	src/rates.py	ok	1 335
1	edit_file	src/rates.py	ok	50
2	run_tests	—	failed	95
3	run_tests	—	ok	1 581
4	read_file	src/rates.py	ok	157
5	write_file	src/rates.py	ok	62
6	run_tests	—	ok	1 581
7	read_file	src/rates.py	ok	157
8	write_file	src/rates.py	ok	61
9	run_tests	—	ok	268
10	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (210 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

12f95490 · cli:cli/claude-sonnet · seed 6 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	src/rates.py	ok	1 335
1	edit_file	src/rates.py	ok	50

STEP	TOOL	PATH	RESULT	CHARS
2	run_tests	—	failed	95
3	run_tests	—	ok	1 581
4	read_file	src/rates.py	ok	157
5	write_file	src/rates.py	ok	62
6	run_tests	—	ok	1 581
7	read_file	src/rates.py	ok	157
8	write_file	src/rates.py	ok	62
9	run_tests	—	ok	268
10	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (214 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

affa9ea3 · cli:cli/claude-sonnet · seed 7 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	src/rates.py	ok	1 335
1	edit_file	src/rates.py	ok	50
2	run_tests	—	failed	95
3	run_tests	—	ok	1 581
4	read_file	src/rates.py	ok	157
5	write_file	src/rates.py	ok	62
6	run_tests	—	ok	1 581
7	write_file	src/rates.py	ok	61
8	run_tests	—	ok	268
9	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (199 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

ab27ff4e · cli:cli/claude-sonnet · seed 8 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	src/rates.py	ok	1 335
1	edit_file	src/rates.py	ok	50
2	run_tests	—	failed	95
3	run_tests	—	ok	1 581
4	read_file	src/rates.py	ok	157
5	write_file	src/rates.py	ok	62
6	run_tests	—	ok	1 581

STEP	TOOL	PATH	RESULT	CHARS
7	read_file	src/rates.py	ok	157
8	write_file	src/rates.py	ok	61
9	run_tests	—	ok	268
10	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (212 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

cd337d97 · cli:cli/claude-sonnet · seed 9 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	src/rates.py	ok	1 335
1	edit_file	src/rates.py	ok	50
2	run_tests	—	failed	95
3	run_tests	—	ok	1 581
4	read_file	src/rates.py	ok	157
5	write_file	src/rates.py	ok	62
6	run_tests	—	ok	1 581
7	write_file	src/rates.py	ok	61
8	run_tests	—	ok	268
9	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (199 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

879eca05 · cli:cli/codex · seed 0 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 633
1	read_file	—	failed	35
2	read_file	src/rates.py	ok	1 335
3	edit_file	src/rates.py	ok	50
4	read_file	src/rates.py	ok	157
5	read_file	src/rates.py	ok	0
6	run_tests	—	failed	95
7	run_tests	—	ok	1 581
8	write_file	src/rates.py	ok	62
9	read_file	src/rates.py	ok	157
10	read_file	src/rates.py	ok	40
11	read_file	src/rates.py	ok	0

STEP	TOOL	PATH	RESULT	CHARS
12	write_file	src/rates.py	ok	62
13	read_file	src/rates.py	ok	148
14	run_tests	—	ok	268
15	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (270 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

fd449898 · cli:cli/codex · seed 1 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	list_files	.	ok	94
1	read_file	TASK.md	ok	1 633
2	read_file	src/rates.py	ok	1 335
3	edit_file	src/rates.py	ok	50
4	read_file	src/rates.py	ok	157
5	read_file	src/rates.py	ok	0
6	read_file	src/rates.py	ok	0
7	run_tests	—	failed	95
8	run_tests	—	ok	1 581
9	write_file	src/rates.py	ok	62
10	read_file	src/rates.py	ok	157
11	read_file	src/rates.py	ok	20
12	read_file	src/rates.py	ok	0
13	run_tests	—	ok	1 581
14	write_file	src/rates.py	ok	61
15	read_file	src/rates.py	ok	892
16	run_tests	—	ok	268
17	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (288 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

5d078aad · cli:cli/codex · seed 2 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	list_files	.	ok	94
1	read_file	TASK.md	ok	1 633
2	read_file	—	failed	35
3	read_file	src/rates.py	ok	1 335

STEP	TOOL	PATH	RESULT	CHARS
4	read_file	tests/test_rates.py	ok	719
5	edit_file	src/rates.py	ok	50
6	read_file	src/rates.py	ok	157
7	read_file	src/rates.py	ok	0
8	read_file	src/rates.py	ok	39
9	read_file	src/rates.py	ok	60
10	read_file	src/rates.py	ok	0
11	run_tests	–	failed	95
12	run_tests	–	ok	1 581
13	write_file	src/rates.py	ok	62
14	read_file	src/rates.py	ok	157
15	run_tests	–	ok	1 581
16	write_file	src/rates.py	ok	61
17	read_file	src/rates.py	ok	787
18	run_tests	–	ok	268
19	submit	–	ok	85

Submission diff withheld by the public-disclosure policy (327 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

4ac68744 · cli:cli/codex · seed 3 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 633
1	read_file	–	failed	35
2	read_file	src/rates.py	ok	1 335
3	edit_file	src/rates.py	ok	50
4	read_file	src/rates.py	ok	157
5	read_file	src/rates.py	ok	0
6	run_tests	–	failed	95
7	read_file	src/rates.py	ok	0
8	run_tests	–	ok	1 581
9	write_file	src/rates.py	ok	62
10	read_file	src/rates.py	ok	157
11	run_tests	–	ok	1 581
12	write_file	src/rates.py	ok	61
13	read_file	src/rates.py	ok	791
14	run_tests	–	ok	268

STEP	TOOL	PATH	RESULT	CHARS
15	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (284 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

7d1b1531 · cli:cli/codex · seed 4 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 633
1	read_file	—	failed	35
2	read_file	src/rates.py	ok	1 335
3	edit_file	src/rates.py	ok	50
4	run_tests	—	failed	95
5	run_tests	—	ok	1 581
6	read_file	src/rates.py	ok	157
7	write_file	src/rates.py	ok	62
8	read_file	src/rates.py	ok	157
9	write_file	src/rates.py	ok	62
10	read_file	src/rates.py	ok	1 278
11	run_tests	—	ok	268
12	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (213 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

662f4a3d · cli:cli/codex · seed 5 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 633
1	read_file	—	failed	35
2	read_file	src/rates.py	ok	1 335
3	read_file	tests/test_rates.py	ok	719
4	edit_file	src/rates.py	ok	50
5	read_file	src/rates.py	ok	157
6	read_file	src/rates.py	ok	0
7	run_tests	—	failed	95
8	run_tests	—	ok	1 581
9	write_file	src/rates.py	ok	62
10	read_file	src/rates.py	ok	157
11	read_file	src/rates.py	ok	20

STEP	TOOL	PATH	RESULT	CHARS
12	write_file	src/rates.py	ok	62
13	read_file	src/rates.py	ok	1 198
14	run_tests	—	ok	268
15	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (246 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

660efecd · cli:cli/codex · seed 6 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 633
1	list_files	.	ok	94
2	read_file	src/rates.py	ok	1 335
3	read_file	tests/test_rates.py	ok	719
4	edit_file	src/rates.py	ok	50
5	read_file	src/rates.py	ok	157
6	run_tests	—	failed	95
7	run_tests	—	ok	1 581
8	write_file	src/rates.py	ok	62
9	read_file	src/rates.py	ok	157
10	read_file	src/rates.py	ok	0
11	read_file	src/rates.py	ok	0
12	write_file	src/rates.py	ok	62
13	run_tests	—	ok	268
14	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (237 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

51c40893 · cli:cli/codex · seed 7 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	list_files	.	ok	94
1	read_file	TASK.md	ok	1 633
2	read_file	src/rates.py	ok	1 335
3	read_file	tests/test_rates.py	ok	719
4	edit_file	src/rates.py	ok	50
5	run_tests	—	failed	95
6	run_tests	—	ok	1 581

STEP	TOOL	PATH	RESULT	CHARS
7	read_file	src/rates.py	ok	157
8	write_file	src/rates.py	ok	62
9	read_file	src/rates.py	ok	157
10	read_file	src/rates.py	ok	20
11	write_file	src/rates.py	ok	61
12	read_file	src/rates.py	ok	892
13	run_tests	—	ok	268
14	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (255 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

899ad596 · cli:cli/codex · seed 8 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 633
1	read_file	—	failed	35
2	read_file	src/rates.py	ok	1 335
3	edit_file	src/rates.py	ok	50
4	read_file	src/rates.py	ok	157
5	read_file	src/rates.py	ok	0
6	read_file	src/rates.py	ok	60
7	run_tests	—	failed	95
8	run_tests	—	ok	1 581
9	write_file	src/rates.py	ok	62
10	read_file	src/rates.py	ok	157
11	read_file	src/rates.py	ok	20
12	run_tests	—	ok	1 581
13	write_file	src/rates.py	ok	61
14	read_file	src/rates.py	ok	892
15	run_tests	—	ok	268
16	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (278 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

e90dd0c8 · cli:cli/codex · seed 9 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 633

STEP	TOOL	PATH	RESULT	CHARS
1	list_files	.	ok	94
2	read_file	src/rates.py	ok	1 335
3	read_file	tests/test_rates.py	ok	719
4	edit_file	src/rates.py	ok	50
5	read_file	src/rates.py	ok	157
6	read_file	src/rates.py	ok	20
7	read_file	src/rates.py	ok	0
8	write_file	src/rates.py	ok	62
9	read_file	src/rates.py	ok	157
10	run_tests	—	failed	95
11	run_tests	—	ok	1 581
12	write_file	src/rates.py	ok	62
13	read_file	src/rates.py	ok	1 168
14	run_tests	—	ok	268
15	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (267 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

0711410a · mistral:codestral-latest · seed 0 · context_exhaustion

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	—	failed	35
1	invalid_tool_call	"}	failed	103
2	list_files	.	ok	94
3	read_file	TASK.md	ok	1 633
4	read_file	src/rates.py	ok	1 335
5	run_tests	—	failed	95
6	edit_file	src/rates.py	ok	50
7	run_tests	—	ok	1 581
8	edit_file	src/rates.py	failed	29
9	edit_file	src/rates.py	failed	29
10	edit_file	src/rates.py	failed	29
11	edit_file	src/rates.py	failed	29
12	edit_file	src/rates.py	failed	29
13	edit_file	src/rates.py	failed	29
14	edit_file	src/rates.py	failed	29
15	edit_file	src/rates.py	failed	29

STEP	TOOL	PATH	RESULT	CHARS
16	edit_file	src/rates.py	failed	29
17	edit_file	src/rates.py	failed	29
18	edit_file	src/rates.py	failed	29
19	edit_file	src/rates.py	failed	29
20	edit_file	src/rates.py	failed	29
21	edit_file	src/rates.py	failed	29
22	edit_file	src/rates.py	failed	29
23	edit_file	src/rates.py	failed	29

Submission diff withheld by the public-disclosure policy (350 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

de0f078f · mistral:codestral-latest · seed 1 · context_exhaustion

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	—	failed	35
1	invalid_tool_call	"}	failed	103
2	list_files	.	ok	94
3	read_file	src/rates.py	ok	1 335
4	edit_file	src/rates.py	ok	50
5	run_tests	—	failed	95
6	run_tests	—	ok	1 581
7	edit_file	src/rates.py	failed	29
8	edit_file	src/rates.py	failed	29
9	write_file	src/rates.py	ok	62
10	run_tests	—	ok	1 581
11	write_file	src/rates.py	ok	62
12	run_tests	—	ok	1 579
13	edit_file	src/rates.py	ok	50
14	run_tests	—	ok	1 579
15	edit_file	src/rates.py	ok	50
16	run_tests	—	ok	1 579
17	edit_file	src/rates.py	ok	50
18	run_tests	—	ok	1 579
19	edit_file	src/rates.py	ok	50
20	run_tests	—	ok	1 579
21	edit_file	src/rates.py	ok	50
22	run_tests	—	ok	1 579

STEP	TOOL	PATH	RESULT	CHARS
23	edit_file	src/rates.py	ok	50

Submission diff withheld by the public-disclosure policy (341 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

5c42bd02 · mistral:codestral-latest · seed 2 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	src/rates.py	ok	1 335
1	run_tests	—	failed	95
2	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (57 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

3498fcf5 · mistral:codestral-latest · seed 3 · context_exhaustion

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	—	failed	35
1	invalid_tool_call	"}	failed	103
2	list_files	.	ok	94
3	read_file	TASK.md	ok	1 633
4	read_file	src/rates.py	ok	1 335
5	edit_file	src/rates.py	ok	50
6	run_tests	—	failed	95
7	run_tests	—	ok	1 581
8	edit_file	src/rates.py	ok	50
9	run_tests	—	ok	1 561
10	edit_file	src/rates.py	ok	50
11	run_tests	—	ok	1 561
12	edit_file	src/rates.py	ok	50
13	run_tests	—	ok	1 570
14	edit_file	src/rates.py	ok	50
15	run_tests	—	ok	1 543
16	edit_file	src/rates.py	ok	50
17	run_tests	—	ok	1 543
18	edit_file	src/rates.py	ok	50
19	run_tests	—	ok	1 543
20	edit_file	src/rates.py	ok	50
21	run_tests	—	ok	1 543

STEP	TOOL	PATH	RESULT	CHARS
22	edit_file	src/rates.py	ok	50
23	run_tests	—	ok	1 543

Submission diff withheld by the public-disclosure policy (337 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

5e744f53 · mistral:codestral-latest · seed 4 · edited_tested_no_submit

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	—	failed	35
1	read_file	src/rates.py	ok	1 335
2	run_tests	—	failed	95
3	edit_file	src/rates.py	ok	50
4	run_tests	—	ok	1 581
5	edit_file	src/rates.py	ok	50
6	run_tests	—	ok	1 561
7	edit_file	src/rates.py	ok	50
8	run_tests	—	ok	2 005
9	edit_file	src/rates.py	ok	50
10	run_tests	—	ok	1 999
11	edit_file	src/rates.py	ok	50
12	run_tests	—	ok	1 999
13	edit_file	src/rates.py	ok	50
14	run_tests	—	ok	1 999
15	edit_file	src/rates.py	ok	50
16	run_tests	—	ok	1 999
17	edit_file	src/rates.py	ok	50
18	run_tests	—	ok	1 999
19	edit_file	src/rates.py	ok	50
20	run_tests	—	ok	1 999
21	edit_file	src/rates.py	ok	50
22	run_tests	—	ok	1 999
23	edit_file	src/rates.py	ok	50

Submission diff withheld by the public-disclosure policy (339 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

b761b8be · mistral:codestral-latest · seed 5 · edited_tested_no_submit

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	—	failed	35
1	invalid_tool_call	"}	failed	103
2	list_files	.	ok	94
3	read_file	src/rates.py	ok	1 335
4	edit_file	src/rates.py	ok	50
5	run_tests	—	failed	95
6	run_tests	—	ok	1 581
7	edit_file	src/rates.py	ok	50
8	run_tests	—	ok	1 561
9	edit_file	src/rates.py	ok	50
10	run_tests	—	ok	1 539
11	edit_file	src/rates.py	ok	50
12	run_tests	—	ok	1 540
13	edit_file	src/rates.py	ok	50
14	run_tests	—	ok	1 540
15	edit_file	src/rates.py	ok	50
16	run_tests	—	ok	1 540
17	edit_file	src/rates.py	ok	50
18	run_tests	—	ok	1 540
19	edit_file	src/rates.py	ok	50
20	run_tests	—	ok	1 540
21	edit_file	src/rates.py	ok	50
22	run_tests	—	ok	1 540
23	edit_file	src/rates.py	ok	50

Submission diff withheld by the public-disclosure policy (338 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

eddbc8c8 · mistral:codestral-latest · seed 6 · edited_tested_no_submit

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	—	failed	35
1	read_file	TASK.md	ok	1 633
2	read_file	src/rates.py	ok	1 335
3	edit_file	src/rates.py	ok	50
4	run_tests	—	failed	95
5	run_tests	—	ok	1 581

STEP	TOOL	PATH	RESULT	CHARS
6	edit_file	src/rates.py	failed	29
7	edit_file	src/rates.py	failed	29
8	edit_file	src/rates.py	failed	29
9	edit_file	src/rates.py	failed	29
10	edit_file	src/rates.py	failed	29
11	edit_file	src/rates.py	failed	29
12	edit_file	src/rates.py	failed	29
13	edit_file	src/rates.py	failed	29
14	edit_file	src/rates.py	failed	29
15	edit_file	src/rates.py	failed	29
16	edit_file	src/rates.py	failed	29
17	edit_file	src/rates.py	failed	29
18	edit_file	src/rates.py	failed	29
19	edit_file	src/rates.py	failed	29
20	edit_file	src/rates.py	failed	29
21	edit_file	src/rates.py	failed	29
22	edit_file	src/rates.py	failed	29
23	edit_file	src/rates.py	failed	29

Submission diff withheld by the public-disclosure policy (355 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

b0c7e737 · mistral:codestral-latest · seed 7 · edited_tested_no_submit

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	—	failed	35
1	invalid_tool_call	"}	failed	103
2	list_files	.	ok	94
3	read_file	src/rates.py	ok	1 335
4	edit_file	src/rates.py	ok	50
5	run_tests	—	failed	95
6	run_tests	—	ok	1 581
7	edit_file	src/rates.py	failed	29
8	edit_file	src/rates.py	failed	29
9	edit_file	src/rates.py	failed	29
10	edit_file	src/rates.py	failed	29
11	edit_file	src/rates.py	failed	29
12	edit_file	src/rates.py	failed	29

STEP	TOOL	PATH	RESULT	CHARS
13	edit_file	src/rates.py	failed	29
14	edit_file	src/rates.py	failed	29
15	edit_file	src/rates.py	failed	29
16	edit_file	src/rates.py	failed	29
17	edit_file	src/rates.py	failed	29
18	edit_file	src/rates.py	failed	29
19	edit_file	src/rates.py	failed	29
20	edit_file	src/rates.py	failed	29
21	edit_file	src/rates.py	failed	29
22	edit_file	src/rates.py	failed	29
23	edit_file	src/rates.py	failed	29

Submission diff withheld by the public-disclosure policy (350 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

3061565d · mistral:codestral-latest · seed 8 · context_exhaustion

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	src/rates.py	ok	1 335
1	edit_file	src/rates.py	ok	50
2	run_tests	—	failed	95
3	run_tests	—	ok	1 581
4	read_file	src/rates.py	ok	157
5	edit_file	src/rates.py	ok	50
6	run_tests	—	ok	1 001
7	read_file	src/rates.py	ok	318
8	edit_file	src/rates.py	ok	50
9	run_tests	—	ok	1 001
10	read_file	src/rates.py	ok	318
11	edit_file	src/rates.py	ok	50
12	run_tests	—	ok	1 001
13	read_file	src/rates.py	ok	318
14	edit_file	src/rates.py	ok	50
15	run_tests	—	ok	1 001
16	read_file	src/rates.py	ok	318
17	edit_file	src/rates.py	ok	50
18	run_tests	—	ok	1 001
19	read_file	src/rates.py	ok	318

STEP	TOOL	PATH	RESULT	CHARS
20	edit_file	src/rates.py	ok	50
21	run_tests	—	ok	1 001
22	read_file	src/rates.py	ok	318
23	edit_file	src/rates.py	ok	50

Submission diff withheld by the public-disclosure policy (343 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

8b5af559 · mistral:codestral-latest · seed 9 · context_exhaustion

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	—	failed	35
1	invalid_tool_call	"}	failed	103
2	list_files	.	ok	94
3	read_file	TASK.md	ok	1 633
4	read_file	src/rates.py	ok	1 335
5	run_tests	—	failed	95
6	run_tests	—	ok	1 583
7	edit_file	src/rates.py	ok	50
8	run_tests	—	ok	1 581
9	edit_file	src/rates.py	ok	50
10	run_tests	—	ok	1 581
11	edit_file	src/rates.py	ok	50
12	run_tests	—	ok	1 561
13	edit_file	src/rates.py	ok	50
14	run_tests	—	ok	1 539
15	edit_file	src/rates.py	ok	50
16	run_tests	—	ok	1 562
17	edit_file	src/rates.py	ok	50
18	run_tests	—	ok	1 540
19	edit_file	src/rates.py	ok	50
20	run_tests	—	ok	1 562
21	edit_file	src/rates.py	ok	50
22	run_tests	—	ok	1 541
23	edit_file	src/rates.py	ok	50

Submission diff withheld by the public-disclosure policy (343 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

12 Exam · grounded-rag-1

vvdex.knowledge.grounded-rag-1 · record fr-20260902-2e7f4171 · record digest 157ec324a925c0b1... · harness clean · containment clean

Executive summary

MODELS TESTED 3	CERTIFIED PASSES 11 of 29 graded rollouts	SUBMITTED 29 of 29 graded rollouts that called submit	NOT GRADED — LANE ERRORS 1 of 30 attempts
---------------------------------------	--	--	--

Purpose

This report presents a controlled evaluation of 3 model lane(s) against the certified exam vvdex.knowledge.grounded-rag-1 (v0.1.0, family rag_grounding). Every lane received the identical frozen workspace, tool contract and limits; grading ran afterwards, in isolation, against hidden tests no lane ever saw. The purpose is to state, with reproducible evidence, what each lane did and where it stopped.

Outcome

11 of 29

graded rollouts passed the independent hidden tests and were certified — 1 of 30 attempts not graded (lane errors)

11 of 29 graded rollouts passed. cli/claude-sonnet, cli/codex completed the full loop (inspect, edit, run the visible tests, submit) and the sealed grader accepted the submission. The most common way to fail was submitted_failed — submitted a change; the independent hidden tests still failed.

Key findings

- F-01 cli/claude-sonnet, cli/codex passed the independent hidden tests
- F-02 submitted_failed × 18 (cli/claude-sonnet, codestral-latest)
- F-03 model_api_failure × 1 (codestral-latest)

Each finding is stated in full, with evidence, in the Findings section; recommendations for acting on them follow it. Elapsed wall clock for the whole engagement: 1500.2s.

Exam definition

In scope. The ability of each configured model lane to complete this one certified task end to end — inspect the workspace, produce a correct edit, verify it with the visible tests, and submit — within 30 tool steps and 7320s of wall clock, at 10 rollout(s) per lane. Behavioural measurements along the way (tool discipline, grounding, overclaim) are in scope because they are read from the recorded trace, not judged by another model.

Out of scope. General model quality, performance on other task families, prompt-tuning on the lane's behalf, and any comparison this run's sample size cannot support. A lane's API failure is reported as infrastructure, never as model capability.

The instrument. Exam `vvdex.knowledge.grounded-rag-1` was certified before this run: its reference solution passes, an empty submission fails, its grader distinguishes near-misses, and its sealed evidence is unchanged by this evaluation (see Certification evidence).

What was measured

EXAM	VERSION	FAMILY	ROLLOUTS
vvdex.knowledge.grounded-rag-1	0.1.0	rag_grounding	30 (10 per model)

Where the defect came from

This exam has no upstream repository to credit: the defect, the reference solution and the hidden suite are VVDex's own.

Certification evidence

4 of 5 certification pillars are recorded in this package, loaded from the receipts sealed into the exam and matching this run's exam fingerprint. What is recorded is shown with its real value; what is not is named below and is not claimed.

FROZEN BASELINE

FAIL

expected — an empty submission over 2 run(s) must not pass

→

GOLD / REFERENCE

PASS

expected — the reference solution over 2 run(s) must pass

GRADER CONTROLS

7 / 7

The grader accepted the reference and rejected every near-miss.

ATTACK PROBES

10 / 10 blocked

0 trivial exploit(s) found.

RUNTIME

sha256:679b36225a1f83238efba8c71b9cde3c24caaeacc9b682ae92819cb55eec328f

network policy: deny

Certification metadata is not included in this report package for:

- sealed evidence bundle

Benchmark results remain valid as run evidence; the pillar(s) above are not substantiated by this document and this report does not claim them.

What the package does carry: the exam fingerprint below is the contract digest this run was executed against — a holder of the certification bundle for this fingerprint can join the two independently.

EXAM FINGERPRINT

4880a527ef97ca2f3fe6b57be026c3dc1be51e9e104d79fc13a3e558331dda92

The contract digest this run was executed against.

Evaluation configuration

REPORT	fr-20260902-2e7f4171
ISSUED	2026-09-02
SUBJECT EXAM	vvdex.knowledge.grounded-rag-1 · v0.1.0
FAMILY	rag_grounding
PREPARED BY	VVDex Forge engine vvdex-env 0.1.0
CLASSIFICATION	Independent model evaluation report — independently verifiable
CERTIFICATION	Exam certification: partially recorded (see Certification evidence)
STATUS	FINAL · report sealed against its own digest

Timeline

WHEN (UTC)	EVENT
2026-09-01 23:53:46	Evaluation run started
2026-09-02 00:18:47	Evaluation run finished
2026-09-02 00:18:47	Report generated from the sealed run evidence

Models under test

MODEL	PROVIDER	ENDPOINT	BILLING
Claude Code CLI (Sonnet, subscription)	cli	local runtime	operator subscription, flat — not billed per token
Codex CLI (GPT, subscription)	cli	local runtime	operator subscription, flat — not billed per token
codestral-latest	mistral	VVDex-configured	VVDex free-quota

A customer endpoint is recorded by origin only — never its port, its path, or its key.

Limits

RUNTIME docker	STEP LIMIT 30	WALL-CLOCK LIMIT 7320s	MAX OUTPUT TOKENS 2048
TEMPERATURE 0.2	RETRY POLICY none	COMMAND BUDGET 40	TOOLS OFFERED list_files, read_file, write_file, edit_file, run_tests, submit, run_command

The evaluated model never saw the hidden tests or the reference solution, and could not change its result. Grading ran in a separate step, after submission, on a container with no network.

Model outcomes

Denominators. Attempts requested: 30. Graded (evaluable) rollouts: 29. Not graded — lane errors: 1. A lane error is a rollout the API or the runtime ended before the hidden grader ran: it is listed, excluded from every rate and pass@k, and never silently dropped. Passes are certified passes: hidden grader exit 0 with integrity verified.

MODEL	ATTEMPTS	GRADED	PASSED	MODEL FAILS	LANE ERRORS	SUBMITTED	MEAN TIME	TYPICAL DEEPEST STAGE
Claude Code CLI (Sonnet, subscription) cli/claude-sonnet	10	10	1 / 10	9	0	10 / 10	45.8s	Graded
Codex CLI (GPT, subscription) cli/codex	10	10	10 / 10	0	0	10 / 10	67.8s	Passed hidden tests
codestral-latest codestral-latest	10	9	0 / 9	9	1	9 / 9	18.0s	Graded

Every rollout, with its own deepest stage and grade, is in Appendix A.

Success rate, confidence and pass@k

Every rate on this page is a measurement of a finite number of rollouts, so every rate carries the interval that measurement actually supports. n is the evaluable count — attempts minus rollouts the lane ended before grading; those are listed under "not graded" and sit in no rate, interval or pass@k.

MODEL	ATTEMPTS	N	NOT GRADED	PASSES	RATE	95% INTERVAL	SEEDS	PASS@1	PASS@3	PASS@5	PASS@10
cli/claude-sonnet	10	10	0	1	10.0%	[2%, 40%]	0–9	10.0%	30.0%	50.0%	100.0%
cli/codex	10	10	0	10	100.0%	[72%, 100%]	0–9	100.0%	100.0%	100.0%	100.0%
codestral-latest	10	9	1	0	0.0%	[0%, 30%]	0–9	0.0%	0.0%	0.0%	—

Interval and pass@k methods are stated in full in Appendix D. pass@k is shown for k = 1, 3, 5, 10; the sealed record carries every k up to n.

Model comparison

MODEL	PASS@1	95% INTERVAL	TOKENS	COST	MEAN ROLLOUT	OVERCLAIM RATE	TOP FAILURE CLASS
cli/claude-sonnet	10.0%	[2%, 40%]	n/a	n/a	45.8s	0.0%	submitted_failed
cli/codex	100.0%	[72%, 100%]	n/a	n/a	67.8s	0.0%	none
codestral-latest	0.0%	[0%, 30%]	225 330	not billed	18.0s	0.0%	submitted_failed

Tokens and cost read n/a where a lane reports no telemetry (a flat-rate CLI) — an absent measurement, not zero. pass@1 and its interval are over evaluable rollouts only.

This is not a leaderboard. Separated by this run (intervals do not overlap): cli/claude-sonnet vs cli/codex; cli/codex vs codestral-latest. Not separated (intervals overlap): cli/claude-sonnet vs codestral-latest. Below the 10-rollout floor for reading a rate as a rate: codestral-latest (9) — read the count, not the percentage. The primary comparison in this report is the stage funnel below — where a model stops is a finding; a percentage over this many rollouts is not.

Stage funnel

The funnel is the primary reading of a run. Where a model stops says more than a single pass percentage does, especially at low N.

STAGE	CLI:CLI/CLAUDE-SONNET	CLI:CLI/CODEX	MISTRAL:CODESTRAL-LATEST
Inspected	10 / 10	10 / 10	9 / 9
Edited	8 / 10	10 / 10	9 / 9
Ran tests	10 / 10	10 / 10	9 / 9
Submitted	10 / 10	10 / 10	9 / 9
Graded	10 / 10	10 / 10	9 / 9
Passed hidden tests	1 / 10	10 / 10	0 / 9

Deepest stage reached

MODEL	DEEPEST (ANY ROLLOUT)	TYPICAL (MOST ROLLOUTS)	HIDDEN-TEST PASSES	NOT GRADED
cli:cli/claude-sonnet	Passed hidden tests	Graded	1 / 10	0
cli:cli/codex	Passed hidden tests	Passed hidden tests	10 / 10	0
mistral:codestral-latest	Graded	Graded	0 / 9	1

Stage definitions are in Appendix D. The matrix counts evaluable rollouts; lane errors are shown beside it, not inside it. The stages are not strictly nested: a model can submit without editing, and the matrix shows that rather than hiding it. Each rollout's own deepest stage is in Appendix A.

Behavioural analysis

Long-horizon behaviour

Counted off the recorded trace of each rollout. A dead end is a step that moved nothing: a re-read of a slice already returned, a repeat of the previous action, or a call to a tool this exam does not have.

MODEL	N	STEPS USED	RAN OUT OF STEPS	COMMANDS	CHANGED THE TREE	REFUSED (OVER BUDGET)	REVISITS	DEAD ENDS / ROLLOUT
cli:cli/claude-sonnet	10	9 / 30	0	0 / 40	0	0	0	0
cli:cli/codex	10	14 / 30	0	0 / 40	0	0	0	0
mistral:codestral-latest	10	8 / 30	0	0 / 40	0	0	0	0

Each column is defined in Appendix D.

Hallucination & overclaim

Three measurements, each with a stated definition. None of them is a judgement of the model's prose by another model — every number below comes from comparing what the model wrote against what was on disk, or against what the independent grader returned.

MODEL	N	GROUNDING FAILURES	KINDS	ROLLOUTS AFFECTED	OVERCLAIMS	OVERCLAIM RATE	RE-READS
cli:cli/claude-sonnet	10	0	none	0 / 10	0 / 10	0.0%	0
cli:cli/codex	10	0	none	0 / 10	0 / 10	0.0%	0
mistral:codestral-latest	10	0	none	0 / 10	0 / 10	0.0%	0

No grounding failures and no false pass claims were detected across these 30 rollout(s). No lane re-read material it had already been given.

Every term above is defined in Appendix D.

Efficiency & telemetry

This run has no per-token price attached — a local model, a free quota, or a customer endpoint we are not billed for — so spend is reported in tokens. Tokens are what was actually measured; a dollar figure here would be invented. Lanes run through a flat-rate local CLI report no token telemetry; their cells say n/a — an absent measurement, not zero. Rollout counts here are evaluable rollouts.

MODEL	ROLLOUTS	TOKENS	TOKENS / ROLLOUT	COST	PASSES	RATE
cli:cli/claude-sonnet	10	n/a	n/a	n/a – telemetry unavailable	1 / 10	10.0%
cli:cli/codex	10	n/a	n/a	n/a – telemetry unavailable	10 / 10	100.0%
mistral:codestral-latest	9	225 330	25 036	not billed	0 / 9	0.0%

Findings

Every finding below is derived from the recorded data of this run — none is an opinion.
Severity: PASS a lane succeeded · ATTENTION a capability gap · LANE / INFRA infrastructure, not capability · NOTE a property of the measurement.

F-01	PASS
cli/claude-sonnet, cli/codex passed the independent hidden tests	
Severity: PASS	Evidence: Model outcomes · Stage funnel · Appendix A
11 of 29 graded rollouts completed the full loop (inspect, edit, run the visible tests, submit) and the sealed grader accepted the submission. This demonstrates the task is solvable under the published limits.	
F-01 · derived from the recorded run data	
F-02	ATTENTION
submitted_failed × 18 (cli/claude-sonnet, codestral-latest)	
Severity: ATTENTION	Evidence: Appendix A · Appendix C
Submitted a change; the independent hidden tests still failed.	
F-02 · derived from the recorded run data	
F-03	LANE / INFRA
model_api_failure × 1 (codestral-latest)	
Severity: LANE / INFRA	Evidence: Appendix A · Appendix C
The model endpoint failed to answer. Not a verdict about the model's ability. These rollouts are excluded from any capability reading.	
F-03 · derived from the recorded run data	

Recommendations

One recommendation per observed failure mode, addressed to the lane it affects. Each traces back to a finding above and to the raw trace in Appendix A.

R-01 PROBLEM
Submitted a change; the independent hidden tests still failed.

AFFECTED LANES
cli/claude-sonnet, codestral-latest

WHAT WE WOULD LOOK AT FIRST
The model completes the loop and produces a plausible near-miss — the named failing tests are the exact behavioral cases distinguishing it from the reference fix.

R-02 PROBLEM
The model endpoint failed to answer. Not a verdict about the model's ability.

AFFECTED LANES
codestral-latest

WHAT WE WOULD LOOK AT FIRST
The lane failed, not the model; these rollouts carry no signal about capability and are labeled accordingly.

Verification

RUN ID

live-20260901T235346Z

EXAM FINGERPRINT

4880a527ef97ca2f3fe6b57be026c3dc1be51e9e104d79fc13a3e558331dda92

The contract digest this run was executed against.

Provenance of this chapter

This chapter was rendered by the campaign generator from the sealed evaluation record; it is not the standalone report reproduced byte for byte. Verify the standalone report with its own digest, verify the record with `vvdex-env records verify`, and verify this campaign document as a whole.

RELATIONSHIP

campaign-rendered chapter

Derived from the same sealed evaluation record as the standalone report; these bytes are authenticated by the campaign artifact that contains them.

SEALED RECORD DIGEST

157ec324a925c0b10fc8769d01f628b6143c40a00fa2dc70ea309a70e55a37a8

sha256 over the record's canonical JSON — the identity every renderer works from.

SOURCE RECORD FILE

a3e742fe73fe3667155e22dfc941bdc9e428a4bb1bd47c5c89a0454e299e7eb0

sha256 of `fr-20260902-2e7f4171.json` as written beside the standalone report.

SOURCE STANDALONE REPORT

564618c51dddcf7f05d9eb4f806c4ee4191a2e99206b3d952e1664cbef81b3da

sha256 of `fr-20260902-2e7f4171.html` — independently verifiable with its own digest.

SOURCE STANDALONE PDF

30de973cceb4b6cdc6799736cc493c32a4b221156136388d27eade5fbd86707c

sha256 of `fr-20260902-2e7f4171.pdf`.

RUN EVIDENCE BUNDLE

2e7f4171f049cced92f7a3aa7ae13b2ee211449e96938e7e6c437f30289db0ff

The digest the run's own bundle computed over itself.

CERTIFIED EVIDENCE SEAL

not available in this package

The exam's sealed evidence, established before this run.

Measurement history

This is the first recorded run on this exam. There is nothing to compare it against yet, and a single measurement is not a trend.

The full artifact-by-artifact digest table, and the reproduction protocol, are in Appendix B.

Appendix A — every rollout

ROLLOUT	MODEL	SEED	OUTCOME	SUBMITTED	DEEPEST STAGE	HIDDEN TESTS	FILES	TOKENS	ELAPSED
8ec5820d	cli/claude-sonnet	0	submitted_failed	yes	Graded	fail	1	n/a	25.2s
72bcf971	cli/claude-sonnet	1	submitted_passed	yes	Passed hidden tests	pass	1	n/a	51.1s
62bceaaa	cli/claude-sonnet	2	submitted_failed	yes	Graded	fail	1	n/a	52.1s
a66f85fd	cli/claude-sonnet	3	submitted_failed	yes	Graded	fail	0	n/a	36.1s
3fcba938	cli/claude-sonnet	4	submitted_failed	yes	Graded	fail	1	n/a	60.0s
70cf5fcc	cli/claude-sonnet	5	submitted_failed	yes	Graded	fail	1	n/a	56.3s
d13f1d97	cli/claude-sonnet	6	submitted_failed	yes	Graded	fail	1	n/a	53.8s
9d47ee0f	cli/claude-sonnet	7	submitted_failed	yes	Graded	fail	1	n/a	24.5s
a8cca2fc	cli/claude-sonnet	8	submitted_failed	yes	Graded	fail	0	n/a	43.8s
153192c8	cli/claude-sonnet	9	submitted_failed	yes	Graded	fail	1	n/a	55.6s
781eb674	cli/codex	0	submitted_passed	yes	Passed hidden tests	pass	1	n/a	76.9s
e71d8908	cli/codex	1	submitted_passed	yes	Passed hidden tests	pass	1	n/a	66.0s
2158c4f5	cli/codex	2	submitted_passed	yes	Passed hidden tests	pass	1	n/a	69.7s
8b6e987c	cli/codex	3	submitted_passed	yes	Passed hidden tests	pass	1	n/a	68.8s
fe164929	cli/codex	4	submitted_passed	yes	Passed hidden tests	pass	1	n/a	68.2s
854f7dae	cli/codex	5	submitted_passed	yes	Passed hidden tests	pass	1	n/a	65.6s
18d03c3a	cli/codex	6	submitted_passed	yes	Passed hidden tests	pass	1	n/a	65.9s
854f8510	cli/codex	7	submitted_passed	yes	Passed hidden tests	pass	1	n/a	66.5s
a111a1ee	cli/codex	8	submitted_passed	yes	Passed hidden tests	pass	1	n/a	65.7s
afdb6ff0	cli/codex	9	submitted_passed	yes	Passed hidden tests	pass	1	n/a	65.2s
ab1db90e	codestral-latest	0	submitted_failed	yes	Graded	fail	1	28 352	18.5s
6ed36c95	codestral-latest	1	submitted_failed	yes	Graded	fail	1	21 309	13.9s
496ff0a6	codestral-latest	2	submitted_failed	yes	Graded	fail	1	24 495	16.5s
e411d189	codestral-latest	3	submitted_failed	yes	Graded	fail	1	36 382	24.3s
acd6942b	codestral-latest	4	model_api_failure	no	Not graded — lane error	not graded	0	9 107	9.8s
ee5a8bdb	codestral-latest	5	submitted_failed	yes	Graded	fail	1	23 906	18.5s
3897e57e	codestral-latest	6	submitted_failed	yes	Graded	fail	1	18 536	14.1s
f23935c8	codestral-latest	7	submitted_failed	yes	Graded	fail	1	21 309	14.2s
9f5bdfb8	codestral-latest	8	submitted_failed	yes	Graded	fail	1	32 417	30.0s
f96b96b6	codestral-latest	9	submitted_failed	yes	Graded	fail	1	18 624	12.2s

Failure taxonomy

CLASS	COUNT	WHAT IT MEANS
submitted_failed	18	Submitted a change; the independent hidden tests still failed.
model_api_failure	1	The model endpoint failed to answer. Not a verdict about the model's ability.

Appendix B — evidence and artifact digests

RUN BUNDLE DIGEST 2e7f4171f049cced92f7a3aa 7ae13b2ee211449e96938e7e 6c437f30289db0ff	RUN ID live- 20260901T235346Z	STARTED 2026-09- 01T23:53:46.828443 +00:00	FINISHED 2026-09- 02T00:18:47.040828 +00:00
--	--	--	---

Artifact digests

ARTIFACT	SHA-256
publicSummary	e4c4792d488475e66d9750cfca4484181c78985538f65d32605e587faad1cb4f
rollouts/153192c8-0932-4da9-bf81-111446fc32a5.json	6e633b09ebe3ee95b5cb4d387fe1356b3a02781bbd741f81e78a85a0e73728af
rollouts/18d03c3a-6629-4250-886e-c73986fdece7.json	ca658ebbd354c8f20e9cfba884d8e3098d4bbc990dfac08d14fc85ed4201f660
rollouts/2158c4f5-9a18-45cc-82e5-aa84a17c46a2.json	963bf7642556ae0d1a961a485be5acb2542e4942e1e79656bacc3463cc5d5c1e
rollouts/3897e57e-353e-402d-bb53-bc68259b1aa2.json	011530ec4616efa6f0b86d2b5a42b5aea82680b86f689f8af69217b80839ade8
rollouts/3fcb938-b19c-41d4-84fe-6ab6fec3c0a5.json	d75c04727307bfd1f50555fc0f892c36a01308a740eeef7706305cb59ce371f
rollouts/496ff0a6-9f53-4b50-8c9e-d1939301419c.json	c0ff5ca7fc1649be66606e9be02f4c8bfbec36583293f0e9aefcbd6c37c2d86c
rollouts/62bceaaa-2bf1-4067-8bef-6393b41f4d42.json	97da73eca72708cf8661d79aa0ac5351166de7f8599395f5253fdccf5627ec2b
rollouts/6ed36c95-91f9-4e3c-88da-fe17c010d584.json	5f72442b16936bf2145ef717e960d09429d418db43ad8c08de1c02dec6fec762
rollouts/70cf5fcc-7581-48bd-a1a9-53b4ce24a3a9.json	b0a5bfffdf0754ff0609c73cc67e5a019e62e0fccf0e1c4630634818ba605ceb3
rollouts/72bcf971-7d06-42c0-9d6f-7f4ab5380c76.json	c95578b9f69a937f89a18d2cf3e70049a371791d10a614a6c87b7ca3ef918796
rollouts/781eb674-aa4a-4b2f-b1bc-502f4710a4fe.json	dbeff75b715e052925a9c65e68a9ea73059a54245df5c87809b82f95aaaa9ff
rollouts/854f7dae-d4bc-441f-be3f-fc48eb324c71.json	77ee126e557bb11b7e88c38abae6bfcd9432c27cc8e9684ae49170cc69acfa1
rollouts/854f8510-cc34-4a46-a50f-771c323e3025.json	630b147f5a7c8a9a8ddc12bb638d1f92d8d54cae19ba7ec3b49a42c479a79045
rollouts/8b6e987c-8ccd-42a0-a076-30800bf7504b.json	51a15b141bcf2a6c0d377058ed74127fb4b2d1b29389638b1ef785e2919506fa
rollouts/8ec5820d-ef88-42c9-9c7f-e5cf9923de8d.json	9ab5fbbd4ef13e7c801e7f8c99343bf26ee7bb9d3f6f38cc077863b7386742b7
rollouts/9d47ee0f-e8c3-4eb8-a5cd-ba455bfb80a3.json	21e2ecff0304fbd74fa806d478ce3e2e4140a6cf2054fc03ea99a9ab3ab96cec
rollouts/9f5bdfb8-69a0-4b45-9c77-2f87b5217ca7.json	b298cff46d8ab12505f0510a561f6ca4a2a58c6dc5ce2af3ab83586acf0728c
rollouts/a111a1ee-81a7-4118-827f-2f8cf436e16c.json	fc831c3cbaaf9dba9f5100de58cbc210c43dd42997d6dfebc088dd4ff75aff5c
rollouts/a66f85fd-827a-4d31-b407-6b0983c085ab.json	204a94e5e1d59d5faa5433a5b740546258c664c75b0e5b6490326b817bcb8aa5
rollouts/a8cca2fc-949e-4473-a79f-ac3027cab760.json	7775ac3ca099f5182d53326959308b3cd425565d5dfb7215e73f014bbe546537
rollouts/ab1db90e-5252-4819-b24b-0770692d26a9.json	d6bbea07e8570149b9f65ea3844ea0eeefc7dc5acc07253e271a2aed94b7628ed
rollouts/acd6942b-ff2c-478b-8fc5-36473631c0db.json	cdee6f5f35965fda4a7cda3d1eaab8aa1e6e22a13c0912a2afdd5811b35cc2b6
rollouts/afdb6ff0-9c18-494d-874a-43fce0ae7de7.json	dc31e5124a7738f55023ce1dea643d88d14cbfc18a1165e9003de6447be90f12
rollouts/d13f1d97-627d-4353-b467-a02faeff3156.json	b3e529077b975790be00a38c23a8f708f019c526165fb0ef68213cd2176a8731
rollouts/e411d189-390c-40ef-9318-3fbbbfb44d1.json	9fe97e4b4806ef362789ff8206bc560946bfc400b1d3046d5789534beb7c94ca
rollouts/e71d8908-6f2a-4362-9b6f-bdaf29b367f9.json	29ad3560920eadc86473f2efefdb0ee2e60c45ab030d4f68b6d95e282a8be013
rollouts/ee5a8bdb-0fe4-4909-a6f1-071b188dee55.json	d3ef142e69e1f9d296d05046099a73cb442e1df77c73463ac027717e0a749102
rollouts/f23935c8-c08a-442e-ba1d-ed39ea664b9d.json	9f3476ef0cbb5bca9e19b0e738ed19bfecb0238ef784503aa1cf0a46b9675c2f

ARTIFACT	SHA-256
rollouts/f96b96b6-e608-44d7-9f4b-8506d41d456d.json	c862c8429a774f5c5abab4442e2c2023466b8dd9eab98c117e8a1672e0209a8c
rollouts/fe164929-6cd7-4f59-acbc-ccb6ad02ddb2.json	97fc6c8b9bc91843bfa5958ffb79146d35a7c921e2f464093a7ff5736c1c6097
spec	0a225ad1c6f73d8104d73d1b5f0717e8516fc4c57f1957fc16447f8551663daa
summary	b77bc1cfc26a8c6e1848f0a23e0379695a67cf5a2db2f0543b8d6ade3a1bd1c1

Reproducing this run

Every rollout can be re-derived: the exam is frozen (fingerprint above), the grader ran in the recorded container image, and each rollout's full trace — every tool call, every result, every word — is in the evidence bundle at its rollouts/<id>.json path. A re-run uses the same job shape: the exam id, your endpoint, n, and the runtime. The record id names this run forever.

Appendix C — redacted transcripts

For each rollout: the tool calls it made, in order — which tool, on which of the model's own paths, whether the call succeeded, and how many characters came back. What a tool returned is not here, and neither is the submission: on this exam the submitted content is the graded answer, so no diff, no answer value and no model claim is printed. A request for the reference solution or the hidden tests is recorded as a request; the path it asked for is replaced by a label. See Appendix D, "What this report withholds and why".

Disclosure class applied to this exam: answer-key disclosure (ownership vvdex_proprietary — engine registry; family rag_grounding): the exam is VVDex evaluation IP and the graded output is an answer or a keyed value, so no submission content, diff, model claim or hidden-assertion name is published.

8ec5820d · cli:cli/claude-sonnet · seed 0 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	list_files	.	ok	287
1	read_file	questions.json	ok	883
2	run_command	—	ok	1 780
3	write_file	answers/answers.json	ok	97
4	run_tests	—	ok	181
5	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (57 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

72bcf971 · cli:cli/claude-sonnet · seed 1 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	list_files	.	ok	287
1	read_file	questions.json	ok	883
2	read_file	corpus/data-protection-policy-v3.txt	ok	293
3	read_file	corpus/retention-addendum-2026.txt	ok	142
4	read_file	corpus/org-chart-2026.txt	ok	183
5	read_file	corpus/incident-report-2026.txt	ok	169
6	read_file	corpus/security-whitepaper.txt	ok	142
7	read_file	corpus/employee-handbook.txt	ok	167
8	read_file	corpus/vendor-list.txt	ok	153
9	write_file	answers/answers.json	ok	69
10	run_tests	—	ok	181
11	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (45 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

62bceaaa · cli:cli/claude-sonnet · seed 2 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	list_files	.	ok	287
1	read_file	questions.json	ok	883
2	read_file	corpus/data-protection-policy-v3.txt	ok	293
3	read_file	corpus/retention-addendum-2026.txt	ok	142
4	read_file	corpus/org-chart-2026.txt	ok	183
5	read_file	corpus/incident-report-2026.txt	ok	169
6	read_file	corpus/security-whitepaper.txt	ok	142
7	read_file	corpus/employee-handbook.txt	ok	167
8	read_file	corpus/vendor-list.txt	ok	153
9	write_file	answers/answers.json	ok	69
10	run_tests	—	ok	181
11	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (45 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

a66f85fd · cli:cli/claude-sonnet · seed 3 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	list_files	.	ok	287
1	read_file	questions.json	ok	883
2	run_command	—	ok	132
3	run_command	—	ok	1 399
4	run_command	—	ok	758
5	invalid_tool_call	(none)	failed	120
6	run_tests	—	ok	181
7	submit	—	ok	85

Submission diff withheld by the public-disclosure policy. Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

3fcba938 · cli:cli/claude-sonnet · seed 4 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	list_files	.	ok	287
1	read_file	questions.json	ok	883
2	read_file	corpus/org-chart-2026.txt	ok	183
3	read_file	corpus/data-protection-policy-v3.txt	ok	293

STEP	TOOL	PATH	RESULT	CHARS
4	read_file	corpus/retention-addendum-2026.txt	ok	142
5	read_file	corpus/incident-report-2026.txt	ok	169
6	read_file	corpus/security-whitepaper.txt	ok	142
7	read_file	corpus/employee-handbook.txt	ok	167
8	read_file	corpus/vendor-list.txt	ok	153
9	write_file	answers/answers.json	ok	69
10	run_tests	—	ok	181
11	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (45 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

70cf5fcc · cli:cli/claude-sonnet · seed 5 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	list_files	.	ok	287
1	read_file	questions.json	ok	883
2	read_file	corpus/org-chart-2026.txt	ok	183
3	read_file	corpus/data-protection-policy-v3.txt	ok	293
4	read_file	corpus/retention-addendum-2026.txt	ok	142
5	read_file	corpus/incident-report-2026.txt	ok	169
6	read_file	corpus/security-whitepaper.txt	ok	142
7	read_file	corpus/employee-handbook.txt	ok	167
8	write_file	answers/answers.json	ok	69
9	run_tests	—	ok	181
10	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (45 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

d13f1d97 · cli:cli/claude-sonnet · seed 6 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	list_files	.	ok	287
1	read_file	questions.json	ok	883
2	read_file	corpus/data-protection-policy-v3.txt	ok	293
3	read_file	corpus/retention-addendum-2026.txt	ok	142
4	read_file	corpus/org-chart-2026.txt	ok	183
5	read_file	corpus/incident-report-2026.txt	ok	169
6	read_file	corpus/security-whitepaper.txt	ok	142

STEP	TOOL	PATH	RESULT	CHARS
7	read_file	corpus/employee-handbook.txt	ok	167
8	read_file	corpus/vendor-list.txt	ok	153
9	write_file	answers/answers.json	ok	69
10	run_tests	—	ok	181
11	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (4 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

9d47ee0f · cli:cli/claude-sonnet · seed 7 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	list_files	.	ok	287
1	read_file	questions.json	ok	883
2	run_command	—	ok	1 780
3	write_file	answers/answers.json	ok	69
4	run_tests	—	ok	181
5	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (4 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

a8cca2fc · cli:cli/claude-sonnet · seed 8 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	list_files	.	ok	287
1	read_file	questions.json	ok	883
2	run_command	—	ok	1 399
3	run_command	—	ok	758
4	run_command	—	ok	414
5	invalid_tool_call	(none)	failed	120
6	run_tests	—	ok	181
7	submit	—	ok	85

Submission diff withheld by the public-disclosure policy. Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

153192c8 · cli:cli/claude-sonnet · seed 9 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	list_files	.	ok	287

STEP	TOOL	PATH	RESULT	CHARS
1	read_file	questions.json	ok	883
2	read_file	corpus/data-protection-policy-v3.txt	ok	293
3	read_file	corpus/org-chart-2026.txt	ok	183
4	read_file	corpus/retention-addendum-2026.txt	ok	142
5	read_file	corpus/incident-report-2026.txt	ok	169
6	read_file	corpus/security-whitepaper.txt	ok	142
7	read_file	corpus/employee-handbook.txt	ok	167
8	read_file	corpus/vendor-list.txt	ok	153
9	write_file	answers/answers.json	ok	69
10	run_tests	—	ok	181
11	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (4 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

781eb674 · cli:cli/codex · seed 0 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	ISSUE.md	ok	1 953
1	list_files	.	ok	287
2	read_file	questions.json	ok	883
3	read_file	corpus/data-protection-policy-v3.txt	ok	293
4	read_file	corpus/data-protection-policy-v1.txt	ok	239
5	read_file	corpus/employee-handbook.txt	ok	167
6	read_file	corpus/incident-report-2026.txt	ok	169
7	read_file	corpus/org-chart-2026.txt	ok	183
8	read_file	corpus/retention-addendum-2026.txt	ok	142
9	read_file	corpus/security-whitepaper.txt	ok	142
10	read_file	corpus/vendor-list.txt	ok	153
11	write_file	answers/answers.json	ok	69
12	run_tests	—	ok	181
13	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (45 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

e71d8908 · cli:cli/codex · seed 1 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	list_files	.	ok	287

STEP	TOOL	PATH	RESULT	CHARS
1	read_file	ISSUE.md	ok	1 953
2	read_file	questions.json	ok	883
3	read_file	corpus/data-protection-policy-v3.txt	ok	293
4	read_file	corpus/data-protection-policy-v1.txt	ok	239
5	read_file	corpus/employee-handbook.txt	ok	167
6	read_file	corpus/incident-report-2026.txt	ok	169
7	read_file	corpus/org-chart-2026.txt	ok	183
8	read_file	corpus/retention-addendum-2026.txt	ok	142
9	read_file	corpus/security-whitepaper.txt	ok	142
10	read_file	corpus/vendor-list.txt	ok	153
11	write_file	answers/answers.json	ok	69
12	run_tests	—	ok	181
13	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (13 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

2158c4f5 · cli:cli/codex · seed 2 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	ISSUE.md	ok	1 953
1	list_files	.	ok	287
2	read_file	questions.json	ok	883
3	read_file	corpus/data-protection-policy-v1.txt	ok	239
4	read_file	corpus/data-protection-policy-v3.txt	ok	293
5	read_file	corpus/employee-handbook.txt	ok	167
6	read_file	corpus/incident-report-2026.txt	ok	169
7	read_file	corpus/org-chart-2026.txt	ok	183
8	read_file	corpus/retention-addendum-2026.txt	ok	142
9	read_file	corpus/security-whitepaper.txt	ok	142
10	read_file	corpus/vendor-list.txt	ok	153
11	write_file	answers/answers.json	ok	69
12	run_tests	—	ok	181
13	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (45 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

8b6e987c · cli:cli/codex · seed 3 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	ISSUE.md	ok	1 953
1	list_files	.	ok	287
2	read_file	questions.json	ok	883
3	read_file	corpus/data-protection-policy-v1.txt	ok	239
4	read_file	corpus/data-protection-policy-v3.txt	ok	293
5	read_file	corpus/employee-handbook.txt	ok	167
6	read_file	corpus/incident-report-2026.txt	ok	169
7	read_file	corpus/org-chart-2026.txt	ok	183
8	read_file	corpus/retention-addendum-2026.txt	ok	142
9	read_file	corpus/security-whitepaper.txt	ok	142
10	read_file	corpus/vendor-list.txt	ok	153
11	write_file	answers/answers.json	ok	69
12	run_tests	—	ok	181
13	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (45 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

fe164929 · cli:cli/codex · seed 4 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	ISSUE.md	ok	1 953
1	list_files	.	ok	287
2	read_file	questions.json	ok	883
3	read_file	corpus/data-protection-policy-v1.txt	ok	239
4	read_file	corpus/data-protection-policy-v3.txt	ok	293
5	read_file	corpus/employee-handbook.txt	ok	167
6	read_file	corpus/incident-report-2026.txt	ok	169
7	read_file	corpus/org-chart-2026.txt	ok	183
8	read_file	corpus/retention-addendum-2026.txt	ok	142
9	read_file	corpus/security-whitepaper.txt	ok	142
10	read_file	corpus/vendor-list.txt	ok	153
11	write_file	answers/answers.json	ok	69
12	run_tests	—	ok	181
13	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (45 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

854f7dae · cli:cli/codex · seed 5 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	ISSUE.md	ok	1 953
1	list_files	.	ok	287
2	read_file	questions.json	ok	883
3	read_file	corpus/data-protection-policy-v1.txt	ok	239
4	read_file	corpus/data-protection-policy-v3.txt	ok	293
5	read_file	corpus/employee-handbook.txt	ok	167
6	read_file	corpus/incident-report-2026.txt	ok	169
7	read_file	corpus/org-chart-2026.txt	ok	183
8	read_file	corpus/retention-addendum-2026.txt	ok	142
9	read_file	corpus/security-whitepaper.txt	ok	142
10	read_file	corpus/vendor-list.txt	ok	153
11	write_file	answers/answers.json	ok	69
12	run_tests	—	ok	181
13	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (13 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

18d03c3a · cli:cli/codex · seed 6 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	list_files	.	ok	287
1	read_file	ISSUE.md	ok	1 953
2	read_file	questions.json	ok	883
3	read_file	corpus/data-protection-policy-v1.txt	ok	239
4	read_file	corpus/data-protection-policy-v3.txt	ok	293
5	read_file	corpus/employee-handbook.txt	ok	167
6	read_file	corpus/incident-report-2026.txt	ok	169
7	read_file	corpus/org-chart-2026.txt	ok	183
8	read_file	corpus/retention-addendum-2026.txt	ok	142
9	read_file	corpus/security-whitepaper.txt	ok	142
10	read_file	corpus/vendor-list.txt	ok	153
11	write_file	answers/answers.json	ok	69
12	run_tests	—	ok	181

STEP	TOOL	PATH	RESULT	CHARS
13	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (45 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

854f8510 · cli:cli/codex · seed 7 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	ISSUE.md	ok	1 953
1	list_files	.	ok	287
2	read_file	questions.json	ok	883
3	read_file	corpus/data-protection-policy-v1.txt	ok	239
4	read_file	corpus/data-protection-policy-v3.txt	ok	293
5	read_file	corpus/employee-handbook.txt	ok	167
6	read_file	corpus/incident-report-2026.txt	ok	169
7	read_file	corpus/org-chart-2026.txt	ok	183
8	read_file	corpus/retention-addendum-2026.txt	ok	142
9	read_file	corpus/security-whitepaper.txt	ok	142
10	read_file	corpus/vendor-list.txt	ok	153
11	write_file	answers/answers.json	ok	69
12	run_tests	—	ok	181
13	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (45 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

a111a1ee · cli:cli/codex · seed 8 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	ISSUE.md	ok	1 953
1	list_files	.	ok	287
2	read_file	questions.json	ok	883
3	read_file	corpus/data-protection-policy-v1.txt	ok	239
4	read_file	corpus/data-protection-policy-v3.txt	ok	293
5	read_file	corpus/employee-handbook.txt	ok	167
6	read_file	corpus/incident-report-2026.txt	ok	169
7	read_file	corpus/org-chart-2026.txt	ok	183
8	read_file	corpus/retention-addendum-2026.txt	ok	142
9	read_file	corpus/security-whitepaper.txt	ok	142
10	read_file	corpus/vendor-list.txt	ok	153

STEP	TOOL	PATH	RESULT	CHARS
11	write_file	answers/answers.json	ok	69
12	run_tests	—	ok	181
13	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (45 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

afdb6ff0 · cli:cli/codex · seed 9 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	ISSUE.md	ok	1 953
1	list_files	.	ok	287
2	read_file	questions.json	ok	883
3	read_file	corpus/data-protection-policy-v1.txt	ok	239
4	read_file	corpus/data-protection-policy-v3.txt	ok	293
5	read_file	corpus/employee-handbook.txt	ok	167
6	read_file	corpus/incident-report-2026.txt	ok	169
7	read_file	corpus/org-chart-2026.txt	ok	183
8	read_file	corpus/retention-addendum-2026.txt	ok	142
9	read_file	corpus/security-whitepaper.txt	ok	142
10	read_file	corpus/vendor-list.txt	ok	153
11	write_file	answers/answers.json	ok	69
12	run_tests	—	ok	181
13	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (45 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

ab1db90e · mistral:codestral-latest · seed 0 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	ISSUE.md	ok	1 953
1	read_file	questions.json	ok	883
2	list_files	corpus	ok	249
3	read_file	corpus/data-protection-policy-v3.txt	ok	293
4	read_file	corpus/org-chart-2026.txt	ok	183
5	run_tests	—	ok	181
6	invalid_tool_call	(none)	failed	120
7	write_file	answers/answers.json	ok	69
8	run_tests	—	ok	181

STEP	TOOL	PATH	RESULT	CHARS
9	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (45 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

6ed36c95 · mistral:codestral-latest · seed 1 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	ISSUE.md	ok	1 953
1	read_file	questions.json	ok	883
2	list_files	answers	ok	13
3	list_files	corpus	ok	249
4	read_file	corpus/data-protection-policy-v3.txt	ok	293
5	write_file	answers/answers.json	ok	69
6	run_tests	—	ok	181
7	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (45 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

496ff0a6 · mistral:codestral-latest · seed 2 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	ISSUE.md	ok	1 953
1	read_file	questions.json	ok	883
2	read_file	—	failed	35
3	list_files	corpus	ok	249
4	read_file	corpus/retention-addendum-2026.txt	ok	142
5	read_file	corpus/data-protection-policy-v3.txt	ok	293
6	write_file	answers/answers.json	ok	69
7	run_tests	—	ok	181
8	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (45 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

e411d189 · mistral:codestral-latest · seed 3 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	ISSUE.md	ok	1 953
1	list_files	answers	ok	13

STEP	TOOL	PATH	RESULT	CHARS
2	read_file	answers/.keep	ok	0
3	list_files	corpus	ok	249
4	read_file	corpus/data-protection-policy-v1.txt	ok	239
5	read_file	corpus/data-protection-policy-v3.txt	ok	293
6	read_file	corpus/employee-handbook.txt	ok	167
7	read_file	questions.json	ok	883
8	run_tests	–	ok	181
9	invalid_tool_call	(none)	failed	120
10	write_file	answers/answers.json	ok	69
11	run_tests	–	ok	181
12	submit	–	ok	85

Submission diff withheld by the public-disclosure policy (45 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

acd6942b · mistral:codestral-latest · seed 4 · model_api_failure

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	ISSUE.md	ok	1 953
1	read_file	questions.json	ok	883
2	list_files	corpus	ok	249
3	invalid_tool_call	(none)	failed	120

Submission diff withheld by the public-disclosure policy. Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

ee5a8bdb · mistral:codestral-latest · seed 5 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	ISSUE.md	ok	1 953
1	read_file	questions.json	ok	883
2	list_files	answers	ok	13
3	read_file	answers/.keep	ok	0
4	list_files	corpus	ok	249
5	read_file	corpus/retention-addendum-2026.txt	ok	142
6	write_file	answers/answers.json	ok	69
7	run_tests	–	ok	181
8	submit	–	ok	85

Submission diff withheld by the public-disclosure policy (45 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

3897e57e · mistral:codestral-latest · seed 6 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	ISSUE.md	ok	1 953
1	read_file	questions.json	ok	883
2	list_files	corpus	ok	249
3	read_file	corpus/retention-addendum-2026.txt	ok	142
4	write_file	answers/answers.json	ok	69
5	run_tests	—	ok	181
6	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (45 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

f23935c8 · mistral:codestral-latest · seed 7 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	ISSUE.md	ok	1 953
1	read_file	questions.json	ok	883
2	list_files	answers	ok	13
3	list_files	corpus	ok	249
4	read_file	corpus/data-protection-policy-v3.txt	ok	293
5	write_file	answers/answers.json	ok	69
6	run_tests	—	ok	181
7	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (45 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

9f5bdfb8 · mistral:codestral-latest · seed 8 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	ISSUE.md	ok	1 953
1	list_files	answers	ok	13
2	read_file	answers/.keep	ok	0
3	list_files	corpus	ok	249
4	read_file	corpus/data-protection-policy-v1.txt	ok	239
5	read_file	corpus/data-protection-policy-v3.txt	ok	293

STEP	TOOL	PATH	RESULT	CHARS
6	write_file	answers/answers.json	ok	69
7	run_tests	—	ok	181
8	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (15 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

f96b96b6 · mistral:codestral-latest · seed 9 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	ISSUE.md	ok	1 953
1	read_file	questions.json	ok	883
2	list_files	corpus	ok	249
3	read_file	corpus/data-protection-policy-v3.txt	ok	293
4	write_file	answers/answers.json	ok	69
5	run_tests	—	ok	181
6	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (45 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

13 Exam · memory-fact-update-1

vvdex.knowledge.memory-fact-update-1 · record fr-20260901-46eb9548 · record digest 1127e8356d42d317... · harness clean · containment clean

Executive summary

MODELS TESTED 3	CERTIFIED PASSES 22 of 30 graded rollouts	SUBMITTED 30 of 30 graded rollouts that called submit	NOT GRADED — LANE ERRORS 0 of 30 attempts
---------------------------------------	--	--	--

Purpose

This report presents a controlled evaluation of 3 model lane(s) against the certified exam vvdex.knowledge.memory-fact-update-1 (v0.1.0, family long_term_memory). Every lane received the identical frozen workspace, tool contract and limits; grading ran afterwards, in isolation, against hidden tests no lane ever saw. The purpose is to state, with reproducible evidence, what each lane did and where it stopped.

Outcome

22 of 30

graded rollouts passed the independent hidden tests and were certified

22 of 30 graded rollouts passed. cli/claude-sonnet, cli/codex, codestral-latest completed the full loop (inspect, edit, run the visible tests, submit) and the sealed grader accepted the submission. The most common way to fail was submitted_failed — submitted a change; the independent hidden tests still failed.

Key findings

F-01 cli/claude-sonnet, cli/codex, codestral-latest passed the independent hidden tests

F-02 submitted_failed × 8 (codestral-latest)

Each finding is stated in full, with evidence, in the Findings section; recommendations for acting on them follow it. Elapsed wall clock for the whole engagement: 695.7s.

Exam definition

In scope. The ability of each configured model lane to complete this one certified task end to end — inspect the workspace, produce a correct edit, verify it with the visible tests, and submit — within 20 tool steps and 4920s of wall clock, at 10 rollout(s) per lane. Behavioural measurements along the way (tool discipline, grounding, overclaim) are in scope because they are read from the recorded trace, not judged by another model.

Out of scope. General model quality, performance on other task families, prompt-tuning on the lane's behalf, and any comparison this run's sample size cannot support. A lane's API failure is reported as infrastructure, never as model capability.

The instrument. Exam `vvdex.knowledge.memory-fact-update-1` was certified before this run: its reference solution passes, an empty submission fails, its grader distinguishes near-misses, and its sealed evidence is unchanged by this evaluation (see Certification evidence).

What was measured

EXAM	VERSION	FAMILY	ROLLOUTS
vvdex.knowledge.memory-fact-update-1	0.1.0	long_term_memory	30 (10 per model)

Where the defect came from

This exam has no upstream repository to credit: the defect, the reference solution and the hidden suite are VVDex's own.

Certification evidence

4 of 5 certification pillars are recorded in this package, loaded from the receipts sealed into the exam and matching this run's exam fingerprint. What is recorded is shown with its real value; what is not is named below and is not claimed.

FROZEN BASELINE

FAIL

expected — an empty submission over 2 run(s) must not pass

→

GOLD / REFERENCE

PASS

expected — the reference solution over 2 run(s) must pass

GRADER CONTROLS

7 / 7

The grader accepted the reference and rejected every near-miss.

ATTACK PROBES

10 / 10 blocked

0 trivial exploit(s) found.

RUNTIME

sha256:679b36225a1f83238efba8c71b9cde3c24caaeacc9b682ae92819cb55eec328f

network policy: deny

Certification metadata is not included in this report package for:

- sealed evidence bundle

Benchmark results remain valid as run evidence; the pillar(s) above are not substantiated by this document and this report does not claim them.

What the package does carry: the exam fingerprint below is the contract digest this run was executed against — a holder of the certification bundle for this fingerprint can join the two independently.

EXAM FINGERPRINT

2052eec0ff2d3044dc269e3599ae1bfaa45d0d8cd5df9554796105538d226c21

The contract digest this run was executed against.

Evaluation configuration

REPORT	fr-20260901-46eb9548
ISSUED	2026-09-01
SUBJECT EXAM	vvdex.knowledge.memory-fact-update-1 · v0.1.0
FAMILY	long_term_memory
PREPARED BY	VVDex Forge engine vvdex-env 0.1.0
CLASSIFICATION	Independent model evaluation report — independently verifiable
CERTIFICATION	Exam certification: partially recorded (see Certification evidence)
STATUS	FINAL · report sealed against its own digest

Timeline

WHEN (UTC)	EVENT
2026-09-01 23:42:10	Evaluation run started
2026-09-01 23:53:46	Evaluation run finished
2026-09-01 23:53:46	Report generated from the sealed run evidence

Models under test

MODEL	PROVIDER	ENDPOINT	BILLING
Claude Code CLI (Sonnet, subscription)	cli	local runtime	operator subscription, flat — not billed per token
Codex CLI (GPT, subscription)	cli	local runtime	operator subscription, flat — not billed per token
codestral-latest	mistral	VVDex-configured	VVDex free-quota

A customer endpoint is recorded by origin only — never its port, its path, or its key.

Limits

RUNTIME docker	STEP LIMIT 20	WALL-CLOCK LIMIT 4920s	MAX OUTPUT TOKENS 2048
TEMPERATURE 0.2	RETRY POLICY none	COMMAND BUDGET 24	TOOLS OFFERED list_files, read_file, write_file, edit_file, run_tests, submit, run_command

The evaluated model never saw the hidden tests or the reference solution, and could not change its result. Grading ran in a separate step, after submission, on a container with no network.

Model outcomes

Denominators. Attempts requested: 30. Graded (evaluable) rollouts: 30. Not graded — lane errors: 0. A lane error is a rollout the API or the runtime ended before the hidden grader ran: it is listed, excluded from every rate and pass@k, and never silently dropped. Passes are certified passes: hidden grader exit 0 with integrity verified.

MODEL	ATTEMPTS	GRADED	PASSED	MODEL FAILS	LANE ERRORS	SUBMITTED	MEAN TIME	TYPICAL DEEPEST STAGE
Claude Code CLI (Sonnet, subscription) cli/claude-sonnet	10	10	10 / 10	0	0	10 / 10	17.0s	Passed hidden tests
Codex CLI (GPT, subscription) cli/codex	10	10	10 / 10	0	0	10 / 10	26.3s	Passed hidden tests
codestral-latest codestral-latest	10	10	2 / 10	8	0	10 / 10	7.4s	Graded

Every rollout, with its own deepest stage and grade, is in Appendix A.

Success rate, confidence and pass@k

Every rate on this page is a measurement of a finite number of rollouts, so every rate carries the interval that measurement actually supports. n is the evaluable count — attempts minus rollouts the lane ended before grading; those are listed under "not graded" and sit in no rate, interval or pass@k.

MODEL	ATTEMPTS	N	NOT GRADED	PASSES	RATE	95% INTERVAL	SEEDS	PASS@1	PASS@3	PASS@5	PASS@10
cli/claude-sonnet	10	10	0	10	100.0%	[72%, 100%]	0–9	100.0%	100.0%	100.0%	100.0%
cli/codex	10	10	0	10	100.0%	[72%, 100%]	0–9	100.0%	100.0%	100.0%	100.0%
codestral-latest	10	10	0	2	20.0%	[6%, 51%]	0–9	20.0%	53.3%	77.8%	100.0%

Interval and pass@k methods are stated in full in Appendix D. pass@k is shown for k = 1, 3, 5, 10; the sealed record carries every k up to n.

Model comparison

MODEL	PASS@1	95% INTERVAL	TOKENS	COST	MEAN ROLLOUT	OVERCLAIM RATE	TOP FAILURE CLASS
cli/claude-sonnet	100.0%	[72%, 100%]	n/a	n/a	17.0s	0.0%	none
cli/codex	100.0%	[72%, 100%]	n/a	n/a	26.3s	0.0%	none
codestral-latest	20.0%	[6%, 51%]	60 737	not billed	7.4s	0.0%	submitted_failed

Tokens and cost read n/a where a lane reports no telemetry (a flat-rate CLI) — an absent measurement, not zero. pass@1 and its interval are over evaluable rollouts only.

This is not a leaderboard. Separated by this run (intervals do not overlap): cli/claude-sonnet vs codestral-latest; cli/codex vs codestral-latest. Not separated (intervals overlap): cli/claude-sonnet vs cli/codex. The primary comparison in this report is the stage funnel below — where a model stops is a finding; a percentage over this many rollouts is not.

Stage funnel

The funnel is the primary reading of a run. Where a model stops says more than a single pass percentage does, especially at low N.

STAGE	CLI:CLI/CLAUDE-SONNET	CLI:CLI/CODEX	MISTRAL:CODESTRAL-LATEST
Inspected	10 / 10	10 / 10	10 / 10
Edited	10 / 10	10 / 10	10 / 10
Ran tests	10 / 10	10 / 10	10 / 10
Submitted	10 / 10	10 / 10	10 / 10
Graded	10 / 10	10 / 10	10 / 10
Passed hidden tests	10 / 10	10 / 10	2 / 10

Deepest stage reached

MODEL	DEEPEST (ANY ROLLOUT)	TYPICAL (MOST ROLLOUTS)	HIDDEN-TEST PASSES	NOT GRADED
cli:cli/claude-sonnet	Passed hidden tests	Passed hidden tests	10 / 10	0
cli:cli/codex	Passed hidden tests	Passed hidden tests	10 / 10	0
mistral:codestral-latest	Passed hidden tests	Graded	2 / 10	0

Stage definitions are in Appendix D. The matrix counts evaluable rollouts; lane errors are shown beside it, not inside it. The stages are not strictly nested: a model can submit without editing, and the matrix shows that rather than hiding it. Each rollout's own deepest stage is in Appendix A.

Behavioural analysis

Long-horizon behaviour

Counted off the recorded trace of each rollout. A dead end is a step that moved nothing: a re-read of a slice already returned, a repeat of the previous action, or a call to a tool this exam does not have.

MODEL	N	STEPS USED	RAN OUT OF STEPS	COMMANDS	CHANGED THE TREE	REFUSED (OVER BUDGET)	REVISITS	DEAD ENDS / ROLLOUT
cli:cli/claude-sonnet	10	5 / 20	0	0 / 24	0	0	0	0
cli:cli/codex	10	6 / 20	0	0 / 24	0	0	0	0
mistral:codestral-latest	10	4 / 20	0	0 / 24	0	0	0	0

Each column is defined in Appendix D.

Hallucination & overclaim

Three measurements, each with a stated definition. None of them is a judgement of the model's prose by another model — every number below comes from comparing what the model wrote against what was on disk, or against what the independent grader returned.

MODEL	N	GROUNDING FAILURES	KINDS	ROLLOUTS AFFECTED	OVERCLAIMS	OVERCLAIM RATE	RE-READS
cli:cli/claude-sonnet	10	0	none	0 / 10	0 / 10	0.0%	0
cli:cli/codex	10	0	none	0 / 10	0 / 10	0.0%	0
mistral:codestral-latest	10	0	none	0 / 10	0 / 10	0.0%	0

No grounding failures and no false pass claims were detected across these 30 rollout(s). No lane re-read material it had already been given.

Every term above is defined in Appendix D.

Efficiency & telemetry

This run has no per-token price attached — a local model, a free quota, or a customer endpoint we are not billed for — so spend is reported in tokens. Tokens are what was actually measured; a dollar figure here would be invented. Lanes run through a flat-rate local CLI report no token telemetry; their cells say n/a — an absent measurement, not zero. Rollout counts here are evaluable rollouts.

MODEL	ROLLOUTS	TOKENS	TOKENS / ROLLOUT	COST	PASSES	RATE
cli:cli/claude-sonnet	10	n/a	n/a	n/a — telemetry unavailable	10 / 10	100.0%
cli:cli/codex	10	n/a	n/a	n/a — telemetry unavailable	10 / 10	100.0%
mistral:codestral-latest	10	60 737	6 073	not billed	2 / 10	20.0%

Findings

Every finding below is derived from the recorded data of this run — none is an opinion.
Severity: PASS a lane succeeded · ATTENTION a capability gap · LANE / INFRA
infrastructure, not capability · NOTE a property of the measurement.

F-01

PASS

cli/claude-sonnet, cli/codex, codestral-latest passed the independent hidden tests

Severity: PASS

Evidence: Model outcomes · Stage funnel · Appendix A

22 of 30 graded rollouts completed the full loop (inspect, edit, run the visible tests, submit) and the sealed grader accepted the submission. This demonstrates the task is solvable under the published limits.

F-01 · derived from the recorded run data

F-02

ATTENTION

submitted_failed × 8 (codestral-latest)

Severity: ATTENTION

Evidence: Appendix A · Appendix C

Submitted a change; the independent hidden tests still failed.

F-02 · derived from the recorded run data

Recommendations

One recommendation per observed failure mode, addressed to the lane it affects. Each traces back to a finding above and to the raw trace in Appendix A.

R-01 PROBLEM

Submitted a change; the independent hidden tests still failed.

AFFECTED LANES

codestral-latest

WHAT WE WOULD LOOK AT FIRST

The model completes the loop and produces a plausible near-miss — the named failing tests are the exact behavioral cases distinguishing it from the reference fix.

Verification

RUN ID

live-20260901T234210Z

EXAM FINGERPRINT

2052eec0ff2d3044dc269e3599ae1bfaa45d0d8cd5df9554796105538d226c21

The contract digest this run was executed against.

Provenance of this chapter

This chapter was rendered by the campaign generator from the sealed evaluation record; it is not the standalone report reproduced byte for byte. Verify the standalone report with its own digest, verify the record with `vvdex-env records verify`, and verify this campaign document as a whole.

RELATIONSHIP

campaign-rendered chapter

Derived from the same sealed evaluation record as the standalone report; these bytes are authenticated by the campaign artifact that contains them.

SEALED RECORD DIGEST

1127e8356d42d31767a3ddc4d6c6cd0a515f5bbf7239ef0a82129e00fc2e6900

sha256 over the record's canonical JSON — the identity every renderer works from.

SOURCE RECORD FILE

d3f91221afd12e368b53c09a12eba41633c305d41855842ca6eef918de09e857

sha256 of `fr-20260901-46eb9548.json` as written beside the standalone report.

SOURCE STANDALONE REPORT

17b7f5bd9f45aa56a19fd1dec4893b9d1ff83b773eace2cc593a0a828e0baf1f

sha256 of `fr-20260901-46eb9548.html` — independently verifiable with its own digest.

SOURCE STANDALONE PDF

0b2cf9181b184b492f780d9f7b9b419093e897329ef9fb98afdfa348c0f88066

sha256 of `fr-20260901-46eb9548.pdf`.

RUN EVIDENCE BUNDLE

46eb954840000beb6e8466f409bb6104c91cd91a0d34efa18a41871246efd142

The digest the run's own bundle computed over itself.

CERTIFIED EVIDENCE SEAL

not available in this package

The exam's sealed evidence, established before this run.

Measurement history

This is the first recorded run on this exam. There is nothing to compare it against yet, and a single measurement is not a trend.

The full artifact-by-artifact digest table, and the reproduction protocol, are in Appendix B.

Appendix A — every rollout

ROLLOUT	MODEL	SEED	OUTCOME	SUBMITTED	DEEPEST STAGE	HIDDEN TESTS	FILES	TOKENS	ELAPSED
c3e59f84	cli/claude-sonnet	0	submitted_passed	yes	Passed hidden tests	pass	1	n/a	20.5s
3f5d33eb	cli/claude-sonnet	1	submitted_passed	yes	Passed hidden tests	pass	1	n/a	16.5s
91cd84ed	cli/claude-sonnet	2	submitted_passed	yes	Passed hidden tests	pass	1	n/a	15.7s
d3ccb200	cli/claude-sonnet	3	submitted_passed	yes	Passed hidden tests	pass	1	n/a	14.7s
a41f48cc	cli/claude-sonnet	4	submitted_passed	yes	Passed hidden tests	pass	1	n/a	18.9s
be0bb482	cli/claude-sonnet	5	submitted_passed	yes	Passed hidden tests	pass	1	n/a	15.6s
295bfb44	cli/claude-sonnet	6	submitted_passed	yes	Passed hidden tests	pass	1	n/a	15.8s
dd33c07a	cli/claude-sonnet	7	submitted_passed	yes	Passed hidden tests	pass	1	n/a	18.3s
c138c3a9	cli/claude-sonnet	8	submitted_passed	yes	Passed hidden tests	pass	1	n/a	15.1s
9d3daa5f	cli/claude-sonnet	9	submitted_passed	yes	Passed hidden tests	pass	1	n/a	18.8s
37ff761f	cli/codex	0	submitted_passed	yes	Passed hidden tests	pass	1	n/a	26.7s
7452d373	cli/codex	1	submitted_passed	yes	Passed hidden tests	pass	1	n/a	26.2s
9d7733ce	cli/codex	2	submitted_passed	yes	Passed hidden tests	pass	1	n/a	29.9s
83cc4801	cli/codex	3	submitted_passed	yes	Passed hidden tests	pass	1	n/a	26.9s
607bc653	cli/codex	4	submitted_passed	yes	Passed hidden tests	pass	1	n/a	26.5s
f495da98	cli/codex	5	submitted_passed	yes	Passed hidden tests	pass	1	n/a	25.9s
e880dfa5	cli/codex	6	submitted_passed	yes	Passed hidden tests	pass	1	n/a	25.0s
730a0c3f	cli/codex	7	submitted_passed	yes	Passed hidden tests	pass	1	n/a	24.3s
59a40ede	cli/codex	8	submitted_passed	yes	Passed hidden tests	pass	1	n/a	25.8s
1b564187	cli/codex	9	submitted_passed	yes	Passed hidden tests	pass	1	n/a	26.3s
d38a9234	codestral-latest	0	submitted_failed	yes	Graded	fail	1	5 764	7.2s
0f23b24f	codestral-latest	1	submitted_failed	yes	Graded	fail	1	5 764	7.2s
17fd7879	codestral-latest	2	submitted_failed	yes	Graded	fail	1	5 764	6.6s
cee5d0af	codestral-latest	3	submitted_passed	yes	Passed hidden tests	pass	1	7 279	8.1s
2e244621	codestral-latest	4	submitted_failed	yes	Graded	fail	1	5 800	7.1s
d85946ef	codestral-latest	5	submitted_failed	yes	Graded	fail	1	5 764	6.6s
1ba66ffd	codestral-latest	6	submitted_failed	yes	Graded	fail	1	5 749	6.4s
217e2513	codestral-latest	7	submitted_failed	yes	Graded	fail	1	5 735	6.9s
4f72002b	codestral-latest	8	submitted_passed	yes	Passed hidden tests	pass	1	7 354	10.0s
036f7a4e	codestral-latest	9	submitted_failed	yes	Graded	fail	1	5 764	7.9s

Failure taxonomy

CLASS	COUNT	WHAT IT MEANS
submitted_failed	8	Submitted a change; the independent hidden tests still failed.

Appendix B — evidence and artifact digests

RUN BUNDLE DIGEST

46eb954840000beb6e8466f4

09bb6104c91cd91a0d34efa1

8a41871246efd142

RUN ID

live-

20260901T234210Z

STARTED

2026-09-

01T23:42:10.844945

+00:00

FINISHED

2026-09-

01T23:53:46.554546

+00:00

Artifact digests

ARTIFACT	SHA-256
publicSummary	f00eecfd8c1b1f59dba8b9a50ea2f63d453b4e0510a0c539a07f29b817333639
rollouts/036f7a4e-95a8-44f4-9416-4a630956bde5.json	d078da23ce6dd43330d34fbca2f32ae906d034258b3aa6ad54082119bffcbbfd4
rollouts/0f23b24f-8799-4da2-8eed-9409735afdfc.json	1fcf3e129ae41a02f38af0d9a98f4363fa1440c365311bef7fac89e35c98d70f
rollouts/17fd7879-a548-4888-b365-393483b721f2.json	e96c42e097096a493d9b8a62c847c7ae08199ace0b64a9c0185742753a99e82b
rollouts/1b564187-de5c-4b9e-bba0-a3497ba2abc2.json	820522a65667acba036bb85c75e50022864897aeadf268460be1a69015d233a4
rollouts/1ba66ffd-229b-40e8-b3b0-2dae920c85b1.json	54313858ba041c81401c4e7730dca90280991e21d7014a191ba8a9dedba0ef2b
rollouts/217e2513-92b3-4f09-9148-da9988b90f7f.json	2405e3e91b701f9be6136075080b29913d9befbd38e5e7293664929941eff90c
rollouts/295bfb44-12e8-48f8-8bdc-f5df9bcecc9b.json	470e991bd1b3c34e333f75659a78b7766c90c4f98c1d01cdd5e61809ff2cb042
rollouts/2e244621-3ca6-446e-9945-ed98537cc81b.json	7b786e0eaa7fa5202436cd67418fc96bec6068b29d0b9a30f53b253a1920894c
rollouts/37ff761f-a2b3-41fa-9b42-26ed00fb4d26.json	0ab83298fae4aca11a763554a35211137c3e0d1fbe46d59466c5a04bfcacf5c43
rollouts/3f5d33eb-55bb-4320-b67c-2c20616d2509.json	0c25ae80d7162cdc131aa3168c52798a84cd18c33ec974100028dcd9cd6d09a4
rollouts/4f72002b-67f3-42b3-928a-79d7022d5180.json	c08d218f7332e24c3008c9c55814e5e29923041a02082d90f7c80defe1e741b3
rollouts/59a40ede-e515-48e5-9aa7-cce5bdd309b1.json	b33d19ff66eac735813648cfd97eeda7dd2de54eb8b82ad3b415cec658cc8f0d
rollouts/607bc653-711f-45d6-a26b-7bdfd01b0142.json	6bba74788ab8b38728a1fb1a72f07888544667fc5052389a6db6add67d782f1f
rollouts/730a0c3f-2c1b-43a0-a6c7-f005a0c7a0d6.json	c31b768016f3e9260afc4c7c8478eecdd4406a35377cf4340a49ee3514eaa5eb8
rollouts/7452d373-a07f-4c66-951b-b8c82c82badd.json	16725edfff8953f43b792cf66a786bf628a0a0e1e062413fa2f6d8327357e420
rollouts/83cc4801-1d9e-432d-81b1-31caa05f6033.json	237bfa5642324e6eac921a7aa6740045e8159a3604c9349dfec9c3551f07ec0c
rollouts/91cd84ed-bdd4-4eb6-9de6-56238bcb1d17.json	05848a847f9a10cb04d4b458675838967c5f43d7ff3dd3f19b23358ff4e0b3ce
rollouts/9d3daa5f-a2fe-4469-87dc-539a6d2d3d45.json	0182e362994ab161f4cf7ce4ed1a5206c64ad1c21a847412635538c8864d8e1b
rollouts/9d7733ce-0c7c-4f9f-8491-2d58b27af9d7.json	872aca186137b7f25f1454c455310ab41a76481ea656e1c49dfccb7b9f21e97d
rollouts/a41f48cc-c287-42e1-81b0-152f4f2017e6.json	a16e127ec3f8e19a5a816a9ae53fd4f343336e65efbb0de94fe28991caacd5b5
rollouts/be0bb482-0aa0-4452-aebc-78d3267da5c6.json	e9969415dbc882c8c8d1a53b1cc7b23adf2525700c97cbdb046f6f291691f10e
rollouts/c138c3a9-d1fc-4396-8437-dfa9356e8f57.json	ffe83adba7a97ad3ef5b2f83a52968f7b382ad2e88f252464e060e78c319edfe
rollouts/c3e59f84-2de4-4f17-9165-32d58a341c89.json	6a323bc275c391f285dd1bb3c2c67cd504190eef3f19636da5339f2caccc16b2
rollouts/cee5d0af-4f81-49fa-9da2-3ec9f3f4c2b1.json	0941fecc01b8577f8aaa411cd47566e51749690ac3280ec036282c6d1753fd65
rollouts/d38a9234-5c00-4282-b6d1-8597b2cbd8b5.json	d89c7fa3b927cb83fac42b3cadf115e0232420df1dd9675804079c2251e5b883
rollouts/d3ccb200-349b-4f7d-83fb-8cdb6453a294.json	d8ffbb6990497de81d41d9c34e29bfcfb6665dba0dfe2ee00c4df4411b1db866f
rollouts/d85946ef-5125-4b10-806a-9d17237fed3f.json	2418001521dae6707aa7154745363f2e55f589a1fbcb04ff66eeda00c924a16a
rollouts/dd33c07a-e79c-4cfa-9367-eb1c99991173.json	62da7725d6f6ee0b5c479b7336dfa4d1aaee713977a6df765ac506f5e0585479

ARTIFACT	SHA-256
rollouts/e880dfa5-8f79-4a13-817f-5be17c6dfba7.json	b33a34cf8dc3133a80c45e33a78b1d8e386c03fe8d9f0f253143db986ae57ddc
rollouts/f495da98-b293-4bd7-8722-bfb4966bc3ce.json	f2a4ffc76691fa6e27d96a489f2410c0c649f91fbd0d41307ae46cbf5abf1d1f
spec	cdaaa1f547234167c055e49c22f93ad3965c87c32220a4bf8c4c7cf94ea5ce50
summary	d5fb8760400394f5e669eff1476bdfccb37e0140a48e65d4f6954001d4cafd99

Reproducing this run

Every rollout can be re-derived: the exam is frozen (fingerprint above), the grader ran in the recorded container image, and each rollout's full trace — every tool call, every result, every word — is in the evidence bundle at its rollouts/<id>.json path. A re-run uses the same job shape: the exam id, your endpoint, n, and the runtime. The record id names this run forever.

Appendix C — redacted transcripts

For each rollout: the tool calls it made, in order — which tool, on which of the model's own paths, whether the call succeeded, and how many characters came back. What a tool returned is not here, and neither is the submission: on this exam the submitted content is the graded answer, so no diff, no answer value and no model claim is printed. A request for the reference solution or the hidden tests is recorded as a request; the path it asked for is replaced by a label. See Appendix D, "What this report withholds and why".

Disclosure class applied to this exam: answer-key disclosure (ownership vvdex_proprietary — engine registry; family long_term_memory): the exam is VVDex evaluation IP and the graded output is an answer or a keyed value, so no submission content, diff, model claim or hidden-assertion name is published.

c3e59f84 · cli:cli/claude-sonnet · seed 0 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	memory/notes.md	ok	135
1	read_file	facts/current.json	ok	65
2	write_file	answers/answer.json	ok	68
3	run_tests	—	ok	181
4	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (9 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

3f5d33eb · cli:cli/claude-sonnet · seed 1 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	memory/notes.md	ok	135
1	read_file	facts/current.json	ok	65
2	write_file	answers/answer.json	ok	68
3	run_tests	—	ok	181
4	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (9 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

91cd84ed · cli:cli/claude-sonnet · seed 2 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	memory/notes.md	ok	135
1	read_file	facts/current.json	ok	65
2	write_file	answers/answer.json	ok	68
3	run_tests	—	ok	181
4	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (9 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

d3ccb200 · cli:cli/claude-sonnet · seed 3 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	memory/notes.md	ok	135
1	read_file	facts/current.json	ok	65
2	write_file	answers/answer.json	ok	96
3	run_tests	—	ok	181
4	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (9 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

a41f48cc · cli:cli/claude-sonnet · seed 4 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	memory/notes.md	ok	135
1	read_file	facts/current.json	ok	65
2	write_file	answers/answer.json	ok	68
3	run_tests	—	ok	181
4	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (9 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

be0bb482 · cli:cli/claude-sonnet · seed 5 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	memory/notes.md	ok	135
1	read_file	facts/current.json	ok	65
2	write_file	answers/answer.json	ok	68
3	run_tests	—	ok	181
4	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (9 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

295bfb44 · cli:cli/claude-sonnet · seed 6 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	memory/notes.md	ok	135
1	read_file	facts/current.json	ok	65
2	write_file	answers/answer.json	ok	68
3	run_tests	—	ok	181
4	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (9 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

dd33c07a · cli:cli/claude-sonnet · seed 7 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	memory/notes.md	ok	135
1	read_file	facts/current.json	ok	65
2	write_file	answers/answer.json	ok	68
3	run_tests	—	ok	181
4	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (9 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

c138c3a9 · cli:cli/claude-sonnet · seed 8 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	memory/notes.md	ok	135
1	read_file	facts/current.json	ok	65
2	write_file	answers/answer.json	ok	68
3	run_tests	—	ok	181
4	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (9 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

9d3daa5f · cli:cli/claude-sonnet · seed 9 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	memory/notes.md	ok	135
1	read_file	facts/current.json	ok	65
2	write_file	answers/answer.json	ok	68

STEP	TOOL	PATH	RESULT	CHARS
3	run_tests	—	ok	181
4	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (9 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

37ff761f · cli:cli/codex · seed 0 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	ISSUE.md	ok	1 220
1	read_file	facts/current.json	ok	65
2	read_file	memory/notes.md	ok	135
3	write_file	answers/answer.json	ok	68
4	run_tests	—	ok	181
5	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (9 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

7452d373 · cli:cli/codex · seed 1 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	ISSUE.md	ok	1 220
1	read_file	facts/current.json	ok	65
2	read_file	memory/notes.md	ok	135
3	write_file	answers/answer.json	ok	68
4	run_tests	—	ok	181
5	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (9 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

9d7733ce · cli:cli/codex · seed 2 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	ISSUE.md	ok	1 220
1	list_files	.	ok	57
2	read_file	facts/current.json	ok	65
3	read_file	memory/notes.md	ok	135
4	write_file	answers/answer.json	ok	68
5	run_tests	—	ok	181

STEP	TOOL	PATH	RESULT	CHARS
6	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (9 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

83cc4801 · cli:cli/codex · seed 3 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	ISSUE.md	ok	1 220
1	list_files	.	ok	57
2	read_file	facts/current.json	ok	65
3	read_file	memory/notes.md	ok	135
4	write_file	answers/answer.json	ok	68
5	run_tests	—	ok	181
6	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (9 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

607bc653 · cli:cli/codex · seed 4 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	ISSUE.md	ok	1 220
1	read_file	facts/current.json	ok	65
2	read_file	memory/notes.md	ok	135
3	write_file	answers/answer.json	ok	68
4	run_tests	—	ok	181
5	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (9 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

f495da98 · cli:cli/codex · seed 5 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	ISSUE.md	ok	1 220
1	read_file	memory/notes.md	ok	135
2	read_file	facts/current.json	ok	65
3	write_file	answers/answer.json	ok	68
4	run_tests	—	ok	181
5	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (9 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

e880dfa5 · cli:cli/codex · seed 6 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	ISSUE.md	ok	1 220
1	read_file	memory/notes.md	ok	135
2	read_file	facts/current.json	ok	65
3	write_file	answers/answer.json	ok	68
4	run_tests	—	ok	181
5	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (9 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

730a0c3f · cli:cli/codex · seed 7 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	ISSUE.md	ok	1 220
1	read_file	memory/notes.md	ok	135
2	read_file	facts/current.json	ok	65
3	write_file	answers/answer.json	ok	68
4	run_tests	—	ok	181
5	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (9 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

59a40ede · cli:cli/codex · seed 8 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	ISSUE.md	ok	1 220
1	read_file	memory/notes.md	ok	135
2	read_file	facts/current.json	ok	65
3	write_file	answers/answer.json	ok	68
4	run_tests	—	ok	181
5	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (9 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

1b564187 · cli:cli/codex · seed 9 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	ISSUE.md	ok	1 220
1	read_file	facts/current.json	ok	65
2	read_file	memory/notes.md	ok	135
3	write_file	answers/answer.json	ok	68
4	run_tests	—	ok	181
5	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (9 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

d38a9234 · mistral:codestral-latest · seed 0 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	memory/notes.md	ok	135
1	write_file	answers/answer.json	ok	68
2	run_tests	—	ok	181
3	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (4 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

0f23b24f · mistral:codestral-latest · seed 1 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	memory/notes.md	ok	135
1	write_file	answers/answer.json	ok	68
2	run_tests	—	ok	181
3	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (4 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

17fd7879 · mistral:codestral-latest · seed 2 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	memory/notes.md	ok	135
1	write_file	answers/answer.json	ok	68
2	run_tests	—	ok	181
3	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (4 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

cee5d0af · mistral:codestral-latest · seed 3 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	facts/current.json	ok	65
1	read_file	memory/notes.md	ok	135
2	write_file	answers/answer.json	ok	68
3	run_tests	—	ok	181
4	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (4 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

2e244621 · mistral:codestral-latest · seed 4 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	memory/notes.md	ok	135
1	write_file	answers/answer.json	ok	68
2	run_tests	—	ok	181
3	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (9 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

d85946ef · mistral:codestral-latest · seed 5 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	memory/notes.md	ok	135
1	write_file	answers/answer.json	ok	68
2	run_tests	—	ok	181
3	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (4 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

1ba66ffd · mistral:codestral-latest · seed 6 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	memory/notes.md	ok	135
1	write_file	answers/answer.json	ok	68

STEP	TOOL	PATH	RESULT	CHARS
2	run_tests	—	ok	181
3	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (4 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

217e2513 · mistral:codestral-latest · seed 7 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	facts/current.json	ok	65
1	write_file	answers/answer.json	ok	68
2	run_tests	—	ok	181
3	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (9 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

4f72002b · mistral:codestral-latest · seed 8 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	facts/current.json	ok	65
1	read_file	memory/notes.md	ok	135
2	run_tests	—	ok	181
3	write_file	answers/answer.json	ok	68
4	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (9 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

036f7a4e · mistral:codestral-latest · seed 9 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	memory/notes.md	ok	135
1	write_file	answers/answer.json	ok	68
2	run_tests	—	ok	181
3	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (4 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

14 Evidence

Every number above is recomputable from the sealed records below; each record verifies itself with `vvdex-env records verify`, and this document was regenerated from them with `vvdex-env records campaign`.

- `public.swe.martinblech-xmltodict-issue-257` → `fr-20260902-6616b188` · `recordDigest` `138425da59aabc18aa208c339aa41919f387e1ecb80fb725cfed8f7d6ba78713`
- `vvdex.browser.order-desk-1` → `fr-20260902-6e618c05` · `recordDigest` `59f94854d78f7ce174866ff114505b2b41ca2333f287f40b604fbb38cc267da5`
- `vvdex.harness.fault-recovery-1` → `fr-20260902-a6f12b64` · `recordDigest` `d5a8264aa0c20dd67c14879856a744e2ff1ab79da35ea8bcea92b4ba1ea14d5f`
- `vvdex.knowledge.grounded-rag-1` → `fr-20260902-2e7f4171` · `recordDigest` `157ec324a925c0b10fc8769d01f628b6143c40a00fa2dc70ea309a70e55a37a8`
- `vvdex.knowledge.memory-fact-update-1` → `fr-20260901-46eb9548` · `recordDigest` `1127e8356d42d31767a3ddc4d6c6cd0a515f5bbf7239ef0a82129e00fc2e6900`

15 Attestation

Certified results appear only where evidence proves them. Withheld, invalid and lane-error cells are never presented as model results, and no historical evidence was modified in producing this document. Counts are exact; below ten rollouts per cell no rate or ranking beyond the counts is asserted.

© VVDex. All rights reserved. Public materials are provided for viewing, evaluation and verification. Reuse, redistribution, derivative benchmark creation, training use, republication or commercial use requires written permission unless a specific file states otherwise. Full terms: vvdexops.com/terms/. Third-party open source reproduced in an exam keeps its own licence.