

Answering questions from a document collection with outdated and distractor sources, and refusing when the evidence is insufficient

EVALUATION REPORT

2026-09-02

Evaluation scope:

Claude Code CLI (Sonnet, subscription)

Codex CLI (GPT, subscription)

codestral-latest

Prepared by:

VVDex Forge — generated directly from sealed run evidence

fr-20260902-2e7f4171 · vvdex.knowledge.grounded-rag-1 v0.1.0 · Independent model evaluation report

About this report

About VVDex Forge

VVDex Forge builds and certifies evaluation environments for AI agents. Every exam is sealed before any model sees it: the reference solution must pass, an empty submission must fail, the grader must reject near-misses, and the hidden tests never enter the agent's tree. Reports and the public catalog live at **vvdexops.com**.

Report integrity

This report is generated from run evidence; its digests verify authenticity. The sole test of a copy is the three-line procedure in Verification — a copy whose digests do not reproduce is not this report.

Scope

One exam · 30 rollout(s) · 30 tool steps and 7320s wall clock per rollout, stated in full in Evaluation configuration.

Boundary

Not a leaderboard. Not a general capability score. A post-seal benchmark: this run stamped no certification gate and rewrote nothing — not the exam, not the reference solution, not the grader, not the certified evidence.

Contents

01.	Executive summary	4
02.	Exam definition	5
03.	Certification evidence	6
04.	Evaluation configuration	7
05.	Model outcomes	8
06.	Stage funnel	9
07.	Behavioural analysis	10
08.	Efficiency & telemetry	11
09.	Findings	12
10.	Recommendations	13
11.	Verification	14

Appendices: A — every rollout · B — evidence digests · C — redacted transcripts · D — metric definitions · Attestation

01 Executive summary

MODELS TESTED 3	CERTIFIED PASSES 11 of 29 graded rollouts	SUBMITTED 29 of 29 graded rollouts that called submit	NOT GRADED — LANE ERRORS 1 of 30 attempts
---------------------------------------	--	--	--

Purpose

This report presents a controlled evaluation of 3 model lane(s) against the certified exam vv dex.knowledge.grounded-rag-1 (v0.1.0, family rag_grounding). Every lane received the identical frozen workspace, tool contract and limits; grading ran afterwards, in isolation, against hidden tests no lane ever saw. The purpose is to state, with reproducible evidence, what each lane did and where it stopped.

Outcome

11 of 29

graded rollouts passed the independent hidden tests and were certified — 1 of 30 attempts not graded (lane errors)

11 of 29 graded rollouts passed. cli/claude-sonnet, cli/codex completed the full loop (inspect, edit, run the visible tests, submit) and the sealed grader accepted the submission. The most common way to fail was submitted_failed — submitted a change; the independent hidden tests still failed.

Key findings

F-01 cli/claude-sonnet, cli/codex passed the independent hidden tests

F-02 submitted_failed × 18 (cli/claude-sonnet, codestral-latest)

F-03 model_api_failure × 1 (codestral-latest)

Each finding is stated in full, with evidence, in the Findings section; recommendations for acting on them follow it. Elapsed wall clock for the whole engagement: 1500.2s.

02 Exam definition

In scope. The ability of each configured model lane to complete this one certified task end to end — inspect the workspace, produce a correct edit, verify it with the visible tests, and submit — within 30 tool steps and 7320s of wall clock, at 10 rollout(s) per lane. Behavioural measurements along the way (tool discipline, grounding, overclaim) are in scope because they are read from the recorded trace, not judged by another model.

Out of scope. General model quality, performance on other task families, prompt-tuning on the lane's behalf, and any comparison this run's sample size cannot support. A lane's API failure is reported as infrastructure, never as model capability.

The instrument. Exam `vvdex.knowledge.grounded-rag-1` was certified before this run: its reference solution passes, an empty submission fails, its grader distinguishes near-misses, and its sealed evidence is unchanged by this evaluation (see Certification evidence).

What was measured

EXAM	VERSION	FAMILY	ROLLOUTS
vvdex.knowledge.grounded-rag-1	0.1.0	rag_grounding	30 (10 per model)

Where the defect came from

This exam has no upstream repository to credit: the defect, the reference solution and the hidden suite are VVDex's own.

03 Certification evidence

4 of 5 certification pillars are recorded in this package, loaded from the receipts sealed into the exam and matching this run's exam fingerprint. What is recorded is shown with its real value; what is not is named below and is not claimed.

FROZEN BASELINE

FAIL

expected — an empty submission over 2 run(s) must not pass

→

GOLD / REFERENCE

PASS

expected — the reference solution over 2 run(s) must pass

GRADER CONTROLS

7 / 7

The grader accepted the reference and rejected every near-miss.

ATTACK PROBES

10 / 10 blocked

0 trivial exploit(s) found.

RUNTIME

sha256:679b36225a1f83238efba8c71b9cde3c24caaeacc9b682ae92819cb55eec328f

network policy: deny

Certification metadata is not included in this report package for:

- sealed evidence bundle

Benchmark results remain valid as run evidence; the pillar(s) above are not substantiated by this document and this report does not claim them.

What the package does carry: the exam fingerprint below is the contract digest this run was executed against — a holder of the certification bundle for this fingerprint can join the two independently.

EXAM FINGERPRINT

4880a527ef97ca2f3fe6b57be026c3dc1be51e9e104d79fc13a3e558331dda92

The contract digest this run was executed against.

04 Evaluation configuration

REPORT	fr-20260902-2e7f4171
ISSUED	2026-09-02
SUBJECT EXAM	vvdex.knowledge.grounded-rag-1 · v0.1.0
FAMILY	rag_grounding
PREPARED BY	VVDex Forge engine vvdex-env 0.1.0
CLASSIFICATION	Independent model evaluation report — independently verifiable
CERTIFICATION	Exam certification: partially recorded (see Certification evidence)
STATUS	FINAL · report sealed against its own digest

Timeline

WHEN (UTC)	EVENT
2026-09-01 23:53:46	Evaluation run started
2026-09-02 00:18:47	Evaluation run finished
2026-09-02 00:18:47	Report generated from the sealed run evidence

Models under test

MODEL	PROVIDER	ENDPOINT	BILLING
Claude Code CLI (Sonnet, subscription)	cli	local runtime	operator subscription, flat — not billed per token
Codex CLI (GPT, subscription)	cli	local runtime	operator subscription, flat — not billed per token
codestral-latest	mistral	VVDex-configured	VVDex free-quota

A customer endpoint is recorded by origin only — never its port, its path, or its key.

Limits

RUNTIME docker	STEP LIMIT 30	WALL-CLOCK LIMIT 7320s	MAX OUTPUT TOKENS 2048
TEMPERATURE 0.2	RETRY POLICY none	COMMAND BUDGET 40	TOOLS OFFERED list_files, read_file, write_file, edit_file, run_tests, submit, run_command

The evaluated model never saw the hidden tests or the reference solution, and could not change its result. Grading ran in a separate step, after submission, on a container with no network.

05 Model outcomes

Denominators. Attempts requested: 30. Graded (evaluable) rollouts: 29. Not graded — lane errors: 1. A lane error is a rollout the API or the runtime ended before the hidden grader ran: it is listed, excluded from every rate and pass@k, and never silently dropped. Passes are certified passes: hidden grader exit 0 with integrity verified.

MODEL	ATTEMPTS	GRADED	PASSED	MODEL FAILS	LANE ERRORS	SUBMITTED	MEAN TIME	TYPICAL DEEPEST STAGE
Claude Code CLI (Sonnet, subscription) cli/claude-sonnet	10	10	1 / 10	9	0	10 / 10	45.8s	Graded
Codex CLI (GPT, subscription) cli/codex	10	10	10 / 10	0	0	10 / 10	67.8s	Passed hidden tests
codestral-latest codestral-latest	10	9	0 / 9	9	1	9 / 9	18.0s	Graded

Every rollout, with its own deepest stage and grade, is in Appendix A.

Success rate, confidence and pass@k

Every rate on this page is a measurement of a finite number of rollouts, so every rate carries the interval that measurement actually supports. n is the evaluable count — attempts minus rollouts the lane ended before grading; those are listed under "not graded" and sit in no rate, interval or pass@k.

MODEL	ATTEMPTS	N	NOT GRADED	PASSES	RATE	95% INTERVAL	SEEDS	PASS@1	PASS@3	PASS@5	PASS@10
cli/claude-sonnet	10	10	0	1	10.0%	[2%, 40%]	0–9	10.0%	30.0%	50.0%	100.0%
cli/codex	10	10	0	10	100.0%	[72%, 100%]	0–9	100.0%	100.0%	100.0%	100.0%
codestral-latest	10	9	1	0	0.0%	[0%, 30%]	0–9	0.0%	0.0%	0.0%	—

Interval and pass@k methods are stated in full in Appendix D. pass@k is shown for k = 1, 3, 5, 10; the sealed record carries every k up to n.

Model comparison

MODEL	PASS@1	95% INTERVAL	TOKENS	COST	MEAN ROLLOUT	OVERCLAIM RATE	TOP FAILURE CLASS
cli/claude-sonnet	10.0%	[2%, 40%]	n/a	n/a	45.8s	0.0%	submitted_failed
cli/codex	100.0%	[72%, 100%]	n/a	n/a	67.8s	0.0%	none
codestral-latest	0.0%	[0%, 30%]	225 330	not billed	18.0s	0.0%	submitted_failed

Tokens and cost read n/a where a lane reports no telemetry (a flat-rate CLI) — an absent measurement, not zero. pass@1 and its interval are over evaluable rollouts only.

This is not a leaderboard. Separated by this run (intervals do not overlap): cli/claude-sonnet vs cli/codex; cli/codex vs codestral-latest. Not separated (intervals overlap): cli/claude-sonnet vs codestral-latest. Below the 10-rollout floor for reading a rate as a rate: codestral-latest (9) — read the count, not the percentage. The primary comparison in this report is the stage funnel below — where a model stops is a finding; a percentage over this many rollouts is not.

06 Stage funnel

The funnel is the primary reading of a run. Where a model stops says more than a single pass percentage does, especially at low N.

STAGE	CLI:CLI/CLAUDE-SONNET	CLI:CLI/CODEX	MISTRAL:CODESTRAL-LATEST
Inspected	10 / 10	10 / 10	9 / 9
Edited	8 / 10	10 / 10	9 / 9
Ran tests	10 / 10	10 / 10	9 / 9
Submitted	10 / 10	10 / 10	9 / 9
Graded	10 / 10	10 / 10	9 / 9
Passed hidden tests	1 / 10	10 / 10	0 / 9

Deepest stage reached

MODEL	DEEPEST (ANY ROLLOUT)	TYPICAL (MOST ROLLOUTS)	HIDDEN-TEST PASSES	NOT GRADED
cli:cli/claude-sonnet	Passed hidden tests	Graded	1 / 10	0
cli:cli/codex	Passed hidden tests	Passed hidden tests	10 / 10	0
mistral:codestral-latest	Graded	Graded	0 / 9	1

Stage definitions are in Appendix D. The matrix counts evaluable rollouts; lane errors are shown beside it, not inside it. The stages are not strictly nested: a model can submit without editing, and the matrix shows that rather than hiding it. Each rollout's own deepest stage is in Appendix A.

07 Behavioural analysis

Long-horizon behaviour

Counted off the recorded trace of each rollout. A dead end is a step that moved nothing: a re-read of a slice already returned, a repeat of the previous action, or a call to a tool this exam does not have.

MODEL	N	STEPS USED	RAN OUT OF STEPS	COMMANDS	CHANGED THE TREE	REFUSED (OVER BUDGET)	REVISITS	DEAD ENDS / ROLLOUT
cli:cli/claude-sonnet	10	9 / 30	0	0 / 40		0	0	0
cli:cli/codex	10	14 / 30	0	0 / 40		0	0	0
mistral:codestral-latest	10	8 / 30	0	0 / 40		0	0	0

Each column is defined in Appendix D.

Hallucination & overclaim

Three measurements, each with a stated definition. None of them is a judgement of the model's prose by another model — every number below comes from comparing what the model wrote against what was on disk, or against what the independent grader returned.

MODEL	N	GROUNDING FAILURES	KINDS	ROLLOUTS AFFECTED	OVERCLAIMS	OVERCLAIM RATE	RE-READS
cli:cli/claude-sonnet	10	0	none	0 / 10	0 / 10	0.0%	0
cli:cli/codex	10	0	none	0 / 10	0 / 10	0.0%	0
mistral:codestral-latest	10	0	none	0 / 10	0 / 10	0.0%	0

No grounding failures and no false pass claims were detected across these 30 rollout(s). No lane re-read material it had already been given.

Every term above is defined in Appendix D.

08 Efficiency & telemetry

This run has no per-token price attached — a local model, a free quota, or a customer endpoint we are not billed for — so spend is reported in tokens. Tokens are what was actually measured; a dollar figure here would be invented. Lanes run through a flat-rate local CLI report no token telemetry; their cells say n/a — an absent measurement, not zero. Rollout counts here are evaluable rollouts.

MODEL	ROLLOUTS	TOKENS	TOKENS / ROLLOUT	COST	PASSES	RATE
cli:cli/claude-sonnet	10	n/a	n/a	n/a – telemetry unavailable	1 / 10	10.0%
cli:cli/codex	10	n/a	n/a	n/a – telemetry unavailable	10 / 10	100.0%
mistral:codestral-latest	9	225 330	25 036	not billed	0 / 9	0.0%

09 Findings

Every finding below is derived from the recorded data of this run — none is an opinion.
Severity: PASS a lane succeeded · ATTENTION a capability gap · LANE / INFRA
infrastructure, not capability · NOTE a property of the measurement.

F-01

PASS

cli/claude-sonnet, cli/codex passed the independent hidden tests

Severity: PASS

Evidence: Model outcomes · Stage funnel · Appendix A

11 of 29 graded rollouts completed the full loop (inspect, edit, run the visible tests, submit) and the sealed grader accepted the submission. This demonstrates the task is solvable under the published limits.

F-01 · derived from the recorded run data

F-02

ATTENTION

submitted_failed × 18 (cli/claude-sonnet, codestral-latest)

Severity: ATTENTION

Evidence: Appendix A · Appendix C

Submitted a change; the independent hidden tests still failed.

F-02 · derived from the recorded run data

F-03

LANE / INFRA

model_api_failure × 1 (codestral-latest)

Severity: LANE / INFRA

Evidence: Appendix A · Appendix C

The model endpoint failed to answer. Not a verdict about the model's ability. These rollouts are excluded from any capability reading.

F-03 · derived from the recorded run data

10 Recommendations

One recommendation per observed failure mode, addressed to the lane it affects. Each traces back to a finding above and to the raw trace in Appendix A.

R-01

PROBLEM

Submitted a change; the independent hidden tests still failed.

AFFECTED LANES

cli/claude-sonnet, codestral-latest

WHAT WE WOULD LOOK AT FIRST

The model completes the loop and produces a plausible near-miss — the named failing tests are the exact behavioral cases distinguishing it from the reference fix.

R-02

PROBLEM

The model endpoint failed to answer. Not a verdict about the model's ability.

AFFECTED LANES

codestral-latest

WHAT WE WOULD LOOK AT FIRST

The lane failed, not the model; these rollouts carry no signal about capability and are labeled accordingly.

11 Verification

RUN ID

live-20260901T235346Z

EXAM FINGERPRINT

4880a527ef97ca2f3fe6b57be026c3dc1be51e9e104d79fc13a3e558331dda92

The contract digest this run was executed against.

Tamper-evidence

Three digests, and how to check each one. If any of them does not reproduce, this page is not the page that was generated for this run.

THIS REPORT

de6657ee45236bc7cdff7c14a99200ccaec9d3465caeaedc8c9e8eb8c6615013

sha256 over the report's UTF-8 bytes with every occurrence of the report digest itself replaced by 64 '0' characters

EVAL RECORD

a3e742fe73fe3667155e22dfc941bdc9e428a4bb1bd47c5c89a0454e299e7eb0

sha256 of fr-20260902-2e7f4171.json as written beside this file.

RUN EVIDENCE BUNDLE

2e7f4171f049cced92f7a3aa7ae13b2ee211449e96938e7e6c437f30289db0ff

The digest the run's own bundle computed over itself.

CERTIFIED EVIDENCE SEAL

not recorded

The exam's sealed evidence, established before this run.

Measurement history

This is the first recorded run on this exam. There is nothing to compare it against yet, and a single measurement is not a trend.

Verifying it, in three lines

1. vvdex-env verify-report --report fr-20260902-2e7f4171.html

2. or by hand: replace every occurrence of the report digest above with 64 zeros, then sha256 the file – it must equal that digest

3. shasum -a 256 fr-20260902-2e7f4171.json must equal the eval-record digest, and the bundle digest must equal bundleSha256 in the run's bundle file

The full artifact-by-artifact digest table, and the reproduction protocol, are in Appendix B.

Appendix A — every rollout

ROLLOUT	MODEL	SEED	OUTCOME	SUBMITTED	DEEPEST STAGE	HIDDEN TESTS	FILES	TOKENS	ELAPSED
8ec5820d	cli/claude-sonnet	0	submitted_failed	yes	Graded	fail	1	n/a	25.2s
72bcf971	cli/claude-sonnet	1	submitted_passed	yes	Passed hidden tests	pass	1	n/a	51.1s
62bceaaa	cli/claude-sonnet	2	submitted_failed	yes	Graded	fail	1	n/a	52.1s
a66f85fd	cli/claude-sonnet	3	submitted_failed	yes	Graded	fail	0	n/a	36.1s
3fcba938	cli/claude-sonnet	4	submitted_failed	yes	Graded	fail	1	n/a	60.0s
70cf5fcc	cli/claude-sonnet	5	submitted_failed	yes	Graded	fail	1	n/a	56.3s
d13f1d97	cli/claude-sonnet	6	submitted_failed	yes	Graded	fail	1	n/a	53.8s
9d47ee0f	cli/claude-sonnet	7	submitted_failed	yes	Graded	fail	1	n/a	24.5s
a8cca2fc	cli/claude-sonnet	8	submitted_failed	yes	Graded	fail	0	n/a	43.8s
153192c8	cli/claude-sonnet	9	submitted_failed	yes	Graded	fail	1	n/a	55.6s
781eb674	cli/codex	0	submitted_passed	yes	Passed hidden tests	pass	1	n/a	76.9s
e71d8908	cli/codex	1	submitted_passed	yes	Passed hidden tests	pass	1	n/a	66.0s
2158c4f5	cli/codex	2	submitted_passed	yes	Passed hidden tests	pass	1	n/a	69.7s
8b6e987c	cli/codex	3	submitted_passed	yes	Passed hidden tests	pass	1	n/a	68.8s
fe164929	cli/codex	4	submitted_passed	yes	Passed hidden tests	pass	1	n/a	68.2s
854f7dae	cli/codex	5	submitted_passed	yes	Passed hidden tests	pass	1	n/a	65.6s
18d03c3a	cli/codex	6	submitted_passed	yes	Passed hidden tests	pass	1	n/a	65.9s
854f8510	cli/codex	7	submitted_passed	yes	Passed hidden tests	pass	1	n/a	66.5s
a111a1ee	cli/codex	8	submitted_passed	yes	Passed hidden tests	pass	1	n/a	65.7s
afdb6ff0	cli/codex	9	submitted_passed	yes	Passed hidden tests	pass	1	n/a	65.2s
ab1db90e	codestral-latest	0	submitted_failed	yes	Graded	fail	1	28 352	18.5s
6ed36c95	codestral-latest	1	submitted_failed	yes	Graded	fail	1	21 309	13.9s
496ff0a6	codestral-latest	2	submitted_failed	yes	Graded	fail	1	24 495	16.5s
e411d189	codestral-latest	3	submitted_failed	yes	Graded	fail	1	36 382	24.3s
acd6942b	codestral-latest	4	model_api_failure	no	Not graded — lane error	not graded	0	9 107	9.8s
ee5a8bdb	codestral-latest	5	submitted_failed	yes	Graded	fail	1	23 906	18.5s
3897e57e	codestral-latest	6	submitted_failed	yes	Graded	fail	1	18 536	14.1s
f23935c8	codestral-latest	7	submitted_failed	yes	Graded	fail	1	21 309	14.2s
9f5bdfb8	codestral-latest	8	submitted_failed	yes	Graded	fail	1	32 417	30.0s
f96b96b6	codestral-latest	9	submitted_failed	yes	Graded	fail	1	18 624	12.2s

Failure taxonomy

CLASS	COUNT	WHAT IT MEANS
submitted_failed	18	Submitted a change; the independent hidden tests still failed.
model_api_failure	1	The model endpoint failed to answer. Not a verdict about the model's ability.

Appendix B — evidence and artifact digests

RUN BUNDLE DIGEST

2e7f4171f049cced92f7a3aa

7ae13b2ee211449e96938e7e

6c437f30289db0ff

RUN ID

live-

20260901T235346Z

STARTED

2026-09-

01T23:53:46.828443

+00:00

FINISHED

2026-09-

02T00:18:47.040828

+00:00

Artifact digests

ARTIFACT	SHA-256
publicSummary	e4c4792d488475e66d9750cfca4484181c78985538f65d32605e587faad1cb4f
rollouts/153192c8-0932-4da9-bf81-111446fc32a5.json	6e633b09ebe3ee95b5cb4d387fe1356b3a02781bbd741f81e78a85a0e73728af
rollouts/18d03c3a-6629-4250-886e-c73986fdece7.json	ca658ebbd354c8f20e9cfba884d8e3098d4bbc990dfac08d14fc85ed4201f660
rollouts/2158c4f5-9a18-45cc-82e5-aa84a17c46a2.json	963bf7642556ae0d1a961a485be5acb2542e4942e1e79656bacc3463cc5d5c1e
rollouts/3897e57e-353e-402d-bb53-bc68259b1aa2.json	011530ec4616efa6f0b86d2b5a42b5aea82680b86f689f8af69217b80839ade8
rollouts/3fcba938-b19c-41d4-84fe-6ab6fec3c0a5.json	d75c04727307bfd1f50555fc0f892c36a01308a740eeef7706305cb59ce371f
rollouts/496ff0a6-9f53-4b50-8c9e-d1939301419c.json	c0ff5ca7fc1649be66606e9be02f4c8bfbec36583293f0e9aefcbd6c37c2d86c
rollouts/62bceaaa-2bf1-4067-8bef-6393b41f4d42.json	97da73eca72708cf8661d79aa0ac5351166de7f8599395f5253fdccf5627ec2b
rollouts/6ed36c95-91f9-4e3c-88da-fe17c010d584.json	5f72442b16936bf2145ef717e960d09429d418db43ad8c08de1c02dec6fec762
rollouts/70cf5fcc-7581-48bd-a1a9-53b4ce24a3a9.json	b0a5bffdff0754ff0609c73cc67e5a019e62e0fccf0e1c4630634818ba605ceb3
rollouts/72bcf971-7d06-42c0-9d6f-7f4ab5380c76.json	c95578b9f69a937f89a18d2cf3e70049a371791d10a614a6c87b7ca3ef918796
rollouts/781eb674-aa4a-4b2f-b1bc-502f4710a4fe.json	dbeff75b715e052925a9c65e68a9ea73059a54245df5c87809b82f95aaaaa9ff
rollouts/854f7dae-d4bc-441f-be3f-fc48eb324c71.json	77ee126e557bb11b7e88c38abae6bfcd9432c27cc8e9684ae49170cc69acfa1
rollouts/854f8510-cc34-4a46-a50f-771c323e3025.json	630b147f5a7c8a9a8ddc12bb638d1f92d8d54cae19ba7ec3b49a42c479a79045
rollouts/8b6e987c-8ccd-42a0-a076-30800bf7504b.json	51a15b141bcf2a6c0d377058ed74127fb4b2d1b29389638b1ef785e2919506fa
rollouts/8ec5820d-ef88-42c9-9c7f-e5cf9923de8d.json	9ab5fbbd4ef13e7c801e7f8c99343bf26ee7bb9d3f6f38cc077863b7386742b7
rollouts/9d47ee0f-e8c3-4eb8-a5cd-ba455bfb80a3.json	21e2ecff0304fbd74fa806d478ce3e2e4140a6cf2054fc03ea99a9ab3ab96cec
rollouts/9f5bdfb8-69a0-4b45-9c77-2f87b5217ca7.json	b298cff46d8ab12505f0510a561f6ca4a2a58c6dcdb5ce2af3ab83586acf0728c
rollouts/a111a1ee-81a7-4118-827f-2f8cf436e16c.json	fc831c3baaf9dba9f5100de58cbc210c43dd42997d6dfebcb088dd4ff75aff5c
rollouts/a66f85fd-827a-4d31-b407-6b0983c085ab.json	204a94e5e1d59d5faa5433a5b740546258c664c75b0e5b6490326b817bcb8aa5
rollouts/a8cca2fc-949e-4473-a79f-ac3027cab760.json	7775ac3ca099f5182d53326959308b3cd425565d5dfb7215e73f014bbe546537
rollouts/ab1db90e-5252-4819-b24b-0770692d26a9.json	d6bbea07e8570149b9f65ea3844ea0eeefc7dc5acc07253e271a2aed94b7628ed
rollouts/acd6942b-ff2c-478b-8fc5-36473631c0db.json	cdee6f5f35965fda4a7cda3d1eaab8aa1e6e22a13c0912a2afdd5811b35cc2b6
rollouts/afdb6ff0-9c18-494d-874a-43fce0ae7de7.json	dc31e5124a7738f55023ce1dea643d88d14cbfc18a1165e9003de6447be90f12
rollouts/d13f1d97-627d-4353-b467-a02faeff3156.json	b3e529077b975790be00a38c23a8f708f019c526165fb0ef68213cd2176a8731
rollouts/e411d189-390c-40ef-9318-3fbbbfbfb44d1.json	9fe97e4b4806ef362789ff8206bc560946bfc400b1d3046d5789534beb7c94ca
rollouts/e71d8908-6f2a-4362-9b6f-bdaf29b367f9.json	29ad3560920eadc86473f2efefdb0ee2e60c45ab030d4f68b6d95e282a8be013

ARTIFACT	SHA-256
rollouts/ee5a8bdb-0fe4-4909-a6f1-071b188dee55.json	d3ef142e69e1f9d296d05046099a73cb442e1df77c73463ac027717e0a749102
rollouts/f23935c8-c08a-442e-ba1d-ed39ea664b9d.json	9f3476ef0cbb5bca9e19b0e738ed19bfecb0238ef784503aa1cf0a46b9675c2f
rollouts/f96b96b6-e608-44d7-9f4b-8506d41d456d.json	c862c8429a774f5c5abab4442e2c2023466b8dd9eab98c117e8a1672e0209a8c
rollouts/fe164929-6cd7-4f59-acbc-ccb6ad02ddb2.json	97fc6c8b9bc91843bfa5958ffb79146d35a7c921e2f464093a7ff5736c1c6097
spec	0a225ad1c6f73d8104d73d1b5f0717e8516fc4c57f1957fc16447f8551663daa
summary	b77bc1cfc26a8c6e1848f0a23e0379695a67cf5a2db2f0543b8d6ade3a1bd1c1

Reproducing this run

Every rollout can be re-derived: the exam is frozen (fingerprint above), the grader ran in the recorded container image, and each rollout's full trace — every tool call, every result, every word — is in the evidence bundle at its rollouts/<id>.json path. A re-run uses the same job shape: the exam id, your endpoint, n, and the runtime. The record id names this run forever.

Appendix C — redacted transcripts

For each rollout: the tool calls it made, in order — which tool, on which of the model's own paths, whether the call succeeded, and how many characters came back. What a tool returned is not here, and neither is the submission: on this exam the submitted content is the graded answer, so no diff, no answer value and no model claim is printed. A request for the reference solution or the hidden tests is recorded as a request; the path it asked for is replaced by a label. See Appendix D, "What this report withholds and why".

Disclosure class applied to this exam: answer-key disclosure (ownership vvdex_proprietary — engine registry; family rag_grounding): the exam is VVDex evaluation IP and the graded output is an answer or a keyed value, so no submission content, diff, model claim or hidden-assertion name is published.

8ec5820d · cli:cli/claude-sonnet · seed 0 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	list_files	.	ok	287
1	read_file	questions.json	ok	883
2	run_command	—	ok	1 780
3	write_file	answers/answers.json	ok	97
4	run_tests	—	ok	181
5	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (57 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

72bcf971 · cli:cli/claude-sonnet · seed 1 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	list_files	.	ok	287
1	read_file	questions.json	ok	883
2	read_file	corpus/data-protection-policy-v3.txt	ok	293
3	read_file	corpus/retention-addendum-2026.txt	ok	142
4	read_file	corpus/org-chart-2026.txt	ok	183
5	read_file	corpus/incident-report-2026.txt	ok	169
6	read_file	corpus/security-whitepaper.txt	ok	142
7	read_file	corpus/employee-handbook.txt	ok	167
8	read_file	corpus/vendor-list.txt	ok	153
9	write_file	answers/answers.json	ok	69
10	run_tests	—	ok	181
11	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (45 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

62bceaaa · cli:cli/claude-sonnet · seed 2 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	list_files	.	ok	287
1	read_file	questions.json	ok	883
2	read_file	corpus/data-protection-policy-v3.txt	ok	293
3	read_file	corpus/retention-addendum-2026.txt	ok	142
4	read_file	corpus/org-chart-2026.txt	ok	183
5	read_file	corpus/incident-report-2026.txt	ok	169
6	read_file	corpus/security-whitepaper.txt	ok	142
7	read_file	corpus/employee-handbook.txt	ok	167
8	read_file	corpus/vendor-list.txt	ok	153
9	write_file	answers/answers.json	ok	69
10	run_tests	—	ok	181
11	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (45 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

a66f85fd · cli:cli/claude-sonnet · seed 3 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	list_files	.	ok	287
1	read_file	questions.json	ok	883
2	run_command	—	ok	132
3	run_command	—	ok	1 399
4	run_command	—	ok	758
5	invalid_tool_call	(none)	failed	120
6	run_tests	—	ok	181
7	submit	—	ok	85

Submission diff withheld by the public-disclosure policy. Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

3fcba938 · cli:cli/claude-sonnet · seed 4 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	list_files	.	ok	287
1	read_file	questions.json	ok	883
2	read_file	corpus/org-chart-2026.txt	ok	183
3	read_file	corpus/data-protection-policy-v3.txt	ok	293
4	read_file	corpus/retention-addendum-2026.txt	ok	142

STEP	TOOL	PATH	RESULT	CHARS
5	read_file	corpus/incident-report-2026.txt	ok	169
6	read_file	corpus/security-whitepaper.txt	ok	142
7	read_file	corpus/employee-handbook.txt	ok	167
8	read_file	corpus/vendor-list.txt	ok	153
9	write_file	answers/answers.json	ok	69
10	run_tests	—	ok	181
11	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (45 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

70cf5fcc · cli:cli/claude-sonnet · seed 5 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	list_files	.	ok	287
1	read_file	questions.json	ok	883
2	read_file	corpus/org-chart-2026.txt	ok	183
3	read_file	corpus/data-protection-policy-v3.txt	ok	293
4	read_file	corpus/retention-addendum-2026.txt	ok	142
5	read_file	corpus/incident-report-2026.txt	ok	169
6	read_file	corpus/security-whitepaper.txt	ok	142
7	read_file	corpus/employee-handbook.txt	ok	167
8	write_file	answers/answers.json	ok	69
9	run_tests	—	ok	181
10	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (45 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

d13f1d97 · cli:cli/claude-sonnet · seed 6 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	list_files	.	ok	287
1	read_file	questions.json	ok	883
2	read_file	corpus/data-protection-policy-v3.txt	ok	293
3	read_file	corpus/retention-addendum-2026.txt	ok	142
4	read_file	corpus/org-chart-2026.txt	ok	183
5	read_file	corpus/incident-report-2026.txt	ok	169
6	read_file	corpus/security-whitepaper.txt	ok	142
7	read_file	corpus/employee-handbook.txt	ok	167
8	read_file	corpus/vendor-list.txt	ok	153

STEP	TOOL	PATH	RESULT	CHARS
9	write_file	answers/answers.json	ok	69
10	run_tests	—	ok	181
11	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (4 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

9d47ee0f · cli:cli/claude-sonnet · seed 7 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	list_files	.	ok	287
1	read_file	questions.json	ok	883
2	run_command	—	ok	1 780
3	write_file	answers/answers.json	ok	69
4	run_tests	—	ok	181
5	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (4 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

a8cca2fc · cli:cli/claude-sonnet · seed 8 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	list_files	.	ok	287
1	read_file	questions.json	ok	883
2	run_command	—	ok	1 399
3	run_command	—	ok	758
4	run_command	—	ok	414
5	invalid_tool_call	(none)	failed	120
6	run_tests	—	ok	181
7	submit	—	ok	85

Submission diff withheld by the public-disclosure policy. Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

153192c8 · cli:cli/claude-sonnet · seed 9 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	list_files	.	ok	287
1	read_file	questions.json	ok	883
2	read_file	corpus/data-protection-policy-v3.txt	ok	293
3	read_file	corpus/org-chart-2026.txt	ok	183
4	read_file	corpus/retention-addendum-2026.txt	ok	142

STEP	TOOL	PATH	RESULT	CHARS
5	read_file	corpus/incident-report-2026.txt	ok	169
6	read_file	corpus/security-whitepaper.txt	ok	142
7	read_file	corpus/employee-handbook.txt	ok	167
8	read_file	corpus/vendor-list.txt	ok	153
9	write_file	answers/answers.json	ok	69
10	run_tests	—	ok	181
11	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (4 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

781eb674 · cli:cli/codex · seed 0 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	ISSUE.md	ok	1 953
1	list_files	.	ok	287
2	read_file	questions.json	ok	883
3	read_file	corpus/data-protection-policy-v3.txt	ok	293
4	read_file	corpus/data-protection-policy-v1.txt	ok	239
5	read_file	corpus/employee-handbook.txt	ok	167
6	read_file	corpus/incident-report-2026.txt	ok	169
7	read_file	corpus/org-chart-2026.txt	ok	183
8	read_file	corpus/retention-addendum-2026.txt	ok	142
9	read_file	corpus/security-whitepaper.txt	ok	142
10	read_file	corpus/vendor-list.txt	ok	153
11	write_file	answers/answers.json	ok	69
12	run_tests	—	ok	181
13	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (45 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

e71d8908 · cli:cli/codex · seed 1 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	list_files	.	ok	287
1	read_file	ISSUE.md	ok	1 953
2	read_file	questions.json	ok	883
3	read_file	corpus/data-protection-policy-v3.txt	ok	293
4	read_file	corpus/data-protection-policy-v1.txt	ok	239
5	read_file	corpus/employee-handbook.txt	ok	167

STEP	TOOL	PATH	RESULT	CHARS
6	read_file	corpus/incident-report-2026.txt	ok	169
7	read_file	corpus/org-chart-2026.txt	ok	183
8	read_file	corpus/retention-addendum-2026.txt	ok	142
9	read_file	corpus/security-whitepaper.txt	ok	142
10	read_file	corpus/vendor-list.txt	ok	153
11	write_file	answers/answers.json	ok	69
12	run_tests	—	ok	181
13	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (13 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

2158c4f5 · cli:cli/codex · seed 2 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	ISSUE.md	ok	1 953
1	list_files	.	ok	287
2	read_file	questions.json	ok	883
3	read_file	corpus/data-protection-policy-v1.txt	ok	239
4	read_file	corpus/data-protection-policy-v3.txt	ok	293
5	read_file	corpus/employee-handbook.txt	ok	167
6	read_file	corpus/incident-report-2026.txt	ok	169
7	read_file	corpus/org-chart-2026.txt	ok	183
8	read_file	corpus/retention-addendum-2026.txt	ok	142
9	read_file	corpus/security-whitepaper.txt	ok	142
10	read_file	corpus/vendor-list.txt	ok	153
11	write_file	answers/answers.json	ok	69
12	run_tests	—	ok	181
13	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (45 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

8b6e987c · cli:cli/codex · seed 3 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	ISSUE.md	ok	1 953
1	list_files	.	ok	287
2	read_file	questions.json	ok	883
3	read_file	corpus/data-protection-policy-v1.txt	ok	239
4	read_file	corpus/data-protection-policy-v3.txt	ok	293

STEP	TOOL	PATH	RESULT	CHARS
5	read_file	corpus/employee-handbook.txt	ok	167
6	read_file	corpus/incident-report-2026.txt	ok	169
7	read_file	corpus/org-chart-2026.txt	ok	183
8	read_file	corpus/retention-addendum-2026.txt	ok	142
9	read_file	corpus/security-whitepaper.txt	ok	142
10	read_file	corpus/vendor-list.txt	ok	153
11	write_file	answers/answers.json	ok	69
12	run_tests	—	ok	181
13	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (45 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

fe164929 · cli:cli/codex · seed 4 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	ISSUE.md	ok	1 953
1	list_files	.	ok	287
2	read_file	questions.json	ok	883
3	read_file	corpus/data-protection-policy-v1.txt	ok	239
4	read_file	corpus/data-protection-policy-v3.txt	ok	293
5	read_file	corpus/employee-handbook.txt	ok	167
6	read_file	corpus/incident-report-2026.txt	ok	169
7	read_file	corpus/org-chart-2026.txt	ok	183
8	read_file	corpus/retention-addendum-2026.txt	ok	142
9	read_file	corpus/security-whitepaper.txt	ok	142
10	read_file	corpus/vendor-list.txt	ok	153
11	write_file	answers/answers.json	ok	69
12	run_tests	—	ok	181
13	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (45 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

854f7dae · cli:cli/codex · seed 5 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	ISSUE.md	ok	1 953
1	list_files	.	ok	287
2	read_file	questions.json	ok	883
3	read_file	corpus/data-protection-policy-v1.txt	ok	239

STEP	TOOL	PATH	RESULT	CHARS
4	read_file	corpus/data-protection-policy-v3.txt	ok	293
5	read_file	corpus/employee-handbook.txt	ok	167
6	read_file	corpus/incident-report-2026.txt	ok	169
7	read_file	corpus/org-chart-2026.txt	ok	183
8	read_file	corpus/retention-addendum-2026.txt	ok	142
9	read_file	corpus/security-whitepaper.txt	ok	142
10	read_file	corpus/vendor-list.txt	ok	153
11	write_file	answers/answers.json	ok	69
12	run_tests	—	ok	181
13	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (13 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

18d03c3a · cli:cli/codex · seed 6 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	list_files	.	ok	287
1	read_file	ISSUE.md	ok	1 953
2	read_file	questions.json	ok	883
3	read_file	corpus/data-protection-policy-v1.txt	ok	239
4	read_file	corpus/data-protection-policy-v3.txt	ok	293
5	read_file	corpus/employee-handbook.txt	ok	167
6	read_file	corpus/incident-report-2026.txt	ok	169
7	read_file	corpus/org-chart-2026.txt	ok	183
8	read_file	corpus/retention-addendum-2026.txt	ok	142
9	read_file	corpus/security-whitepaper.txt	ok	142
10	read_file	corpus/vendor-list.txt	ok	153
11	write_file	answers/answers.json	ok	69
12	run_tests	—	ok	181
13	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (45 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

854f8510 · cli:cli/codex · seed 7 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	ISSUE.md	ok	1 953
1	list_files	.	ok	287
2	read_file	questions.json	ok	883

STEP	TOOL	PATH	RESULT	CHARS
3	read_file	corpus/data-protection-policy-v1.txt	ok	239
4	read_file	corpus/data-protection-policy-v3.txt	ok	293
5	read_file	corpus/employee-handbook.txt	ok	167
6	read_file	corpus/incident-report-2026.txt	ok	169
7	read_file	corpus/org-chart-2026.txt	ok	183
8	read_file	corpus/retention-addendum-2026.txt	ok	142
9	read_file	corpus/security-whitepaper.txt	ok	142
10	read_file	corpus/vendor-list.txt	ok	153
11	write_file	answers/answers.json	ok	69
12	run_tests	—	ok	181
13	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (45 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

all1alee · cli:cli/codex · seed 8 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	ISSUE.md	ok	1 953
1	list_files	.	ok	287
2	read_file	questions.json	ok	883
3	read_file	corpus/data-protection-policy-v1.txt	ok	239
4	read_file	corpus/data-protection-policy-v3.txt	ok	293
5	read_file	corpus/employee-handbook.txt	ok	167
6	read_file	corpus/incident-report-2026.txt	ok	169
7	read_file	corpus/org-chart-2026.txt	ok	183
8	read_file	corpus/retention-addendum-2026.txt	ok	142
9	read_file	corpus/security-whitepaper.txt	ok	142
10	read_file	corpus/vendor-list.txt	ok	153
11	write_file	answers/answers.json	ok	69
12	run_tests	—	ok	181
13	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (45 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

afdb6ff0 · cli:cli/codex · seed 9 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	ISSUE.md	ok	1 953
1	list_files	.	ok	287

STEP	TOOL	PATH	RESULT	CHARS
2	read_file	questions.json	ok	883
3	read_file	corpus/data-protection-policy-v1.txt	ok	239
4	read_file	corpus/data-protection-policy-v3.txt	ok	293
5	read_file	corpus/employee-handbook.txt	ok	167
6	read_file	corpus/incident-report-2026.txt	ok	169
7	read_file	corpus/org-chart-2026.txt	ok	183
8	read_file	corpus/retention-addendum-2026.txt	ok	142
9	read_file	corpus/security-whitepaper.txt	ok	142
10	read_file	corpus/vendor-list.txt	ok	153
11	write_file	answers/answers.json	ok	69
12	run_tests	—	ok	181
13	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (45 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

ab1db90e · mistral:codestral-latest · seed 0 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	ISSUE.md	ok	1 953
1	read_file	questions.json	ok	883
2	list_files	corpus	ok	249
3	read_file	corpus/data-protection-policy-v3.txt	ok	293
4	read_file	corpus/org-chart-2026.txt	ok	183
5	run_tests	—	ok	181
6	invalid_tool_call	(none)	failed	120
7	write_file	answers/answers.json	ok	69
8	run_tests	—	ok	181
9	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (45 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

6ed36c95 · mistral:codestral-latest · seed 1 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	ISSUE.md	ok	1 953
1	read_file	questions.json	ok	883
2	list_files	answers	ok	13
3	list_files	corpus	ok	249
4	read_file	corpus/data-protection-policy-v3.txt	ok	293

STEP	TOOL	PATH	RESULT	CHARS
5	write_file	answers/answers.json	ok	69
6	run_tests	—	ok	181
7	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (45 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

496ff0a6 · mistral:codestral-latest · seed 2 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	ISSUE.md	ok	1 953
1	read_file	questions.json	ok	883
2	read_file	—	failed	35
3	list_files	corpus	ok	249
4	read_file	corpus/retention-addendum-2026.txt	ok	142
5	read_file	corpus/data-protection-policy-v3.txt	ok	293
6	write_file	answers/answers.json	ok	69
7	run_tests	—	ok	181
8	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (45 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

e411d189 · mistral:codestral-latest · seed 3 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	ISSUE.md	ok	1 953
1	list_files	answers	ok	13
2	read_file	answers/.keep	ok	0
3	list_files	corpus	ok	249
4	read_file	corpus/data-protection-policy-v1.txt	ok	239
5	read_file	corpus/data-protection-policy-v3.txt	ok	293
6	read_file	corpus/employee-handbook.txt	ok	167
7	read_file	questions.json	ok	883
8	run_tests	—	ok	181
9	invalid_tool_call	(none)	failed	120
10	write_file	answers/answers.json	ok	69
11	run_tests	—	ok	181
12	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (45 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

acd6942b · mistral:codestral-latest · seed 4 · model_api_failure

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	ISSUE.md	ok	1 953
1	read_file	questions.json	ok	883
2	list_files	corpus	ok	249
3	invalid_tool_call	(none)	failed	120

Submission diff withheld by the public-disclosure policy. Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

ee5a8bdb · mistral:codestral-latest · seed 5 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	ISSUE.md	ok	1 953
1	read_file	questions.json	ok	883
2	list_files	answers	ok	13
3	read_file	answers/.keep	ok	0
4	list_files	corpus	ok	249
5	read_file	corpus/retention-addendum-2026.txt	ok	142
6	write_file	answers/answers.json	ok	69
7	run_tests	—	ok	181
8	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (45 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

3897e57e · mistral:codestral-latest · seed 6 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	ISSUE.md	ok	1 953
1	read_file	questions.json	ok	883
2	list_files	corpus	ok	249
3	read_file	corpus/retention-addendum-2026.txt	ok	142
4	write_file	answers/answers.json	ok	69
5	run_tests	—	ok	181
6	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (45 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

f23935c8 · mistral:codestral-latest · seed 7 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	ISSUE.md	ok	1 953
1	read_file	questions.json	ok	883
2	list_files	answers	ok	13
3	list_files	corpus	ok	249
4	read_file	corpus/data-protection-policy-v3.txt	ok	293
5	write_file	answers/answers.json	ok	69
6	run_tests	—	ok	181
7	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (45 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

9f5bdfb8 · mistral:codestral-latest · seed 8 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	ISSUE.md	ok	1 953
1	list_files	answers	ok	13
2	read_file	answers/.keep	ok	0
3	list_files	corpus	ok	249
4	read_file	corpus/data-protection-policy-v1.txt	ok	239
5	read_file	corpus/data-protection-policy-v3.txt	ok	293
6	write_file	answers/answers.json	ok	69
7	run_tests	—	ok	181
8	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (15 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

f96b96b6 · mistral:codestral-latest · seed 9 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	ISSUE.md	ok	1 953
1	read_file	questions.json	ok	883
2	list_files	corpus	ok	249
3	read_file	corpus/data-protection-policy-v3.txt	ok	293
4	write_file	answers/answers.json	ok	69
5	run_tests	—	ok	181
6	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (45 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is

the graded answer, so publishing an excerpt of it would publish the answer key.

Appendix D — metric definitions and methods

Statistics, stage funnel and efficiency in this document are derived at render time from the sealed rollout rows and their raw tool events: lane errors leave every capability denominator; a rollout's stages are booleans read off its own event sequence (inspected = a successful list or read, or a successful browser open/snapshot, edited = the workspace changed, tested = run_tests invoked, submitted = submit invoked, graded = the independent grader returned a verdict, passed = certified pass); absent telemetry is n/a. The record's own summary blocks are the values computed when it was sealed; the rows and events are the evidence.

Denominators

n counts evaluable rollouts only: attempts minus lane errors. A rollout the lane ended (API or runtime failure) is reported but never graded, so it is in no rate, interval or pass@k.

Interval method

Wilson score interval, not the normal approximation. The normal approximation collapses to a zero-width interval at 0 successes — it would print a 0/2 run as '0%, certain'. For k successes in n trials with $p = k/n$: centre = $(p + z^2/2n) / (1 + z^2/n)$, half-width = $z/(1 + z^2/n) \cdot \sqrt{(p(1-p)/n + z^2/4n^2)}$, clamped to [0, 1], $z = 1.9600$.

pass@k

Unbiased estimator: $\text{pass@k} = 1 - C(n-c, k) / C(n, k)$, for $k \leq n$ only — reported only up to the number of rollouts actually run. It answers "given k independent attempts at this task, how often would at least one pass" — nothing more.

Reading a failed rollout

A failed rollout is a measurement, not a defect report. Model API failures and runtime failures are classified separately precisely because they say nothing about the model's ability on this task.

Stage definitions

Inspected

Listed or read at least one file in the repository — or, on a browser exam, opened or snapshotted the page.

Edited

Wrote at least one file that actually changed on disk.

Ran tests

Invoked the visible-test tool at least once.

Submitted

Called submit, ending the rollout on its own terms.

Graded

The independent hidden grader ran on the submission and returned a verdict.

Passed hidden tests

The hidden grader returned exit 0 and the rollout was certified.

Long-horizon terms

commandsRun

Shell commands executed inside the model's own container, against the command budget the contract declares. Only exams that advertise run_command have any.

deadEnds

Steps that moved nothing: a re-read, a repeat of the previous action, or a call to a tool this exam does not have. Counted, not judged.

hitStepLimit

Rollouts that ran out of steps without submitting. On a long task this is the number that separates 'could not finish' from 'finished wrong'.

revisits

Re-reads of a slice the rollout had already been given. On a long task, re-reading is how a model compensates for losing its own earlier context.

stepsUsed

Tool calls the rollout actually spent, out of the step limit the contract declares. A rollout that submitted early spent fewer, and that is a result, not a shortfall.

Grounding and overclaim terms

groundingFailure

A file path the model wrote in its own reasoning that does not exist in the tree it was given, or that names a directory it is not allowed to see. Detected by matching source-file paths in the model's text against the agent box on disk — not by judging the prose.

nonexistent_path

The model named a source file that is not in the repository it was given. It was reasoning about a file that does not exist.

overclaim

The model stated the tests pass or the issue is fixed, and the independent hidden tests then failed. Counted per rollout: claimed-and-wrong, or not.

reReadChurn

How many times the model read a file it had already read in the same rollout. It is a cost and a coherence signal, not a failure: re-reading is sometimes right.

requested_hidden_path

The model named a path under the reference solution, the hidden grader or the hidden test directory. Those are not in its box; the request could not be served.

Failure classes

model_api_failure

The model endpoint failed to answer. Not a verdict about the model's ability.

submitted_failed

Submitted a change; the independent hidden tests still failed.

submitted_passed

Submitted a change that passed the independent hidden tests.

What this report withholds, and why

On this exam the model's submitted content IS the graded answer, so nothing derived from a submission is published: no diff excerpt, no answer value or answer JSON body, no reference or gold value, no expected hidden output, none of the model's own words, and not the names of the hidden assertions a rollout failed — a grader's name can carry the value it asserts. What Appendix C does publish for each rollout: its id, its lane, its seed, the tool it called at each step, the public-safe path it called the tool on, whether the call succeeded, how many characters came back, the stage it reached and its outcome. A rollout the grader accepted is reported as submitted and accepted — the numbers in this report are unaffected, because the verdict is the grader's, not the reader's.

Rule applied to this exam: answer-key disclosure (ownership vvdex_proprietary — engine registry; family rag_grounding): the exam is VVDex evaluation IP and the graded output is an answer or a keyed value, so no submission content, diff, model claim or hidden-assertion name is published. The rule is a property of the exam family, not a judgement made per document, and the same rule governs every edition of this report, the campaign report that embeds it as a chapter, this exam's public page and every published JSON derived from the same sealed record.

This exam revision is retired for future evaluation

Retired for future evaluation on 2026-09-02: post-run public disclosure of answer-bearing submission content. Historical sealed results on this revision are unaffected: the disclosure happened after those runs, so it could not have influenced them.

Attestation

This report was generated end to end by VVDex Forge from the sealed evidence of the evaluation run described above. Measurement values were populated from recorded artifacts, not manually entered. The digests in Verification let any holder of this file confirm that the report, the eval record and the run evidence are the ones this run produced — without contacting VVDex.

VVDex Forge · vvdexops.com	2026-09-02	vvdex-env 0.1.0
Issuer — certified evaluation environments	Date of issue	Engine · verifiers pin c521e4146026

No upstream repository: this exam is VVDex's own.
Post-seal benchmark. It does not stamp any certifying gate and did not rewrite the exam, the gold solution, the grader or the certified evidence.

— END OF REPORT · fr-20260902-2e7f4171 —