

Correct one order in an admin application, through a real browser

EVALUATION REPORT

2026-09-02

Evaluation scope:

- Claude Code CLI (Sonnet, subscription)
- Codex CLI (GPT, subscription)
- codestral-latest

Prepared by:

VVDex Forge — generated directly from sealed run evidence

fr-20260902-6e618c05 · vvdex.browser.order-desk-1 v0.1.0 · Independent model evaluation report

About this report

About VVDex Forge

VVDex Forge builds and certifies evaluation environments for AI agents. Every exam is sealed before any model sees it: the reference solution must pass, an empty submission must fail, the grader must reject near-misses, and the hidden tests never enter the agent's tree. Reports and the public catalog live at **vvdexops.com**.

Report integrity

This report is generated from run evidence; its digests verify authenticity. The sole test of a copy is the three-line procedure in Verification — a copy whose digests do not reproduce is not this report.

Scope

One exam · 30 rollout(s) · 40 tool steps and 9720s wall clock per rollout, stated in full in Evaluation configuration.

Boundary

Not a leaderboard. Not a general capability score. A post-seal benchmark: this run stamped no certification gate and rewrote nothing — not the exam, not the reference solution, not the grader, not the certified evidence.

Contents

01.	Executive summary	4
02.	Exam definition	5
03.	Certification evidence	6
04.	Evaluation configuration	7
05.	Model outcomes	9
06.	Stage funnel	10
07.	Behavioural analysis	11
08.	Efficiency & telemetry	12
09.	Findings	13
10.	Recommendations	14
11.	Verification	15

Appendices: A — every rollout · B — evidence digests · C — redacted transcripts · D — metric definitions · Attestation

01 Executive summary

MODELS TESTED 3	CERTIFIED PASSES 21 of 29 graded rollouts	SUBMITTED 29 of 29 graded rollouts that called submit	NOT GRADED — LANE ERRORS 1 of 30 attempts
---------------------------------------	--	--	--

Purpose

This report presents a controlled evaluation of 3 model lane(s) against the certified exam `vvdex.browser.order-desk-1` (v0.1.0, family browser). Every lane received the identical frozen workspace, tool contract and limits; grading ran afterwards, in isolation, against hidden tests no lane ever saw. The purpose is to state, with reproducible evidence, what each lane did and where it stopped.

Outcome

21 of 29

graded rollouts passed the independent hidden tests and were certified — 1 of 30 attempts not graded (lane errors)

21 of 29 graded rollouts passed. `cli/claude-sonnet`, `cli/codex`, `codestral-latest` completed the required browser workflow, submitted, and were accepted by the independent grader. The most common way to fail was `submitted_failed` — submitted a final application state that the independent grader rejected.

Key findings

- F-01 `cli/claude-sonnet`, `cli/codex`, `codestral-latest` passed the independent hidden tests
- F-02 `submitted_failed` × 8 (`codestral-latest`)
- F-03 `model_api_failure` × 1 (`codestral-latest`)

Each finding is stated in full, with evidence, in the Findings section; recommendations for acting on them follow it. Elapsed wall clock for the whole engagement: 3292.4s.

02 Exam definition

In scope. The ability of each configured model lane to complete this one certified task end to end — inspect and open the application, drive it through the declared browser actions to the required state, submit, and be accepted by the independent grader — within 40 tool steps and 9720s of wall clock, at 10 rollout(s) per lane. Behavioural measurements along the way (tool discipline, grounding, overclaim) are in scope because they are read from the recorded trace, not judged by another model.

Out of scope. General model quality, performance on other task families, prompt-tuning on the lane's behalf, and any comparison this run's sample size cannot support. A lane's API failure is reported as infrastructure, never as model capability.

The instrument. Exam `vvdex.browser.order-desk-1` was certified before this run: its reference solution passes, an empty submission fails, its grader distinguishes near-misses, and its sealed evidence is unchanged by this evaluation (see Certification evidence).

What was measured

EXAM	VERSION	FAMILY	ROLLOUTS
vvdex.browser.order-desk-1	0.1.0	browser	30 (10 per model)

Where the defect came from

This exam has no upstream repository to credit: the defect, the reference solution and the hidden suite are VVDex's own.

03 Certification evidence

4 of 5 certification pillars are recorded in this package, loaded from the receipts sealed into the exam and matching this run's exam fingerprint. What is recorded is shown with its real value; what is not is named below and is not claimed.

FROZEN BASELINE

FAIL

expected — an empty submission over 2 run(s) must not pass

→

GOLD / REFERENCE

PASS

expected — the reference solution over 2 run(s) must pass

GRADER CONTROLS

9 / 9

The grader accepted the reference and rejected every near-miss.

ATTACK PROBES

10 / 10 blocked

0 trivial exploit(s) found.

RUNTIME

sha256:679b36225a1f83238efba8c71b9cde3c24caaeacc9b682ae92819cb55eec328f

network policy: deny

Certification metadata is not included in this report package for:

- sealed evidence bundle

Benchmark results remain valid as run evidence; the pillar(s) above are not substantiated by this document and this report does not claim them.

What the package does carry: the exam fingerprint below is the contract digest this run was executed against — a holder of the certification bundle for this fingerprint can join the two independently.

EXAM FINGERPRINT

8d404ee5d5ab71456de3d3df0b8ca69271d8810767a694f71c6e45e72fc347bf

The contract digest this run was executed against.

04 Evaluation configuration

REPORT	fr-20260902-6e618c05
ISSUED	2026-09-02
SUBJECT EXAM	vvdex.browser.order-desk-1 · v0.1.0
FAMILY	browser
PREPARED BY	VVDex Forge engine vvdex-env 0.1.0
CLASSIFICATION	Independent model evaluation report — independently verifiable
CERTIFICATION	Exam certification: partially recorded (see Certification evidence)
STATUS	FINAL · report sealed against its own digest

Timeline

WHEN (UTC)	EVENT
2026-09-02 14:49:56	Evaluation run started
2026-09-02 15:44:48	Evaluation run finished
2026-09-02 15:44:48	Report generated from the sealed run evidence

Models under test

MODEL	PROVIDER	ENDPOINT	BILLING
Claude Code CLI (Sonnet, subscription)	cli	local runtime	operator subscription, flat — not billed per token
Codex CLI (GPT, subscription)	cli	local runtime	operator subscription, flat — not billed per token
codestral-latest	mistral	VVDex-configured	VVDex free-quota

A customer endpoint is recorded by origin only — never its port, its path, or its key.

Limits

RUNTIME docker	STEP LIMIT 40	WALL-CLOCK LIMIT 9720s	MAX OUTPUT TOKENS 2048
TEMPERATURE 0.2	RETRY POLICY none	COMMAND BUDGET 0	TOOLS OFFERED list_files, read_file, submit, browser_open, browser_snapshot, browser_click, browser_fill, browser_select, browser_press, browser_back

The evaluated model never saw the hidden tests or the reference solution, and could not change its result. Grading ran in a separate step, after submission, on a container with no network.

05 Model outcomes

Denominators. Attempts requested: 30. Graded (evaluable) rollouts: 29. Not graded — lane errors: 1. A lane error is a rollout the API or the runtime ended before the hidden grader ran: it is listed, excluded from every rate and pass@k, and never silently dropped. Passes are certified passes: hidden grader exit 0 with integrity verified.

MODEL	ATTEMPTS	GRADED	PASSED	MODEL FAILS	LANE ERRORS	SUBMITTED	MEAN TIME	TYPICAL DEEPEST STAGE
Claude Code CLI (Sonnet, subscription) cli/claude-sonnet	10	10	10 / 10	0	0	10 / 10	89.4s	Passed hidden tests
Codex CLI (GPT, subscription) cli/codex	10	10	10 / 10	0	0	10 / 10	125.2s	Passed hidden tests
codestral-latest codestral-latest	10	9	1 / 9	8	1	9 / 9	84.4s	Graded

Every rollout, with its own deepest stage and grade, is in Appendix A.

Success rate, confidence and pass@k

Every rate on this page is a measurement of a finite number of rollouts, so every rate carries the interval that measurement actually supports. n is the evaluable count — attempts minus rollouts the lane ended before grading; those are listed under "not graded" and sit in no rate, interval or pass@k.

MODEL	ATTEMPTS	N	NOT GRADED	PASSES	RATE	95% INTERVAL	SEEDS	PASS@1	PASS@3	PASS@5	PASS@10
cli/claude-sonnet	10	10	0	10	100.0%	[72%, 100%]	0–9	100.0%	100.0%	100.0%	100.0%
cli/codex	10	10	0	10	100.0%	[72%, 100%]	0–9	100.0%	100.0%	100.0%	100.0%
codestral-latest	10	9	1	1	11.1%	[2%, 43%]	0–9	11.1%	33.3%	55.6%	—

Interval and pass@k methods are stated in full in Appendix D. pass@k is shown for k = 1, 3, 5, 10; the sealed record carries every k up to n.

Model comparison

MODEL	PASS@1	95% INTERVAL	TOKENS	COST	MEAN ROLLOUT	OVERCLAIM RATE	TOP FAILURE CLASS
cli/claude-sonnet	100.0%	[72%, 100%]	n/a	n/a	89.4s	0.0%	none
cli/codex	100.0%	[72%, 100%]	n/a	n/a	125.2s	0.0%	none
codestral-latest	11.1%	[2%, 43%]	857 790	not billed	84.4s	0.0%	submitted_failed

Tokens and cost read n/a where a lane reports no telemetry (a flat-rate CLI) — an absent measurement, not zero. pass@1 and its interval are over evaluable rollouts only.

This is not a leaderboard. Separated by this run (intervals do not overlap): cli/claude-sonnet vs codestral-latest; cli/codex vs codestral-latest. Not separated (intervals overlap): cli/claude-sonnet vs cli/codex. Below the 10-rollout floor for reading a rate as a rate: codestral-latest (9) — read the count, not the percentage. The primary comparison in this report is the stage funnel below — where a model stops is a finding; a percentage over this many rollouts is not.

06 Stage funnel

The funnel is the primary reading of a run. Where a model stops says more than a single pass percentage does, especially at low N.

STAGE	CLI:CLI/CLAUDE-SONNET	CLI:CLI/CODEX	MISTRAL:CODESTRAL-LATEST
Inspected	10 / 10	10 / 10	9 / 9
Edited	N/A	N/A	N/A
Ran tests	N/A	N/A	N/A
Submitted	10 / 10	10 / 10	9 / 9
Graded	10 / 10	10 / 10	9 / 9
Passed hidden tests	10 / 10	10 / 10	1 / 9

Deepest stage reached

MODEL	DEEPEST (ANY ROLLOUT)	TYPICAL (MOST ROLLOUTS)	HIDDEN-TEST PASSES	NOT GRADED
cli:cli/claude-sonnet	Passed hidden tests	Passed hidden tests	10 / 10	0
cli:cli/codex	Passed hidden tests	Passed hidden tests	10 / 10	0
mistral:codestral-latest	Passed hidden tests	Graded	1 / 9	1

Stage definitions are in Appendix D. The matrix counts evaluable rollouts; lane errors are shown beside it, not inside it. The stages are not strictly nested: a model can submit without editing, and the matrix shows that rather than hiding it. N/A marks a stage this exam does not offer (Edited, Ran tests): not reached, not failed — not asked. Each rollout's own deepest stage is in Appendix A.

07 Behavioural analysis

Long-horizon behaviour

Counted off the recorded trace of each rollout. A dead end is a step that moved nothing: a re-read of a slice already returned, a repeat of the previous action, or a call to a tool this exam does not have.

MODEL	N	STEPS USED	RAN OUT OF STEPS	COMMANDS	CHANGED THE TREE	REFUSED (OVER BUDGET)	REVISITS	DEAD ENDS / ROLLOUT
cli:cli/claude-sonnet	10	18 / 40	0	0 / 0		0	0	3
cli:cli/codex	10	20 / 40	0	0 / 0		0	0	2
mistral:codestral-latest	10	22 / 40	0	0 / 0		0	0	1

Each column is defined in Appendix D.

Hallucination & overclaim

Three measurements, each with a stated definition. None of them is a judgement of the model's prose by another model — every number below comes from comparing what the model wrote against what was on disk, or against what the independent grader returned.

MODEL	N	GROUNDING FAILURES	KINDS	ROLLOUTS AFFECTED	OVERCLAIMS	OVERCLAIM RATE	RE-READS
cli:cli/claude-sonnet	10	0	none	0 / 10	0 / 10	0.0%	0
cli:cli/codex	10	0	none	0 / 10	0 / 10	0.0%	0
mistral:codestral-latest	10	0	none	0 / 10	0 / 10	0.0%	0

No grounding failures and no false pass claims were detected across these 30 rollout(s). No lane re-read material it had already been given.

Every term above is defined in Appendix D.

08 Efficiency & telemetry

This run has no per-token price attached — a local model, a free quota, or a customer endpoint we are not billed for — so spend is reported in tokens. Tokens are what was actually measured; a dollar figure here would be invented. Lanes run through a flat-rate local CLI report no token telemetry; their cells say n/a — an absent measurement, not zero. Rollout counts here are evaluable rollouts.

MODEL	ROLLOUTS	TOKENS	TOKENS / ROLLOUT	COST	PASSES	RATE
cli:cli/claude-sonnet	10	n/a	n/a	n/a – telemetry unavailable	10 / 10	100.0%
cli:cli/codex	10	n/a	n/a	n/a – telemetry unavailable	10 / 10	100.0%
mistral:codestral-latest	9	857 790	95 310	not billed	1 / 9	11.1%

09 Findings

Every finding below is derived from the recorded data of this run — none is an opinion.
Severity: PASS a lane succeeded · ATTENTION a capability gap · LANE / INFRA
infrastructure, not capability · NOTE a property of the measurement.

F-01

PASS

cli/claude-sonnet, cli/codex, codestral-latest passed the independent hidden tests

Severity: PASS

Evidence: Model outcomes · Stage funnel · Appendix A

21 of 29 graded rollouts completed the required browser workflow, submitted, and were accepted by the independent grader. This demonstrates the task is solvable under the published limits.

F-01 · derived from the recorded run data

F-02

ATTENTION

submitted_failed × 8 (codestral-latest)

Severity: ATTENTION

Evidence: Appendix A · Appendix C

Submitted a final application state that the independent grader rejected.

F-02 · derived from the recorded run data

F-03

LANE / INFRA

model_api_failure × 1 (codestral-latest)

Severity: LANE / INFRA

Evidence: Appendix A · Appendix C

The model endpoint failed to answer. Not a verdict about the model's ability. These rollouts are excluded from any capability reading.

F-03 · derived from the recorded run data

10 Recommendations

One recommendation per observed failure mode, addressed to the lane it affects. Each traces back to a finding above and to the raw trace in Appendix A.

R-01 PROBLEM
Submitted a final application state that the independent grader rejected.

AFFECTED LANES
codestral-latest

WHAT WE WOULD LOOK AT FIRST
The model drives the application and submits a final state the grader rejects — the named failing axes say which part of the required state was wrong (the target record, a collateral change, a skipped confirmation, an uncorrected input).

R-02 PROBLEM
The model endpoint failed to answer. Not a verdict about the model's ability.

AFFECTED LANES
codestral-latest

WHAT WE WOULD LOOK AT FIRST
The lane failed, not the model; these rollouts carry no signal about capability and are labeled accordingly.

11 Verification

RUN ID

live-20260902T144956Z

EXAM FINGERPRINT

8d404ee5d5ab71456de3d3df0b8ca69271d8810767a694f71c6e45e72fc347bf

The contract digest this run was executed against.

Tamper-evidence

Three digests, and how to check each one. If any of them does not reproduce, this page is not the page that was generated for this run.

THIS REPORT

3b0659c3a9103065ac1e2388f00f3f624cf566b0abbe4951272f0bf619be1afa

sha256 over the report's UTF-8 bytes with every occurrence of the report digest itself replaced by 64 '0' characters

EVAL RECORD

251730136fa0efbdd157d980d037447c967af4be4aaef1f04ce2c7a07a605bea

sha256 of fr-20260902-6e618c05.json as written beside this file.

RUN EVIDENCE BUNDLE

6e618c0542eb41677ef2556126365aea5538ce4f4a959245e1e09e1732ae8620

The digest the run's own bundle computed over itself.

CERTIFIED EVIDENCE SEAL

not recorded

The exam's sealed evidence, established before this run.

Measurement history

Previous run on this exam: fr-20260902-c05474ed. Model progress over time on this exam is the difference between that record and this one — same frozen tree, same hidden grader, same limits. Anything else that moved between them moved outside this exam.

Verifying it, in three lines

1. vvdex-env verify-report --report fr-20260902-6e618c05.html

2. or by hand: replace every occurrence of the report digest above with 64 zeros, then sha256 the file – it must equal that digest

3. shasum -a 256 fr-20260902-6e618c05.json must equal the eval-record digest, and the bundle digest must equal bundleSha256 in the run's bundle file

The full artifact-by-artifact digest table, and the reproduction protocol, are in Appendix B.

Appendix A — every rollout

ROLLOUT	MODEL	SEED	OUTCOME	SUBMITTED	DEEPEST STAGE	HIDDEN TESTS	FILES	TOKENS	ELAPSED	BROWSER BOX
532e1db3	cli/claude-sonnet	0	submitted_passed	yes	Passed hidden tests	pass	0	n/a	137.9s	docker sha256: c5587afc e772...
d641f68a	cli/claude-sonnet	1	submitted_passed	yes	Passed hidden tests	pass	0	n/a	96.5s	docker sha256: c5587afc e772...
b4f9ae08	cli/claude-sonnet	2	submitted_passed	yes	Passed hidden tests	pass	0	n/a	89.4s	docker sha256: c5587afc e772...
da5c8bda	cli/claude-sonnet	3	submitted_passed	yes	Passed hidden tests	pass	0	n/a	79.5s	docker sha256: c5587afc e772...
a116d4fa	cli/claude-sonnet	4	submitted_passed	yes	Passed hidden tests	pass	0	n/a	86.6s	docker sha256: c5587afc e772...
d6987896	cli/claude-sonnet	5	submitted_passed	yes	Passed hidden tests	pass	0	n/a	75.1s	docker sha256: c5587afc e772...
e52d496d	cli/claude-sonnet	6	submitted_passed	yes	Passed hidden tests	pass	0	n/a	80.6s	docker sha256: c5587afc e772...
7f40f719	cli/claude-sonnet	7	submitted_passed	yes	Passed hidden tests	pass	0	n/a	85.9s	docker sha256: c5587afc e772...
9b8db9ca	cli/claude-sonnet	8	submitted_passed	yes	Passed hidden tests	pass	0	n/a	79.3s	docker sha256: c5587afc e772...
38d7cace	cli/claude-sonnet	9	submitted_passed	yes	Passed hidden tests	pass	0	n/a	83.1s	docker sha256: c5587afc e772...
6bd00d3c	cli/codex	0	submitted_passed	yes	Passed hidden tests	pass	0	n/a	109.6s	docker sha256: c5587afc e772...
b3f86ec4	cli/codex	1	submitted_passed	yes	Passed hidden tests	pass	0	n/a	120.7s	docker sha256: c5587afc e772...
0d1dabed	cli/codex	2	submitted_passed	yes	Passed hidden tests	pass	0	n/a	147.3s	docker sha256: c5587afc e772...
e415a6c4	cli/codex	3	submitted_passed	yes	Passed hidden tests	pass	0	n/a	137.7s	docker sha256: c5587afc e772...
e771f4b2	cli/codex	4	submitted_passed	yes	Passed hidden tests	pass	0	n/a	130.1s	docker sha256: c5587afc e772...

ROLLOUT	MODEL	SEED	OUTCOME	SUBMITTED	DEEPEST STAGE	HIDDEN TESTS	FILES	TOKENS	ELAPSED	BROWSER BOX
c8c5c1f3	cli/codex	5	submitted_passed	yes	Passed hidden tests	pass	0	n/a	115.9s	docker sha256: c5587afc e772...
c20efdb4	cli/codex	6	submitted_passed	yes	Passed hidden tests	pass	0	n/a	125.2s	docker sha256: c5587afc e772...
6015f268	cli/codex	7	submitted_passed	yes	Passed hidden tests	pass	0	n/a	128.9s	docker sha256: c5587afc e772...
bbef3c65	cli/codex	8	submitted_passed	yes	Passed hidden tests	pass	0	n/a	121.4s	docker sha256: c5587afc e772...
3aa7a74b	cli/codex	9	submitted_passed	yes	Passed hidden tests	pass	0	n/a	114.7s	docker sha256: c5587afc e772...
cc461d13	codestral-latest	0	submitted_passed	yes	Passed hidden tests	pass	0	64 558	53.5s	docker sha256: c5587afc e772...
26427e2b	codestral-latest	1	submitted_failed	yes	Graded	fail	0	246 174	233.0s	docker sha256: c5587afc e772...
b4a95e09	codestral-latest	2	model_api_failure	no	Not graded — lane error	not graded	0	48 578	126.2s	docker sha256: c5587afc e772...
4df83755	codestral-latest	3	submitted_failed	yes	Graded	fail	0	31 078	31.6s	docker sha256: c5587afc e772...
3d19affd	codestral-latest	4	submitted_failed	yes	Graded	fail	0	110 075	77.3s	docker sha256: c5587afc e772...
7ee376ca	codestral-latest	5	submitted_failed	yes	Graded	fail	0	72 643	76.5s	docker sha256: c5587afc e772...
5ca4b95a	codestral-latest	6	submitted_failed	yes	Graded	fail	0	72 643	59.4s	docker sha256: c5587afc e772...
fe44cb82	codestral-latest	7	submitted_failed	yes	Graded	fail	0	96 982	67.5s	docker sha256: c5587afc e772...
86e8cb58	codestral-latest	8	submitted_failed	yes	Graded	fail	0	90 994	96.3s	docker sha256: c5587afc e772...
95636cf6	codestral-latest	9	submitted_failed	yes	Graded	fail	0	72 643	64.3s	docker sha256: c5587afc e772...

Failure taxonomy

CLASS	COUNT	WHAT IT MEANS
submitted_failed	8	Submitted a final application state that the independent grader rejected.
model_api_failure	1	The model endpoint failed to answer. Not a verdict about the model's ability.

Appendix B — evidence and artifact digests

RUN BUNDLE DIGEST

6e618c0542eb41677ef25561

26365aea5538ce4f4a959245

e1e09e1732ae8620

RUN ID

live-

20260902T144956Z

STARTED

2026-09-

02T14:49:56.156491

+00:00

FINISHED

2026-09-

02T15:44:48.577359

+00:00

Artifact digests

ARTIFACT	SHA-256
publicSummary	b2d0aedc0fb60c85fd4d643f4400fcfa02fd7774b09ccf4260be46c53ad93a56
rollouts/0d1dabed-f1a4-464d-8573-6298c4456eba.json	fb3d28ab7a5e999a43d50283dac14d6d6a892b208afcacf8d50718d410c1ca98
rollouts/26427e2b-2379-4a9f-891d-9fac9ff5783d.json	e480c99256a5a11253e10cd41035917212805b1af6cdf6aa7d78a61a5c9d1599
rollouts/38d7cace-518a-47c2-9678-318d7791b65f.json	a1da0a8316e97be2789ae1391260ac73ba67d97e7a4ebceb5c81cc37dded1469
rollouts/3aa7a74b-9c98-4c6e-814d-f759615903f5.json	559cc093d11a57d222156438f4e07f0e79914cece32ac6e3869394ecb8070bd
rollouts/3d19affd-3e09-4a9b-be1d-1bc702a87a51.json	3279d55e621bda9ecad8bcab0d2cb6584e833f1e1ae5392c25da0704f6349e94
rollouts/4df83755-7cf3-4d09-939f-44245af49dd0.json	e265d305de0acd9a29b6e9599c44785b51f5449dc375b9fb9ebae70773651b41
rollouts/532e1db3-0b8f-4765-808d-9341d00aee8f.json	53564320a7e4555050a2993db03fe77e59b6eb970fa3d7a94682a139b67e2e8b
rollouts/5ca4b95a-34c3-47a6-92aa-64987001944c.json	818ad5941d717436ba85cfc0f5d2729f61db4d525a5e0238cdc1ff7ab372ed98
rollouts/6015f268-cbda-4c89-a337-78b12414d270.json	eb1bea2ad651ae35bef64554e00b24a2eb4fa9a3f7ab3aaa6576036cc93d95a0
rollouts/6bd00d3c-1064-4c32-8a9c-121373dcfe98.json	ac2013c954cd90cc51718fabdc56d1a082a15a764d8c6f27ccd36cea99978256
rollouts/7ee376ca-7679-44f3-b428-5e3ad0d7e6f0.json	491b3dce54e59a85d2218733d08f9cb04d7afe7b3154473faae9b1cbbaa2cd6a
rollouts/7f40f719-4f86-464e-a1f9-cadfeeffa559.json	9e66bccf1b28e14a3621c81a47a593078f99a13e08b164b01139b66baa1deb74
rollouts/86e8cb58-b83d-4b8d-83f6-0003ac366fda.json	488baffcef7b005cb6fa47c621aa9fb6d26ecdac29560bb48cfe7f751ddd8f41
rollouts/95636cf6-da8f-48bd-9730-35964ac99fdf.json	92e873fb223f3f6b29ad197bc53140a9135955ec5a51b180ba4d791e1a60ec56
rollouts/9b8db9ca-f53f-4811-bfa9-1a2c618d8fd6.json	4593c01d72ce73d83cd476f4e3f89a964a0adcc118b699d1cfbeef6d24b1df22
rollouts/a116d4fa-48b8-47d4-83d9-0d300be62a8b.json	f19a9026c6aab52630d8f5a80ed314f78480ef2602410aa91bd5f3597ed6a5d1
rollouts/b3f86ec4-60d4-40f1-9482-b51bb3f0b2b7.json	957aa6a8fe5ef17f00a06ffdc0c5fa5c6c2b038f1b632a6cf74efb99baa4ee31b
rollouts/b4a95e09-f935-48d0-b79f-69c45c5dc341.json	4e7206169f78c5a8b99e8fd3c674d0a05455d9857d3ad9d5908f701519c7dc24
rollouts/b4f9ae08-95c3-41f4-848c-ed0edd82b145.json	75edd1a44f3b8b5ed6a07f37096a4ee55ca7a7b2fcd636c6ae79344674401451
rollouts/bbef3c65-99be-4a8a-83db-ea794108c317.json	a7a9815bf129d98da564ff9bf2a93634e959eea510cdb9cc1bc0a7a509df7b78
rollouts/c20efdb4-48a5-42f1-af8a-3fc6df3560ce.json	5e297e565b112b644a8f7363c6442ff7ab79f8f30893dcd351b6773b20101365
rollouts/c8c5c1f3-0c97-4022-8451-127d48f6953a.json	0956e3caf54e2596183f40906fd7a388cdab8015b12184528a71a80116121a46
rollouts/cc461d13-0b1b-40a2-8903-2d17c8ffe681.json	6eb1859ed1f57ca4cedb10414fd862ce818d2f3126c1681d4cd96a71ec4d2ea3
rollouts/d641f68a-d3a6-45d0-8228-0e8ff30ea176.json	f81f17004b13489e33073f2b92a786353dee80586510f96eb77d3efcee24cdca
rollouts/d6987896-0e83-4e30-832a-35ae28f5e0e8.json	76aab57eef5293c1270d57a9d370164ff1e5a00fbe27801edf9b700f51232fcf
rollouts/da5c8bda-9910-4ec7-893d-e4fda9fb617c.json	0da0f242816d1d1a3e254ae079a4d9a28367bf69010a80d86d284098214dfa6c

ARTIFACT	SHA-256
rollouts/e415a6c4-b070-41d7-be54-c68e5dcefaaa.json	4fcb78cb3f748628d95f62a10451d7ef2131dae262aaa123f12327f629b7a6f2
rollouts/e52d496d-c202-4c0f-9fbc-4c4328dcf0bc.json	8c1b0975dd0a6fbedbac9b54643b2f841eefd62ce07c48d0ddd7841cf444db81
rollouts/e771f4b2-674d-46db-8f64-ac620c935e3d.json	b5142923e8c761389e644df882ec7c94f14f3cfafbc1eed05cb4d92954030f8e
rollouts/fe44cb82-d2b1-4a7e-ac94-c9c086a65d01.json	23e1bc1d203a2f9db695033b15a360baa6e8d99659bd16a264f6f884becb306d
spec	79647b29130d3013f9061b0089ff43a26c9952ef5a5f3332fee8d5d702e755ae
summary	1fc2604a382037b95e760a3aeac37e70b462eb99d33a14dda05035ebf033c2db

Reproducing this run

Every rollout can be re-derived: the exam is frozen (fingerprint above), the grader ran in the recorded container image, and each rollout's full trace — every tool call, every result, every word — is in the evidence bundle at its rollouts/<id>.json path. A re-run uses the same job shape: the exam id, your endpoint, n, and the runtime. The record id names this run forever.

Appendix C — redacted transcripts

For each rollout: the tool calls it made, in order — which tool, on which of the model's own paths, whether the call succeeded, and how many characters came back. What a tool returned is not here, and neither is the submission: on this exam the submitted content is the graded answer, so no diff, no answer value and no model claim is printed. A request for the reference solution or the hidden tests is recorded as a request; the path it asked for is replaced by a label. See Appendix D, "What this report withholds and why".

Disclosure class applied to this exam: state disclosure (ownership vvdex_proprietary — engine registry; family browser): the exam is VVDex evaluation IP and the graded output is the application's own state and audit trail, so no submission content, diff or state is published.

532e1db3 · cli:cli/claude-sonnet · seed 0 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 975
1	browser_open	—	ok	62
2	browser_snapshot	—	ok	1 298
3	browser_click	—	ok	81
4	browser_snapshot	—	ok	1 328
5	browser_click	—	ok	65
6	browser_snapshot	—	ok	605
7	browser_fill	—	ok	36
8	browser_fill	—	ok	36
9	browser_fill	—	ok	36
10	browser_click	—	ok	65
11	browser_click	—	ok	65
12	browser_snapshot	—	ok	715
13	browser_fill	—	ok	36
14	browser_click	—	ok	65
15	browser_snapshot	—	ok	775
16	browser_click	—	ok	73
17	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (46 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

d641f68a · cli:cli/claude-sonnet · seed 1 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 975
1	browser_open	—	ok	62

STEP	TOOL	PATH	RESULT	CHARS
2	browser_snapshot	—	ok	1 298
3	browser_click	—	ok	81
4	browser_snapshot	—	ok	1 328
5	browser_click	—	ok	65
6	browser_snapshot	—	ok	605
7	browser_fill	—	ok	36
8	browser_fill	—	ok	36
9	browser_fill	—	ok	36
10	browser_click	—	ok	65
11	browser_click	—	ok	65
12	browser_snapshot	—	ok	715
13	browser_fill	—	ok	36
14	browser_click	—	ok	65
15	browser_snapshot	—	ok	775
16	browser_click	—	ok	73
17	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (46 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

b4f9ae08 · cli:cli/claude-sonnet · seed 2 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 975
1	browser_open	—	ok	62
2	browser_snapshot	—	ok	1 298
3	browser_click	—	ok	81
4	browser_snapshot	—	ok	1 328
5	browser_click	—	ok	65
6	browser_snapshot	—	ok	605
7	browser_fill	—	ok	36
8	browser_fill	—	ok	36
9	browser_fill	—	ok	36
10	browser_click	—	ok	65
11	browser_click	—	ok	65
12	browser_snapshot	—	ok	715
13	browser_fill	—	ok	36
14	browser_click	—	ok	65
15	browser_snapshot	—	ok	775

STEP	TOOL	PATH	RESULT	CHARS
16	browser_click	—	ok	73
17	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (46 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

da5c8bda · cli:cli/claude-sonnet · seed 3 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 975
1	browser_open	—	ok	62
2	browser_snapshot	—	ok	1 298
3	browser_click	—	ok	81
4	browser_snapshot	—	ok	1 328
5	browser_click	—	ok	65
6	browser_snapshot	—	ok	605
7	browser_fill	—	ok	36
8	browser_fill	—	ok	36
9	browser_fill	—	ok	36
10	browser_click	—	ok	65
11	browser_click	—	ok	65
12	browser_snapshot	—	ok	715
13	browser_fill	—	ok	36
14	browser_click	—	ok	65
15	browser_snapshot	—	ok	775
16	browser_click	—	ok	73
17	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (46 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

a116d4fa · cli:cli/claude-sonnet · seed 4 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 975
1	browser_open	—	ok	62
2	browser_snapshot	—	ok	1 298
3	browser_click	—	ok	81
4	browser_snapshot	—	ok	1 328
5	browser_click	—	ok	65
6	browser_snapshot	—	ok	605

STEP	TOOL	PATH	RESULT	CHARS
7	browser_fill	—	ok	36
8	browser_fill	—	ok	36
9	browser_fill	—	ok	36
10	browser_click	—	ok	65
11	browser_click	—	ok	65
12	browser_snapshot	—	ok	715
13	browser_fill	—	ok	36
14	browser_click	—	ok	65
15	browser_snapshot	—	ok	775
16	browser_click	—	ok	73
17	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (46 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

d6987896 · cli:cli/claude-sonnet · seed 5 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 975
1	browser_open	—	ok	62
2	browser_snapshot	—	ok	1 298
3	browser_click	—	ok	81
4	browser_snapshot	—	ok	1 328
5	browser_click	—	ok	65
6	browser_snapshot	—	ok	605
7	browser_fill	—	ok	36
8	browser_fill	—	ok	36
9	browser_fill	—	ok	36
10	browser_click	—	ok	65
11	browser_click	—	ok	65
12	browser_snapshot	—	ok	715
13	browser_fill	—	ok	36
14	browser_click	—	ok	65
15	browser_snapshot	—	ok	775
16	browser_click	—	ok	73
17	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (46 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

e52d496d · cli:cli/claude-sonnet · seed 6 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 975
1	browser_open	—	ok	62
2	browser_snapshot	—	ok	1 298
3	browser_click	—	ok	81
4	browser_snapshot	—	ok	1 328
5	browser_click	—	ok	65
6	browser_snapshot	—	ok	605
7	browser_fill	—	ok	36
8	browser_fill	—	ok	36
9	browser_fill	—	ok	36
10	browser_click	—	ok	65
11	browser_click	—	ok	65
12	browser_snapshot	—	ok	715
13	browser_fill	—	ok	36
14	browser_click	—	ok	65
15	browser_snapshot	—	ok	775
16	browser_click	—	ok	73
17	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (46 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

7f40f719 · cli:cli/claude-sonnet · seed 7 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 975
1	browser_open	—	ok	62
2	browser_snapshot	—	ok	1 298
3	browser_click	—	ok	81
4	browser_snapshot	—	ok	1 328
5	browser_click	—	ok	65
6	browser_snapshot	—	ok	605
7	browser_fill	—	ok	36
8	browser_fill	—	ok	36
9	browser_fill	—	ok	36
10	browser_click	—	ok	65
11	browser_click	—	ok	65
12	browser_snapshot	—	ok	715

STEP	TOOL	PATH	RESULT	CHARS
13	browser_fill	—	ok	36
14	browser_click	—	ok	65
15	browser_snapshot	—	ok	775
16	browser_click	—	ok	73
17	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (46 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

9b8db9ca · cli:cli/claude-sonnet · seed 8 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 975
1	browser_open	—	ok	62
2	browser_snapshot	—	ok	1 298
3	browser_click	—	ok	81
4	browser_snapshot	—	ok	1 328
5	browser_click	—	ok	65
6	browser_snapshot	—	ok	605
7	browser_fill	—	ok	36
8	browser_fill	—	ok	36
9	browser_fill	—	ok	36
10	browser_click	—	ok	65
11	browser_click	—	ok	65
12	browser_snapshot	—	ok	715
13	browser_fill	—	ok	36
14	browser_click	—	ok	65
15	browser_snapshot	—	ok	775
16	browser_click	—	ok	73
17	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (46 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

38d7cace · cli:cli/claude-sonnet · seed 9 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 975
1	browser_open	—	ok	62
2	browser_snapshot	—	ok	1 298
3	browser_click	—	ok	81

STEP	TOOL	PATH	RESULT	CHARS
4	browser_snapshot	—	ok	1 328
5	browser_click	—	ok	65
6	browser_snapshot	—	ok	605
7	browser_fill	—	ok	36
8	browser_fill	—	ok	36
9	browser_fill	—	ok	36
10	browser_click	—	ok	65
11	browser_click	—	ok	65
12	browser_snapshot	—	ok	715
13	browser_fill	—	ok	36
14	browser_click	—	ok	65
15	browser_snapshot	—	ok	775
16	browser_click	—	ok	73
17	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (46 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

6bd00d3c · cli:cli/codex · seed 0 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 975
1	browser_open	—	ok	62
2	browser_snapshot	—	ok	1 298
3	browser_click	—	ok	81
4	browser_snapshot	—	ok	1 328
5	browser_click	—	ok	65
6	browser_snapshot	—	ok	605
7	browser_fill	—	ok	36
8	browser_fill	—	ok	36
9	browser_fill	—	ok	36
10	browser_fill	—	ok	36
11	browser_click	—	ok	65
12	browser_click	—	ok	65
13	browser_snapshot	—	ok	715
14	browser_fill	—	ok	36
15	browser_click	—	ok	65
16	browser_snapshot	—	ok	775
17	browser_click	—	ok	73

STEP	TOOL	PATH	RESULT	CHARS
18	browser_snapshot	—	ok	635
19	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (48 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

b3f86ec4 · cli:cli/codex · seed 1 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 975
1	browser_open	—	ok	62
2	browser_snapshot	—	ok	1 298
3	browser_fill	—	ok	36
4	browser_click	—	ok	87
5	browser_snapshot	—	ok	621
6	browser_click	—	ok	65
7	browser_snapshot	—	ok	605
8	browser_fill	—	ok	36
9	browser_fill	—	ok	36
10	browser_fill	—	ok	36
11	browser_click	—	ok	65
12	browser_click	—	ok	65
13	browser_snapshot	—	ok	715
14	browser_fill	—	ok	36
15	browser_click	—	ok	65
16	browser_snapshot	—	ok	775
17	browser_click	—	ok	73
18	browser_snapshot	—	ok	635
19	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (48 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

0d1dabed · cli:cli/codex · seed 2 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 975
1	browser_open	—	ok	62
2	browser_snapshot	—	ok	1 298
3	browser_fill	—	ok	36
4	browser_click	—	ok	87

STEP	TOOL	PATH	RESULT	CHARS
5	browser_snapshot	—	ok	621
6	browser_click	—	ok	65
7	browser_snapshot	—	ok	605
8	browser_fill	—	ok	36
9	browser_fill	—	ok	36
10	browser_fill	—	ok	36
11	browser_click	—	ok	65
12	browser_click	—	ok	65
13	browser_snapshot	—	ok	715
14	browser_fill	—	ok	36
15	browser_click	—	ok	65
16	browser_snapshot	—	ok	775
17	browser_click	—	ok	73
18	browser_snapshot	—	ok	635
19	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (48 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

e415a6c4 · cli:cli/codex · seed 3 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 975
1	browser_open	—	ok	62
2	browser_snapshot	—	ok	1 298
3	browser_fill	—	ok	36
4	browser_click	—	ok	87
5	browser_snapshot	—	ok	621
6	browser_click	—	ok	65
7	browser_snapshot	—	ok	605
8	browser_fill	—	ok	36
9	browser_fill	—	ok	36
10	browser_fill	—	ok	36
11	browser_click	—	ok	65
12	browser_snapshot	—	ok	614
13	browser_click	—	ok	65
14	browser_snapshot	—	ok	715
15	browser_fill	—	ok	36
16	browser_click	—	ok	65

STEP	TOOL	PATH	RESULT	CHARS
17	browser_snapshot	—	ok	775
18	browser_click	—	ok	73
19	browser_snapshot	—	ok	635
20	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (49 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

e771f4b2 · cli:cli/codex · seed 4 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 975
1	browser_open	—	ok	62
2	browser_snapshot	—	ok	1 298
3	browser_fill	—	ok	36
4	browser_click	—	ok	87
5	browser_snapshot	—	ok	621
6	browser_click	—	ok	65
7	browser_snapshot	—	ok	605
8	browser_fill	—	ok	36
9	browser_fill	—	ok	36
10	browser_fill	—	ok	36
11	browser_click	—	ok	65
12	browser_snapshot	—	ok	614
13	browser_click	—	ok	65
14	browser_snapshot	—	ok	715
15	browser_fill	—	ok	36
16	browser_click	—	ok	65
17	browser_snapshot	—	ok	775
18	browser_click	—	ok	73
19	browser_snapshot	—	ok	635
20	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (49 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

c8c5c1f3 · cli:cli/codex · seed 5 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 975
1	browser_open	—	ok	62

STEP	TOOL	PATH	RESULT	CHARS
2	browser_snapshot	—	ok	1 298
3	browser_fill	—	ok	36
4	browser_click	—	ok	87
5	browser_snapshot	—	ok	621
6	browser_click	—	ok	65
7	browser_snapshot	—	ok	605
8	browser_fill	—	ok	36
9	browser_fill	—	ok	36
10	browser_fill	—	ok	36
11	browser_click	—	ok	65
12	browser_click	—	ok	65
13	browser_snapshot	—	ok	715
14	browser_fill	—	ok	36
15	browser_click	—	ok	65
16	browser_snapshot	—	ok	775
17	browser_click	—	ok	73
18	browser_snapshot	—	ok	635
19	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (48 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

c20efdb4 · cli:cli/codex · seed 6 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 975
1	browser_open	—	ok	62
2	browser_snapshot	—	ok	1 298
3	browser_fill	—	ok	36
4	browser_click	—	ok	87
5	browser_snapshot	—	ok	621
6	browser_click	—	ok	65
7	browser_snapshot	—	ok	605
8	browser_fill	—	ok	36
9	browser_fill	—	ok	36
10	browser_fill	—	ok	36
11	browser_click	—	ok	65
12	browser_click	—	ok	65
13	browser_snapshot	—	ok	715

STEP	TOOL	PATH	RESULT	CHARS
14	browser_fill	—	ok	36
15	browser_click	—	ok	65
16	browser_snapshot	—	ok	775
17	browser_click	—	ok	73
18	browser_snapshot	—	ok	635
19	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (48 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

6015f268 · cli:cli/codex · seed 7 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 975
1	browser_open	—	ok	62
2	browser_snapshot	—	ok	1 298
3	browser_fill	—	ok	36
4	browser_click	—	ok	87
5	browser_snapshot	—	ok	621
6	browser_click	—	ok	65
7	browser_snapshot	—	ok	605
8	browser_fill	—	ok	36
9	browser_fill	—	ok	36
10	browser_fill	—	ok	36
11	browser_click	—	ok	65
12	browser_snapshot	—	ok	614
13	browser_click	—	ok	65
14	browser_snapshot	—	ok	715
15	browser_fill	—	ok	36
16	browser_snapshot	—	ok	712
17	browser_click	—	ok	65
18	browser_snapshot	—	ok	775
19	browser_click	—	ok	73
20	browser_snapshot	—	ok	635
21	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (50 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

bbef3c65 · cli:cli/codex · seed 8 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 975
1	browser_open	—	ok	62
2	browser_snapshot	—	ok	1 298
3	browser_fill	—	ok	36
4	browser_click	—	ok	87
5	browser_snapshot	—	ok	621
6	browser_click	—	ok	65
7	browser_snapshot	—	ok	605
8	browser_fill	—	ok	36
9	browser_fill	—	ok	36
10	browser_fill	—	ok	36
11	browser_click	—	ok	65
12	browser_click	—	ok	65
13	browser_snapshot	—	ok	715
14	browser_fill	—	ok	36
15	browser_click	—	ok	65
16	browser_snapshot	—	ok	775
17	browser_click	—	ok	73
18	browser_snapshot	—	ok	635
19	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (48 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

3aa7a74b · cli:cli/codex · seed 9 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	read_file	TASK.md	ok	1 975
1	browser_open	—	ok	62
2	browser_snapshot	—	ok	1 298
3	browser_fill	—	ok	36
4	browser_click	—	ok	87
5	browser_snapshot	—	ok	621
6	browser_click	—	ok	65
7	browser_snapshot	—	ok	605
8	browser_fill	—	ok	36
9	browser_fill	—	ok	36
10	browser_fill	—	ok	36

STEP	TOOL	PATH	RESULT	CHARS
11	browser_click	—	ok	65
12	browser_click	—	ok	65
13	browser_snapshot	—	ok	715
14	browser_fill	—	ok	36
15	browser_click	—	ok	65
16	browser_snapshot	—	ok	775
17	browser_click	—	ok	73
18	browser_snapshot	—	ok	635
19	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (48 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

cc461d13 · mistral:codestral-latest · seed 0 · submitted_passed

STEP	TOOL	PATH	RESULT	CHARS
0	browser_open	—	ok	62
1	browser_snapshot	—	ok	1 298
2	browser_click	—	ok	65
3	browser_snapshot	—	ok	594
4	browser_back	—	ok	61
5	browser_snapshot	—	ok	1 298
6	browser_click	—	ok	81
7	browser_snapshot	—	ok	1 328
8	browser_click	—	ok	65
9	browser_snapshot	—	ok	605
10	browser_fill	—	ok	36
11	browser_fill	—	ok	36
12	browser_fill	—	ok	36
13	browser_click	—	ok	65
14	browser_click	—	ok	65
15	browser_snapshot	—	ok	715
16	browser_fill	—	ok	36
17	browser_click	—	ok	65
18	browser_snapshot	—	ok	775
19	browser_click	—	ok	73
20	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (50 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

26427e2b · mistral:codestral-latest · seed 1 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	browser_open	—	ok	62
1	browser_snapshot	—	ok	1 298
2	browser_click	—	ok	65
3	browser_snapshot	—	ok	594
4	browser_back	—	ok	61
5	browser_snapshot	—	ok	1 298
6	browser_click	—	ok	65
7	browser_snapshot	—	ok	600
8	browser_back	—	ok	61
9	browser_snapshot	—	ok	1 298
10	browser_click	—	ok	65
11	browser_snapshot	—	ok	599
12	browser_back	—	ok	61
13	browser_snapshot	—	ok	1 298
14	browser_click	—	ok	65
15	browser_snapshot	—	ok	599
16	browser_fill	—	ok	36
17	browser_fill	—	ok	36
18	browser_fill	—	ok	36
19	browser_fill	—	ok	36
20	browser_click	—	ok	65
21	browser_snapshot	—	ok	602
22	browser_click	—	ok	65
23	browser_snapshot	—	ok	703
24	browser_fill	—	ok	36
25	browser_click	—	ok	65
26	browser_snapshot	—	ok	781
27	browser_click	—	ok	73
28	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (59 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

b4a95e09 · mistral:codestral-latest · seed 2 · model_api_failure

STEP	TOOL	PATH	RESULT	CHARS
0	browser_open	—	ok	62
1	browser_snapshot	—	ok	1 298
2	browser_click	—	ok	65
3	browser_snapshot	—	ok	594
4	browser_back	—	ok	61
5	browser_snapshot	—	ok	1 298
6	browser_click	—	ok	81
7	browser_snapshot	—	ok	1 328
8	browser_click	—	ok	65
9	browser_snapshot	—	ok	605
10	browser_fill	—	ok	36
11	browser_fill	—	ok	36
12	browser_click	—	ok	65
13	browser_snapshot	—	ok	616
14	browser_click	—	ok	65
15	browser_snapshot	—	ok	717
16	browser_fill	—	ok	36

Submission diff withheld by the public-disclosure policy (24 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

4df83755 · mistral:codestral-latest · seed 3 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	browser_open	—	ok	62
1	browser_snapshot	—	ok	1 298
2	browser_click	—	ok	65
3	browser_snapshot	—	ok	594
4	browser_fill	—	ok	36
5	browser_fill	—	ok	36
6	browser_fill	—	ok	36
7	browser_click	—	ok	65
8	browser_snapshot	—	ok	706
9	browser_fill	—	ok	36
10	browser_click	—	ok	65
11	browser_snapshot	—	ok	774
12	browser_click	—	ok	73
13	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (38 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

3d19affd · mistral:codestral-latest · seed 4 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	browser_open	—	ok	62
1	browser_snapshot	—	ok	1 298
2	browser_click	—	ok	65
3	browser_snapshot	—	ok	594
4	browser_back	—	ok	61
5	browser_snapshot	—	ok	1 298
6	browser_click	—	ok	65
7	browser_snapshot	—	ok	600
8	browser_back	—	ok	61
9	browser_snapshot	—	ok	1 298
10	browser_back	—	ok	44
11	browser_open	—	ok	62
12	browser_snapshot	—	ok	1 298
13	browser_click	—	ok	81
14	browser_snapshot	—	ok	1 328
15	browser_click	—	ok	65
16	browser_snapshot	—	ok	605
17	browser_fill	—	ok	36
18	browser_fill	—	ok	36
19	browser_click	—	ok	65
20	browser_snapshot	—	ok	616
21	browser_click	—	ok	65
22	browser_snapshot	—	ok	717
23	browser_fill	—	ok	36
24	browser_click	—	ok	65
25	browser_snapshot	—	ok	777
26	browser_click	—	ok	73
27	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (54 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

7ee376ca · mistral:codestral-latest · seed 5 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	browser_open	—	ok	62
1	browser_snapshot	—	ok	1 298
2	browser_click	—	ok	65
3	browser_snapshot	—	ok	594
4	browser_back	—	ok	61
5	browser_snapshot	—	ok	1 298
6	browser_click	—	ok	81
7	browser_snapshot	—	ok	1 328
8	browser_click	—	ok	65
9	browser_snapshot	—	ok	605
10	browser_fill	—	ok	36
11	browser_fill	—	ok	36
12	browser_click	—	ok	65
13	browser_snapshot	—	ok	616
14	browser_click	—	ok	65
15	browser_snapshot	—	ok	717
16	browser_fill	—	ok	36
17	browser_click	—	ok	65
18	browser_snapshot	—	ok	777
19	browser_click	—	ok	73
20	browser_snapshot	—	ok	637
21	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (48 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

5ca4b95a · mistral:codestral-latest · seed 6 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	browser_open	—	ok	62
1	browser_snapshot	—	ok	1 298
2	browser_click	—	ok	65
3	browser_snapshot	—	ok	594
4	browser_back	—	ok	61
5	browser_snapshot	—	ok	1 298
6	browser_click	—	ok	81
7	browser_snapshot	—	ok	1 328
8	browser_click	—	ok	65

STEP	TOOL	PATH	RESULT	CHARS
9	browser_snapshot	—	ok	605
10	browser_fill	—	ok	36
11	browser_fill	—	ok	36
12	browser_click	—	ok	65
13	browser_snapshot	—	ok	616
14	browser_click	—	ok	65
15	browser_snapshot	—	ok	717
16	browser_fill	—	ok	36
17	browser_click	—	ok	65
18	browser_snapshot	—	ok	777
19	browser_click	—	ok	73
20	browser_snapshot	—	ok	637
21	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (48 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

fe44cb82 · mistral:codestral-latest · seed 7 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	browser_open	—	ok	62
1	browser_snapshot	—	ok	1 298
2	browser_click	—	ok	65
3	browser_snapshot	—	ok	594
4	browser_back	—	ok	61
5	browser_snapshot	—	ok	1 298
6	browser_click	—	ok	65
7	browser_snapshot	—	ok	600
8	browser_back	—	ok	61
9	browser_snapshot	—	ok	1 298
10	browser_back	—	ok	44
11	browser_open	—	ok	62
12	browser_snapshot	—	ok	1 298
13	browser_click	—	ok	81
14	browser_snapshot	—	ok	1 328
15	browser_click	—	ok	65
16	browser_snapshot	—	ok	605
17	browser_fill	—	ok	36
18	browser_fill	—	ok	36

STEP	TOOL	PATH	RESULT	CHARS
19	browser_click	—	ok	65
20	browser_snapshot	—	ok	616
21	browser_click	—	ok	65
22	browser_snapshot	—	ok	717
23	browser_fill	—	ok	36
24	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (31 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

86e8cb58 · mistral:codestral-latest · seed 8 · submitted_failed

STEP	TOOL	PATH	RESULT	CHARS
0	browser_open	—	ok	62
1	browser_snapshot	—	ok	1 298
2	browser_click	—	ok	65
3	browser_snapshot	—	ok	594
4	browser_back	—	ok	61
5	browser_snapshot	—	ok	1 298
6	browser_click	—	ok	81
7	browser_snapshot	—	ok	1 328
8	browser_click	—	ok	65
9	browser_snapshot	—	ok	605
10	browser_fill	—	ok	36
11	browser_fill	—	ok	36
12	browser_click	—	ok	65
13	browser_snapshot	—	ok	616
14	browser_click	—	ok	65
15	browser_snapshot	—	ok	717
16	browser_fill	—	ok	36
17	browser_click	—	ok	65
18	browser_snapshot	—	ok	777
19	browser_click	—	ok	73
20	browser_snapshot	—	ok	637
21	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (48 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

STEP	TOOL	PATH	RESULT	CHARS
0	browser_open	—	ok	62
1	browser_snapshot	—	ok	1 298
2	browser_click	—	ok	65
3	browser_snapshot	—	ok	594
4	browser_back	—	ok	61
5	browser_snapshot	—	ok	1 298
6	browser_click	—	ok	81
7	browser_snapshot	—	ok	1 328
8	browser_click	—	ok	65
9	browser_snapshot	—	ok	605
10	browser_fill	—	ok	36
11	browser_fill	—	ok	36
12	browser_click	—	ok	65
13	browser_snapshot	—	ok	616
14	browser_click	—	ok	65
15	browser_snapshot	—	ok	717
16	browser_fill	—	ok	36
17	browser_click	—	ok	65
18	browser_snapshot	—	ok	777
19	browser_click	—	ok	73
20	browser_snapshot	—	ok	637
21	submit	—	ok	85

Submission diff withheld by the public-disclosure policy (48 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

Appendix D — metric definitions and methods

Statistics, stage funnel and efficiency in this document are derived at render time from the sealed rollout rows and their raw tool events: lane errors leave every capability denominator; a rollout's stages are booleans read off its own event sequence (inspected = a successful list or read, or a successful browser open/snapshot, edited = the workspace changed, tested = run_tests invoked, submitted = submit invoked, graded = the independent grader returned a verdict, passed = certified pass); absent telemetry is n/a. The record's own summary blocks are the values computed when it was sealed; the rows and events are the evidence.

Denominators

n counts evaluable rollouts only: attempts minus lane errors. A rollout the lane ended (API or runtime failure) is reported but never graded, so it is in no rate, interval or pass@k.

Interval method

Wilson score interval, not the normal approximation. The normal approximation collapses to a zero-width interval at 0 successes — it would print a 0/2 run as '0%, certain'. For k successes in n trials with $p = k/n$: centre = $(p + z^2/2n) / (1 + z^2/n)$, half-width = $z/(1 + z^2/n) \cdot \sqrt{(p(1-p)/n + z^2/4n^2)}$, clamped to [0, 1], $z = 1.9600$.

pass@k

Unbiased estimator: $\text{pass@k} = 1 - C(n-c, k) / C(n, k)$, for $k \leq n$ only — reported only up to the number of rollouts actually run. It answers "given k independent attempts at this task, how often would at least one pass" — nothing more.

Reading a failed rollout

A failed rollout is a measurement, not a defect report. Model API failures and runtime failures are classified separately precisely because they say nothing about the model's ability on this task.

Stage definitions

Inspected

Listed or read at least one file in the repository — or, on a browser exam, opened or snapshotted the page.

Edited

Wrote at least one file that actually changed on disk. Not applicable to this exam: the capability this stage measures is not offered by its tool contract, so it renders N/A.

Ran tests

Invoked the visible-test tool at least once. Not applicable to this exam: the capability this stage measures is not offered by its tool contract, so it renders N/A.

Submitted

Called submit, ending the rollout on its own terms.

Graded

The independent hidden grader ran on the submission and returned a verdict.

Passed hidden tests

The hidden grader returned exit 0 and the rollout was certified.

Long-horizon terms

commandsRun

Shell commands executed inside the model's own container, against the command budget the contract declares. Only exams that advertise `run_command` have any.

deadEnds

Steps that moved nothing: a re-read, a repeat of the previous action, or a call to a tool this exam does not have. Counted, not judged.

hitStepLimit

Rollouts that ran out of steps without submitting. On a long task this is the number that separates 'could not finish' from 'finished wrong'.

revisits

Re-reads of a slice the rollout had already been given. On a long task, re-reading is how a model compensates for losing its own earlier context.

stepsUsed

Tool calls the rollout actually spent, out of the step limit the contract declares. A rollout that submitted early spent fewer, and that is a result, not a shortfall.

Grounding and overclaim terms

groundingFailure

A file path the model wrote in its own reasoning that does not exist in the tree it was given, or that names a directory it is not allowed to see. Detected by matching source-file paths in the model's text against the agent box on disk — not by judging the prose.

nonexistent_path

The model named a source file that is not in the repository it was given. It was reasoning about a file that does not exist.

overclaim

The model stated the tests pass or the issue is fixed, and the independent hidden tests then failed. Counted per rollout: claimed-and-wrong, or not.

reReadChurn

How many times the model read a file it had already read in the same rollout. It is a cost and a coherence signal, not a failure: re-reading is sometimes right.

requested_hidden_path

The model named a path under the reference solution, the hidden grader or the hidden test directory. Those are not in its box; the request could not be served.

Failure classes

model_api_failure

The model endpoint failed to answer. Not a verdict about the model's ability.

submitted_failed

Submitted a final application state that the independent grader rejected.

submitted_passed

Submitted a final application state that the independent grader accepted.

What this report withholds, and why

On this exam the model's submitted content IS the graded answer, so nothing derived from a submission is published: no diff excerpt, no answer value or answer JSON body, no reference or gold value, no expected hidden output, none of the model's own words, and not the names of the hidden assertions a rollout failed — a grader's name can carry the value it asserts. What Appendix C does publish for each rollout: its id, its lane, its seed, the tool it called at each step, the public-safe path it called the tool on, whether the call succeeded, how many characters came back, the stage it reached and its outcome. A rollout the grader accepted is reported as submitted and accepted — the numbers in this report are unaffected, because the verdict is the grader's, not the reader's.

Rule applied to this exam: state disclosure (ownership `vvdex_proprietary` — engine registry; family browser): the exam is VVDex evaluation IP and the graded output is the application's own state and audit trail, so no submission content, diff or state is published. The rule is a property of the exam family, not a judgement made per document, and the same rule

governs every edition of this report, the campaign report that embeds it as a chapter, this exam's public page and every published JSON derived from the same sealed record.

Attestation

This report was generated end to end by VVDex Forge from the sealed evidence of the evaluation run described above. Measurement values were populated from recorded artifacts, not manually entered. The digests in Verification let any holder of this file confirm that the report, the eval record and the run evidence are the ones this run produced — without contacting VVDex.

VVDex Forge · vvdexops.com	2026-09-02	vvdex-env 0.1.0
Issuer — certified evaluation environments	Date of issue	Engine · verifiers pin c521e4146026

No upstream repository: this exam is VVDex's own.
Post-seal benchmark. It does not stamp any certifying gate and did not rewrite the exam, the gold solution, the grader or the certified evidence.

— END OF REPORT · fr-20260902-6e618c05 —