

# Streaming items must equal what the full parse would have built

## EVALUATION REPORT

2026-09-02

*Evaluation scope:*

**Claude Code CLI (Sonnet, subscription)**

**Codex CLI (GPT, subscription)**

**codestral-latest**

*Prepared by:*

**VVDex Forge** — generated directly from sealed run evidence

---

fr-20260902-6616b188 · public.swe.martinblech-xmltodict-issue-257 v0.1.0 · Independent model evaluation report

# About this report

---

## About VVDex Forge

VVDex Forge builds and certifies evaluation environments for AI agents. Every exam is sealed before any model sees it: the reference solution must pass, an empty submission must fail, the grader must reject near-misses, and the hidden tests never enter the agent's tree. Reports and the public catalog live at **vvdexops.com**.

## Report integrity

This report is generated from run evidence; its digests verify authenticity. The sole test of a copy is the three-line procedure in Verification — a copy whose digests do not reproduce is not this report.

## Scope

One exam · 30 rollout(s) · 40 tool steps and 9720s wall clock per rollout, stated in full in Evaluation configuration.

## Boundary

Not a leaderboard. Not a general capability score. A post-seal benchmark: this run stamped no certification gate and rewrote nothing — not the exam, not the reference solution, not the grader, not the certified evidence.

# Contents

|     |                        |   |
|-----|------------------------|---|
| 01. | Executive summary      |   |
| 02. | Exam definition        | 4 |
| 03. | Certification evidence | 5 |
| 04. | Model outcomes         | 6 |
| 05. | Stage funnel           | 7 |
| 06. | Findings               | 7 |
| 07. | Recommendations        | 8 |
| 08. | Verification           | 8 |

Followed by: companion evidence report reference - Attestation

## 01 Executive summary

MODELS TESTED

3

CERTIFIED PASSES

19

of 30 graded rollouts

SUBMITTED

30 of 30

graded rollouts that called submit

NOT GRADED — LANE ERRORS

0

of 30 attempts

### Purpose

This report presents a controlled evaluation of 3 model lane(s) against the certified exam public.swe.martinblech-xmltodict-issue-257 (v0.1.0, family public\_oss\_swe). Every lane received the identical frozen workspace, tool contract and limits; grading ran afterwards, in isolation, against hidden tests no lane ever saw. The purpose is to state, with reproducible evidence, what each lane did and where it stopped.

Outcome

19 of 30

graded rollouts passed the independent hidden tests and were certified

19 of 30 graded rollouts passed. cli/claude-sonnet, cli/codex completed the full loop (inspect, edit, run the visible tests, submit) and the sealed grader accepted the submission. The most common way to fail was submitted\_failed — submitted a change; the independent hidden tests still failed.

Key findings

F-01 cli/claude-sonnet, cli/codex passed the independent hidden tests

F-02 submitted\_failed × 11 (cli/claude-sonnet, codestral-latest)

Each finding is stated in full, with evidence, in the Findings section; recommendations for acting on them follow it. Elapsed wall clock for the whole engagement: 2327.3s.

02 Exam definition

**In scope.** The ability of each configured model lane to complete this one certified task end to end — inspect the workspace, produce a correct edit, verify it with the visible tests, and submit — within 40 tool steps and 9720s of wall clock, at 10 rollout(s) per lane. Behavioural measurements along the way (tool discipline, grounding, overclaim) are in scope because they are read from the recorded trace, not judged by another model.

**Out of scope.** General model quality, performance on other task families, prompt-tuning on the lane's behalf, and any comparison this run's sample size cannot support. A lane's API failure is reported as infrastructure, never as model capability.

**The instrument.** Exam public.swe.martinblech-xmltodict-issue-257 was certified before this run: its reference solution passes, an empty submission fails, its grader distinguishes near-misses, and its sealed evidence is unchanged by this evaluation (see Certification evidence).

What was measured

EXAM

public.swe.martinblech-xmltodict-issue-257

VERSION

0.1.0

FAMILY

public\_oss\_swe

ROLLOUTS

30 (10 per model)

Where the defect came from

Source: martinblech/xmltodict (MIT) · issue  
https://github.com/martinblech/xmltodict/issues/257 · upstream fix  
https://github.com/martinblech/xmltodict/pull/421

The defect and its upstream fix are public and unmodified. The exam around them — the frozen parent tree, the independently written hidden suite, the control probes and the isolation — is VVDex's, and it is what this report measures.

03 Certification evidence

4 of 5 certification pillars are recorded in this package, loaded from the receipts sealed into the exam and matching this run's exam fingerprint. What is recorded is shown with its real value; what is not is named below and is not claimed.

FROZEN BASELINE

**FAIL**

expected — an empty submission over 2 run(s) must not pass

→

GOLD / REFERENCE

**PASS**

expected — the reference solution over 2 run(s) must pass

GRADER CONTROLS

**7 / 7**

The grader accepted the reference and rejected every near-miss.

ATTACK PROBES

**10 / 10 blocked**

0 trivial exploit(s) found.

RUNTIME

**sha256:679b36225a1f83238efba8c71b9cde3c24caaeacc9b682ae92819cb55eec328f**

network policy: deny

Certification metadata is not included in this report package for:

- sealed evidence bundle

Benchmark results remain valid as run evidence; the pillar(s) above are not substantiated by this document and this report does not claim them.

What the package does carry: the exam fingerprint below is the contract digest this run was executed against — a holder of the certification bundle for this fingerprint can join the two independently.

EXAM FINGERPRINT

**aa41977120e6ecdddb5c4fc603e8ab23cb272726a86750119969ff4dfc90e297**

The contract digest this run was executed against.

04 Model outcomes

**Denominators.** Attempts requested: 30. Graded (evaluable) rollouts: 30. Not graded — lane errors: 0. A lane error is a rollout the API or the runtime ended before the hidden grader ran: it is listed, excluded from every rate and pass@k, and never silently dropped. Passes are certified passes: hidden grader exit 0 with integrity verified.

| MODEL                                                       | ATTEMPTS | GRADED | PASSED  | MODEL FAILS | LANE ERRORS | SUBMITTED | MEAN TIME | TYPICAL DEEPEST STAGE |
|-------------------------------------------------------------|----------|--------|---------|-------------|-------------|-----------|-----------|-----------------------|
| Claude Code CLI (Sonnet, subscription)<br>cli/claude-sonnet | 10       | 10     | 9 / 10  | 1           | 0           | 10 / 10   | 159.8s    | Passed hidden tests   |
| Codex CLI (GPT, subscription)<br>cli/codex                  | 10       | 10     | 10 / 10 | 0           | 0           | 10 / 10   | 42.4s     | Passed hidden tests   |
| codestral-latest<br>codestral-latest                        | 10       | 10     | 0 / 10  | 10          | 0           | 10 / 10   | 11.5s     | Graded                |

Every rollout, with its own deepest stage and grade, is in Appendix A.

Success rate, confidence and pass@k

Every rate on this page is a measurement of a finite number of rollouts, so every rate carries the interval that measurement actually supports. n is the evaluable count — attempts minus rollouts the lane ended before grading; those are listed under "not graded" and sit in no rate, interval or pass@k.

| MODEL             | ATTEMPTS | N  | NOT GRADED | PASSES | RATE   | 95% INTERVAL | SEEDS | PASS@1 | PASS@3 | PASS@5 | PASS@10 |
|-------------------|----------|----|------------|--------|--------|--------------|-------|--------|--------|--------|---------|
| cli/claude-sonnet | 10       | 10 | 0          | 9      | 90.0%  | [60%, 98%]   | 0–9   | 90.0%  | 100.0% | 100.0% | 100.0%  |
| cli/codex         | 10       | 10 | 0          | 10     | 100.0% | [72%, 100%]  | 0–9   | 100.0% | 100.0% | 100.0% | 100.0%  |
| codestral-latest  | 10       | 10 | 0          | 0      | 0.0%   | [0%, 28%]    | 0–9   | 0.0%   | 0.0%   | 0.0%   | 0.0%    |

Interval and pass@k methods are stated in full in Appendix D. pass@k is shown for k = 1, 3, 5, 10; the sealed record carries every k up to n.

Model comparison

| MODEL             | PASS@1 | 95% INTERVAL | TOKENS  | COST       | MEAN ROLLOUT | OVERCLAIM RATE | TOP FAILURE CLASS |
|-------------------|--------|--------------|---------|------------|--------------|----------------|-------------------|
| cli/claude-sonnet | 90.0%  | [60%, 98%]   | n/a     | n/a        | 159.8s       | 0.0%           | submitted_failed  |
| cli/codex         | 100.0% | [72%, 100%]  | n/a     | n/a        | 42.4s        | 0.0%           | none              |
| codestral-latest  | 0.0%   | [0%, 28%]    | 172 890 | not billed | 11.5s        | 0.0%           | submitted_failed  |

Tokens and cost read n/a where a lane reports no telemetry (a flat-rate CLI) — an absent measurement, not zero. pass@1 and its interval are over evaluable rollouts only.

This is not a leaderboard. Separated by this run (intervals do not overlap): cli/claude-sonnet vs codestral-latest; cli/codex vs codestral-latest. Not separated (intervals overlap): cli/claude-sonnet vs cli/codex. The primary comparison in this report is the stage funnel below — where a model stops is a finding; a percentage over this many rollouts is not.

05 Stage funnel

The funnel is the primary reading of a run. Where a model stops says more than a single pass percentage does, especially at low N.

| STAGE               | CLI:CLI/CLAUDE-SONNET | CLI:CLI/CODEX | MISTRAL:CODESTRAL-LATEST |
|---------------------|-----------------------|---------------|--------------------------|
| Inspected           | 10 / 10               | 10 / 10       | 10 / 10                  |
| Edited              | 9 / 10                | 10 / 10       | 10 / 10                  |
| Ran tests           | 10 / 10               | 10 / 10       | 10 / 10                  |
| Submitted           | 10 / 10               | 10 / 10       | 10 / 10                  |
| Graded              | 10 / 10               | 10 / 10       | 10 / 10                  |
| Passed hidden tests | 9 / 10                | 10 / 10       | 0 / 10                   |

Deepest stage reached

| MODEL                    | DEEPEST (ANY ROLLOUT) | TYPICAL (MOST ROLLOUTS) | HIDDEN-TEST PASSES | NOT GRADED |
|--------------------------|-----------------------|-------------------------|--------------------|------------|
| cli:cli/claude-sonnet    | Passed hidden tests   | Passed hidden tests     | 9 / 10             | 0          |
| cli:cli/codex            | Passed hidden tests   | Passed hidden tests     | 10 / 10            | 0          |
| mistral:codestral-latest | Graded                | Graded                  | 0 / 10             | 0          |

Stage definitions are in Appendix D. The matrix counts evaluable rollouts; lane errors are shown beside it, not inside it. The stages are not strictly nested: a model can submit without editing, and the matrix shows that rather than hiding it. Each rollout's own deepest stage is in Appendix A.

06 Findings

Every finding below is derived from the recorded data of this run — none is an opinion.  
Severity: PASS a lane succeeded · ATTENTION a capability gap · LANE / INFRA infrastructure, not capability · NOTE a property of the measurement.

F-01

PASS

cli/claude-sonnet, cli/codex passed the independent hidden tests

Severity: PASS

Evidence: Model outcomes · Stage funnel · Appendix A

19 of 30 graded rollouts completed the full loop (inspect, edit, run the visible tests, submit) and the sealed grader accepted the submission. This demonstrates the task is solvable under the published limits.

F-01 · derived from the recorded run data

F-02

ATTENTION

submitted\_failed × 11 (cli/claude-sonnet, codestral-latest)

Severity: ATTENTION

Evidence: Appendix A · Appendix C

Submitted a change; the independent hidden tests still failed.

F-02 · derived from the recorded run data

07 Recommendations

One recommendation per observed failure mode, addressed to the lane it affects. Each traces back to a finding above and to the raw trace in Appendix A.

R-01 PROBLEM

Submitted a change; the independent hidden tests still failed.

AFFECTED LANES

cli/claude-sonnet, codestral-latest

WHAT WE WOULD LOOK AT FIRST

The model completes the loop and produces a plausible near-miss — the named failing tests are the exact behavioral cases distinguishing it from the reference fix.

08 Verification

RUN ID

live-20260902T001847Z

EXAM FINGERPRINT

aa41977120e6ecddb5c4fc603e8ab23cb272726a8675011  
9969ff4dfc90e297

The contract digest this run was executed against.

Tamper-evidence

Three digests, and how to check each one. If any of them does not reproduce, this page is not the page that was generated for this run.

THIS REPORT

ac4e25e8f4b54f9f881b2fbd1226d  
711789e219fa894902d545c41fd62  
2f86cb

sha256 over the report's UTF-8 bytes with every occurrence of the report digest itself replaced by 64 '0' characters

EVAL RECORD

02b4f83ee15251a5808a83153c9c4  
590936af23c58cc94d172556e9f98  
3d1909

sha256 of fr-20260902-6616b188.json as written beside this file.

RUN EVIDENCE BUNDLE

6616b188f015f52e157bdfb257e8d  
0f819e1f6cb9bcd3b024b340ac077  
1e0fd8

The digest the run's own bundle computed over itself.

#### CERTIFIED EVIDENCE SEAL

##### **not recorded**

The exam's sealed evidence,  
established before this run.

## Measurement history

This is the first recorded run on this exam. There is nothing to compare it against yet, and a single measurement is not a trend.

## Verifying it, in three lines

1. `vvdex-env verify-report --report fr-20260902-6616b188.html`
2. or by hand: replace every occurrence of the report digest above with 64 zeros, then sha256 the file – it must equal that digest
3. `shasum -a 256 fr-20260902-6616b188.json` must equal the eval-record digest, and the bundle digest must equal `bundleSha256` in the run's bundle file

The full artifact-by-artifact digest table, and the reproduction protocol, are in Appendix B.

This is the executive report. The full verifiable evidence report for the same run — every rollout, artifact digests, redacted transcripts and methods — is `fr-20260902-6616b188.pdf`, generated from exactly the same sealed record.

Attestation

This report was generated end to end by VVDex Forge from the sealed evidence of the evaluation run described above. Measurement values were populated from recorded artifacts, not manually entered. The digests in Verification let any holder of this file confirm that the report, the eval record and the run evidence are the ones this run produced — without contacting VVDex.

|                                            |                   |                                     |
|--------------------------------------------|-------------------|-------------------------------------|
| <b>VVDex Forge · vvdexops.com</b>          | <b>2026-09-02</b> | <b>vvdex-env 0.1.0</b>              |
| Issuer — certified evaluation environments | Date of issue     | Engine · verifiers pin c521e4146026 |

Source: martinblech/xmltodict (MIT)  
Post-seal benchmark. It does not stamp any certifying gate and did not rewrite the exam, the gold solution, the grader or the certified evidence.

— END OF REPORT · fr-20260902-6616b188 —