

GPT-5.5 annotation evaluation — ten attempts on each of four tasks

4 fixed tasks · 40 sealed attempt records · 40 graded attempts · 25 certified passes

CURRENT

CAMPAIGN REPORT — EXECUTIVE EDITION

2026-09-06

Evaluation scope:

cli/codex-gpt-5.5

Prepared by:

VVDex Forge — regenerated from verified sealed records only

fc-ec21239c00aa · 40 record(s) · every value recomputable

Contents

01.	Executive summary	3
02.	Per-exam results	5
03.	Environment families	6
04.	Campaign aggregate across these four exams	7
05.	Harness status	8
06.	Cells excluded from capability	9
07.	What the outcomes mean	10
08.	What this report withholds, and why	11
09.	Evidence	12
10.	Attestation	14

01 Executive summary

Release status: CURRENT. This is the campaign the published featured index carries; its exams' public pages state these numbers.

FIXED TASKS

4

SEALED RECORDS

40

MODEL LANES

1

ATTEMPTS

40

GRADED ATTEMPTS

40

CERTIFIED PASSES

25

LANE ERRORS (NOT GRADED)

0

WITHHELD / INVALID

0

Purpose

This report presents a controlled evaluation of 1 model lane(s) against 4 fixed certified annotation tasks, ten sealed attempt records on each task. For each task, every lane received that task's frozen workspace, tool contract and limits; grading ran afterwards, in isolation, against hidden tests no lane ever saw; and every rollout carries three separate outcomes — grader, integrity, certified — of which only the certified outcome is counted. The executive edition omits record-level chapters; the full report renders each sealed attempt record as a chapter derived from the same record as its standalone report. The full report authenticates the chapter bytes, while the standalone report carries its own digest.

Annotation campaign scope. This campaign evaluates four fixed annotation tasks with ten sealed attempt records on each task (40 sealed attempts in total). The sealed tasks cover image, structured, text, video inputs. Video was evaluated through the sampled still frames recorded in every video attempt; these results do not establish native-video capability. Audio is N/A because no audio task appears in this campaign's sealed records. CLI temperature and stochastic-seed controls are not established by these records; attempt indices are bookkeeping. Token use and per-attempt cost are N/A because every recorded lane result marks token telemetry unavailable and no cost measurement is attached. Quality measurements are reported per task and remain separate from the strict certified-pass gate.

25

certified pass(es) across 40 graded attempts (40 attempts, 0 lane error(s) not graded) in 4 cells — counted only where the grader accepted the submission AND containment was verified AND provenance is complete AND no integrity invariant or verdict-bearing harness suspicion applies.

fc-ec21239c00aa · generated from 40 verified sealed record(s) · EXECUTIVE EDITION

Key findings across the campaign

- `image-object-1` · `cli/codex-gpt-5.5` · **ATTENTION** — 5 certified passes / 10 graded attempts; 5 model failures; 0 excluded (0 lane errors, 0 withheld, 0 invalid).
- `robot-video-1` · `cli/codex-gpt-5.5` · **ATTENTION** — 0 certified passes / 10 graded attempts; 10 model failures; 0 excluded (0 lane errors, 0 withheld, 0 invalid).
- `structured-data-1` · `cli/codex-gpt-5.5` · **PASS** — 10 certified passes / 10 graded attempts; 0 model failures; 0 excluded (0 lane errors, 0 withheld, 0 invalid).
- `text-labeling-1` · `cli/codex-gpt-5.5` · **PASS** — 10 certified passes / 10 graded attempts; 0 model failures; 0 excluded (0 lane errors, 0 withheld, 0 invalid).

02 Per-exam results

MODEL LANE	ANNOTATION · IMAGE-OBJECT-1 VVDEX.ANNOTATI ON.IMAGE-OBJEC T-1	ANNOTATION · ROBOT-VIDEO-1 VVDEX.ANNOTAT ION.ROBOT-VID EO-1	ANNOTATION · STRUCTURED-DATA-1 VVDEX.ANNOTATIO N.STRUCTURED-DA TA-1	ANNOTATION · TEXT- LABELING-1 VVDEX.ANNOTATI ON.TEXT-LABELI NG-1
cli/codex-gpt-5.5	5/10	0/10	10/10	10/10

Each cell is *certified passes / graded attempts* for that lane on that exam. An attempt the lane ended before grading is reported beside the cell — 1/9 + 1 lane error, never 1/10 — and is outside the denominator; withheld and invalid attempts are reported the same way. No cell here carries a percentage: these are separate exams, not one score.

03 Environment families

This campaign covers four exam(s) drawn from one environment family(ies). The sentence in each row is the family's own description of what the exam tests; nothing is claimed here that the exam does not.

FAMILY	EXAM	WHAT IT TESTS
Multimodal annotation	<code>vvdex.annotation.image-object-1</code>	A constructed image is reviewed for classification, object boxes, attributes and keypoints.
Multimodal annotation	<code>vvdex.annotation.robot-video-1</code>	A robot-catching video is reviewed for temporal events, frame-level object boxes, track identity and occlusion.
Multimodal annotation	<code>vvdex.annotation.structured-data-1</code>	Invented tabular records are reviewed for normalization, validity, duplicates, anomalies and relationships.
Multimodal annotation	<code>vvdex.annotation.text-labeling-1</code>	An invented message is reviewed for classification and entity spans with consistent offsets.

04 Campaign aggregate across these four exams

MODEL LANE	CERTIFIED PASSES	GRADED ATTEMPTS	LANE ERRORS	WITHHELD	INVALID
cli/codex-gpt-5.5	25	40	0	0	0

Raw counts first. *Graded attempts* = certified passes + model fails; lane errors, withheld and invalid cells are outside capability arithmetic in both directions. These four fixed annotation tasks measure different work, so no combined rate or confidence interval is reported. The totals are accounting counts only; the per-task results above carry the quality claims.

05 Harness status

No record in this campaign fired a harness self-suspicion rule. Every countable cell is a grader outcome under verified containment.

06 Cells excluded from capability

An attempt is excluded from capability only when the lane, not the model, ended it (lane error), when the run cannot prove its own conditions (withheld), or when an integrity invariant fired (INVALID). A submission the hidden grader failed is a model result and is counted, never listed here.

No attempt in this campaign was excluded: every attempt produced a countable model result.

07 What the outcomes mean

- **CERTIFIED PASS:** the grader accepted the submission, containment was verified, provenance is complete and no integrity invariant fired.
- **model fail:** the submission failed the hidden tests under verified containment; a real, attributable model result.
- **lane error:** the provider or CLI lane failed (rate limit, quota, transport) before an attributable measurement existed; not a model result.
- **withheld:** the run cannot prove its own conditions (containment unproven, provenance incomplete, or a harness suspicion the record cannot clear); neither a pass nor a model failure.
- **INVALID:** an integrity invariant fired (isolation breach, evidence mismatch); the row is not a model result.
- **harness suspect:** the run's own self-check fired; each rule is classified as verdict-bearing or diagnostic, and verdict-bearing suspicion withholds the affected cells.

08 What this report withholds, and why

What a public document may print about a rollout is decided by what the exam GRADES, not case by case. Where the graded artefact is a change to a public upstream repository, a bounded redacted excerpt of the model's own diff is published. Where the graded output is an answer, a keyed value or the state of an application, nothing derived from the submission is published — no diff, no answer value, no model claim, and not the names of the hidden assertions a rollout failed, because a grader's name can carry the value it asserts. Withheld on every exam without exception: the contents of tool results, the hidden tests, the reference solution, host paths and keys. The counts, verdicts and intervals in this report are the grader's and are unaffected by what is withheld.

EXAM	DISCLOSURE CLASS	RULE APPLIED
image-object-1	answer_key	Applied consistently to 10 sealed records. answer-key disclosure (ownership vvdex_proprietary — no declaration — proprietary by default; family annotation): the exam is VVDex evaluation IP and the graded output is an answer or a keyed value, so no submission content, diff, model claim or hidden-assertion name is published.
robot-video-1	answer_key	Applied consistently to 10 sealed records. answer-key disclosure (ownership vvdex_proprietary — no declaration — proprietary by default; family annotation): the exam is VVDex evaluation IP and the graded output is an answer or a keyed value, so no submission content, diff, model claim or hidden-assertion name is published.
structured-data-1	answer_key	Applied consistently to 10 sealed records. answer-key disclosure (ownership vvdex_proprietary — no declaration — proprietary by default; family annotation): the exam is VVDex evaluation IP and the graded output is an answer or a keyed value, so no submission content, diff, model claim or hidden-assertion name is published.
text-labeling-1	answer_key	Applied consistently to 10 sealed records. answer-key disclosure (ownership vvdex_proprietary — no declaration — proprietary by default; family annotation): the exam is VVDex evaluation IP and the graded output is an answer or a keyed value, so no submission content, diff, model claim or hidden-assertion name is published.

09 Evidence

Every number above is recomputable from the sealed records below; each record verifies itself with `vvdex-env records verify`, and this document was regenerated from them with `vvdex-env records campaign`.

- `vvdex.annotation.image-object-1` → `fr-20260906-2d00c77f` · `recordDigest` `648e95b3d571109148c78137d9202e454bec7f03164f20e1ac4d24ca823395e0`
- `vvdex.annotation.image-object-1` → `fr-20260906-470625ed` · `recordDigest` `d1c1e3f5eaa431b0b2ae53bccae13895768f3d52b8a7d98ddb2f7f1ee6a53e74`
- `vvdex.annotation.image-object-1` → `fr-20260906-4a88a9bf` · `recordDigest` `6947e974ea555f5a7d8f4664650e33c2badab9922358dbac8ee209231413d495`
- `vvdex.annotation.image-object-1` → `fr-20260906-62fb1341` · `recordDigest` `7a95898a5f2cd0f0b73c4e885cbbff22d2d90754e7602bb646a68df5f8f5a1e6`
- `vvdex.annotation.image-object-1` → `fr-20260906-6dc38b98` · `recordDigest` `a66b51f5da540d99f9ecf6ff3602d204399934d919da8d7a9859684c63d583c9`
- `vvdex.annotation.image-object-1` → `fr-20260906-884851f3` · `recordDigest` `37a729b8483f10b49bdd80e7ba918abe57e5ddc3d6ba8be593d23d816e082d9c`
- `vvdex.annotation.image-object-1` → `fr-20260906-ad6f18b9` · `recordDigest` `ad3ad73c931070d584cb6d9cb4c645a177b416c9758efe1b62d1bfee32b293b6`
- `vvdex.annotation.image-object-1` → `fr-20260906-aef86705` · `recordDigest` `cde2fc42622e4ef27c1a9505f9a81637fcf0e48bb634045dc0dc9f33e4f87ba5`
- `vvdex.annotation.image-object-1` → `fr-20260906-af2506a7` · `recordDigest` `b7b6915c0f62d8d5942e7d4d50d86b3967f5e9e4ab77908f010e7f79ecba5497`
- `vvdex.annotation.image-object-1` → `fr-20260906-fd180ec1` · `recordDigest` `6a5c91191a02c8ada9bb6fdd380d95e2df6559de255866d09d3af7b2cec1eb23`
- `vvdex.annotation.robot-video-1` → `fr-20260906-08cf41e3` · `recordDigest` `4cb0e73aa85083f8e0ef903fd4d28f277a177b2774a66767e5635c84b75f8213`
- `vvdex.annotation.robot-video-1` → `fr-20260906-12f8c6af` · `recordDigest` `5fdd4fe7918f8689e1ad975912f7241d661247fdccbc3b30aca8b96499b7a427`
- `vvdex.annotation.robot-video-1` → `fr-20260906-55c7c9bc` · `recordDigest` `c50dca52eb4e3d7351d38238b8f67be3139846d39eae70d6fb9f3e7edef1d8da`
- `vvdex.annotation.robot-video-1` → `fr-20260906-57451e50` · `recordDigest` `bde0ebe1b4f14b7989a966b5a11157dc19cb88391c5937324b23d81d9da1a61b`
- `vvdex.annotation.robot-video-1` → `fr-20260906-5921101f` · `recordDigest` `cafc99313b275331a9ae86c70be3f042b7ab7b8d0b442a9d5100f51c3ba81e93`
- `vvdex.annotation.robot-video-1` → `fr-20260906-7b7f8fb2` · `recordDigest` `1879a1fa6919d32cd484a4e39530ec262f54bdc00e7b0b1015ed4d2032245145`
- `vvdex.annotation.robot-video-1` → `fr-20260906-7f86cd6e` · `recordDigest` `6637ce529278ae9c381a7407303cca0aa6aadfa60671a0efdedf2ba76e576f18`
- `vvdex.annotation.robot-video-1` → `fr-20260906-c5386852` · `recordDigest` `014e24edff39b7c34708f0bc5e8f1c826ba7e8ec940bf705cbb0f96f22d7a832`
- `vvdex.annotation.robot-video-1` → `fr-20260906-ea40e0a7` · `recordDigest` `6b7b5285b5fc6dd7c66338dc973344ad8a80f87314711b7aefe95f783f8ad1b5`
- `vvdex.annotation.robot-video-1` → `fr-20260906-ee230324` · `recordDigest` `f041ab53c0abd3fbd8155b14173cb956bb0ccdb8f3abe59c2b0ec260e46b29ed`

- vvdex.annotation.structured-data-1 → fr-20260906-0524b448 · recordDigest 6c7ef1ad19e86c265db5ca3ae26545f32a3aa7610be592847156999eaad0a613
- vvdex.annotation.structured-data-1 → fr-20260906-53c54cfc · recordDigest 04d2e9b4b28fa343d7827d6cab7a2a40a2b7d60e2a7ebbe6a19df7d214bb1de4
- vvdex.annotation.structured-data-1 → fr-20260906-6acfa16c · recordDigest 79cfd5852be5991c5c7a8e00c6855b01deda4cdd257fa08bb9e1eeb114640a12
- vvdex.annotation.structured-data-1 → fr-20260906-81881665 · recordDigest 6d4694530740ae02b13144a46a80351a3fe98fa31d89d3424737da904225ab37
- vvdex.annotation.structured-data-1 → fr-20260906-8a3e7827 · recordDigest 1307b28564b6fd0bdaabd143ff60207f7ea1c12dbfe5c387f37bb3f691d44b3a
- vvdex.annotation.structured-data-1 → fr-20260906-8e07c82f · recordDigest b6767314873b3ff13f6ce7f53682b9c383c88196633b2cdcae9e044e06b1358
- vvdex.annotation.structured-data-1 → fr-20260906-a9e8fca4 · recordDigest c3be15a7ee24dc46e4f61be830b3111cbf48792324238e28295020c33354e496
- vvdex.annotation.structured-data-1 → fr-20260906-b9a3013f · recordDigest 3753722b773d4e4aeee317ea729b2f76946a55af57cdfe71aa5f64727a2c9c1d
- vvdex.annotation.structured-data-1 → fr-20260906-ebf6f550 · recordDigest 98cf7408f3accbb7029809751c4723e435e720a77c22a94beb98bf611d5406de
- vvdex.annotation.structured-data-1 → fr-20260906-f1d50a1b · recordDigest 5fec386d9f78f894df537589ca1c3a9470ec13dfbf92b360d28240b91023ee65
- vvdex.annotation.text-labeling-1 → fr-20260906-09f49314 · recordDigest ad912cc0b1e879202be847d6ff47326a8b769d3913c33c89df9263326287cb48
- vvdex.annotation.text-labeling-1 → fr-20260906-1ac877d5 · recordDigest b5b01baf361a364761324340a0e7f776d1de9aa5edea6652a2f9bedfe9f7d379
- vvdex.annotation.text-labeling-1 → fr-20260906-3876bd90 · recordDigest 87e2b2d954245b76650d144b5e2ed2cc666cb17d73af83e76c5b3f6a23d0bdbb
- vvdex.annotation.text-labeling-1 → fr-20260906-38dcc889 · recordDigest 003299dc0eca735f784408e749762c887a90832c70bc66439503b7da10333eee
- vvdex.annotation.text-labeling-1 → fr-20260906-49710bdf · recordDigest 33104a999a0428902401b7ffea929b2de537af21b71863d2de01b5868fe21f9
- vvdex.annotation.text-labeling-1 → fr-20260906-50564102 · recordDigest 827e57535e98985d0fbe512136942a78163cd6d2123acab263ca597bdd883077
- vvdex.annotation.text-labeling-1 → fr-20260906-5c3a23d2 · recordDigest ec06840685b25beaf6b6df8775cc34a377131af1c61026b278359b8f422db04f
- vvdex.annotation.text-labeling-1 → fr-20260906-64a9e783 · recordDigest 314ad299e8c74d9f76a9615eb5054c16c73489e4a87ef0e3506ae7c7bff8993b
- vvdex.annotation.text-labeling-1 → fr-20260906-80719317 · recordDigest cb2f2d25f460da6d9620639467f41d4824c1c8013c8e7cd679da84a34be70cec
- vvdex.annotation.text-labeling-1 → fr-20260906-8a0a838e · recordDigest 40b107e6f5a2bc46cb85916a0758419edb7567c61980979e348def9a35e3b9b9

10 Attestation

Certified results appear only where evidence proves them. Withheld, invalid and lane-error cells are never presented as model results, and no historical evidence was modified in producing this document. Counts are exact; below ten rollouts per cell no rate or ranking beyond the counts is asserted.

© VVDex. All rights reserved. Public materials are provided for viewing, evaluation and verification. Reuse, redistribution, derivative benchmark creation, training use, republication or commercial use requires written permission unless a specific file states otherwise. Full terms: vvdexops.com/terms/. Third-party open source reproduced in an exam keeps its own licence.