

# Object geometry annotation

## EVALUATION REPORT

2026-09-06

*Evaluation scope:*

**GPT-5.5 via Codex CLI (subscription, image input)**

*Prepared by:*

**VVDex Forge** — generated directly from sealed run evidence

---

fr-20260906-62fb1341 · vvdex.annotation.image-object-1 v0.1.0 · Independent model evaluation report

# About this report

---

## About VVDex Forge

VVDex Forge builds and certifies evaluation environments for AI agents. Every exam is sealed before any model sees it: the reference solution must pass, an empty submission must fail, the grader must reject near-misses, and the hidden tests never enter the agent's tree. Reports and the public catalog live at **vvdexops.com**.

## Report integrity

This report is generated from run evidence; its digests verify authenticity. The sole test of a copy is the three-line procedure in Verification — a copy whose digests do not reproduce is not this report.

## Scope

One exam · 1 rollout(s) · 24 tool steps and 3000s wall clock per rollout, stated in full in Evaluation configuration.

## Boundary

Not a leaderboard. Not a general capability score. A post-seal benchmark: this run stamped no certification gate and rewrote nothing — not the exam, not the reference solution, not the grader, not the certified evidence.

# Contents

|     |                        |   |
|-----|------------------------|---|
| 01. | Annotation evidence    |   |
| 02. | Executive summary      | 4 |
| 03. | Exam definition        | 5 |
| 04. | Certification evidence | 6 |
| 05. | Model outcomes         | 7 |
| 06. | Stage funnel           | 8 |
| 07. | Findings               | 8 |
| 08. | Recommendations        | 9 |
| 09. | Verification           | 9 |

Followed by: companion evidence report reference - Attestation

## 01 Annotation evidence

**Modality:** image. **Rubric:** 1.0.

A constructed image is reviewed for classification, object boxes, attributes and keypoints.

### Submission aaff1bb0-dd17-45f4-9e5b-d47dc07237ae

| Score | Items | Valid / invalid | Types                         |
|-------|-------|-----------------|-------------------------------|
| 1     | 4     | 4 / 0           | box, classification, keypoint |

### Measured metrics

|                  |   |                         |   |
|------------------|---|-------------------------|---|
| annotationCount  | 4 | boxF1                   | 1 |
| boxIoU           | 1 | boxPrecision            | 1 |
| boxRecall        | 1 | classificationAccuracy  | 1 |
| classificationF1 | 1 | classificationPrecision | 1 |

|                         |   |                   |   |
|-------------------------|---|-------------------|---|
| classificationRecall    | 1 | keypointF1        | 1 |
| keypointNormalizedError | 0 | keypointPrecision | 1 |
| keypointRecall          | 1 | matchedCount      | 4 |
| mismatchCount           | 0 |                   |   |

Track consistency depends on successful box matching. A zero can reflect inaccurate boxes even when track IDs are correct.

Validation error codes

None recorded

Disagreement

Agreement reviewNot measured

Measured grader aggregates only. Missing metrics or disagreement are not measured, never zero. Reference answers and item-level labels are excluded. Certification and the evidence digest are recorded in this report's verification sections.

02 Executive summary

MODELS TESTED

1

CERTIFIED PASSES

1

of 1 graded rollouts

SUBMITTED

1 of 1

graded rollouts that called submit

NOT GRADED — LANE ERRORS

0

of 1 attempts

## Purpose

This report presents a controlled evaluation of 1 model lane(s) against the certified exam `vvdex.annotation.image-object-1` (v0.1.0, family annotation). Every lane received the identical frozen workspace, tool contract and limits; grading ran afterwards, in isolation, against hidden tests no lane ever saw. The purpose is to state, with reproducible evidence, what each lane did and where it stopped.

## Outcome

**1 of 1** graded rollouts passed the independent hidden tests and were certified

1 of 1 graded rollouts passed. `cli/codex-gpt-5.5` completed the applicable annotation workflow and were accepted by the independent annotation grader. **Low-N:** 1 rollout(s) is below the 10-rollout floor for reading a rate as a rate — read the per-rollout outcomes, not the percentages.

## Key findings

**F-01** `cli/codex-gpt-5.5` passed the independent hidden tests

**F-02** Statistical floor — 1 rollout(s) is below the 10-rollout rate floor

Each finding is stated in full, with evidence, in the Findings section; recommendations for acting on them follow it. Elapsed wall clock for the whole engagement: 47.2s.

## 03 Exam definition

**In scope.** The ability of each configured model lane to complete this one certified task end to end — inspect the asset, annotate against the declared rubric, validate, submit, and be accepted by the independent annotation grader — within 24 tool steps and 3000s of wall clock, at 1 rollout(s) per lane. Behavioural measurements along the way (tool discipline, grounding, overclaim) are in scope because they are read from the recorded trace, not judged by another model.

**Out of scope.** General model quality, performance on other task families, prompt-tuning on the lane's behalf, and any comparison this run's sample size cannot support. A lane's API failure is reported as infrastructure, never as model capability.

**The instrument.** Exam `vvdex.annotation.image-object-1` was certified before this run: its reference solution passes, an empty submission fails, its grader distinguishes near-misses, and its sealed evidence is unchanged by this evaluation (see Certification evidence).

What was measured

EXAM

**vvdex.annotation.image-object-1**

VERSION

**0.1.0**

FAMILY

**annotation**

ROLLOUTS

**1 (1 per model)**

Where the defect came from

This exam has no upstream repository to credit: the defect, the reference solution and the hidden suite are VVDex's own.

04 Certification evidence

**Why this exam is trustworthy.** Every pillar below was established before any model saw the exam, loaded from the sealed evidence matching this run's exam fingerprint, and unchanged by this evaluation.

FROZEN BASELINE

**FAIL**

expected — an empty submission over 2 run(s)  
must not pass

→

GOLD / REFERENCE

**PASS**

expected — the reference solution over 2 run(s)  
must pass

GRADER CONTROLS

**19 / 19**

The grader accepted the reference and rejected every near-miss.

ATTACK PROBES

**9 / 9 blocked**

0 trivial exploit(s) found.

SEALED EVIDENCE

**verified**

Sealed 2026-09-06 over 13 artifact(s).

RUNTIME

sha256:679b36225a1f  
83238efba8c71b9cde  
3c24caaeacc9b682ae  
92819cb55eec328f

network policy: deny

CERTIFICATION BUNDLE DIGEST

c f945b4838a266664300b32fdd5fed0c4e781ab95e  
aaebfde31a0ce8ae2d7295

EXAM FINGERPRINT

422753713aeb9f41abcc4135...

Full value in Appendix B; the contract digest this run was executed against.

05 Model outcomes

**Denominators.** Attempts requested: 1. Graded (evaluable) rollouts: 1. Not graded — lane errors: 0. A lane error is a rollout the API or the runtime ended before the hidden grader ran: it is listed, excluded from every rate and pass@k, and never silently dropped. Passes are certified passes: hidden grader exit 0 with integrity verified.

| MODEL                                                                                   | ATTEMPTS | GRADED | PASSED | MODEL FAILS | LANE ERRORS | SUBMITTED | MEAN TIME | TYPICAL DEEPEST STAGE |
|-----------------------------------------------------------------------------------------|----------|--------|--------|-------------|-------------|-----------|-----------|-----------------------|
| GPT-5.5 via<br>Codex CLI<br>(subscription,<br>image input)<br><br>cli/codex-gpt-5.<br>5 | 1        | 1      | 1 / 1  | 0           | 0           | 1 / 1     | 45.8s     | Passed                |

Every rollout, with its own deepest stage and grade, is in Appendix A.

Success rate, confidence and pass@k

Every rate on this page is a measurement of a finite number of rollouts, so every rate carries the interval that measurement actually supports. n is the evaluable count — attempts minus rollouts the lane ended before grading; those are listed under "not graded" and sit in no rate, interval or pass@k.

| MODEL             | ATTEMPTS | N | NOT GRADED | PASSES | RATE   | 95% INTERVAL | SEEDS | PASS@1 |
|-------------------|----------|---|------------|--------|--------|--------------|-------|--------|
| cli/codex-gpt-5.5 | 1        | 1 | 0          | 1      | 100.0% | [21%, 100%]  | 0     | 100.0% |

Interval and pass@k methods are stated in full in Appendix D. pass@k is shown for k = 1; the sealed record carries every k up to n.

Model comparison

One model ran in this record, so there is nothing to compare it with. A comparison table across models appears here when a run covers more than one.

06 Stage funnel

Recorded annotation actions and independent grader verdicts. Missing capabilities are N/A; each stage is measured independently.

| STAGE    | CLI:CLI/CODEX-GPT-5.5 |
|----------|-----------------------|
| Inspect  | 1 / 1                 |
| Annotate | 1 / 1                 |
| Validate | 1 / 1                 |
| Submit   | 1 / 1                 |
| Graded   | 1 / 1                 |
| Passed   | 1 / 1                 |

Deepest stage reached

| MODEL                 | DEEPEST (ANY ATTEMPT) | TYPICAL (MOST ATTEMPTS) | ANNOTATION QA PASSES | NOT GRADED |
|-----------------------|-----------------------|-------------------------|----------------------|------------|
| cli:cli/codex-gpt-5.5 | Passed                | Passed                  | 1 / 1                | 0          |

The matrix counts evaluable annotation attempts; lane errors are shown separately. Successful recorded actions establish each stage independently.

07 Findings

Every finding below is derived from the recorded data of this run — none is an opinion. Severity: PASS a lane succeeded · ATTENTION a capability gap · LANE / INFRA infrastructure, not capability · NOTE a property of the measurement.

F-01

PASS

**cli/codex-gpt-5.5 passed the independent hidden tests**

Severity: PASS

Evidence: Model outcomes · Stage funnel · Appendix A

1 of 1 graded rollouts completed the applicable annotation workflow and were accepted by the independent annotation grader. This demonstrates the task is solvable under the published limits.

F-01 · derived from the recorded run data

F-02

NOTE

**Statistical floor — 1 rollout(s) is below the 10-rollout rate floor**

Severity: NOTE

Evidence: Appendix D

The counts are exact; the percentages are anecdotes with decimal points. Per-rollout outcomes, not rates, are the reliable reading of this run.

F-02 · derived from the recorded run data

## 08 Recommendations

One recommendation per observed failure mode, addressed to the lane it affects. Each traces back to a finding above and to the raw trace in Appendix A.

**R-01 For anyone reading the percentages**

Commission a run at  $n \geq 10$  rollouts per lane before treating any rate here as a rate; this run's per-rollout outcomes are exact, its percentages are not yet stable.

## 09 Verification

RUN ID

**live-20260906T032642Z**

EXAM FINGERPRINT

**422753713aeb9f41abcc4135d637e3064cae1e770f  
c8d22ec9dacf260b35a3ef**

The contract digest this run was executed against.

## Tamper-evidence

Three digests, and how to check each one. If any of them does not reproduce, this page is not the page that was generated for this run.

### THIS REPORT

**3867366760468740587830edb7  
b94751a07cea17c9ce14cedd0c  
94d2bb7bc0dc**

sha256 over the report's UTF-8 bytes with every occurrence of the report digest itself replaced by 64 '0' characters

### EVAL RECORD

**14a04a7582760fe0819c4872bf  
3c8a1cc07b4b4a45f334fb740d  
08cbb4ac17cb**

sha256 of fr-20260906-62fb1341.json as written beside this file.

### RUN EVIDENCE BUNDLE

**62fb1341ce4fa28ba2c710132b  
84cfba742df77f008d6ea3877e  
233b3c5a5a17**

The digest the run's own bundle computed over itself.

### CERTIFIED EVIDENCE SEAL

**cf945b4838a266664300b32fdd  
5fed0c4e781ab95eaaebfde31a  
0ce8ae2d7295**

The exam's sealed evidence, established before this run.

## Measurement history

This is the first recorded run on this exam. There is nothing to compare it against yet, and a single measurement is not a trend.

## Verifying it, in three lines

1. `vvdex-env verify-report --report fr-20260906-62fb1341.html`
2. or by hand: replace every occurrence of the report digest above with 64 zeros, then sha256 the file – it must equal that digest
3. `shasum -a 256 fr-20260906-62fb1341.json` must equal the eval-record digest, and the bundle digest must equal `bundleSha256` in the run's bundle file

The full artifact-by-artifact digest table, and the reproduction protocol, are in Appendix B.

This is the executive report. The full verifiable evidence report for the same run — every rollout, artifact digests, redacted transcripts and methods — is `fr-20260906-62fb1341.pdf`, generated from exactly the same sealed record.

Attestation

This report was generated end to end by VVDex Forge from the sealed evidence of the evaluation run described above. Measurement values were populated from recorded artifacts, not manually entered. The digests in Verification let any holder of this file confirm that the report, the eval record and the run evidence are the ones this run produced — without contacting VVDex.

|                                            |                   |                                        |
|--------------------------------------------|-------------------|----------------------------------------|
| <b>VVDex Forge · vvdexops.com</b>          | <b>2026-09-06</b> | <b>vvdex-env 0.1.0</b>                 |
| Issuer — certified evaluation environments | Date of issue     | Engine · verifiers pin<br>c521e4146026 |

No upstream repository: this exam is VVDex's own.

Post-seal benchmark. It does not stamp any certifying gate and did not rewrite the exam, the gold solution, the grader or the certified evidence.

— END OF REPORT · fr-20260906-62fb1341 —