

# Object geometry annotation

## EVALUATION REPORT

2026-09-06

*Evaluation scope:*

**GPT-5.5 via Codex CLI (subscription, image input)**

*Prepared by:*

**VVDex Forge** — generated directly from sealed run evidence

---

fr-20260906-ad6f18b9 · vvdex.annotation.image-object-1 v0.1.0 · Independent model evaluation report

## About this report

---

### About VVDex Forge

VVDex Forge builds and certifies evaluation environments for AI agents. Every exam is sealed before any model sees it: the reference solution must pass, an empty submission must fail, the grader must reject near-misses, and the hidden tests never enter the agent's tree. Reports and the public catalog live at [vvdexops.com](https://vvdexops.com).

### Report integrity

This report is generated from run evidence; its digests verify authenticity. The sole test of a copy is the three-line procedure in Verification — a copy whose digests do not reproduce is not this report.

### Scope

One exam · 1 rollout(s) · 24 tool steps and 3000s wall clock per rollout, stated in full in Evaluation configuration.

### Boundary

Not a leaderboard. Not a general capability score. A post-seal benchmark: this run stamped no certification gate and rewrote nothing — not the exam, not the reference solution, not the grader, not the certified evidence.

# Contents

|     |                          |    |
|-----|--------------------------|----|
| 01. | Executive summary        | 4  |
| 02. | Exam definition          | 5  |
| 03. | Certification evidence   | 6  |
| 04. | Evaluation configuration | 7  |
| 05. | Model outcomes           | 8  |
| 06. | Stage funnel             | 9  |
| 07. | Annotation evidence      | 10 |
| 08. | Quality axes             | 11 |
| 09. | Behavioural analysis     | 12 |
| 10. | Findings                 | 13 |
| 11. | Recommendations          | 14 |
| 12. | Verification             | 15 |

Appendices: A — every rollout · B — evidence digests · C — redacted transcripts · D — metric definitions · Attestation

01 Executive summary

|                                       |                                                                          |                                                                                      |                                                                           |
|---------------------------------------|--------------------------------------------------------------------------|--------------------------------------------------------------------------------------|---------------------------------------------------------------------------|
| <div>MODELS TESTED</div> <div>1</div> | <div>CERTIFIED PASSES</div> <div>0</div> <div>of 1 graded rollouts</div> | <div>SUBMITTED</div> <div>1 of 1</div> <div>graded rollouts that called submit</div> | <div>NOT GRADED — LANE ERRORS</div> <div>0</div> <div>of 1 attempts</div> |
|---------------------------------------|--------------------------------------------------------------------------|--------------------------------------------------------------------------------------|---------------------------------------------------------------------------|

Purpose

This report presents a controlled evaluation of 1 model lane(s) against the certified exam `vv dex.annotation.image-object-1` (v0.1.0, family annotation). Every lane received the identical frozen workspace, tool contract and limits; grading ran afterwards, in isolation, against hidden tests no lane ever saw. The purpose is to state, with reproducible evidence, what each lane did and where it stopped.

Outcome

0 of 1

graded rollouts passed the independent hidden tests and were certified

0 of 1 graded rollouts passed the hidden tests. The exam's certification establishes the task is winnable, so every failure below is a finding about a lane, not about the instrument. **Low-N:** 1 rollout(s) is below the 10-rollout floor for reading a rate as a rate — read the per-rollout outcomes, not the percentages. The most common way to fail was `submitted_failed` — submitted annotations rejected by the independent grader.

Key findings

- F-01 No evaluated lane passed the hidden tests in this run
- F-02 `submitted_failed` × 1 (cli/codex-gpt-5.5)
- F-03 Statistical floor — 1 rollout(s) is below the 10-rollout rate floor

Each finding is stated in full, with evidence, in the Findings section; recommendations for acting on them follow it. Elapsed wall clock for the whole engagement: 47.8s.

02 Exam definition

**In scope.** The ability of each configured model lane to complete this one certified task end to end — inspect the asset, annotate against the declared rubric, validate, submit, and be accepted by the independent annotation grader — within 24 tool steps and 3000s of wall clock, at 1 rollout(s) per lane. Behavioural measurements along the way (tool discipline, grounding, overclaim) are in scope because they are read from the recorded trace, not judged by another model.

**Out of scope.** General model quality, performance on other task families, prompt-tuning on the lane's behalf, and any comparison this run's sample size cannot support. A lane's API failure is reported as infrastructure, never as model capability.

**The instrument.** Exam `vvdex.annotation.image-object-1` was certified before this run: its reference solution passes, an empty submission fails, its grader distinguishes near-misses, and its sealed evidence is unchanged by this evaluation (see Certification evidence).

What was measured

|                                 |         |            |                 |
|---------------------------------|---------|------------|-----------------|
| EXAM                            | VERSION | FAMILY     | ROLLOUTS        |
| vvdex.annotation.image-object-1 | 0.1.0   | annotation | 1 (1 per model) |

Where the defect came from

This exam has no upstream repository to credit: the defect, the reference solution and the hidden suite are VVDex's own.

03 Certification evidence

**Why this exam is trustworthy.** Every pillar below was established before any model saw the exam, loaded from the sealed evidence matching this run's exam fingerprint, and unchanged by this evaluation.

FROZEN BASELINE

FAIL

expected — an empty submission over 2 run(s) must not pass

→

GOLD / REFERENCE

PASS

expected — the reference solution over 2 run(s) must pass

GRADER CONTROLS

19 / 19

The grader accepted the reference and rejected every near-miss.

ATTACK PROBES

9 / 9 blocked

0 trivial exploit(s) found.

SEALED EVIDENCE

verified

Sealed 2026-09-06 over 13 artifact(s).

RUNTIME

sha256:679b36225a1f83238efba8c71b9cde3c24caaeacc9b682ae92819cb55eec328f

network policy: deny

CERTIFICATION BUNDLE DIGEST

cf945b4838a266664300b32fdd5fed0c4e781ab95eaaebfde31a0ce8ae2d7295

EXAM FINGERPRINT

422753713aeb9f41abcc4135...

Full value in Appendix B; the contract digest this run was executed against.

04 Evaluation configuration

|                |                                                                       |
|----------------|-----------------------------------------------------------------------|
| REPORT         | fr-20260906-ad6f18b9                                                  |
| ISSUED         | 2026-09-06                                                            |
| SUBJECT EXAM   | vvdex.annotation.image-object-1 · v0.1.0                              |
| FAMILY         | annotation                                                            |
| PREPARED BY    | VVDex Forge engine vvdex-env 0.1.0                                    |
| CLASSIFICATION | Independent model evaluation report — independently verifiable        |
| CERTIFICATION  | Exam certification: recorded and passing (see Certification evidence) |
| STATUS         | FINAL · report sealed against its own digest                          |

Timeline

| WHEN (UTC)          | EVENT                                                |
|---------------------|------------------------------------------------------|
| 2026-09-06 00:01:41 | Exam evidence sealed (before any model saw the exam) |
| 2026-09-06 03:32:07 | Evaluation run started                               |
| 2026-09-06 03:32:55 | Evaluation run finished                              |
| 2026-09-06 03:32:55 | Report generated from the sealed run evidence        |

Models under test

| MODEL                                             | PROVIDER | ENDPOINT      | BILLING                                            |
|---------------------------------------------------|----------|---------------|----------------------------------------------------|
| GPT-5.5 via Codex CLI (subscription, image input) | cli      | local runtime | operator subscription, flat — not billed per token |

A customer endpoint is recorded by origin only — never its port, its path, or its key.

Limits

|                                       |                                         |                                              |                                                                                                     |
|---------------------------------------|-----------------------------------------|----------------------------------------------|-----------------------------------------------------------------------------------------------------|
| <div>RUNTIME</div> <div>docker</div>  | <div>STEP LIMIT</div> <div>24</div>     | <div>WALL-CLOCK LIMIT</div> <div>3000s</div> | <div>MAX OUTPUT TOKENS</div> <div>2048</div>                                                        |
| <div>TEMPERATURE</div> <div>0.2</div> | <div>RETRY POLICY</div> <div>none</div> | <div>COMMAND BUDGET</div> <div>0</div>       | <div>TOOLS OFFERED</div> <div>list_files, read_file, write_file, edit_file, run_tests, submit</div> |

The evaluated model never saw the hidden tests or the reference solution, and could not change its result. Grading ran in a separate step, after submission, on a container with no network.

05 Model outcomes

**Denominators.** Attempts requested: 1. Graded (evaluable) rollouts: 1. Not graded — lane errors: 0. A lane error is a rollout the API or the runtime ended before the hidden grader ran: it is listed, excluded from every rate and pass@k, and never silently dropped. Passes are certified passes: hidden grader exit 0 with integrity verified.

| MODEL                                                                     | ATTEMPTS | GRADED | PASSED | MODEL FAILS | LANE ERRORS | SUBMITTED | MEAN TIME | TYPICAL DEEPEST STAGE |
|---------------------------------------------------------------------------|----------|--------|--------|-------------|-------------|-----------|-----------|-----------------------|
| GPT-5.5 via Codex CLI<br>(subscription, image input)<br>cli/codex-gpt-5.5 | 1        | 1      | 0 / 1  | 1           | 0           | 1 / 1     | 46.5s     | Graded                |

Every rollout, with its own deepest stage and grade, is in Appendix A.

Success rate, confidence and pass@k

Every rate on this page is a measurement of a finite number of rollouts, so every rate carries the interval that measurement actually supports. n is the evaluable count — attempts minus rollouts the lane ended before grading; those are listed under "not graded" and sit in no rate, interval or pass@k.

| MODEL             | ATTEMPTS | N | NOT GRADED | PASSES | RATE | 95% INTERVAL | SEEDS | PASS@1 |
|-------------------|----------|---|------------|--------|------|--------------|-------|--------|
| cli/codex-gpt-5.5 | 1        | 1 | 0          | 0      | 0.0% | [0%, 79%]    | 0     | 0.0%   |

Interval and pass@k methods are stated in full in Appendix D. pass@k is shown for k = 1; the sealed record carries every k up to n.

Model comparison

One model ran in this record, so there is nothing to compare it with. A comparison table across models appears here when a run covers more than one.



06 Stage funnel

Recorded annotation actions and independent grader verdicts. Missing capabilities are N/A; each stage is measured independently.

| STAGE    | CLI:CLI/CODEX-GPT-5.5 |
|----------|-----------------------|
| Inspect  | 1 / 1                 |
| Annotate | 1 / 1                 |
| Validate | 1 / 1                 |
| Submit   | 1 / 1                 |
| Graded   | 1 / 1                 |
| Passed   | 0 / 1                 |

Deepest stage reached

| MODEL                 | DEEPEST (ANY ATTEMPT) | TYPICAL (MOST ATTEMPTS) | ANNOTATION QA PASSES | NOT GRADED |
|-----------------------|-----------------------|-------------------------|----------------------|------------|
| cli:cli/codex-gpt-5.5 | Graded                | Graded                  | 0 / 1                | 0          |

The matrix counts evaluable annotation attempts; lane errors are shown separately. Successful recorded actions establish each stage independently.

07 Annotation evidence

Modality: image. Rubric: 1.0.

A constructed image is reviewed for classification, object boxes, attributes and keypoints.

Submission a470601f-1c98-499f-9b16-b5759fd1239b

| Score | Items | Valid / invalid | Types                         |
|-------|-------|-----------------|-------------------------------|
| 0     | 4     | 4 / 0           | box, classification, keypoint |

Measured metrics

|                  |   |                         |   |                        |   |
|------------------|---|-------------------------|---|------------------------|---|
| annotationCount  | 4 | boxF1                   | 1 | boxIoU                 | 0 |
| boxPrecision     | 1 | boxRecall               | 1 | classificationAccuracy | 1 |
| classificationF1 | 1 | classificationPrecision | 1 | classificationRecall   | 1 |
| keypointF1       | 1 | keypointNormalizedError | 0 | keypointPrecision      | 1 |
| keypointRecall   | 1 | matchedCount            | 4 | mismatchCount          | 1 |

Validation error codes

None recorded

Disagreement

|                  |              |
|------------------|--------------|
| Agreement review | Not measured |
|------------------|--------------|

Measured grader aggregates only. Missing metrics or disagreement are not measured, never zero. Reference answers and item-level labels are excluded. Certification and the evidence digest are recorded in this report's verification sections.

08 Quality axes

Each axis was measured inside the grade box, on the submitted workspace, by the same deterministic script for every rollout. The reward is the conjunction of the required axes; the score is a weighted reading beside it, not the gate.

| AXIS                    | GATE     | ROLLOUTS PASSING | MEASURED | BUDGET | WEIGHT | WHAT IT MEASURES |
|-------------------------|----------|------------------|----------|--------|--------|------------------|
| boxPrecision            | recorded | N/A              | 1        | –      | –      |                  |
| boxRecall               | recorded | N/A              | 1        | –      | –      |                  |
| boxF1                   | recorded | N/A              | 1        | –      | –      |                  |
| boxIoU                  | recorded | N/A              | 0        | –      | –      |                  |
| classificationPrecision | recorded | N/A              | 1        | –      | –      |                  |
| classificationRecall    | recorded | N/A              | 1        | –      | –      |                  |
| classificationF1        | recorded | N/A              | 1        | –      | –      |                  |
| classificationAccuracy  | recorded | N/A              | 1        | –      | –      |                  |
| keypointPrecision       | recorded | N/A              | 1        | –      | –      |                  |
| keypointRecall          | recorded | N/A              | 1        | –      | –      |                  |
| keypointF1              | recorded | N/A              | 1        | –      | –      |                  |

QUALITY SCORE (MEAN)

0

Weighted reading over the axes. It is not the reward.

REWARD COMPOSITION

Withheld from the public report.

09 Behavioural analysis

Long-horizon behaviour

Counted off the recorded trace of each rollout. A dead end is a step that moved nothing: a re-read of a slice already returned, a repeat of the previous action, or a call to a tool this exam does not have.

| MODEL                 | N | STEPS USED | RAN OUT OF STEPS | COMMANDS | CHANGED THE TREE | REFUSED (OVER BUDGET) | REVISITS | DEAD ENDS / ROLLOUT |
|-----------------------|---|------------|------------------|----------|------------------|-----------------------|----------|---------------------|
| cli:cli/codex-gpt-5.5 | 1 | 8 / 24     | 0                | 0 / 0    | 0                | 0                     | 0        | 0                   |

Each column is defined in Appendix D.

Hallucination & overclaim

Three measurements, each with a stated definition. None of them is a judgement of the model's prose by another model — every number below comes from comparing what the model wrote against what was on disk, or against what the independent grader returned.

| MODEL                 | N | GROUNDING FAILURES | KINDS | ROLLOUTS AFFECTED | OVERCLAIMS | OVERCLAIM RATE | RE-READS |
|-----------------------|---|--------------------|-------|-------------------|------------|----------------|----------|
| cli:cli/codex-gpt-5.5 | 1 | 0                  | none  | 0 / 1             | 0 / 1      | 0.0%           | 0        |

No grounding failures and no false pass claims were detected across these 1 rollout(s). No lane re-read material it had already been given.

Every term above is defined in Appendix D.

10 Findings

Every finding below is derived from the recorded data of this run — none is an opinion.

Severity: PASS a lane succeeded · ATTENTION a capability gap · LANE / INFRA infrastructure, not capability · NOTE a property of the measurement.

F-01

ATTENTION

No evaluated lane passed the hidden tests in this run

Severity: ATTENTIONEvidence: Model outcomes · Certification evidence

Every rollout either stopped before submitting or submitted an annotation submission the independent grader rejected. The exam's own certification (gold solution passes, empty submission fails) establishes the task itself is winnable.

F-01 · derived from the recorded run data

F-02

ATTENTION

submitted\_failed × 1 (cli/codex-gpt-5.5)

Severity: ATTENTIONEvidence: Appendix A · Appendix C

Submitted annotations rejected by the independent grader.

F-02 · derived from the recorded run data

F-03

NOTE

Statistical floor — 1 rollout(s) is below the 10-rollout rate floor

Severity: NOTEEvidence: Appendix D

The counts are exact; the percentages are anecdotes with decimal points. Per-rollout outcomes, not rates, are the reliable reading of this run.

F-03 · derived from the recorded run data

# 11 Recommendations

One recommendation per observed failure mode, addressed to the lane it affects. Each traces back to a finding above and to the raw trace in Appendix A.

R-01

PROBLEM

Submitted annotations rejected by the independent grader.

AFFECTED LANES

cli/codex-gpt-5.5

WHAT WE WOULD LOOK AT FIRST

Review the measured annotation metrics and validator error codes against the declared rubric.

R-02

For anyone reading the percentages

Commission a run at  $n \geq 10$  rollouts per lane before treating any rate here as a rate; this run's per-rollout outcomes are exact, its percentages are not yet stable.

12 Verification

RUN ID

live-20260906T033207Z

EXAM FINGERPRINT

422753713aeb9f41abcc4135d637e3064cae1e770fc8d22ec9dacf260b35a3ef

The contract digest this run was executed against.

Tamper-evidence

Three digests, and how to check each one. If any of them does not reproduce, this page is not the page that was generated for this run.

THIS REPORT

17048b9ee39ed1a78e39b1fd1800a5ea7a0f61a10b0be5a8812f4820d7974c16

sha256 over the report's UTF-8 bytes with every occurrence of the report digest itself replaced by 64 '0' characters

EVAL RECORD

dffbda1adb2d1182d71ed36c155f5474096029901eb676388ed02d13466fc69a

sha256 of fr-20260906-ad6f18b9.json as written beside this file.

RUN EVIDENCE BUNDLE

ad6f18b99b25bac17364e7af8e60b2f4dea5b1acf1c1b3efe112e13b1088c3ff

The digest the run's own bundle computed over itself.

CERTIFIED EVIDENCE SEAL

cf945b4838a266664300b32fdd5fed0c4e781ab95eaaebfde31a0ce8ae2d7295

The exam's sealed evidence, established before this run.

Measurement history

Previous run on this exam: fr-20260906-470625ed. Model progress over time on this exam is the difference between that record and this one — same frozen tree, same hidden grader, same limits. Anything else that moved between them moved outside this exam.

Verifying it, in three lines

1. vvdex-env verify-report --report fr-20260906-ad6f18b9.html

2. or by hand: replace every occurrence of the report digest above with 64 zeros, then sha256 the file – it must equal that digest

3. shasum -a 256 fr-20260906-ad6f18b9.json must equal the eval-record digest, and the bundle digest must equal bundleSha256 in the run's bundle file

The full artifact-by-artifact digest table, and the reproduction protocol, are in Appendix B.

## Appendix A — every rollout

| ROLLOUT  | MODEL             | SEED | OUTCOME          | SUBMITTED | DEEPEST STAGE | HIDDEN TESTS | FILES | TOKENS | ELAPSED |
|----------|-------------------|------|------------------|-----------|---------------|--------------|-------|--------|---------|
| a470601f | cli/codex-gpt-5.5 | 0    | submitted_failed | yes       | Graded        | fail         | 1     | n/a    | 46.5s   |

### Failure taxonomy

| CLASS            | COUNT | WHAT IT MEANS                                             |
|------------------|-------|-----------------------------------------------------------|
| submitted_failed | 1     | Submitted annotations rejected by the independent grader. |



## Appendix B — evidence and artifact digests

|                                                                                                                    |                                                              |                                                                          |                                                                           |
|--------------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------|--------------------------------------------------------------------------|---------------------------------------------------------------------------|
| <div><div>RUN BUNDLE DIGEST</div><div>ad6f18b99b25bac17364e7af8e60b2f4dea5b1acf1c1b3efe112e13b1088c3ff</div></div> | <div><div>RUN ID</div><div>live-20260906T033207Z</div></div> | <div><div>STARTED</div><div>2026-09-06T03:32:07.981657+00:00</div></div> | <div><div>FINISHED</div><div>2026-09-06T03:32:55.745591+00:00</div></div> |
|--------------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------|--------------------------------------------------------------------------|---------------------------------------------------------------------------|

### Artifact digests

| ARTIFACT                                           | SHA-256                                                          |
|----------------------------------------------------|------------------------------------------------------------------|
| publicSummary                                      | 452d38c30b5e31215b25a4342f9fe04d0e7587b2015abbd717a4f8867c607361 |
| rollouts/a470601f-1c98-499f-9b16-b5759fd1239b.json | c9ca1ac48d8068c3065ebd1dc4717af4f3eb394007e5bc639459f22278501725 |
| spec                                               | c3ff978ed9d1442006d00396b9e887a623d1eb0c4183777e8372d5caa235d405 |
| summary                                            | d213668dd56ef59a605bc39da74700e11bae32e5e6def3022fa74dee479e2f94 |

### Reproducing this run

Every rollout can be re-derived: the exam is frozen (fingerprint above), the grader ran in the recorded container image, and each rollout's full trace — every tool call, every result, every word — is in the evidence bundle at its rollouts/<id>.json path. A re-run uses the same job shape: the exam id, your endpoint, n, and the runtime. The record id names this run forever.

## Appendix C — redacted transcripts

For each rollout: the tool calls it made, in order — which tool, on which of the model's own paths, whether the call succeeded, and how many characters came back. What a tool returned is not here, and neither is the submission: on this exam the submitted content is the graded answer, so no diff, no answer value and no model claim is printed. A request for the reference solution or the hidden tests is recorded as a request; the path it asked for is replaced by a label. See Appendix D, "What this report withholds and why".

Disclosure class applied to this exam: answer-key disclosure (ownership vvdex\_proprietary — no declaration — proprietary by default; family annotation): the exam is VVDex evaluation IP and the graded output is an answer or a keyed value, so no submission content, diff, model claim or hidden-assertion name is published.

**a470601f · cli:cli/codex-gpt-5.5 · seed 0 · submitted\_failed**

| STEP | TOOL       | PATH                 | RESULT | CHARS |
|------|------------|----------------------|--------|-------|
| 0    | list_files | .                    | ok     | 80    |
| 1    | read_file  | TASK.md              | ok     | 443   |
| 2    | read_file  | annotation-task.json | ok     | 3 093 |
| 3    | read_file  | assets/objects.png   | ok     | 397   |
| 4    | read_file  | submission.json      | ok     | 74    |
| 5    | write_file | submission.json      | ok     | 64    |
| 6    | run_tests  | —                    | ok     | 186   |
| 7    | submit     | —                    | ok     | 85    |

Submission diff withheld by the public-disclosure policy (48 line(s) changed). Submission content withheld — proprietary evaluation material. On this exam the submitted change is the graded answer, so publishing an excerpt of it would publish the answer key.

## Appendix D — metric definitions and methods

Statistics, stage funnel and efficiency in this document are derived at render time from the sealed rollout rows and their raw tool events: lane errors leave every capability denominator; a rollout's stages are booleans read off its own event sequence (inspected = a successful list or read, or a successful browser open/snapshot, edited = the workspace changed, tested = run\_tests invoked, submitted = submit invoked, graded = the independent grader returned a verdict, passed = certified pass); absent telemetry is n/a. The record's own summary blocks are the values computed when it was sealed; the rows and events are the evidence.

### Denominators

n counts evaluable rollouts only: attempts minus lane errors. A rollout the lane ended (API or runtime failure) is reported but never graded, so it is in no rate, interval or pass@k.

### Interval method

Wilson score interval, not the normal approximation. The normal approximation collapses to a zero-width interval at 0 successes — it would print a 0/2 run as '0%, certain'. For k successes in n trials with  $p = k/n$ : centre =  $(p + z^2/2n) / (1 + z^2/n)$ , half-width =  $z/(1 + z^2/n) \cdot \sqrt{(p(1-p))/n + z^2/4n^2}$ , clamped to [0, 1],  $z = 1.9600$ .

### pass@k

Unbiased estimator:  $\text{pass@k} = 1 - C(n-c, k) / C(n, k)$ , for  $k \leq n$  only — reported only up to the number of rollouts actually run. It answers "given k independent attempts at this task, how often would at least one pass" — nothing more.

### Reading a failed rollout

A failed rollout is a measurement, not a defect report. Model API failures and runtime failures are classified separately precisely because they say nothing about the model's ability on this task.

This run has 1 rollout(s), below the 10 treated as the floor for reading a rate as a rate: the counts are exact, the percentages are anecdotes with decimal points.

### Stage definitions

#### Inspect

A successful annotation asset inspection event was recorded.

#### Annotate

A successful annotation write event was recorded.

#### Validate

A successful annotation validation event was recorded.

#### Submit

A successful annotation submission event was recorded.

#### Graded

The independent annotation grader returned a verdict.

#### Passed

The independent grader accepted the annotation submission. Local submission QA does not certify an environment or measure a model.

### Long-horizon terms

#### commandsRun

Shell commands executed inside the model's own container, against the command budget the contract declares. Only exams that advertise `run_command` have any.

#### **deadEnds**

Steps that moved nothing: a re-read, a repeat of the previous action, or a call to a tool this exam does not have. Counted, not judged.

#### **hitStepLimit**

Rollouts that ran out of steps without submitting. On a long task this is the number that separates 'could not finish' from 'finished wrong'.

#### **revisits**

Re-reads of a slice the rollout had already been given. On a long task, re-reading is how a model compensates for losing its own earlier context.

#### **stepsUsed**

Tool calls the rollout actually spent, out of the step limit the contract declares. A rollout that submitted early spent fewer, and that is a result, not a shortfall.

### **Grounding and overclaim terms**

#### **groundingFailure**

A file path the model wrote in its own reasoning that does not exist in the tree it was given, or that names a directory it is not allowed to see. Detected by matching source-file paths in the model's text against the agent box on disk — not by judging the prose.

#### **nonexistent\_path**

The model named a source file that is not in the repository it was given. It was reasoning about a file that does not exist.

#### **overclaim**

The model stated the tests pass or the issue is fixed, and the independent hidden tests then failed. Counted per rollout: claimed-and-wrong, or not.

#### **reReadChurn**

How many times the model read a file it had already read in the same rollout. It is a cost and a coherence signal, not a failure: re-reading is sometimes right.

#### **requested\_hidden\_path**

The model named a path under the reference solution, the hidden grader or the hidden test directory. Those are not in its box; the request could not be served.

### **Failure classes**

#### **submitted\_failed**

Submitted annotations rejected by the independent grader.

#### **submitted\_passed**

Submitted annotations accepted by the independent grader.

### **What this report withholds, and why**

On this exam the model's submitted content IS the graded answer, so nothing derived from a submission is published: no diff excerpt, no answer value or answer JSON body, no reference or gold value, no expected hidden output, none of the model's own words, and not the names of the hidden assertions a rollout failed — a grader's name can carry the value it asserts. What Appendix C does publish for each rollout: its id, its lane, its seed, the tool it called at each step, the public-safe path it called the tool on, whether the call succeeded, how many characters came back, the stage it reached and its outcome. A rollout the grader accepted is reported as submitted and accepted — the numbers in this report are unaffected, because the verdict is the grader's, not the reader's.

Rule applied to this exam: answer-key disclosure (ownership `vvdex_proprietary` — no declaration — proprietary by default; family annotation): the exam is VVDex evaluation IP and the graded output is an answer or a keyed value, so no submission content, diff, model claim or hidden-assertion name is published. The rule is a property of the exam family, not a judgement made per document, and the same rule governs every edition of this report, the campaign report that embeds it as a chapter, this exam's public page and every published JSON derived from the same sealed record.

Attestation

This report was generated end to end by VVDex Forge from the sealed evidence of the evaluation run described above. Measurement values were populated from recorded artifacts, not manually entered. The digests in Verification let any holder of this file confirm that the report, the eval record and the run evidence are the ones this run produced — without contacting VVDex.

|                                            |                   |                                     |
|--------------------------------------------|-------------------|-------------------------------------|
| <b>VVDex Forge · vvdexops.com</b>          | <b>2026-09-06</b> | <b>vvdex-env 0.1.0</b>              |
| Issuer — certified evaluation environments | Date of issue     | Engine · verifiers pin c521e4146026 |

No upstream repository: this exam is VVDex's own.  
Post-seal benchmark. It does not stamp any certifying gate and did not rewrite the exam, the gold solution, the grader or the certified evidence.

— END OF REPORT · fr-20260906-ad6f18b9 —