

Certified campaign — one rollout per lane

5 sealed evaluations · 16 graded rollouts · 8
certified passes

HISTORICAL — SUPERSEDED

CAMPAIGN REPORT — EXECUTIVE EDITION

2026-09-01

Evaluation scope:

cli/claude-sonnet

cli/codex

codestral-latest

openai/gpt-oss-120b

Prepared by:

VVDex Forge — regenerated from verified sealed records only

fc-8626f712e26f · 5 record(s) · every value recomputable

Contents

01.	Executive summary	3
02.	Per-exam results	5
03.	Environment families	6
04.	Campaign aggregate across these five exams	7
05.	Harness status	8
06.	Cells excluded from capability	9
07.	What the outcomes mean	10
08.	Evidence	11
09.	Attestation	12

01 Executive summary

Release status: HISTORICAL — superseded. A newer campaign carries these exams. This report is published so the earlier measurement stays checkable; its exams are NOT on the public featured index, and the current numbers are in the campaign the index names.

EXAMS 5	MODEL LANES 4	ATTEMPTS 20
GRADED ROLLOUTS 16	CERTIFIED PASSES 8	LANE ERRORS (NOT GRADED) 4
WITHHELD / INVALID 0		

Purpose

This report presents a controlled evaluation of 4 model lane(s) against 5 certified exam(s), one rollout per lane. Every lane received the identical frozen workspace, tool contract and limits per exam; grading ran afterwards, in isolation, against hidden tests no lane ever saw; and every rollout carries three separate outcomes — grader, integrity, certified — of which only the certified outcome is counted. Each exam is rendered below as a chapter derived from the same sealed evaluation record as its standalone report: the chapter's bytes are authenticated by this campaign document, the standalone report by its own digest.

Outcome

8

certified pass(es) across 16 graded rollouts (20 attempts, 4 lane error(s) not graded) in 20 cells — counted only where the grader accepted the submission AND containment was verified AND provenance is complete AND no integrity invariant or verdict-bearing harness suspicion applies.

Key findings across the campaign

- martinblech-xmltodict-issue-257 · **PASS** — cli/claude-sonnet, cli/codex passed the independent hidden tests
- martinblech-xmltodict-issue-257 · **LANE** — model_api_failure × 1 (openai/gpt-oss-120b)
- martinblech-xmltodict-issue-257 · **ATTENTION** — submitted_failed × 1 (codestral-latest)
- martinblech-xmltodict-issue-257 · **INFO** — Statistical floor — 4 rollout(s) is below the 10-rollout rate floor
- pytoolz-toolz-issue-602 · **PASS** — cli/claude-sonnet passed the independent hidden tests
- pytoolz-toolz-issue-602 · **ATTENTION** — submitted_failed × 2 (cli/codex, codestral-latest)
- pytoolz-toolz-issue-602 · **LANE** — model_api_failure × 1 (openai/gpt-oss-120b)
- pytoolz-toolz-issue-602 · **INFO** — Statistical floor — 4 rollout(s) is below the 10-rollout rate floor
- r1chardj0n3s-parse-issue-102 · **PASS** — cli/claude-sonnet passed the independent hidden tests
- r1chardj0n3s-parse-issue-102 · **ATTENTION** — submitted_failed × 2 (cli/codex, codestral-latest)
- r1chardj0n3s-parse-issue-102 · **LANE** — model_api_failure × 1 (openai/gpt-oss-120b)
- r1chardj0n3s-parse-issue-102 · **INFO** — Statistical floor — 4 rollout(s) is below the 10-rollout rate floor
- grounded-rag-1 · **PASS** — cli/codex passed the independent hidden tests
- grounded-rag-1 · **ATTENTION** — submitted_failed × 2 (cli/claude-sonnet, codestral-latest)

02 Per-exam results

MODEL LANE	SWE · XMLTODICT-257 PUBLIC.SWE.M ARTINBLECH-X MLTODICT-ISS UE-257	SWE · TOOLZ-602 PUBLIC.SWE. PYTOOLZ-T00 LZ-ISSUE-60 2	SWE · PARSE-102 PUBLIC.SWE.R 1CHARDJ0N3S- PARSE-ISSUE- 102	RAG VVDEX.KNOWL EDGE.GROUND ED-RAG-1	MEMORY VVDEX. KNOWLE DGE.ME MORY-F ACT-UP DATE-1
cli/claude-sonnet	1/1	1/1	1/1	0/1	1/1
cli/codex	1/1	0/1	0/1	1/1	1/1
openai/gpt-oss-120b	0/0 + 1 lane error	0/0 + 1 lane error	0/0 + 1 lane error	0/0 + 1 lane error	1/1
codestral-latest	0/1	0/1	0/1	0/1	0/1

Each cell is *certified passes / graded rollouts* for that lane on that exam. An attempt the lane ended before grading is reported beside the cell — 1/9 + 1 lane error, never 1/10 — and is outside the denominator; withheld and invalid attempts are reported the same way. No cell here carries a percentage: these are separate exams, not one score.

03 Environment families

This campaign covers five exam(s) drawn from three environment family(ies). The sentence in each row is the family's own description of what the exam tests; nothing is claimed here that the exam does not.

FAMILY	EXAM	WHAT IT TESTS
Software repair (public OSS SWE)	<code>public.swe.martinblech-xmltodict-issue-257</code>	A real software-engineering task taken from a public open-source repository: the model is given the codebase and a written issue, and must change the code so that the behaviour described in the issue is fixed.
Software repair (public OSS SWE)	<code>public.swe.pytoolz-toolz-issue-602</code>	A real software-engineering task taken from a public open-source repository: the model is given the codebase and a written issue, and must change the code so that the behaviour described in the issue is fixed.
Software repair (public OSS SWE)	<code>public.swe.r1chardj0n3s-parse-issue-102</code>	A real software-engineering task taken from a public open-source repository: the model is given the codebase and a written issue, and must change the code so that the behaviour described in the issue is fixed.
Grounded retrieval (RAG)	<code>vvdex.knowledge.grounded-rag-1</code>	Grounded question answering over a small document collection.
Cross-session memory update	<code>vvdex.knowledge.memory-fact-update-1</code>	A state-and-memory task across sessions: a fact was recorded earlier, an authoritative source has since changed it, and the model must answer from the newer authoritative roster rather than the stale carried memory, and acknowledge the change.

04 Campaign aggregate across these five exams

MODEL LANE	CERTIFIED PASSES	GRADED ROLLOUTS	LANE ERRORS	WITHHELD	INVALID	CERTIFIED RATE (95% CI)
cli/claude-sonnet	4	5	0	0	0	—
cli/codex	3	5	0	0	0	—
openai/gpt-oss-120b	1	1	4	0	0	—
codestral-latest	0	5	0	0	0	—

Raw counts first. *Graded rollouts* = certified passes + model fails; lane errors, withheld and invalid cells are outside capability arithmetic in both directions. Denominators differ between lanes, so no rate is shown and none should be inferred — a lane with one countable attempt has one data point, not a percentage.

05 Harness status

No record in this campaign fired a harness self-suspicion rule. Every countable cell is a grader outcome under verified containment.

06 Cells excluded from capability

An attempt is excluded from capability only when the lane, not the model, ended it (lane error), when the run cannot prove its own conditions (withheld), or when an integrity invariant fired (INVALID). A submission the hidden grader failed is a model result and is counted, never listed here.

EXAM	LANE	SEED	EXCLUSION	RECORDED REASON
<code>martinblech-xmltodict-issue-257</code>	<code>openai/gpt-oss-120b</code>	0	lane error	HTTP 429: {"error":{"message":"Rate limit reached for model `openai/gpt-oss-120b` in organization [redacted] service tier `on_demand` on tokens per day (TPD): ...
<code>pytoolz-toolz-issue-602</code>	<code>openai/gpt-oss-120b</code>	0	lane error	HTTP 429: {"error":{"message":"Rate limit reached for model `openai/gpt-oss-120b` in organization [redacted] service tier `on_demand` on tokens per day (TPD): ...
<code>r1chardj0n3s-parse-issue-102</code>	<code>openai/gpt-oss-120b</code>	0	lane error	HTTP 429: {"error":{"message":"Rate limit reached for model `openai/gpt-oss-120b` in organization [redacted] service tier `on_demand` on tokens per day (TPD): ...
<code>grounded-rag-1</code>	<code>openai/gpt-oss-120b</code>	0	lane error	HTTP 429: {"error":{"message":"Rate limit reached for model `openai/gpt-oss-120b` in organization [redacted] service tier `on_demand` on tokens per day (TPD): ...

07 What the outcomes mean

- **CERTIFIED PASS:** the grader accepted the submission, containment was verified, provenance is complete and no integrity invariant fired.
- **model fail:** the submission failed the hidden tests under verified containment; a real, attributable model result.
- **lane error:** the provider or CLI lane failed (rate limit, quota, transport) before an attributable measurement existed; not a model result.
- **withheld:** the run cannot prove its own conditions (containment unproven, provenance incomplete, or a harness suspicion the record cannot clear); neither a pass nor a model failure.
- **INVALID:** an integrity invariant fired (isolation breach, evidence mismatch); the row is not a model result.
- **harness suspect:** the run's own self-check fired; each rule is classified as verdict-bearing or diagnostic, and verdict-bearing suspicion withholds the affected cells.

08 Evidence

Every number above is recomputable from the sealed records below; each record verifies itself with `vvdex-env records verify`, and this document was regenerated from them with `vvdex-env records campaign`.

- `public.swe.martinblech-xmltodict-issue-257` → `fr-20260901-a27316a3` · recordDigest `96f8b631c6d65f03b21e4a9a94e86c82874e1506fac38c9a093d2f8c618dd2dc`
- `public.swe.pytoolz-toolz-issue-602` → `fr-20260901-88c3cdf1` · recordDigest `f49a31336e4865c1dcefccf5141e2f6926110c72f03d6cb3d645a1761fa09ac8`
- `public.swe.r1chardj0n3s-parse-issue-102` → `fr-20260901-4b1e3251` · recordDigest `d94c313144f8b36db4448440b873462bc9372e1783b1164451695b7b8497816a`
- `vvdex.knowledge.grounded-rag-1` → `fr-20260901-718db345` · recordDigest `a6426b0b7932b6c0f679f230c44c4b75a30df10a17f7e145b567ac9849d63a36`
- `vvdex.knowledge.memory-fact-update-1` → `fr-20260901-109cfa3c` · recordDigest `6354cff6fc47eca13b596d285f805a9f3994d048d69b446166acfeb7f8659929`

09 Attestation

Certified results appear only where evidence proves them. Withheld, invalid and lane-error cells are never presented as model results, and no historical evidence was modified in producing this document. Counts are exact; below ten rollouts per cell no rate or ranking beyond the counts is asserted.